

# Table of Contents

## C++ Language Reference

Welcome Back to C++ (Modern C++)

C++ Type System (Modern C++)

Uniform Initialization and Delegating Constructors

Object Lifetime And Resource Management (Modern C++)

Objects Own Resources (RAII)

Smart Pointers (Modern C++)

Pimpl For Compile-Time Encapsulation (Modern C++)

Containers (Modern C++)

Algorithms (Modern C++)

String and I-O Formatting (Modern C++)

Errors and Exception Handling (Modern C++)

Portability At ABI Boundaries (Modern C++)

## Lexical Conventions

Overview of File Translation

Character Sets2

Tokens (C++)

Comments (C++)

Identifiers (C++)

Keywords (C++)

Punctuators (C++)

Numeric, Boolean and Pointer Literals (C++)

String and Character Literals (C++)

User-Defined Literals (C++)

## Basic Concepts (C++)

Declarations and Definitions (C++)

Scope (Visual C++)

Program and Linkage (C++)

Startup and Termination (C++)

Lvalues and Rvalues (Visual C++)

Temporary Objects

Alignment (C++ Declarations)

`alignof` and `alignas` (C++)

Trivial, standard-layout, and POD types

Fundamental Types (C++)

Data Type Ranges

`nullptr`

`void` (C++)

`bool` (C++)

`false` (C++)

`true` (C++)

`char`, `wchar_t`, `char16_t`, `char32_t`

`_int8`, `_int16`, `_int32`, `_int64`

`_m64`

`_m128`

`_m128d`

`_m128i`

`_ptr32`, `_ptr64`

Numerical Limits (C++)

Standard Conversions

C++ Built-in Operators, Precedence and Associativity

`_alignof` Operator

`_uuidof` Operator

Additive Operators: `+` and `-`

Address-of Operator: `&`

Assignment Operators

Bitwise AND Operator: `&`

Bitwise Exclusive OR Operator: `^`

Bitwise Inclusive OR Operator: `|`

Cast Operator: `()`

Comma Operator: `,`

Conditional Operator: ?:

delete Operator (C++)

Equality Operators: == and !=

Explicit Type Conversion Operator: ()

Function Call Operator: ()

Indirection Operator: \*

Left Shift and Right Shift Operators (>> and <<)

Logical AND Operator: &&

Logical Negation Operator: !

Logical OR Operator: ||

Member Access Operators: . and ->

Multiplicative Operators and the Modulus Operator

new Operator (C++)

One's Complement Operator: ~

Pointer-to-Member Operators: .\* and ->\*

Postfix Increment and Decrement Operators: ++ and --

Prefix Increment and Decrement Operators: ++ and --

Relational Operators: <, >, <=, and >=

Scope Resolution Operator: ::

sizeof Operator

Subscript Operator:

typeid Operator

Unary Plus and Negation Operators: + and -

Expressions (C++)

Types of Expressions

Semantics of Expressions

Casting

Statements (C++)

Overview of C++ Statements

Labeled Statements

Expression Statement

Compound Statements (Blocks)

- [Selection Statements \(C++\)](#)
- [Iteration Statements \(C++\)](#)
- [Jump Statements \(C++\)](#)
- [Namespaces \(C++\)](#)
- [Enumerations \(C++\)](#)
- [Unions](#)
- [Functions \(C++\)](#)
  - [Functions with Variable Argument Lists \(C++\)](#)
  - [Function Overloading](#)
  - [Explicitly Defaulted and Deleted Functions](#)
  - [Argument-Dependent Name \(Koenig\) Lookup on Functions](#)
  - [Default Arguments](#)
  - [Inline Functions \(C++\)](#)
- [Operator Overloading](#)
  - [General Rules for Operator Overloading](#)
  - [Overloading Unary Operators](#)
  - [Binary Operators](#)
  - [Assignment](#)
  - [Function Call \(C++\)](#)
  - [Subscripting](#)
  - [Member Access](#)
- [Classes and Structs \(C++\)](#)
  - [class \(C++\)](#)
  - [struct \(C++\)](#)
  - [Class Member Overview](#)
  - [Member Access Control \(C++\)](#)
  - [Initializing classes and structs without constructors \(C++\)](#)
  - [Constructors \(C++\)](#)
  - [Destructors \(C++\)](#)
  - [Overview of Member Functions](#)
  - [Inheritance \(C++\)](#)
  - [Special Member Functions](#)

[Static Members \(C++\)](#)

[User-Defined Type Conversions \(C++\)](#)

[Mutable Data Members \(C++\)](#)

[Nested Class Declarations](#)

[Anonymous Class Types](#)

[Pointers to Members](#)

[this Pointer](#)

[C++ Bit Fields](#)

[Lambda Expressions in C++](#)

[Lambda Expression Syntax](#)

[Examples of Lambda Expressions](#)

[constexpr Lambda Expressions](#)

[Arrays \(C++\)](#)

[Using Arrays \(C++\)](#)

[Initializing Arrays](#)

[Arrays in Expressions](#)

[Interpretation of Subscript Operator](#)

[Indirection on Array Types](#)

[Ordering of C++ Arrays](#)

[References \(C++\)](#)

[Lvalue Reference Declarator: &](#)

[Rvalue Reference Declarator: &&](#)

[Reference-Type Function Arguments](#)

[Reference-Type Function Returns](#)

[References to Pointers](#)

[Pointers \(C++\)](#)

[const and volatile Pointers](#)

[new and delete Operators](#)

[Exception Handling in Visual C++](#)

[C++ Exception Handling](#)

[Structured Exception Handling \(C/C++\)](#)

[Transporting Exceptions Between Threads](#)

## Assertion and User-Supplied Messages (C++)

`static_assert`

## Templates (C++)

`typename`

`Class Templates`

`Template Friends`

`Function Templates`

`Template Specialization (C++)`

`Templates and Name Resolution`

`Source code organization (C++ Templates)`

## Event Handling

`_event`

`_hook`

`_raise`

`_unhook`

`Event Handling in Native C++`

`Event Handling in COM`

## Microsoft-Specific Modifiers

`Based Addressing`

`Calling Conventions`

`restrict (C++ AMP)`

`tile_static Keyword`

`_declspec`

`_restrict`

`_sptr, _uptr`

`_unaligned`

`_w64`

`_func_`

## Compiler COM Support

`Compiler COM Global Functions`

`Compiler COM Support Classes`

## Microsoft Extensions

## Nonstandard Behavior

### Compiler Limits

# C++ Language Reference

12/21/2017 • 2 min to read • [Edit Online](#)

This reference explains the C++ programming language as implemented in Microsoft Visual C++. The organization is based on *The Annotated C++ Reference Manual* by Margaret Ellis and Bjarne Stroustrup and on the ANSI/ISO C++ International Standard (ISO/IEC FDIS 14882). Microsoft-specific implementations of C++ language features are included.

For an overview of Modern C++ programming practices, see [Welcome Back to C++](#).

See the following tables to quickly find a keyword or operator:

- [C++ Keywords](#)
- [C++ Operators](#)

## In This Section

### [Lexical Conventions](#)

Fundamental lexical elements of a C++ program: tokens, comments, operators, keywords, punctuators, literals. Also, file translation, operator precedence/associativity.

### [Basic Concepts](#)

Scope, linkage, program startup and termination, storage classes, and types.

### [Standard Conversions](#)

Type conversions between built-in, or "fundamental," types. Also, arithmetic conversions and conversions among pointer, reference, and pointer-to-member types.

### [Operators, Precedence and Associativity](#)

The operators in C++.

### [Expressions](#)

Types of expressions, semantics of expressions, reference topics on operators, casting and casting operators, run-time type information.

### [Lambda Expressions](#)

A programming technique that implicitly defines a function object class and constructs a function object of that class type.

### [Statements](#)

Expression, null, compound, selection, iteration, jump, and declaration statements.

### [Declarations and Definitions](#)

Storage-class specifiers, function definitions, initializations, enumerations, class, struct, and union declarations, and typedef declarations. Also, inline functions, const keyword, namespaces.

### [Declarators](#)

The portion of a declaration statement that names an object, type, or function. Abstract declarators, type names, initializers, function declarations and definitions, arrays, references.

### [Classes, Structures, and Unions](#)

Introduction to classes, structures, and unions. Also, member functions, special member functions, data members, bit fields, this pointer, nested classes.

## Derived Classes

Single and multiple inheritance, virtual functions, multiple base classes, abstract classes, scope rules. Also, the `_super` and `_interface` keywords.

## Member-Access Control

Controlling access to class members: public, private, and protected keywords. Friend functions and classes.

## Overloading

Overloaded operators, rules for operator overloading.

## Exception Handling

C++ exception handling, structured exception handling (SEH), keywords used in writing exception handling statements.

## Assertion and User-Supplied Messages

`#error` directive, the `static_assert` keyword, the `assert` macro.

## Templates

Template specifications, function templates, class templates, typename keyword, templates vs. macros, templates and smart pointers.

## Event Handling

Declaring events and event handlers.

## Microsoft-Specific Modifiers

Modifiers specific to Microsoft C++. Memory addressing, calling conventions, naked functions, extended storage-class attributes (`_declspec`), `_w64`.

## Inline Assembler

Using assembly language and C++ in `_asm` blocks.

## Compiler COM Support

A reference to Microsoft-specific classes and global functions used to support COM types.

## Microsoft Extensions

Microsoft extensions to C++.

## Nonstandard Behavior

Information about nonstandard behavior of the Visual C++ compiler.

**Welcome Back to C++** An overview of modern C++ programming practices for writing safe, correct and efficient programs.

# Related Sections

## Component Extensions for Runtime Platforms

Reference material on using Visual C++ to target the common language runtime.

## C/C++ Building Reference

Compiler options, linker options, and other build tools.

## C/C++ Preprocessor Reference

Reference material on pragmas, preprocessor directives, predefined macros, and the preprocessor.

## Visual C++ Libraries

A list of links to the reference start pages for the various Visual C++ libraries.

# See Also



# Welcome Back to C++ (Modern C++)

1/3/2018 • 5 min to read • [Edit Online](#)

C++ is one of the most widely used programming languages in the world. Well-written C++ programs are fast and efficient. The language is more flexible than other languages because you can use it to create a wide range of apps—from fun and exciting games, to high-performance scientific software, to device drivers, embedded programs, and Windows client apps. For more than 20 years, C++ has been used to solve problems like these and many others. What you might not know is that an increasing number of C++ programmers have folded up the dowdy C-style programming of yesterday and have donned modern C++ instead.

One of the original requirements for C++ was backward compatibility with the C language. Since then, C++ has evolved through several iterations—C with Classes, then the original C++ language specification, and then the many subsequent enhancements. Because of this heritage, C++ is often referred to as a multi-paradigm programming language. In C++, you can do purely procedural C-style programming that involves raw pointers, arrays, null-terminated character strings, custom data structures, and other features that may enable great performance but can also spawn bugs and complexity. Because C-style programming is fraught with perils like these, one of the founding goals for C++ was to make programs both type-safe and easier to write, extend, and maintain. Early on, C++ embraced programming paradigms such as object-oriented programming. Over the years, features have been added to the language, together with highly-tested standard libraries of data structures and algorithms. It's these additions that have made the modern C++ style possible.

Modern C++ emphasizes:

- Stack-based scope instead of heap or static global scope.
- Auto type inference instead of explicit type names.
- Smart pointers instead of raw pointers.
- `std::string` and `std::wstring` types (see [`<string>`](#)) instead of raw `char[]` arrays.
- C++ Standard Library containers like `vector`, `list`, and `map` instead of raw arrays or custom containers. See [`<vector>`](#), [`<list>`](#), and [`<map>`](#).
- C++ Standard Library [algorithms](#) instead of manually coded ones.
- Exceptions, to report and handle error conditions.
- Lock-free inter-thread communication using C++ Standard Library `std::atomic<>` (see [`<atomic>`](#)) instead of other inter-thread communication mechanisms.
- Inline [lambda functions](#) instead of small functions implemented separately.
- Range-based for loops to write more robust loops that work with arrays, C++ Standard Library containers, and Windows Runtime collections in the form `for ( for-range-declaration : expression )`. This is part of the Core Language support. For more information, see [Range-based for Statement \(C++\)](#).

The C++ language itself has also evolved. Compare the following code snippets. This one shows how things used to be in C++:

```

// circle and shape are user-defined types
circle* p = new circle( 42 );
vector<shape*> v = load_shapes();

for( vector<circle*>::iterator i = v.begin(); i != v.end(); ++i ) {
    if( *i && **i == *p )
        cout << **i << " is a match\n";
}

for( vector<circle*>::iterator i = v.begin();
      i != v.end(); ++i ) {
    delete *i; // not exception safe
}

delete p;

```

Here's how the same thing is accomplished in modern C++:

```

#include <memory>
#include <vector>
// ...
// circle and shape are user-defined types
auto p = make_shared<circle>( 42 );
vector<shared_ptr<shape>> v = load_shapes();

for( auto& s : v ) {
    if( s && *s == *p )
        cout << *s << " is a match\n";
}

```

In modern C++, you don't have to use `new/delete` or explicit exception handling because you can use smart pointers instead. When you use the `auto` type deduction and [lambda function](#), you can write code quicker, tighten it, and understand it better. And `for_each` is cleaner, easier to use, and less prone to unintended errors than a `for` loop. You can use boilerplate together with minimal lines of code to write your app. And you can make that code exception-safe and memory-safe, and have no allocation/deallocation or error codes to deal with.

Modern C++ incorporates two kinds of polymorphism: compile-time, through templates, and run-time, through inheritance and virtualization. You can mix the two kinds of polymorphism to great effect. The C++ Standard Library template `shared_ptr` uses internal virtual methods to accomplish its apparently effortless type erasure. But don't over-use virtualization for polymorphism when a template is the better choice. Templates can be very powerful.

If you're coming to C++ from another language, especially from a managed language in which most of the types are reference types and very few are value types, know that C++ classes are value types by default. But you can specify them as reference types to enable polymorphic behavior that supports object-oriented programming. A helpful perspective: value types are more about memory and layout control, reference types are more about base classes and virtual functions to support polymorphism. By default, value types are copyable—they each have a copy constructor and a copy assignment operator. When you specify a reference type, make the class non-copyable—disable the copy constructor and copy assignment operator—and use a virtual destructor, which supports the polymorphism. Value types are also about the contents, which, when they are copied, give you two independent values that you can modify separately. But reference types are about identity—what kind of object it is—and for this reason are sometimes referred to as polymorphic types.

C++ is experiencing a renaissance because power is king again. Languages like Java and C# are good when programmer productivity is important, but they show their limitations when power and performance are paramount. For high efficiency and power, especially on devices that have limited hardware, nothing beats modern C++.

Not only the language is modern, the development tools are, too. Visual Studio makes all parts of the development cycle robust and efficient. It includes Application Lifecycle Management (ALM) tools, IDE enhancements like IntelliSense, tool-friendly mechanisms like XAML, and building, debugging, and many other tools.

The articles in this part of the documentation provide high-level guidelines and best practices for the most important features and techniques for writing modern C++ programs.

- [C++ Type System](#)
- [Uniform Initialization and Delegating Constructors](#)
- [Object Lifetime And Resource Management](#)
- [Objects Own Resources \(RAII\)](#)
- [Smart Pointers](#)
- [Pimpl For Compile-Time Encapsulation](#)
- [Containers](#)
- [Algorithms](#)
- [String and I/O Formatting \(Modern C++\)](#)
- [Errors and Exception Handling](#)
- [Portability At ABI Boundaries](#)

For more information, see the StackOverflow article [what C++ idioms are deprecated in C++11](#)

## See Also

- [C++ Language Reference](#)
- [Lambda Expressions](#)
- [C++ Standard Library](#)
- [Visual C++ language conformance](#)

# C++ Type System (Modern C++)

12/21/2017 • 13 min to read • [Edit Online](#)

The concept of *type* is very important in C++. Every variable, function argument, and function return value must have a type in order to be compiled. Also, every expression (including literal values) is implicitly given a type by the compiler before it is evaluated. Some examples of types include `int` to store integral values, `double` to store floating-point values (also known as *scalar* data types), or the Standard Library class `std::basic_string` to store text. You can create your own type by defining a `class` or `struct`. The type specifies the amount of memory that will be allocated for the variable (or expression result), the kinds of values that may be stored in that variable, how those values (as bit patterns) are interpreted, and the operations that can be performed on it. This article contains an informal overview of the major features of the C++ type system.

## Terminology

**Variable:** The symbolic name of a quantity of data so that the name can be used to access the data it refers to throughout the scope of the code where it is defined. In C++, *variable* is generally used to refer to instances of scalar data types, whereas instances of other types are usually called *objects*.

**Object:** For simplicity and consistency, this article uses the term *object* to refer to any instance of a class or structure, and when it is used in the general sense includes all types, even scalar variables.

**POD type** (plain old data): This informal category of data types in C++ refers to types that are scalar (see the Fundamental types section) or are *POD classes*. A POD class has no static data members that aren't also PODs, and has no user-defined constructors, user-defined destructors, or user-defined assignment operators. Also, a POD class has no virtual functions, no base class, and no private or protected non-static data members. POD types are often used for external data interchange, for example with a module written in the C language (which has POD types only).

## Specifying variable and function types

C++ is a *strongly typed* language and it is also *statically-typed*; every object has a type and that type never changes (not to be confused with static data objects).

**When you declare a variable** in your code, you must either specify its type explicitly, or use the `auto` keyword to instruct the compiler to deduce the type from the initializer.

**When you declare a function** in your code, you must specify the type of each argument and its return value, or `void` if no value is returned by the function. The exception is when you are using function templates, which allow for arguments of arbitrary types.

After you first declare a variable, you cannot change its type at some later point. However, you can copy the variable's value or a function's return value into another variable of a different type. Such operations are called *type conversions*, which are sometimes necessary but are also potential sources of data loss or incorrectness.

When you declare a variable of POD type, we strongly recommend you initialize it, which means to give it an initial value. Until you initialize a variable, it has a "garbage" value that consists of whatever bits happened to be in that memory location previously. This is an important aspect of C++ to remember, especially if you are coming from another language that handles initialization for you. When declaring a variable of non-POD class type, the constructor handles initialization.

The following example shows some simple variable declarations with some descriptions for each. The example also shows how the compiler uses type information to allow or disallow certain subsequent operations on the variable.

```

int result = 0;           // Declare and initialize an integer.
double coefficient = 10.8; // Declare and initialize a floating
                          // point value.
auto name = "Lady G.";   // Declare a variable and let compiler
                          // deduce the type.
auto address;            // error. Compiler cannot deduce a type
                          // without an initializing value.
age = 12;                // error. Variable declaration must
                          // specify a type or use auto!
result = "Kenny G.";    // error. Can't assign text to an int.
string result = "zero";  // error. Can't redefine a variable with
                          // new type.
int maxValue;            // Not recommended! maxValue contains
                          // garbage bits until it is initialized.

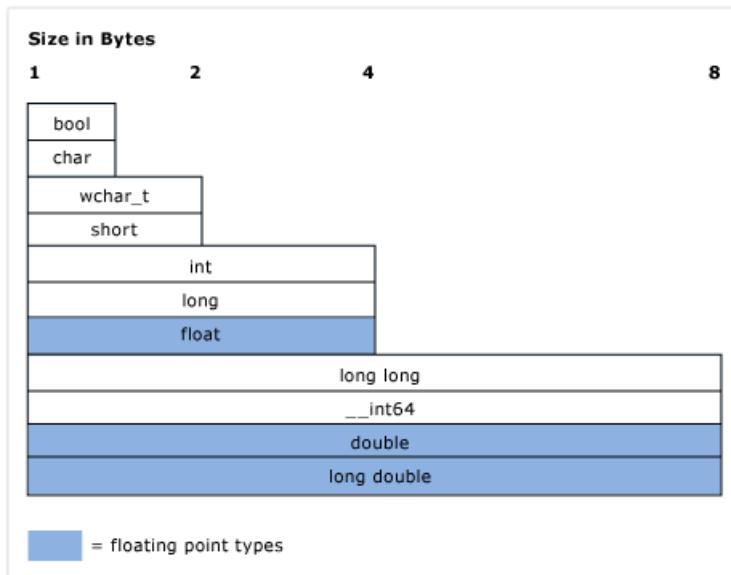
```

## Fundamental (built-in) types

Unlike some languages, C++ has no universal base type from which all other types are derived. The Visual C++ implementation of the language includes many *fundamental types*, also known as *built-in types*. This includes numeric types such as `int`, `double`, `long`, `bool`, plus the `char` and `wchar_t` types for ASCII and UNICODE characters, respectively. Most fundamental types (except `bool`, `double`, `wchar_t` and related types) all have unsigned versions, which modify the range of values that the variable can store. For example, an `int`, which stores a 32-bit signed integer, can represent a value from -2,147,483,648 to 2,147,483,647. An `unsigned int`, which is also stored as 32-bits, can store a value from 0 to 4,294,967,295. The total number of possible values in each case is the same; only the range is different.

The fundamental types are recognized by the compiler, which has built-in rules that govern what operations you can perform on them, and how they can be converted to other fundamental types. For a complete list of built-in types and their size and numeric limits, see [Fundamental Types](#).

The following illustration shows the relative sizes of the built-in types:



The following table lists the most frequently used fundamental types:

TYPE	SIZE	COMMENT
<code>int</code>	4 bytes	The default choice for integral values.

TYPE	SIZE	COMMENT
double	8 bytes	The default choice for floating point values.
bool	1 byte	Represents values that can be either true or false.
char	1 byte	Use for ASCII characters in older C-style strings or std::string objects that will never have to be converted to UNICODE.
wchar_t	2 bytes	Represents "wide" character values that may be encoded in UNICODE format (UTF-16 on Windows, other operating systems may differ). This is the character type that is used in strings of type std::wstring .
unsigned char	1 byte	C++ has no built-in byte type. Use unsigned char to represent a byte value.
unsigned int	4 bytes	Default choice for bit flags.
long long	8 bytes	Represents very large integer values.

## The void type

The void type is a special type; you cannot declare a variable of type void , but you can declare a variable of type void \* (pointer to void ), which is sometimes necessary when allocating raw (un-typed) memory. However, pointers to void are not type-safe and generally their use is strongly discouraged in modern C++. In a function declaration, a void return value means that the function does not return a value; this is a common and acceptable use of void . While the C language required functions that have zero parameters to declare void in the parameter list, for example, fou(void) , this practice is discouraged in modern C++ and should be declared fou() . For more information, see [Type Conversions and Type Safety](#).

## const type qualifier

Any built-in or user-defined type may be qualified by the const keyword. Additionally, member functions may be const -qualified and even const -overloaded. The value of a const type cannot be modified after it is initialized.

```
const double PI = 3.1415;
PI = .75 //Error. Cannot modify const variable.
```

The const qualifier is used extensively in function and variable declarations and "const correctness" is an important concept in C++; essentially it means to use const to guarantee, at compile time, that values are not modified unintentionally. For more information, see [const](#).

A const type is distinct from its non-const version; for example, const int is a distinct type from int . You can use the C++ const\_cast operator on those rare occasions when you must remove const-ness from a variable. For more information, see [Type Conversions and Type Safety](#).

## String types

Strictly speaking, the C++ language has no built-in string type; `char` and `wchar_t` store single characters - you must declare an array of these types to approximate a string, adding a terminating null value (for example, ASCII `'\0'`) to the array element one past the last valid character (also called a *C-style string*). C-style strings required much more code to be written or the use of external string utility library functions. But in modern C++, we have the Standard Library types `std::string` (for 8-bit `char`-type character strings) or `std::wstring` (for 16-bit `wchar_t`-type character strings). These C++ Standard Library containers can be thought of as native string types because they are part of the standard libraries that are included in any compliant C++ build environment. Simply use the `#include <string>` directive to make these types available in your program. (If you are using MFC or ATL, the `CString` class is also available, but is not part of the C++ standard.) The use of null-terminated character arrays (the C-style strings previously mentioned) is strongly discouraged in modern C++.

## User-defined types

When you define a `class`, `struct`, `union`, or `enum`, that construct is used in the rest of your code as if it were a fundamental type. It has a known size in memory, and certain rules about how it can be used apply to it for compile-time checking and, at runtime, for the life of your program. The primary differences between the fundamental built-in types and user-defined types are as follows:

- The compiler has no built-in knowledge of a user-defined type. It learns of the type when it first encounters the definition during the compilation process.
- You specify what operations can be performed on your type, and how it can be converted to other types, by defining (through overloading) the appropriate operators, either as class members or non-member functions. For more information, see [Function Overloading](#).
- They do not have to be statically typed (the rule that an object's type never changes). Through the mechanisms of *inheritance* and *polymorphism*, a variable declared as a user-defined type of class (referred to as an object instance of a class) might have a different type at run-time than at compile time. For more information, see [Inheritance](#).

## Pointer types

Dating back to the earliest versions of the C language, C++ continues to let you declare a variable of a pointer type by using the special declarator `*` (asterisk). A pointer type stores the address of the location in memory where the actual data value is stored. In modern C++, these are referred to as *raw pointers*, and are accessed in your code through special operators `*` (asterisk) or `->` (dash with greater-than). This is called *dereferencing*, and which one that you use depends on whether you are dereferencing a pointer to a scalar or a pointer to a member in an object. Working with pointer types has long been one of the most challenging and confusing aspects of C and C++ program development. This section outlines some facts and practices to help use raw pointers if you want to, but in modern C++ it's no longer required (or recommended) to use raw pointers for object ownership at all, due to the evolution of the [smart pointer](#) (discussed more at the end of this section). It is still useful and safe to use raw pointers for observing objects, but if you must use them for object ownership, you should do so with caution and very careful consideration of how the objects owned by them are created and destroyed.

The first thing that you should know is declaring a raw pointer variable will allocate only the memory that is required to store an address of the memory location that the pointer will be referring to when it is dereferenced. Allocation of the memory for the data value itself (also called *backing store*) is not yet allocated. In other words, by declaring a raw pointer variable, you are creating a memory address variable, not an actual data variable. Dereferencing a pointer variable before making sure that it contains a valid address to a backing store will cause undefined behavior (usually a fatal error) in your program. The following example demonstrates this kind of error:

```
int* pNumber;      // Declare a pointer-to-int variable.  
*pNumber = 10;    // error. Although this may compile, it is  
// a serious error. We are dereferencing an  
// uninitialized pointer variable with no  
// allocated memory to point to.
```

The example dereferences a pointer type without having any memory allocated to store the actual integer data or a valid memory address assigned to it. The following code corrects these errors:

```
int number = 10;          // Declare and initialize a local integer  
// variable for data backing store.  
int* pNumber = &number;   // Declare and initialize a local integer  
// pointer variable to a valid memory  
// address to that backing store.  
...  
*pNumber = 41;           // Dereference and store a new value in  
// the memory pointed to by  
// pNumber, the integer variable called  
// "number". Note "number" was changed, not  
// "pNumber".
```

The corrected code example uses local stack memory to create the backing store that `pNumber` points to. We use a fundamental type for simplicity. In practice, the backing store for pointers are most often user-defined types that are dynamically-allocated in an area of memory called the *heap* (or *free store*) by using a `new` keyword expression (in C-style programming, the older `malloc()` C runtime library function was used). Once allocated, these variables are usually referred to as objects, especially if they are based on a class definition. Memory that is allocated with `new` must be deleted by a corresponding `delete` statement (or, if you used the `malloc()` function to allocate it, the C runtime function `free()`).

However, it is easy to forget to delete a dynamically-allocated object- especially in complex code, which causes a resource bug called a *memory leak*. For this reason, the use of raw pointers is strongly discouraged in modern C++. It is almost always better to wrap a raw pointer in a *smart pointer*, which will automatically release the memory when its destructor is invoked (when the code goes out of scope for the smart pointer); by using smart pointers you virtually eliminate a whole class of bugs in your C++ programs. In the following example, assume `MyClass` is a user-defined type that has a public method `DoSomeWork();`

```
void someFunction() {  
    unique_ptr<MyClass> pMc(new MyClass);  
    pMc->DoSomeWork();  
}  
// No memory leak. Out-of-scope automatically calls the destructor  
// for the unique_ptr, freeing the resource.
```

For more information about smart pointers, see [Smart Pointers](#).

For more information about pointer conversions, see [Type Conversions and Type Safety](#).

For more information about pointers in general, see [Pointers](#).

## Windows data types

In classic Win32 programming for C and C++, most functions use Windows-specific `typedefs` and `#define` macros (defined in `winddef.h`) to specify the types of parameters and return values. These Windows data types are mostly just special names (aliases) given to C/C++ built-in types. For a complete list of these `typedefs` and `preprocessor`

definitions, see [Windows Data Types](#). Some of these typedefs, such as HRESULT and LCID, are useful and descriptive. Others, such as INT, have no special meaning and are just aliases for fundamental C++ types. Other Windows data types have names that are retained from the days of C programming and 16-bit processors, and have no purpose or meaning on modern hardware or operating systems. There are also special data types associated with the Windows Runtime Library, listed as [Windows Runtime base data types](#). In modern C++, the general guideline is to prefer the C++ fundamental types unless the Windows type communicates some additional meaning about how the value is to be interpreted.

## More Information

For more information about the C++ type system, see the following topics.

<a href="#">Value Types</a>	Describes <i>value types</i> along with issues relating to their use.
<a href="#">Type Conversions and Type Safety</a>	Describes common type conversion issues and shows how to avoid them.

## See Also

- [Welcome Back to C++](#)
- [C++ Language Reference](#)
- [C++ Standard Library](#)

# Value Types (Modern C++)

12/21/2017 • 3 min to read • [Edit Online](#)

C++ classes are by default value types. This topic provides an introductory overview of value types and issues relating to their use.

## Value vs. reference types

As previously stated, C++ classes are by default value types. They can be specified as reference types, which enable polymorphic behavior to support object-oriented programming. Value types are sometimes viewed from the perspective of memory and layout control, whereas reference types are about base classes and virtual functions for polymorphic purposes. By default, value types are copyable, which means there is always a copy constructor and a copy assignment operator. For reference types, you make the class non-copyable (disable the copy constructor and copy assignment operator) and use a virtual destructor, which supports their intended polymorphism. Value types are also about the contents, which, when they are copied, always give you two independent values that can be modified separately. Reference types are about identity - what kind of object is it? For this reason, "reference types" are also referred to as "polymorphic types".

If you really want a reference-like type (base class, virtual functions), you need to explicitly disable copying, as shown in the `MyRefType` class in the following code.

```
// cl /EHsc /nologo /W4

class MyRefType {
private:
    MyRefType & operator=(const MyRefType &);
    MyRefType(const MyRefType &);

public:
    MyRefType () {}
};

int main()
{
    MyRefType Data1, Data2;
    // ...
    Data1 = Data2;
}
```

Compiling the above code will result in the following error:

```
test.cpp(15) : error C2248: 'MyRefType::operator =' : cannot access private member declared in class
'MyRefType'
        meow.cpp(5) : see declaration of 'MyRefType::operator ='
        meow.cpp(3) : see declaration of 'MyRefType'
```

## Value types and move efficiency

Copy allocation overhead is avoided due to new copy optimizations. For example, when you insert a string in the middle of a vector of strings, there will be no copy re-allocation overhead, only a move- even if it results in a grow of the vector itself. This also applies to other operations, for instance performing an add operation on two very large objects. How do you enable these value operation optimizations? In some C++ compilers, the compiler will enable this for you implicitly, much like copy constructors can be automatically generated by the compiler.

However, in Visual C++, your class will need to "opt-in" to move assignment and constructors by declaring it in your class definition. This is accomplished by using the double ampersand (`&&`) rvalue reference in the appropriate member function declarations and defining move constructor and move assignment methods. You also need to insert the correct code to "steal the guts" out of the source object.

How do you decide if you need move enabled? If you already know you need copy construction enabled, you probably want move enabled if it can be cheaper than a deep copy. However, if you know you need move support, it doesn't necessarily mean you want copy enabled. This latter case would be called a "move-only type". An example already in the standard library is `unique_ptr`. As a side note, the old `auto_ptr` is deprecated, and was replaced by `unique_ptr` precisely due to the lack of move semantics support in the previous version of C++.

By using move semantics you can return-by-value or insert-in-middle. Move is an optimization of copy. There is need for heap allocation as a workaround. Consider the following pseudocode:

```
#include <set>
#include <vector>
#include <string>
using namespace std;

//...
set<widget> LoadHugeData() {
    set<widget> ret;
    // ... load data from disk and populate ret
    return ret;
}
//...
widgets = LoadHugeData(); // efficient, no deep copy

vector<string> v = IfIHadAMillionStrings();
v.insert( begin(v)+v.size()/2, "scott" ); // efficient, no deep copy-shuffle
v.insert( begin(v)+v.size()/2, "Andrei" ); // (just 1M ptr/len assignments)
//...
HugeMatrix operator+(const HugeMatrix&, const HugeMatrix& );
HugeMatrix operator+(const HugeMatrix&, HugeMatrix&& );
HugeMatrix operator+( HugeMatrix&&, const HugeMatrix& );
HugeMatrix operator+( HugeMatrix&&, HugeMatrix&& );
//...
hm5 = hm1+hm2+hm3+hm4+hm5; // efficient, no extra copies
```

## Enabling move for appropriate value types

For a value-like class where move can be cheaper than a deep copy, enable move construction and move assignment for efficiency. Consider the following pseudocode:

```
#include <memory>
#include <stdexcept>
using namespace std;
// ...
class my_class {
    unique_ptr<BigHugeData> data;
public:
    my_class( my_class&& other ) // move construction
        : data( move( other.data ) ) { }
    my_class& operator=( my_class&& other ) // move assignment
    { data = move( other.data ); return *this; }
    // ...
    void method() { // check (if appropriate)
        if( !data )
            throw std::runtime_error("RUNTIME ERROR: Insufficient resources!");
    }
};
```

If you enable copy construction/assignment, also enable move construction/assignment if it can be cheaper than a deep copy.

Some *non-value* types are move-only, such as when you can't clone a resource, only transfer ownership. Example:  
`unique_ptr`.

## Section

Content

## See Also

[C++ Type System](#)

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# Type Conversions and Type Safety (Modern C++)

12/21/2017 • 9 min to read • [Edit Online](#)

This document identifies common type conversion problems and describes how you can avoid them in your C++ code.

When you write a C++ program, it's important to ensure that it's type-safe. This means that every variable, function argument, and function return value is storing an acceptable kind of data, and that operations that involve values of different types "make sense" and don't cause data loss, incorrect interpretation of bit patterns, or memory corruption. A program that never explicitly or implicitly converts values from one type to another is type-safe by definition. However, type conversions, even unsafe conversions, are sometimes required. For example, you might have to store the result of a floating point operation in a variable of type `int`, or you might have to pass the value in an unsigned `int` to a function that takes a signed `int`. Both examples illustrate unsafe conversions because they may cause data loss or re-interpretation of a value.

When the compiler detects an unsafe conversion, it issues either an error or a warning. An error stops compilation; a warning allows compilation to continue but indicates a possible error in the code. However, even if your program compiles without warnings, it still may contain code that leads to implicit type conversions that produce incorrect results. Type errors can also be introduced by explicit conversions, or casts, in the code.

## Implicit type conversions

When an expression contains operands of different built-in types, and no explicit casts are present, the compiler uses built-in *standard conversions* to convert one of the operands so that the types match. The compiler tries the conversions in a well-defined sequence until one succeeds. If the selected conversion is a promotion, the compiler does not issue a warning. If the conversion is a narrowing, the compiler issues a warning about possible data loss. Whether actual data loss occurs depends on the actual values involved, but we recommend that you treat this warning as an error. If a user-defined type is involved, then the compiler tries to use the conversions that you have specified in the class definition. If it can't find an acceptable conversion, the compiler issues an error and does not compile the program. For more information about the rules that govern the standard conversions, see [Standard Conversions](#). For more information about user-defined conversions, see [User-Defined Conversions \(C++/CLI\)](#).

### Widening conversions (promotion)

In a widening conversion, a value in a smaller variable is assigned to a larger variable with no loss of data. Because widening conversions are always safe, the compiler performs them silently and does not issue warnings. The following conversions are widening conversions.

FROM	TO
Any signed or unsigned integral type except <code>long long</code> or <code>_int64</code>	<code>double</code>
<code>bool</code> or <code>char</code>	Any other built-in type
<code>short</code> or <code>wchar_t</code>	<code>int</code> , <code>long</code> , <code>long long</code>
<code>int</code> , <code>long</code>	<code>long long</code>
<code>float</code>	<code>double</code>

## Narrowing conversions (coercion)

The compiler performs narrowing conversions implicitly, but it warns you about potential data loss. Take these warnings very seriously. If you are certain that no data loss will occur because the values in the larger variable will always fit in the smaller variable, then add an explicit cast so that the compiler will no longer issue a warning. If you are not sure that the conversion is safe, add to your code some kind of runtime check to handle possible data loss so that it does not cause your program to produce incorrect results.

Any conversion from a floating point type to an integral type is a narrowing conversion because the fractional portion of the floating point value is discarded and lost.

The following code example shows some implicit narrowing conversions, and the warnings that the compiler issues for them.

```
int i = INT_MAX + 1; //warning C4307:'+'integral constant overflow
wchar_t wch = 'A'; //OK
char c = wch; // warning C4244:'initializing':conversion from 'wchar_t'
               // to 'char', possible loss of data
unsigned char c2 = 0xffff; //warning C4305:'initializing':truncation from
                         // 'int' to 'unsigned char'
int j = 1.9f; // warning C4244:'initializing':conversion from 'float' to
               // 'int', possible loss of data
int k = 7.7; // warning C4244:'initializing':conversion from 'double' to
               // 'int', possible loss of data
```

## Signed - unsigned conversions

A signed integral type and its unsigned counterpart are always the same size, but they differ in how the bit pattern is interpreted for value transformation. The following code example demonstrates what happens when the same bit pattern is interpreted as a signed value and as an unsigned value. The bit pattern stored in both `num` and `num2` never changes from what is shown in the earlier illustration.

```
using namespace std;
unsigned short num = numeric_limits<unsigned short>::max(); // #include <limits>
short num2 = num;
cout << "unsigned val = " << num << " signed val = " << num2 << endl;
// Prints: unsigned val = 65535 signed val = -1

// Go the other way.
num2 = -1;
num = num2;
cout << "unsigned val = " << num << " signed val = " << num2 << endl;
// Prints: unsigned val = 65535 signed val = -1
```

Notice that values are reinterpreted in both directions. If your program produces odd results in which the sign of the value seems inverted from what you expect, look for implicit conversions between signed and unsigned integral types. In the following example, the result of the expression `(0 - 1)` is implicitly converted from `int` to `unsigned int` when it's stored in `num`. This causes the bit pattern to be reinterpreted.

```
unsigned int u3 = 0 - 1;
cout << u3 << endl; // prints 4294967295
```

The compiler does not warn about implicit conversions between signed and unsigned integral types. Therefore, we recommend that you avoid signed-to-unsigned conversions altogether. If you can't avoid them, then add to your code a runtime check to detect whether the value being converted is greater than or equal to zero and less than or equal to the maximum value of the signed type. Values in this range will transfer from signed to unsigned or from unsigned to signed without being reinterpreted.

## Pointer conversions

In many expressions, a C-style array is implicitly converted to a pointer to the first element in the array, and constant conversions can happen silently. Although this is convenient, it's also potentially error-prone. For example, the following badly designed code example seems nonsensical, and yet it will compile in Visual C++ and produces a result of 'p'. First, the "Help" string constant literal is converted to a `char*` that points to the first element of the array; that pointer is then incremented by three elements so that it now points to the last element 'p'.

```
char* s = "Help" + 3;
```

## Explicit conversions (casts)

By using a cast operation, you can instruct the compiler to convert a value of one type to another type. The compiler will raise an error in some cases if the two types are completely unrelated, but in other cases it will not raise an error even if the operation is not type-safe. Use casts sparingly because any conversion from one type to another is a potential source of program error. However, casts are sometimes required, and not all casts are equally dangerous. One effective use of a cast is when your code performs a narrowing conversion and you know that the conversion is not causing your program to produce incorrect results. In effect, this tells the compiler that you know what you are doing and to stop bothering you with warnings about it. Another use is to cast from a pointer-to-derived class to a pointer-to-base class. Another use is to cast away the `const`-ness of a variable to pass it to a function that requires a non-`const` argument. Most of these cast operations involve some risk.

In C-style programming, the same C-style cast operator is used for all kinds of casts.

```
(int) x; // old-style cast, old-style syntax  
int(x); // old-style cast, functional syntax
```

The C-style cast operator is identical to the call operator () and is therefore inconspicuous in code and easy to overlook. Both are bad because they're difficult to recognize at a glance or search for, and they're disparate enough to invoke any combination of `static`, `const`, and `reinterpret_cast`. Figuring out what an old-style cast actually does can be difficult and error-prone. For all these reasons, when a cast is required, we recommend that you use one of the following C++ cast operators, which in some cases are significantly more type-safe, and which express much more explicitly the programming intent:

- `static_cast`, for casts that are checked at compile time only. `static_cast` returns an error if the compiler detects that you are trying to cast between types that are completely incompatible. You can also use it to cast between pointer-to-base and pointer-to-derived, but the compiler can't always tell whether such conversions will be safe at runtime.

```
double d = 1.58947;  
int i = d; // warning C4244 possible loss of data  
int j = static_cast<int>(d); // No warning.  
string s = static_cast<string>(d); // Error C2440:cannot convert from  
// double to std::string  
  
// No error but not necessarily safe.  
Base* b = new Base();  
Derived* d2 = static_cast<Derived*>(b);
```

For more information, see [static\\_cast](#).

- `dynamic_cast`, for safe, runtime-checked casts of pointer-to-base to pointer-to-derived. A `dynamic_cast` is safer than a `static_cast` for downcasts, but the runtime check incurs some overhead.

```

Base* b = new Base();

// Run-time check to determine whether b is actually a Derived*
Derived* d3 = dynamic_cast<Derived*>(b);

// If b was originally a Derived*, then d3 is a valid pointer.
if(d3)
{
    // Safe to call Derived method.
    cout << d3->DoSomethingMore() << endl;
}
else
{
    // Run-time check failed.
    cout << "d3 is null" << endl;
}

//Output: d3 is null;

```

For more information, see [dynamic\\_cast](#).

- `const_cast`, for casting away the `const`-ness of a variable, or converting a non-`const` variable to be `const`. Casting away `const`-ness by using this operator is just as error-prone as is using a C-style cast, except that with `const-cast` you are less likely to perform the cast accidentally. Sometimes you have to cast away the `const`-ness of a variable, for example, to pass a `const` variable to a function that takes a non-`const` parameter. The following example shows how to do this.

```

void Func(double& d) { ... }
void ConstCast()
{
    const double pi = 3.14;
    Func(const_cast<double&>(pi)); //No error.
}

```

For more information, see [const\\_cast](#).

- `reinterpret_cast`, for casts between unrelated types such as `pointer` to `int`.

#### **NOTE**

This cast operator is not used as often as the others, and it's not guaranteed to be portable to other compilers.

The following example illustrates how `reinterpret_cast` differs from `static_cast`.

```

const char* str = "hello";
int i = static_cast<int>(str); //error C2440: 'static_cast' : cannot
                                // convert from 'const char *' to 'int'
int j = (int)str; // C-style cast. Did the programmer really intend
                  // to do this?
int k = reinterpret_cast<int>(str); // Programming intent is clear.
                                    // However, it is not 64-bit safe.

```

For more information, see [reinterpret\\_cast Operator](#).

## See Also

[C++ Type System](#)

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# Uniform Initialization and Delegating Constructors

12/21/2017 • 5 min to read • [Edit Online](#)

In modern C++, you can use *brace initialization* for any type, without the equals sign. Also, you can use delegating constructors to simplify your code when you have multiple constructors that perform similar work.

## Brace Initialization

You can use brace initialization for any class, struct, or union. If a type has a default constructor, either implicitly or explicitly declared, you can use default brace initialization (with empty braces). For example, the following class may be initialized by using both default and non-default brace initialization:

```
#include <string>
using namespace std;

class class_a {
public:
    class_a()
    class_a(string str) : m_string{ str } {}
    class_a(string str, double dbl) : m_string{ str }, m_double{ dbl } {}
double m_double;
string m_string;
};

int main()
{
    class_a c1{};
    class_a c1_1;

    class_a c2{ "ww" };
    class_a c2_1("xx");

    // order of parameters is the same as the constructor
    class_a c3{ "yy", 4.4 };
    class_a c3_1("zz", 5.5);
}
```

If a class has non-default constructors, the order in which class members appear in the brace initializer is the order in which the corresponding parameters appear in the constructor, not the order in which the members are declared (as with `class_a` in the previous example). Otherwise, if the type has no declared constructor, the order in which the members appear in the brace initializer is the same as the order in which they are declared; in this case, you can initialize as many of the public members as you wish, but you cannot skip any member. The following example shows the order that's used in brace initialization when there is no declared constructor:

```

class class_d {
public:
    float m_float;
    string m_string;
    wchar_t m_char;
};

int main()
{
    class_d d1{};
    class_d d1{ 4.5 };
    class_d d2{ 4.5, "string" };
    class_d d3{ 4.5, "string", 'c' };

    class_d d4{ "string", 'c' }; // compiler error
    class_d d5("string", 'c', 2.0 ); // compiler error
}

```

If the default constructor is explicitly declared but marked as deleted, default brace initialization cannot be used:

```

class class_f {
public:
    class_f() = delete;
    class_f(string x): m_string { x } {}
    string m_string;
};
int main()
{
    class_f cf{ "hello" };
    class_f cf1{}; // compiler error C2280: attempting to reference a deleted function
}

```

You can use brace initialization anywhere you would typically do initialization—for example, as a function parameter or a return value, or with the `new` keyword:

```

class_d* cf = new class_d{4.5};
kr->add_d({ 4.5 });
return { 4.5 };

```

## initializer\_list Constructors

The [initializer\\_list Class](#) represents a list of objects of a specified type that can be used in a constructor, and in other contexts. You can construct an `initializer_list` by using brace initialization:

```
initializer_list<int> int_list{5, 6, 7};
```

### IMPORTANT

To use this class, you must include the `<initializer_list>` header.

An `initializer_list` can be copied. In this case, the members of the new list are references to the members of the original list:

```
initializer_list<int> ilist1{ 5, 6, 7 };
initializer_list<int> ilist2( ilist1 );
if (ilist1.begin() == ilist2.begin())
    cout << "yes" << endl; // expect "yes"
```

The standard library container classes, and also `string`, `wstring`, and `regex`, have `initializer_list` constructors. The following examples show how to do brace initialization with these constructors:

```
vector<int> v1{ 9, 10, 11 };
map<int, string> m1{ {1, "a"}, {2, "b"} };
string s{ 'a', 'b', 'c' };
regex rgx{'x', 'y', 'z'};
```

## Delegating Constructors

Many classes have multiple constructors that do similar things—for example, validate parameters:

```
class class_c {
public:
    int max;
    int min;
    int middle;

    class_c()
    class_c(int my_max) {
        max = my_max > 0 ? my_max : 10;
    }
    class_c(int my_max, int my_min) {
        max = my_max > 0 ? my_max : 10;
        min = my_min > 0 && my_min < max ? my_min : 1;
    }
    class_c(int my_max, int my_min, int my_middle) {
        max = my_max > 0 ? my_max : 10;
        min = my_min > 0 && my_min < max ? my_min : 1;
        middle = my_middle < max && my_middle > min ? my_middle : 5;
    }
};
```

You could reduce the repetitive code by adding a function that does all of the validation, but the code for `class_c` would be easier to understand and maintain if one constructor could delegate some of the work to another one. To add delegating constructors, use the `constructor (. . .) : constructor (. . .)` syntax:

```

class class_c {
public:
    int max;
    int min;
    int middle;

    class_c(int my_max) {
        max = my_max > 0 ? my_max : 10;
    }
    class_c(int my_max, int my_min) : class_c(my_max) {
        min = my_min > 0 && my_min < max ? my_min : 1;
    }
    class_c(int my_max, int my_min, int my_middle) : class_c (my_max, my_min){
        middle = my_middle < max && my_middle > min ? my_middle : 5;
    }
};
int main() {

    class_c c1{ 1, 3, 2 };
}

```

As you step through the previous example, notice that the constructor `class_c(int, int, int)` first calls the constructor `class_c(int, int)`, which in turn calls `class_c(int)`. Each of the constructors performs only the work that is not performed by the other constructors.

The first constructor that's called initializes the object so that all of its members are initialized at that point. You can't do member initialization in a constructor that delegates to another constructor, as shown here:

```

class class_a {
public:
    class_a() {}
    // member initialization here, no delegate
    class_a(string str) : m_string{ str } {}

    //can't do member initialization here
    // error C3511: a call to a delegating constructor shall be the only member-initializer
    class_a(string str, double dbl) : class_a(str) , m_double{ dbl } {}

    // only member assignment
    class_a(string str, double dbl) : class_a(str) { m_double = dbl; }
    double m_double{ 1.0 };
    string m_string;
};

```

The next example shows the use of non-static data-member initializers. Notice that if a constructor also initializes a given data member, the member initializer is overridden:

```

class class_a {
public:
    class_a() {}
    class_a(string str) : m_string{ str } {}
    class_a(string str, double dbl) : class_a(str) { m_double = dbl; }
    double m_double{ 1.0 };
    string m_string{ m_double < 10.0 ? "alpha" : "beta" };
};

int main() {
    class_a a{ "hello", 2.0 }; //expect a.m_double == 2.0, a.m_string == "hello"
    int y = 4;
}

```

The constructor delegation syntax doesn't prevent the accidental creation of constructor recursion—Constructor1 calls Constructor2 which calls Constructor1—and no errors are thrown until there is a stack overflow. It's your responsibility to avoid cycles.

```
class class_f{
public:
    int max;
    int min;

    // don't do this
    class_f() : class_f(6, 3){ }
    class_f(int my_max, int my_min) : class_f() { }
};
```

# Object Lifetime And Resource Management (Modern C++)

12/21/2017 • 4 min to read • [Edit Online](#)

Unlike managed languages, C++ doesn't have garbage collection (GC), which automatically releases no-longer-used memory resources as a program runs. In C++, resource management is directly related to object lifetime. This document describes the factors that affect object lifetime in C++ and how to manage it.

C++ doesn't have GC primarily because it doesn't handle non-memory resources. Only deterministic destructors like those in C++ can handle memory and non-memory resources equally. GC also has other problems, like higher overhead in memory and CPU consumption, and locality. But universality is a fundamental problem that can't be mitigated through clever optimizations.

## Concepts

An important thing in object-lifetime management is the encapsulation—whatever's using an object doesn't have to know what resources that object owns, or how to get rid of them, or even whether it owns any resources at all. It just has to destroy the object. The C++ core language is designed to ensure that objects are destroyed at the correct times, that is, as blocks are exited, in reverse order of construction. When an object is destroyed, its bases and members are destroyed in a particular order. The language automatically destroys objects, unless you do special things like heap allocation or placement new. For example, [smart pointers](#) like `unique_ptr` and `shared_ptr`, and C++ Standard Library containers like `vector`, encapsulate `new / delete` and `new[] / delete[]` in objects, which have destructors. That's why it's so important to use smart pointers and C++ Standard Library containers.

Another important concept in lifetime management: destructors. Destructors encapsulate resource release. (The commonly used mnemonic is RRID, Resource Release Is Destruction.) A resource is something that you get from "the system" and have to give back later. Memory is the most common resource, but there are also files, sockets, textures, and other non-memory resources. "Owning" a resource means you can use it when you need it but you also have to release it when you're finished with it. When an object is destroyed, its destructor releases the resources that it owned.

The final concept is the DAG (Directed Acyclic Graph). The structure of ownership in a program forms a DAG. No object can own itself—that's not only impossible but also inherently meaningless. But two objects can share ownership of a third object. Several kinds of links are possible in a DAG like this: A is a member of B (B owns A), C stores a `vector<D>` (C owns each D element), E stores a `shared_ptr<F>` (E shares ownership of F, possibly with other objects), and so forth. As long as there are no cycles and every link in the DAG is represented by an object that has a destructor (instead of a raw pointer, handle, or other mechanism), then resource leaks are impossible because the language prevents them. Resources are released immediately after they're no longer needed, without a garbage collector running. The lifetime tracking is overhead-free for stack scope, bases, members, and related cases, and inexpensive for `shared_ptr`.

### Heap-based lifetime

For heap object lifetime, use [smart pointers](#). Use `shared_ptr` and `make_shared` as the default pointer and allocator. Use `weak_ptr` to break cycles, do caching, and observe objects without affecting or assuming anything about their lifetimes.

```

void func() {

    auto p = make_shared<widget>(); // no leak, and exception safe
    ...
    p->draw();

} // no delete required, out-of-scope triggers smart pointer destructor

```

Use `unique_ptr` for unique ownership, for example, in the *pimpl* idiom. (See [Pimpl For Compile-Time Encapsulation](#).) Make a `unique_ptr` the primary target of all explicit `new` expressions.

```
unique_ptr<widget> p(new widget());
```

You can use raw pointers for non-ownership and observation. A non-owning pointer may dangle, but it can't leak.

```

class node {
    ...
    vector<unique_ptr<node>> children; // node owns children
    node* parent; // node observes parent, which is not a concern
    ...
};

node::node() : parent(...) { children.emplace_back(new node(...)); }

```

When performance optimization is required, you might have to use *well-encapsulated* owning pointers and explicit calls to delete. An example is when you implement your own low-level data structure.

### Stack-based lifetime

In modern C++, *stack-based scope* is a powerful way to write robust code because it combines automatic *stack lifetime* and *data member lifetime* with high efficiency—lifetime tracking is essentially free of overhead. Heap object lifetime requires diligent manual management and can be the source of resource leaks and inefficiencies, especially when you are working with raw pointers. Consider this code, which demonstrates stack-based scope:

```

class widget {
private:
    gadget g; // lifetime automatically tied to enclosing object
public:
    void draw();
};

void functionUsingWidget () {
    widget w; // lifetime automatically tied to enclosing scope
    // constructs w, including the w.g gadget member
    // ...
    w.draw();
    // ...
} // automatic destruction and deallocation for w and w.g
// automatic exception safety,
// as if "finally { w.dispose(); w.g.dispose(); }"

```

Use static lifetime sparingly (global static, function local static) because problems can arise. What happens when the constructor of a global object throws an exception? Typically, the app faults in a way that can be difficult to debug. Construction order is problematic for static lifetime objects, and is not concurrency-safe. Not only is object construction a problem, destruction order can be complex, especially where polymorphism is involved. Even if your object or variable isn't polymorphic and doesn't have complex construction/destruction ordering, there's still the issue of thread-safe concurrency. A multithreaded app can't safely modify the data in static objects without having thread-local storage, resource locks, and other special precautions.

## See Also

[Welcome Back to C++](#)  
[C++ Language Reference](#)  
[C++ Standard Library](#)

# Objects Own Resources (RAII)

12/21/2017 • 1 min to read • [Edit Online](#)

Make sure that objects own resources. This principle is also known as "resource acquisition is initialization" or "RAII."

## Example

Pass every "new" object as a constructor argument to another named object that owns it (almost always unique\_ptr).

```
void f() {
    unique_ptr<widget> p( new widget() );
    my_class x( new widget() );
    // ...
} // automatic destruction and deallocation for both widget objects
// automatic exception safety, as if "finally { p->dispose(); x.w.dispose(); }"
```

Always immediately pass any new resource to another object that owns it.

```
void g() {
    other_class y( OpenFile() );
    // ...
} // automatic closing and release for file resource
// automatic exception safety, as if "finally { y.file.dispose(); }"
```

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# Smart Pointers (Modern C++)

12/21/2017 • 7 min to read • [Edit Online](#)

In modern C++ programming, the Standard Library includes *smart pointers*, which are used to help ensure that programs are free of memory and resource leaks and are exception-safe.

## Uses for smart pointers

Smart pointers are defined in the `std` namespace in the `<memory>` header file. They are crucial to the **RAll** or *Resource Acquisition Is Initialization* programming idiom. The main goal of this idiom is to ensure that resource acquisition occurs at the same time that the object is initialized, so that all resources for the object are created and made ready in one line of code. In practical terms, the main principle of RAll is to give ownership of any heap-allocated resource—for example, dynamically-allocated memory or system object handles—to a stack-allocated object whose destructor contains the code to delete or free the resource and also any associated cleanup code.

In most cases, when you initialize a raw pointer or resource handle to point to an actual resource, pass the pointer to a smart pointer immediately. In modern C++, raw pointers are only used in small code blocks of limited scope, loops, or helper functions where performance is critical and there is no chance of confusion about ownership.

The following example compares a raw pointer declaration to a smart pointer declaration.

```
void UseRawPointer()
{
    // Using a raw pointer -- not recommended.
    Song* pSong = new Song(L"Nothing on You", L"Bruno Mars");

    // Use pSong...

    // Don't forget to delete!
    delete pSong;
}

void UseSmartPointer()
{
    // Declare a smart pointer on stack and pass it the raw pointer.
    unique_ptr<Song> song2(new Song(L"Nothing on You", L"Bruno Mars"));

    // Use song2...
    wstring s = song2->duration_;
    //...

} // song2 is deleted automatically here.
```

As shown in the example, a smart pointer is a class template that you declare on the stack, and initialize by using a raw pointer that points to a heap-allocated object. After the smart pointer is initialized, it owns the raw pointer. This means that the smart pointer is responsible for deleting the memory that the raw pointer specifies. The smart pointer destructor contains the call to delete, and because the smart pointer is declared on the stack, its destructor is invoked when the smart pointer goes out of scope, even if an exception is thrown somewhere further up the stack.

Access the encapsulated pointer by using the familiar pointer operators, `->` and `*`, which the smart pointer

class overloads to return the encapsulated raw pointer.

The C++ smart pointer idiom resembles object creation in languages such as C#: you create the object and then let the system take care of deleting it at the correct time. The difference is that no separate garbage collector runs in the background; memory is managed through the standard C++ scoping rules so that the runtime environment is faster and more efficient.

#### IMPORTANT

Always create smart pointers on a separate line of code, never in a parameter list, so that a subtle resource leak won't occur due to certain parameter list allocation rules.

The following example shows how a `unique_ptr` smart pointer type from the C++ Standard Library could be used to encapsulate a pointer to a large object.

```
class LargeObject
{
public:
    void DoSomething(){}
};

void ProcessLargeObject(const LargeObject& lo){}
void SmartPointerDemo()
{
    // Create the object and pass it to a smart pointer
    std::unique_ptr<LargeObject> pLarge(new LargeObject());

    //Call a method on the object
    pLarge->DoSomething();

    // Pass a reference to a method.
    ProcessLargeObject(*pLarge);

} //pLarge is deleted automatically when function block goes out of scope.
```

The example demonstrates the following essential steps for using smart pointers.

1. Declare the smart pointer as an automatic (local) variable. (Do not use the `new` or `malloc` expression on the smart pointer itself.)
2. In the type parameter, specify the pointed-to type of the encapsulated pointer.
3. Pass a raw pointer to a `new`-ed object in the smart pointer constructor. (Some utility functions or smart pointer constructors do this for you.)
4. Use the overloaded `->` and `*` operators to access the object.
5. Let the smart pointer delete the object.

Smart pointers are designed to be as efficient as possible both in terms of memory and performance. For example, the only data member in `unique_ptr` is the encapsulated pointer. This means that `unique_ptr` is exactly the same size as that pointer, either four bytes or eight bytes. Accessing the encapsulated pointer by using the smart pointer overloaded \* and -> operators is not significantly slower than accessing the raw pointers directly.

Smart pointers have their own member functions, which are accessed by using "dot" notation. For example, some C++ Standard Library smart pointers have a reset member function that releases ownership of the pointer. This is useful when you want to free the memory owned by the smart pointer

before the smart pointer goes out of scope, as shown in the following example.

```
void SmartPointerDemo2()
{
    // Create the object and pass it to a smart pointer
    std::unique_ptr<LargeObject> pLarge(new LargeObject());

    // Call a method on the object
    pLarge->DoSomething();

    // Free the memory before we exit function block.
    pLarge.reset();

    // Do some other work...
}
```

Smart pointers usually provide a way to access their raw pointer directly. C++ Standard Library smart pointers have a `get` member function for this purpose, and `CComPtr` has a public `p` class member. By providing direct access to the underlying pointer, you can use the smart pointer to manage memory in your own code and still pass the raw pointer to code that does not support smart pointers.

```
void SmartPointerDemo4()
{
    // Create the object and pass it to a smart pointer
    std::unique_ptr<LargeObject> pLarge(new LargeObject());

    // Call a method on the object
    pLarge->DoSomething();

    // Pass raw pointer to a legacy API
    LegacyLargeObjectFunction(pLarge.get());
}
```

## Kinds of Smart Pointers

The following section summarizes the different kinds of smart pointers that are available in the Windows programming environment and describes when to use them.

### C++ Standard Library Smart Pointers

Use these smart pointers as a first choice for encapsulating pointers to plain old C++ objects (POCO).

- `unique_ptr`

Allows exactly one owner of the underlying pointer. Use as the default choice for POCO unless you know for certain that you require a `shared_ptr`. Can be moved to a new owner, but not copied or shared.

Replaces `auto_ptr`, which is deprecated. Compare to `boost::scoped_ptr`. `unique_ptr` is small and efficient; the size is one pointer and it supports rvalue references for fast insertion and retrieval from C++ Standard Library collections. Header file: `<memory>`. For more information, see [How to: Create and Use unique\\_ptr Instances](#) and [unique\\_ptr Class](#).

- `shared_ptr`

Reference-counted smart pointer. Use when you want to assign one raw pointer to multiple owners, for example, when you return a copy of a pointer from a container but want to keep the original. The raw pointer is not deleted until all `shared_ptr` owners have gone out of scope or have otherwise given up ownership. The size is two pointers; one for the object and one for the shared control block that contains the reference count. Header file: `<memory>`. For more information, see [How to: Create and Use shared\\_ptr Instances](#) and [shared\\_ptr Class](#).

- `weak_ptr`

Special-case smart pointer for use in conjunction with `shared_ptr`. A `weak_ptr` provides access to an object that is owned by one or more `shared_ptr` instances, but does not participate in reference counting. Use when you want to observe an object, but do not require it to remain alive. Required in some cases to break circular references between `shared_ptr` instances. Header file: `<memory>`. For more information, see [How to: Create and Use `weak\_ptr` Instances and `weak\_ptr` Class](#).

## Smart Pointers for COM Objects (Classic Windows Programming)

When you work with COM objects, wrap the interface pointers in an appropriate smart pointer type. The Active Template Library (ATL) defines several smart pointers for various purposes. You can also use the `_com_ptr_t` smart pointer type, which the compiler uses when it creates wrapper classes from .tlb files. It's the best choice when you do not want to include the ATL header files.

### [CComPtr Class](#)

Use this unless you cannot use ATL. Performs reference counting by using the `AddRef` and `Release` methods. For more information, see [How to: Create and Use `CComPtr` and `CComQIPtr` Instances](#).

### [CComQIPtr Class](#)

Resembles `CComPtr` but also provides simplified syntax for calling `QueryInterface` on COM objects. For more information, see [How to: Create and Use `CComPtr` and `CComQIPtr` Instances](#).

### [CComHeapPtr Class](#)

Smart pointer to objects that use `CoTaskMemFree` to free memory.

### [CComGIPtr Class](#)

Smart pointer for interfaces that are obtained from the global interface table (GIT).

### [\\_com\\_ptr\\_t Class](#)

Resembles `CComQIPtr` in functionality but does not depend on ATL headers.

## ATL Smart Pointers for POCO Objects

In addition to smart pointers for COM objects, ATL also defines smart pointers, and collections of smart pointers, for plain old C++ objects. In classic Windows programming, these types are useful alternatives to the C++ Standard Library collections, especially when code portability is not required or when you do not want to mix the programming models of the C++ Standard Library and ATL.

### [CAutoPtr Class](#)

Smart pointer that enforces unique ownership by transferring ownership on copy. Comparable to the deprecated `std::auto_ptr` Class.

### [CHeapPtr Class](#)

Smart pointer for objects that are allocated by using the C `malloc` function.

### [CAutoVectorPtr Class](#)

Smart pointer for arrays that are allocated by using `new[]`.

### [CAutoPtrArray Class](#)

Class that encapsulates an array of `CAutoPtr` elements.

### [CAutoPtrList Class](#)

Class that encapsulates methods for manipulating a list of `CAutoPtr` nodes.

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)



# How to: Create and Use unique\_ptr Instances

12/21/2017 • 2 min to read • [Edit Online](#)

A `unique_ptr` does not share its pointer. It cannot be copied to another `unique_ptr`, passed by value to a function, or used in any C++ Standard Library algorithm that requires copies to be made. A `unique_ptr` can only be moved. This means that the ownership of the memory resource is transferred to another `unique_ptr` and the original `unique_ptr` no longer owns it. We recommend that you restrict an object to one owner, because multiple ownership adds complexity to the program logic. Therefore, when you need a smart pointer for a plain C++ object, use `unique_ptr`, and when you construct a `unique_ptr`, use the `make_unique` helper function.

The following diagram illustrates the transfer of ownership between two `unique_ptr` instances.

```
auto ptrA = make_unique<Song>(L"Diana Krall", L"The Look of Love");
```



```
auto ptrB = std::move(ptrA);
```



`unique_ptr` is defined in the `<memory>` header in the C++ Standard Library. It is exactly as efficient as a raw pointer and can be used in C++ Standard Library containers. The addition of `unique_ptr` instances to C++ Standard Library containers is efficient because the move constructor of the `unique_ptr` eliminates the need for a copy operation.

## Example

The following example shows how to create `unique_ptr` instances and pass them between functions.

```
unique_ptr<Song> SongFactory(const std::wstring& artist, const std::wstring& title)
{
    // Implicit move operation into the variable that stores the result.
    return make_unique<Song>(artist, title);
}

void MakeSongs()
{
    // Create a new unique_ptr with a new object.
    auto song = make_unique<Song>(L"Mr. Children", L"Namonaki Uta");

    // Use the unique_ptr.
    vector<wstring> titles = { song->title };

    // Move raw pointer from one unique_ptr to another.
    unique_ptr<Song> song2 = std::move(song);

    // Obtain unique_ptr from function that returns by value.
    auto song3 = SongFactory(L"Michael Jackson", L"Beat It");
}
```

These examples demonstrate this basic characteristic of `unique_ptr`: it can be moved, but not copied. "Moving" transfers ownership to a new `unique_ptr` and resets the old `unique_ptr`.

## Example

The following example shows how to create `unique_ptr` instances and use them in a vector.

```
void SongVector()
{
    vector<unique_ptr<Song>> songs;

    // Create a few new unique_ptr<Song> instances
    // and add them to vector using implicit move semantics.
    songs.push_back(make_unique<Song>(L"B'z", L"Juice"));
    songs.push_back(make_unique<Song>(L"Namie Amuro", L"Funky Town"));
    songs.push_back(make_unique<Song>(L"Kome Kome Club", L"Kimi ga Iru Dake de"));
    songs.push_back(make_unique<Song>(L"Ayumi Hamasaki", L"Poker Face"));

    // Pass by const reference when possible to avoid copying.
    for (const auto& song : songs)
    {
        wcout << L"Artist: " << song->artist << L"    Title: " << song->title << endl;
    }
}
```

In the range for loop, notice that the `unique_ptr` is passed by reference. If you try to pass by value here, the compiler will throw an error because the `unique_ptr` copy constructor is deleted.

## Example

The following example shows how to initialize a `unique_ptr` that is a class member.

```
class MyClass
{
private:
    // MyClass owns the unique_ptr.
    unique_ptr<ClassFactory> factory;
public:

    // Initialize by using make_unique with ClassFactory default constructor.
    MyClass() : factory ( make_unique<ClassFactory>())
    {

    }

    void MakeClass()
    {
        factory->DoSomething();
    }
};
```

## Example

You can use `make_unique` to create a `unique_ptr` to an array, but you cannot use `make_unique` to initialize the array elements.

```
// Create a unique_ptr to an array of 5 integers.  
auto p = make_unique<int[]>(5);  
  
// Initialize the array.  
for (int i = 0; i < 5; ++i)  
{  
    p[i] = i;  
    wcout << p[i] << endl;  
}
```

For more examples, see [make\\_unique](#).

## See Also

[Smart Pointers](#)

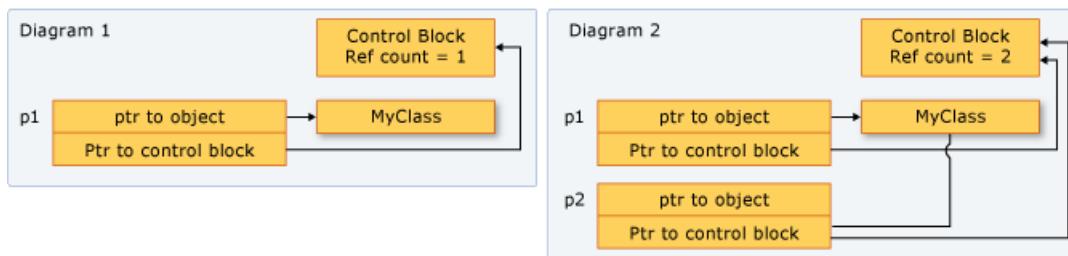
[make\\_unique](#)

# How to: Create and Use `shared_ptr` Instances

12/21/2017 • 6 min to read • [Edit Online](#)

The `shared_ptr` type is a smart pointer in the C++ standard library that is designed for scenarios in which more than one owner might have to manage the lifetime of the object in memory. After you initialize a `shared_ptr`, you can copy it, pass it by value in function arguments, and assign it to other `shared_ptr` instances. All the instances point to the same object, and share access to one "control block" that increments and decrements the reference count whenever a new `shared_ptr` is added, goes out of scope, or is reset. When the reference count reaches zero, the control block deletes the memory resource and itself.

The following illustration shows several `shared_ptr` instances that point to one memory location.



## Example

Whenever possible, use the `make_shared` function to create a `shared_ptr` when the memory resource is created for the first time. `make_shared` is exception-safe. It uses the same call to allocate the memory for the control block and the resource, and thereby reduces the construction overhead. If you do not use `make_shared`, then you have to use an explicit new expression to create the object before you pass it to the `shared_ptr` constructor. The following example shows various ways to declare and initialize a `shared_ptr` together with a new object.

```
// Use make_shared function when possible.
auto sp1 = make_shared<Song>(L"The Beatles", L"Im Happy Just to Dance With You");

// Ok, but slightly less efficient.
// Note: Using new expression as constructor argument
// creates no named variable for other code to access.
shared_ptr<Song> sp2(new Song(L"Lady Gaga", L"Just Dance"));

// When initialization must be separate from declaration, e.g. class members,
// initialize with nullptr to make your programming intent explicit.
shared_ptr<Song> sp5(nullptr);
//Equivalent to: shared_ptr<Song> sp5;
//...
sp5 = make_shared<Song>(L"Elton John", L"I'm Still Standing");
```

## Example

The following example shows how to declare and initialize `shared_ptr` instances that take on shared ownership of an object that has already been allocated by another `shared_ptr`. Assume that `sp2` is an initialized `shared_ptr`.

```

//Initialize with copy constructor. Increments ref count.
auto sp3(sp2);

//Initialize via assignment. Increments ref count.
auto sp4 = sp2;

//Initialize with nullptr. sp7 is empty.
shared_ptr<Song> sp7(nullptr);

// Initialize with another shared_ptr. sp1 and sp2
// swap pointers as well as ref counts.
sp1.swap(sp2);

```

## Example

`shared_ptr` is also helpful in C++ Standard Library containers when you are using algorithms that copy elements. You can wrap elements in a `shared_ptr`, and then copy it into other containers with the understanding that the underlying memory is valid as long as you need it, and no longer. The following example shows how to use the `replace_copy_if` algorithm on `shared_ptr` instances in a vector.

```

vector<shared_ptr<Song>> v;

v.push_back(make_shared<Song>(L"Bob Dylan", L"The Times They Are A Changing"));
v.push_back(make_shared<Song>(L"Aretha Franklin", L"Bridge Over Troubled Water"));
v.push_back(make_shared<Song>(L"Thalía", L"Entre El Mar y Una Estrella"));

vector<shared_ptr<Song>> v2;
remove_copy_if(v.begin(), v.end(), back_inserter(v2), [] (shared_ptr<Song> s)
{
    return s->artist.compare(L"Bob Dylan") == 0;
});

for (const auto& s : v2)
{
    wcout << s->artist << L":" << s->title << endl;
}

```

## Example

You can use `dynamic_pointer_cast`, `static_pointer_cast`, and `const_pointer_cast` to cast a `shared_ptr`. These functions resemble the `dynamic_cast`, `static_cast`, and `const_cast` operators. The following example shows how to test the derived type of each element in a vector of `shared_ptr` of base classes, and then copy the elements and display information about them.

```

vector<shared_ptr<MediaAsset>> assets;

assets.push_back(shared_ptr<Song>(new Song(L"Himesh Reshammiya", L"Tera Surroor")));
assets.push_back(shared_ptr<Song>(new Song(L"Penaz Masani", L"Tu Dil De De")));
assets.push_back(shared_ptr<Photo>(new Photo(L"2011-04-06", L"Redmond, WA", L"Soccer field at Microsoft.")));

vector<shared_ptr<MediaAsset>> photos;

copy_if(assets.begin(), assets.end(), back_inserter(photos), [] (shared_ptr<MediaAsset> p) -> bool
{
    // Use dynamic_pointer_cast to test whether
    // element is a shared_ptr<Photo>.
    shared_ptr<Photo> temp = dynamic_pointer_cast<Photo>(p);
    return temp.get() != nullptr;
});

for (const auto& p : photos)
{
    // We know that the photos vector contains only
    // shared_ptr<Photo> objects, so use static_cast.
    wcout << "Photo location: " << (static_pointer_cast<Photo>(p))->location_ << endl;
}

```

## Example

You can pass a `shared_ptr` to another function in the following ways:

- Pass the `shared_ptr` by value. This invokes the copy constructor, increments the reference count, and makes the callee an owner. There is a small amount of overhead in this operation, which may be significant depending on how many `shared_ptr` objects you are passing. Use this option when the code contract (implied or explicit) between the caller and callee requires that the callee be an owner.
- Pass the `shared_ptr` by reference or const reference. In this case, the reference count is not incremented, and the callee can access the pointer as long as the caller does not go out of scope. Or, the callee can decide to create a `shared_ptr` based on the reference, and thereby become a shared owner. Use this option when the caller has no knowledge of the callee, or when you must pass a `shared_ptr` and want to avoid the copy operation for performance reasons.
- Pass the underlying pointer or a reference to the underlying object. This enables the callee to use the object, but does not enable it to share ownership or extend the lifetime. If the callee creates a `shared_ptr` from the raw pointer, the new `shared_ptr` is independent from the original, and does not control the underlying resource. Use this option when the contract between the caller and callee clearly specifies that the caller retains ownership of the `shared_ptr` lifetime.
- When you are deciding how to pass a `shared_ptr`, determine whether the callee has to share ownership of the underlying resource. An "owner" is an object or function that can keep the underlying resource alive for as long as it needs it. If the caller has to guarantee that the callee can extend the life of the pointer beyond its (the function's) lifetime, use the first option. If you don't care whether the callee extends the lifetime, then pass by reference and let the callee copy it or not.
- If you have to give a helper function access to the underlying pointer, and you know that the helper function will just use the pointer and return before the calling function returns, then that function does not have to share ownership of the underlying pointer. It just has to access the pointer within the lifetime of the caller's `shared_ptr`. In this case, it is safe to pass the `shared_ptr` by reference, or pass the raw pointer or a reference to the underlying object. Passing this way provides a small performance benefit, and may also help you express your programming intent.
- Sometimes, for example in a `std::vector<shared_ptr<T>>`, you may have to pass each `shared_ptr` to a

lambda expression body or named function object. If the lambda or function is not storing the pointer, then pass the `shared_ptr` by reference to avoid invoking the copy constructor for each element.

## Example

The following example shows how `shared_ptr` overloads various comparison operators to enable pointer comparisons on the memory that is owned by the `shared_ptr` instances.

```
// Initialize two separate raw pointers.  
// Note that they contain the same values.  
auto song1 = new Song(L"Village People", L"YMCA");  
auto song2 = new Song(L"Village People", L"YMCA");  
  
// Create two unrelated shared_ptrs.  
shared_ptr<Song> p1(song1);  
shared_ptr<Song> p2(song2);  
  
// Unrelated shared_ptrs are never equal.  
wcout << "p1 < p2 = " << std::boolalpha << (p1 < p2) << endl;  
wcout << "p1 == p2 = " << std::boolalpha << (p1 == p2) << endl;  
  
// Related shared_ptr instances are always equal.  
shared_ptr<Song> p3(p2);  
wcout << "p3 == p2 = " << std::boolalpha << (p3 == p2) << endl;
```

## See Also

[Smart Pointers](#)

# How to: Create and Use weak\_ptr Instances

12/21/2017 • 3 min to read • [Edit Online](#)

Sometimes an object must store a way to access the underlying object of a `shared_ptr` without causing the reference count to be incremented. Typically, this situation occurs when you have cyclic references between `shared_ptr` instances.

The best design is to avoid shared ownership of pointers whenever you can. However, if you must have shared ownership of `shared_ptr` instances, avoid cyclic references between them. When cyclic references are unavoidable, or even preferable for some reason, use `weak_ptr` to give one or more of the owners a weak reference to another `shared_ptr`. By using a `weak_ptr`, you can create a `shared_ptr` that joins to an existing set of related instances, but only if the underlying memory resource is still valid. A `weak_ptr` itself does not participate in the reference counting, and therefore, it cannot prevent the reference count from going to zero. However, you can use a `weak_ptr` to try to obtain a new copy of the `shared_ptr` with which it was initialized. If the memory has already been deleted, a **bad\_weak\_ptr** exception is thrown. If the memory is still valid, the new shared pointer increments the reference count and guarantees that the memory will be valid as long as the `shared_ptr` variable stays in scope.

## Example

The following code example shows a case where `weak_ptr` is used to ensure proper deletion of objects that have circular dependencies. As you examine the example, assume that it was created only after alternative solutions were considered. The `Controller` objects represent some aspect of a machine process, and they operate independently. Each controller must be able to query the status of the other controllers at any time, and each one contains a private `vector<weak_ptr<Controller>>` for this purpose. Each vector contains a circular reference, and therefore, `weak_ptr` instances are used instead of `shared_ptr`.

```
#include <iostream>
#include <memory>
#include <string>
#include <vector>
#include <algorithm>

using namespace std;

class Controller
{
public:
    int Num;
    wstring Status;
    vector<weak_ptr<Controller>> others;
    explicit Controller(int i) : Num(i), Status(L"On")
    {
        wcout << L"Creating Controller" << Num << endl;
    }

    ~Controller()
    {
        wcout << L"Destroying Controller" << Num << endl;
    }

    // Demonstrates how to test whether the
    // pointed-to memory still exists or not.
    void CheckStatuses() const
    {
        for (auto& other : others)
        {
            if (other.expired())
                wcout << L"Object destroyed" << endl;
            else
                wcout << L"Object still exists" << endl;
        }
    }
}
```

```

    for_each(others.begin(), others.end(), [] (weak_ptr<Controller> wp)
    {
        try
        {
            auto p = wp.lock();
            wcout << L"Status of " << p->Num << " = " << p->Status << endl;
        }

        catch (bad_weak_ptr b)
        {
            wcout << L"Null object" << endl;
        }
    });
}

};

void RunTest()
{
    vector<shared_ptr<Controller>> v;

    v.push_back(shared_ptr<Controller>(new Controller(0)));
    v.push_back(shared_ptr<Controller>(new Controller(1)));
    v.push_back(shared_ptr<Controller>(new Controller(2)));
    v.push_back(shared_ptr<Controller>(new Controller(3)));
    v.push_back(shared_ptr<Controller>(new Controller(4)));

    // Each controller depends on all others not being deleted.
    // Give each controller a pointer to all the others.
    for (int i = 0 ; i < v.size(); ++i)
    {
        for_each(v.begin(), v.end(), [v,i] (shared_ptr<Controller> p)
        {
            if(p->Num != i)
            {
                v[i]->others.push_back(weak_ptr<Controller>(p));
                wcout << L"push_back to v[" << i << "]: " << p->Num << endl;
            }
        });
    }

    for_each(v.begin(), v.end(), [](shared_ptr<Controller>& p)
    {
        wcout << L"use_count = " << p.use_count() << endl;
        p->CheckStatuses();
    });
}

int main()
{
    RunTest();
    wcout << L"Press any key" << endl;
    char ch;
    cin.getline(&ch, 1);
}

```

Creating Controller0Creating Controller1Creating Controller2Creating Controller3Creating Controller4push\_back to v[0]: 1push\_back to v[0]: 2push\_back to v[0]: 3push\_back to v[0]: 4push\_back to v[1]: 0push\_back to v[1]: 2push\_back to v[1]: 3push\_back to v[1]: 4push\_back to v[2]: 0push\_back to v[2]: 1push\_back to v[2]: 3push\_back to v[2]: 4push\_back to v[3]: 0push\_back to v[3]: 1push\_back to v[3]: 2push\_back to v[3]: 4push\_back to v[4]: 0push\_back to v[4]: 1push\_back to v[4]: 2push\_back to v[4]: 3use\_count = 1Status of 1 = OnStatus of 2 = OnStatus of 3 = OnStatus of 4 = Onuse\_count = 1Status of 0 = OnStatus of 1 = OnStatus of 3 = OnStatus of 4 = Onuse\_count = 1Status of 0 = OnStatus of 1 = OnStatus of 2 = OnStatus of 4 = Onuse\_count = 1Status of 0 = OnStatus of 1 = OnStatus of 2 = OnStatus of 3 = OnDestroying Controller0Destroying Controller1Destroying Controller2Destroying Controller3Destroying Controller4Press any key

As an experiment, modify the vector `others` to be a `vector<shared_ptr<Controller>>`, and then in the output, notice that no destructors are invoked when `TestRun` returns.

## See Also

[Smart Pointers](#)

# How to: Create and Use CComPtr and CComQIPtr Instances

12/21/2017 • 3 min to read • [Edit Online](#)

In classic Windows programming, libraries are often implemented as COM objects (or more precisely, as COM servers). Many Windows operating system components are implemented as COM servers, and many contributors provide libraries in this form. For information about the basics of COM, see [Component Object Model \(COM\)](#).

When you instantiate a Component Object Model (COM) object, store the interface pointer in a COM smart pointer, which performs the reference counting by using calls to `AddRef` and `Release` in the destructor. If you are using the Active Template Library (ATL) or the Microsoft Foundation Class Library (MFC), then use the `CComPtr` smart pointer. If you are not using ATL or MFC, then use `_com_ptr_t`. Because there is no COM equivalent to `std::unique_ptr`, use these smart pointers for both single-owner and multiple-owner scenarios. Both `CComPtr` and `ComQIPtr` support move operations that have rvalue references.

## Example

The following example shows how to use `CComPtr` to instantiate a COM object and obtain pointers to its interfaces. Notice that the `CComPtr::CoCreateInstance` member function is used to create the COM object, instead of the Win32 function that has the same name.

```

void CComPtrDemo()
{
    HRESULT hr = CoInitialize(NULL);

    // Declare the smart pointer.
    CComPtr pGraph;

    // Use its member function CoCreateInstance to
    // create the COM object and obtain the IGraphBuilder pointer.
    hr = pGraph.CoCreateInstance(CLSID_FilterGraph);
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Use the overloaded -> operator to call the interface methods.
    hr = pGraph->RenderFile(L"C:\\\\Users\\\\Public\\\\Music\\\\Sample Music\\\\Sleep Away.mp3", NULL);
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Declare a second smart pointer and use it to
    // obtain another interface from the object.
    CComPtr pControl;
    hr = pGraph->QueryInterface(IID_PPV_ARGS(&pControl));
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Obtain a third interface.
    CComPtr pEvent;
    hr = pGraph->QueryInterface(IID_PPV_ARGS(&pEvent));
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Use the second interface.
    hr = pControl->Run();
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Use the third interface.
    long evCode = 0;
    pEvent->WaitForCompletion(INFINITE, &evCode);

    CoUninitialize();

    // Let the smart pointers do all reference counting.
}

```

`CComPtr` and its relatives are part of the ATL and are defined in `atlcomcli.h`. `_com_ptr_t` is declared in `comip.h`. The compiler creates specializations of `_com_ptr_t` when it generates wrapper classes for type libraries.

## Example

ATL also provides `CComQIPtr`, which has a simpler syntax for querying a COM object to retrieve an additional interface. However, we recommend `CComPtr` because it does everything that `CComQIPtr` can do and is semantically more consistent with raw COM interface pointers. If you use a `CComPtr` to query for an interface, the new interface pointer is placed in an out parameter. If the call fails, an HRESULT is returned, which is the typical COM pattern. With `CComQIPtr`, the return value is the pointer itself, and if the call fails, the internal HRESULT return value cannot be accessed. The following two lines show how the error handling mechanisms in `CComPtr` and `CComQIPtr` differ.

```

// CComPtr with error handling:
CComPtr< IMFMediaControl> pControl;
hr = pGraph->QueryInterface(IID_PPV_ARGS(&pControl));
if(FAILED(hr)){ /*... handle hr error*/ }

// CComQIPtr with error handling
CComQIPtr< IMFMediaEvent> pEvent = pControl;
if(!pEvent){ /*... handle NULL pointer error*/ }

// Use the second interface.
hr = pControl->Run();
if(FAILED(hr)){ /*... handle hr error*/ }

```

## Example

`CComPtr` provides a specialization for `IDispatch` that enables it to store pointers to COM automation components and invoke the methods on the interface by using late binding. `CComDispatchDriver` is a typedef for `CComQIPtr< IDispatch, &IID_IDispatch>`, which is implicitly convertible to `CComPtr< IDispatch>`. Therefore, when any of these three names appears in code, it is equivalent to `CComPtr< IDispatch>`. The following example shows how to obtain a pointer to the Microsoft Word object model by using a `CComPtr< IDispatch>`.

```

void COMAutomationSmartPointerDemo()
{
    CComPtr< IDispatch> pWord;
    CComQIPtr< IDispatch, &IID_IDispatch> pqi = pWord;
    CComDispatchDriver pDriver = pqi;

    HRESULT hr;
    _variant_t pOutVal;

    CoInitialize(NULL);
    hr = pWord.CoCreateInstance(L"Word.Application", NULL, CLSCTX_LOCAL_SERVER);
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Make Word visible.
    hr = pWord.PutPropertyByName(_bstr_t("Visible"), &_variant_t(1));
    if(FAILED(hr)){ /*... handle hr error*/ }

    // Get the Documents collection and store it in new CComPtr
    hr = pWord.GetPropertyByName(_bstr_t("Documents"), &pOutVal);
    if(FAILED(hr)){ /*... handle hr error*/ }

    CComPtr< IDispatch> pDocuments = pOutVal.pdispVal;

    // Use Documents to open a document
    hr = pDocuments.Invoke1 (_bstr_t("Open"),
    &_variant_t("c:\\users\\public\\documents\\sometext.txt"),&pOutVal);
    if(FAILED(hr)){ /*... handle hr error*/ }

    CoUninitialize();
}

```

## See Also

[Smart Pointers](#)

# Pimpl For Compile-Time Encapsulation (Modern C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The *pimpl idiom* is a modern C++ technique to hide implementation, to minimize coupling, and to separate interfaces. Pimpl is short for "pointer to implementation." You may already be familiar with the concept but know it by other names like Cheshire Cat or Compiler Firewall idiom.

## Why use pimpl?

Here's how the pimpl idiom can improve the software development lifecycle:

- Minimization of compilation dependencies.
- Separation of interface and implementation.
- Portability.

## Pimpl header

```
// my_class.h
class my_class {
    // ... all public and protected stuff goes here ...
private:
    class impl; unique_ptr<impl> pimpl; // opaque type here
};
```

The pimpl idiom avoids rebuild cascades and brittle object layouts. It's well suited for (transitively) popular types.

## Pimpl implementation

Define the `impl` class in the .cpp file.

```
// my_class.cpp
class my_class::impl { // defined privately here
    // ... all private data and functions: all of these
    //      can now change without recompiling callers ...
};

my_class::my_class(): pimpl( new impl )
{
    // ... set impl values ...
}
```

## Best practices

Consider whether to add support for non-throwing swap specialization.

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)



# Containers (Modern C++)

12/21/2017 • 1 min to read • [Edit Online](#)

By default, use `vector` as the preferred sequential container in C++. This is equivalent to `List<T>` in .NET languages.

```
vector<string> apples;
apples.push_back("Granny Smith");
```

Use `map` (not `unordered_map`) as the default associative container. Use `set`, `mymap`, and `multiset` for degenerate & multi cases.

```
map<string, string> apple_color;
// ...
apple_color["Granny Smith"] = "Green";
```

When performance optimization is needed, consider using:

1. The `array` type when embedding is important, for example, as a class member.
2. Unordered associative containers such as [unordered\_map](./standard-library/unordered-map-class.md). These have lower per-element overhead and constant-time lookup, but they can be harder to use correctly and efficiently.
3. Sorted `vector`. For more information, see [Algorithms](#).

Don't use C-style arrays. For older APIs that need direct access to the data, use accessor methods such as

```
f(vec.data(), vec.size());
```

For more information about containers, see [C++ Standard Library Containers](#).

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# Algorithms (Modern C++)

12/21/2017 • 2 min to read • [Edit Online](#)

For modern C++ programming, we recommend that you use the algorithms in the [C++ Standard Library](#). Here are some important examples:

- `for_each`, which is the default traversal algorithm. (Also `transform` for not-in-place semantics.)
- `find_if`, which is the default search algorithm.
- `sort`, `lower_bound`, and the other default sorting and searching algorithms.

To write a comparator, use strict `<` and use *named lambdas* when you can.

```
auto comp = [](const widget& w1, const widget& w2)
    { return w1.weight() < w2.weight(); }

sort( v.begin(), v.end(), comp );

auto i = lower_bound( v.begin(), v.end(), comp );
```

## Loops

When possible, use range-based `for` loops or algorithm calls, or both, instead of hand-written loops. `copy`, `transform`, `count_if`, `remove_if`, and others like them are much better than handwritten loops because their intent is obvious and they make it easier to write bug-free code. Also, many C++ Standard Library algorithms have implementation optimizations that make them more efficient.

Instead of old C++ like this:

```
for ( auto i = strings.begin(); i != strings.end(); ++i ) {
    /* ... */
}

auto i = v.begin();

for ( ; i != v.end(); ++i ) {
    if (*i > x && *i < y) break;
}
```

Use modern C++ like this:

```
for_each( begin(strings), end(strings), [](string& s) {
    // ...
} );

auto i = find_if( begin(v), end(v), [=](int i) { return i > x && i < y; } );
```

### Range-based for loops

The range-based `for` loop is a C++11 language feature, not a C++ Standard Library algorithm. But it deserves mention in this discussion about loops. Range-based `for` loops are an extension of the `for` keyword and provide a convenient and efficient way to write loops that iterate over a range of values. C++ Standard Library containers,

strings, and arrays are ready-made for range-based `for` loops. To enable this new iteration syntax for your user-defined type, add the following support:

- A `begin` method that returns an iterator to the beginning of the structure and an `end` method that returns an iterator to the end of the structure.
- Support in the iterator for these methods: `operator*`, `operator!=`, and `operator++` (prefix version).

These methods can be either members or stand-alone functions.

## Random Numbers

It's no secret that the old CRT `rand()` function has many flaws, which have been discussed at length in the C++ community. In modern C++, you don't have to deal with those shortcomings—nor do you have to invent your own uniformly distributed random number generator—because the tools for quickly and easily creating them are available in the C++ Standard Library, as shown in `<random>`.

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# String and I/O Formatting (Modern C++)

12/21/2017 • 1 min to read • [Edit Online](#)

C++ [iostreams](#) are capable of formatted string I/O. For example, the following code shows how to set cout to format an integer to output in hexadecimal, first saving off the current state and re-setting afterwards, because once state formatting is passed to cout, it stays that way until changed, not just for the one line of code.

```
#include <iostream>
#include <iomanip>

using namespace std;

int main()
{
    ios state(nullptr);

    cout << "The answer in decimal is: " << 42 << endl;

    state.copyfmt(cout); // save current formatting
    cout << "In hex: 0x" // now load up a bunch of formatting modifiers
        << hex
        << uppercase
        << setw(8)
        << setfill('0')
        << 42           // the actual value we wanted to print out
        << endl;
    cout.copyfmt(state); // restore previous formatting
}
```

This can be entirely too cumbersome in many cases. As an alternative, you can use Boost.Format from the Boost C++ libraries, even though it's nonstandard. You can download any Boost library from the [Boost](#) website.

Some advantages of Boost.Format are:

- **Safe:** Type-safe, and throws an exception for errors—for example, the specification of too few or too many items.
- **Extensible:** Works for any type that can be streamed.
- **Convenient:** Standard Posix and similar format strings.

Although Boost.Format is built on C++ [iostreams](#), which are safe and extensible, they aren't performance-optimized. When you require performance optimization, consider C [printf](#) and [sprintf](#), which are fast and easy to use. However, they are not extensible or safe from vulnerabilities. (Safe versions exist, but they incur a slight performance penalty. For more information, see [printf\\_s](#), [\\_printf\\_s\\_l](#), [wprintf\\_s](#), [\\_wprintf\\_s\\_l](#) and [sprintf\\_s](#), [\\_sprintf\\_s\\_l](#), [swprintf\\_s](#), [\\_swprintf\\_s\\_l](#)).

The following code demonstrates some of the Boost formatting features.

```
string s = str( format("%2% %2% %1%\n") % "world" % "hello" );
// s contains "hello hello world"

for( auto i = 0; i < names.size(); ++i )
    cout << format("%1% %2% %|40t|%3%\n") % first[i] % last[i] % tel[i];
// Georges Benjamin Clemenceau          +33 (0) 123 456 789
// Jean de Lattre de Tassigny          +33 (0) 987 654 321
```

## See Also

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

[<iostream>](#)

[<limits>](#)

[<iomanip>](#)

# Errors and Exception Handling (Modern C++)

12/21/2017 • 6 min to read • [Edit Online](#)

In modern C++, in most scenarios, the preferred way to report and handle both logic errors and runtime errors is to use exceptions. This is especially true when the stack might contain several function calls between the function that detects the error and the function that has the context to know how to handle it. Exceptions provide a formal, well-defined way for code that detects errors to pass the information up the call stack.

Program errors are generally divided into two categories: logic errors that are caused by programming mistakes, for example, an "index out of range" error, and runtime errors that are beyond the control of programmer, for example, a "network service unavailable" error. In C-style programming and in COM, error reporting is managed either by returning a value that represents an error code or a status code for a particular function, or by setting a global variable that the caller may optionally retrieve after every function call to see whether errors were reported. For example, COM programming uses the HRESULT return value to communicate errors to the caller, and the Win32 API has the GetLastError function to retrieve the last error that was reported by the call stack. In both of these cases, it's up to the caller to recognize the code and respond to it appropriately. If the caller doesn't explicitly handle the error code, the program might crash without warning, or continue to execute with bad data and produce incorrect results.

Exceptions are preferred in modern C++ for the following reasons:

- An exception forces calling code to recognize an error condition and handle it. Unhandled exceptions stop program execution.
- An exception jumps to the point in the call stack that can handle the error. Intermediate functions can let the exception propagate. They do not have to coordinate with other layers.
- The exception stack-unwinding mechanism destroys all objects in scope according to well-defined rules after an exception is thrown.
- An exception enables a clean separation between the code that detects the error and the code that handles the error.

The following simplified example shows the necessary syntax for throwing and catching exceptions in C++.

```

#include <stdexcept>
#include <limits>
#include <iostream>

using namespace std;
class MyClass
{
public:
    void MyFunc(char c)
    {
        if(c < numeric_limits<char>::max())
            throw invalid_argument("MyFunc argument too large.");
        //...
    }
};

int main()
{
    try
    {
        MyFunc(256); //cause an exception to throw
    }

    catch(invalid_argument& e)
    {
        cerr << e.what() << endl;
        return -1;
    }
    //...
    return 0;
}

```

Exceptions in C++ resemble those in languages such as C# and Java. In the `try` block, if an exception is *thrown* it will be *caught* by the first associated `catch` block whose type matches that of the exception. In other words, execution jumps from the `throw` statement to the `catch` statement. If no usable catch block is found, `std::terminate` is invoked and the program exits. In C++, any type may be thrown; however, we recommend that you throw a type that derives directly or indirectly from `std::exception`. In the previous example, the exception type, `invalid_argument`, is defined in the standard library in the `<stdexcept>` header file. C++ does not provide, and does not require, a `finally` block to make sure that all resources are released if an exception is thrown. The resource acquisition is initialization (RAII) idiom, which uses smart pointers, provides the required functionality for resource cleanup. For more information, see [How to: Design for Exception Safety](#). For information about the C++ stack-unwinding mechanism, see [Exceptions and Stack Unwinding](#).

## Basic guidelines

Robust error handling is challenging in any programming language. Although exceptions provide several features that support good error handling, they can't do all the work for you. To realize the benefits of the exception mechanism, keep exceptions in mind as you design your code.

- Use asserts to check for errors that should never occur. Use exceptions to check for errors that might occur, for example, errors in input validation on parameters of public functions. For more information, see the section titled **Exceptions vs. Assertions**.
- Use exceptions when the code that handles the error might be separated from the code that detects the error by one or more intervening function calls. Consider whether to use error codes instead in performance-critical loops when code that handles the error is tightly-coupled to the code that detects it.
- For every function that might throw or propagate an exception, provide one of the three exception guarantees: the strong guarantee, the basic guarantee, or the nothrow (noexcept) guarantee. For more information, see [How to: Design for Exception Safety](#).

- Throw exceptions by value, catch them by reference. Don't catch what you can't handle.
- Don't use exception specifications, which are deprecated in C++11. For more information, see the section titled **Exception specifications and noexcept**.
- Use standard library exception types when they apply. Derive custom exception types from the [exception Class](#) hierarchy.
- Don't allow exceptions to escape from destructors or memory-deallocation functions.

## Exceptions and performance

The exception mechanism has a very minimal performance cost if no exception is thrown. If an exception is thrown, the cost of the stack traversal and unwinding is roughly comparable to the cost of a function call. Additional data structures are required to track the call stack after a `try` block is entered, and additional instructions are required to unwind the stack if an exception is thrown. However, in most scenarios, the cost in performance and memory footprint is not significant. The adverse effect of exceptions on performance is likely to be significant only on very memory-constrained systems, or in performance-critical loops where an error is likely to occur regularly and the code to handle it is tightly coupled to the code that reports it. In any case, it's impossible to know the actual cost of exceptions without profiling and measuring. Even in those rare cases when the cost is significant, you can weigh it against the increased correctness, easier maintainability, and other advantages that are provided by a well-designed exception policy.

## Exceptions vs. assertions

Exceptions and asserts are two distinct mechanisms for detecting run-time errors in a program. Use asserts to test for conditions during development that should never be true if all your code is correct. There is no point in handling such an error by using an exception because the error indicates that something in the code has to be fixed, and doesn't represent a condition that the program has to recover from at run time. An assert stops execution at the statement so that you can inspect the program state in the debugger; an exception continues execution from the first appropriate catch handler. Use exceptions to check error conditions that might occur at run time even if your code is correct, for example, "file not found" or "out of memory." You might want to recover from these conditions, even if the recovery just outputs a message to a log and ends the program. Always check arguments to public functions by using exceptions. Even if your function is error-free, you might not have complete control over arguments that a user might pass to it.

## C++ exceptions versus Windows SEH exceptions

Both C and C++ programs can use the structured exception handling (SEH) mechanism in the Windows operating system. The concepts in SEH resemble those in C++ exceptions, except that SEH uses the `__try`, `__except`, and `__finally` constructs instead of `try` and `catch`. In Visual C++, C++ exceptions are implemented for SEH. However, when you write C++ code, use the C++ exception syntax.

For more information about SEH, see [Structured Exception Handling \(C/C++\)](#).

## Exception specifications and noexcept

Exception specifications were introduced in C++ as a way to specify the exceptions that a function might throw. However, exception specifications proved problematic in practice, and are deprecated in the C++11 draft standard. We recommend that you do not use exception specifications except for `throw()`, which indicates that the function allows no exceptions to escape. If you must use exception specifications of the type `throw( type )`, be aware that Visual C++ departs from the standard in certain ways. For more information, see [Exception Specifications \(throw\)](#). The `noexcept` specifier is introduced in C++11 as the preferred alternative to `throw()`.

## See Also

[How to: Interface Between Exceptional and Non-Exceptional Code](#)

[Welcome Back to C++](#)

[C++ Language Reference](#)

[C++ Standard Library](#)

# How to: Design for Exception Safety

12/21/2017 • 6 min to read • [Edit Online](#)

One of the advantages of the exception mechanism is that execution, together with data about the exception, jumps directly from the statement that throws the exception to the first catch statement that handles it. The handler may be any number of levels up in the call stack. Functions that are called between the try statement and the throw statement are not required to know anything about the exception that is thrown. However, they have to be designed so that they can go out of scope "unexpectedly" at any point where an exception might propagate up from below, and do so without leaving behind partially created objects, leaked memory, or data structures that are in unusable states.

## Basic Techniques

A robust exception-handling policy requires careful thought and should be part of the design process. In general, most exceptions are detected and thrown at the lower layers of a software module, but typically these layers do not have enough context to handle the error or expose a message to end users. In the middle layers, functions can catch and rethrow an exception when they have to inspect the exception object, or they have additional useful information to provide for the upper layer that ultimately catches the exception. A function should catch and "swallow" an exception only if it is able to completely recover from it. In many cases, the correct behavior in the middle layers is to let an exception propagate up the call stack. Even at the highest layer, it might be appropriate to let an unhandled exception terminate a program if the exception leaves the program in a state in which its correctness cannot be guaranteed.

No matter how a function handles an exception, to help guarantee that it is "exception-safe," it must be designed according to the following basic rules.

### Keep Resource Classes Simple

When you encapsulate manual resource management in classes, use a class that does nothing else to manage each resource; otherwise, you might introduce leaks. Use [smart pointers](#) when possible, as shown in the following example. This example is intentionally artificial and simplistic to highlight the differences when `shared_ptr` is used.

```

// old-style new/delete version
class NDResourceClass {
private:
    int* m_p;
    float* m_q;
public:
    NDResourceClass() : m_p(0), m_q(0) {
        m_p = new int;
        m_q = new float;
    }

    ~NDResourceClass() {
        delete m_p;
        delete m_q;
    }
    // Potential leak! When a constructor emits an exception,
    // the destructor will not be invoked.
};

// shared_ptr version
#include <memory>

using namespace std;

class SPResourceClass {
private:
    shared_ptr<int> m_p;
    shared_ptr<float> m_q;
public:
    SPResourceClass() : m_p(new int), m_q(new float) { }
    // Implicitly defined dtor is OK for these members,
    // shared_ptr will clean up and avoid leaks regardless.
};

// A more powerful case for shared_ptr

class Shape {
    // ...
};

class Circle : public Shape {
    // ...
};

class Triangle : public Shape {
    // ...
};

class SPSHapeResourceClass {
private:
    shared_ptr<Shape> m_p;
    shared_ptr<Shape> m_q;
public:
    SPSHapeResourceClass() : m_p(new Circle), m_q(new Triangle) { }
};

```

## Use the RAII Idiom to Manage Resources

To be exception-safe, a function must ensure that objects that it has allocated by using `malloc` or `new` are destroyed, and all resources such as file handles are closed or released even if an exception is thrown. The *Resource Acquisition Is Initialization* (RAII) idiom ties management of such resources to the lifespan of automatic variables. When a function goes out of scope, either by returning normally or because of an exception, the destructors for all fully-constructed automatic variables are invoked. An RAII wrapper object such as a smart pointer calls the appropriate delete or close function in its destructor. In exception-safe code, it is critically important to pass ownership of each resource immediately to some kind of RAII object. Note that the `vector`,

`string`, `make_shared`, `fstream`, and similar classes handle acquisition of the resource for you. However, `unique_ptr` and traditional `shared_ptr` constructions are special because resource acquisition is performed by the user instead of the object; therefore, they count as *Resource Release Is Destruction* but are questionable as RAII.

## The Three Exception Guarantees

Typically, exception safety is discussed in terms of the three exception guarantees that a function can provide: the *no-fail guarantee*, the *strong guarantee*, and the *basic guarantee*.

### No-fail Guarantee

The no-fail (or, "no-throw") guarantee is the strongest guarantee that a function can provide. It states that the function will not throw an exception or allow one to propagate. However, you cannot reliably provide such a guarantee unless (a) you know that all the functions that this function calls are also no-fail, or (b) you know that any exceptions that are thrown are caught before they reach this function, or (c) you know how to catch and correctly handle all exceptions that might reach this function.

Both the strong guarantee and the basic guarantee rely on the assumption that the destructors are no-fail. All containers and types in the Standard Library guarantee that their destructors do not throw. There is also a converse requirement: The Standard Library requires that user-defined types that are given to it—for example, as template arguments—must have non-throwing destructors.

### Strong Guarantee

The strong guarantee states that if a function goes out of scope because of an exception, it will not leak memory and program state will not be modified. A function that provides a strong guarantee is essentially a transaction that has commit or rollback semantics: either it completely succeeds or it has no effect.

### Basic Guarantee

The basic guarantee is the weakest of the three. However, it might be the best choice when a strong guarantee is too expensive in memory consumption or in performance. The basic guarantee states that if an exception occurs, no memory is leaked and the object is still in a usable state even though the data might have been modified.

## Exception-Safe Classes

A class can help ensure its own exception safety, even when it is consumed by unsafe functions, by preventing itself from being partially constructed or partially destroyed. If a class constructor exits before completion, then the object is never created and its destructor will never be called. Although automatic variables that are initialized prior to the exception will have their destructors invoked, dynamically allocated memory or resources that are not managed by a smart pointer or similar automatic variable will be leaked.

The built-in types are all no-fail, and the Standard Library types support the basic guarantee at a minimum. Follow these guidelines for any user-defined type that must be exception-safe:

- Use smart pointers or other RAII-type wrappers to manage all resources. Avoid resource management functionality in your class destructor, because the destructor will not be invoked if the constructor throws an exception. However, if the class is a dedicated resource manager that controls just one resource, then it's acceptable to use the destructor to manage resources.
- Understand that an exception thrown in a base class constructor cannot be swallowed in a derived class constructor. If you want to translate and re-throw the base class exception in a derived constructor, use a function try block.
- Consider whether to store all class state in a data member that is wrapped in a smart pointer, especially if a class has a concept of "initialization that is permitted to fail." Although C++ allows for uninitialized data members, it does not support uninitialized or partially initialized class instances. A constructor must either succeed or fail; no object is created if the constructor does not run to completion.

- Do not allow any exceptions to escape from a destructor. A basic axiom of C++ is that destructors should never allow an exception to propagate up the call stack. If a destructor must perform a potentially exception-throwing operation, it must do so in a try catch block and swallow the exception. The standard library provides this guarantee on all destructors it defines.

## See Also

[Errors and Exception Handling](#)

[How to: Interface Between Exceptional and Non-Exceptional Code](#)

# How to: Interface Between Exceptional and Non-Exceptional Code

12/21/2017 • 6 min to read • [Edit Online](#)

This article describes how to implement consistent exception-handling in a C++ module, and also how to translate those exceptions to and from error codes at the exception boundaries.

Sometimes a C++ module has to interface with code that doesn't use exceptions (non-exceptional code). Such an interface is known as an *exception boundary*. For example, you may want to call the Win32 function `CreateFile` in your C++ program. `CreateFile` doesn't throw exceptions; instead it sets error codes that can be retrieved by the `GetLastError` function. If your C++ program is non-trivial, then in it you probably prefer to have a consistent exception-based error-handling policy. And you probably don't want to abandon exceptions just because you interface with non-exceptional code, and neither do you want to mix exception-based and non-exception-based error policies in your C++ module.

## Calling Non-Exceptional Functions from C++

When you call a non-exceptional function from C++, the idea is to wrap that function in a C++ function that detects any errors and then possibly throws an exception. When you design such a wrapper function, first decide which type of exception guarantee to provide: no-throw, strong, or basic. Second, design the function so that all resources, for example, file handles, are correctly released if an exception is thrown. Typically, this means that you use smart pointers or similar resource managers to own the resources. For more information about design considerations, see [How to: Design for Exception Safety](#).

### Example

The following example shows C++ functions that use the Win32 `CreateFile` and `ReadFile` functions internally to open and read two files. The `File` class is a resource acquisition is initialization (RAII) wrapper for the file handles. Its constructor detects a "file not found" condition and throws an exception to propagate the error up the call stack of the C++ module (in this example, the `main()` function). If an exception is thrown after a `File` object is fully constructed, the destructor automatically calls `CloseHandle` to release the file handle. (If you prefer, you can use the Active Template Library (ATL) `CHandle` class for this same purpose, or a `unique_ptr` together with a custom deleter.) The functions that call Win32 and CRT APIs detect errors and then throw C++ exceptions using the locally-defined `ThrowLastErrorIf` function, which in turn uses the `Win32Exception` class, derived from the `runtime_error` class. All functions in this example provide a strong exception guarantee; if an exception is thrown at any point in these functions, no resources are leaked and no program state is modified.

```
// compile with: /EHsc
#include <Windows.h>
#include <stdlib.h>
#include <vector>
#include <iostream>
#include <string>
#include <limits>
#include <stdexcept>

using namespace std;

string FormatErrorMessage(DWORD error, const string& msg)
{
    static const int BUFFERLENGTH = 1024;
    vector<char> buf(BUFFERLENGTH);
    FormatMessageA(FORMAT_MESSAGE_FROM_SYSTEM, 0, error, 0, buf.data(),
```



```

ThrowLastErrorIf(result == FALSE, "ReadFile failed: " + filename);

cout << filename << " file size: " << filesize << ", bytesRead: "
    << bytesRead << endl;

return readbuffer;
}

bool IsFileDiff(const string& filename1, const string& filename2)
{
    return ReadFileVector(filename1) != ReadFileVector(filename2);
}

#include <iomanip>

int main ( int argc, char* argv[] )
{
    string filename1("file1.txt");
    string filename2("file2.txt");

    try
    {
        if(argc > 2) {
            filename1 = argv[1];
            filename2 = argv[2];
        }

        cout << "Using file names " << filename1 << " and " << filename2 << endl;

        if (IsFileDiff(filename1, filename2)) {
            cout << "*** Files are different." << endl;
        } else {
            cout << "*** Files match." << endl;
        }
    }
    catch(const Win32Exception& e)
    {
        ios state(nullptr);
        state.copyfmt(cout);
        cout << e.what() << endl;
        cout << "Error code: 0x" << hex << uppercase << setw(8) << setfill('0')
            << e.GetErrorCode() << endl;
        cout.copyfmt(state); // restore previous formatting
    }
}

```

## Calling Exceptional Code from Non-Exceptional Code

C++ functions that are declared as "extern C" can be called by C programs. C++ COM servers can be consumed by code written in any of a number of different languages. When you implement public exception-aware functions in C++ to be called by non-exceptional code, the C++ function must not allow any exceptions to propagate back to the caller. Therefore, the C++ function must specifically catch every exception that it knows how to handle and, if appropriate, convert the exception to an error code that the caller understands. If not all potential exceptions are known, the C++ function should have a `catch(...)` block as the last handler. In such a case, it's best to report a fatal error to the caller, because your program might be in an unknown state.

The following example shows a function that assumes that any exception that might be thrown is either a `Win32Exception` or an exception type derived from `std::exception`. The function catches any exception of these types and propagates the error information as a Win32 error code to the caller.

```

BOOL DiffFiles2(const string& file1, const string& file2)
{
    try
    {
        File f1(file1);
        File f2(file2);
        if (IsTextFileDiff(f1, f2))
        {
            SetLastError(MY_APPLICATION_ERROR_FILE_MISMATCH);
            return FALSE;
        }
        return TRUE;
    }
    catch(Win32Exception& e)
    {
        SetLastError(e.GetErrorCode());
    }

    catch(std::exception& e)
    {
        SetLastError(MY_APPLICATION_GENERAL_ERROR);
    }
    return FALSE;
}

```

When you convert from exceptions to error codes, one potential issue is that error codes often don't contain the richness of information that an exception can store. To address this, you can provide a `catch` block for each specific exception type that might be thrown, and perform logging to record the details of the exception before it is converted to an error code. This approach can create a lot of code repetition if multiple functions all use the same set of `catch` blocks. A good way to avoid code repetition is by refactoring those blocks into one private utility function that implements the `try` and `catch` blocks and accepts a function object that is invoked in the `try` block. In each public function, pass the code to the utility function as a lambda expression.

```

template<typename Func>
bool Win32ExceptionBoundary(Func&& f)
{
    try
    {
        return f();
    }
    catch(Win32Exception& e)
    {
        SetLastError(e.GetErrorCode());
    }
    catch(const std::exception& e)
    {
        SetLastError(MY_APPLICATION_GENERAL_ERROR);
    }
    return false;
}

```

The following example shows how to write the lambda expression that defines the functor. When a functor is defined "inline" by using a lambda expression, it is often easier to read than it would be if it were written as a named function object.

```
bool DiffFiles3(const string& file1, const string& file2)
{
    return Win32ExceptionBoundary([&]() -> bool
    {
        File f1(file1);
        File f2(file2);
        if (IsTextFileDiff(f1, f2))
        {
            SetLastError(MY_APPLICATION_ERROR_FILE_MISMATCH);
            return false;
        }
        return true;
    });
}
```

For more information about lambda expressions, see [Lambda Expressions](#).

## See Also

[Errors and Exception Handling](#)

[How to: Design for Exception Safety](#)

# Portability At ABI Boundaries (Modern C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Use sufficiently portable types and conventions at binary interface boundaries. A “portable type” is a C built-in type or a struct that contains only C built-in types. Class types can only be used when caller and callee agree on layout, calling convention, etc. This is only possible when both are compiled with the same compiler and compiler settings.

## How to flatten a class for C portability

When callers may be compiled with another compiler/language, then “flatten” to an **extern "C"** API with a specific calling convention:

```
// class widget {  
//     widget();  
//     ~widget();  
//     double method( int, gadget& );  
// };  
extern "C" {      // functions using explicit "this"  
    struct widget; // opaque type (forward declaration only)  
    widget* STDCALL widget_create();      // constructor creates new "this"  
    void STDCALL widget_destroy(widget*); // destructor consumes "this"  
    double STDCALL widget_method(widget*, int, gadget*); // method uses "this"  
}
```

## See Also

[Welcome Back to C++](#)  
[C++ Language Reference](#)  
[C++ Standard Library](#)

# Lexical Conventions

12/21/2017 • 1 min to read • [Edit Online](#)

This section introduces the fundamental elements of a C++ program. You use these elements, called "lexical elements" or "tokens" to construct statements, definitions, declarations, and so on, which are used to construct complete programs. The following lexical elements are discussed in this section:

- [Character Sets](#)
- [Tokens](#)
- [Comments](#)
- [Identifiers](#)
- [Keywords](#)
- [Punctuators](#)
- [Numeric, Boolean, and Pointer Literals](#)
- [String and Character Literals](#)
- [User-Defined Literals](#)

This section also has an [Overview of File Translation](#) topic.

## See Also

[C++ Language Reference](#)

# Overview of File Translation

12/21/2017 • 1 min to read • [Edit Online](#)

C++ programs, like C programs, consist of one or more files. Each of these files is translated in the following conceptual order (the actual order follows the "as if" rule: translation must occur as if these steps had been followed):

1. Lexical tokenizing. Character mapping and trigraph processing, line splicing, and tokenizing are performed in this translation phase.
2. Preprocessing. This translation phase brings in ancillary source files referenced by `#include` directives, handles "stringizing" and "charizing" directives, and performs token pasting and macro expansion (see [Preprocessor Directives](#) in the *Preprocessor Reference* for more information). The result of the preprocessing phase is a sequence of tokens that, taken together, define a "translation unit."

Preprocessor directives always begin with the number-sign (#) character (that is, the first nonwhite-space character on the line must be a number sign). Only one preprocessor directive can appear on a given line. For example:

```
#include <iostream> // Include text of iostream in
                  // translation unit.
#define NDEBUG      // Define NDEBUG (NDEBUG contains empty
                  // text string).
```

3. Code generation. This translation phase uses the tokens generated in the preprocessing phase to generate object code.

During this phase, syntactic and semantic checking of the source code is performed.

See [Phases of Translation](#) in the *Preprocessor Reference* for more information.

The C++ preprocessor is a strict superset of the ANSI C preprocessor, but the C++ preprocessor differs in a few instances. The following list describes several differences between the ANSI C and the C++ preprocessors:

- Single-line comments are supported. See [Comments](#) for more information.
- One predefined macro, `_cplusplus`, is defined only for C++. See [Predefined Macros](#) in the *Preprocessor Reference* for more information.
- The C preprocessor does not recognize the C++ operators: `.>`, `->`, and `::`. See [Operators](#) and [Expressions](#), for more information about operators.

## See Also

[Lexical Conventions](#)

# Character Sets

12/21/2017 • 3 min to read • [Edit Online](#)

The text of a C++ program is stored in source files that use a particular character encoding. The C++ standard specifies a basic source character set for source files and a basic execution character set for compiled files. Visual C++ allows an additional set of locale-specific characters to be used in source files and compiled files.

## Character sets

The C++ standard specifies a *basic source character set* that may be used in source files. To represent characters outside of this set, additional characters can be specified by using a *universal character name*. When compiled, the *basic execution character set* and *basic execution wide-character set* represent the characters and strings that can appear in a program. The Visual C++ implementation allows additional characters in source code and compiled code.

### Basic source character set

The *basic source character set* consists of 96 characters that may be used in source files. This set includes the space character, horizontal tab, vertical tab, form feed and new-line control characters, and this set of graphical characters:

a b c d e f g h i j k l m n o p q r s t u v w x y z

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

0 1 2 3 4 5 6 7 8 9

\_ { } [ ] # ( ) < > % : ; . ? \* + - / ^ & | ~ ! = , \ " '

### Microsoft Specific

Visual C++ includes the `$` character as a member of the basic source character set. Visual C++ also allows an additional set of characters to be used in source files, based on the file encoding. By default, Visual Studio stores source files by using the default codepage. When source files are saved by using a locale-specific codepage or a Unicode codepage, Visual C++ allows you to use any of the characters of that code page in your source code, except for the control codes not explicitly allowed in the basic source character set. For example, you can put Japanese characters in comments, identifiers, or string literals if you save the file using a Japanese codepage. Visual C++ does not allow character sequences that cannot be translated into valid multibyte characters or Unicode code points. Depending on compiler options, not all allowed characters may appear in identifiers. For more information, see [Identifiers](#).

### END Microsoft Specific

### Universal character names

Because C++ programs can use many more characters than the ones specified in the basic source character set, you can specify these characters in a portable way by using *universal character names*. A universal character name consists of a sequence of characters that represent a Unicode code point. These take two forms. Use `\UNNNNNNNN` to represent a Unicode code point of the form U+NNNNNNNN, where NNNNNNNN is the eight-digit hexadecimal code point number. Use four-digit `\uNNNN` to represent a Unicode code point of the form U+0000NNNN.

Universal character names can be used in identifiers and in string and character literals. A universal character name cannot be used to represent a surrogate code point in the range 0xD800-0xDFFF. Instead, use the desired code point; the compiler automatically generates any required surrogates. Additional restrictions apply to the

universal character names that can be used in identifiers. For more information, see [Identifiers](#) and [String and Character Literals](#).

## **Microsoft Specific**

The Visual C++ compiler treats a character in universal character name form and literal form interchangeably. For example, you can declare an identifier using universal character name form, and use it in literal form:

```
auto \u30AD = 42; // \u30AD is 'ヰ'  
if (\u30AD == 42) return true; // \u30AD and ウ are the same to the compiler
```

The format of extended characters on the Windows clipboard is specific to application locale settings. Cutting and pasting these characters into your code from another application may introduce unexpected character encodings. This can result in parsing errors with no visible cause in your code. We recommend that you set your source file encoding to a Unicode codepage before pasting extended characters. We also recommend that you use an IME or the Character Map app to generate extended characters.

## **END Microsoft Specific**

### **Basic execution character set**

The *basic execution character set* and the *basic execution wide-character set* consist of all the characters in the basic source character set, and the control characters that represent alert, backspace, carriage return, and the null character. The *execution character set* and *execution wide-character set* are supersets of the basic sets. They include the implementation-defined source characters outside the basic source character set. The execution character set has a locale-specific representation.

# Tokens (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

A token is the smallest element of a C++ program that is meaningful to the compiler. The C++ parser recognizes these kinds of tokens: identifiers, keywords, literals, operators, punctuators, and other separators. A stream of these tokens makes up a translation unit.

Tokens are usually separated by *white space*. White space can be one or more:

- Blanks
- Horizontal or vertical tabs
- New lines
- Formfeeds
- Comments

The parser recognizes keywords, identifiers, literals, operators, and punctuators. For information on specific token types, see [Keywords](#), [Identifiers](#), [Numeric, Boolean and Pointer Literals](#), [String and Character Literals](#), [User-Defined Literals](#), [C++ Built-in Operators](#), [Precedence and Associativity](#), and [Punctuators](#). White space is ignored, except as required to separate tokens.

Preprocessing tokens are used in the preprocessing phases to generate the token stream passed to the compiler. The preprocessing token categories are header names, identifiers, preprocessing numbers, character literals, string literals, preprocessing operators and punctuators, and single non-white-space characters that do not match one of the other categories. Character and string literals can be user-defined literals. Preprocessing tokens can be separated by white space or comments.

The parser separates tokens from the input stream by creating the longest token possible using the input characters in a left-to-right scan. Consider this code fragment:

```
a = i+++j;
```

The programmer who wrote the code might have intended either of these two statements:

```
a = i + (++j)  
a = (i++) + j
```

Because the parser creates the longest token possible from the input stream, it chooses the second interpretation, making the tokens `i++`, `+`, and `j`.

## See Also

[Lexical Conventions](#)

# Comments (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

A comment is text that the compiler ignores but that is useful for programmers. Comments are normally used to annotate code for future reference. The compiler treats them as white space. You can use comments in testing to make certain lines of code inactive; however, `#if` / `#endif` preprocessor directives work better for this because you can surround code that contains comments but you cannot nest comments.

A C++ comment is written in one of the following ways:

- The `/*` (slash, asterisk) characters, followed by any sequence of characters (including new lines), followed by the `*/` characters. This syntax is the same as ANSI C.
- The `//` (two slashes) characters, followed by any sequence of characters. A new line not immediately preceded by a backslash terminates this form of comment. Therefore, it is commonly called a "single-line comment."

The comment characters (`/*`, `*/`, and `//`) have no special meaning within a character constant, string literal, or comment. Comments using the first syntax, therefore, cannot be nested.

## See Also

[Lexical Conventions](#)

# Identifiers (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

An identifier is a sequence of characters used to denote one of the following:

- Object or variable name
- Class, structure, or union name
- Enumerated type name
- Member of a class, structure, union, or enumeration
- Function or class-member function
- typedef name
- Label name
- Macro name
- Macro parameter

The following characters are allowed as any character of an identifier:

```
_ a b c d e f g h i j k l m  
n o p q r s t u v w x y z  
A B C D E F G H I J K L M  
N O P Q R S T U V W X Y Z
```

Certain ranges of universal character names are also allowed in an identifier. A universal character name in an identifier cannot designate a control character or a character in the basic source character set. For more information, see [Character Sets](#). These Unicode code point number ranges are allowed as universal character names for any character in an identifier:

- 00A8, 00AA, 00AD, 00AF, 00B2-00B5, 00B7-00BA, 00BC-00BE, 00C0-00D6, 00D8-00F6, 00F8-00FF, 0100-02FF, 0370-167F, 1681-180D, 180F-1DBF, 1E00-1FFF, 200B-200D, 202A-202E, 203F-2040, 2054, 2060-206F, 2070-20CF, 2100-218F, 2460-24FF, 2776-2793, 2C00-2DFF, 2E80-2FFF, 3004-3007, 3021-302F, 3031-303F, 3040-D7FF, F900-FD3D, FD40-FDCF, FDF0-FE1F, FE30-FE44, FE47-FFFD, 10000-1FFFD, 20000-2FFFD, 30000-3FFFD, 40000-4FFFD, 50000-5FFFD, 60000-6FFFD, 70000-7FFFD, 80000-8FFFD, 90000-9FFFD, A0000-AFFFD, B0000-BFFFD, C0000-CFFFD, D0000-DFFFD, E0000-EFFFD

The following characters are allowed as any character in an identifier except the first:

```
0 1 2 3 4 5 6 7 8 9
```

These Unicode code point number ranges are also allowed as universal character names for any character in an identifier except the first:

- 0300-036F, 1DC0-1DFF, 20D0-20FF, FE20-FE2F

## Microsoft Specific

Only the first 2048 characters of Microsoft C++ identifiers are significant. Names for user-defined types are "decorated" by the compiler to preserve type information. The resultant name, including the type

information, cannot be longer than 2048 characters. (See [Decorated Names](#) for more information.) Factors that can influence the length of a decorated identifier are:

- Whether the identifier denotes an object of user-defined type or a type derived from a user-defined type.
- Whether the identifier denotes a function or a type derived from a function.
- The number of arguments to a function.

The dollar sign `$` is a valid identifier character in Visual C++. Visual C++ also allows you to use the actual characters represented by the allowed ranges of universal character names in identifiers. To use these characters, you must save the file by using a file encoding codepage that includes them. This example shows how both extended characters and universal character names can be used interchangeably in your code.

```
// extended_identifier.cpp
// In Visual Studio, use File, Advanced Save Options to set
// the file encoding to Unicode codepage 1200
struct テスト          // Japanese 'test'
{
    void トスト() {}   // Japanese 'toast'
};

int main()
{
    テスト \u30D1\u30F3; // Japanese パン 'bread' in UCN form
    パン.トスト();      // compiler recognizes UCN or literal form
}
```

The range of characters allowed in an identifier is less restrictive when compiling C++/CLI code. Identifiers in code compiled by using /clr should follow [Standard ECMA-335: Common Language Infrastructure \(CLI\)](#).

#### **END Microsoft Specific**

The first character of an identifier must be an alphabetic character, either uppercase or lowercase, or an underscore (`_`). Because C++ identifiers are case sensitive, `fileName` is different from `FileName`.

Identifiers cannot be exactly the same spelling and case as keywords. Identifiers that contain keywords are legal. For example, `Pint` is a legal identifier, even though it contains `int`, which is a keyword.

Use of two sequential underscore characters (`__`) at the beginning of an identifier, or a single leading underscore followed by a capital letter, is reserved for C++ implementations in all scopes. You should avoid using one leading underscore followed by a lowercase letter for names with file scope because of possible conflicts with current or future reserved identifiers.

## See Also

[Lexical Conventions](#)

# Keywords (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Keywords are predefined reserved identifiers that have special meanings. They cannot be used as identifiers in your program. The following keywords are reserved for Microsoft C++. Names with leading underscores are Microsoft extensions.

<code>_abstract</code> <sup>2</sup>	<code>_alignof Operator</code>	<code>_asm</code>
<code>_assume</code>	<code>_based</code>	<code>_box</code> <sup>2</sup>
<code>_cdecl</code>	<code>_declspec</code>	<code>_delegate</code> <sup>2</sup>
<code>_event</code>	<code>_except</code>	<code>_fastcall</code>
<code>_finally</code>	<code>_forceinline</code>	<code>_gc</code> <sup>2</sup>
<code>_hook</code> <sup>3</sup>	<code>_identifier</code>	<code>_if_exists</code>
<code>_if_not_exists</code>	<code>_inline</code>	<code>_int16</code>
<code>_int32</code>	<code>_int64</code>	<code>_int8</code>
<code>_interface</code>	<code>_leave</code>	<code>_m128</code>
<code>_m128d</code>	<code>_m128i</code>	<code>_m64</code>
<code>_multiple_inheritance</code>	<code>_nogc</code> <sup>2</sup>	<code>_noop</code>
<code>_pin</code> <sup>2</sup>	<code>_property</code> <sup>2</sup>	<code>_raise</code>
<code>_sealed</code> <sup>2</sup>	<code>_single_inheritance</code>	<code>_stdcall</code>
<code>_super</code>	<code>_thiscall</code>	<code>_try/_except, _try/_finally</code>
<code>_try_cast</code> <sup>2</sup>	<code>_unaligned</code>	<code>_unhook</code> <sup>3</sup>
<code>_uuidof</code>	<code>_value</code> <sup>2</sup>	<code>_virtual_inheritance</code>
<code>_w64</code>	<code>_wchar_t, wchar_t</code>	<code>abstract</code>
<code>array</code>	<code>auto</code>	<code>bool</code>
<code>break</code>	<code>case</code>	<code>catch</code>
<code>char</code>	<code>class</code>	<code>const</code>

const_cast	continue	decltype
default	delegate	delete
deprecated <sup>1</sup>	dllexport <sup>1</sup>	dllimport <sup>1</sup>
do	double	dynamic_cast
else	enum	enum class
enum struct	event	explicit
extern	false	finally
float	for	for each, in
friend	friend_as	gcnew
generic	goto	if
initonly	inline	int
interface class	interface struct	interior_ptr
literal	long	mutable
naked <sup>1</sup>	namespace	new
new	noinline <sup>1</sup>	noreturn <sup>1</sup>
nothrow <sup>1</sup>	novtable <sup>1</sup>	nullptr
operator	private	property
property <sup>1</sup>	protected	public
ref class	ref struct	register
reinterpret_cast	return	safecast
sealed	selectany <sup>1</sup>	short
signed	sizeof	static
static_assert	static_cast	struct
switch	template	this
thread <sup>1</sup>	throw	true
try	typedef	typeid

<code>typeid</code>	<code>typename</code>	<code>union</code>
<code>unsigned</code>	<code>using</code> declaration, <code>using</code> directive	<code>uuid</code> <sup>1</sup>
<code>value class</code>	<code>value struct</code>	<code>virtual</code>
<code>void</code>	<code>volatile</code>	<code>while</code>

1 Extended attributes for the `__declspec` keyword.

2 Applicable to Managed Extensions for C++ only. This syntax is now deprecated.

3 Intrinsic function used in event handling.

See [Component Extensions for Runtime Platforms](#) for more information.

## Microsoft Specific

In Microsoft C++, identifiers with two leading underscores are reserved for compiler implementations. Therefore, the Microsoft convention is to precede Microsoft-specific keywords with double underscores. These words cannot be used as identifier names.

Microsoft extensions are enabled by default. To ensure that your programs are fully portable, you can disable Microsoft extensions by specifying the ANSI-compatible `/Za` command-line option (compile for ANSI compatibility) during compilation. When you do this, Microsoft-specific keywords are disabled.

When Microsoft extensions are enabled, you can use the Microsoft-specific keywords in your programs. For ANSI compliance, these keywords are prefaced by a double underscore. For backward compatibility, single-underscore versions of all the double-underscored keywords except `__except`, `__finally`, `__leave`, and `__try` are supported. In addition, `__cdecl` is available with no leading underscore.

The `__asm` keyword replaces C++ `asm` syntax. `asm` is reserved for compatibility with other C++ implementations, but not implemented. Use `__asm`.

The `__based` keyword has limited uses for 32-bit and 64-bit target compilations.

### END Microsoft Specific

## See Also

[Lexical Conventions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

# auto Keyword

12/21/2017 • 1 min to read • [Edit Online](#)

The `auto` keyword is a declaration specifier. However, the C++ standard defines an original and a revised meaning for this keyword. Before Visual C++ 2010, the `auto` keyword declares a variable in the *automatic* storage class; that is, a variable that has a local lifetime. Starting with Visual C++ 2010, the `auto` keyword declares a variable whose type is deduced from the initialization expression in its declaration. The `/Zc:auto[-]` compiler option controls the meaning of the `auto` keyword.

## Syntax

```
auto declarator ;  
auto declarator initializer;
```

## Remarks

The definition of the `auto` keyword changes in the C++ programming language, but not in the C programming language.

The following topics describe the `auto` keyword and the corresponding compiler option:

- [auto](#) describes the new definition of the `auto` keyword.
- [/Zc:auto \(Deduce Variable Type\)](#) describes the compiler option that tells the compiler which definition of the `auto` keyword to use.

## See Also

[Keywords](#)

# Punctuators (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Punctuators in C++ have syntactic and semantic meaning to the compiler but do not, of themselves, specify an operation that yields a value. Some punctuators, either alone or in combination, can also be C++ operators or be significant to the preprocessor.

Any of the following characters are considered punctuators:

```
! % ^ & * ( ) - + = { } | ~  
[ ] \ ; ' : " < > ? , . / #
```

The punctuators `[]`, `( )`, and `{ }` must appear in pairs after [translation phase 4](#).

## See Also

[Lexical Conventions](#)

# Numeric, Boolean and Pointer Literals (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

A literal is a program element that directly represents a value. This article covers literals of type integer, floating-point, boolean and pointer. For information about string and character literals, see [String and Character Literals \(C++\)](#). You can also define your own literals based on any of these categories; for more information see [User-Defined Literals \(C++\)](#)

You can use literals in many contexts, but most commonly to initialize named variables and to pass arguments to functions:

```
const int answer = 42; // integer literal
double d = sin(108.87); //floating point literal passed to sin function
bool b = true; // boolean literal
MyClass* mc = nullptr; // pointer literal
```

Sometimes it's important to tell the compiler how to interpret a literal, or what specific type to give to it. You do this by appending prefixes or suffixes to the literal. For example, the prefix `0x` tells the compiler to interpret the number that follows it as a hexadecimal value, for example `0x35`. The `ULL` suffix tells the compiler to treat the value as an `unsigned long long` type, as in `5894345ULL`. See the following sections for the complete list of prefixes and suffixes for each literal type.

## Syntax

### Integer literals

Integer literals begin with a digit and have no fractional parts or exponents. You can specify integer literals in decimal, octal, or hexadecimal form. They can specify signed or unsigned types and long or short types.

When no prefix or suffix is present, the compiler will give an integral literal value type `int` (32 bits), if the value will fit, otherwise it will give it type `long long` (64 bits).

To specify a decimal integral literal, begin the specification with a nonzero digit. For example:

```
int i = 157; // Decimal literal
int j = 0198; // Not a decimal number; erroneous octal literal
int k = 0365; // Leading zero specifies octal literal, not decimal
int m = 36'000'000 // digit separators make large values more readable
int
```

To specify an octal integral literal, begin the specification with `0`, followed by a sequence of digits in the range `0` through `7`. The digits `8` and `9` are errors in specifying an octal literal. For example:

```
int i = 0377; // Octal literal
int j = 0397; // Error: 9 is not an octal digit
```

To specify a hexadecimal integral literal, begin the specification with `0x` or `0X` (the case of the "x" does not matter), followed by a sequence of digits in the range `0` through `9` and `a` (or `A`) through `f` (or `F`). Hexadecimal digits `a` (or `A`) through `f` (or `F`) represent values in the range `10` through `15`. For example:

```
int i = 0x3fff; // Hexadecimal literal
int j = 0X3FFF; // Equal to i
```

To specify an unsigned type, use either the **u** or **U** suffix. To specify a long type, use either the **l** or **L** suffix. To specify a 64-bit integral type, use the **LL**, or **ll** suffix. The **i64** suffix is still supported but should be avoided because it is specific to Microsoft and is not portable. For example:

```
unsigned val_1 = 328u; // Unsigned value
long val_2 = 0x7FFFFFFL; // Long value specified
// as hex literal
unsigned long val_3 = 0776745ul; // Unsigned long value
auto val_4 = 108LL; // signed long long
auto val_4 = 0x800000000000000ULL << 16; // unsigned long long
```

**Digit separators:** You can use the single-quote character (apostrophe) to separate place values in larger numbers to make them easier for humans to read. Separators have no effect on compilation.

```
long long i = 24'847'458'121
```

## Floating point literals

Floating-point literals specify values that must have a fractional part. These values contain decimal points (.) and can contain exponents.

Floating-point literals have a "mantissa," which specifies the value of the number, an "exponent," which specifies the magnitude of the number, and an optional suffix that specifies the literal's type. The mantissa is specified as a sequence of digits followed by a period, followed by an optional sequence of digits representing the fractional part of the number. For example:

```
18.46
38.
```

The exponent, if present, specifies the magnitude of the number as a power of 10, as shown in the following example:

```
18.46e0 // 18.46
18.46e1 // 184.6
```

The exponent may be specified using **e** or **E**, which have the same meaning, followed by an optional sign (+ or -) and a sequence of digits. If an exponent is present, the trailing decimal point is unnecessary in whole numbers such as `18E0`.

Floating-point literals default to type **double**. By using the suffixes **f** or **l** (or **F** or **L** — the suffix is not case sensitive), the literal can be specified as **float** or `long double`, respectively.

Although `long double` and **double** have the same representation, they are not the same type. For example, you can have overloaded functions like

```
void func( double );
```

and

```
void func( long double );
```

## Boolean literals

The boolean literals are `true` and `false`.

## Pointer literal (C++11)

C++ introduces the `nullptr` literal to specify a zero-initialized pointer. In portable code, `nullptr` should be used instead of integral-type zero or macros such as NULL.

## Binary literals (C++14)

A binary literal can be specified by the use of the `0B` or `0b` prefix, followed by a sequence of 1's and 0's:

```
auto x = 0B001101 ; // int
auto y = 0b000001 ; // int
```

## Avoid using literals as "magic constants"

You can use literals directly in expressions and statements although it's not always good programming practice:

```
if (num < 100)
    return "Success";
```

In the previous example, it might be better to use a named constant that conveys a clear meaning, for example "MAXIMUM\_ERROR\_THRESHOLD". And if the return value "Success" is seen by end users, then it might be better to use a named string constant that can be stored in a single location in a file from where it can be localized into other languages. Using named constants helps others as well as yourself to understand the intent of the code.

## See Also

[Lexical Conventions](#)

[C++ String Literals](#)

[C++ User-Defined Literals](#)

# String and Character Literals (C++)

1/3/2018 • 17 min to read • [Edit Online](#)

C++ supports various string and character types, and provides ways to express literal values of each of these types. In your source code, you express the content of your character and string literals using a character set. Universal character names and escape characters allow you to express any string using only the basic source character set. A raw string literal enables you to avoid using escape characters, and can be used to express all types of string literals. You can also create std::string literals without having to perform extra construction or conversion steps.

```
#include <string>
using namespace std::string_literals; // enables s-suffix for std::string literals

int main()
{
    // Character literals
    auto c0 = 'A'; // char
    auto c1 = u8'A'; // char
    auto c2 = L'A'; // wchar_t
    auto c3 = u'A'; // char16_t
    auto c4 = U'A'; // char32_t

    // String literals
    auto s0 = "hello"; // const char*
    auto s1 = u8"hello"; // const char*, encoded as UTF-8
    auto s2 = L"hello"; // const wchar_t*
    auto s3 = u"hello"; // const char16_t*, encoded as UTF-16
    auto s4 = U"hello"; // const char32_t*, encoded as UTF-32

    // Raw string literals containing unescaped \ and "
    auto R0 = R"(Hello \ world)"; // const char*
    auto R1 = u8R"(Hello \ world)"; // const char*, encoded as UTF-8
    auto R2 = LR"(Hello \ world)"; // const wchar_t*
    auto R3 = uR"(Hello \ world)"; // const char16_t*, encoded as UTF-16
    auto R4 = UR"(Hello \ world)"; // const char32_t*, encoded as UTF-32

    // Combining string literals with standard s-suffix
    auto S0 = "hello"s; // std::string
    auto S1 = u8"hello"s; // std::string
    auto S2 = L"hello"s; // std::wstring
    auto S3 = u"hello"s; // std::u16string
    auto S4 = U"hello"s; // std::u32string

    // Combining raw string literals with standard s-suffix
    auto S5 = R"(Hello \ world)"s; // std::string from a raw const char*
    auto S6 = u8R"(Hello \ world)"s; // std::string from a raw const char*, encoded as UTF-8
    auto S7 = LR"(Hello \ world)"s; // std::wstring from a raw const wchar_t*
    auto S8 = uR"(Hello \ world)"s; // std::u16string from a raw const char16_t*, encoded as UTF-16
    auto S9 = UR"(Hello \ world)"s; // std::u32string from a raw const char32_t*, encoded as UTF-32
}
```

String literals can have no prefix, or `u8`, `L`, `u`, and `U` prefixes to denote narrow character (single-byte or multi-byte), UTF-8, wide character (UCS-2 or UTF-16), UTF-16 and UTF-32 encodings, respectively. A raw string literal can have `R`, `u8R`, `LR`, `uR` and `UR` prefixes for the raw version equivalents of these encodings. To create temporary or static std::string values, you can use string literals or raw string literals with an `s` suffix. For more information, see the String literals section below. For more information on the basic source character set, universal character names, and using characters from extended codepages in your source code, see [Character](#)

## Character literals

A *character literal* is composed of a constant character. It is represented by the character surrounded by single quotation marks. There are five kinds of character literals:

- Ordinary character literals of type `char`, for example `'a'`
- UTF-8 character literals of type `char`, for example `u8'a'`
- Wide-character literals of type `wchar_t`, for example `L'a'`
- UTF-16 character literals of type `char16_t`, for example `u'a'`
- UTF-32 character literals of type `char32_t`, for example `U'a'`

The character used for a character literal may be any character, except for the reserved characters backslash ('\''), single quotation mark ('), or new line. Reserved characters can be specified by using an escape sequence. Characters may be specified by using universal character names, as long as the type is large enough to hold the character.

### Encoding

Character literals are encoded differently based their prefix.

- A character literal without a prefix is an ordinary character literal. The value of an ordinary character literal containing a single character, escape sequence, or universal character name that can be represented in the execution character set has a value equal to the numerical value of its encoding in the execution character set. An ordinary character literal that contains more than one character, escape sequence, or universal character name is a *multicharacter literal*. A multicharacter literal or an ordinary character literal that can't be represented in the execution character set is conditionally-supported, has type `int`, and its value is implementation-defined.
- A character literal that begins with the L prefix is a wide-character literal. The value of a wide-character literal containing a single character, escape sequence, or universal character name has a value equal to the numerical value of its encoding in the execution wide-character set unless the character literal has no representation in the execution wide-character set, in which case the value is implementation-defined. The value of a wide-character literal containing multiple characters, escape sequences, or universal character names is implementation-defined.
- A character literal that begins with the u8 prefix is a UTF-8 character literal. The value of a UTF-8 character literal containing a single character, escape sequence, or universal character name has a value equal to its ISO 10646 code point value if it can be represented by a single UTF-8 code unit (corresponding to the C0 Controls and Basic Latin Unicode block). If the value can't be represented by a single UTF-8 code unit, the program is ill-formed. A UTF-8 character literal containing more than one character, escape sequence, or universal character name is ill-formed.
- A character literal that begins with the u prefix is a UTF-16 character literal. The value of a UTF-16 character literal containing a single character, escape sequence, or universal character name has a value equal to its ISO 10646 code point value if it can be represented by a single UTF-16 code unit (corresponding to the basic multi-lingual plane). If the value can't be represented by a single UTF-16 code unit, the program is ill-formed. A UTF-16 character literal containing more than one character, escape sequence, or universal character name is ill-formed.
- A character literal that begins with the U prefix is a UTF-32 character literal. The value of a UTF-32 character literal containing a single character, escape sequence, or universal character name has a value equal to its ISO 10646 code point value. A UTF-8 character literal containing more than one character, escape

sequence, or universal character name is ill-formed.

## Escape Sequences

There are three kinds of escape sequences: simple, octal, and hexadecimal. Escape sequences may be any of the following:

VALUE	ESCAPE SEQUENCE	VALUE	ESCAPE SEQUENCE
newline	\n	backslash	\
horizontal tab	\t	question mark	? or \?
vertical tab	\v	single quote	'
backspace	\b	double quote	"
carriage return	\r	the null character	\0
form feed	\f	octal	\ooo
alert (bell)	\a	hexadecimal	\xhhh

The following code shows some examples of escaped characters using ordinary character literals. The same escape sequence syntax is valid for the other character literal types.

```
#include <iostream>
using namespace std;

int main() {
    char newline = '\n';
    char tab = '\t';
    char backspace = '\b';
    char backslash = '\\';
    char nullChar = '\0';

    cout << "Newline character: " << newline << "ending" << endl; // Newline character:
                                                // ending
    cout << "Tab character: " << tab << "ending" << endl; // Tab character : ending
    cout << "Backspace character: " << backspace << "ending" << endl; // Backspace character : ending
    cout << "Backslash character: " << backslash << "ending" << endl; // Backslash character : \ending
    cout << "Null character: " << nullChar << "ending" << endl; // Null character: ending
}
```

## Microsoft Specific

To create a value from an ordinary character literal (those without a prefix), the compiler converts the character or character sequence between single quotes into 8-bit values within a 32-bit integer. Multiple characters in the literal fill corresponding bytes as needed from high-order to low-order. To create a `char` value, the compiler takes the low-order byte. To create a `wchar_t` or `char16_t` value, the compiler takes the low-order word. The compiler warns that the result is truncated if any bits are set above the assigned byte or word.

```
char c0      = 'abcd';    // C4305, C4309, truncates to 'd'
wchar_t w0 = 'abcd';    // C4305, C4309, truncates to '\x6364'
```

An octal escape sequence is a backslash followed by a sequence of up to 3 octal digits. The behavior of an octal escape sequence that appears to contain more than three digits is treated as a 3-digit octal sequence followed by

the subsequent digits as characters; this can give surprising results. For example:

```
char c1 = '\100'; // '@'  
char c2 = '\1000'; // C4305, C4309, truncates to '0'
```

Escape sequences that appear to contain non-octal characters are evaluated as an octal sequence up to the last octal character, followed by the remaining characters. For example:

```
char c3 = '\009'; // '9'  
char c4 = '\089'; // C4305, C4309, truncates to '9'  
char c5 = '\qrs'; // C4129, C4305, C4309, truncates to 's'
```

A hexadecimal escape sequence is a backslash followed by the character `x`, followed by a sequence of hexadecimal digits. An escape sequence that contains no hexadecimal digits causes compiler error C2153: "hex literals must have at least one hex digit". Leading zeroes are ignored. An escape sequence that appears to have hexadecimal and non-hexadecimal characters is evaluated as a hexadecimal escape sequence up to the last hexadecimal character, followed by the non-hexadecimal characters. In an ordinary or u8-prefixed character literal, the highest hexadecimal value is 0xFF. In an L-prefixed or u-prefixed wide character literal, the highest hexadecimal value is 0xFFFF. In a U-prefixed wide character literal, the highest hexadecimal value is 0xFFFFFFFF.

```
char c6 = '\x0050'; // 'P'  
char c7 = '\x0pqr'; // C4305, C4309, truncates to 'r'
```

If a wide character literal prefixed with `L` contains more than one character, the value is taken from the first character. Subsequent characters are ignored, unlike the behavior of the equivalent ordinary character literal.

```
wchar_t w1 = L'\100'; // L'@'  
wchar_t w2 = L'\1000'; // C4066 L'@', 0 ignored  
wchar_t w3 = L'\009'; // C4066 L'\0', 9 ignored  
wchar_t w4 = L'\089'; // C4066 L'\0', 89 ignored  
wchar_t w5 = L'\qrs'; // C4129, C4066 L'q' escape, rs ignored  
wchar_t w6 = L'\x0050'; // L'P'  
wchar_t w7 = L'\x0pqr'; // C4066 L'\0', pqr ignored
```

## END Microsoft Specific

The backslash character (`\`) is a line-continuation character when it is placed at the end of a line. If you want a backslash character to appear as a character literal, you must type two backslashes in a row (`\\\`). For more information about the line continuation character, see [Phases of Translation](#).

## Universal character names

In character literals and native (non-raw) string literals, any character may be represented by a universal character name. Universal character names are formed by a prefix `\U` followed by an eight-digit Unicode code point, or by a prefix `\u` followed by a four digit Unicode code point. All eight or four digits, respectively, must be present to make a well-formed universal character name.

```
char u1 = 'A'; // 'A'  
char u2 = '\101'; // octal, 'A'  
char u3 = '\x41'; // hexadecimal, 'A'  
char u4 = '\u0041'; // \u UCN 'A'  
char u5 = '\U00000041'; // \U UCN 'A'
```

## Surrogate Pairs

Universal character names cannot encode values in the surrogate code point range D800-DFFF. For Unicode surrogate pairs, specify the universal character name by using `\UNNNNNNNN`, where NNNNNNNN is the eight-digit code point for the character. The compiler generates a surrogate pair if required.

In C++03, the language only allowed a subset of characters to be represented by their universal character names, and allowed some universal character names that didn't actually represent any valid Unicode characters. This was fixed in the C++11 standard. In C++11, both character and string literals and identifiers can use universal character names. For more information on universal character names, see [Character Sets](#). For more information about Unicode, see [Unicode](#). For more information about surrogate pairs, see [Surrogate Pairs and Supplementary Characters](#).

## String literals

A string literal represents a sequence of characters that together form a null-terminated string. The characters must be enclosed between double quotation marks. There are the following kinds of string literals:

### Narrow String Literals

A narrow string literal is a non-prefixed, double-quote delimited, null-terminated array of type `const char[n]`, where n is the length of the array in bytes. A narrow string literal may contain any graphic character except the double quotation mark (" ), backslash (\ ), or newline character. A narrow string literal may also contain the escape sequences listed above, and universal character names that fit in a byte.

```
const char *narrow = "abcd";  
  
// represents the string: yes\nno  
const char *escaped = "yes\\no";
```

### UTF-8 encoded strings

A UTF-8 encoded string is a u8-prefixed, double-quote delimited, null-terminated array of type `const char[n]`, where n is the length of the encoded array in bytes. A u8-prefixed string literal may contain any graphic character except the double quotation mark (" ), backslash (\ ), or newline character. A u8-prefixed string literal may also contain the escape sequences listed above, and any universal character name.

```
const char* str1 = u8"Hello World";  
const char* str2 = u8"\U0001F607 is 0:-)";
```

### Wide String Literals

A wide string literal is a null-terminated array of constant `wchar_t` that is prefixed by 'L' and contains any graphic character except the double quotation mark (" ), backslash (\ ), or newline character. A wide string literal may contain the escape sequences listed above and any universal character name.

```
const wchar_t* wide = L"zyxw";  
const wchar_t* newline = L"hello\ngoodbye";
```

### char16\_t and char32\_t (C++11)

C++11 introduces the portable `char16_t` (16-bit Unicode) and `char32_t` (32-bit Unicode) character types:

```
auto s3 = u"hello"; // const char16_t*  
auto s4 = U"hello"; // const char32_t*
```

### Raw String Literals (C++11)

A raw string literal is a null-terminated array—of any character type—that contains any graphic character,

including the double quotation mark ("), backslash (\), or newline character. Raw string literals are often used in regular expressions that use character classes, and in HTML strings and XML strings. For examples, see the following article: [Bjarne Stroustrup's FAQ on C++11](#).

```
// represents the string: An unescaped \ character
const char* raw_narrow = R"(An unescaped \ character)";
const wchar_t* raw_wide = LR"(An unescaped \ character)";
const char* raw_utf8 = u8R"(An unescaped \ character)";
const char16_t* raw_utf16 = uR"(An unescaped \ character)";
const char32_t* raw_utf32 = UR"(An unescaped \ character);
```

A delimiter is a user-defined sequence of up to 16 characters that immediately precedes the opening parenthesis of a raw string literal and immediately follows its closing parenthesis. For example, in `R"abc(Hello"\()abc"` the delimiter sequence is `abc` and the string content is `Hello"\()`. You can use a delimiter to disambiguate raw strings that contain both double quotation marks and parentheses. This causes a compiler error:

```
// meant to represent the string: )"
const char* bad_parens = R"()""; // error C2059
```

But a delimiter resolves it:

```
const char* good_parens = R"xyz()"xyz";
```

You can construct a raw string literal in which there is a newline (not the escaped character) in the source:

```
// represents the string: hello
//goodbye
const wchar_t* newline = LR"(hello
goodbye)";
```

### std::string Literals (C++14)

`std::string` literals are Standard Library implementations of user-defined literals (see below) that are represented as "xyx"s (with a `s` suffix). This kind of string literal produces a temporary object of type `std::string`, `std::wstring`, `std::u32string` or `std::u16string` depending on the prefix that is specified. When no prefix is used, as above, a `std::string` is produced. `L"xyz"s` produces a `std::wstring`. `u"xyz"s` produces a `std::u16string`, and `U"xyz"s` produces a `std::u32string`.

```
//#include <string>
//using namespace std::string_literals;
string str{ "hello"s };
string str2{ u8"Hello World" };
wstring str3{ L"hello"s };
u16string str4{ u"hello"s };
u32string str5{ U"hello"s };
```

The `s` suffix may also be used on raw string literals:

```
u32string str6{ UR"(She said \"hello.\")"s };
```

`std::string` literals are defined in the namespace `std::literals::string_literals` in the `<string>` header file. Because `std::literals::string_literals`, and `std::literals` are both declared as [inline namespaces](#), `std::literals::string_literals` is automatically treated as if it belonged directly in namespace `std`.

## Size of String Literals

For ANSI `char*` strings and other single-byte encodings (not UTF-8), the size (in bytes) of a string literal is the number of characters plus 1 for the terminating null character. For all other string types, the size is not strictly related to the number of characters. UTF-8 uses up to four `char` elements to encode some *code units*, and `char16_t` or `wchar_t` encoded as UTF-16 may use two elements (for a total of four bytes) to encode a single *code unit*. This example shows the size of a wide string literal in bytes:

```
const wchar_t* str = L"Hello!";
const size_t byteSize = (wcslen(str) + 1) * sizeof(wchar_t);
```

Notice that `strlen()` and `wcslen()` do not include the size of the terminating null character, whose size is equal to the element size of the string type: one byte on a `char*` string, two bytes on `wchar_t*` or `char16_t*` strings, and four bytes on `char32_t*` strings.

The maximum length of a string literal is 65535 bytes. This limit applies to both narrow string literals and wide string literals.

## Modifying String Literals

Because string literals (not including `std::string` literals) are constants, trying to modify them—for example, `str[2] = 'A'`—causes a compiler error.

### Microsoft Specific

In Visual C++ you can use a string literal to initialize a pointer to non-`const` `char` or `wchar_t`. This is allowed in C99 code, but is deprecated in C++98 and removed in C++11. An attempt to modify the string causes an access violation, as in this example:

```
wchar_t* str = L"hello";
str[2] = L'a'; // run-time error: access violation
```

You can cause the compiler to emit an error when a string literal is converted to a `non_const` character pointer when you set the [/Zc:strictStrings \(Disable string literal type conversion\)](#) compiler option. We recommend it for standards-compliant portable code. It is also a good practice to use the `auto` keyword to declare string literal-initialized pointers, because it resolves to the correct (`const`) type. For example, this code example catches an attempt to write to a string literal at compile time:

```
auto str = L"hello";
str[2] = L'a'; // C3892: you cannot assign to a variable that is const.
```

In some cases, identical string literals may be pooled to save space in the executable file. In string-literal pooling, the compiler causes all references to a particular string literal to point to the same location in memory, instead of having each reference point to a separate instance of the string literal. To enable string pooling, use the [/GF](#) compiler option.

### End Microsoft Specific

## Concatenating adjacent string literals

Adjacent wide or narrow string literals are concatenated. This declaration:

```
char str[] = "12" "34";
```

is identical to this declaration:

```
char atr[] = "1234";
```

and to this declaration:

```
char atr[] = "12\  
34";
```

Using embedded hexadecimal escape codes to specify string literals can cause unexpected results. The following example seeks to create a string literal that contains the ASCII 5 character, followed by the characters f, i, v, and e:

```
"\x05five"
```

The actual result is a hexadecimal 5F, which is the ASCII code for an underscore, followed by the characters i, v, and e. To get the correct result, you can use one of these:

```
"\005five"      // Use octal literal.  
"\x05" "five"   // Use string splicing.
```

`std::string` literals, because they are `std::string` types, can be concatenated with the `+` operator that is defined for [basic\\_string](#) types. They can also be concatenated in the same way as adjacent string literals. In both cases, the string encoding and the suffix must match:

```
auto x1 = "hello" " " " world"; // OK  
auto x2 = U"hello" " " L"world"; // C2308: disagree on prefix  
auto x3 = u8"hello" " "s u8"world"s; // OK, agree on prefixes and suffixes  
auto x4 = u8"hello" " "s u8"world"z; // C3688, disagree on suffixes
```

### String literals with universal character names

Native (non-raw) string literals may use universal character names to represent any character, as long as the universal character name can be encoded as one or more characters in the string type. For example, a universal character name representing an extended character cannot be encoded in a narrow string using the ANSI code page, but it can be encoded in narrow strings in some multi-byte code pages, or in UTF-8 strings, or in a wide string. In C++11, Unicode support is extended by the `char16_t*` and `char32_t*` string types:

```
// ASCII smiling face  
const char* s1 = ":-)";  
  
// UTF-16 (on Windows) encoded WINKING FACE (U+1F609)  
const wchar_t* s2 = L"\u263a = \u0001F609 is ;-)"  
  
// UTF-8 encoded SMILING FACE WITH HALO (U+1F607)  
const char* s3 = u8"\u263a = \u0001F607 is O:-)"  
  
// UTF-16 encoded SMILING FACE WITH OPEN MOUTH (U+1F603)  
const char16_t* s4 = u"\u263a = \u0001F603 is :-D"  
  
// UTF-32 encoded SMILING FACE WITH SUNGLASSES (U+1F60E)  
const char32_t* s5 = U"\u263a = \u0001F60E is B-)"
```

## See Also

[Character Sets](#)

[Numeric, Boolean and Pointer Literals](#)

## User-Defined Literals

# User-Defined Literals (C++)

12/21/2017 • 7 min to read • [Edit Online](#)

There are five major categories of literals: integer, character, floating-point, string, boolean and pointer. Starting in C++ 11 you can define your own literals based on these categories to provide syntactic shortcuts for common idioms and increase type safety. For example, let's say you have a Distance class. You could define a literal for kilometers and another one for miles, and encourage the user to be explicit about the units of measure by simply writing: `auto d = 42.0_km` or `auto d = 42.0_mi`. There is no performance advantage or disadvantage to user-defined literals; they are primarily for convenience or for compile-time type deduction. The Standard Library has user-defined literals for `std::string`, for `std::complex`, and for units in time and duration operations in the `<chrono>` header:

```
Distance d = 36.0_mi + 42.0_km;           // Custom UDL (see below)
std::string str = "hello"s + "World"s;    // Standard Library <string> UDL
complex<double> num =
    (2.0 + 3.01i) * (5.0 + 4.3i);      // Standard Library <complex> UDL
auto duration = 15ms + 42h;              // Standard Library <chrono> UDLs
```

## User-defined literal operator signatures

You implement a user-defined literal by defining an `operator""` at namespace scope with one of the following forms:

```
ReturnType operator "" _a(unsigned long long int);    // Literal operator for user-defined INTEGRAL literal
ReturnType operator "" _b(long double);                // Literal operator for user-defined FLOATING literal
ReturnType operator "" _c(char);                       // Literal operator for user-defined CHARACTER literal
ReturnType operator "" _d(wchar_t);                    // Literal operator for user-defined CHARACTER literal
ReturnType operator "" _e(char16_t);                   // Literal operator for user-defined CHARACTER literal
ReturnType operator "" _f(char32_t);                   // Literal operator for user-defined CHARACTER literal
ReturnType operator "" _g(const     char*, size_t);   // Literal operator for user-defined STRING literal
ReturnType operator "" _h(const wchar_t*, size_t);    // Literal operator for user-defined STRING literal
ReturnType operator "" _i(const char16_t*, size_t);  // Literal operator for user-defined STRING literal
ReturnType operator "" _g(const char32_t*, size_t);  // Literal operator for user-defined STRING literal
ReturnType operator "" _r(const char*);                // Raw literal operator
template<char...> ReturnType operator "" _t();        // Literal operator template
```

The operator names in the previous example are placeholders for whatever name you provide; however, the leading underscore is required. (Only the Standard Library is allowed to define literals without the underscore.) The return type is where you customize the conversion or other operation that the literal performs. Also, any of these operators can be defined as `constexpr`.

## Cooked literals

In source code any literal whether user-defined or not is essentially a sequence of alphanumeric characters, such as `101`, or `54.7`, or `"hello"` or `true`. The compiler interprets the sequence as an integer, float, `const char*` string, and so on. A user-defined literal that accepts as input whatever type the compiler assigned to the literal value is informally known as a *cooked literal*. All the operators above except `_r` and `_t` are cooked literals. For example, a literal `42.0_km` would bind to an operator named `_km` that had a signature similar to `_b` and the literal `42_km` would bind to an operator with a signature similar to `_a`.

The following example shows how user-defined literals can encourage callers to be explicit about their input. To

construct a `Distance`, the user must explicitly specify kilometers or miles by using the appropriate user-defined literal. Of course you can also achieve the same result in other ways, but user-defined literals are less verbose than the alternatives.

```
struct Distance
{
private:
    explicit Distance(long double val) : kilometers(val)
    {}

    friend Distance operator"" _km(long double val);
    friend Distance operator"" _mi(long double val);
    long double kilometers{ 0 };

public:
    long double get_kilometers() { return kilometers; }
    Distance operator+(Distance& other)
    {
        return Distance(get_kilometers() + other.get_kilometers());
    }
};

Distance operator"" _km(long double val)
{
    return Distance(val);
}

Distance operator"" _mi(long double val)
{
    return Distance(val * 1.6);
}

int main(int argc, char* argv[])
{
    // Must have a decimal point to bind to the operator we defined!
    Distance d{ 402.0_km }; // construct using kilometers
    cout << "Kilometers in d: " << d.get_kilometers() << endl; // 402

    Distance d2{ 402.0_mi }; // construct using miles
    cout << "Kilometers in d2: " << d2.get_kilometers() << endl; // 643.2

    // add distances constructed with different units
    Distance d3 = 36.0_mi + 42.0_km;
    cout << "d3 value = " << d3.get_kilometers() << endl; // 99.6

    // Distance d4(90.0); // error constructor not accessible

    string s;
    getline(cin, s);
    return 0;
}
```

Note that the literal number must use a decimal, otherwise the number would be interpreted as an integer and the type would not be compatible with the operator. Also note that for floating point input, the type must be `long double`, and for integral types it must be `long long`.

## Raw literals

In a raw user-defined literal, the operator that you define accepts the literal as a sequence of char values and it is up to you to interpret that sequence as a number or string or other type. In the list of operators shown earlier in this page, `_r` and `_t` can be used to define raw literals:

```

ReturnType operator "" _r(const char*);           // Raw literal operator
template<char...> ReturnType operator "" _t();      // Literal operator template

```

You can use raw literals to provide a custom interpretation of an input sequence that is different than what the compiler would perform. For example, you could define a literal that converts the sequence [4.75987](#) into a custom Decimal type instead of an IEEE 754 floating point type. Raw literals, like cooked literals, can also be used to perform compile-time validation of input sequences.

## Example

### Limitations of raw literals

The raw literal operator and literal operator template only work for integral and floating-point user-defined literals, as shown by the following example:

```

#include <cstddef>
#include <cstdio>

void operator "" _dump(unsigned long long int lit) { printf("operator \"\" _dump(unsigned long long int) : "
==>%llu<===\n", lit); } // Literal operator for user-defined INTEGRAL literal
void operator "" _dump(long double lit) { printf("operator \"\" _dump(long double) : "
==>%Lf<===\n", lit); } // Literal operator for user-defined FLOATING literal
void operator "" _dump(char lit) { printf("operator \"\" _dump(char) : "
==>%c<===\n", lit); } // Literal operator for user-defined CHARACTER literal
void operator "" _dump(wchar_t lit) { printf("operator \"\" _dump(wchar_t) : "
==>%d<===\n", lit); } // Literal operator for user-defined CHARACTER literal
void operator "" _dump(char16_t lit) { printf("operator \"\" _dump(char16_t) : "
==>%d<===\n", lit); } // Literal operator for user-defined CHARACTER literal
void operator "" _dump(char32_t lit) { printf("operator \"\" _dump(char32_t) : "
==>%d<===\n", lit); } // Literal operator for user-defined CHARACTER literal
void operator "" _dump(const char* lit, size_t) { printf("operator \"\" _dump(const char*, size_t): "
==>%s<===\n", lit); } // Literal operator for user-defined STRING literal
void operator "" _dump(const wchar_t* lit, size_t) { printf("operator \"\" _dump(const wchar_t*, size_t): "
==>%ls<===\n", lit); } // Literal operator for user-defined STRING literal
void operator "" _dump(const char16_t* lit, size_t) { printf("operator \"\" _dump(const char16_t*, "
size_t):\n" ); } // Literal operator for user-defined STRING literal
void operator "" _dump(const char32_t* lit, size_t) { printf("operator \"\" _dump(const char32_t*, "
size_t):\n" ); } // Literal operator for user-defined STRING literal
void operator "" _dump_raw(const char* lit) { printf("operator \"\" _dump_raw(const char*) : "
==>%s<===\n", lit); } // Raw literal operator

template<char...> void operator "" _dump_template(); // Literal operator template

int main(int argc, const char* argv[])
{
    42_dump;
    3.1415926_dump;
    3.14e+25_dump;
    'A'_dump;
    L'B'_dump;
    u'C'_dump;
    U'D'_dump;
    "Hello World"_dump;
    L"Wide String"_dump;
    u8"UTF-8 String"_dump;
    u"UTF-16 String"_dump;
    U"UTF-32 String"_dump;

    42_dump_raw;
    3.1415926_dump_raw;
    3.14e+25_dump_raw;
    // 'A'_dump_raw; // There is no raw literal operator or literal operator template support
    on this type
    //L'B'_dump_raw; // There is no raw literal operator or literal operator template support on
    this type
}

```

```
//u'C'_dump_raw;           // There is no raw literal operator or literal operator template support on
this type
//U'D'_dump_raw;           // There is no raw literal operator or literal operator template support on
this type
// "Hello World"_dump_raw; // There is no raw literal operator or literal operator template support on
this type
// L"Wide String"_dump_raw; // There is no raw literal operator or literal operator template support on
this type
//u8"UTF-8 String"_dump_raw; // There is no raw literal operator or literal operator template support on
this type
// u"UTF-16 String"_dump_raw; // There is no raw literal operator or literal operator template support on
this type
// U"UTF-32 String"_dump_raw; // There is no raw literal operator or literal operator template support on
this type
}
*****
Output:
operator "" _dump(unsigned long long int) : ===>42<===
operator "" _dump(long double)          : ===>3.141593<===
operator "" _dump(long double)          : ===>3139999999999998506827776.000000<===
operator "" _dump(char)                : ===>A<===
operator "" _dump(wchar_t)              : ===>66<===
operator "" _dump(char16_t)             : ===>67<===
operator "" _dump(char32_t)              : ===>68<===
operator "" _dump(const char*, size_t): ===>Hello World<===
operator "" _dump(const wchar_t*, size_t): ===>Wide String<===
operator "" _dump(const char*, size_t): ===>UTF-8 String<===
operator "" _dump(const char16_t*, size_t):
operator "" _dump(const char32_t*, size_t):
operator "" _dump_raw(const char*)      : ===>42<===
operator "" _dump_raw(const char*)      : ===>3.1415926<===
operator "" _dump_raw(const char*)      : ===>3.14e+25<===
*****/
```

# Basic Concepts (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This section explains concepts that are critical to understanding C++. C programmers will be familiar with many of these concepts, but there are some subtle differences that can cause unexpected program results. The following topics are included:

- [Declarations and definitions](#)
- [Scope of a C++ object or function](#)
- [Program definition and linkage rules](#)
- [Startup and termination](#)
- [L-values and r-values](#)
- [Temporary Objects](#)
- [Alignment](#)
- [alignof and alignas](#)
- [Trivial, standard-layout and POD types](#)

## See Also

[C++ Language Reference](#)

# Declarations and Definitions (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Declarations introduce names in a program, for example the names of variables, namespaces, functions and classes. Declarations also specify type information as well as other characteristics of the object that is being declared. A name must be declared before it can be used; in C++ the point at which a name is declared determines whether it is visible to the compiler. You cannot refer to a function or class that is declared at some later point in the compilation unit; you can use *forward declarations* to get around this limitation.

Definitions specify what code or data the name describes. The compiler needs the definition in order to allocate storage space for the thing that is being declared.

## Declarations

A declaration introduces one or more names into a program. Declarations can occur more than once in a program. Therefore, classes, structures, enumerated types, and other user-defined types can be declared for each compilation unit. The constraint on this multiple declaration is that all declarations must be identical. Declarations also serve as definitions, except when the declaration:

1. Is a function prototype (a function declaration with no function body).
2. Contains the `extern` specifier but no initializer (objects and variables) or function body (functions). This signifies that the definition is not necessarily in the current translation unit and gives the name external linkage.
3. Is of a static data member inside a class declaration.

Because static class data members are discrete variables shared by all objects of the class, they must be defined and initialized outside the class declaration. (For more information about classes and class members, see [Classes](#).)

4. Is a class name declaration with no following definition, such as `class T;`.
5. Is a `typedef` statement.

Examples of declarations that are also definitions are:

```
// Declare and define int variables i and j.  
int i;  
int j = 10;  
  
// Declare enumeration suits.  
enum suits { Spades = 1, Clubs, Hearts, Diamonds };  
  
// Declare class CheckBox.  
class CheckBox : public Control  
{  
public:  
    Boolean IsChecked();  
    virtual int ChangeState() = 0;  
};
```

Some declarations that are not definitions are:

```
extern int i;
char *strchr( const char *Str, const char Target );
```

A name is considered to be declared immediately after its declarator but before its (optional) initializer. For more information, see [Point of Declaration](#).

Declarations occur in a *scope*. The scope controls the visibility of the name declared and the duration of the object defined, if any. For more information about how scope rules interact with declarations, see [Scope](#).

An object declaration is also a definition unless it contains the `extern` storage-class specifier described in [Storage classes](#). A function declaration is also a definition unless it is a prototype. A prototype is a function header without a defining function body. The definition of an object causes allocation of storage and appropriate initializations for that object.

## Definitions

A definition is a unique specification of an object or variable, function, class, or enumerator. Because definitions must be unique, a program can contain only one definition for a given program element. There can be a many-to-one correspondence between declarations and definitions. There are two cases in which a program element can be declared and not defined:

1. A function is declared but never referenced with a function call or with an expression that takes the function's address.
2. A class is used only in a way that does not require its definition be known. However, the class must be declared. The following code illustrates such a case:

```
// definitions.cpp
class WindowCounter; // Forward declaration; no definition

class Window
{
    // Definition of WindowCounter not required
    static WindowCounter windowCounter;
};

int main()
{
}
```

## See Also

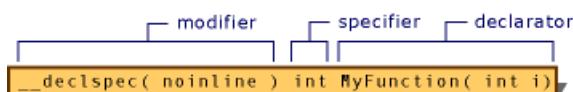
[Basic Concepts](#)

[Point of Declaration](#)

# Overview of Declarators

12/21/2017 • 5 min to read • [Edit Online](#)

Declarators are the components of a declaration that specify names of objects or functions. Declarators also specify whether or not the named object is an object, pointer, reference or array. While declarators do not specify the base type, they do modify the type information in the basic type to specify derived types such as pointers, references, and arrays. Applied to functions, the declarator works with the type specifier to fully specify the return type of a function to be an object, pointer, or reference. (Specifiers, discussed in [Declarations and Definitions](#), convey properties such as type and storage class. Modifiers, discussed in this section and in [Microsoft-Specific Modifiers](#), modify declarators.) The following figure shows a complete declaration of `MyFunction`, and calls out the components of the declaration.



Specifiers, Modifiers, and Declarators

## Microsoft Specific

Most Microsoft extended keywords can be used as modifiers to form derived types; they are not specifiers or declarators. (See [Microsoft-Specific Modifiers](#).)

## END Microsoft Specific

Declarators appear in the declaration syntax after an optional list of specifiers. These specifiers are discussed in [Declarations](#). A declaration can contain more than one declarator, but each declarator declares only one name.

The following sample declaration shows how specifiers and declarators are combined to form a complete declaration:

```
const char *pch, ch;
```

In the preceding declaration, the keywords `const` and `char` make up the list of specifiers. Two declarators are listed: `*pch` and `ch`. A declaration that declares multiple entities consists of a type specifier followed by a comma-separated list of declarators, terminated with a semicolon.

## Declarators for Simple Objects

The declarator of a simple object such as an int or double is simply its name, with optional parentheses.

```
int i; // declarator is i
```

```
int (i); // declarator is (i)
```

## Declarators for Pointers, References and Arrays

Pointer operators inserted in front of the name cause the object to be a pointer or reference. The `\*` operator declares the name as a pointer; the `\&` operator declares it as a reference.

```
int *i; // declarator is *i
int **i; // declarator is **i;
int &i = x; // declarator is &i
```

Appending `const` or `volatile` give the pointer these special properties. The use of these specifiers in a declarator

(as opposed to in the type specifier) modifies the properties of the pointer, not the object pointed to:

```
char *const cpc; // const pointer to char
const char *pcc; // pointer to const char
const char *const cpcc; // const pointer to const char
```

Further information may be found in [const and volatile Pointers](#).

A pointer to a member of a class or struct is declared with the appropriate nested name specifier:

```
int X::* pIntMember;
int ::X::* pIntMember; // the initial :: specifies X is in global scope
char Outer::Inner::* pIntMember; // pointer to char in a nested class
```

Brackets enclosing an optional constant expression after the name cause the object to be an array. Successive brackets declare additional dimensions to the array.

```
int i[5]; // array with five elements of type int numbered from 0 to 4
int i[]; // array of unknown size
char *s[4]; // array of pointers to char
int i[2][2]; // two dimensional array
```

## Declarators for Functions

Parentheses containing the argument list are used after the name to declare a function. The following declares a function of return type `int` and three arguments of type `int`.

```
int f(int a, int b, int c);
```

Pointers and references to functions are declared by prepending the pointer or reference operator to the function name as shown below. Parentheses, normally optional, are required to distinguish a pointer to a function from a function that returns a pointer:

```
int (*pf)(int); // pointer to function returning int
int *f(int i); // function returning pointer to int
int (&pf)(int); // reference to function
```

Pointers to member functions are distinguished by nested name specifiers:

```
int (X::* pmf)(); // pointer to member function of X returning int
int* (X::* pmf)(); // pointer to member function returning pointer to int
```

See also [Pointers to Members](#).

## Functions and objects in the same declaration

Functions and objects may be declared in the same declaration as follows:

```
int i, *j, f(int k); // int, pointer to int, function returning int
```

The syntax may be misleading in some circumstances. The following declaration

```
int* i, f(int k); // pointer to int, function returning int (not int*)
```

may look like the declaration of an `int` pointer and a function returning `int*`, but it is not. That's because the `*` is part of the declarator for `i`, not part of the declarator for `f`.

### Simplifying declarator syntax with `typedef`

A better technique, however, is to use a `typedef` or a combination of parentheses and the `typedef` keyword.

Consider declaring an array of pointers to functions:

```
// Function returning type int that takes one
// argument of type char *.
typedef int (*PIFN)( char * );
// Declare an array of 7 pointers to functions
// returning int and taking one argument of type
// char *.
PIFN pifnDispatchArray[7];
```

The equivalent declaration can be written without the `typedef` declaration, but it is so complicated that the potential for error exceeds any benefits:

```
int ( *pifnDispatchArray[7] )( char * );
```

For more information on `typedef`, see [Aliases and `typedefs`](#).

Pointers, references, arrays of a single base type can be combined in a single declaration (separated by commas) as

```
int a, *b, c[5], **d, &e=a;
```

### More complex declarator syntax

- Pointer, reference, array, and function declarators may be combined to specify such objects as arrays of pointers to functions, pointers to arrays, etc.
- The following recursive grammar describes pointer declarator syntax fully.
- A `declarator` is defined as one of:

1. identifier
2. qualified-name
3. declarator ( argument-list ) [cv-qualifiers] [exception-spec]
4. declarator [ [ constant-expression ] ]
5. pointer-operator declarator
6. ( declarator )

- and `pointer-operator` is one of:

```
* [cv-qualifiers]
& [cv-qualifiers]
::nested-name-specifier * [cv-qualifiers]
```

Because a declarator may contain declarators, it is possible to construct the more complex derived types such as arrays of pointers, functions returning arrays of function pointers, by using the above rules. To form each step of

the construction, start with the identifier representing the base data type and apply the syntax rule above with the previous expression as the `declarator`. The order that you apply the syntax rules should be the reverse of the way the expression is stated in English. If applying the *pointer-operator* syntax rule to an array or function expression, use parentheses if you want a pointer to the array or function, as in the last row in the table below.

The following example shows the construction of "pointer to array of 10 pointers to int".

VERBAL EXPRESSION	DECLARATOR	SYNTAX RULE APPLIED
	i	1
pointer(s) to	*i	5
array of 10	(*i)[10]	4
pointer to	*((*i)[10])	6 and then 5

When multiple pointer, reference, array or function modifiers are used, declarators may become quite complicated. The topic [Interpreting More Complex Declarators](#) describes how to read more complex declarator syntax. The topic is applicable to both C and C++, although in C++, anywhere the \* is used to indicate a pointer, a qualified name such as MyClass::\* may be used to specify a pointer to a member of a class.

# Specifiers

12/21/2017 • 1 min to read • [Edit Online](#)

This topic describes the *decl-specifiers* (declaration specifiers) component of a [declaration](#).

The following placeholders and language keywords are declaration specifiers:

*storage-class-specifier*

*type-specifier*

*function-specifier*

`friend`

`typedef`

`_declspec ( [ extended-decl-modifier-seq ] )`

## Remarks

The *decl-specifiers* part of a declaration is the longest sequence of *decl-specifiers* that can be taken to mean a type name, not including the pointer or reference modifiers. The remainder of the declaration is the *declarator*, which includes the name introduced.

The following table lists four declarations, and then lists each declaration's *decl-specifiers* and *declarator* component separately.

DECLARATION	DECL-SPECIFIERS	DECLARATOR
<code>char *lpszAppName;</code>	<code>char</code>	<code>*lpszAppName</code>
<code>typedef char * LPSTR;</code>	<code>char</code>	<code>*LPSTR</code>
<code>const int func1();</code>	<code>const int</code>	<code>func1</code>
<code>volatile void *pvvObj;</code>	<code>volatile void</code>	<code>*pvvObj</code>

Because `signed`, `unsigned`, `long`, and `short` all imply `int`, a `typedef` name following one of these keywords is taken to be a member of *declarator-list*, not of *decl-specifiers*.

### NOTE

Because a name can be redeclared, its interpretation is subject to the most recent declaration in the current scope.

Redeclaration can affect how names are interpreted by the compiler, especially `typedef` names.

## See Also

[Declarations and Definitions](#)

# Point of declaration in C++

12/21/2017 • 1 min to read • [Edit Online](#)

A name is considered to be declared immediately after its declarator but before its (optional) initializer. (For more information on declarators, see [Declarations and definitions](#).)

Consider this example:

```
// point_of_declaration1.cpp
// compile with: /W1
double dVar = 7.0;
int main()
{
    double dVar = dVar;    // C4700
}
```

If the point of declaration were *after* the initialization, then the local `dVar` would be initialized to 7.0, the value of the global variable `dVar`. However, since that is not the case, `dVar` is initialized to an undefined value.

## See Also

[Scope](#)

# Initializers

12/21/2017 • 12 min to read • [Edit Online](#)

An initializer specifies the initial value of a variable. You can initialize variables in these contexts:

- In the definition of a variable:

```
int i = 3;
Point p1{ 1, 2 };
```

- As one of the parameters of a function:

```
set_point(Point{ 5, 6 });
```

- As the return value of a function:

```
Point get_new_point(int x, int y) { return { x, y }; }
Point get_new_point(int x, int y) { return Point{ x, y }; }
```

Initializers may take these forms:

- An expression (or a comma-separated list of expressions) in parentheses:

```
Point p1(1, 2);
```

- An equals sign followed by an expression:

```
string s = "hello";
```

- A braced initializer list. The list may be empty or may consist of a set of lists, as in the following example:

```
struct Point{
    int x;
    int y;
};
class PointConsumer{
public:
    void set_point(Point p){};
    void set_points(initializer_list<Point> my_list){};
};
int main() {
    PointConsumer pc{};
    pc.set_point({ });
    pc.set_point({ 3, 4 });
    pc.set_points({ { 3, 4 }, { 5, 6 } });
}
```

## Kinds of initialization

There are several kinds of initialization, which may occur at different points in program execution. Different kinds

of initialization are not mutually exclusive—for example, list initialization can trigger value initialization and in other circumstances, it can trigger aggregate initialization.

## Zero initialization

Zero initialization is the setting of a variable to a zero value implicitly converted to the type:

- Numeric variables are initialized to 0 (or 0.0, or 0.0000000000, etc.).
- Char variables are initialized to `'\0'`.
- Pointers are initialized to `nullptr`.
- Arrays, [POD](#) classes, structs, and unions have their members initialized to a zero value.

Zero initialization is performed at different times:

- At program startup, for all named variables that have static duration. These variables may later be initialized again.
- During value initialization, for scalar types and POD class types that are initialized by using empty braces.
- For arrays that have only a subset of their members initialized.

Here are some examples of zero initialization:

```
struct my_struct{
    int i;
    char c;
};

int i0;           // zero-initialized to 0
int main() {
    static float f1; // zero-initialized to 0.00000000
    double d{};     // zero-initialized to 0.0000000000000000
    int* ptr{};    // initialized to nullptr
    char s_array[3]{'a', 'b'}; // the third char is initialized to '\0'
    int int_array[5] = { 8, 9, 10 }; // the fourth and fifth ints are initialized to 0
    my_struct a_struct{}; // i = 0, c = '\0'
}
```

## Default initialization

Default initialization for classes, structs, and unions is initialization with a default constructor. The default constructor can be called with no initialization expression or with the `new` keyword:

```
MyClass mc1;
MyClass* mc3 = new MyClass;
```

If the class, struct, or union does not have a default constructor, the compiler emits an error.

Scalar variables are default initialized when they are defined with no initialization expression. They have indeterminate values.

```
int i1;
float f;
char c;
```

Arrays are default initialized when they are defined with no initialization expression. When an array is default initialized, its members are default initialized and have indeterminate values, as in the following example:

```
int int_arr[3];
```

If the array members do not have a default constructor, the compiler emits an error.

#### Default initialization of constant variables

Constant variables must be declared together with an initializer. If they are scalar types they cause a compiler error, and if they are class types that have a default constructor they cause a warning:

```
class MyClass{};  
int main() {  
    //const int i2;    // compiler error C2734: const object must be initialized if not extern  
    //const char c2;  // same error  
    const MyClass mc1; // compiler error C4269: 'const automatic data initialized with compiler generated  
    default constructor produces unreliable results  
}
```

#### Default initialization of static variables

Static variables that are declared with no initializer are initialized to 0 (implicitly converted to the type).

```
class MyClass {  
private:  
    int m_int;  
    char m_char;  
};  
  
int main() {  
    static int int1;      // 0  
    static char char1;   // '\0'  
    static bool bool1;   // false  
    static MyClass mc1;  // {0, '\0'}  
}
```

For more information about initialization of global static objects, see [Additional Startup Considerations](#).

#### Value initialization

Value initialization occurs in the following cases:

- a named value is initialized using empty brace initialization
- an anonymous temporary object is initialized using empty parentheses or braces
- an object is initialized with the `new` keyword plus empty parentheses or braces

Value initialization does the following:

- for classes with at least one public constructor, the default constructor is called
- for non-union classes with no declared constructors, the object is zero-initialized and the default constructor is called
- for arrays, every element is value-initialized
- in all other cases, the variable is zero initialized

```

class BaseClass {
private:
    int m_int;
};

int main() {
    BaseClass bc{};      // class is initialized
    BaseClass* bc2 = new BaseClass(); // class is initialized, m_int value is 0
    int int_arr[3]{};   // value of all members is 0
    int a{};           // value of a is 0
    double b{};         // value of b is 0.0000000000000000
}

```

## Copy initialization

Copy initialization is the initialization of one object using a different object. It occurs in the following cases:

- a variable is initialized using an equals sign
- an argument is passed to a function
- an object is returned from a function
- an exception is thrown or caught
- a non-static data member is initialized using an equals sign
- class, struct, and union members are initialized by copy initialization during aggregate initialization. See [Aggregate initialization](#) for examples.

The following code shows several examples of copy initialization:

```

#include <iostream>
using namespace std;

class MyClass{
public:
    MyClass(int myInt) {}
    void set_int(int myInt) { m_int = myInt; }
    int get_int() const { return m_int; }
private:
    int m_int = 7; // copy initialization of m_int

};

class MyException : public exception{};

int main() {
    int i = 5;           // copy initialization of i
    MyClass mc1{ i };
    MyClass mc2 = mc1;    // copy initialization of mc2 from mc1
    MyClass mc1.set_int(i); // copy initialization of parameter from i
    int i2 = mc2.get_int(); // copy initialization of i2 from return value of get_int()

    try{
        throw MyException();
    }
    catch (MyException ex){ // copy initialization of ex
        cout << ex.what();
    }
}

```

Copy initialization cannot invoke explicit constructors.

```

vector<int> v = 10; // the constructor is explicit; compiler error C2440: cannot convert from 'int' to
'std::vector<int,std::allocator<_Ty>>'
regex r = "a.*b"; // the constructor is explicit; same error
shared_ptr<int> sp = new int(1729); // the constructor is explicit; same error

```

In some cases, if the copy constructor of the class is deleted or inaccessible, copy initialization causes a compiler error.

### Direct initialization

Direct initialization is initialization using (non-empty) braces or parentheses. Unlike copy initialization, it can invoke explicit constructors. It occurs in the following cases:

- a variable is initialized with non-empty braces or parentheses
- a variable is initialized with the `new` keyword plus non-empty braces or parentheses
- a variable is initialized with `static_cast`
- in a constructor, base classes and non-static members are initialized with an initializer list
- in the copy of a captured variable inside a lambda expression

The following code shows some examples of direct initialization:

```

class BaseClass{
public:
    BaseClass(int n) :m_int(n){} // m_int is direct initialized
private:
    int m_int;
};

class DerivedClass : public BaseClass{
public:
    // BaseClass and m_char are direct initialized
    DerivedClass(int n, char c) : BaseClass(n), m_char(c) {}
private:
    char m_char;
};
int main(){
    BaseClass bc1(5);
    DerivedClass dc1{ 1, 'c' };
    BaseClass* bc2 = new BaseClass(7);
    BaseClass bc3 = static_cast<BaseClass>(dc1);

    int a = 1;
    function<int()> func = [a](){ return a + 1; }; // a is direct initialized
    int n = func();
}

```

### List initialization

List initialization occurs when a variable is initialized using a braced initializer list. Braced initializer lists can be used in the following cases:

- a variable is initialized
- a class is initialized with the `new` keyword
- an object is returned from a function
- an argument passed to a function
- one of the arguments in a direct initialization

- in a non-static data member initializer
- in a constructor initializer list

The following code shows some examples of list initialization:

```

class MyClass {
public:
    MyClass(int myInt, char myChar) {}
private:
    int m_int[]{ 3 };
    char m_char;
};

class MyClassConsumer{
public:
    void set_class(MyClass c) {}
    MyClass get_class() { return MyClass{ 0, '\0' }; }
};

struct MyStruct{
    int my_int;
    char my_char;
    MyClass my_class;
};

int main() {
    MyClass mc1{ 1, 'a' };
    MyClass* mc2 = new MyClass{ 2, 'b' };
    MyClass mc3 = { 3, 'c' };

    MyClassConsumer mcc;
    mcc.set_class(MyClass{ 3, 'c' });
    mcc.set_class({ 4, 'd' });

    MyStruct ms1{ 1, 'a', { 2, 'b' } };
}

```

## Aggregate initialization

Aggregate initialization is a form of list initialization for arrays or class types (often structs or unions) that have:

- no private or protected members
- no user-provided constructors, except for explicitly defaulted or deleted constructors
- no base classes
- no virtual member functions

### NOTE

In Visual Studio 2015 and earlier, an aggregate is not allowed to have brace-or-equal initializers for non-static members. This restriction was removed in the C++14 standard and implemented in Visual Studio 2017.

Aggregate initializers consist of a braced initialization list, with or without an equals sign, as in the following example:

```

#include <iostream>
using namespace std;

struct MyAggregate{
    int myInt;
    char myChar;
};

int main() {
    MyAggregate agg1{ 1, 'c' };

    cout << "agg1: " << agg1.myChar << ":" << agg1.myInt << endl;
    cout << "agg2: " << agg2.myChar << ":" << agg2.myInt << endl;

    int myArr1[]{ 1, 2, 3, 4 };
    int myArr2[3] = { 5, 6, 7 };
    int myArr3[5] = { 8, 9, 10 };

    cout << "myArr1: ";
    for (int i : myArr1){
        cout << i << " ";
    }
    cout << endl;

    cout << "myArr3: ";
    for (auto const &i : myArr3) {
        cout << i << " ";
    }
    cout << endl;
}

```

You should see the following output:

```

agg1: c: 1
agg2: d: 2
myArr1: 1 2 3 4
myArr3: 8 9 10 0 0

```

### IMPORTANT

Array members that are declared but not explicitly initialized during aggregate initialization are zero-initialized, as in `myArr3` above.

### Initializing unions and structs

If a union does not have a constructor, you can initialize it with a single value (or with another instance of a union). The value is used to initialize the first non-static field. This is different from struct initialization, in which the first value in the initializer is used to initialize the first field, the second to initialize the second field, and so on. Compare the initialization of unions and structs in the following example:

```

struct MyStruct {
    int myInt;
    char myChar;
};

union MyUnion {
    int my_int;
    char my_char;
    bool my_bool;
    MyStruct my_struct;
};

int main() {
    MyUnion mu1{ 'a' }; // my_int = 97, my_char = 'a', my_bool = true, {myInt = 97, myChar = '\0'}
    MyUnion mu2{ 1 }; // my_int = 1, my_char = 'x1', my_bool = true, {myInt = 1, myChar = '\0'}
    MyUnion mu3{}; // my_int = 0, my_char = '\0', my_bool = false, {myInt = 0, myChar = '\0'}
    MyUnion mu4 = mu3; // my_int = 0, my_char = '\0', my_bool = false, {myInt = 0, myChar = '\0'}
    //MyUnion mu5{ 1, 'a', true }; // compiler error: C2078: too many initializers
    //MyUnion mu6 = 'a'; // compiler error: C2440: cannot convert from 'char' to 'MyUnion'
    //MyUnion mu7 = 1; // compiler error: C2440: cannot convert from 'int' to 'MyUnion'

    MyStruct ms1{ 'a' }; // myInt = 97, myChar = '\0'
    MyStruct ms2{ 1 }; // myInt = 1, myChar = '\0'
    MyStruct ms3{}; // myInt = 0, myChar = '\0'
    MyStruct ms4{1, 'a'}; // myInt = 1, myChar = 'a'
    MyStruct ms5 = { 2, 'b' }; // myInt = 2, myChar = 'b'
}

```

### Initializing aggregates that contain aggregates

Aggregate types can contain other aggregate types, for example arrays of arrays, arrays of structs, and so on. These types are initialized by using nested sets of braces, for example:

```

struct MyStruct {
    int myInt;
    char myChar;
};

int main() {
    int intArr1[2][2]{{ 1, 2 }, { 3, 4 }};
    int intArr3[2][2] = {1, 2, 3, 4};
    MyStruct structArr[]{{ 1, 'a' }, { 2, 'b' }, { 3, 'c' } };
}

```

### Reference initialization

Variables of reference type must be initialized with an object of the type from which the reference type is derived, or with an object of a type that can be converted to the type from which the reference type is derived. For example:

```

// initializing_references.cppint
int iVar;
long lVar;
int main()
{   long& LongRef1 = lVar; // No conversion required.
    long& LongRef2 = iVar; // C2440
    const long& LongRef3 = iVar; // OK
    LongRef1 = 23L; // Change lVar through a reference.
    LongRef2 = 11L; // Change iVar through a reference.
    LongRef3 = 11L; // C3892}

```

The only way to initialize a reference with a temporary object is to initialize a constant temporary object. Once initialized, a reference-type variable always points to the same object; it cannot be modified to point to another object.

Although the syntax can be the same, initialization of reference-type variables and assignment to reference-type

variables are semantically different. In the preceding example, the assignments that change `iVar` and `lVar` look similar to the initializations, but have different effects. The initialization specifies the object to which the reference-type variable points; the assignment assigns to the referred-to object through the reference.

Because both passing an argument of reference type to a function and returning a value of reference type from a function are initializations, the formal arguments to a function are initialized correctly, as are the references returned.

Reference-type variables can be declared without initializers only in the following:

- Function declarations (prototypes). For example:

```
int func( int& );
```

- Function-return type declarations. For example:

```
int& func( int& );
```

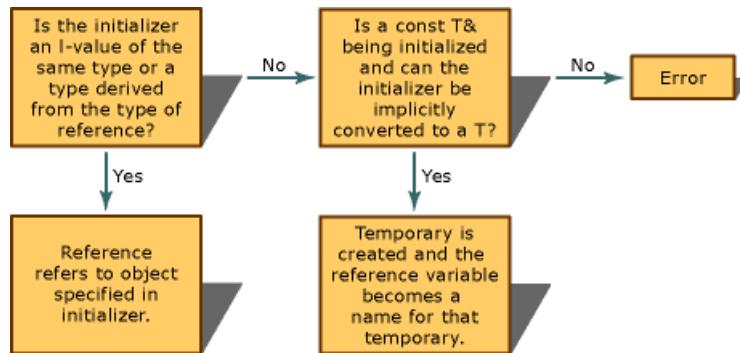
- Declaration of a reference-type class member. For example:

```
class c {public:    int& i;};
```

- Declaration of a variable explicitly specified as `extern`. For example:

```
extern int& iVal;
```

When initializing a reference-type variable, the compiler uses the decision graph shown in the following figure to select between creating a reference to an object or creating a temporary object to which the reference points.



Decision Graph for Initialization of Reference Types

References to `volatile` types (declared as `volatile typename& identifier`) can be initialized with `volatile` objects of the same type or with objects that have not been declared as `volatile`. They cannot, however, be initialized with `const` objects of that type. Similarly, references to `const` types (declared as `const typename& identifier`) can be initialized with `const` objects of the same type (or anything that has a conversion to that type) or with objects that have not been declared as `const`. They cannot, however, be initialized with `volatile` objects of that type.

References that are not qualified with either the `const` or `volatile` keyword can be initialized only with objects declared as neither `const` nor `volatile`.

## Initialization of external variables

Declarations of automatic, static, and external variables can contain initializers. However, declarations of external

variables can contain initializers only if the variables are not declared as `extern`.

# Aliases and typedefs (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

You can use an *alias declaration* to declare a name to use as a synonym for a previously declared type. (This mechanism is also referred to informally as a *type alias*). You can also use this mechanism to create an *alias template*, which can be particularly useful for custom allocators.

## Syntax

```
using identifier = type;
```

## Remarks

`identifier`

The name of the alias.

`type`

The type identifier you are creating an alias for.

An alias does not introduce a new type and cannot change the meaning of an existing type name.

The simplest form of an alias is equivalent to the `typedef` mechanism from C++03:

```
// C++11
using counter = long;

// C++03 equivalent:
// typedef long counter;
```

Both of these enable the creation of variables of type "counter". Something more useful would be a type alias like this one for `std::ios_base::fmtflags`:

```
// C++11
using fmtfl = std::ios_base::fmtflags;

// C++03 equivalent:
// typedef std::ios_base::fmtflags fmtfl;

fmtfl fl_orig = std::cout.flags();
fmtfl fl_hex = (fl_orig & ~std::cout.basefield) | std::cout.showbase | std::cout.hex;
// ...
std::cout.flags(fl_hex);
```

Aliases also work with function pointers, but are much more readable than the equivalent `typedef`:

```
// C++11
using func = void()(int);

// C++03 equivalent:
// typedef void (*func)(int);

// func can be assigned to a function pointer value
void actual_function(int arg) { /* some code */ }
func fptr = &actual_function;
```

A limitation of the `typedef` mechanism is that it doesn't work with templates. However, the type alias syntax in C++11 enables the creation of alias templates:

```
template<typename T> using ptr = T*;

// the name 'ptr<T>' is now an alias for pointer to T
ptr<int> ptr_int;
```

## Example

The following example demonstrates how to use an alias template with a custom allocator—in this case, an integer vector type. You can substitute any type for `int` to create a convenient alias to hide the complex parameter lists in your main functional code. By using the custom allocator throughout your code you can improve readability and reduce the risk of introducing bugs caused by typos.

```

#include <stdlib.h>
#include <new>

template <typename T> struct MyAlloc {
    typedef T value_type;

    MyAlloc() { }
    template <typename U> MyAlloc(const MyAlloc<U>&) { }

    bool operator==(const MyAlloc&) const { return true; }
    bool operator!=(const MyAlloc&) const { return false; }

    T * allocate(const size_t n) const {
        if (n == 0) {
            return nullptr;
        }

        if (n > static_cast<size_t>(-1) / sizeof(T)) {
            throw std::bad_array_new_length();
        }

        void * const pv = malloc(n * sizeof(T));

        if (!pv) {
            throw std::bad_alloc();
        }

        return static_cast<T *>(pv);
    }

    void deallocate(T * const p, size_t) const {
        free(p);
    }
};

};


```

```

#include <vector>
using MyIntVector = std::vector<int, MyAlloc<int>>;

#include <iostream>

int main ()
{
    MyIntVector foov = { 1701, 1764, 1664 };

    for (auto a: foov) std::cout << a << " ";
    std::cout << "\n";

    return 0;
}

```

```
1701 1764 1664
```

## Typedefs

A `typedef` declaration introduces a name that, within its scope, becomes a synonym for the type given by the *type-declaration* portion of the declaration.

You can use `typedef` declarations to construct shorter or more meaningful names for types already defined by the language or for types that you have declared. `Typedef` names allow you to encapsulate implementation details that may change.

In contrast to the `class`, `struct`, `union`, and `enum` declarations, `typedef` declarations do not introduce new types

— they introduce new names for existing types.

Names declared using `typedef` occupy the same namespace as other identifiers (except statement labels). Therefore, they cannot use the same identifier as a previously declared name, except in a class-type declaration. Consider the following example:

```
// typedef_names1.cpp
// C2377 expected
typedef unsigned long UL;    // Declare a typedef name, UL.
int UL;                      // C2377: redefined.
```

The name-hiding rules that pertain to other identifiers also govern the visibility of names declared using `typedef`. Therefore, the following example is legal in C++:

```
// typedef_names2.cpp
typedef unsigned long UL;    // Declare a typedef name, UL
int main()
{
    unsigned int UL;    // Redefinition hides typedef name
}

// typedef UL back in scope
```

```
// typedef_specifier1.cpp
typedef char FlagType;

int main()
{

void myproc( int )
{
    int FlagType;
}
```

When declaring a local-scope identifier by the same name as a `typedef`, or when declaring a member of a structure or union in the same scope or in an inner scope, the type specifier must be specified. For example:

```
typedef char FlagType;
const FlagType x;
```

To reuse the `FlagType` name for an identifier, a structure member, or a union member, the type must be provided:

```
const int FlagType; // Type specifier required
```

It is not sufficient to say

```
const FlagType;      // Incomplete specification
```

because the `FlagType` is taken to be part of the type, not an identifier that is being redeclared. This declaration is taken to be an illegal declaration like

```
int; // Illegal declaration
```

You can declare any type with `typedef`, including pointer, function, and array types. You can declare a `typedef` name for a pointer to a structure or union type before you define the structure or union type, as long as the definition has the same visibility as the declaration.

## Examples

One use of `typedef` declarations is to make declarations more uniform and compact. For example:

```
typedef char CHAR;           // Character type.  
typedef CHAR * PSTR;         // Pointer to a string (char *).  
PSTR strchr( PSTR source, CHAR target );  
typedef unsigned long ulong;  
ulong ul;      // Equivalent to "unsigned long ul;"
```

To use `typedef` to specify fundamental and derived types in the same declaration, you can separate declarators with commas. For example:

```
typedef char CHAR, *PSTR;
```

The following example provides the type `DRAWF` for a function returning no value and taking two int arguments:

```
typedef void DRAWF( int, int );
```

After the above `typedef` statement, the declaration

```
DRAWF box;
```

would be equivalent to the declaration

```
void box( int, int );
```

`typedef` is often combined with `struct` to declare and name user-defined types:

```
// typedef_specifier2.cpp  
#include <stdio.h>  
  
typedef struct mystructtag  
{  
    int i;  
    double f;  
} mystruct;  
  
int main()  
{  
    mystruct ms;  
    ms.i = 10;  
    ms.f = 0.99;  
    printf_s("%d %f\n", ms.i, ms.f);  
}
```

```
10 0.990000
```

## Re-declaration of `typedefs`

The `typedef` declaration can be used to redeclare the same name to refer to the same type. For example:

```

// FILE1.H
typedef char CHAR;

// FILE2.H
typedef char CHAR;

// PROG.CPP
#include "file1.h"
#include "file2.h" // OK

```

The program PROG.CPP includes two header files, both of which contain `typedef` declarations for the name `CHAR`. As long as both declarations refer to the same type, such redeclaration is acceptable.

A `typedef` cannot redefine a name that was previously declared as a different type. Therefore, if FILE2.H contains

```

// FILE2.H
typedef int CHAR; // Error

```

the compiler issues an error because of the attempt to redeclare the name `CHAR` to refer to a different type. This extends to constructs such as:

```

typedef char CHAR;
typedef CHAR CHAR; // OK: redeclared as same type

typedef union REGS // OK: name REGS redeclared
{
    // by typedef name with the
    struct wordregs x; // same meaning.
    struct byteregs h;
} REGS;

```

## typedefs in C++ vs. C

Use of the `typedef` specifier with class types is supported largely because of the ANSI C practice of declaring unnamed structures in `typedef` declarations. For example, many C programmers use the following:

```

// typedef_with_class_types1.cpp
// compile with: /c
typedef struct { // Declare an unnamed structure and give it the
    // typedef name POINT.
    unsigned x;
    unsigned y;
} POINT;

```

The advantage of such a declaration is that it enables declarations like:

```
POINT ptOrigin;
```

instead of:

```
struct point_t ptOrigin;
```

In C++, the difference between `typedef` names and real types (declared with the `class`, `struct`, `union`, and `enum` keywords) is more distinct. Although the C practice of declaring a nameless structure in a `typedef` statement still works, it provides no notational benefits as it does in C.

```
// typedef_with_class_types2.cpp
// compile with: /c /W1
typedef struct {
    int POINT();
    unsigned x;
    unsigned y;
} POINT;
```

The preceding example declares a class named `POINT` using the unnamed class `typedef` syntax. `POINT` is treated as a class name; however, the following restrictions apply to names introduced this way:

- The name (the synonym) cannot appear after a `class`, `struct`, or `union` prefix.
- The name cannot be used as constructor or destructor names within a class declaration.

In summary, this syntax does not provide any mechanism for inheritance, construction, or destruction.

# using Declaration

12/21/2017 • 5 min to read • [Edit Online](#)

The using declaration introduces a name into the declarative region in which the using declaration appears.

## Syntax

```
using [typename] nested-name-specifier unqualified-id ;  
using declarator-list ;
```

### Parameters

#### *nested-name-specifier*

A sequence of namespace, class, or enumeration names and scope resolution operators (::), terminated by a scope resolution operator. A single scope resolution operator may be used to introduce a name from the global namespace. The keyword `typename` is optional and may be used to resolve dependent names when introduced into a class template from a base class.

#### *unqualified-id*

An unqualified id-expression, which may be an identifier, an overloaded operator name, a user-defined literal operator or conversion function name, a class destructor name, or a template name and argument list.

#### *declarator-list*

A comma-separated list of `[ typename ] nested-name-specifier unqualified-id` declarators, followed optionally by an ellipsis.

## Remarks

A using declaration introduces an unqualified name as a synonym for an entity declared elsewhere. It allows a single name from a specific namespace to be used without explicit qualification in the declaration region in which it appears. This is in contrast to the [using directive](#), which allows *all* the names in a namespace to be used without qualification. The `using` keyword is also used for [type aliases](#).

## Example

A using declaration can be used in a class definition.

```

// using_declaration1.cpp
#include <stdio.h>
class B {
public:
    void f(char) {
        printf_s("In B::f()\n");
    }

    void g(char) {
        printf_s("In B::g()\n");
    }
};

class D : B {
public:
    using B::f;      // B::f(char) is now visible as D::f(char)
    using B::g;      // B::g(char) is now visible as D::g(char)
    void f(int) {
        printf_s("In D::f()\n");
        f('c');      // Invokes B::f(char) instead of recursing
    }

    void g(int) {
        printf_s("In D::g()\n");
        g('c');      // Invokes B::g(char) instead of recursing
    }
};

int main() {
    D myD;
    myD.f(1);
    myD.g('a');
}

```

```

In D::f()
In B::f()
In B::g()

```

## Example

When used to declare a member, a using declaration must refer to a member of a base class.

```
// using_declaration2.cpp
#include <stdio.h>

class B {
public:
    void f(char) {
        printf_s("In B::f()\n");
    }

    void g(char) {
        printf_s("In B::g()\n");
    }
};

class C {
public:
    int g();
};

class D2 : public B {
public:
    using B::f; // ok: B is a base of D2
    // using C::g; // error: C isn't a base of D2
};

int main() {
    D2 MyD2;
    MyD2.f('a');
}
```

```
In B::f()
```

## Example

Members declared by using a using declaration can be referenced by using explicit qualification. The `::` prefix refers to the global namespace.

```

// using_declaration3.cpp
#include <stdio.h>

void f() {
    printf_s("In f\n");
}

namespace A {
    void g() {
        printf_s("In A::g\n");
    }
}

namespace X {
    using ::f;    // global f is also visible as X::f
    using A::g;   // A's g is now visible as X::g
}

void h() {
    printf_s("In h\n");
    X::f();      // calls ::f
    X::g();      // calls A::g
}

int main() {
    h();
}

```

```

In h
In f
In A::g

```

## Example

When a using declaration is made, the synonym created by the declaration refers only to definitions that are valid at the point of the using declaration. Definitions added to a namespace after the using declaration are not valid synonyms.

A name defined by a `using` declaration is an alias for its original name. It does not affect the type, linkage or other attributes of the original declaration.

```

// post_declaratoin_namespace_additions.cpp
// compile with: /c
namespace A {
    void f(int) {}
}

using A::f;    // f is a synonym for A::f(int) only

namespace A {
    void f(char) {}
}

void f() {
    f('a');    // refers to A::f(int), even though A::f(char) exists
}

void b() {
    using A::f;    // refers to A::f(int) AND A::f(char)
    f('a');    // calls A::f(char);
}

```

## Example

With respect to functions in namespaces, if a set of local declarations and using declarations for a single name are given in a declarative region, they must all refer to the same entity, or they must all refer to functions.

```
// functions_in_namespaces1.cpp
// C2874 expected
namespace B {
    int i;
    void f(int);
    void f(double);
}

void g() {
    int i;
    using B::i;    // error: i declared twice
    void f(char);
    using B::f;    // ok: each f is a function
}
```

In the example above, the `using B::i` statement causes a second `int i` to be declared in the `g()` function. The `using B::f` statement does not conflict with the `f(char)` function because the function names introduced by `B::f` have different parameter types.

## Example

A local function declaration cannot have the same name and type as a function introduced by using declaration. For example:

```
// functions_in_namespaces2.cpp
// C2668 expected
namespace B {
    void f(int);
    void f(double);
}

namespace C {
    void f(int);
    void f(double);
    void f(char);
}

void h() {
    using B::f;          // introduces B::f(int) and B::f(double)
    using C::f;          // C::f(int), C::f(double), and C::f(char)
    f('h');              // calls C::f(char)
    f(1);                // C2668 ambiguous: B::f(int) or C::f(int)?
    void f(int);         // C2883 conflicts with B::f(int) and C::f(int)
}
```

## Example

With respect to inheritance, when a using declaration introduces a name from a base class into a derived class scope, member functions in the derived class override virtual member functions with the same name and argument types in the base class.

```

// using_declaration_inheritance1.cpp
#include <stdio.h>
struct B {
    virtual void f(int) {
        printf_s("In B::f(int)\n");
    }

    virtual void f(char) {
        printf_s("In B::f(char)\n");
    }

    void g(int) {
        printf_s("In B::g\n");
    }

    void h(int);
};

struct D : B {
    using B::f;
    void f(int) { // ok: D::f(int) overrides B::f(int)
        printf_s("In D::f(int)\n");
    }

    using B::g;
    void g(char) { // ok: there is no B::g(char)
        printf_s("In D::g(char)\n");
    }

    using B::h;
    void h(int) {} // Note: D::h(int) hides non-virtual B::h(int)
};

void f(D* pd) {
    pd->f(1); // calls D::f(int)
    pd->f('a'); // calls B::f(char)
    pd->g(1); // calls B::g(int)
    pd->g('a'); // calls D::g(char)
}

int main() {
    D * myd = new D();
    f(myd);
}

```

```

In D::f(int)
In B::f(char)
In B::g
In D::g(char)

```

## Example

All instances of a name mentioned in a using declaration must be accessible. In particular, if a derived class uses a using declaration to access a member of a base class, the member name must be accessible. If the name is that of an overloaded member function, then all functions named must be accessible.

For more information on accessibility of members, see [Member-Access Control](#).

```
// using_declaration_inheritance2.cpp
// C2876 expected
class A {
private:
    void f(char);
public:
    void f(int);
protected:
    void g();
};

class B : public A {
    using A::f;    // C2876: A::f(char) is inaccessible
public:
    using A::g;    // B::g is a public synonym for A::g
};
```

## See Also

[Namespaces](#)

[Keywords](#)

# Resolving ambiguous declarations (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

To perform explicit conversions from one type to another, you must use casts, specifying the desired type name. Some type casts result in syntactic ambiguity. The following function-style type cast is ambiguous:

```
char *aName( String( s ) );
```

It is unclear whether it is a function declaration or an object declaration with a function-style cast as the initializer: It could declare a function returning type `char \*` that takes one argument of type `String`, or it could declare the object `aName` and initialize it with the value of `s` cast to type `String`.

If a declaration can be considered a valid function declaration, it is treated as such. Only if it cannot possibly be a function declaration — that is, if it would be syntactically incorrect — is a statement examined to see if it is a function-style type cast. Therefore, the compiler considers the statement to be a declaration of a function and ignores the parentheses around the identifier `s`. On the other hand, the statements:

```
char *aName( (String)s );
```

and

```
char *aName = String( s );
```

are clearly declarations of objects, and a user-defined conversion from type `String` to type `char \*` is invoked to perform the initialization of `aName`.

## See Also

# Storage classes (C++)

12/21/2017 • 8 min to read • [Edit Online](#)

A *storage class* in the context of C++ variable declarations is a type specifier that governs the lifetime, linkage, and memory location of objects. A given object can have only one storage class. Variables defined within a block have automatic storage unless otherwise specified using the `extern`, `static`, or `thread_local` specifiers. Automatic objects and variables have no linkage; they are not visible to code outside the block.

## Notes

1. The `mutable` keyword may be considered a storage class specifier. However, it is only available in the member list of a class definition.
2. **Visual C++ 2010 and later:** The `auto` keyword is no longer a C++ storage-class specifier, and the `register` keyword is deprecated. **Visual Studio 2017 version 15.3 and later:** (available with `/std:c++17`): The `register` keyword is no longer a supported storage class. The keyword is still reserved in the standard for future use.

```
register int val; // warning C5033: 'register' is no longer a supported storage class
```

## In this section:

- [static](#)
- [extern](#)
- [thread\\_local](#)

## static

The `static` keyword can be used to declare variables and functions at global scope, namespace scope, and class scope. Static variables can also be declared at local scope.

Static duration means that the object or variable is allocated when the program starts and is deallocated when the program ends. External linkage means that the name of the variable is visible from outside the file in which the variable is declared. Conversely, internal linkage means that the name is not visible outside the file in which the variable is declared. By default, an object or variable that is defined in the global namespace has static duration and external linkage. The `static` keyword can be used in the following situations.

1. When you declare a variable or function at file scope (global and/or namespace scope), the `static` keyword specifies that the variable or function has internal linkage. When you declare a variable, the variable has static duration and the compiler initializes it to 0 unless you specify another value.
2. When you declare a variable in a function, the `static` keyword specifies that the variable retains its state between calls to that function.
3. When you declare a data member in a class declaration, the `static` keyword specifies that one copy of the member is shared by all instances of the class. A static data member must be defined at file scope. An integral data member that you declare as `const static` can have an initializer.
4. When you declare a member function in a class declaration, the `static` keyword specifies that the function is shared by all instances of the class. A static member function cannot access an instance member because the function does not have an implicit `this` pointer. To access an instance member, declare the function

with a parameter that is an instance pointer or reference.

5. You cannot declare the members of a union as static. However, a globally declared anonymous union must be explicitly declared `static`.

This example shows how a variable declared `static` in a function retains its state between calls to that function.

```
// static1.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
void showstat( int curr ) {
    static int nStatic;      // Value of nStatic is retained
                            // between each function call
    nStatic += curr;
    cout << "nStatic is " << nStatic << endl;
}

int main() {
    for ( int i = 0; i < 5; i++ )
        showstat( i );
}
```

```
nStatic is 0
nStatic is 1
nStatic is 3
nStatic is 6
nStatic is 10
```

This example shows the use of `static` in a class.

```
// static2.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class CMyClass {
public:
    static int m_i;
};

int CMyClass::m_i = 0;
CMyClass myObject1;
CMyClass myObject2;

int main() {
    cout << myObject1.m_i << endl;
    cout << myObject2.m_i << endl;

    myObject1.m_i = 1;
    cout << myObject1.m_i << endl;
    cout << myObject2.m_i << endl;

    myObject2.m_i = 2;
    cout << myObject1.m_i << endl;
    cout << myObject2.m_i << endl;

    CMyClass::m_i = 3;
    cout << myObject1.m_i << endl;
    cout << myObject2.m_i << endl;
}
```

```
0  
0  
1  
1  
2  
2  
3  
3
```

This example shows a local variable declared `static` in a member function. The static variable is available to the whole program; all instances of the type share the same copy of the static variable.

```
// static3.cpp  
// compile with: /EHsc  
#include <iostream>  
using namespace std;  
struct C {  
    void Test(int value) {  
        static int var = 0;  
        if (var == value)  
            cout << "var == value" << endl;  
        else  
            cout << "var != value" << endl;  
  
        var = value;  
    }  
};  
  
int main() {  
    C c1;  
    C c2;  
    c1.Test(100);  
    c2.Test(100);  
}
```

```
var != value  
var == value
```

Starting in C++11, a static local variable initialization is guaranteed to be thread-safe. This feature is sometimes called *magic statics*. However, in a multithreaded application all subsequent assignments must be synchronized. The thread-safe static initialization feature can be disabled by using the [/Zc:threadSafeInit-](#) flag to avoid taking a dependency on the CRT.

## extern

Objects and variables declared as `extern` declare an object that is defined in another translation unit or in an enclosing scope as having external linkage.

Declaration of `const` variables with the `extern` storage class forces the variable to have external linkage. An initialization of an `extern const` variable is allowed in the defining translation unit. Initializations in translation units other than the defining translation unit produce undefined results. For more information, see [Using extern to Specify Linkage](#)

The following code shows two `extern` declarations, `DefinedElsewhere` (which refers to a name defined in a different translation unit) and `DefinedHere` (which refers to a name defined in an enclosing scope):

```
// external.cpp
// DefinedElsewhere is defined in another translation unit
extern int DefinedElsewhere;
int main() {
    int DefinedHere;
    {
        // refers to DefinedHere in the enclosing scope
        extern int DefinedHere;
    }
}
```

## thread\_local (C++11)

A variable declared with the `thread_local` specifier is accessible only on the thread on which it is created. The variable is created when the thread is created, and destroyed when the thread is destroyed. Each thread has its own copy of the variable. On Windows, `thread_local` is functionally equivalent to the Microsoft-specific `_declspec(thread)` attribute.

```
thread_local float f = 42.0; // Global namespace. Not implicitly static.

struct S // cannot be applied to type definition
{
    thread_local int i; // Illegal. The member must be static.
    thread_local static char buf[10]; // OK
};

void DoSomething()
{
    // Apply thread_local to a local variable.
    // Implicitly "thread_local static S my_struct".
    thread_local S my_struct;
}
```

Things to note about the `thread_local` specifier:

- Dynamically initialized thread-local variables in DLLs may not be correctly initialized on all calling threads. For more information, see [thread](#).
- The `thread_local` specifier may be combined with `static` or `extern`.
- You can apply `thread_local` only to data declarations and definitions; `thread_local` cannot be used on function declarations or definitions.
- You can specify `thread_local` only on data items with static storage duration. This includes global data objects (both `static` and `extern`), local static objects, and static data members of classes. Any local variable declared `thread_local` is implicitly static if no other storage class is provided; in other words, at block scope `thread_local` is equivalent to `thread_local static`.
- You must specify `thread_local` for both the declaration and the definition of a thread local object, whether the declaration and definition occur in the same file or separate files.

On Windows, `thread_local` is functionally equivalent to `_declspec(thread)` except that `_declspec(thread)` can be applied to a type definition and is valid in C code. Whenever possible, use `thread_local` because it is part of the C++ standard and is therefore more portable.

## register

**Visual Studio 2017 version 15.3 and later** (available with `/std:c++17`): The `register` keyword is no longer a

supported storage class. The keyword is still reserved in the standard for future use.

```
register int val; // warning C5033: 'register' is no longer a supported storage class
```

## Example: automatic vs. static initialization

A local automatic object or variable is initialized every time the flow of control reaches its definition. A local static object or variable is initialized the first time the flow of control reaches its definition.

Consider the following example, which defines a class that logs initialization and destruction of objects and then defines three objects, `I1`, `I2`, and `I3`:

```

// initialization_of_objects.cpp
// compile with: /EHsc
#include <iostream>
#include <string.h>
using namespace std;

// Define a class that logs initializations and destructions.
class InitDemo {
public:
    InitDemo( const char *szWhat );
    ~InitDemo();

private:
    char *szObjName;
    size_t sizeofObjName;
};

// Constructor for class InitDemo
InitDemo::InitDemo( const char *szWhat ) :
    szObjName(NULL), sizeofObjName(0) {
    if ( szWhat != 0 && strlen( szWhat ) > 0 ) {
        // Allocate storage for szObjName, then copy
        // initializer szWhat into szObjName, using
        // secured CRT functions.
        sizeofObjName = strlen( szWhat ) + 1;

        szObjName = new char[ sizeofObjName ];
        strcpy_s( szObjName, sizeofObjName, szWhat );

        cout << "Initializing: " << szObjName << "\n";
    }
    else {
        szObjName = 0;
    }
}

// Destructor for InitDemo
InitDemo::~InitDemo() {
    if( szObjName != 0 ) {
        cout << "Destroying: " << szObjName << "\n";
        delete szObjName;
    }
}

// Enter main function
int main() {
    InitDemo I1( "Auto I1" );
    cout << "In block.\n";
    InitDemo I2( "Auto I2" );
    static InitDemo I3( "Static I3" );
}
cout << "Exited block.\n";
}

```

```

Initializing: Auto I1
In block.
Initializing: Auto I2
Initializing: Static I3
Destroying: Auto I2
Exited block.
Destroying: Auto I1
Destroying: Static I3

```

This example demonstrates how and when the objects `I1`, `I2`, and `I3` are initialized and when they are

destroyed.

There are several points to note about the program:

- First, `I1` and `I2` are automatically destroyed when the flow of control exits the block in which they are defined.
- Second, in C++, it is not necessary to declare objects or variables at the beginning of a block. Furthermore, these objects are initialized only when the flow of control reaches their definitions. (`I2` and `I3` are examples of such definitions.) The output shows exactly when they are initialized.
- Finally, static local variables such as `I3` retain their values for the duration of the program, but are destroyed as the program terminates.

## See Also

[Declarations and Definitions](#)

# const (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

When modifying a data declaration, the **const** keyword specifies that the object or variable is not modifiable.

## Syntax

```
const declaration ;
member-function const ;
```

## const values

The **const** keyword specifies that a variable's value is constant and tells the compiler to prevent the programmer from modifying it.

```
// constant_values1.cpp
int main() {
    const int i = 5;
    i = 10;    // C3892
    i++;      // C2105
}
```

In C++, you can use the **const** keyword instead of the **#define** preprocessor directive to define constant values. Values defined with **const** are subject to type checking, and can be used in place of constant expressions. In C++, you can specify the size of an array with a **const** variable as follows:

```
// constant_values2.cpp
// compile with: /c
const int maxarray = 255;
char store_char[maxarray]; // allowed in C++; not allowed in C
```

In C, constant values default to external linkage, so they can appear only in source files. In C++, constant values default to internal linkage, which allows them to appear in header files.

The **const** keyword can also be used in pointer declarations.

```
// constant_values3.cpp
int main() {
    char *mybuf = 0, *yourbuf;
    char *const aptr = mybuf;
    *aptr = 'a';    // OK
    aptr = yourbuf; // C3892
}
```

A pointer to a variable declared as **const** can be assigned only to a pointer that is also declared as **const**.

```
// constant_values4.cpp
#include <stdio.h>
int main() {
    const char *mybuf = "test";
    char *yourbuf = "test2";
    printf_s("%s\n", mybuf);

    const char *bptr = mybuf; // Pointer to constant data
    printf_s("%s\n", bptr);

    // *bptr = 'a'; // Error
}
```

You can use pointers to constant data as function parameters to prevent the function from modifying a parameter passed through a pointer.

For objects that are declared as **const**, you can only call constant member functions. This ensures that the constant object is never modified.

```
birthday.getMonth(); // Okay
birthday.setMonth( 4 ); // Error
```

You can call either constant or nonconstant member functions for a nonconstant object. You can also overload a member function using the **const** keyword; this allows a different version of the function to be called for constant and nonconstant objects.

You cannot declare constructors or destructors with the **const** keyword.

## const member functions

Declaring a member function with the **const** keyword specifies that the function is a "read-only" function that does not modify the object for which it is called. A constant member function cannot modify any non-static data members or call any member functions that aren't constant. To declare a constant member function, place the **const** keyword after the closing parenthesis of the argument list. The **const** keyword is required in both the declaration and the definition.

```

// constant_member_function.cpp
class Date
{
public:
    Date( int mn, int dy, int yr );
    int getMonth() const;      // A read-only function
    void setMonth( int mn );   // A write function; can't be const
private:
    int month;
};

int Date::getMonth() const
{
    return month;           // Doesn't modify anything
}
void Date::setMonth( int mn )
{
    month = mn;            // Modifies data member
}
int main()
{
    Date MyDate( 7, 4, 1998 );
    const Date BirthDate( 1, 18, 1953 );
    MyDate.setMonth( 4 );    // Okay
    BirthDate.getMonth();   // Okay
    BirthDate.setMonth( 4 ); // C2662 Error
}

```

## C and C++ const Differences

When you declare a variable as **const** in a C source code file, you do so as:

```
const int i = 2;
```

You can then use this variable in another module as follows:

```
extern const int i;
```

But to get the same behavior in C++, you must declare your **const** variable as:

```
extern const int i = 2;
```

If you wish to declare an **extern** variable in a C++ source code file for use in a C source code file, use:

```
extern "C" const int x=10;
```

to prevent name mangling by the C++ compiler.

## Remarks

When following a member function's parameter list, the **const** keyword specifies that the function does not modify the object for which it is invoked.

For more information on **const**, see the following topics:

- [const and volatile Pointers](#)

- [Type Qualifiers \(C Language Reference\)](#)
- [volatile](#)
- [#define](#).

## See Also

[Keywords](#)

# constexpr (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

The keyword `constexpr` was introduced in C++11 and improved in C++14. It means *constant expression*. Like `const`, it can be applied to variables so that a compiler error will be raised if any code attempts to modify the value. Unlike `const`, `constexpr` can also be applied to functions and class constructors. `constexpr` indicates that the value, or return value, is constant and, if possible, will be computed at compile time. A `constexpr` integral value can be used wherever a `const` integer is required, such as in template arguments and array declarations. And when a value can be computed at compile time instead of run time, it can help your program run faster and use less memory.

## Syntax

```
constexpr literal-type identifier = constant-expression;
constexpr literal-type identifier { constant-expression };
constexpr literal-type identifier(params);
constexpr ctor (params);
```

### Parameters

`params`

One or more parameters which must be a literal type (as listed below) and must itself be a constant expression.

## Return Value

A `constexpr` variable or function must return one of the literal types, as listed below.

## Literal types

To limit the complexity of computing compile time constants, and their potential impacts of compilation time, the C++14 standard requires that the types involved in constant expressions be restricted to literal types. A literal type is one whose layout can be determined at compile time. The following are the literal types:

1. `void`
2. scalar types
3. references
4. Arrays of `void`, scalar types or references
5. A class that has a trivial destructor, and one or more `constexpr` constructors that are not move or copy constructors. Additionally, all its non-static data members and base classes must be literal types and not volatile.

## constexpr variables

The primary difference between `const` and `constexpr` variables is that the initialization of a `const` variable can be deferred until run time whereas a `constexpr` variable must be initialized at compile time. All `constexpr` variables are `const`.

- A variable can be declared with `constexpr`, if it has a literal type and is initialized. If the initialization is

performed by a constructor, the constructor must be declared as `constexpr`.

- A reference may be declared as `constexpr` if the object that it references has been initialized by a constant expression and any implicit conversions that are invoked during initialization are also constant expressions.
- All declarations of a `constexpr` variable or function must have the `constexpr` specifier.

```
constexpr float x = 42.0;
constexpr float y{108};
constexpr float z = exp(5, 3);
constexpr int i; // Error! Not initialized
int j = 0;
constexpr int k = j + 1; //Error! j not a constant expression
```

## constexpr functions

A `constexpr` function is one whose return value can be computed at compile time when consuming code requires it. When its arguments are `constexpr` values, and consuming code requires the return value at compile time, for example to initialize a `constexpr` variable or provide a non-type template argument, it produces a compile-time constant. When called with non-`constexpr` arguments, or when its value is not required at compile-time, it produces a value at run time like a regular function. (This dual behavior saves you from having to write `constexpr` and non-`constexpr` versions of the same function.)

A `constexpr` function or constructor is implicitly `inline`.

The following rules apply to `constexpr` functions:

- A `constexpr` function must accept and return only literal types.
- A `constexpr` function can be recursive.
- It cannot be `virtual`. A constructor cannot be defined as `constexpr` if the enclosing class has any virtual base classes.
- The body can be defined as `= default` or `= delete`.
- The body can contain no `goto` statements or try blocks.
- An explicit specialization of a non-`constexpr` template can be declared as `constexpr`:
- An explicit specialization of a `constexpr` template does not have to also be `constexpr`:

The following rules apply to `constexpr` functions in Visual Studio 2017 and later:

- It may contain if and switch statements, and all looping statements including for, range-based for, while, and do-while
- It may contain local variable declarations, but the variable must be initialized, must be a literal type, and cannot be static or thread-local. The locally-declared variable is not required to be `const` and may mutate.
- A `constexpr` non-static member function is not required to be implicitly `const`.

```
constexpr float exp(float x, int n)
{
    return n == 0 ? 1 :
        n % 2 == 0 ? exp(x * x, n / 2) :
        exp(x * x, (n - 1) / 2) * x;
};
```

**TIP**

Note: In the Visual Studio debugger, you can tell whether a `constexpr` function is being evaluated at compile time by putting a breakpoint inside it. If the breakpoint is hit, the function was called at run-time. If not, then the function was called at compile time.

## Example

The following example shows `constexpr` variables, functions and a user-defined type. Note that in the last statement in main(), the `constexpr` member function `GetValue()` is a run-time call because the value is not required to be known at compile time.

```

#include <iostream>

using namespace std;

// Pass by value
constexpr float exp(float x, int n)
{
    return n == 0 ? 1 :
        n % 2 == 0 ? exp(x * x, n / 2) :
        exp(x * x, (n - 1) / 2) * x;
};

// Pass by reference
constexpr float exp2(const float& x, const int& n)
{
    return n == 0 ? 1 :
        n % 2 == 0 ? exp2(x * x, n / 2) :
        exp2(x * x, (n - 1) / 2) * x;
};

// Compile time computation of array length
template<typename T, int N>
constexpr int length(const T(&ary)[N])
{
    return N;
}

// Recursive constexpr function
constexpr int fac(int n)
{
    return n == 1 ? 1 : n*fac(n - 1);
}

// User-defined type
class Foo
{
public:
    constexpr explicit Foo(int i) : _i(i) {}
    constexpr int GetValue()
    {
        return _i;
    }
private:
    int _i;
};

int main()
{
    //foo is const:
    constexpr Foo foo(5);
    // foo = Foo(6); //Error!

    //Compile time:
    constexpr float x = exp(5, 3);
    constexpr float y { exp(2, 5) };
    constexpr int val = foo.GetValue();
    constexpr int f5 = fac(5);
    const int nums[] { 1, 2, 3, 4 };
    const int nums2[length(nums) * 2] { 1, 2, 3, 4, 5, 6, 7, 8 };

    //Run time:
    cout << "The value of foo is " << foo.GetValue() << endl;
}

```

# Requirements

Visual Studio 2015

## See Also

[Declarations and Definitions](#)

`const`

# volatile (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

A type qualifier that you can use to declare that an object can be modified in the program by the hardware.

## Syntax

```
volatile declarator ;
```

## Remarks

You can use the [/volatile](#) compiler switch to modify how the compiler interprets this keyword.

Visual Studio interprets the `volatile` keyword differently depending on the target architecture. For ARM, if no **/volatile** compiler option is specified, the compiler performs as if **/volatile:iso** were specified. For architectures other than ARM, if no **/volatile** compiler option is specified, the compiler performs as if **/volatile:ms** were specified; therefore, for architectures other than ARM we strongly recommend that you specify **/volatile:iso**, and use explicit synchronization primitives and compiler intrinsics when you are dealing with memory that is shared across threads.

You can use the `volatile` qualifier to provide access to memory locations that are used by asynchronous processes such as interrupt handlers.

When `volatile` is used on a variable that also has the [\\_\\_restrict](#) keyword, `volatile` takes precedence.

If a `struct` member is marked as `volatile`, then `volatile` is propagated to the whole structure. If a structure does not have a length that can be copied on the current architecture by using one instruction, `volatile` may be completely lost on that structure.

The `volatile` keyword may have no effect on a field if one of the following conditions is true:

- The length of the volatile field exceeds the maximum size that can be copied on the current architecture by using one instruction.
- The length of the outermost containing `struct`—or if it's a member of a possibly nested `struct`—exceeds the maximum size that can be copied on the current architecture by using one instruction.

Although the processor does not reorder un-cacheable memory accesses, un-cacheable variables must be marked as `volatile` to guarantee that the compiler does not reorder the memory accesses.

Objects that are declared as `volatile` are not used in certain optimizations because their values can change at any time. The system always reads the current value of a volatile object when it is requested, even if a previous instruction asked for a value from the same object. Also, the value of the object is written immediately on assignment.

## ISO Compliant

If you are familiar with the C# `volatile` keyword, or familiar with the behavior of `volatile` in earlier versions of Visual C++, be aware that the C++11 ISO Standard `volatile` keyword is different and is supported in Visual Studio when the [/volatile:iso](#) compiler option is specified. (For ARM, it's specified by default). The `volatile`

keyword in C++11 ISO Standard code is to be used only for hardware access; do not use it for inter-thread communication. For inter-thread communication, use mechanisms such as `std::atomic<T>` from the [C++ Standard Library](#).

## End of ISO Compliant

## Microsoft Specific

When the `/volatile:ms` compiler option is used—by default when architectures other than ARM are targeted—the compiler generates extra code to maintain ordering among references to volatile objects in addition to maintaining ordering to references to other global objects. In particular:

- A write to a volatile object (also known as volatile write) has Release semantics; that is, a reference to a global or static object that occurs before a write to a volatile object in the instruction sequence will occur before that volatile write in the compiled binary.
- A read of a volatile object (also known as volatile read) has Acquire semantics; that is, a reference to a global or static object that occurs after a read of volatile memory in the instruction sequence will occur after that volatile read in the compiled binary.

This enables volatile objects to be used for memory locks and releases in multithreaded applications.

### NOTE

When it relies on the enhanced guarantee that's provided when the `/volatile:ms` compiler option is used, the code is non-portable.

## END Microsoft Specific

## See Also

[Keywords](#)

[const](#)

[const and volatile Pointers](#)

# auto (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

Deduces the type of a declared variable from its initialization expression.

## Syntax

```
auto declarator initializer;
```

```
[](auto param1, auto param2) {};
```

## Remarks

The `auto` keyword directs the compiler to use the initialization expression of a declared variable, or lambda expression parameter, to deduce its type.

We recommend that you use the `auto` keyword for most situations—unless you really want a conversion—because it provides these benefits:

- **Robustness:** If the expression's type is changed—this includes when a function return type is changed—it just works.
- **Performance:** You're guaranteed that there will be no conversion.
- **Usability:** You don't have to worry about type name spelling difficulties and typos.
- **Efficiency:** Your coding can be more efficient.

Conversion cases in which you might not want to use `auto`:

- When you want a specific type and nothing else will do.
- Expression template helper types—for example, `(valarray+valarray)`.

To use the `auto` keyword, use it instead of a type to declare a variable, and specify an initialization expression. In addition, you can modify the `auto` keyword by using specifiers and declarators such as `const`, `volatile`, pointer (`*`), reference (`&`), and rvalue reference (`&&`). The compiler evaluates the initialization expression and then uses that information to deduce the type of the variable.

The initialization expression can be an assignment (equal-sign syntax), a direct initialization (function-style syntax), an `operator new` expression, or the initialization expression can be the *for-range-declaration* parameter in a `Range-based for Statement (C++)` statement. For more information, see [Initializers](#) and the code examples later in this document.

The `auto` keyword is a placeholder for a type, but it is not itself a type. Therefore, the `auto` keyword cannot be used in casts or operators such as `sizeof` and `typeid`.

## Usefulness

The `auto` keyword is a simple way to declare a variable that has a complicated type. For example, you can use `auto` to declare a variable where the initialization expression involves templates, pointers to functions, or pointers

to members.

You can also use `auto` to declare and initialize a variable to a lambda expression. You can't declare the type of the variable yourself because the type of a lambda expression is known only to the compiler. For more information, see [Examples of Lambda Expressions](#).

## Trailing Return Types

You can use `auto`, together with the `decltype` type specifier, to help write template libraries. Use `auto` and `decltype` to declare a template function whose return type depends on the types of its template arguments. Or, use `auto` and `decltype` to declare a template function that wraps a call to another function, and then returns whatever is the return type of that other function. For more information, see [decltype](#).

## References and cv-qualifiers

Note that using `auto` drops references, const qualifiers, and volatile qualifiers. Consider the following example:

```
// cl.exe /analyze /EHsc /W4
#include <iostream>

using namespace std;

int main( )
{
    int count = 10;
    int& countRef = count;
    auto myAuto = countRef;

    countRef = 11;
    cout << count << " ";

    myAuto = 12;
    cout << count << endl;
}
```

In the previous example, `myAuto` is an `int`, not an `int` reference, so the output is `11 11`, not `11 12` as would be the case if the reference qualifier had not been dropped by `auto`.

## Type deduction with braced initializers (C++14)

The following code example shows how to initialize an `auto` variable using braces. Note the difference between B and C and between A and E.

```

#include <initializer_list>

int main()
{
    // std::initializer_list<int>
    auto A = { 1, 2 };

    // std::initializer_list<int>
    auto B = { 3 };

    // int
    auto C{ 4 };

    // C3535: cannot deduce type for 'auto' from initializer list'
    auto D = { 5, 6.7 };

    // C3518 in a direct-list-initialization context the type for 'auto'
    // can only be deduced from a single initializer expression
    auto E{ 8, 9 };

    return 0;
}

```

## Restrictions and Error Messages

The following table lists the restrictions on the use of the `auto` keyword, and the corresponding diagnostic error message that the compiler emits.

ERROR NUMBER	DESCRIPTION
C3530	The <code>auto</code> keyword cannot be combined with any other type-specifier.
C3531	A symbol that is declared with the <code>auto</code> keyword must have an initializer.
C3532	You incorrectly used the <code>auto</code> keyword to declare a type. For example, you declared a method return type or an array.
C3533, C3539	A parameter or template argument cannot be declared with the <code>auto</code> keyword.
C3535	A method or template parameter cannot be declared with the <code>auto</code> keyword.
C3536	A symbol cannot be used before it is initialized. In practice, this means that a variable cannot be used to initialize itself.
C3537	You cannot cast to a type that is declared with the <code>auto</code> keyword.
C3538	All the symbols in a declarator list that is declared with the <code>auto</code> keyword must resolve to the same type. For more information, see <a href="#">Declarations and Definitions</a> .
C3540, C3541	The <code>sizeof</code> and <code>typeid</code> operators cannot be applied to a symbol that is declared with the <code>auto</code> keyword.

## Examples

These code fragments illustrate some of the ways in which the `auto` keyword can be used.

The following declarations are equivalent. In the first statement, variable `j` is declared to be type `int`. In the second statement, variable `k` is deduced to be type `int` because the initialization expression (0) is an integer.

```
int j = 0; // Variable j is explicitly type int.  
auto k = 0; // Variable k is implicitly type int because 0 is an integer.
```

The following declarations are equivalent, but the second declaration is simpler than the first. One of the most compelling reasons to use the `auto` keyword is simplicity.

```
map<int,list<string>>::iterator i = m.begin();  
auto i = m.begin();
```

The following code fragment declares the type of variables `iter` and `elem` when the `for` and range `for` loops start.

```
// cl /EHsc /nologo /W4  
#include <deque>  
using namespace std;  
  
int main()  
{  
    deque<double> dqDoubleData(10, 0.1);  
  
    for (auto iter = dqDoubleData.begin(); iter != dqDoubleData.end(); ++iter)  
    { /* ... */ }  
  
    // prefer range-for loops with the following information in mind  
    // (this applies to any range-for with auto, not just deque)  
  
    for (auto elem : dqDoubleData) // COPIES elements, not much better than the previous examples  
    { /* ... */ }  
  
    for (auto& elem : dqDoubleData) // observes and/or modifies elements IN-PLACE  
    { /* ... */ }  
  
    for (const auto& elem : dqDoubleData) // observes elements IN-PLACE  
    { /* ... */ }  
}
```

The following code fragment uses the `new` operator and pointer declaration to declare pointers.

```
double x = 12.34;  
auto *y = new auto(x), **z = new auto(&x);
```

The next code fragment declares multiple symbols in each declaration statement. Notice that all of the symbols in each statement resolve to the same type.

```
auto x = 1, *y = &x, **z = &y; // Resolves to int.  
auto a(2.01), *b (&a); // Resolves to double.  
auto c = 'a', *d(&c); // Resolves to char.  
auto m = 1, &n = m; // Resolves to int.
```

This code fragment uses the conditional operator (`? :`) to declare variable `x` as an integer that has a value of 200:

```
int v1 = 100, v2 = 200;
auto x = v1 > v2 ? v1 : v2;
```

The following code fragment initializes variable `x` to type `int`, variable `y` to a reference to type `const int`, and variable `fp` to a pointer to a function that returns type `int`.

```
int f(int x) { return x; }
int main()
{
    auto x = f(0);
    const auto & y = f(1);
    int (*p)(int x);
    p = f;
    auto fp = p;
    //...
}
```

## See Also

[auto Keyword](#)  
[Keywords](#)  
[/Zc:auto \(Deduce Variable Type\)](#)  
[sizeof Operator](#)  
[typeid](#)  
[operator new](#)  
[Declarations and Definitions](#)  
[Examples of Lambda Expressions](#)  
[Initializers](#)  
[decltype](#)

# decltype (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

The `decltype` type specifier yields the type of a specified expression. The `decltype` type specifier, together with the `auto` keyword, is useful primarily to developers who write template libraries. Use `auto` and `decltype` to declare a template function whose return type depends on the types of its template arguments. Or, use `auto` and `decltype` to declare a template function that wraps a call to another function, and then returns the return type of the wrapped function.

## Syntax

```
decltype( expression )
```

### Parameters

PARAMETER	DESCRIPTION
<code>expression</code>	An expression. For more information, see <a href="#">Expressions</a> .

## Return Value

The type of the `expression` parameter.

## Remarks

The `decltype` type specifier is supported in Visual C++ 2010 or later versions, and can be used with native or managed code. `decltype(auto)` (C++14) is supported in Visual Studio 2015 and later.

The compiler uses the following rules to determine the type of the `expression` parameter.

- If the `expression` parameter is an identifier or a [class member access](#), `decltype(expression)` is the type of the entity named by `expression`. If there is no such entity or the `expression` parameter names a set of overloaded functions, the compiler yields an error message.
- If the `expression` parameter is a call to a function or an overloaded operator function, `decltype(expression)` is the return type of the function. Parentheses around an overloaded operator are ignored.
- If the `expression` parameter is an [rvalue](#), `decltype(expression)` is the type of `expression`. If the `expression` parameter is an [lvalue](#), `decltype(expression)` is an [lvalue reference](#) to the type of `expression`.

The following code example demonstrates some uses of the `decltype` type specifier. First, assume that you have coded the following statements.

```
int var;
const int&& fx();
struct A { double x; }
const A* a = new A();
```

Next, examine the types that are returned by the four `decltype` statements in the following table.

STATEMENT	TYPE	NOTES
<code>decltype(fx());</code>	<code>const int&amp;&amp;</code>	An <a href="#">rvalue reference</a> to a <code>const int</code> .
<code>decltype(var);</code>	<code>int</code>	The type of variable <code>var</code> .
<code>decltype(a-&gt;x);</code>	<code>double</code>	The type of the member access.
<code>decltype((a-&gt;x));</code>	<code>const double&amp;</code>	The inner parentheses cause the statement to be evaluated as an expression instead of a member access. And because <code>a</code> is declared as a <code>const</code> pointer, the type is a reference to <code>const double</code> .

## Decltype and Auto

In C++14, you can use `decltype(auto)` with no trailing return type to declare a template function whose return type depends on the types of its template arguments.

In C++11, you can use the `decltype` type specifier on a trailing return type, together with the `auto` keyword, to declare a template function whose return type depends on the types of its template arguments. For example, consider the following code example in which the return type of the template function depends on the types of the template arguments. In the code example, the *UNKNOWN* placeholder indicates that the return type cannot be specified.

```
template<typename T, typename U>
UNKNOWN func(T&& t, U&& u){ return t + u; };
```

The introduction of the `decltype` type specifier enables a developer to obtain the type of the expression that the template function returns. Use the *alternative function declaration syntax* that is shown later, the `auto` keyword, and the `decltype` type specifier to declare a *late-specified* return type. The late-specified return type is determined when the declaration is compiled, instead of when it is coded.

The following prototype illustrates the syntax of an alternative function declaration. Note that the `const` and `volatile` qualifiers, and the `throw` exception specification are optional. The *function\_body* placeholder represents a compound statement that specifies what the function does. As a best coding practice, the *expression* placeholder in the `decltype` statement should match the expression specified by the `return` statement, if any, in the *function\_body*.

```
auto function_name ( parametersopt ) constopt volatileopt -> decltype( expression ) throwopt { function_body };
```

In the following code example, the late-specified return type of the `myFunc` template function is determined by the types of the `t` and `u` template arguments. As a best coding practice, the code example also uses rvalue references and the `forward` function template, which support *perfect forwarding*. For more information, see [Rvalue Reference Declarator: &&](#).

```
//C++11
template<typename T, typename U>
auto myFunc(T&& t, U&& u) -> decltype (forward<T>(t) + forward<U>(u))
    { return forward<T>(t) + forward<U>(u); }

//C++14
template<typename T, typename U>
decltype(auto) myFunc(T&& t, U&& u)
    { return forward<T>(t) + forward<U>(u); }
```

## Decltype and Forwarding Functions (C++11)

Forwarding functions wrap calls to other functions. Consider a function template that forwards its arguments, or the results of an expression that involves those arguments, to another function. Furthermore, the forwarding function returns the result of calling the other function. In this scenario, the return type of the forwarding function should be the same as the return type of the wrapped function.

In this scenario, you cannot write an appropriate type expression without the `decltype` type specifier. The `decltype` type specifier enables generic forwarding functions because it does not lose required information about whether a function returns a reference type. For a code example of a forwarding function, see the previous `myFunc` template function example.

## Example

The following code example declares the late-specified return type of template function `Plus()`. The `Plus` function processes its two operands with the `operator+` overload. Consequently, the interpretation of the plus operator (+) and the return type of the `Plus` function depends on the types of the function arguments.

```

// decltype_1.cpp
// compile with: cl /EHsc decltype_1.cpp

#include <iostream>
#include <string>
#include <utility>
#include <iomanip>

using namespace std;

template<typename T1, typename T2>
auto Plus(T1&& t1, T2&& t2) ->
    decltype(forward<T1>(t1) + forward<T2>(t2))
{
    return forward<T1>(t1) + forward<T2>(t2);
}

class X
{
    friend X operator+(const X& x1, const X& x2)
    {
        return X(x1.m_data + x2.m_data);
    }

public:
    X(int data) : m_data(data) {}
    int Dump() const { return m_data; }
private:
    int m_data;
};

int main()
{
    // Integer
    int i = 4;
    cout <<
        "Plus(i, 9) = " <<
        Plus(i, 9) << endl;

    // Floating point
    float dx = 4.0;
    float dy = 9.5;
    cout <<
        setprecision(3) <<
        "Plus(dx, dy) = " <<
        Plus(dx, dy) << endl;

    // String
    string hello = "Hello, ";
    string world = "world!";
    cout << Plus(hello, world) << endl;

    // Custom type
    X x1(20);
    X x2(22);
    X x3 = Plus(x1, x2);
    cout <<
        "x3.Dump() = " <<
        x3.Dump() << endl;
}

```

```

Plus(i, 9) = 13
Plus(dx, dy) = 13.5
Hello, world!
x3.Dump() = 42

```

## Example

**Visual Studio 2017 and later:** The compiler parses decltype arguments when the templates are declared rather than instantiated. Consequently, if a non-dependent specialization is found in the decltype argument, it will not be deferred to instantiation-time and will be processed immediately and any resulting errors will be diagnosed at that time.

The following example shows such a compiler error that is raised at the point of declaration:

```
#include <utility>
template <class T, class ReturnT, class... ArgsT> class IsCallable
{
public:
    struct BadType {};
    template <class U>
    static decltype(std::declval<T>()(std::declval<ArgsT>(...)) Test(int); //C2064. Should be declval<U>
    template <class U>
    static BadType Test(...);
    static constexpr bool value = std::is_convertible<decltype(Test<T>()), ReturnT>::value;
};

constexpr bool test1 = IsCallable<int(), int>::value;
static_assert(test1, "PASS1");
constexpr bool test2 = !IsCallable<int*, int>::value;
static_assert(test2, "PASS2");
```

## Requirements

Visual C++ 2010 or later versions.

`decltype(auto)` requires Visual Studio 2015 or later.

# Attributes in C++

10/31/2017 • 4 min to read • [Edit Online](#)

The C++ Standard defines a set of attributes and also allows compiler vendors to define their own attributes (within a vendor-specific namespace), but compilers are required to recognize only those attributes defined in the standard.

In some cases, standard attributes overlap with compiler-specific declspec parameters. In Visual C++, you can use the `[[deprecated]]` attribute instead of using `declspec(deprecated)` and the attribute will be recognized by any conformant compiler. For all other declspec parameters such as `dllimport` and `dllexport`, there is as yet no attribute equivalent so you must continue to use declspec syntax. Attributes do not affect the type system, and they don't change the meaning of a program. Compilers ignore attribute values they don't recognize.

**Visual Studio 2017 version 15.3 and later** (available with `/std:c++17`): In the scope of an attribute list, you can specify the namespace for all names with a single `using` introducer:

```
void g() {
    [[using rpr:: kernel, target(cpu,gpu)]] // equivalent to [[ rpr::kernel, rpr::target(cpu,gpu) ]]
    do task();
}
```

## C++ Standard Attributes

In C++11, attributes provide a standardized way to annotate C++ constructs (including but not limited to classes, functions, variables, and blocks) with additional information that may or may not be vendor-specific. A compiler can use this information to generate informational messages, or to apply special logic when compiling the attributed code. The compiler ignores any attributes that it does not recognize, which means that you cannot define your own custom attributes using this syntax. Attributes are enclosed by double square brackets:

```
[[deprecated]]
void Foo(int);
```

Attributes represent a standardized alternative to vendor-specific extensions such as `#pragma` directives, `_declspec()` (Visual C++), or `_attribute_` (GNU). However, you will still need to use the vendor-specific constructs for most purposes. The standard currently specifies the following attributes that a conforming compiler should recognize:

- `[[noreturn]]` Specifies that a function never returns; in other words it always throws an exception. The compiler can adjust its compilation rules for `[[noreturn]]` entities.
- `[[carries_dependency]]` Specifies that the function propagates data dependency ordering with respect to thread synchronization. The attribute can be applied to one or more parameters, to specify that the passed-in argument carries a dependency into the function body. The attribute can be applied to the function itself, to specify that the return value carries a dependency out of the function. The compiler can use this information to generate more efficient code.
- `[[deprecated]]` **Visual Studio 2015 and later:** Specifies that a function is not intended to be used, and might not exist in future versions of a library interface. The compiler can use this to generate an informational message when client code attempts to call the function. Can be applied to declaration of a class, a `typedef`-name, a variable, a non-static data member, a function, a namespace, an enumeration, an

enumerator, or a template specialization.

- **[[fallthrough]] Visual Studio 2017 and later:** (available with [/std:c++17](#)) The `[[fallthrough]]` attribute can be used in the context of `switch` statements as a hint to the compiler (or anyone reading the code) that the fallthrough behavior is intended. The Visual C++ compiler currently does not warn on fallthrough behavior, so this attribute has no effect compiler behavior.
- **[[nodiscard]] Visual Studio 2017 version 15.3 and later:** (available with [/std:c++17](#)) Specifies that a function's return value is not intended to be discarded. Raises warning C4834, as shown in this example:

```
[[nodiscard]]
int foo(int i) { return i * i; }

int main()
{
    foo(42); //warning C4834: discarding return value of function with 'nodiscard' attribute
    return 0;
}
```

- **[[maybe\_unused]] Visual Studio 2017 version 15.3 and later:** (available with [/std:c++17](#)) Specifies that a variable, function, class, typedef, non-static data member, enum, or template specialization may intentionally not be used. The compiler does not warn when an entity marked `[[maybe_unused]]` is not used. An entity that is declared without the attribute can later be redeclared with the attribute and vice versa. An entity is considered marked after its first declaration that is marked is analyzed, and for the remainder of translation of the current translation unit.

## Microsoft-specific attributes

- **[[gsl::suppress(rules)]]** This Microsoft-specific attribute is used for suppressing warnings from checkers that enforce [Guidelines Support Library \(GSL\)](#) rules in code. For example, consider this code snippet:

```
void main()
{
    int arr[10]; // GSL warning 26494 will be fired
    int* p = arr; // GSL warning 26485 will be fired
    [[gsl::suppress(bounds.1)]] // This attribute suppresses Bounds rule #1
    {
        int* q = p + 1; // GSL warning 26481 suppressed
        p = q--; // GSL warning 26481 suppressed
    }
}
```

The example raises these warnings:

- 26494 (Type Rule 5: Always initialize an object.)
- 26485 (Bounds Rule 3: No array to pointer decay.)
- 26481 (Bounds Rule 1: Don't use pointer arithmetic. Use span instead.)

The first two warnings fire when you compile this code with the CppCoreCheck code analysis tool installed and activated. But the third warning doesn't fire because of the attribute. You can suppress the entire bounds profile by writing `[[gsl::suppress(bounds)]]` without including a specific rule number. The C++ Core Guidelines are designed to help you write better and safer code. The suppress attribute makes it easy to turn off the warnings when they are not wanted.

# Scope (Visual C++)

12/21/2017 • 4 min to read • [Edit Online](#)

C++ names can be used only in certain regions of a program. This area is called the "scope" of the name. Scope determines the "lifetime" of a name that does not denote an object of static extent. Scope also determines the visibility of a name, when class constructors and destructors are called, and when variables local to the scope are initialized. (For more information, see [Constructors](#) and [Destructors](#).) There are five kinds of scope:

- **Local scope** A name declared within a block is accessible only within that block and blocks enclosed by it, and only after the point of declaration. The names of formal arguments to a function in the scope of the outermost block of the function have local scope, as if they had been declared inside the block enclosing the function body. Consider the following code fragment:

```
{  
    int i;  
}
```

Because the declaration of `i` is in a block enclosed by curly braces, `i` has local scope and is never accessible because no code accesses it before the closing curly brace.

- **Function scope** Labels are the only names that have function scope. They can be used anywhere within a function, but are not accessible outside that function. Formal arguments (arguments specified in function definitions) to functions are considered to be in the scope of the outermost block of the function body.
- **File scope** Any name declared outside all blocks or classes has file scope. It is accessible anywhere in the translation unit after its declaration. Names with file scope that do not declare static objects are often called global names.

In C++, file scope is also known as namespace scope.

- **Class scope** Names of class members have class scope. Class member functions can be accessed only by using the member-selection operators (`.` or `->`) or pointer-to-member operators (`.*` or `->*`) on an object or pointer to an object of that class; nonstatic class member data is considered local to the object of that class. Consider the following class declaration:

```
class Point  
{  
    int x;  
    int y;  
};
```

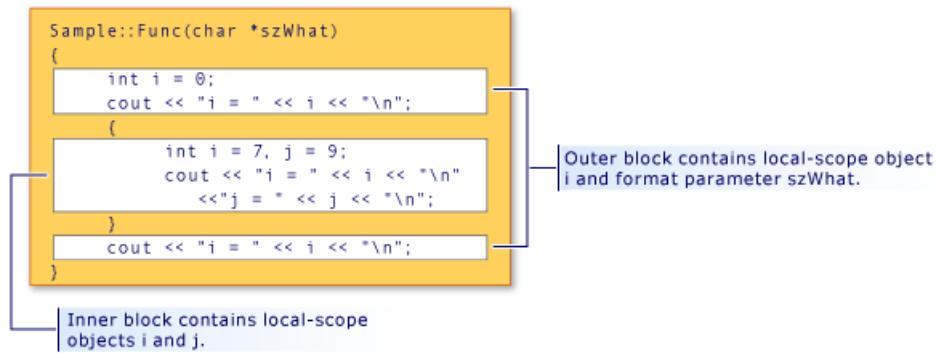
The class members `x` and `y` are considered to be in the scope of class `Point`.

- **Prototype scope** Names declared in a function prototype are visible only until the end of the prototype. The following prototype declares three names (`strDestination`, `numberOfElements`, and `strSource`); these names go out of scope at the end of the prototype:

```
errno_t strcpy_s( char *strDestination, size_t numberOfElements, const char *strSource );
```

## Hiding Names

You can hide a name by declaring it in an enclosed block. In the following figure, `i` is redeclared within the inner block, thereby hiding the variable associated with `i` in the outer block scope.



## Block Scope and Name Hiding

The output from the program shown in the figure is:

```
i = 0
i = 7
j = 9
i = 0
```

### NOTE

The argument `szWhat` is considered to be in the scope of the function. Therefore, it is treated as if it had been declared in the outermost block of the function.

## Hiding class names

You can hide class names by declaring a function, object or variable, or enumerator in the same scope. However, the class name can still be accessed when prefixed by the keyword `class`.

```

// hiding_class_names.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

// Declare class Account at file scope.
class Account
{
public:
    Account( double InitialBalance )
        { balance = InitialBalance; }
    double GetBalance()
        { return balance; }
private:
    double balance;
};

double Account = 15.37;           // Hides class name Account

int main()
{
    class Account Checking( Account ); // Qualifies Account as
                                         // class name

    cout << "Opening account with balance of: "
        << Checking.GetBalance() << "\n";
}
//Output: Opening account with balance of: 15.37

```

#### NOTE

Any place the class name (`Account`) is called for, the keyword `class` must be used to differentiate it from the file-scoped variable `Account`. This rule does not apply when the class name occurs on the left side of the scope-resolution operator (`::`). Names on the left side of the scope-resolution operator are always considered class names.

The following example demonstrates how to declare a pointer to an object of type `Account` using the `class` keyword:

```
class Account *Checking = new class Account( Account );
```

The `Account` in the initializer (in parentheses) in the preceding statement has file scope; it is of type `double`.

#### NOTE

The reuse of identifier names as shown in this example is considered poor programming style.

For more information about pointers, see [Derived Types](#). For information about declaration and initialization of class objects, see [Classes, Structures, and Unions](#). For information about using the `new` and `delete` free-store operators, see [new and delete operators](#).

## Hiding names with file scope

You can hide names with file scope by explicitly declaring the same name in block scope. However, file-scope names can be accessed using the scope-resolution operator (`::`).

```
// file_scopes.cpp
// compile with: /EHsc
#include <iostream>

int i = 7;    // i has file scope, outside all blocks
using namespace std;

int main( int argc, char *argv[] ) {
    int i = 5;    // i has block scope, hides i at file scope
    cout << "Block-scoped i has the value: " << i << "\n";
    cout << "File-scoped i has the value: " << ::i << "\n";
}
```

```
Block-scoped i has the value: 5
File-scoped i has the value: 7
```

## See Also

[Basic Concepts](#)

# Program and Linkage (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

A program consists of one or more translation units linked together. Execution (conceptually) begins in the translation unit that contains the function **main**. (For more information on translation units, see [Phases of Translation](#), in the *Preprocessor Reference*.) For more information about the **main** function, see [Program Startup: the main Function](#).)

## See Also

[Basic Concepts](#)

# Types of Linkage

12/21/2017 • 1 min to read • [Edit Online](#)

The way the names of objects and functions are shared between translation units is called linkage. These names can have:

- Internal linkage, in which case they refer only to program elements inside their own translation units; they are not shared with other translation units.

The same name in another translation unit may refer to a different object or a different class. Names with internal linkage are sometimes referred to as being local to their translation units.

An example declaration of a name with internal linkage is:

```
static int i; // The static keyword ensures internal linkage.
```

- External linkage, in which case they can refer to program elements in any translation unit in the program — the program element is shared among the translation units.

The same name in another translation unit is guaranteed to refer to the same object or class. Names with external linkage are sometimes referred to as being global.

An example declaration of a name with external linkage is:

```
extern int i;
```

- No linkage, in which case they refer to unique entities. The same name in another scope may not refer to the same object. An example is an enumeration. (Note, however, that you can pass a pointer to an object with no linkage. This makes the object accessible in other translation units.)

## See Also

[Program and Linkage](#)

# Linkage in Names with File Scope

12/21/2017 • 1 min to read • [Edit Online](#)

The following linkage rules apply to names (other than `typedef` and enumerator names) with file scope:

- If a name is explicitly declared as **static**, it has internal linkage and identifies a program element inside its own translation unit.
- Enumerator names and `typedef` names have no linkage.
- All other names with file scope have external linkage.

## Microsoft Specific

- If a function name with file scope is explicitly declared as **inline**, it has external linkage if it is instantiated or its address is referenced. Therefore, it is possible for a function with file scope to have either internal or external linkage.

## END Microsoft Specific

A class has internal linkage if it:

- Uses no C++ functionality (for example, member-access control, member functions, constructors, destructors, and so on).
- Is not used in the declaration of another name that has external linkage. This constraint means that objects of class type that are passed to functions with external linkage cause the class to have external linkage.

## See Also

[Program and Linkage](#)

# Linkage in Names with Class Scope

12/21/2017 • 1 min to read • [Edit Online](#)

The following linkage rules apply to names with class scope:

- Static class members have external linkage.
- Class member functions have external linkage.
- Enumerators and `typedef` names have no linkage.

## Microsoft Specific

- Functions declared as `friend` functions must have external linkage. Declaring a static function as a `friend` generates an error.

## END Microsoft Specific

## See Also

[Program and Linkage](#)

# Linkage in Names with Block Scope

12/21/2017 • 1 min to read • [Edit Online](#)

The following linkage rules apply to names with block scope (local names):

- Names declared as `extern` have external linkage unless they were previously declared as `static`.
- All other names with block scope have no linkage.

## See Also

[Program and Linkage](#)

# Names with No Linkage

12/21/2017 • 1 min to read • [Edit Online](#)

The only names that have no linkage are:

- Function parameters.
- Block-scoped names not declared as `extern` or `static`.
- Enumerators.
- Names declared in a `typedef` statement. An exception is when the `typedef` statement is used to provide a name for an unnamed class type. The name may then have external linkage if the class has external linkage.

The following example shows a situation in which a `typedef` name has external linkage:

```
// names_with_no_linkage.cpp
typedef struct
{
    short x;
    short y;
} POINT;

extern int MoveTo( POINT pt );

int main()
{}
```

The `typedef` name, `POINT`, becomes the class name for the unnamed structure. It is then used to declare a function with external linkage.

Because `typedef` names have no linkage, their definitions can differ between translation units. Because the compilations take place discretely, there is no way for the compiler to detect these differences. As a result, errors of this kind are not detected until link time. Consider the following case:

```
// Translation unit 1
typedef int INT

INT myInt;
...

// Translation unit 2
typedef short INT

extern INT myInt;
...
```

The preceding code generates an "unresolved external" error at link time.

## Example

C++ functions can be defined only in file or class scope. The following example illustrates how to define functions and shows an erroneous function definition:

```
// function_definitions.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
void ShowChar( char ch );      // Declare function ShowChar.

void ShowChar( char ch )      // Define function ShowChar.
{
                           // Function has file scope.
    cout << ch;
}

struct Char              // Define class Char.
{
    char Show();          // Declare Show function.
    char Get();           // Declare Get function.
    char ch;
};

char Char::Show()          // Define Show function
{                         // with class scope.
    cout << ch;
    return ch;
}

void GoodFuncDef( char ch ) // Define GoodFuncDef
{                         // with file scope.
    int BadFuncDef( int i ) // C2601, Erroneous attempt to
    {                     // nest functions.
        return i * 7;
    }
    for( int i = 0; i < BadFuncDef( 2 ); ++i )
        cout << ch;
    cout << "\n";
}
```

## See Also

[Program and Linkage](#)

# Using `extern` to Specify Linkage

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
extern string-literal { declaration-list }
extern string-literal declaration
```

## Remarks

The `extern` keyword declares a variable or function and specifies that it has external linkage (its name is visible from files other than the one in which it's defined). When modifying a variable, `extern` specifies that the variable has static duration (it is allocated when the program begins and deallocated when the program ends). The variable or function may be defined in another source file, or later in the same file. Declarations of variables and functions at file scope are external by default.

## Example

```
// specifying_linkage1.cpp
int i = 1;
void other();

int main() {
    // Reference to i, defined above:
    extern int i;
}

void other() {
    // Address of global i assigned to pointer variable:
    static int *external_i = &i;

    // i will be redefined; global i no longer visible:
    // int i = 16;
}
```

In C++, when used with a string, `extern` specifies that the linkage conventions of another language are being used for the declarator(s). C functions and data can be accessed only if they are previously declared as having C linkage. However, they must be defined in a separately compiled translation unit.

Microsoft C++ supports the strings "`C`" and "`C++`" in the *string-literal* field. All of the standard include files use the `extern "C"` syntax to allow the run-time library functions to be used in C++ programs.

## Example

The following example shows alternative ways to declare names that have C linkage:

```

// specifying_linkage2.cpp
// compile with: /c
// Declare printf with C linkage.
extern "C" int printf( const char *fmt, ... );

// Cause everything in the specified header files
// to have C linkage.
extern "C" {
    // add your #include statements here
    #include <stdio.h>
}

// Declare the two functions ShowChar and GetChar
// with C linkage.
extern "C" {
    char ShowChar( char ch );
    char GetChar( void );
}

// Define the two functions ShowChar and GetChar
// with C linkage.
extern "C" char ShowChar( char ch ) {
    putchar( ch );
    return ch;
}

extern "C" char GetChar( void ) {
    char ch;

    ch = getchar();
    return ch;
}

// Declare a global variable, errno, with C linkage.
extern "C" int errno;

```

If a function has more than one linkage specification, they must agree; it is an error to declare functions as having both C and C++ linkage. Furthermore, if two declarations for a function occur in a program — one with a linkage specification and one without — the declaration with the linkage specification must be first. Any redundant declarations of functions that already have linkage specification are given the linkage specified in the first declaration. For example:

```

extern "C" int CFunc1();
...
int CFunc1();           // Redeclaration is benign; C linkage is
                       // retained.

int CFunc2();
...
extern "C" int CFunc2(); // Error: not the first declaration of
                       // CFunc2; cannot contain linkage
                       // specifier.

```

Functions and objects explicitly declared as **static** within the body of a compound linkage specifier (`{ }`) are treated as static functions or objects; the linkage specifier is ignored. Other functions and objects behave as if declared using the `extern` keyword. (See [Using `extern` to Specify Linkage](#) for details about the `extern` keyword.)

## See Also

[Keywords](#)

[extern Storage-Class Specifier](#)

## Behavior of Identifiers

### Linkage

# Startup and Termination (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Program startup and termination are facilitated by using two functions: `main` and `exit`. Other startup and termination code may be executed.

## See Also

[Basic Concepts](#)

# main: Program Startup

12/21/2017 • 1 min to read • [Edit Online](#)

A special function named `main` is the starting point of execution for all C and C++ programs. If you are writing code that adheres to the Unicode programming model, you can use `wmain`, which is the wide-character version of `main`.

The `main` function is not predefined by the compiler. It must be supplied in the program text.

The declaration syntax for `main` is

```
int main();
```

or, optionally,

```
int main(int argc, char *argv[], char *envp[]);
```

## Microsoft Specific

The declaration syntax for `wmain` is as follows:

```
int wmain( );
```

or, optionally,

```
int wmain(int argc, wchar_t *argv[], wchar_t *envp[]);
```

You can also use `_tmain`, which is defined in TCHAR.h. `_tmain` resolves to `main` unless `_UNICODE` is defined. In that case, `_tmain` resolves to `wmain`.

Alternatively, the `main` and `wmain` functions can be declared as returning `void` (no return value). If you declare `main` or `wmain` as returning `void`, you cannot return an exit code to the parent process or operating system by using a `return` statement. To return an exit code when `main` or `wmain` is declared as `void`, you must use the `exit` function.

### END Microsoft Specific

The types for `argc` and `argv` are defined by the language. The names `argc`, `argv`, and `envp` are traditional, but are not required by the compiler. For more information and an example, see [Argument Definitions](#).

## See Also

[Keywords](#)

[Using wmain Instead of main](#)

[main Function Restrictions](#)

# Using wmain Instead of main

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

In the Unicode programming model, you can define a wide-character version of the **main** function. Use **wmain** instead of **main** if you want to write portable code that adheres to the Unicode specification.

You declare formal parameters to **wmain** using a similar format to **main**. You can then pass wide-character arguments and, optionally, a wide-character environment pointer to the program. The `argv` and `envp` parameters to **wmain** are of type `wchar_t*`.

If your program uses a **main** function, the multibyte-character environment is created by the operating system at program startup. A wide-character copy of the environment is created only when needed (for example, by a call to the `_wgetenv` or `_wpurplenv` functions). On the first call to `_wpurplenv`, or on the first call to `_wgetenv` if an MBCS environment already exists, a corresponding wide-character string environment is created and is then pointed to by the `_wenviron` global variable, which is a wide-character version of the `_environ` global variable. At this point, two copies of the environment (MBCS and Unicode) exist simultaneously and are maintained by the operating system throughout the life of the program.

Similarly, if your program uses a **wmain** function, an MBCS (ASCII) environment is created on the first call to `_putenv` or `getenv`, and is pointed to by the `_environ` global variable.

For more information on the MBCS environment, see [Single-byte and Multibyte Character Sets](#) in the *Run-Time Library Reference*.

## END Microsoft Specific

## See Also

[main: Program Startup](#)

# Argument Definitions

12/21/2017 • 2 min to read • [Edit Online](#)

The arguments in the prototype

```
int main( int  
argc[ ,char*argv[] [,char*envp[] ] ] );intwmain(intargc[ ,wchar_t*argv[] [,wchar_t*envp[] ] ] );
```

allow convenient command-line parsing of arguments and, optionally, access to environment variables. The argument definitions are as follows:

`argc`

An integer that contains the count of arguments that follow in `argv`. The `argc` parameter is always greater than or equal to 1.

`argv`

An array of null-terminated strings representing command-line arguments entered by the user of the program. By convention, `argv [0]` is the command with which the program is invoked, `argv [1]` is the first command-line argument, and so on, until `argv [ argc ]`, which is always **NULL**. See [Customizing Command Line Processing](#) for information on suppressing command-line processing.

The first command-line argument is always `argv [1]` and the last one is `argv [ argc - 1 ]`.

## NOTE

By convention, `argv [0]` is the command with which the program is invoked. However, it is possible to spawn a process using [CreateProcess](#) and if you use both the first and second arguments (`lpApplicationName` and `lpCommandLine`), `argv [0]` may not be the executable name; use [GetModuleFileName](#) to retrieve the executable name, and its fully-qualified path.

## Microsoft Specific

`envp`

The `envp` array, which is a common extension in many UNIX systems, is used in Microsoft C++. It is an array of strings representing the variables set in the user's environment. This array is terminated by a **NULL** entry. It can be declared as an array of pointers to `char (char *envp[ ])` or as a pointer to pointers to `char (char **envp)`. If your program uses `wmain` instead of `main`, use the `wchar_t` data type instead of `char`. The environment block passed to `main` and `wmain` is a "frozen" copy of the current environment. If you subsequently change the environment via a call to `putenv` or `_wputenv`, the current environment (as returned by `getenv` / `_wgetenv` and the `_environ` / `_wenviron` variable) will change, but the block pointed to by `envp` will not change. See [Customizing Command Line Processing](#) for information on suppressing environment processing. This argument is ANSI compatible in C, but not in C++.

## END Microsoft Specific

## Example

The following example shows how to use the `argc`, `argv`, and `envp` arguments to `main`:

```
// argument_definitions.cpp
// compile with: /EHsc
#include <iostream>
#include <string.h>

using namespace std;
int main( int argc, char *argv[], char *envp[] ) {
    int iNumberLines = 0;      // Default is no line numbers.

    // If /n is passed to the .exe, display numbered listing
    // of environment variables.

    if ( (argc == 2) && _stricmp( argv[1], "/n" ) == 0 )
        iNumberLines = 1;

    // Walk through list of strings until a NULL is encountered.
    for( int i = 0; envp[i] != NULL; ++i ) {
        if( iNumberLines )
            cout << i << ": " << envp[i] << "\n";
    }
}
```

## See Also

[main: Program Startup](#)

# Wildcard Expansion

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

You can use wildcards — the question mark (?) and asterisk (\*) — to specify filename and path arguments on the command-line.

Command-line arguments are handled by a routine called `_setargv` (or `_wsetargv` in the wide-character environment), which by default does not expand wildcards into separate strings in the `argv` string array. For more information on enabling wildcard expansion, refer to [Expanding Wildcard Arguments](#).

**END Microsoft Specific**

## See Also

[main: Program Startup](#)

# Parsing C++ Command-Line Arguments

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Microsoft C/C++ startup code uses the following rules when interpreting arguments given on the operating system command line:

- Arguments are delimited by white space, which is either a space or a tab.
- The caret character (^) is not recognized as an escape character or delimiter. The character is handled completely by the command-line parser in the operating system before being passed to the `argv` array in the program.
- A string surrounded by double quotation marks ("string") is interpreted as a single argument, regardless of white space contained within. A quoted string can be embedded in an argument.
- A double quotation mark preceded by a backslash (\") is interpreted as a literal double quotation mark character ("").
- Backslashes are interpreted literally, unless they immediately precede a double quotation mark.
- If an even number of backslashes is followed by a double quotation mark, one backslash is placed in the `argv` array for every pair of backslashes, and the double quotation mark is interpreted as a string delimiter.
- If an odd number of backslashes is followed by a double quotation mark, one backslash is placed in the `argv` array for every pair of backslashes, and the double quotation mark is "escaped" by the remaining backslash, causing a literal double quotation mark ("") to be placed in `argv`.

## Example

The following program demonstrates how command-line arguments are passed:

```
// command_line_arguments.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
int main( int argc,      // Number of strings in array argv
          char *argv[],    // Array of command-line argument strings
          char *envp[] )   // Array of environment variable strings
{
    int count;

    // Display each command-line argument.
    cout << "\nCommand-line arguments:\n";
    for( count = 0; count < argc; count++ )
        cout << " " argv[count] << " "
              << argv[count] << "\n";
}
```

The following table shows example input and expected output, demonstrating the rules in the preceding list.

## Results of Parsing Command Lines

COMMAND-LINE INPUT	ARGV[1]	ARGV[2]	ARGV[3]
"abc" d e	abc	d	e
a\\b d"e f"g h	a\\b	de fg	h
a\\\"b c d	a\"b	c	d
a\\\\\"b c" d e	a\\b c	d	e

**END Microsoft Specific**

## See Also

[main: Program Startup](#)

# Customizing C++ Command-Line Processing

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

If your program does not take command-line arguments, you can save a small amount of space by suppressing use of the library routine that performs command-line processing. This routine is called `_setargv` and is described in [Wildcard Expansion](#). To suppress its use, define a routine that does nothing in the file containing the `main` function, and name it `_setargv`. The call to `_setargv` is then satisfied by your definition of `_setargv`, and the library version is not loaded.

Similarly, if you never access the environment table through the `envp` argument, you can provide your own empty routine to be used in place of `_setenvp`, the environment-processing routine. Just as with the `_setargv` function, `_setenvp` must be declared as `extern "C"`.

Your program might make calls to the `spawn` or `exec` family of routines in the C run-time library. If this is the case, you should not suppress the environment-processing routine, since this routine is used to pass an environment from the parent process to the child process.

**END Microsoft Specific**

## See Also

[main: Program Startup](#)

# main Function Restrictions

12/21/2017 • 1 min to read • [Edit Online](#)

Several restrictions apply to the **main** function that do not apply to any other C++ functions. The **main** function:

- Cannot be overloaded (see [Function Overloading](#)).
- Cannot be declared as **inline**.
- Cannot be declared as **static**.
- Cannot have its address taken.
- Cannot be called.

## See Also

[main: Program Startup](#)

# Program Termination

12/21/2017 • 1 min to read • [Edit Online](#)

In C++, there are several ways to exit a program:

- Call the [exit](#) function.
- Call the [abort](#) function.
- Execute a [return](#) statement from **main**.

## See Also

[main: Program Startup](#)

# exit Function

12/21/2017 • 1 min to read • [Edit Online](#)

The **exit** function, declared in the standard include file STDLIB.H, terminates a C++ program.

The value supplied as an argument to **exit** is returned to the operating system as the program's return code or exit code. By convention, a return code of zero means that the program completed successfully.

## NOTE

You can use the constants `EXIT_FAILURE` and `EXIT_SUCCESS`, defined in STDLIB.H, to indicate success or failure of your program.

Issuing a `return` statement from the **main** function is equivalent to calling the **exit** function with the return value as its argument.

For more information, see [exit](#) in the *Run-Time Library Reference*.

## See Also

[Program Termination](#)

# abort Function

12/21/2017 • 1 min to read • [Edit Online](#)

The **abort** function, also declared in the standard include file STDLIB.H, terminates a C++ program. The difference between **exit** and **abort** is that **exit** allows the C++ run-time termination processing to take place (global object destructors will be called), whereas **abort** terminates the program immediately. For more information, see **abort** in the *Run-Time Library Reference*.

## See also

[Program Termination](#)

# return Statement in Program Termination (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Issuing a `return` statement from `main` is functionally equivalent to calling the `exit` function. Consider the following example:

```
// return_statement.cpp
#include <stdlib.h>
int main()
{
    exit( 3 );
    return 3;
}
```

The `exit` and `return` statements in the preceding example are functionally identical. However, C++ requires that functions that have return types other than `void` return a value. The `return` statement allows you to return a value from `main`.

## See Also

[Program Termination](#)

# Additional Startup Considerations

12/21/2017 • 1 min to read • [Edit Online](#)

In C++, object construction and destruction can involve executing user code. Therefore, it is important to understand which initializations happen before entry to **main** and which destructors are invoked after exit from **main**. (For detailed information about construction and destruction of objects, see [Constructors](#) and [Destructors](#).)

The following initializations take place prior to entry to **main**:

- Default initialization of static data to zero. All static data without explicit initializers are set to zero prior to executing any other code, including run-time initialization. Static data members must still be explicitly defined.
- Initialization of global static objects in a translation unit. This may occur either before entry to **main** or before the first use of any function or object in the object's translation unit.

## Microsoft Specific

In Microsoft C++, global static objects are initialized before entry to **main**.

## END Microsoft Specific

Global static objects that are mutually interdependent but in different translation units may cause incorrect behavior.

## See Also

[Startup and Termination](#)

# Additional Termination Considerations

12/21/2017 • 1 min to read • [Edit Online](#)

You can terminate a C++ program by using `exit`, `return`, or `abort`. You can add exit processing using the `atexit` function. These are discussed in the following sections.

## See Also

[Startup and Termination](#)

# Using exit or return

12/21/2017 • 1 min to read • [Edit Online](#)

When you call **exit** or execute a `return` statement from **main**, static objects are destroyed in the reverse order of their initialization. The following example shows how such initialization and cleanup works.

## Example

```
// using_exit_or_return1.cpp
#include <stdio.h>
class ShowData {
public:
    // Constructor opens a file.
    ShowData( const char *szDev ) {
        errno_t err;
        err = fopen_s(&OutputDev, szDev, "w" );
    }

    // Destructor closes the file.
    ~ShowData() { fclose( OutputDev ); }

    // Disp function shows a string on the output device.
    void Disp( char *szData ) {
        fputs( szData, OutputDev );
    }
private:
    FILE *OutputDev;
};

// Define a static object of type ShowData. The output device
// selected is "CON" -- the standard output device.
ShowData sd1 = "CON";

// Define another static object of type ShowData. The output
// is directed to a file called "HELLO.DAT"
ShowData sd2 = "hello.dat";

int main() {
    sd1.Disp( "hello to default device\n" );
    sd2.Disp( "hello to file hello.dat\n" );
}
```

In the preceding example, the static objects `sd1` and `sd2` are created and initialized before entry to `main`. After this program terminates using the `return` statement, first `sd2` is destroyed and then `sd1`. The destructor for the `ShowData` class closes the files associated with these static objects.

Another way to write this code is to declare the `ShowData` objects with block scope, allowing them to be destroyed when they go out of scope:

```
int main() {
    ShowData sd1, sd2( "hello.dat" );

    sd1.Disp( "hello to default device\n" );
    sd2.Disp( "hello to file hello.dat\n" );
}
```

## See Also

[Additional Termination Considerations](#)

# Using atexit

12/21/2017 • 1 min to read • [Edit Online](#)

With the `atexit` function, you can specify an exit-processing function that executes prior to program termination. No global static objects initialized prior to the call to `atexit` are destroyed prior to execution of the exit-processing function.

## See Also

[Additional Termination Considerations](#)

# Using abort

12/21/2017 • 1 min to read • [Edit Online](#)

Calling the `abort` function causes immediate termination. It bypasses the normal destruction process for initialized global static objects. It also bypasses any special processing that was specified using the `atexit` function.

## See Also

[Additional Termination Considerations](#)

# Lvalues and Rvalues (Visual C++)

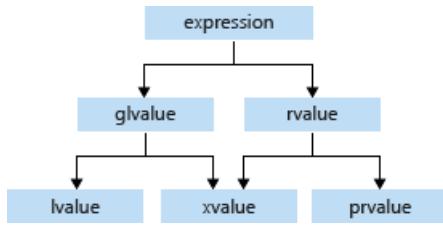
12/21/2017 • 2 min to read • [Edit Online](#)

Every C++ expression has a type, and belongs to a *value category*. The value categories are the basis for rules that compilers must follow when creating, copying, and moving temporary objects during expression evaluation.

The C++17 standard defines expression value categories as follows:

- A *glvalue* is an expression whose evaluation determines the identity of an object, bit-field, or function.
- A *prvalue* is an expression whose evaluation initializes an object or a bit-field, or computes the value of the operand of an operator, as specified by the context in which it appears.
- An *xvalue* is a glvalue that denotes an object or bit-field whose resources can be reused (usually because it is near the end of its lifetime). [ Example: Certain kinds of expressions involving rvalue references (8.3.2) yield xvalues, such as a call to a function whose return type is an rvalue reference or a cast to an rvalue reference type. ]
- An *lvalue* is a glvalue that is not an xvalue.
- An *rvalue* is a prvalue or an xvalue.

The following diagram illustrates the relationships between the categories:



An *lvalue* has an address that your program can access. Examples of *lvalue* expressions include variable names, including `const` variables, array elements, function calls that return an *lvalue* reference, bit-fields, unions, and class members.

A *prvalue* expression has no address that is accessible by your program. Examples of *prvalue* expressions include literals, function calls that return a non-reference type, and temporary objects that are created during expression evalution but accessible only by the compiler.

An *xvalue* expression has no address but can be used to initialize an *rvalue* reference, which provides access to the expression. Examples include function calls that return an *rvalue* reference, and the array subscript, member and pointer to member expressions where the array or object is an *rvalue* reference.

The following example demonstrates several correct and incorrect usages of *lvalues* and *rvalues*:

```
// lvalues_and_rvalues2.cpp
int main()
{
    int i, j, *p;

    // Correct usage: the variable i is an lvalue and the literal 7 is a prvalue.
    i = 7;

    // Incorrect usage: The left operand must be an lvalue (C2106). `j * 4` is a prvalue.
    7 = i; // C2106
    j * 4 = 7; // C2106

    // Correct usage: the dereferenced pointer is an lvalue.
    *p = i;

    const int ci = 7;
    // Incorrect usage: the variable is a non-modifiable lvalue (C3892).
    ci = 9; // C3892

    // Correct usage: the conditional operator returns an lvalue.
    ((i < 3) ? i : j) = 7;
}
```

#### NOTE

The examples in this topic illustrate correct and incorrect usage when operators are not overloaded. By overloading operators, you can make an expression such as `j * 4` an lvalue.

The terms *lvalue* and *rvalue* are often used when you refer to object references. For more information about references, see [Lvalue Reference Declarator: &](#) and [Rvalue Reference Declarator: &&](#).

## See Also

[Basic Concepts](#)

[Lvalue Reference Declarator: &](#)

[Rvalue Reference Declarator: &&](#)

# Temporary Objects

12/21/2017 • 2 min to read • [Edit Online](#)

In some cases, it is necessary for the compiler to create temporary objects. These temporary objects can be created for the following reasons:

- To initialize a `const` reference with an initializer of a type different from that of the underlying type of the reference being initialized.
- To store the return value of a function that returns a user-defined type. These temporaries are created only if your program does not copy the return value to an object. For example:

```
UDT Func1();    // Declare a function that returns a user-defined
                 // type.

...

Func1();         // Call Func1, but discard return value.
                 // A temporary object is created to store the return
                 // value.
```

Because the return value is not copied to another object, a temporary object is created. A more common case where temporaries are created is during the evaluation of an expression where overloaded operator functions must be called. These overloaded operator functions return a user-defined type that often is not copied to another object.

Consider the expression `ComplexResult = Complex1 + Complex2 + Complex3`. The expression `Complex1 + Complex2` is evaluated, and the result is stored in a temporary object. Next, the expression `temporary + Complex3` is evaluated, and the result is copied to `ComplexResult` (assuming the assignment operator is not overloaded).

- To store the result of a cast to a user-defined type. When an object of a given type is explicitly converted to a user-defined type, that new object is constructed as a temporary object.

Temporary objects have a lifetime that is defined by their point of creation and the point at which they are destroyed. Any expression that creates more than one temporary object eventually destroys them in the reverse order in which they were created. The points at which destruction occurs are shown in the following table.

## Destruction Points for Temporary Objects

REASON TEMPORARY CREATED	DESTRUCTION POINT
Result of expression evaluation	All temporaries created as a result of expression evaluation are destroyed at the end of the expression statement (that is, at the semicolon), or at the end of the controlling expressions for <code>for</code> , <code>if</code> , <code>while</code> , <code>do</code> , and <code>switch</code> statements.
Initializing <code>const</code> references	If an initializer is not an l-value of the same type as the reference being initialized, a temporary of the underlying object type is created and initialized with the initialization expression. This temporary object is destroyed immediately after the reference object to which it is bound is destroyed.

# Alignment (C++ Declarations)

12/21/2017 • 3 min to read • [Edit Online](#)

One of the low-level features of C++ is the ability to specify the precise alignment of objects in memory to take maximum advantage of a specific hardware architecture. By default, the compiler aligns class and struct members on their size value: bool and char are aligned one one byte boundaries, short on two byte, int on four bytes, long long, double and long double on eight bytes. In most scenarios you never have to be concerned with alignment because the default alignment is already optimal. In some cases however, you can achieve significant performance improvements, or memory savings, by specifying a custom alignment for your data structures. Prior to Visual Studio 2015 you could use the Microsoft-specific keywords `_alignof` and `declspec(alignas)` to specify an alignment greater than the default. Starting in Visual Studio 2015 you should use the C++11 standard keywords [alignof](#) and [alignas](#) for maximum code portability. The new keywords behave in the same way under the hood as the Microsoft-specific extensions, and the documentation for those extensions also applies to the new keywords. See [\\_alignof Operator](#) and [align](#) for more information. The C++ standard does not specify packing behavior for aligning on boundaries smaller than the compiler default for the target platform, so you still need to use the Microsoft #pragma `pack` in that case.

The C++ standard library provides the [aligned\\_storage Class](#) for allocating memory for data structures with custom alignments, and the [aligned\\_union Class](#) for specifying alignment for unions with non-trivial constructors or destructors.

## About Alignment

Alignment is a property of a memory address, expressed as the numeric address modulo a power of 2. For example, the address 0x0001103F modulo 4 is 3; that address is said to be aligned to  $4n+3$ , where 4 indicates the chosen power of 2. The alignment of an address depends on the chosen power of two. The same address modulo 8 is 7. An address is said to be aligned to X if its alignment is  $Xn+0$ .

CPUs execute instructions that operate on data stored in memory, and the data are identified by their addresses in memory. In addition to its address, a single datum also has a size. A datum is called naturally aligned if its address is aligned to its size, and misaligned otherwise. For example, an 8-byte floating-point datum is naturally aligned if the address used to identify it is aligned to 8.

Compiler handling of data alignmentDevice compilers attempt to allocate data in a way that prevents data misalignment.

For simple data types, the compiler assigns addresses that are multiples of the size in bytes of the data type. Thus, the compiler assigns addresses to variables of type long that are multiples of four, setting the bottom two bits of the address to zero.

In addition, the compiler pads structures in a way that naturally aligns each element of the structure. Consider the structure struct `x_` in the following code example:

```
struct x_
{
    char a;      // 1 byte
    int b;       // 4 bytes
    short c;    // 2 bytes
    char d;      // 1 byte
} MyStruct;
```

The compiler pads this structure to enforce alignment naturally.

The following code example shows how the compiler places the padded structure in memory:Copy

```
// Shows the actual memory layout
struct x_
{
    char a;           // 1 byte
    char _pad0[3];   // padding to put 'b' on 4-byte boundary
    int b;           // 4 bytes
    short c;         // 2 bytes
    char d;          // 1 byte
    char _pad1[1];   // padding to make sizeof(x_) multiple of 4
}
```

1. Both declarations return sizeof(struct x\_) as 12 bytes.
2. The second declaration includes two padding elements:
  3. char \_pad0[3] to align the int b member on a four-byte boundary
  4. char \_pad1[1] to align the array elements of the structure struct \_x bar[3];
5. The padding aligns the elements of bar[3] in a way that allows natural access.

The following code example shows the bar[3] array layout:

adr	offset	element
0x0000		---
0x0000	char a;	// bar[0]
0x0001	char _pad0[3];	
0x0004	int b;	
0x0008	short c;	
0x000a	char d;	
0x000b	char _pad1[1];	
0x000c		---
0x000c	char a;	// bar[1]
0x000d	char _pad0[3];	
0x0010	int b;	
0x0014	short c;	
0x0016	char d;	
0x0017	char _pad1[1];	
0x0018		---
0x0018	char a;	// bar[2]
0x0019	char _pad0[3];	
0x001c	int b;	
0x0020	short c;	
0x0022	char d;	
0x0023	char _pad1[1];	

## See Also

[Data Structure Alignment](#)

# alignof and alignas (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The `alignas` type specifier is a portable, C++ standard way to specify custom alignment of variables and user defined types. The `alignof` operator is likewise a standard, portable way to obtain the alignment of a specified type or variable.

## Example

You can use `alignas` on a class, struct or union, or on individual members. When multiple `alignas` specifiers are encountered, the compiler will choose the strictest one, (the one with the largest value).

```
// alignas_alignof.cpp
// compile with: cl /EHsc alignas_alignof.cpp
#include <iostream>

struct alignas(16) Bar
{
    int i;          // 4 bytes
    int n;          // 4 bytes
    alignas(4) char arr[3];
    short s;        // 2 bytes
};

int main()
{
    std::cout << alignof(Bar) << std::endl; // output: 16
}
```

## See Also

[Alignment](#)

# Trivial, standard-layout, and POD types

8/31/2017 • 5 min to read • [Edit Online](#)

The term *layout* refers to how the members of an object of class, struct or union type are arranged in memory. In some cases, the layout is well-defined by the language specification. But when a class or struct contains certain C++ language features such as virtual base classes, virtual functions, members with different access control, then the compiler is free to choose a layout. That layout may vary depending on what optimizations are being performed and in many cases object might not even occupy a contiguous area of memory. For example, if a class has virtual functions, all the instances of that class might share a single virtual function table. Such types are of course very useful, but they also have limitations. Because the layout is undefined they cannot be passed to programs written in other languages, such as C, and because they might be non-contiguous they cannot be reliably copied with fast low-level functions such as `memcpy` or serialized over a network.

To enable compilers as well as C++ programs and metaprograms to reason about the suitability of any given type for operations that depend on a particular memory layout, C++14 introduced three categories of simple classes and structs: *trivial*, *standard-layout*, and *POD* or Plain Old Data. The Standard Library has the function templates `is_trivial<T>`, `is_standard_layout<T>` and `is_pod<T>` that determine whether a given type belongs to a given category.

## Trivial types

When a class or struct in C++ has compiler-provided or explicitly defaulted special member functions, then it is a trivial type. It occupies a contiguous memory area. It can have members with different access specifiers. In C++, the compiler is free to choose how to order members in this situation. Therefore, you can `memcpy` such objects but you cannot reliably consume them from a C program. A trivial type T can be copied into an array of `char` or `unsigned char`, and safely copied back into a T variable. Note that because of alignment requirements, there might be padding bytes between type members.

Trivial types have a trivial default constructor, trivial copy constructor, trivial copy assignment operator and trivial destructor. In each case, *trivial* means the constructor/operator/destructor is not user-provided and belongs to a class that has

- no virtual functions or virtual base classes,
- no base classes with a corresponding non-trivial constructor/operator/destructor
- no data members of class type with a corresponding non-trivial constructor/operator/destructor

The following examples show trivial types. In `Trivial2`, the presence of the `Trivial2(int a, int b)` constructor requires that you provide a default constructor. For the type to qualify as trivial, you must explicitly default that constructor.

```

struct Trivial
{
    int i;
private:
    int j;
};

struct Trivial2
{
    int i;
    Trivial2(int a, int b) : i(a), j(b) {}
    Trivial2() = default;
private:
    int j; // Different access control
};

```

## Standard layout types

When a class or struct does not contain certain C++ language features such as virtual functions which are not found in the C language, and all members have the same access control, it is a standard-layout type. It is memcpyable and the layout is sufficiently defined that it can be consumed by C programs. Standard-layout types can have user-defined special member functions. In addition, standard layout types have these characteristics:

- no virtual functions or virtual base classes
- all non-static data members have the same access control
- all non-static members of class type are standard-layout
- any base classes are standard-layout
- has no base classes of the same type as the first non-static data member.
- meets one of these conditions:
  - no non-static data member in the most-derived class and no more than one base class with non-static data members, or
  - has no base classes with non-static data members

The following code shows one example of a standard-layout type:

```

struct SL
{
    // All members have same access:
    int i;
    int j;
    SL(int a, int b) : i(a), j(b) {} // User-defined constructor OK
};

```

The last two requirements can perhaps be better illustrated with code. In the next example, even though `Base` is standard-layout, `Derived` is not standard layout because both it (the most derived class) and `Base` have non-static data members:

```
struct Base
{
    int i;
    int j;
};

// std::is_standard_layout<<Derived> == false!
struct Derived : public Base
{
    int x;
    int y;
};
```

In this example `Derived` is standard-layout because `Base` has no non-static data members:

```
struct Base
{
    void Foo() {}
};

// std::is_standard_layout<<Derived> == true
struct Derived : public Base
{
    int x;
    int y;
};
```

`Derived` would also be standard-layout if `Base` had the data members and `Derived` had only member functions.

## POD types

When a class or struct is both trivial and standard-layout, it is a POD (Plain Old Data) type. The memory layout of POD types is therefore contiguous and each member has a higher address than the member that was declared before it, so that byte for byte copies and binary I/O can be performed on these types. Scalar types such as `int` are also POD types. POD types that are classes can have only POD types as non-static data members.

## Example

The following example shows the distinctions between trivial, standard-layout, and POD types:

```

#include <type_traits>
#include <iostream>

using namespace std;

struct B
{
protected:
    virtual void Foo() {}
};

// Neither trivial nor standard-layout
struct A : B
{
    int a;
    int b;
    void Foo() override {} // Virtual function
};

// Trivial but not standard-layout
struct C
{
    int a;
private:
    int b; // Different access control
};

// Standard-layout but not trivial
struct D
{
    int a;
    int b;
    D() {} //User-defined constructor
};

struct POD
{
    int a;
    int b;
};

int main()
{
    cout << boolalpha;
    cout << "A is trivial is " << is_trivial<A>() << endl; // false
    cout << "A is standard-layout is " << is_standard_layout<A>() << endl; // false

    cout << "C is trivial is " << is_trivial<C>() << endl; // true
    cout << "C is standard-layout is " << is_standard_layout<C>() << endl; // false

    cout << "D is trivial is " << is_trivial<D>() << endl; // false
    cout << "D is standard-layout is " << is_standard_layout<D>() << endl; // true

    cout << "POD is trivial is " << is_trivial<POD>() << endl; // true
    cout << "POD is standard-layout is " << is_standard_layout<POD>() << endl; // true

    return 0;
}

```

## See Also

[Basic Concepts](#)

# Fundamental Types (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Fundamental types in C++ are divided into three categories: integral, floating point, and void. Integral types are capable of handling whole numbers. Floating point types are capable of specifying values that may have fractional parts.

The `void` type describes an empty set of values. No variable of type `void` can be specified — it is used primarily to declare functions that return no values or to declare generic pointers to untyped or arbitrarily typed data. Any expression can be explicitly converted or cast to type `void`. However, such expressions are restricted to the following uses:

- An expression statement. (See [Expressions](#), for more information.)
- The left operand of the comma operator. (See [Comma Operator](#) for more information.)
- The second or third operand of the conditional operator (`? :`). (See [Expressions with the Conditional Operator](#) for more information.)

The following table explains the restrictions on type sizes. These restrictions are independent of the Microsoft implementation.

## Fundamental Types of the C++ Language

CATEGORY	TYPE	CONTENTS
Integral	<code>char</code>	Type <code>char</code> is an integral type that usually contains members of the basic execution character set — By default, this is ASCII in Microsoft C++.  The C++ compiler treats variables of type <code>char</code> , <code>signed char</code> , and <code>unsigned char</code> as having different types. Variables of type <code>char</code> are promoted to <code>int</code> as if they are type <code>signed char</code> by default, unless the /J compilation option is used. In this case they are treated as type <code>unsigned char</code> and are promoted to <code>int</code> without sign extension.
	<code>bool</code>	Type <code>bool</code> is an integral type that can have one of the two values <code>true</code> or <code>false</code> . Its size is unspecified.

CATEGORY	TYPE	CONTENTS
	<code>short</code>	<p>Type <code>short</code> (<code>int</code> or simply <code>short</code>) is an integral type that is larger than or equal to the size of type <code>char</code>, and shorter than or equal to the size of type <code>int</code>.</p> <p>Objects of type <code>short</code> can be declared as <code>signed short</code> or <code>unsigned short</code>. <code>Signed short</code> is a synonym for <code>short</code>.</p>
	<code>int</code>	<p>Type <code>int</code> is an integral type that is larger than or equal to the size of type <code>short</code> <code>int</code>, and shorter than or equal to the size of type <code>long</code>.</p> <p>Objects of type <code>int</code> can be declared as <code>signed int</code> or <code>unsigned int</code>. <code>Signed int</code> is a synonym for <code>int</code>.</p>
	<code>__int8</code> , <code>__int16</code> , <code>__int32</code> , <code>__int64</code>	<p>Sized integer <code>__int n</code>, where <code>n</code> is the size, in bits, of the integer variable. <code>__int8</code>, <code>__int16</code>, <code>__int32</code> and <code>__int64</code> are Microsoft-specific keywords. Not all types are available on all architectures. (<code>__int128</code> is not supported.)</p>
	<code>long</code>	<p>Type <code>long</code> (or <code>long int</code>) is an integral type that is larger than or equal to the size of type <code>int</code>.</p> <p>Objects of type <code>long</code> can be declared as <code>signed long</code> or <code>unsigned long</code>. <code>Signed long</code> is a synonym for <code>long</code>.</p>
	<code>long long</code>	<p>Larger than an unsigned <code>long</code>.</p> <p>Objects of type <code>long long</code> can be declared as <code>signed long long</code> or <code>unsigned long long</code>. <code>Signed long long</code> is a synonym for <code>long long</code>.</p>

CATEGORY	TYPE	CONTENTS
	<code>wchar_t</code> , <code>__wchar_t</code>	A variable of type <code>wchar_t</code> designates a wide-character or multibyte character type. By default, <code>wchar_t</code> is a native type, but you can use <code>/Zc:wchar_t-</code> to make <code>wchar_t</code> a typedef for <code>unsigned short</code> . The <code>__wchar_t</code> type is a Microsoft-specific synonym for the native <code>wchar_t</code> type.  Use the L prefix before a character or string literal to designate the wide-character type.
Floating point	<code>float</code>	Type <code>float</code> is the smallest floating point type.
	<code>double</code>	Type <code>double</code> is a floating point type that is larger than or equal to type <code>float</code> , but shorter than or equal to the size of type <code>long double</code> .  Microsoft specific: The representation of <code>long double</code> and <code>double</code> is identical. However, <code>long double</code> and <code>double</code> are separate types.
	<code>long double</code>	Type <code>long double</code> is a floating point type that is larger than or equal to type <code>double</code> .

## Microsoft Specific

The following table lists the amount of storage required for fundamental types in Microsoft C++.

### Sizes of Fundamental Types

TYPE	SIZE
<code>bool</code> , <code>char</code> , <code>unsigned char</code> , <code>signed char</code> , <code>__int8</code>	1 byte
<code>__int16</code> , <code>short</code> , <code>unsigned short</code> , <code>wchar_t</code> , <code>__wchar_t</code>	2 bytes
<code>float</code> , <code>__int32</code> , <code>int</code> , <code>unsigned int</code> , <code>long</code> , <code>unsigned long</code>	4 bytes
<code>double</code> , <code>__int64</code> , <code>long double</code> , <code>long long</code>	8 bytes

## END Microsoft Specific

See [Data Type Ranges](#) for a summary of the range of values of each type.

For more information about type conversion, see [Standard Conversions](#).

## See Also

[Data Type Ranges](#)

# Data Type Ranges

12/21/2017 • 2 min to read • [Edit Online](#)

Visual C++ 32-bit and 64-bit compilers recognize the types in the table later in this article.

- `int` (`unsigned int`)
- `_int8` (`unsigned __int8`)
- `_int16` (`unsigned __int16`)
- `_int32` (`unsigned __int32`)
- `_int64` (`unsigned __int64`)
- `short` (`unsigned short`)
- `long` (`unsigned long`)
- `long long` (`unsigned long long`)

If its name begins with two underscores (`__`), a data type is non-standard.

The ranges that are specified in the following table are inclusive-inclusive.

TYPE NAME	BYTES	OTHER NAMES	RANGE OF VALUES
int	4	signed	-2,147,483,648 to 2,147,483,647
unsigned int	4	unsigned	0 to 4,294,967,295
_int8	1	char	-128 to 127
unsigned __int8	1	unsigned char	0 to 255
_int16	2	short, short int, signed short int	-32,768 to 32,767
unsigned __int16	2	unsigned short, unsigned short int	0 to 65,535
_int32	4	signed, signed int, int	-2,147,483,648 to 2,147,483,647
unsigned __int32	4	unsigned, unsigned int	0 to 4,294,967,295
_int64	8	long long, signed long long	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
unsigned __int64	8	unsigned long long	0 to 18,446,744,073,709,551,615

TYPE NAME	BYTES	OTHER NAMES	RANGE OF VALUES
bool	1	none	false or true
char	1	none	-128 to 127 by default 0 to 255 when compiled by using <a href="#">/_U</a>
signed char	1	none	-128 to 127
unsigned char	1	none	0 to 255
short	2	short int, signed short int	-32,768 to 32,767
unsigned short	2	unsigned short int	0 to 65,535
long	4	long int, signed long int	-2,147,483,648 to 2,147,483,647
unsigned long	4	unsigned long int	0 to 4,294,967,295
long long	8	none (but equivalent to __int64)	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
unsigned long long	8	none (but equivalent to unsigned __int64)	0 to 18,446,744,073,709,551,615
enum	varies	none	
float	4	none	3.4E +/- 38 (7 digits)
double	8	none	1.7E +/- 308 (15 digits)
long double	same as double	none	Same as double
wchar_t	2	__wchar_t	0 to 65,535

Depending on how it's used, a variable of `__wchar_t` designates either a wide-character type or multibyte-character type. Use the `L` prefix before a character or string constant to designate the wide-character-type constant.

`signed` and `unsigned` are modifiers that you can use with any integral type except `bool`. Note that `char`, `signed char`, and `unsigned char` are three distinct types for the purposes of mechanisms like overloading and templates.

The `int` and `unsigned int` types have a size of four bytes. However, portable code should not depend on the size of `int` because the language standard allows this to be implementation-specific.

C/C++ in Visual Studio also supports sized integer types. For more information, see [\\_\\_int8](#), [\\_\\_int16](#), [\\_\\_int32](#), [\\_\\_int64](#) and [Integer Limits](#).

For more information about the restrictions of the sizes of each type, see [Fundamental Types](#).

The range of enumerated types varies depending on the language context and specified compiler flags. For more information, see [C Enumeration Declarations](#) and [Enumerations](#).

## See Also

[Keywords](#)

[Fundamental Types](#)

# nullptr

12/21/2017 • 1 min to read • [Edit Online](#)

Designates a null pointer constant of type `std::nullptr_t`, which is convertible to any raw pointer type. Although you can use the keyword `nullptr` without including any headers, if your code uses the type `std::nullptr_t`, then you must define it by including the header `<cstddef>`.

## NOTE

The `nullptr` keyword is also defined in C++/CLI for managed code applications and is not interchangeable with the ISO Standard C++ keyword. If your code might be compiled by using the `/clr` compiler option, which targets managed code, then use `_nullptr` in any line of code where you must guarantee that the compiler uses the native C++ interpretation. For more information, see [nullptr](#).

## Remarks

Avoid using `NULL` or zero (`0`) as a null pointer constant; `nullptr` is less vulnerable to misuse and works better in most situations. For example, given `func(std::pair<const char *, double>)`, then calling

`func(std::make_pair(NULL, 3.14))` causes a compiler error. The macro `NULL` expands to `0`, so that the call `std::make_pair(0, 3.14)` returns `std::pair<int, double>`, which is not convertible to `func()`'s `std::pair<const char *, double>` parameter type. Calling `func(std::make_pair(nullptr, 3.14))` successfully compiles because `std::make_pair(nullptr, 3.14)` returns `std::pair<std::nullptr_t, double>`, which is convertible to `std::pair<const char *, double>`.

## See Also

[Keywords](#)

[nullptr](#)

# void (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

When used as a function return type, the `void` keyword specifies that the function does not return a value. When used for a function's parameter list, `void` specifies that the function takes no parameters. When used in the declaration of a pointer, `void` specifies that the pointer is "universal."

If a pointer's type is `void`, the pointer can point to any variable that is not declared with the `const` or `volatile` keyword. A void pointer cannot be dereferenced unless it is cast to another type. A void pointer can be converted into any other type of data pointer.

A void pointer can point to a function, but not to a class member in C++.

You cannot declare a variable of type `void`.

## Example

```
// void.cpp
void vobject;    // C2182
void *pv;        // okay
int *pint; int i;
int main() {
    pv = &i;
    // Cast optional in C required in C++
    pint = (int *)pv;
}
```

## See Also

[Keywords](#)

[Fundamental Types](#)

# bool (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This keyword is a built-in type. A variable of this type can have values **true** and **false**. Conditional expressions have the type `bool` and so have values of type `bool`. For example, `i!=0` now has **true** or **false** depending on the value of `i`.

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): The operand of a postfix or prefix increment or decrement operator may not be of type `bool`.

The values **true** and **false** have the following relationship:

```
!false == true  
!true == false
```

In the following statement:

```
if (condexpr1) statement1;
```

If `condexpr1` is **true**, `statement1` is always executed; if `condexpr1` is **false**, `statement1` is never executed.

When a postfix or prefix `++` operator is applied to a variable of type `bool`, the variable is set to **true**. **Visual Studio 2017 version 15.3 and later**: operator`++` for `bool` was removed from the language and is no longer supported.

The postfix or prefix `--` operator cannot be applied to a variable of this type.

The `bool` type participates in integral promotions. An r-value of type `bool` can be converted to an r-value of type `int`, with **false** becoming zero and **true** becoming one. As a distinct type, `bool` participates in overload resolution.

## See Also

[Keywords](#)

[Fundamental Types](#)

# false (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The keyword is one of the two values for a variable of type `bool` or a conditional expression (a conditional expression is now a `true` Boolean expression). For example, if `i` is a variable of type `bool`, the `i = false;` statement assigns `false` to `i`.

## Example

```
// bool_false.cpp
#include <stdio.h>

int main()
{
    bool bb = true;
    printf_s("%d\n", bb);
    bb = false;
    printf_s("%d\n", bb);
}
```

```
1
0
```

## See Also

[Keywords](#)

# true (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
bool-identifier = true ;
bool-expression logical-operator true ;
```

## Remarks

This keyword is one of the two values for a variable of type `bool` or a conditional expression (a conditional expression is now a true boolean expression). If `i` is of type `bool`, then the statement `i = true;` assigns `true` to `i`.

## Example

```
// bool_true.cpp
#include <stdio.h>
int main()
{
    bool bb = true;
    printf_s("%d\n", bb);
    bb = false;
    printf_s("%d\n", bb);
}
```

```
1
0
```

## See Also

[Keywords](#)

# char, wchar\_t, char16\_t, char32\_t

12/21/2017 • 1 min to read • [Edit Online](#)

The types `char`, `wchar_t`, `char16_t` and `char32_t` are built-in types that represent alphanumeric characters as well as non-alphanumeric glyphs and non-printing characters. `char` is eight bits in size, `wchar_t` and `char16_t` are 16 bits in size, and `char32_t` is 32 bits.

## Syntax

```
char      ch1{ 'a' };
wchar_t   ch2{ 'a' }; // or {L'a'}
char16_t  ch3{ L'a' };// or {L'a'}
char32_t  ch4{ L'a' };// or {L'a'}
```

## Remarks

The `char` type was the original character type in C and C++. It can be used to store characters from the ASCII character set or any of the ISO-8859 character sets, or the UTF-8 character set. The type `unsigned char` is often used to represent a *byte* which is not a built-in type in C++. The `char` type is not suitable for text in many languages. In general, modern programs should use one of the wide character types to represent text. Unicode is the

In the C++ standard library, the `basic_string` type is specialized for both narrow and wide strings. Use `std::string` when the characters are of type `char`, and `std::wstring` when the characters are of type `wchar_t`. Other types that represent text, including `std::stringstream` and `std::cout` have specializations for narrow and wide strings.

# `_int8, _int16, _int32, _int64`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Microsoft C/C++ features support for sized integer types. You can declare 8-, 16-, 32-, or 64-bit integer variables by using the `_intn` type specifier, where *n* is 8, 16, 32, or 64.

The following example declares one variable for each of these types of sized integers:

```
_int8 nSmall;      // Declares 8-bit integer
_int16 nMedium;    // Declares 16-bit integer
_int32 nLarge;     // Declares 32-bit integer
_int64 nHuge;      // Declares 64-bit integer
```

The types `_int8`, `_int16`, and `_int32` are synonyms for the ANSI types that have the same size, and are useful for writing portable code that behaves identically across multiple platforms. The `_int8` data type is synonymous with type `char`, `_int16` is synonymous with type `short`, and `_int32` is synonymous with type `int`. The `_int64` type is synonymous with type `long long`.

## Example

The following sample shows that an `_intxx` parameter will be promoted to `int`:

```
// sized_int_types.cpp

#include <stdio.h>

void func(int i) {
    printf_s("%s\n", __FUNCTION__);
}

int main()
{
    _int8 i8 = 100;
    func(i8);    // no void func(_int8 i8) function
                 // _int8 will be promoted to int
}
```

```
func
```

## END Microsoft Specific

## See Also

[Keywords](#)

[Fundamental Types](#)

[Data Type Ranges](#)

# `_m64`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_m64` data type is for use with the MMX and 3DNow! intrinsics, and is defined in `xmmmintrin.h`.

```
// data_types_m64.cpp
#include <xmmmintrin.h>
int main()
{
    _m64 x;
}
```

## Remarks

You should not access the `_m64` fields directly. You can, however, see these types in the debugger. A variable of type `_m64` maps to the MM[0-7] registers.

Variables of type `_m64` are automatically aligned on 8-byte boundaries.

The `_m64` data type is not supported on x64 processors. Applications that use `_m64` as part of MMX intrinsics must be rewritten to use equivalent SSE and SSE2 intrinsics.

## END Microsoft Specific

## See Also

[Keywords](#)

[Fundamental Types](#)

[Data Type Ranges](#)

# `_m128`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_m128` data type, for use with the Streaming SIMD Extensions and Streaming SIMD Extensions 2 instructions intrinsics, is defined in `xmmmintrin.h`.

```
// data_types_m128.cpp
#include <xmmmintrin.h>
int main() {
    _m128 x;
}
```

## Remarks

You should not access the `_m128` fields directly. You can, however, see these types in the debugger. A variable of type `_m128` maps to the XMM[0-7] registers.

Variables of type `_m128` are automatically aligned on 16-byte boundaries.

The `_m128` data type is not supported on ARM processors.

## END Microsoft Specific

## See Also

[Keywords](#)

[Fundamental Types](#)

[Data Type Ranges](#)

# `_m128d`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_m128d` data type, for use with the Streaming SIMD Extensions 2 instructions intrinsics, is defined in `emmintrin.h`.

```
// data_types_m128d.cpp
#include <emmintrin.h>
int main() {
    _m128d x;
}
```

## Remarks

You should not access the `_m128d` fields directly. You can, however, see these types in the debugger. A variable of type `_m128` maps to the XMM[0-7] registers.

Variables of type `_m128d` are automatically aligned on 16-byte boundaries.

The `_m128d` data type is not supported on ARM processors.

## END Microsoft Specific

## See Also

[Keywords](#)

[Fundamental Types](#)

[Data Type Ranges](#)

# `_m128i`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_m128i` data type, for use with the Streaming SIMD Extensions 2 (SSE2) instructions intrinsics, is defined in `emmintrin.h`.

```
// data_types_m128i.cpp
#include <emmintrin.h>
int main() {
    _m128i x;
}
```

## Remarks

You should not access the `_m128i` fields directly. You can, however, see these types in the debugger. A variable of type `_m128i` maps to the XMM[0-7] registers.

Variables of type `_m128i` are automatically aligned on 16-byte boundaries.

### NOTE

Using variables of type `_m128i` will cause the compiler to generate the SSE2 **movdqa** instruction. This instruction does not cause a fault on Pentium III processors but will result in silent failure, with possible side effects caused by whatever instructions **movdqa** translates into on Pentium III processors.

The `_m128i` data type is not supported on ARM processors.

## END Microsoft Specific

## See Also

[Keywords](#)

[Fundamental Types](#)

[Data Type Ranges](#)

# `_ptr32`, `_ptr64`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

`_ptr32` represents a native pointer on a 32-bit system, while `_ptr64` represents a native pointer on a 64-bit system.

The following example shows how to declare each of these pointer types:

```
int * __ptr32 p32;
int * __ptr64 p64;
```

On a 32-bit system, a pointer declared with `_ptr64` is truncated to a 32-bit pointer. On a 64-bit system, a pointer declared with `_ptr32` is coerced to a 64-bit pointer.

### NOTE

You cannot use `_ptr32` or `_ptr64` when compiling with `/clr:pure`. Otherwise, `Compiler Error C2472` will be generated. The `/clr:pure` and `/clr:safe` compiler options are deprecated in Visual Studio 2015.

## Example

The following example shows how to declare and allocate pointers with the `_ptr32` and `_ptr64` keywords.

```
#include <cstdlib>
#include <iostream>

int main()
{
    using namespace std;

    int * __ptr32 p32;
    int * __ptr64 p64;

    p32 = (int * __ptr32)malloc(4);
    *p32 = 32;
    cout << *p32 << endl;

    p64 = (int * __ptr64)malloc(4);
    *p64 = 64;
    cout << *p64 << endl;
}
```

32  
64

## END Microsoft Specific

## See Also

[Fundamental Types](#)

# Numerical Limits (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The two standard include files, LIMITS.H and FLOAT.H, define the numerical limits, or minimum and maximum values that a variable of a given type can hold. These minimums and maximums are guaranteed to be portable to any C++ compiler that uses the same data representation as ANSI C. The LIMITS.H include file defines the [numerical limits for integral types](#), and FLOAT.H defines the [numerical limits for floating types](#).

## See Also

[Basic Concepts](#)

# Integer Limits

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The limits for integer types are listed in the following table. These limits are also defined in the standard header file LIMITS.H.

### Limits on Integer Constants

CONSTANT	MEANING	VALUE
<b>CHAR_BIT</b>	Number of bits in the smallest variable that is not a bit field.	8
<b>SCHAR_MIN</b>	Minimum value for a variable of type <b>signed char</b> .	-128
<b>SCHAR_MAX</b>	Maximum value for a variable of type <b>signed char</b> .	127
<b>UCHAR_MAX</b>	Maximum value for a variable of type <b>unsigned char</b> .	255 (0xff)
<b>CHAR_MIN</b>	Minimum value for a variable of type <b>char</b> .	-128; 0 if /J option used
<b>CHAR_MAX</b>	Maximum value for a variable of type <b>char</b> .	127; 255 if /J option used
<b>MB_LEN_MAX</b>	Maximum number of bytes in a multicharacter constant.	5
<b>SHRT_MIN</b>	Minimum value for a variable of type <b>short</b> .	-32768
<b>SHRT_MAX</b>	Maximum value for a variable of type <b>short</b> .	32767
<b>USHRT_MAX</b>	Maximum value for a variable of type <b>unsigned short</b> .	65535 (0xffff)
<b>INT_MIN</b>	Minimum value for a variable of type <b>int</b> .	-2147483648
<b>INT_MAX</b>	Maximum value for a variable of type <b>int</b> .	2147483647
<b>UINT_MAX</b>	Maximum value for a variable of type <b>unsigned int</b> .	4294967295 (0xffffffff)

CONSTANT	MEANING	VALUE
<b>LONG_MIN</b>	Minimum value for a variable of type <b>long</b> .	-2147483648
<b>LONG_MAX</b>	Maximum value for a variable of type <b>long</b> .	2147483647
<b>ULONG_MAX</b>	Maximum value for a variable of type <b>unsigned long</b> .	4294967295 (0xffffffff)
<b>_I64_MIN</b>	Minimum value for a variable of type <b>__int64</b>	-9223372036854775808
<b>_I64_MAX</b>	Maximum value for a variable of type <b>__int64</b>	9223372036854775807
<b>_UI64_MAX</b>	Maximum value for a variable of type <b>unsigned __int64</b>	18446744073709551615 (0xffffffffffffffffffff)

If a value exceeds the largest integer representation, the Microsoft compiler generates an error.

#### END Microsoft Specific

## See Also

[Floating Limits](#)

# Floating Limits

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The following table lists the limits on the values of floating-point constants. These limits are also defined in the standard header file FLOAT.H.

### Limits on Floating-Point Constants

CONSTANT	MEANING	VALUE
FLT_DIG DBL_DIG LDBL_DIG	Number of digits, q, such that a floating-point number with q decimal digits can be rounded into a floating-point representation and back without loss of precision.	6 15 15
FLT_EPSILON DBL_EPSILON LDBL_EPSILON	Smallest positive number x, such that x + 1.0 is not equal to 1.0.	1.192092896e-07F 2.2204460492503131e-016 2.2204460492503131e-016
FLT_GUARD		0
FLT_MANT_DIG DBL_MANT_DIG LDBL_MANT_DIG	Number of digits in the radix specified by FLT_RADIX in the floating-point significand. The radix is 2; hence these values specify bits.	24 53 53
FLT_MAX DBL_MAX LDBL_MAX	Maximum representable floating-point number.	3.402823466e+38F 1.7976931348623158e+308 1.7976931348623158e+308
FLT_MAX_10_EXP DBL_MAX_10_EXP LDBL_MAX_10_EXP	Maximum integer such that 10 raised to that number is a representable floating-point number.	38 308 308
FLT_MAX_EXP DBL_MAX_EXP LDBL_MAX_EXP	Maximum integer such that FLT_RADIX raised to that number is a representable floating-point number.	128 1024 1024
FLT_MIN DBL_MIN LDBL_MIN	Minimum positive value.	1.175494351e-38F 2.2250738585072014e-308 2.2250738585072014e-308
FLT_MIN_10_EXP DBL_MIN_10_EXP LDBL_MIN_10_EXP	Minimum negative integer such that 10 raised to that number is a representable floating-point number.	-37 -307 -307

CONSTANT	MEANING	VALUE
FLT_MIN_EXP DBL_MIN_EXP LDBL_MIN_EXP	Minimum negative integer such that FLT_RADIX raised to that number is a representable floating-point number.	-125 -1021 -1021
FLT_NORMALIZE		0
FLT_RADIX _DBL_RADIX _LDBL_RADIX	Radix of exponent representation.	2 2 2
FLT_ROUNDS _DBL_ROUNDS _LDBL_ROUNDS	Rounding mode for floating-point addition.	1 (near) 1 (near) 1 (near)

**NOTE**

The information in the table may differ in future versions of the product.

## END Microsoft Specific

## See Also

[Integer Limits](#)

# Standard Conversions

12/21/2017 • 12 min to read • [Edit Online](#)

The C++ language defines conversions between its fundamental types. It also defines conversions for pointer, reference, and pointer-to-member derived types. These conversions are called "standard conversions." (For more information about types, standard types, and derived types, see [Types](#).)

This section discusses the following standard conversions:

- Integral promotions
- Integral conversions
- Floating conversions
- Floating and integral conversions
- Arithmetic conversions
- Pointer conversions
- Reference conversions
- Pointer-to-member conversions

## NOTE

User-defined types can specify their own conversions. Conversion of user-defined types is covered in [Constructors](#) and [Conversions](#).

The following code causes conversions (in this example, integral promotions):

```
long long_num1, long_num2;
int int_num;

// int_num promoted to type long prior to assignment.
long_num1 = int_num;

// int_num promoted to type long prior to multiplication.
long_num2 = int_num * long_num1;
```

The result of a conversion is an l-value only if it produces a reference type. For example, a user-defined conversion declared as `operator int&()` returns a reference and is an l-value. However, a conversion declared as `operator int()` returns an object and is not an l-value.

## Integral promotions

Objects of an integral type can be converted to another wider integral type (that is, a type that can represent a larger set of values). This widening type of conversion is called "integral promotion." With integral promotion, you can use the following in an expression wherever another integral type can be used:

- Objects, literals, and constants of type `char` and `short int`
- Enumeration types

- `int` bit fields
- Enumerators

C++ promotions are "value-preserving." That is, the value after the promotion is guaranteed to be the same as the value before the promotion. In value-preserving promotions, objects of shorter integral types (such as bit fields or objects of type `char`) are promoted to type `int` if `int` can represent the full range of the original type. If `int` cannot represent the full range of values, then the object is promoted to type `unsigned int`. Although this strategy is the same as that used by ANSI C, value-preserving conversions do not preserve the "signedness" of the object.

Value-preserving promotions and promotions that preserve signedness normally produce the same results. However, they can produce different results if the promoted object is one of the following:

- An operand of `/`, `%`, `/=`, `%=`, `<`, `<=`, `>`, or `>=`

These operators rely on sign for determining the result. Therefore, value-preserving and sign-preserving promotions produce different results when applied to these operands.

- The left operand of `>>` or `>>=`

These operators treat signed and unsigned quantities differently when performing a shift operation. For signed quantities, shifting a quantity right causes the sign bit to be propagated into the vacated bit positions. For unsigned quantities, the vacated bit positions are zero-filled.

- An argument to an overloaded function or operand of an overloaded operator that depends on the signedness of the type of that operand for argument matching. (See [Overloaded Operators](#) for more about defining overloaded operators.)

## Integral conversions

Integral conversions are performed between integral types. The integral types are `char`, `int`, and `long` (and the `short`, `signed`, and `unsigned` versions of these types).

### Signed to unsigned

Objects of signed integral types can be converted to corresponding unsigned types. When these conversions occur, the actual bit pattern does not change; however, the interpretation of the data changes. Consider this code:

```
#include <iostream>

using namespace std;
int main()
{
    short i = -3;
    unsigned short u;

    cout << (u = i) << "\n";
}
// Output: 65533
```

In the preceding example, a `signed short`, `i`, is defined and initialized to a negative number. The expression `(u = i)` causes `i` to be converted to an `unsigned short` prior to the assignment to `u`.

### Unsigned to signed

Objects of unsigned integral types can be converted to corresponding signed types. However, such a conversion

can cause misinterpretation of data if the value of the unsigned object is outside the range representable by the signed type, as demonstrated in the following example:

```
#include <iostream>

using namespace std;
int main()
{
    short i;
    unsigned short u = 65533;

    cout << (i = u) << "\n";
}
//Output: -3
```

In the preceding example, `u` is an `unsigned short` integral object that must be converted to a signed quantity to evaluate the expression `(i = u)`. Because its value cannot be properly represented in a `signed short`, the data is misinterpreted as shown.

## Floating point conversions

An object of a floating type can be safely converted to a more precise floating type — that is, the conversion causes no loss of significance. For example, conversions from `float` to `double` or from `double` to `long double` are safe, and the value is unchanged.

An object of a floating type can also be converted to a less precise type, if it is in a range representable by that type. (See [Floating Limits](#) for the ranges of floating types.) If the original value cannot be represented precisely, it can be converted to either the next higher or the next lower representable value. If no such value exists, the result is undefined. Consider the following example:

```
cout << (float)1E300 << endl;
```

The maximum value representable by type `float` is 3.402823466E38 — a much smaller number than 1E300. Therefore, the number is converted to infinity, and the result is 1.#INF.

## Conversions between integral and floating point types

Certain expressions can cause objects of floating type to be converted to integral types, or vice versa. When an object of integral type is converted to a floating type and the original value cannot be represented exactly, the result is either the next higher or the next lower representable value.

When an object of floating type is converted to an integral type, the fractional part is truncated. No rounding takes place in the conversion process. Truncation means that a number like 1.3 is converted to 1, and -1.3 is converted to -1.

## Arithmetic conversions

Many binary operators (discussed in [Expressions with Binary Operators](#)) cause conversions of operands and yield results the same way. The way these operators cause conversions is called "usual arithmetic conversions." Arithmetic conversions of operands of different native types are performed as shown in the following table. Typedef types behave according to their underlying native types.

### Conditions for Type Conversion

CONDITIONS MET	CONVERSION
Either operand is of type <b>long double</b> .	Other operand is converted to type <b>long double</b> .
Preceding condition not met and either operand is of type <b>double</b> .	Other operand is converted to type <b>double</b> .
Preceding conditions not met and either operand is of type <b>float</b> .	Other operand is converted to type <b>float</b> .
Preceding conditions not met (none of the operands are of floating types).	<p>Integral promotions are performed on the operands as follows:</p> <ul style="list-style-type: none"> <li>- If either operand is of type <code>unsigned long</code>, the other operand is converted to type <code>unsigned long</code>.</li> <li>- If preceding condition not met, and if either operand is of type <b>long</b> and the other of type <code>unsigned int</code>, both operands are converted to type <code>unsigned long</code>.</li> <li>- If the preceding two conditions are not met, and if either operand is of type <b>long</b>, the other operand is converted to type <b>long</b>.</li> <li>- If the preceding three conditions are not met, and if either operand is of type <code>unsigned int</code>, the other operand is converted to type <code>unsigned int</code>.</li> <li>- If none of the preceding conditions are met, both operands are converted to type <code>int</code>.</li> </ul>

The following code illustrates the conversion rules described in the table:

```

double dVal;
float fVal;
int iVal;
unsigned long ulVal;

int main() {
    // iVal converted to unsigned long
    // result of multiplication converted to double
    dVal = iVal * ulVal;

    // ulVal converted to float
    // result of addition converted to double
    dVal = ulVal + fVal;
}

```

The first statement in the preceding example shows multiplication of two integral types, `iVal` and `ulVal`. The condition met is that neither operand is of floating type and one operand is of type `unsigned int`. Therefore, the other operand, `iVal`, is converted to type `unsigned int`. The result is assigned to `dVal`. The condition met is that one operand is of type **double**; therefore, the `unsigned int` result of the multiplication is converted to type **double**.

The second statement in the preceding example shows addition of a **float** and an integral type, `fVal` and `ulVal`. The `ulVal` variable is converted to type **float** (third condition in the table). The result of the addition is converted to type **double** (second condition in the table) and assigned to `dVal`.

## Pointer conversions

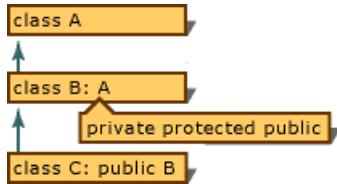
Pointers can be converted during assignment, initialization, comparison, and other expressions.

## Pointer to classes

There are two cases in which a pointer to a class can be converted to a pointer to a base class.

The first case is when the specified base class is accessible and the conversion is unambiguous. (See [Multiple Base Classes](#) for more information about ambiguous base-class references.)

Whether a base class is accessible depends on the kind of inheritance used in derivation. Consider the inheritance illustrated in the following figure.



Inheritance Graph for Illustration of Base-Class Accessibility

The following table shows the base-class accessibility for the situation illustrated in the figure.

### Base-Class Accessibility

TYPE OF FUNCTION	DERIVATION	CONVERSION FROM B* TO A* LEGAL?
External (not class-scoped) function	Private	No
	Protected	No
	Public	Yes
B member function (in B scope)	Private	Yes
	Protected	Yes
	Public	Yes
C member function (in C scope)	Private	No
	Protected	Yes
	Public	Yes

The second case in which a pointer to a class can be converted to a pointer to a base class is when you use an explicit type conversion. (See [Expressions with Explicit Type Conversions](#) for more information about explicit type conversions.)

The result of such a conversion is a pointer to the "subobject," the portion of the object that is completely described by the base class.

The following code defines two classes, `A` and `B`, where `B` is derived from `A`. (For more information on inheritance, see [Derived Classes](#).) It then defines `bObject`, an object of type `B`, and two pointers (`pA` and `pB`) that point to the object.

```

// C2039 expected
class A
{
public:
    int AComponent;
    int AMemberFunc();
};

class B : public A
{
public:
    int BComponent;
    int BMemberFunc();
};

int main()
{
    B bObject;
    A *pA = &bObject;
    B *pB = &bObject;

    pA->AMemberFunc(); // OK in class A
    pB->AMemberFunc(); // OK: inherited from class A
    pA->BMemberFunc(); // Error: not in class A
}

```

The pointer `pA` is of type `A *`, which can be interpreted as meaning "pointer to an object of type `A`." Members of `bObject` (such as `BComponent` and `BMemberFunc`) are unique to type `B` and are therefore inaccessible through `pA`. The `pA` pointer allows access only to those characteristics (member functions and data) of the object that are defined in class `A`.

## Pointer to function

A pointer to a function can be converted to type `void \*`, if type `void \*` is large enough to hold that pointer.

## Pointer to void

Pointers to type `void` can be converted to pointers to any other type, but only with an explicit type cast (unlike in C). (See [Expressions with Explicit Type Conversions](#) for more information about type casts.) A pointer to any type can be converted implicitly to a pointer to type `void`. A pointer to an incomplete object of a type can be converted to a pointer to `void` (implicitly) and back (explicitly). The result of such a conversion is equal to the value of the original pointer. An object is considered incomplete if it is declared, but there is insufficient information available to determine its size or base class.

A pointer to any object that is not `const` or `volatile` can be implicitly converted to a pointer of type `void \*`.

## const and volatile pointers

C++ does not supply a standard conversion from a `const` or `volatile` type to a type that is not `const` or `volatile`. However, any sort of conversion can be specified using explicit type casts (including conversions that are unsafe).

### NOTE

C++ pointers to members, except pointers to static members, are different from normal pointers and do not have the same standard conversions. Pointers to static members are normal pointers and have the same conversions as normal pointers.

## null pointer conversions

An integral constant expression that evaluates to zero, or such an expression cast to a pointer type, is converted to a pointer called the "null pointer." This pointer is guaranteed to compare unequal to a pointer to any valid

object or function (except for pointers to based objects, which can have the same offset and still point to different objects).

In C++11 the `nullptr` type should be preferred to the C-style null pointer.

### Pointer expression conversions

Any expression with an array type can be converted to a pointer of the same type. The result of the conversion is a pointer to the first array element. The following example demonstrates such a conversion:

```
char szPath[_MAX_PATH]; // Array of type char.  
char *pszPath = szPath; // Equals &szPath[0].
```

An expression that results in a function returning a particular type is converted to a pointer to a function returning that type, except when:

- The expression is used as an operand to the address-of operator (`&`).
- The expression is used as an operand to the function-call operator.

## Reference conversions

A reference to a class can be converted to a reference to a base class in the following cases:

- The specified base class is accessible.
- The conversion is unambiguous. (See [Multiple Base Classes](#) for more information about ambiguous base-class references.)

The result of the conversion is a pointer to the subobject that represents the base class.

## Pointer to member

Pointers to class members can be converted during assignment, initialization, comparison, and other expressions. This section describes the following pointer-to-member conversions:

### Pointer to base class member

A pointer to a member of a base class can be converted to a pointer to a member of a class derived from it, when the following conditions are met:

- The inverse conversion, from pointer to derived class to base-class pointer, is accessible.
- The derived class does not inherit virtually from the base class.

When the left operand is a pointer to member, the right operand must be of pointer-to-member type or be a constant expression that evaluates to 0. This assignment is valid only in the following cases:

- The right operand is a pointer to a member of the same class as the left operand.
- The left operand is a pointer to a member of a class derived publicly and unambiguously from the class of the right operand.

## Integral constant conversions

An integral constant expression that evaluates to zero is converted to a pointer called the "null pointer." This pointer is guaranteed to compare unequal to a pointer to any valid object or function (except for pointers to based objects, which can have the same offset and still point to different objects).

The following code illustrates the definition of a pointer to member `i` in class `A`. The pointer, `pai`, is initialized to 0, which is the null pointer.

```
class A
{
public:
    int i;
};

int A::*pai = 0;
```

## See Also

[C++ Language Reference](#)

# C++ Built-in Operators, Precedence and Associativity

12/21/2017 • 2 min to read • [Edit Online](#)

The C++ language includes all C operators and adds several new operators. Operators specify an evaluation to be performed on one or more operands.

Operator *precedence* specifies the order of operations in expressions that contain more than one operator. Operator *associativity* specifies whether, in an expression that contains multiple operators with the same precedence, an operand is grouped with the one on its left or the one on its right. The following table shows the precedence and associativity of C++ operators (from highest to lowest precedence). Operators with the same precedence number have equal precedence unless another relationship is explicitly forced by parentheses.

## C++ Operator Precedence and Associativity

OPERATOR DESCRIPTION	OPERATOR
<b>Group 1 precedence, no associativity</b>	
Scope resolution	::
<b>Group 2 precedence, left to right associativity</b>	
Member selection (object or pointer)	. or ->
Array subscript	[]
Function call	()
Postfix increment	++
Postfix decrement	--
Type name	typeid
Constant type conversion	const_cast
Dynamic type conversion	dynamic_cast
Reinterpreted type conversion	reinterpret_cast
Static type conversion	static_cast
<b>Group 3 precedence, right to left associativity</b>	
Size of object or type	sizeof
Prefix increment	++

OPERATOR DESCRIPTION	OPERATOR
Prefix decrement	--
One's complement	~
Logical not	!
Unary negation	-
Unary plus	+
Address-of	&
Indirection	*
Create object	new
Destroy object	delete
Cast	0
<b>Group 4 precedence, left to right associativity</b>	
Pointer-to-member (objects or pointers)	.* or ->*
<b>Group 5 precedence, left to right associativity</b>	
Multiplication	*
Division	/
Modulus	%
<b>Group 6 precedence, left to right associativity</b>	
Addition	+
Subtraction	-
<b>Group 7 precedence, left to right associativity</b>	
Left shift	<<
Right shift	>>
<b>Group 8 precedence, left to right associativity</b>	
Less than	<
Greater than	>

OPERATOR DESCRIPTION	OPERATOR
Less than or equal to	<code>&lt;=</code>
Greater than or equal to	<code>&gt;=</code>
<b>Group 9 precedence, left to right associativity</b>	
Equality	<code>==</code>
Inequality	<code>!=</code>
<b>Group 10 precedence left to right associativity</b>	
Bitwise AND	<code>&amp;</code>
<b>Group 11 precedence, left to right associativity</b>	
Bitwise exclusive OR	<code>^</code>
<b>Group 12 precedence, left to right associativity</b>	
Bitwise inclusive OR	<code> </code>
<b>Group 13 precedence, left to right associativity</b>	
Logical AND	<code>&amp;&amp;</code>
<b>Group 14 precedence, left to right associativity</b>	
Logical OR	<code>  </code>
<b>Group 15 precedence, right to left associativity</b>	
Conditional	<code>? :</code>
<b>Group 16 precedence, right to left associativity</b>	
Assignment	<code>=</code>
Multiplication assignment	<code>*=</code>
Division assignment	<code>/=</code>
Modulus assignment	<code>%=</code>
Addition assignment	<code>+ =</code>
Subtraction assignment	<code>- =</code>
Left-shift assignment	<code>&lt;&lt;=</code>

OPERATOR DESCRIPTION	OPERATOR
Right-shift assignment	<code>&gt;&gt;=</code>
Bitwise AND assignment	<code>&amp;=</code>
Bitwise inclusive OR assignment	<code> =</code>
Bitwise exclusive OR assignment	<code>^=</code>
<b>Group 17 precedence, right to left associativity</b>	
throw expression	<code>throw</code>
<b>Group 18 precedence, left to right associativity</b>	
Comma	<code>,</code>

## See Also

[Operator Overloading](#)

# \_\_alignof Operator

12/21/2017 • 1 min to read • [Edit Online](#)

C++11 introduces the  `operator that returns the alignment, in bytes, of the specified type. For maximum portability, you should use the  operator instead of the Microsoft-specific __alignof operator.`

## Microsoft Specific

Returns a value of type `size_t` that is the alignment requirement of the type.

## Syntax

```
__alignof(  
    type  
)
```

## Remarks

For example:

EXPRESSION	VALUE
<code>__alignof( char )</code>	1
<code>__alignof( short )</code>	2
<code>__alignof( int )</code>	4
<code>__alignof( __int64 )</code>	8
<code>__alignof( float )</code>	4
<code>__alignof( double )</code>	8
<code>__alignof( char* )</code>	4

The `__alignof` value is the same as the value for `sizeof` for basic types. Consider, however, this example:

```
typedef struct { int a; double b; } S;  
// __alignof(S) == 8
```

In this case, the `__alignof` value is the alignment requirement of the largest element in the structure.

Similarly, for

```
typedef __declspec(align(32)) struct { int a; } S;
```

`__alignof(S)` is equal to `32`.

One use for `__alignof` would be as a parameter to one of your own memory-allocation routines. For example, given the following defined structure `s`, you could call a memory-allocation routine named `aligned_malloc` to allocate memory on a particular alignment boundary.

```
typedef __declspec(align(32)) struct { int a; double b; } S;
int n = 50; // array size
S* p = (S*)aligned_malloc(n * sizeof(S), __alignof(S));
```

For more information on modifying alignment, see:

- [pack](#)
- [align](#)
- [\\_\\_unaligned](#)
- [/Zp \(Struct Member Alignment\)](#)
- [Examples of Structure Alignment](#) (x64 specific)

For more information on differences in alignment in code for x86 and x64, see:

- [Conflicts with the x86 Compiler](#)

## END Microsoft Specific

## See Also

[Expressions with Unary Operators](#)

[Keywords](#)

# `_uuidof` Operator

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Retrieves the GUID attached to the expression.

## Syntax

```
_uuidof (
    expression
)
```

## Remarks

The *expression* can be a type name, pointer, reference, or array of that type, a template specialized on these types, or a variable of these types. The argument is valid as long as the compiler can use it to find the attached GUID.

A special case of this intrinsic is when either **0** or **NULL** is supplied as the argument. In this case, `_uuidof` will return a GUID made up of zeros.

Use this keyword to extract the GUID attached to:

- An object by the `uuid` extended attribute.
- A library block created with the `module` attribute.

### NOTE

In a debug build, `_uuidof` always initializes an object dynamically (at runtime). In a release build, `_uuidof` can statically (at compile time) initialize an object.

## Example

The following code (compiled with ole32.lib) will display the uuid of a library block created with the `module` attribute:

```
// expre_uuidof.cpp
// compile with: ole32.lib
#include <stdio.h>
#include <windows.h>

[emitidl];
[module(name="MyLib")];
[export]
struct stuff {
    int i;
};

int main() {
    LPOLESTR lpolestr;
    StringFromCLSID(__uuidof(MyLib), &lpolestr);
    wprintf_s(L"%s", lpolestr);
    CoTaskMemFree(lpolestr);
}
```

## Comments

In cases where the library name is no longer in scope, you can use `_LIBID\` instead of `__uuidof`. For example:

```
StringFromCLSID(__LIBID_, &lpolestr);
```

### END Microsoft Specific

## See Also

[Expressions with Unary Operators](#)

[Keywords](#)

# Additive Operators: + and -

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
expression + expression  
expression - expression
```

## Remarks

The additive operators are:

- Addition (+)
- Subtraction (-)

These binary operators have left-to-right associativity.

The additive operators take operands of arithmetic or pointer types. The result of the addition (+) operator is the sum of the operands. The result of the subtraction (-) operator is the difference between the operands. If one or both of the operands are pointers, they must be pointers to objects, not to functions. If both operands are pointers, the results are not meaningful unless both are pointers to objects in the same array.

Additive operators take operands of *arithmetic*, *integral*, and *scalar* types. These are defined in the following table.

### Types Used with Additive Operators

TYPE	MEANING
<i>arithmetic</i>	Integral and floating types are collectively called "arithmetic" types.
<i>integral</i>	Types char and int of all sizes (long, short) and enumerations are "integral" types.
<i>scalar</i>	Scalar operands are operands of either arithmetic or pointer type.

The legal combinations for these operators are:

*arithmetic* + *arithmetic*

*scalar* + *integral*

*integral* + *scalar*

*arithmetic* - *arithmetic*

*scalar* - *scalar*

Note that addition and subtraction are not equivalent operations.

If both operands are of arithmetic type, the conversions covered in [Standard Conversions](#) are applied to the operands, and the result is of the converted type.

## Example

```
// expre_Additive_Operators.cpp
// compile with: /EHsc
#include <iostream>
#define SIZE 5
using namespace std;
int main() {
    int i = 5, j = 10;
    int n[SIZE] = { 0, 1, 2, 3, 4 };
    cout << "5 + 10 = " << i + j << endl
        << "5 - 10 = " << i - j << endl;

    // use pointer arithmetic on array

    cout << "n[3] = " << *( n + 3 ) << endl;
}
```

## Pointer addition

If one of the operands in an addition operation is a pointer to an array of objects, the other must be of integral type. The result is a pointer that is of the same type as the original pointer and that points to another array element. The following code fragment illustrates this concept:

```
short IntArray[10]; // Objects of type short occupy 2 bytes
short *pIntArray = IntArray;

for( int i = 0; i < 10; ++i )
{
    *pIntArray = i;
    cout << *pIntArray << "\n";
    pIntArray = pIntArray + 1;
}
```

Although the integral value 1 is added to `pIntArray`, it does not mean "add 1 to the address"; rather it means "adjust the pointer to point to the next object in the array" that happens to be 2 bytes (or `sizeof( int )`) away.

### NOTE

Code of the form `pIntArray = pIntArray + 1` is rarely found in C++ programs; to perform an increment, these forms are preferable: `pIntArray++` or `pIntArray += 1`.

## Pointer subtraction

If both operands are pointers, the result of subtraction is the difference (in array elements) between the operands. The subtraction expression yields a signed integral result of type `ptrdiff_t` (defined in the standard include file `STDDEF.H`).

One of the operands can be of integral type, as long as it is the second operand. The result of the subtraction is of the same type as the original pointer. The value of the subtraction is a pointer to the  $(n - i)$ th array element, where  $n$  is the element pointed to by the original pointer and  $i$  is the integral value of the second operand.

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Additive Operators](#)

# Address-of Operator: &

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
& cast-expression
```

## Remarks

The unary address-of operator (**&**) takes the address of its operand. The operand of the address-of operator can be either a function designator or an l-value that designates an object that is not a bit field and is not declared with the **register** storage-class specifier.

The address-of operator can only be applied to variables with fundamental, structure, class, or union types that are declared at the file-scope level, or to subscripted array references. In these expressions, a constant expression that does not include the address-of operator can be added to or subtracted from the address-of expression.

When applied to functions or l-values, the result of the expression is a pointer type (an r-value) derived from the type of the operand. For example, if the operand is of type `char`, the result of the expression is of type pointer to `char`. The address-of operator, applied to **const** or **volatile** objects, evaluates to **const** `type \*` or **volatile** `type \*`, where `type` is the type of the original object.

When the address-of operator is applied to a **qualified name**, the result depends on whether the *qualified-name* specifies a static member. If so, the result is a pointer to the type specified in the declaration of the member. If the member is not static, the result is a pointer to the member *name* of the class indicated by *qualified-class-name*. (See [Primary Expressions](#) for more about *qualified-class-name*.) The following code fragment shows how the result differs, depending on whether the member is static:

```
// expre_Address_Of_Operator.cpp
// C2440 expected
class PTM {
public:
    int iValue;
    static float fValue;
};

int main() {
    int PTM::*piValue = &PTM::iValue; // OK: non-static
    float PTM::*pfValue = &PTM::fValue; // C2440 error: static
    float *spfValue     = &PTM::fValue; // OK
}
```

In this example, the expression `&PTM::fValue` yields type `float *` instead of type `float PTM::*` because `fValue` is a static member.

The address of an overloaded function can be taken only when it is clear which version of the function is being referenced. See [Function Overloading](#) for information about how to obtain the address of a particular overloaded function.

Applying the address-of operator to a reference type gives the same result as applying the operator to the object to which the reference is bound. For example:

## Example

```
// expre_Address_Of_Operator2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
int main() {
    double d;           // Define an object of type double.
    double& rd = d;   // Define a reference to the object.

    // Obtain and compare their addresses
    if( &d == &rd )
        cout << "&d equals &rd" << endl;
}
```

## Output

```
&d equals &rd
```

The following example uses the address-of operator to pass a pointer argument to a function:

```
// expre_Address_Of_Operator3.cpp
// compile with: /EHsc
// Demonstrate address-of operator &

#include <iostream>
using namespace std;

// Function argument is pointer to type int
int square( int *n ) {
    return (*n) * (*n);
}

int main() {
    int mynum = 5;
    cout << square( &mynum ) << endl;    // pass address of int
}
```

## Output

```
25
```

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Lvalue Reference Declarator: &](#)

[Indirection and Address-of Operators](#)

# Assignment Operators

12/21/2017 • 5 min to read • [Edit Online](#)

## Syntax

```
expression assignment-operator expression
assignment-operator : one of
= *= /= %= += -= <<= >>= &= ^= |=
```

## Remarks

Assignment operators store a value in the object designated by the left operand. There are two kinds of assignment operations: simple assignment, in which the value of the second operand is stored in the object specified by the first operand, and compound assignment, in which an arithmetic, shift, or bitwise operation is performed prior to storing the result. All assignment operators in the following table except the = operator are compound assignment operators.

### Assignment Operators

OPERATOR	MEANING
=	Store the value of the second operand in the object specified by the first operand (simple assignment).
*=	Multiply the value of the first operand by the value of the second operand; store the result in the object specified by the first operand.
/=	Divide the value of the first operand by the value of the second operand; store the result in the object specified by the first operand.
%=	Take modulus of the first operand specified by the value of the second operand; store the result in the object specified by the first operand.
+=	Add the value of the second operand to the value of the first operand; store the result in the object specified by the first operand.
-=	Subtract the value of the second operand from the value of the first operand; store the result in the object specified by the first operand.
<<=	Shift the value of the first operand left the number of bits specified by the value of the second operand; store the result in the object specified by the first operand.
>>=	Shift the value of the first operand right the number of bits specified by the value of the second operand; store the result in the object specified by the first operand.

OPERATOR	MEANING
<b>&amp;=</b>	Obtain the bitwise AND of the first and second operands; store the result in the object specified by the first operand.
<b>^=</b>	Obtain the bitwise exclusive OR of the first and second operands; store the result in the object specified by the first operand.
<b>&amp;#124;=</b>	Obtain the bitwise inclusive OR of the first and second operands; store the result in the object specified by the first operand.

## Operator Keywords

Three of the compound assignment operators have text equivalents. They are:

OPERATOR	EQUIVALENT
<b>&amp;=</b>	<code>and_eq</code>
<b>&amp;#124;=</b>	<code>or_eq</code>
<b>^=</b>	<code>xor_eq</code>

There are two ways to access these operator keywords in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// exre_Assignment_Operators.cpp
// compile with: /EHsc
// Demonstrate assignment operators
#include <iostream>
using namespace std;
int main() {
    int a = 3, b = 6, c = 10, d = 0xAAAA, e = 0x5555;

    a += b;      // a is 9
    b %= a;      // b is 6
    c >>= 1;     // c is 5
    d |= e;      // Bitwise--d is 0xFFFF

    cout << "a = 3, b = 6, c = 10, d = 0xAAAA, e = 0x5555" << endl
        << "a += b yields " << a << endl
        << "b %= a yields " << b << endl
        << "c >>= 1 yields " << c << endl
        << "d |= e yields " << hex << d << endl;
}
```

## Simple assignment

The simple assignment operator (=) causes the value of the second operand to be stored in the object specified by the first operand. If both objects are of arithmetic types, the right operand is converted to the type of the left, prior to storing the value.

Objects of const and volatile types can be assigned to l-values of types that are just volatile or that are neither

const nor volatile.

Assignment to objects of class type (struct, union, and class types) is performed by a function named `operator=`. The default behavior of this operator function is to perform a bitwise copy; however, this behavior can be modified using overloaded operators. (See [Overloaded Operators](#) for more information.)

An object of any unambiguously derived class from a given base class can be assigned to an object of the base class. The reverse is not true because there is an implicit conversion from derived class to base class but not from base class to derived class. For example:

```
// expre_SimpleAssignment.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
class ABase
{
public:
    ABase() { cout << "constructing ABase\n"; }

class ADerived : public ABase
{
public:
    ADerived() { cout << "constructing ADerived\n"; }

int main()
{
    ABase aBase;
    ADerived aDerived;

    aBase = aDerived; // OK
    aDerived = aBase; // C2679
}
```

Assignments to reference types behave as if the assignment were being made to the object to which the reference points.

For class-type objects, assignment is different from initialization. To illustrate how different assignment and initialization can be, consider the code

```
UserType1 A;
UserType2 B = A;
```

The preceding code shows an initializer; it calls the constructor for `UserType2` that takes an argument of type `UserType1`. Given the code

```
UserType1 A;
UserType2 B;

B = A;
```

the assignment statement

```
B = A;
```

can have one of the following effects:

- Call the function `operator=` for `UserType2`, provided `operator=` is provided with a `UserType1` argument.
- Call the explicit conversion function `UserType1::operator UserType2`, if such a function exists.
- Call a constructor `UserType2::UserType2`, provided such a constructor exists, that takes a `UserType1` argument and copies the result.

## Compound assignment

The compound assignment operators, shown in the table in [Assignment Operators](#), are specified in the form  $e1 \boxed{op} = e2$ , where  $e1$  is a modifiable l-value not of const type and  $e2$  is one of the following:

- An arithmetic type
- A pointer, if  $\boxed{op}$  is + or -

The  $e1 \boxed{op} = e2$  form behaves as  $e1 = e1 \boxed{op} e2$ , but  $e1$  is evaluated only once.

Compound assignment to an enumerated type generates an error message. If the left operand is of a pointer type, the right operand must be of a pointer type or it must be a constant expression that evaluates to 0. If the left operand is of an integral type, the right operand must not be of a pointer type.

## Result of assignment operators

The assignment operators return the value of the object specified by the left operand after the assignment. The resultant type is the type of the left operand. The result of an assignment expression is always an l-value. These operators have right-to-left associativity. The left operand must be a modifiable l-value.

In ANSI C, the result of an assignment expression is not an l-value. Therefore, the legal C++ expression `(a += b) += c` is illegal in C.

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Assignment Operators](#)

# Bitwise AND Operator: &

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
expression  
&  
expression
```

## Remarks

The expressions may be other and-expressions, or (subject to the type restrictions mentioned below) equality expressions, relational expressions, additive expressions, multiplicative expressions, pointer to member expressions, cast expressions, unary expressions, postfix expressions, or primary expressions.

The bitwise AND operator (**&**) compares each bit of the first operand to the corresponding bit of the second operand. If both bits are 1, the corresponding result bit is set to 1. Otherwise, the corresponding result bit is set to 0.

Both operands to the bitwise AND operator must be of integral types. The usual arithmetic conversions covered in [Standard Conversions](#), are applied to the operands.

## Operator Keyword for &

The `bitand` operator is the text equivalent of **&**. There are two ways to access the `bitand` operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Bitwise_AND_Operator.cpp  
// compile with: /EHsc  
// Demonstrate bitwise AND  
#include <iostream>  
using namespace std;  
int main() {  
    unsigned short a = 0xFFFF;      // pattern 1111 ...  
    unsigned short b = 0xAAAA;      // pattern 1010 ...  
  
    cout << hex << ( a & b ) << endl;  // prints "aaaa", pattern 1010 ...  
}
```

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Bitwise Operators](#)

# Bitwise Exclusive OR Operator: ^

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
expression ^ expression
```

## Remarks

The bitwise exclusive OR operator (^) compares each bit of its first operand to the corresponding bit of its second operand. If one bit is 0 and the other bit is 1, the corresponding result bit is set to 1. Otherwise, the corresponding result bit is set to 0.

Both operands to the bitwise exclusive OR operator must be of integral types. The usual arithmetic conversions covered in [Standard Conversions](#) are applied to the operands.

## Operator Keyword for ^

The **xor** operator is the text equivalent of ^. There are two ways to access the **xor** operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Bitwise_Exclusive_OR_Operator.cpp
// compile with: /EHsc
// Demonstrate bitwise exclusive OR
#include <iostream>
using namespace std;
int main() {
    unsigned short a = 0x5555;      // pattern 0101 ...
    unsigned short b = 0xFFFF;      // pattern 1111 ...

    cout << hex << ( a ^ b ) << endl; // prints "aaaa" pattern 1010 ...
}
```

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

# Bitwise Inclusive OR Operator: |

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
expression  
|  
expression
```

## Remarks

The bitwise inclusive OR operator (|) compares each bit of its first operand to the corresponding bit of its second operand. If either bit is 1, the corresponding result bit is set to 1. Otherwise, the corresponding result bit is set to 0.

Both operands to the bitwise inclusive OR operator must be of integral types. The usual arithmetic conversions covered in [Standard Conversions](#) are applied to the operands.

## Operator Keyword for |

The `bitor` operator is the text equivalent of `|`. There are two ways to access the `bitor` operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Bitwise_Inclusive_OR_Operator.cpp  
// compile with: /EHsc  
// Demonstrate bitwise inclusive OR  
#include <iostream>  
using namespace std;  
  
int main() {  
    unsigned short a = 0x5555;      // pattern 0101 ...  
    unsigned short b = 0xFFFF;      // pattern 1010 ...  
  
    cout << hex << ( a | b ) << endl;  // prints "ffff" pattern 1111 ...  
}
```

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

[C Bitwise Operators](#)

# Cast Operator: ()

12/21/2017 • 1 min to read • [Edit Online](#)

A type cast provides a method for explicit conversion of the type of an object in a specific situation.

## Syntax

```
unary-expression  
( type-name ) cast-expression
```

## Remarks

Any unary expression is considered a cast expression.

The compiler treats *cast-expression* as type *type-name* after a type cast has been made. Casts can be used to convert objects of any scalar type to or from any other scalar type. Explicit type casts are constrained by the same rules that determine the effects of implicit conversions. Additional restraints on casts may result from the actual sizes or representation of specific types.

## Example

```
// expre_CastOperator.cpp  
// compile with: /EHsc  
// Demonstrate cast operator  
#include <iostream>  
  
using namespace std;  
  
int main()  
{  
    double x = 3.1;  
    int i;  
    cout << "x = " << x << endl;  
    i = (int)x;    // assign i the integer part of x  
    cout << "i = " << i << endl;  
}
```

## Example

```

// expre_CastOperator2.cpp
// The following sample shows how to define and use a cast operator.
#include <string.h>
#include <stdio.h>

class CountedAnsiString
{
public:
    // Assume source is not null terminated
    CountedAnsiString(const char *pStr, size_t nSize) :
        m_nSize(nSize)
    {
        m_pStr = new char[sizeOfBuffer];

        strncpy_s(m_pStr, sizeOfBuffer, pStr, m_nSize);
        memset(&m_pStr[m_nSize], '!', 9); // for demonstration purposes.
    }

    // Various string-like methods...

    const char *GetRawBytes() const
    {
        return(m_pStr);
    }

    //
    // operator to cast to a const char *
    //
    operator const char *()
    {
        m_pStr[m_nSize] = '\0';
        return(m_pStr);
    }

    enum
    {
        sizeOfBuffer = 20
    } size;
}

private:
    char *m_pStr;
    const size_t m_nSize;
};

int main()
{
    const char *kStr = "Excitinggg";
    CountedAnsiString myStr(kStr, 8);

    const char *pRaw = myStr.GetRawBytes();
    printf_s("RawBytes truncated to 10 chars:  %.10s\n", pRaw);

    const char *pCast = myStr; // or (const char *)myStr;
    printf_s("Casted Bytes:  %s\n", pCast);

    puts("Note that the cast changed the raw internal string");
    printf_s("Raw Bytes after cast:  %s\n", pRaw);
}

```

```

RawBytes truncated to 10 chars:  Exciting!!
Casted Bytes:  Exciting
Note that the cast changed the raw internal string
Raw Bytes after cast:  Exciting

```

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Explicit Type Conversion Operator: \(\)](#)

[Casting Operators](#)

[Cast Operators](#)

# Comma Operator: ,

12/21/2017 • 1 min to read • [Edit Online](#)

Allows grouping two statements where one is expected.

## Syntax

```
expression , expression
```

## Remarks

The comma operator has left-to-right associativity. Two expressions separated by a comma are evaluated left to right. The left operand is always evaluated, and all side effects are completed before the right operand is evaluated.

Commas can be used as separators in some contexts, such as function argument lists. Do not confuse the use of the comma as a separator with its use as an operator; the two uses are completely different.

Consider the expression

$e1, e2$

The type and value of the expression are the type and value of  $e2$ ; the result of evaluating  $e1$  is discarded. The result is an l-value if the right operand is an l-value.

Where the comma is normally used as a separator (for example in actual arguments to functions or aggregate initializers), the comma operator and its operands must be enclosed in parentheses. For example:

```
func_one( x, y + 2, z );
func_two( (x--, y + 2), z );
```

In the function call to `func_one` above, three arguments, separated by commas, are passed: `x`, `y + 2`, and `z`. In the function call to `func_two`, parentheses force the compiler to interpret the first comma as the sequential-evaluation operator. This function call passes two arguments to `func_two`. The first argument is the result of the sequential-evaluation operation `(x--, y + 2)`, which has the value and type of the expression `y + 2`; the second argument is `z`.

## Example

```
// cpp_comma_operator.cpp
#include <stdio.h>
int main () {
    int i = 10, b = 20, c= 30;
    i = b, c;
    printf("%i\n", i);

    i = (b, c);
    printf("%i\n", i);
}
```

20

30

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Sequential-Evaluation Operator](#)

# Conditional Operator: ?:

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
expression ? expression : expression
```

## Remarks

The conditional operator (`?:`) is a ternary operator (it takes three operands). The conditional operator works as follows:

- The first operand is implicitly converted to `bool`. It is evaluated and all side effects are completed before continuing.
- If the first operand evaluates to `true` (1), the second operand is evaluated.
- If the first operand evaluates to `false` (0), the third operand is evaluated.

The result of the conditional operator is the result of whichever operand is evaluated — the second or the third. Only one of the last two operands is evaluated in a conditional expression.

Conditional expressions have right-to-left associativity. The first operand must be of integral or pointer type. The following rules apply to the second and third operands:

- If both operands are of the same type, the result is of that type.
- If both operands are of arithmetic or enumeration types, the usual arithmetic conversions (covered in [Standard Conversions](#)) are performed to convert them to a common type.
- If both operands are of pointer types or if one is a pointer type and the other is a constant expression that evaluates to 0, pointer conversions are performed to convert them to a common type.
- If both operands are of reference types, reference conversions are performed to convert them to a common type.
- If both operands are of type `void`, the common type is type `void`.
- If both operands are of the same user-defined type, the common type is that type.
- If the operands have different types and at least one of the operands has user-defined type then the language rules are used to determine the common type. (See warning below.)

Any combinations of second and third operands not in the preceding list are illegal. The type of the result is the common type, and it is an l-value if both the second and third operands are of the same type and both are l-values.

## WARNING

If the types of the second and third operands are not identical, then complex type conversion rules, as specified in the C++ Standard, are invoked. These conversions may lead to unexpected behavior including construction and destruction of temporary objects. For this reason, we strongly advise you to either (1) avoid using user-defined types as operands with the conditional operator or (2) if you do use user-defined types, then explicitly cast each operand to a common type.

## Example

```
// expre_Expressions_with_the_Conditional_Operator.cpp
// compile with: /EHsc
// Demonstrate conditional operator
#include <iostream>
using namespace std;
int main() {
    int i = 1, j = 2;
    cout << ( i > j ? i : j ) << " is greater." << endl;
}
```

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

[Conditional-Expression Operator](#)

# delete Operator (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Deallocates a block of memory.

## Syntax

```
[::] delete cast-expression  
[::] delete [ ] cast-expression
```

## Remarks

The *cast-expression* argument must be a pointer to a block of memory previously allocated for an object created with the [new operator](#). The **delete** operator has a result of type `void` and therefore does not return a value. For example:

```
CDialog* MyDialog = new CDialog;  
// use MyDialog  
delete MyDialog;
```

Using **delete** on a pointer to an object not allocated with **new** gives unpredictable results. You can, however, use **delete** on a pointer with the value 0. This provision means that, when **new** returns 0 on failure, deleting the result of a failed **new** operation is harmless. See [The new and delete Operators](#) for more information.

The **new** and **delete** operators can also be used for built-in types, including arrays. If `pointer` refers to an array, place empty brackets before `pointer`:

```
int* set = new int[100];  
//use set[]  
delete [] set;
```

Using the **delete** operator on an object deallocates its memory. A program that dereferences a pointer after the object is deleted can have unpredictable results or crash.

When **delete** is used to deallocate memory for a C++ class object, the object's destructor is called before the object's memory is deallocated (if the object has a destructor).

If the operand to the **delete** operator is a modifiable l-value, its value is undefined after the object is deleted.

## Using delete

There are two syntactic variants for the [delete operator](#): one for single objects and the other for arrays of objects. The following code fragment shows how these differ:

```

// expre_Using_delete.cpp
struct UDTtype
{
};

int main()
{
    // Allocate a user-defined object, UDObject, and an object
    // of type double on the free store using the
    // new operator.
    UDTtype *UDObject = new UDTtype;
    double *dObject = new double;
    // Delete the two objects.
    delete UDObject;
    delete dObject;
    // Allocate an array of user-defined objects on the
    // free store using the new operator.
    UDTtype (*UDArr)[7] = new UDTtype[5][7];
    // Use the array syntax to delete the array of objects.
    delete [] UDArr;
}

```

The following two cases produce undefined results: using the array form of delete (`delete [ ]`) on an object and using the nonarray form of delete on an array.

## Example

For examples of using **delete**, see [new operator](#).

## How delete works

The delete operator invokes the function **operator delete**.

For objects not of class type ([class](#), [struct](#), or [union](#)), the global delete operator is invoked. For objects of class type, the name of the deallocation function is resolved in global scope if the delete expression begins with the unary scope resolution operator (`::`). Otherwise, the delete operator invokes the destructor for an object prior to deallocating memory (if the pointer is not null). The delete operator can be defined on a per-class basis; if there is no such definition for a given class, the global operator `delete` is invoked. If the delete expression is used to deallocate a class object whose static type has a virtual destructor, the deallocation function is resolved through the virtual destructor of the dynamic type of the object.

## See Also

[Expressions with Unary Operators](#)

[Keywords](#)

[new and delete Operators](#)

# Equality Operators: == and !=

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
expression == expression
expression != expression
```

## Remarks

The binary equality operators compare their operands for strict equality or inequality.

The equality operators, equal to (`==`) and not equal to (`!=`), have lower precedence than the relational operators, but they behave similarly. The result type for these operators is `bool`.

The equal-to operator (`==`) returns **true** (1) if both operands have the same value; otherwise, it returns **false** (0).

The not-equal-to operator (`!=`) returns **true** if the operands do not have the same value; otherwise, it returns **false**.

## Operator Keyword for !=

The `not_eq` operator is the text equivalent of `!=`. There are two ways to access the `not_eq` operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Equality_Operators.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;

int main() {
    cout << boolalpha
        << "The true expression 3 != 2 yields: "
        << (3 != 2) << endl
        << "The false expression 20 == 10 yields: "
        << (20 == 10) << endl;
}
```

Equality operators can compare pointers to members of the same type. In such a comparison, pointer-to-member conversions are performed. Pointers to members can also be compared to a constant expression that evaluates to 0.

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Relational and Equality Operators](#)

# Explicit Type Conversion Operator: ()

12/21/2017 • 2 min to read • [Edit Online](#)

C++ allows explicit type conversion using syntax similar to the function-call syntax.

## Syntax

```
simple-type-name ( expression-list )
```

## Remarks

A *simple-type-name* followed by an *expression-list* enclosed in parentheses constructs an object of the specified type using the specified expressions. The following example shows an explicit type conversion to type int:

```
int i = int( d );
```

The following example shows a `Point` class.

## Example

```

// expre_Explicit_Type_Conversion_Operator.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class Point
{
public:
    // Define default constructor.
    Point() { _x = _y = 0; }
    // Define another constructor.
    Point( int X, int Y ) { _x = X; _y = Y; }

    // Define "accessor" functions as
    // reference types.
    unsigned& x() { return _x; }
    unsigned& y() { return _y; }
    void Show() { cout << "x = " << _x << ", "
                  << "y = " << _y << "\n"; }

private:
    unsigned _x;
    unsigned _y;
};

int main()
{
    Point Point1, Point2;

    // Assign Point1 the explicit conversion
    // of ( 10, 10 ).
    Point1 = Point( 10, 10 );

    // Use x() as an l-value by assigning an explicit
    // conversion of 20 to type unsigned.
    Point1.x() = unsigned( 20 );
    Point1.Show();

    // Assign Point2 the default Point object.
    Point2 = Point();
    Point2.Show();
}

```

## Output

```

x = 20, y = 10
x = 0, y = 0

```

Although the preceding example demonstrates explicit type conversion using constants, the same technique works to perform these conversions on objects. The following code fragment demonstrates this:

```

int i = 7;
float d;

d = float( i );

```

Explicit type conversions can also be specified using the "cast" syntax. The previous example, rewritten using the cast syntax, is:

```

d = (float)i;

```

Both cast and function-style conversions have the same results when converting from single values. However, in the function-style syntax, you can specify more than one argument for conversion. This difference is important for user-defined types. Consider a `Point` class and its conversions:

```
struct Point
{
    Point( short x, short y ) { _x = x; _y = y; }
    ...
    short _x, _y;
};

...
Point pt = Point( 3, 10 );
```

The preceding example, which uses function-style conversion, shows how to convert two values (one for `x` and one for `y`) to the user-defined type `Point`.

**Caution**

Use the explicit type conversions with care, since they override the C++ compiler's built-in type checking.

The `cast` notation must be used for conversions to types that do not have a *simple-type-name* (pointer or reference types, for example). Conversion to types that can be expressed with a *simple-type-name* can be written in either form. See [Type Specifiers](#) for more information about what constitutes a *simple-type-name*.

Type definition within casts is illegal.

## See Also

[Postfix Expressions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

# Function Call Operator: ()

12/21/2017 • 2 min to read • [Edit Online](#)

A postfix-expression followed by the function-call operator, (), specifies a function call.

## Syntax

```
postfix-expression  
( [argument-expression-list] )
```

## Remarks

The arguments to the function-call operator are zero or more expressions separated by commas — the actual arguments to the function.

The *postfix-expression* must evaluate to a function address (for example, a function identifier or the value of a function pointer), and *argument-expression-list* is a list of expressions (separated by commas) whose values (the arguments) are passed to the function. The *argument-expression-list* argument can be empty.

The *postfix-expression* must be of one of these types:

- Function returning type `T`. An example declaration is

```
T func( int i )
```

- Pointer to a function returning type `T`. An example declaration is

```
T (*func)( int i )
```

- Reference to a function returning type `T`. An example declaration is

```
T (&func)(int i)
```

- Pointer-to-member function dereference returning type `T`. Example function calls are

```
(pObject->*pmf)();  
(Object.*pmf)();
```

## Example

The following example calls the standard library function `strcat_s` with three arguments:

```

// expre_Function_Call_Operator.cpp
// compile with: /EHsc

#include <iostream>
#include <string>

// C++ Standard Library name space
using namespace std;

int main()
{
    enum
    {
        sizeOfBuffer = 20
    };

    char s1[ sizeOfBuffer ] = "Welcome to ";
    char s2[ ] = "C++";

    strcat_s( s1, sizeOfBuffer, s2 );

    cout << s1 << endl;
}

```

Welcome to C++

## Function call results

A function call evaluates to an r-value unless the function is declared as a reference type. Functions with reference return type evaluate to l-values, and can be used on the left side of an assignment statement as follows:

```

// expre_Function_Call_Results.cpp
// compile with: /EHsc
#include <iostream>
class Point
{
public:
    // Define "accessor" functions as
    // reference types.
    unsigned& x() { return _x; }
    unsigned& y() { return _y; }
private:
    unsigned _x;
    unsigned _y;
};

using namespace std;
int main()
{
    Point ThePoint;

    ThePoint.x() = 7;           // Use x() as an l-value.
    unsigned y = ThePoint.y(); // Use y() as an r-value.

    // Use x() and y() as r-values.
    cout << "x = " << ThePoint.x() << "\n"
        << "y = " << ThePoint.y() << "\n";
}

```

The preceding code defines a class called `Point`, which contains private data objects that represent *x* and *y* coordinates. These data objects must be modified and their values retrieved. This program is only one of several

designs for such a class; use of the `GetX` and `SetX` or `GetX` and `SetY` functions is another possible design.

Functions that return class types, pointers to class types, or references to class types can be used as the left operand to member-selection operators. Therefore, the following code is legal:

```
// exprre_Function_Results2.cpp
class A {
public:
    A() {}
    A(int i) {}
    int SetA( int i ) {
        return (I = i);
    }

    int GetA() {
        return I;
    }

private:
    int I;
};

A func1() {
    A a = 0;
    return a;
}

A* func2() {
    A *a = new A();
    return a;
}

A& func3() {
    A *a = new A();
    A &b = *a;
    return b;
}

int main() {
    int iResult = func1().GetA();
    func2()->SetA( 3 );
    func3().SetA( 7 );
}
```

Functions can be called recursively. For more information about function declarations, see [Functions](#). Related material is in [Program and Linkage](#).

## See Also

[Postfix Expressions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Function Call](#)

# Indirection Operator: \*

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
* cast-expression
```

## Remarks

The unary indirection operator (`\*`) dereferences a pointer; that is, it converts a pointer value to an l-value. The operand of the indirection operator must be a pointer to a type. The result of the indirection expression is the type from which the pointer type is derived. The use of the `\*` operator in this context is different from its meaning as a binary operator, which is multiplication.

If the operand points to a function, the result is a function designator. If it points to a storage location, the result is an l-value designating the storage location.

The indirection operator may be used cumulatively to dereference pointers to pointers. For example:

```
// exprre_Indirection_Operator.cpp
// compile with: /EHsc
// Demonstrate indirection operator
#include <iostream>
using namespace std;
int main() {
    int n = 5;
    int *pn = &n;
    int **ppn = &pn;

    cout << "Value of n:\n"
        << "direct value: " << n << endl
        << "indirect value: " << *pn << endl
        << "doubly indirect value: " << **ppn << endl
        << "address of n: " << pn << endl
        << "address of n via indirection: " << *ppn << endl;
}
```

If the pointer value is invalid, the result is undefined. The following list includes some of the most common conditions that invalidate a pointer value.

- The pointer is a null pointer.
- The pointer specifies the address of a local item that is not visible at the time of the reference.
- The pointer specifies an address that is inappropriately aligned for the type of the object pointed to.
- The pointer specifies an address not used by the executing program.

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Address-of Operator: &](#)

## Indirection and Address-of Operators

# Left Shift and Right Shift Operators (>> and <<)

12/21/2017 • 5 min to read • [Edit Online](#)

The bitwise shift operators are the right-shift operator (>>), which moves the bits of *shift-expression* to the right, and the left-shift operator (<<), which moves the bits of *shift-expression* to the left.<sup>1</sup>

## Syntax

```
shift-expression << additive-expression  
shift-expression >> additive-expression
```

## Remarks

### IMPORTANT

The following descriptions and examples are valid on Windows for X86 and x64 architectures. The implementation of left-shift and right-shift operators is significantly different on Windows RT for ARM devices. For more information, see the "Shift Operators" section of the [Hello ARM](#) blog post.

## Left Shifts

The left-shift operator causes the bits in *shift-expression* to be shifted to the left by the number of positions specified by *additive-expression*. The bit positions that have been vacated by the shift operation are zero-filled. A left shift is a logical shift (the bits that are shifted off the end are discarded, including the sign bit). For more information about the kinds of bitwise shifts, see [Bitwise shifts](#).

The following example shows left-shift operations using unsigned numbers. The example shows what is happening to the bits by representing the value as a bitset. For more information, see [bitset Class](#).

```
#include <iostream>  
#include <bitset>  
using namespace std;  
  
int main() {  
    unsigned short short1 = 4;  
    bitset<16> bitset1{short1}; // the bitset representation of 4  
    cout << bitset1 << endl; // 000000000000100  
  
    unsigned short short2 = short1 << 1; // 4 left-shifted by 1 = 8  
    bitset<16> bitset2{short2};  
    cout << bitset2 << endl; // 0000000000001000  
  
    unsigned short short3 = short1 << 2; // 4 left-shifted by 2 = 16  
    bitset<16> bitset3{short3};  
    cout << bitset3 << endl; // 00000000000010000  
}
```

If you left-shift a signed number so that the sign bit is affected, the result is undefined. The following example shows what happens in Visual C++ when a 1 bit is left-shifted into the sign bit position.

```

#include <iostream>
#include <bitset>
using namespace std;

int main() {
    short short1 = 16384;
    bitset<16> bitset1{short2};
    cout << bitset1 << endl; // 0100000000000000

    short short3 = short1 << 1;
    bitset<16> bitset3{short3}; // 16384 left-shifted by 1 = -32768
    cout << bitset3 << endl; // 1000000000000000

    short short4 = short1 << 14;
    bitset<16> bitset4{short4}; // 4 left-shifted by 14 = 0
    cout << bitset4 << endl; // 0000000000000000
}

```

## Right Shifts

The right-shift operator causes the bit pattern in *shift-expression* to be shifted to the right by the number of positions specified by *additive-expression*. For unsigned numbers, the bit positions that have been vacated by the shift operation are zero-filled. For signed numbers, the sign bit is used to fill the vacated bit positions. In other words, if the number is positive, 0 is used, and if the number is negative, 1 is used.

### IMPORTANT

The result of a right-shift of a signed negative number is implementation-dependent. Although Visual C++ uses the sign bit to fill vacated bit positions, there is no guarantee that other implementations also do so.

This example shows right-shift operations using unsigned numbers:

```

#include <iostream>
#include <bitset>
using namespace std;

int main() {
    unsigned short short11 = 1024;
    bitset<16> bitset11{short11};
    cout << bitset11 << endl; // 0000010000000000

    unsigned short short12 = short11 >> 1; // 512
    bitset<16> bitset12{short12};
    cout << bitset12 << endl; // 0000001000000000

    unsigned short short13 = short11 >> 10; // 1
    bitset<16> bitset13{short13};
    cout << bitset13 << endl; // 0000000000000001

    unsigned short short14 = short11 >> 11; // 0
    bitset<16> bitset14{short14};
    cout << bitset14 << endl; // 0000000000000000
}

```

The next example shows right-shift operations with positive signed numbers.

```

#include <iostream>
#include <bitset>
using namespace std;

int main() {
    short short1 = 1024;
    bitset<16> bitset1{short1};
    cout << bitset1 << endl;      // 0000010000000000

    short short2 = short1 >> 1; // 512
    bitset<16> bitset2{short2};
    cout << bitset2 << endl;      // 0000001000000000

    short short3 = short1 >> 11; // 0
    bitset<16> bitset3{short3};
    cout << bitset3 << endl;      // 0000000000000000
}

```

The next example shows right-shift operations with negative signed integers.

```

#include <iostream>
#include <bitset>
using namespace std;

int main() {
    short neg1 = -16;
    bitset<16> bn1{neg1};
    cout << bn1 << endl; // 1111111111110000

    short neg2 = neg1 >> 1; // -8
    bitset<16> bn2{neg2};
    cout << bn2 << endl; // 1111111111110000

    short neg3 = neg1 >> 2; // -4
    bitset<16> bn3{neg3};
    cout << bn3 << endl; // 1111111111110000

    short neg4 = neg1 >> 4; // -1
    bitset<16> bn4{neg4};
    cout << bn4 << endl; // 1111111111111111

    short neg5 = neg1 >> 5; // -1
    bitset<16> bn5{neg5};
    cout << bn5 << endl; // 1111111111111111
}

```

## Shifts and Promotions

The expressions on both sides of a shift operator must be integral types. Integral promotions are performed according to the rules described in [Standard Conversions](#). The type of the result is the same as the type of the promoted *shift-expression*.

In the following example, a variable of type `char` is promoted to an `int`.

```

#include <iostream>
#include <typeinfo>

using namespace std;

int main() {
    char char1 = 'a';

    auto promoted1 = char1 << 1; // 194
    cout << typeid(promoted1).name() << endl; // int

    auto promoted2 = char1 << 10; // 99328
    cout << typeid(promoted2).name() << endl; // int
}

```

## Additional Details

The result of a shift operation is undefined if *additive-expression* is negative or if *additive-expression* is greater than or equal to the number of bits in the (promoted) *shift-expression*. No shift operation is performed if *additive-expression* is 0.

```

#include <iostream>
#include <bitset>
using namespace std;

int main() {
    unsigned int int1 = 4;
    bitset<32> b1{int1};
    cout << b1 << endl; // 000000000000000000000000000000100

    unsigned int int2 = int1 << -3; // C4293: '<<' : shift count negative or too big, undefined behavior
    unsigned int int3 = int1 >> -3; // C4293: '>>' : shift count negative or too big, undefined behavior

    unsigned int int4 = int1 << 32; // C4293: '<<' : shift count negative or too big, undefined behavior

    unsigned int int5 = int1 >> 32; // C4293: '>>' : shift count negative or too big, undefined behavior

    unsigned int int6 = int1 << 0;
    bitset<32> b6{int6};
    cout << b6 << endl; // 000000000000000000000000000000100 (no change)
}

```

## Footnotes

<sup>1</sup> The following is the description of the shift operators in the C++11 ISO specification (INCITS/ISO/IEC 14882-2011[2012]), sections 5.8.2 and 5.8.3.

The value of **E1**  $\ll$  **E2** is **E1** left-shifted **E2** bit positions; vacated bits are zero-filled. If **E1** has an unsigned type, the value of the result is **E1**  $\times$   $2^{E2}$ , reduced modulo one more than the maximum value representable in the result type. Otherwise, if **E1** has a signed type and non-negative value, and **E1**  $\times$   $2^{E2}$  is representable in the corresponding unsigned type of the result type, then that value, converted to the result type, is the resulting value; otherwise, the behavior is undefined.

The value of **E1**  $\gg$  **E2** is **E1** right-shifted **E2** bit positions. If **E1** has an unsigned type or if **E1** has a signed type and a non-negative value, the value of the result is the integral part of the quotient of **E1**/ $2^{E2}$ . If **E1** has a signed type and a negative value, the resulting value is implementation-defined.

## See Also

## Expressions with Binary Operators

### C++ Built-in Operators, Precedence and Associativity

# Logical AND Operator: &&

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
expression  
&&  
expression
```

## Remarks

The logical AND operator (**&&**) returns the boolean value **true** if both operands are **true** and returns **false** otherwise. The operands are implicitly converted to type `bool` prior to evaluation, and the result is of type `bool`. Logical AND has left-to-right associativity.

The operands to the logical AND operator need not be of the same type, but they must be of integral or pointer type. The operands are commonly relational or equality expressions.

The first operand is completely evaluated and all side effects are completed before continuing evaluation of the logical AND expression.

The second operand is evaluated only if the first operand evaluates to true (nonzero). This evaluation eliminates needless evaluation of the second operand when the logical AND expression is false. You can use this short-circuit evaluation to prevent null-pointer dereferencing, as shown in the following example:

```
char *pch = 0;  
...  
(pch) && (*pch = 'a');
```

If `pch` is null (0), the right side of the expression is never evaluated. Therefore, the assignment through a null pointer is impossible.

## Operator Keyword for &&

The **and** operator is the text equivalent of **&&**. There are two ways to access the **and** operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Logical_AND_Operator.cpp
// compile with: /EHsc
// Demonstrate logical AND
#include <iostream>

using namespace std;

int main() {
    int a = 5, b = 10, c = 15;
    cout << boolalpha
        << "The true expression "
        << "a < b && b < c yields "
        << (a < b && b < c) << endl
        << "The false expression "
        << "a > b && b < c yields "
        << (a > b && b < c) << endl;
}
```

## See Also

[C++ Built-in Operators Precedence and Associativity](#) [C++ Built-in Operators, Precedence and Associativity](#)  
[C Logical Operators](#)

# Logical Negation Operator: !

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
! cast-expression
```

## Remarks

The logical negation operator (!) reverses the meaning of its operand. The operand must be of arithmetic or pointer type (or an expression that evaluates to arithmetic or pointer type). The operand is implicitly converted to type `bool`. The result is **true** if the converted operand is **false**; the result is **false** if the converted operand is **true**. The result is of type `bool`.

For an expression `e`, the unary expression `!e` is equivalent to the expression `(e == 0)`, except where overloaded operators are involved.

## Operator Keyword for !

The **not** operator is the text equivalent of `!`. There are two ways to access the **not** operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Logical_NOT_Operator.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

int main() {
    int i = 0;
    if (!i)
        cout << "i is zero" << endl;
}
```

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Unary Arithmetic Operators](#)

# Logical OR Operator: ||

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
logical-or-expression  
||  
logical-and-expression
```

## Remarks

The logical OR operator (`||`) returns the boolean value **true** if either or both operands is **true** and returns **false** otherwise. The operands are implicitly converted to type `bool` prior to evaluation, and the result is of type `bool`. Logical OR has left-to-right associativity.

The operands to the logical OR operator need not be of the same type, but they must be of integral or pointer type. The operands are commonly relational or equality expressions.

The first operand is completely evaluated and all side effects are completed before continuing evaluation of the logical OR expression.

The second operand is evaluated only if the first operand evaluates to false (0). This eliminates needless evaluation of the second operand when the logical OR expression is true.

```
printf( "%d" , (x == w || x == y || x == z) );
```

In the above example, if `x` is equal to either `w`, `y`, or `z`, the second argument to the `printf` function evaluates to true and the value 1 is printed. Otherwise, it evaluates to false and the value 0 is printed. As soon as one of the conditions evaluates to true, evaluation ceases.

## Operator Keyword for ||

The **or** operator is the text equivalent of `||`. There are two ways to access the **or** operator in your programs: include the header file `iso646.h`, or compile with the `/Za` (Disable language extensions) compiler option.

## Example

```
// expre_Logical_OR_Operator.cpp
// compile with: /EHsc
// Demonstrate logical OR
#include <iostream>
using namespace std;
int main() {
    int a = 5, b = 10, c = 15;
    cout << boolalpha
        << "The true expression "
        << "a < b || b > c yields "
        << (a < b || b > c) << endl
        << "The false expression "
        << "a > b || b > c yields "
        << (a > b || b > c) << endl;
}
```

## See Also

[C++ Built-in Operators Precedence and Associativity](#) [C++ Built-in Operators, Precedence and Associativity](#)  
[C Logical Operators](#)

# Member Access Operators: . and ->

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
postfix-expression . name  
postfix-expression -> name
```

## Remarks

The member access operators `.` and `->` are used to refer to members of structures, unions, and classes. Member access expressions have the value and type of the selected member.

There are two forms of member access expressions:

1. In the first form, *postfix-expression* represents a value of struct, class, or union type, and *name* names a member of the specified structure, union, or class. The value of the operation is that of *name* and is an l-value if *postfix-expression* is an l-value.
2. In the second form, *postfix-expression* represents a pointer to a structure, union, or class, and *name* names a member of the specified structure, union, or class. The value is that of *name* and is an l-value. The `->` operator dereferences the pointer. Therefore, the expressions `e-> member` and `(*e). member` (where `*e` represents a pointer) yield identical results (except when the operators `->` or `*` are overloaded).

## Example

The following example demonstrates both forms of the member access operator.

```
// expr_Selection_Operator.cpp  
// compile with: /EHsc  
#include <iostream>  
using namespace std;  
  
struct Date {  
    Date(int i, int j, int k) : day(i), month(j), year(k){}  
    int month;  
    int day;  
    int year;  
};  
  
int main() {  
    Date mydate(1,1,1900);  
    mydate.month = 2;  
    cout << mydate.month << "/" << mydate.day  
        << "/" << mydate.year << endl;  
  
    Date *mydate2 = new Date(1,1,2000);  
    mydate2->month = 2;  
    cout << mydate2->month << "/" << mydate2->day  
        << "/" << mydate2->year << endl;  
    delete mydate2;  
}
```

2/1/1900

2/1/2000

## See Also

[Postfix Expressions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Classes and Structs](#)

[Structure and Union Members](#)

# Multiplicative Operators and the Modulus Operator

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
expression * expression  
expression / expression  
expression % expression
```

## Remarks

The multiplicative operators are:

- Multiplication (`\*`)
- Division (`/`)
- Modulus (remainder from division) (`\%`)

These binary operators have left-to-right associativity.

The multiplicative operators take operands of arithmetic types. The modulus operator (`\%`) has a stricter requirement in that its operands must be of integral type. (To get the remainder of a floating-point division, use the run-time function, `fmod`.) The conversions covered in [Standard Conversions](#) are applied to the operands, and the result is of the converted type.

The multiplication operator yields the result of multiplying the first operand by the second.

The division operator yields the result of dividing the first operand by the second.

The modulus operator yields the remainder given by the following expression, where  $e1$  is the first operand and  $e2$  is the second:  $e1 - (e1 / e2) * e2$ , where both operands are of integral types.

Division by 0 in either a division or a modulus expression is undefined and causes a run-time error.

Therefore, the following expressions generate undefined, erroneous results:

```
i % 0  
f / 0.0
```

If both operands to a multiplication, division, or modulus expression have the same sign, the result is positive. Otherwise, the result is negative. The result of a modulus operation's sign is implementation-defined.

### NOTE

Since the conversions performed by the multiplicative operators do not provide for overflow or underflow conditions, information may be lost if the result of a multiplicative operation cannot be represented in the type of the operands after conversion.

## Microsoft Specific

In Microsoft C++, the result of a modulus expression is always the same as the sign of the first operand.

## END Microsoft Specific

If the computed division of two integers is inexact and only one operand is negative, the result is the largest integer (in magnitude, disregarding the sign) that is less than the exact value the division operation would yield. For example, the computed value of  $-11 / 3$  is  $-3.666666666$ . The result of that integral division is  $-3$ .

The relationship between the multiplicative operators is given by the identity  $(e1 / e2) * e2 + e1 \% e2 == e1$ .

## Example

The following program demonstrates the multiplicative operators. Note that either operand of `10 / 3` must be explicitly cast to type `float` to avoid truncation so that both operands are of type `float` before division.

```
// expre_Multiplicative_Operators.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
int main() {
    int x = 3, y = 6, z = 10;
    cout << "3 * 6 is " << x * y << endl
        << "6 / 3 is " << y / x << endl
        << "10 % 3 is " << z % x << endl
        << "10 / 3 is " << (float) z / x << endl;
}
```

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Multiplicative Operators](#)

# new Operator (C++)

12/21/2017 • 7 min to read • [Edit Online](#)

Allocates memory for an object or array of objects of *type-name* from the free store and returns a suitably typed, nonzero pointer to the object.

## NOTE

Microsoft C++ Component Extensions provides support for the `new` keyword to add vtable slot entries. For more information, see [new \(new slot in vtable\)](#)

## Syntax

```
[::] new [placement] new-type-name [new-initializer]  
[::] new [placement] ( type-name ) [new-initializer]
```

## Remarks

If unsuccessful, **new** returns zero or throws an exception; see [The new and delete Operators](#) for more information. You can change this default behavior by writing a custom exception-handling routine and calling the [\\_set\\_new\\_handler](#) run-time library function with your function name as its argument.

For information on how to create an object on the managed heap, see [gcnew](#).

When **new** is used to allocate memory for a C++ class object, the object's constructor is called after the memory is allocated.

Use the **delete** operator to deallocate the memory allocated with the **new** operator.

The following example allocates and then frees a two-dimensional array of characters of size `dim` by 10. When allocating a multidimensional array, all dimensions except the first must be constant expressions that evaluate to positive values; the leftmost array dimension can be any expression that evaluates to a positive value. When allocating an array using the **new** operator, the first dimension can be zero — the **new** operator returns a unique pointer.

```
char (*pchar)[10] = new char[dim][10];  
delete [] pchar;
```

The *type-name* cannot contain **const**, **volatile**, class declarations, or enumeration declarations. Therefore, the following expression is illegal:

```
volatile char *vch = new volatile char[20];
```

The **new** operator does not allocate reference types because they are not objects.

The **new** operator cannot be used to allocate a function, but it can be used to allocate pointers to functions. The following example allocates and then frees an array of seven pointers to functions that return integers.

```
int (**p) () = new (int (*[7]) ());
delete *p;
```

If you use the operator **new** without any extra arguments, and compile with the [/GX](#), [/EHs](#), or [/EHs](#) option, the compiler will generate code to call operator **delete** if the constructor throws an exception.

The following list describes the grammar elements of **new**:

*placement*

Provides a way of passing additional arguments if you overload **new**.

*type-name*

Specifies type to be allocated; it can be either a built-in or user-defined type. If the type specification is complicated, it can be surrounded by parentheses to force the order of binding.

*initializer*

Provides a value for the initialized object. Initializers cannot be specified for arrays. The **new** operator will create arrays of objects only if the class has a default constructor.

## Example

The following code example allocates a character array and an object of class `CName` and then frees them.

```
// expre_new_Operator.cpp
// compile with: /EHsc
#include <string.h>

class CName {
public:
    enum {
        sizeOfBuffer = 256
    };

    char m_szFirst[sizeOfBuffer];
    char m_szLast[sizeOfBuffer];

public:
    void SetName(char* pszFirst, char* pszLast) {
        strcpy_s(m_szFirst, sizeOfBuffer, pszFirst);
        strcpy_s(m_szLast, sizeOfBuffer, pszLast);
    }
};

int main() {
    // Allocate memory for the array
    char* pCharArray = new char[CName::sizeOfBuffer];
    strcpy_s(pCharArray, CName::sizeOfBuffer, "Array of characters");

    // Deallocate memory for the array
    delete [] pCharArray;
    pCharArray = NULL;

    // Allocate memory for the object
    CName* pName = new CName;
    pName->SetName("Firstname", "Lastname");

    // Deallocate memory for the object
    delete pName;
    pName = NULL;
}
```

## Example

If you use the placement new form of the **new** operator, the form with arguments in addition to the size of the allocation, the compiler does not support a placement form of the **delete** operator if the constructor throws an exception. For example:

```
// expre_new_Operator2.cpp
// C2660 expected
class A {
public:
    A(int) { throw "Fail!"; }
};

void F(void) {
    try {
        // heap memory pointed to by pa1 will be deallocated
        // by calling ::operator delete(void*).
        A* pa1 = new A(10);
    } catch (...) {
    }
    try {
        // This will call ::operator new(size_t, char*, int).
        // When A::A(int) does a throw, we should call
        // ::operator delete(void*, char*, int) to deallocate
        // the memory pointed to by pa2. Since
        // ::operator delete(void*, char*, int) has not been implemented,
        // memory will be leaked when the deallocation cannot occur.

        A* pa2 = new(__FILE__, __LINE__) A(20);
    } catch (...) {
    }
}

int main() {
    A a;
}
```

## Initializing object allocated with new

An optional *initializer* field is included in the grammar for the **new** operator. This allows new objects to be initialized with user-defined constructors. For more information about how initialization is done, see [Initializers](#). The following example illustrates how to use an initialization expression with the **new** operator:

```
// expre_Initializing_Objects_Allocated_with_new.cpp
class Acct
{
public:
    // Define default constructor and a constructor that accepts
    // an initial balance.
    Acct() { balance = 0.0; }
    Acct( double init_balance ) { balance = init_balance; }

private:
    double balance;
};

int main()
{
    Acct *CheckingAcct = new Acct;
    Acct *SavingsAcct = new Acct ( 34.98 );
    double *HowMuch = new double ( 43.0 );
    // ...
}
```

In this example, the object `CheckingAcct` is allocated using the `new` operator, but no default initialization is specified. Therefore, the default constructor for the class, `Acct()`, is called. Then the object `SavingsAcct` is allocated the same way, except that it is explicitly initialized to 34.98. Because 34.98 is of type `double`, the constructor that takes an argument of that type is called to handle the initialization. Finally, the nonclass type `HowMuch` is initialized to 43.0.

If an object is of a class type and that class has constructors (as in the preceding example), the object can be initialized by the `new` operator only if one of these conditions is met:

- The arguments provided in the initializer agree with those of a constructor.
- The class has a default constructor (a constructor that can be called with no arguments).

No explicit per-element initialization can be done when allocating arrays using the `new` operator; only the default constructor, if present, is called. See [Default Arguments](#) for more information.

If the memory allocation fails (`operator new` returns a value of 0), no initialization is performed. This protects against attempts to initialize data that does not exist.

As with function calls, the order in which initialized expressions are evaluated is not defined. Furthermore, you should not rely on these expressions being completely evaluated before the memory allocation is performed. If the memory allocation fails and the `new` operator returns zero, some expressions in the initializer may not be completely evaluated.

## Lifetime of objects allocated with new

Objects allocated with the `new` operator are not destroyed when the scope in which they are defined is exited. Because the `new` operator returns a pointer to the objects it allocates, the program must define a pointer with suitable scope to access those objects. For example:

```
// expre_Lifetime_of_Objects_Allocated_with_new.cpp
// C2541 expected
int main()
{
    // Use new operator to allocate an array of 20 characters.
    char *AnArray = new char[20];

    for( int i = 0; i < 20; ++i )
    {
        // On the first iteration of the loop, allocate
        // another array of 20 characters.
        if( i == 0 )
        {
            char *AnotherArray = new char[20];
        }
    }

    delete [] AnotherArray; // Error: pointer out of scope.
    delete [] AnArray;     // OK: pointer still in scope.
}
```

Once the pointer `AnotherArray` goes out of scope in the example, the object can no longer be deleted.

## How new works

The *allocation-expression* — the expression containing the `new` operator — does three things:

- Locates and reserves storage for the object or objects to be allocated. When this stage is complete, the correct amount of storage is allocated, but it is not yet an object.

- Initializes the object(s). Once initialization is complete, enough information is present for the allocated storage to be an object.
- Returns a pointer to the object(s) of a pointer type derived from *new-type-name* or *type-name*. The program uses this pointer to access the newly allocated object.

The **new** operator invokes the function `operator new`. For arrays of any type, and for objects that are not of **class**, **struct**, or **union** types, a global function, `::operator new`, is called to allocate storage. Class-type objects can define their own `operator new` static member function on a per-class basis.

When the compiler encounters the **new** operator to allocate an object of type `type`, it issues a call to `type ::operator new( sizeof( type ) )` or, if no user-defined `operator new` is defined, `::operator new( sizeof( type ) )`. Therefore, the **new** operator can allocate the correct amount of memory for the object.

#### **NOTE**

The argument to `operator new` is of type **size\_t**. This type is defined in DIRECT.H, MALLOC.H, MEMORY.H, SEARCH.H, STDDEF.H, STUDIO.H, STDLIB.H, STRING.H, and TIME.H.

An option in the grammar allows specification of *placement* (see the Grammar for [new Operator](#)). The *placement* parameter can be used only for user-defined implementations of `operator new`; it allows extra information to be passed to `operator new`. An expression with a *placement* field such as `T *TObject = new ( 0x0040 ) T;` is translated to `T *TObject = T::operator new( sizeof( T ), 0x0040 );` if class T has member `operator new`, otherwise to `T *TObject = ::operator new( sizeof( T ), 0x0040 );`.

The original intention of the *placement* field was to allow hardware-dependent objects to be allocated at user-specified addresses.

#### **NOTE**

Although the preceding example shows only one argument in the *placement* field, there is no restriction on how many extra arguments can be passed to `operator new` this way.

Even when `operator new` has been defined for a class type, the global operator can be used by using the form of this example:

```
T *TObject = ::new TObject;
```

The scope-resolution operator (`::`) forces use of the global **new** operator.

## See Also

[Expressions with Unary Operators](#)

[Keywords](#)

[new and delete operators](#)

# One's Complement Operator: ~

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
~ cast-expression
```

## Remarks

The one's complement operator (`~`), sometimes called the "bitwise complement" operator, yields a bitwise one's complement of its operand. That is, every bit that is 1 in the operand is 0 in the result. Conversely, every bit that is 0 in the operand is 1 in the result. The operand to the one's complement operator must be an integral type.

## Operator Keyword for ~

The `compl` operator is the text equivalent of `~`. There are two ways to access the `compl` operator in your programs: include the header file `iso646.h`, or compile with `/Za`.

## Example

```
// expre_One_Complement_Operator.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;

int main () {
    unsigned short y = 0xFFFF;
    cout << hex << y << endl;
    y = ~y;    // Take one's complement
    cout << hex << y << endl;
}
```

In this example, the new value assigned to `y` is the one's complement of the unsigned value `0xFFFF`, or `0x0000`.

Integral promotion is performed on integral operands, and the resultant type is the type to which the operand is promoted. See [Standard Conversions](#) for more information on how the promotion is done.

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Unary Arithmetic Operators](#)

# Pointer-to-Member Operators: `.*` and `->*`

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
expression .* expression
expression ->* expression
```

## Remarks

The pointer-to-member operators, `*` and `->*`, return the value of a specific class member for the object specified on the left side of the expression. The right side must specify a member of the class. The following example shows how to use these operators:

```
// expre_Expressions_with_Pointer_Member_Operators.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;

class Testpm {
public:
    void m_func1() { cout << "m_func1\n"; }
    int m_num;
};

// Define derived types pmfn and pmd.
// These types are pointers to members m_func1() and
// m_num, respectively.
void (Testpm::*pmfn)() = &Testpm::m_func1;
int Testpm::*pmd = &Testpm::m_num;

int main() {
    Testpm ATestpm;
    Testpm *pTestpm = new Testpm;

    // Access the member function
    (ATestpm.*pmfn)();
    (pTestpm->*pmfn)(); // Parentheses required since * binds
                          // less tightly than the function call.

    // Access the member data
    ATestpm.*pmd = 1;
    pTestpm->*pmd = 2;

    cout << ATestpm.*pmd << endl
        << pTestpm->*pmd << endl;
    delete pTestpm;
}
```

## Output

```
m_func1
m_func1
1
2
```

In the preceding example, a pointer to a member, `pmfn`, is used to invoke the member function `m_func1`. Another pointer to a member, `pmd`, is used to access the `m_num` member.

The binary operator `.*` combines its first operand, which must be an object of class type, with its second operand, which must be a pointer-to-member type.

The binary operator `->*` combines its first operand, which must be a pointer to an object of class type, with its second operand, which must be a pointer-to-member type.

In an expression containing the `.*` operator, the first operand must be of the class type of, and be accessible to, the pointer to member specified in the second operand or of an accessible type unambiguously derived from and accessible to that class.

In an expression containing the `->*` operator, the first operand must be of the type "pointer to the class type" of the type specified in the second operand, or it must be of a type unambiguously derived from that class.

## Example

Consider the following classes and program fragment:

```
// expre_Expressions_with_Pointer_Member_Operators2.cpp
// C2440 expected
class BaseClass {
public:
    BaseClass(); // Base class constructor.
    void Func1();
};

// Declare a pointer to member function Func1.
void (BaseClass::*pmfnFunc1)() = &BaseClass::Func1;

class Derived : public BaseClass {
public:
    Derived(); // Derived class constructor.
    void Func2();
};

// Declare a pointer to member function Func2.
void (Derived::*pmfnFunc2)() = &Derived::Func2;

int main() {
    BaseClass ABase;
    Derived ADerived;

    (ABase.*pmfnFunc1)(); // OK: defined for BaseClass.
    (ABase.*pmfnFunc2)(); // Error: cannot use base class to
                          // access pointers to members of
                          // derived classes.

    (ADerived.*pmfnFunc1)(); // OK: Derived is unambiguously
                            // derived from BaseClass.
    (ADerived.*pmfnFunc2)(); // OK: defined for Derived.
}
```

The result of the `.*` or `->*` pointer-to-member operators is an object or function of the type specified in the declaration of the pointer to member. So, in the preceding example, the result of the expression

`ADerived.*pmfnFunc1()` is a pointer to a function that returns void. This result is an l-value if the second operand is an l-value.

**NOTE**

If the result of one of the pointer-to-member operators is a function, then the result can be used only as an operand to the function call operator.

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

# Postfix Increment and Decrement Operators: `++` and `--`

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
postfix-expression ++
postfix-expression --
```

## Remarks

C++ provides prefix and postfix increment and decrement operators; this section describes only the postfix increment and decrement operators. (For more information, see [Prefix Increment and Decrement Operators](#).) The difference between the two is that in the postfix notation, the operator appears after *postfix-expression*, whereas in the prefix notation, the operator appears before *expression*. The following example shows a postfix-increment operator:

```
i++;
```

The effect of applying the postfix increment operator (`++`) is that the operand's value is increased by one unit of the appropriate type. Similarly, the effect of applying the postfix decrement operator (`--`) is that the operand's value is decreased by one unit of the appropriate type.

It is important to note that a postfix increment or decrement expression evaluates to the value of the expression **prior to** application of the respective operator. The increment or decrement operation occurs **after** the operand is evaluated. This issue arises only when the postfix increment or decrement operation occurs in the context of a larger expression.

When a postfix operator is applied to a function argument, the value of the argument is not guaranteed to be incremented or decremented before it is passed to the function. See section 1.9.17 in the C++ standard for more information.

Applying the postfix increment operator to a pointer to an array of objects of type **long** actually adds four to the internal representation of the pointer. This behavior causes the pointer, which previously referred to the *n*th element of the array, to refer to the *(n+1)*th element.

The operands to postfix increment and postfix decrement operators must be modifiable (not **const**) l-values of arithmetic or pointer type. The type of the result is the same as that of the *postfix-expression*, but it is no longer an l-value.

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): The operand of a postfix increment or decrement operator may not be of type `bool`.

The following code illustrates the postfix increment operator:

```
// expre_Postfix_Increment_and_Decrement_Operators.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

int main() {
    int i = 10;
    cout << i++ << endl;
    cout << i << endl;
}
```

Postincrement and postdecrement operations on enumerated types are not supported:

```
enum Compass { North, South, East, West };
Compass myCompass;
for( myCompass = North; myCompass != West; myCompass++ ) // Error
```

## See Also

[Postfix Expressions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Postfix Increment and Decrement Operators](#)

# Prefix Increment and Decrement Operators: `++` and `--`

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
++ unary-expression  
-- unary-expression
```

## Remarks

The prefix increment operator (`++`) adds one to its operand; this incremented value is the result of the expression. The operand must be an l-value not of type **const**. The result is an l-value of the same type as the operand.

The prefix decrement operator (`--`) is analogous to the prefix increment operator, except that the operand is decremented by one and the result is this decremented value.

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): The operand of an increment or decrement operator may not be of type `bool`.

Both the prefix and postfix increment and decrement operators affect their operands. The key difference between them is the order in which the increment or decrement takes place in the evaluation of an expression. (For more information, see [Postfix Increment and Decrement Operators](#).) In the prefix form, the increment or decrement takes place before the value is used in expression evaluation, so the value of the expression is different from the value of the operand. In the postfix form, the increment or decrement takes place after the value is used in expression evaluation, so the value of the expression is the same as the value of the operand. For example, the following program prints "`++i = 6`":

```
// expre_Increment_and_Decrement_Operators.cpp  
// compile with: /EHsc  
#include <iostream>  
  
using namespace std;  
  
int main() {  
    int i = 5;  
    cout << "++i = " << ++i << endl;  
}
```

An operand of integral or floating type is incremented or decremented by the integer value 1. The type of the result is the same as the operand type. An operand of pointer type is incremented or decremented by the size of the object it addresses. An incremented pointer points to the next object; a decremented pointer points to the previous object.

Because increment and decrement operators have side effects, using expressions with increment or decrement operators in a [preprocessor macro](#) can have undesirable results. Consider this example:

```
// expre_Increment_and_Decrement_Operators2.cpp
#define max(a,b) ((a)<(b))?(b):(a)

int main()
{
    int i = 0, j = 0, k;
    k = max( ++i, j );
}
```

The macro expands to:

```
k = (((++i)<(j))?(j):(++i));
```

If `i` is greater than or equal to `j` or less than `j` by 1, it will be incremented twice.

#### NOTE

C++ inline functions are preferable to macros in many cases because they eliminate side effects such as those described here, and allow the language to perform more complete type checking.

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Prefix Increment and Decrement Operators](#)

# Relational Operators: <, >, <=, and >=

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
expression < expression
expression > expression
expression <= expression
expression >= expression
```

## Remarks

The binary relational operators determine the following relationships:

- Less than (<)
- Greater than (>)
- Less than or equal to (<=)
- Greater than or equal to (>=)

The relational operators have left-to-right associativity. Both operands of relational operators must be of arithmetic or pointer type. They yield values of type `bool`. The value returned is **false** (0) if the relationship in the expression is false; otherwise, the value returned is **true** (1).

## Example

```
// expre_Relational_Operators.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;

int main() {
    cout << "The true expression 3 > 2 yields: "
        << (3 > 2) << endl
        << "The false expression 20 < 10 yields: "
        << (20 < 10) << endl;
}
```

The expressions in the preceding example must be enclosed in parentheses because the stream insertion operator (`<<`) has higher precedence than the relational operators. Therefore, the first expression without the parentheses would be evaluated as:

```
(cout << "The true expression 3 > 2 yields: " << 3) < (2 << "\n");
```

The usual arithmetic conversions covered in [Standard Conversions](#) are applied to operands of arithmetic types.

## Comparing pointers

When two pointers to objects of the same type are compared, the result is determined by the location of the

objects pointed to in the program's address space. Pointers can also be compared to a constant expression that evaluates to 0 or to a pointer of type void \*. If a pointer comparison is made against a pointer of type void \*, the other pointer is implicitly converted to type void \*. Then the comparison is made.

Two pointers of different types cannot be compared unless:

- One type is a class type derived from the other type.
- At least one of the pointers is explicitly converted (cast) to type void \*. (The other pointer is implicitly converted to type void \* for the conversion.)

Two pointers of the same type that point to the same object are guaranteed to compare equal. If two pointers to nonstatic members of an object are compared, the following rules apply:

- If the class type is not a union, and if the two members are not separated by an *access-specifier*, such as public, protected, or private, the pointer to the member declared last will compare greater than the pointer to the member declared earlier.
- If the two members are separated by an *access-specifier*, the results are undefined.
- If the class type is a union, pointers to different data members in that union compare equal.

If two pointers point to elements of the same array or to the element one beyond the end of the array, the pointer to the object with the higher subscript compares higher. Comparison of pointers is guaranteed valid only when the pointers refer to objects in the same array or to the location one past the end of the array.

## See Also

[Expressions with Binary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[C Relational and Equality Operators](#)

# Scope Resolution Operator: ::

12/21/2017 • 1 min to read • [Edit Online](#)

The scope resolution operator `::` is used to identify and disambiguate identifiers used in different scopes. For more information about scope, see [Scope](#).

## Syntax

```
:: identifier  
class-name :: identifier  
namespace :: identifier  
enum class :: identifier  
enum struct :: identifier
```

## Remarks

The `identifier` can be a variable, a function, or an enumeration value.

## With Classes and Namespaces

The following example shows how the scope resolution operator is used with namespaces and classes:

```
namespace NamespaceA{  
    int x;  
    class ClassA {  
        public:  
            int x;  
    };  
}  
  
int main() {  
  
    // A namespace name used to disambiguate  
    NamespaceA::x = 1;  
  
    // A class name used to disambiguate  
    NamespaceA::ClassA a1;  
    a1.x = 2;  
}
```

A scope resolution operator without a scope qualifier refers to the global namespace.

```

namespace NamespaceA{
    int x;
}

int x;

int main() {
    int x;

    // the x in main()
    x = 0;
    // The x in the global namespace
    ::x = 1;

    // The x in the A namespace
    NamespaceA::x = 2;
}

```

You can use the scope resolution operator to identify a member of a namespace, or to identify a namespace that nominates the member's namespace in a using-directive. In the example below, you can use `NamespaceC` to qualify `ClassB`, even though `ClassB` was declared in namespace `NamespaceB`, because `NamespaceB` was nominated in `NamespaceC` by a using directive.

```

namespace NamespaceB {
    class ClassB {
        public:
            int x;
    };
}

namespace NamespaceC{
    using namespace B;

}

int main() {
    NamespaceB::ClassB c_b;
    NamespaceC::ClassB c_c;

    c_b.x = 3;
    c_c.x = 4;
}

```

You can use chains of scope resolution operators. In the following example, `NamespaceD::NamespaceD1` identifies the nested namespace `NamespaceD1`, and `NamespaceE::ClassE::ClassE1` identifies the nested class `ClassE1`.

```

namespace NamespaceD{
    namespace NamespaceD1{
        int x;
    }
}

namespace NamespaceE{

    class ClassE{
    public:
        class ClassE1{
        public:
            int x;
        };
    };
}

int main() {
    NamespaceD:: NamespaceD1::x = 6;
    NamespaceE::ClassE::ClassE1 e1;
    e1.x = 7 ;
}

```

## With Static Members

You must use the scope resolution operator to call static members of classes.

```

class ClassG {
public:
    static int get_x() { return x; }
    static int x;
};

int ClassG::x = 6;

int main() {

    int gx1 = ClassG::x;
    int gx2 = ClassG::get_x();
}

```

## With Scoped Enumerations

The scoped resolution operator is also used with the values of a scoped enumeration [Enumeration Declarations](#), as in the following example:

```

enum class EnumA{
    First,
    Second,
    Third
};

int main() {

    EnumA enum_value = EnumA::First;
}

```

## See Also

C++ Built-in Operators, Precedence and Associativity

Namespaces

# sizeof Operator

12/21/2017 • 2 min to read • [Edit Online](#)

Yields the size of its operand with respect to the size of type `char`.

## NOTE

For information about the `sizeof ...` operator, see [Ellipses and Variadic Templates](#).

## Syntax

```
sizeof unary-expression  
sizeof ( type-name )
```

## Remarks

The result of the `sizeof` operator is of type `size_t`, an integral type defined in the include file STDDEF.H. This operator allows you to avoid specifying machine-dependent data sizes in your programs.

The operand to `sizeof` can be one of the following:

- A type name. To use `sizeof` with a type name, the name must be enclosed in parentheses.
- An expression. When used with an expression, `sizeof` can be specified with or without the parentheses. The expression is not evaluated.

When the `sizeof` operator is applied to an object of type `char`, it yields 1. When the `sizeof` operator is applied to an array, it yields the total number of bytes in that array, not the size of the pointer represented by the array identifier. To obtain the size of the pointer represented by the array identifier, pass it as a parameter to a function that uses `sizeof`. For example:

## Example

```
#include <iostream>  
using namespace std;  
  
size_t getPtrSize( char *ptr )  
{  
    return sizeof( ptr );  
}  
  
int main()  
{  
    char szHello[] = "Hello, world!";  
  
    cout << "The size of a char is: "  
        << sizeof( char )  
        << "\nThe length of " << szHello << " is: "  
        << sizeof szHello  
        << "\nThe size of the pointer is "  
        << getPtrSize( szHello ) << endl;  
}
```

## Sample Output

```
The size of a char is: 1
The length of Hello, world! is: 14
The size of the pointer is 4
```

When the `sizeof` operator is applied to a `class`, `struct`, or `union` type, the result is the number of bytes in an object of that type, plus any padding added to align members on word boundaries. The result does not necessarily correspond to the size calculated by adding the storage requirements of the individual members. The `/Zp` compiler option and the `pack` pragma affect alignment boundaries for members.

The `sizeof` operator never yields 0, even for an empty class.

The `sizeof` operator cannot be used with the following operands:

- Functions. (However, `sizeof` can be applied to pointers to functions.)
- Bit fields.
- Undefined classes.
- The type `void`.
- Dynamically allocated arrays.
- External arrays.
- Incomplete types.
- Parenthesized names of incomplete types.

When the `sizeof` operator is applied to a reference, the result is the same as if `sizeof` had been applied to the object itself.

If an unsized array is the last element of a structure, the `sizeof` operator returns the size of the structure without the array.

The `sizeof` operator is often used to calculate the number of elements in an array using an expression of the form:

```
sizeof array / sizeof array[0]
```

## See Also

[Expressions with Unary Operators](#)

[Keywords](#)

# Subscript Operator:

12/21/2017 • 3 min to read • [Edit Online](#)

## Syntax

```
postfix-expression [ expression ]
```

## Remarks

A postfix expression (which can also be a primary expression) followed by the subscript operator, `[ ]`, specifies array indexing.

For information about managed arrays, see [Arrays](#).

Usually, the value represented by *postfix-expression* is a pointer value, such as an array identifier, and *expression* is an integral value (including enumerated types). However, all that is required syntactically is that one of the expressions be of pointer type and the other be of integral type. Thus the integral value could be in the *postfix-expression* position and the pointer value could be in the brackets in the *expression* or subscript position. Consider the following code fragment:

```
int nArray[5] = { 0, 1, 2, 3, 4 };
cout << nArray[2] << endl;           // prints "2"
cout << 2[nArray] << endl;          // prints "2"
```

In the preceding example, the expression `nArray[2]` is identical to `2[nArray]`. The reason is that the result of a subscript expression `e1[e2]` is given by:

`*((e2) + (e1))`

The address yielded by the expression is not *e2* bytes from the address *e1*. Rather, the address is scaled to yield the next object in the array *e2*. For example:

```
double aDb1[2];
```

The addresses of `aDb[0]` and `aDb[1]` are 8 bytes apart — the size of an object of type **double**. This scaling according to object type is done automatically by the C++ language and is defined in [Additive Operators](#) where addition and subtraction of operands of pointer type is discussed.

A subscript expression can also have multiple subscripts, as follows:

`expression1 [expression2] [expression3]...`

Subscript expressions associate from left to right. The leftmost subscript expression, `expression1[expression2]`, is evaluated first. The address that results from adding `expression1` and `expression2` forms a pointer expression; then `expression3` is added to this pointer expression to form a new pointer expression, and so on until the last subscript expression has been added. The indirection operator (`\*`) is applied after the last subscripted expression is evaluated, unless the final pointer value addresses an array type.

Expressions with multiple subscripts refer to elements of multidimensional arrays. A multidimensional array is an array whose elements are arrays. For example, the first element of a three-dimensional array is an array with two

dimensions. The following example declares and initializes a simple two-dimensional array of characters:

```
// expre_Subscript_Operator.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
#define MAX_ROWS 2
#define MAX_COLS 2

int main() {
    char c[ MAX_ROWS ][ MAX_COLS ] = { { 'a', 'b' }, { 'c', 'd' } };
    for ( int i = 0; i < MAX_ROWS; i++ )
        for ( int j = 0; j < MAX_COLS; j++ )
            cout << c[ i ][ j ] << endl;
}
```

## Positive and negative subscripts

The first element of an array is element 0. The range of a C++ array is from *array*[0] to *array*[*size* - 1]. However, C++ supports positive and negative subscripts. Negative subscripts must fall within array boundaries; if they do not, the results are unpredictable. The following code shows positive and negative array subscripts:

```
#include <iostream>
using namespace std;

int main() {
    int intArray[1024];
    for (int i = 0, j = 0; i < 1024; i++)
    {
        intArray[i] = j++;
    }

    cout << intArray[512] << endl;// 512

    int *midArray = &intArray[512]; // pointer to the middle of the array

    cout << midArray[-256] << endl; // 256

    cout << intArray[-256] << endl; // unpredictable
}
```

The negative subscript in the lasta line can produce a run-time error because it points an address 256 bytes lower in memory than the origin of the array. The pointer `midArray` is initialized to the middle of `intArray`; it is therefore possible to use both positive and negative array indices on it. Array subscript errors do not generate compile-time errors, but they yield unpredictable results.

The subscript operator is commutative. Therefore, the expressions *array*[*index*] and *array*[*array*] are guaranteed to be equivalent as long as the subscript operator is not overloaded (see [Overloaded Operators](#)). The first form is the most common coding practice, but either works.

## See Also

[Postfix Expressions](#)

[C++ Built-in Operators, Precedence and Associativity](#)

[Arrays](#)

[One-Dimensional Arrays](#)

[Multidimensional Arrays](#)

# typeid Operator

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
typeid(  
    type-id  
)  
typeid( expression )
```

## Remarks

The `typeid` operator allows the type of an object to be determined at run time.

The result of `typeid` is a **const type\_info&**. The value is a reference to a **type\_info** object that represents either the *type-id* or the type of the *expression*, depending on which form of `typeid` is used. See [type\\_info Class](#) for more information.

The `typeid` operator does not work with managed types (abstract declarators or instances), see [typeid](#) for information on getting the [Type](#) of a specified type.

The `typeid` operator does a run-time check when applied to an l-value of a polymorphic class type, where the true type of the object cannot be determined by the static information provided. Such cases are:

- A reference to a class
- A pointer, dereferenced with \*
- A subscripted pointer (i.e. [ ]). (Note that it is generally not safe to use a subscript with a pointer to a polymorphic type.)

If the *expression* points to a base class type, yet the object is actually of a type derived from that base class, a **type\_info** reference for the derived class is the result. The *expression* must point to a polymorphic type (a class with virtual functions). Otherwise, the result is the **type\_info** for the static class referred to in the *expression*. Further, the pointer must be dereferenced so that the object it points to is used. Without dereferencing the pointer, the result will be the **type\_info** for the pointer, not what it points to. For example:

```

// expre_typeid_Operator.cpp
// compile with: /GR /EHsc
#include <iostream>
#include <typeinfo.h>

class Base {
public:
    virtual void vfunc() {}
};

class Derived : public Base {};

using namespace std;
int main() {
    Derived* pd = new Derived;
    Base* pb = pd;
    cout << typeid( pb ).name() << endl;    //prints "class Base *"
    cout << typeid( *pb ).name() << endl;    //prints "class Derived"
    cout << typeid( pd ).name() << endl;    //prints "class Derived *"
    cout << typeid( *pd ).name() << endl;    //prints "class Derived"
    delete pd;
}

```

If the *expression* is dereferencing a pointer, and that pointer's value is zero, **typeid** throws a [bad\\_typeid exception](#).

If the pointer does not point to a valid object, a [\\_\\_non\\_rtti\\_object](#) exception is thrown, indicating an attempt to analyze the RTTI that triggered a fault (like access violation), because the object is somehow invalid (bad pointer or the code wasn't compiled with [/GR](#)).

If the *expression* is neither a pointer nor a reference to a base class of the object, the result is a **type\_info** reference representing the static type of the *expression*. The *static type* of an expression refers to the type of an expression as it is known at compile time. Execution semantics are ignored when evaluating the static type of an expression. Furthermore, references are ignored when possible when determining the static type of an expression:

```

// expre_typeid_Operator_2.cpp
#include <typeinfo>

int main()
{
    typeid(int) == typeid(int&); // evaluates to true
}

```

**typeid** can also be used in templates to determine the type of a template parameter:

```

// expre_typeid_Operator_3.cpp
// compile with: /c
#include <typeinfo>
template < typename T >
T max( T arg1, T arg2 ) {
    cout << typeid( T ).name() << "s compared." << endl;
    return ( arg1 > arg2 ? arg1 : arg2 );
}

```

## See Also

[Run-Time Type Information](#)

[Keywords](#)

# Unary Plus and Negation Operators: + and -

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
+ cast-expression
```

```
- cast-expression
```

### + operator

The result of the unary plus operator (+) is the value of its operand. The operand to the unary plus operator must be of an arithmetic type.

Integral promotion is performed on integral operands. The resultant type is the type to which the operand is promoted. Thus, the expression `+ch`, where `ch` is of type `char`, results in type `int`; the value is unmodified. See [Standard Conversions](#) for more information about how the promotion is done.

### - operator

The unary negation operator (-) produces the negative of its operand. The operand to the unary negation operator must be an arithmetic type.

Integral promotion is performed on integral operands, and the resultant type is the type to which the operand is promoted. See [Standard Conversions](#) for more information about how the promotion is performed.

## Microsoft specific

Unary negation of unsigned quantities is performed by subtracting the value of the operand from  $2^n$ , where n is the number of bits in an object of the given unsigned type. (Microsoft C++ runs on processors that utilize two's-complement arithmetic. On other processors, the algorithm for negation can differ.)

## See Also

[Expressions with Unary Operators](#)

[C++ Built-in Operators, Precedence and Associativity](#)

# Expressions (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This section describes C++ expressions. Expressions are sequences of operators and operands that are used for one or more of these purposes:

- Computing a value from the operands.
- Designating objects or functions.
- Generating "side effects." (Side effects are any actions other than the evaluation of the expression — for example, modifying the value of an object.)

In C++, operators can be overloaded and their meanings can be user-defined. However, their precedence and the number of operands they take cannot be modified. This section describes the syntax and semantics of operators as they are supplied with the language, not overloaded. In addition to [types of expressions](#) and [semantics of expressions](#), the following topics are covered:

- [Primary expressions](#)
- [Scope resolution operator](#)
- [Postfix expressions](#)
- [Expressions with unary operators](#)
- [Expressions with binary operators](#)
- [Conditional operator](#)
- [Constant expressions](#)
- [Expressions with explicit type conversions](#)
- [Casting operators](#)
- [Run-time type information](#)

Topics on operators in other sections:

- [C++ Built-in Operators, Precedence and Associativity](#)
- [Overloaded operators](#)
- [typeid](#)

## NOTE

Operators for built-in types cannot be overloaded; their behavior is predefined.

## See Also

[C++ Language Reference](#)

# Types of Expressions

12/21/2017 • 1 min to read • [Edit Online](#)

C++ expressions are divided into several categories:

- [Primary expressions](#). These are the building blocks from which all other expressions are formed.
- [Postfix expressions](#). These are primary expressions followed by an operator — for example, the array subscript or postfix increment operator.
- [Expressions formed with unary operators](#). Unary operators act on only one operand in an expression.
- [Expressions formed with binary operators](#). Binary operators act on two operands in an expression.
- [Expressions with the conditional operator](#). The conditional operator is a ternary operator — the only such operator in the C++ language — and takes three operands.
- [Constant expressions](#). Constant expressions are formed entirely of constant data.
- [Expressions with explicit type conversions](#). Explicit type conversions, or "casts," can be used in expressions.
- [Expressions with pointer-to-member operators](#).
- [Casting](#). Type-safe "casts" can be used in expressions.
- [Run-Time Type Information](#). Determine the type of an object during program execution.

## See Also

[Expressions](#)

# Primary Expressions

12/21/2017 • 2 min to read • [Edit Online](#)

Primary expressions are the building blocks of more complex expressions. They are literals, names, and names qualified by the scope-resolution operator ( `::` ). A primary expression may have any of the following forms:

```
literal  
this  
:: namename( expression )
```

A *literal* is a constant primary expression. Its type depends on the form of its specification. See [Literals](#) for complete information about specifying literals.

The **this** keyword is a pointer to a class object. It is available within nonstatic member functions and points to the instance of the class for which the function was invoked. The **this** keyword cannot be used outside the body of a class-member function.

The type of the **this** pointer is `type *const` (where `type` is the class name) within functions not specifically modifying the **this** pointer. The following example shows member function declarations and the types of **this**:

```
// expre_Primary_Expressions.cpp  
// compile with: /LD  
class Example  
{  
public:  
    void Func();           // * const this  
    void Func() const;    // const * const this  
    void Func() volatile; // volatile * const this  
};
```

See [this Pointer](#) for more information about modifying the type of the **this** pointer.

The scope-resolution operator ( `::` ) followed by a name constitutes a primary expression. Such names must be names at global scope, not member names. The type of this expression is determined by the declaration of the name. It is an l-value (that is, it can appear on the left hand side of an assignment operator expression) if the declaring name is an l-value. The scope-resolution operator allows a global name to be referred to, even if that name is hidden in the current scope. See [Scope](#) for an example of how to use the scope-resolution operator.

An expression enclosed in parentheses is a primary expression whose type and value are identical to those of the unparenthesized expression. It is an l-value if the unparenthesized expression is an l-value.

In the context of the primary expression syntax given above, *name* means anything in the syntax described for [Names](#), although when using the scope-resolution operator before the name, types of names that can only occur in a class are not allowed. This includes user-defined conversion function names, and destructor names.

Examples of primary expressions include:

```
100 // literal
'c' // literal
this // in a member function, a pointer to the class instance
::func // a global function
::operator + // a global operator function
::A::B // a global qualified name
( i + 1 ) // a parenthesized expression
```

The examples below are all considered *names*, and hence primary expressions, in various forms:

```
MyClass // a identifier
MyClass::f // a qualified name
operator = // an operator function name
operator char* // a conversion operator function name
~MyClass // a destructor name
A::B // a qualified name
A<int> // a template id
```

## See Also

[Types of Expressions](#)

# Ellipses and Variadic Templates

12/21/2017 • 3 min to read • [Edit Online](#)

This article shows how to use the ellipsis (`...`) with C++ variadic templates. The ellipsis has had many uses in C and C++. These include variable argument lists for functions. The `printf()` function from the C Runtime Library is one of the most well-known examples.

A *variadic template* is a class or function template that supports an arbitrary number of arguments. This mechanism is especially useful to C++ library developers because you can apply it to both class templates and function templates, and thereby provide a wide range of type-safe and non-trivial functionality and flexibility.

## Syntax

An ellipsis is used in two ways by variadic templates. To the left of the parameter name, it signifies a *parameter pack*, and to the right of the parameter name, it expands the parameter packs into separate names.

Here's a basic example of *variadic template class definition* syntax:

```
template<typename... Arguments> class classname;
```

For both parameter packs and expansions, you can add whitespace around the ellipsis, based on your preference, as shown in these examples:

```
template<typename ...Arguments> class classname;
```

Or this:

```
template<typename ... Arguments> class classname;
```

Notice that this article uses the convention that's shown in the first example (the ellipsis is attached to `typename`).

In the preceding examples, `Arguments` is a parameter pack. The class `classname` can accept a variable number of arguments, as in these examples:

```
template<typename... Arguments> class vtclass;  
  
vtclass< > vtinstance1;  
vtclass<int> vtinstance2;  
vtclass<float, bool> vtinstance3;  
vtclass<long, std::vector<int>, std::string> vtinstance4;
```

By using a variadic template class definition, you can also require at least one parameter:

```
template <typename First, typename... Rest> class classname;
```

Here's a basic example of *variadic template function* syntax:

```
template <typename... Arguments> returntype functionname(Arguments... args);
```

The `Arguments` parameter pack is then expanded for use, as shown in the next section, [Understanding variadic templates](#).

Other forms of variadic template function syntax are possible—including, but not limited to, these examples:

```
template <typename... Arguments> returntype functionname(Arguments&... args);
template <typename... Arguments> returntype functionname(Arguments&&... args);
template <typename... Arguments> returntype functionname(Arguments*... args);
```

Specifiers like `const` are also allowed:

```
template <typename... Arguments> returntype functionname(const Arguments&... args);
```

As with variadic template class definitions, you can make functions that require at least one parameter:

```
template <typename First, typename... Rest> returntype functionname(const First& first, const Rest&... args);
```

Variadic templates use the `sizeof...()` operator (unrelated to the older `sizeof()` operator):

```
template<typename... Arguments>
void tfunc(const Arguments&... args)
{
    constexpr auto numargs{ sizeof...(Arguments) };

    X xobj[numargs]; // array of some previously defined type X

    helper_func(xobj, args...);
}
```

## More about ellipsis placement

Previously, this article described ellipsis placement that defines parameter packs and expansions as "to the left of the parameter name, it signifies a parameter pack, and to the right of the parameter name, it expands the parameter packs into separate names". This is technically true but can be confusing in translation to code. Consider:

- In a template-parameter-list (`template <parameter-list>`), `typename...` introduces a template parameter pack.
- In a parameter-declaration-clause (`func(parameter-list)`), a "top-level" ellipsis introduces a function parameter pack, and the ellipsis positioning is important:

```
// v1 is NOT a function parameter pack:
template <typename... Types> void func1(std::vector<Types...> v1);

// v2 IS a function parameter pack:
template <typename... Types> void func2(std::vector<Types>... v2);
```

- Where the ellipsis appears immediately after a parameter name, you have a parameter pack expansion.

## Example

A good way to illustrate the variadic template function mechanism is to use it in a re-write of some of the functionality of `printf`:

```

#include <iostream>

using namespace std;

void print() {
    cout << endl;
}

template <typename T> void print(const T& t) {
    cout << t << endl;
}

template <typename First, typename... Rest> void print(const First& first, const Rest&... rest) {
    cout << first << ", ";
    print(rest...); // recursive call using pack expansion syntax
}

int main()
{
    print(); // calls first overload, outputting only a newline
    print(1); // calls second overload

    // these call the third overload, the variadic template,
    // which uses recursion as needed.
    print(10, 20);
    print(100, 200, 300);
    print("first", 2, "third", 3.14159);
}

```

## Output

```

1
10, 20
100, 200, 300
first, 2, third, 3.14159

```

### NOTE

Most implementations that incorporate variadic template functions use recursion of some form, but it's slightly different from traditional recursion. Traditional recursion involves a function calling itself by using the same signature. (It may be overloaded or templated, but the same signature is chosen each time.) Variadic recursion involves calling a variadic function template by using differing (almost always decreasing) numbers of arguments, and thereby stamping out a different signature every time. A "base case" is still required, but the nature of the recursion is different.

# Postfix Expressions

12/21/2017 • 6 min to read • [Edit Online](#)

Postfix expressions consist of primary expressions or expressions in which postfix operators follow a primary expression. The postfix operators are listed in the following table.

## Postfix Operators

OPERATOR NAME	OPERATOR NOTATION
Subscript operator	[ ]
Function call operator	( )
Explicit type conversion operator	<i>type-name</i> ( )
Member access operator	. or ->
Postfix increment operator	++
Postfix decrement operator	--

The following syntax describes possible postfix expressions:

```
primary-expression
postfix-expression[expression]postfix-expression(expression-list)simple-type-name(expression-list)postfix-
expression.namepostfix-expression->namepostfix-expression++postfix-expression--cast-keyword < typename >
(expression )typeid ( typename )
```

The *postfix-expression* above may be a primary expression or another postfix expression. See **primary expressions**. Postfix expressions group left to right, thus allowing the expressions to be chained together as follows:

```
func(1)->GetValue()++
```

In the above expression, func is a primary expression, func(1) is a function postfix expression, func(1)->GetData is a postfix expression specifying a member of the class, func(1)->GetData() is another function postfix expression, and the entire expression is a postfix expression incrementing the return value of GetData. The meaning of the expression as a whole is "call func passing 1 as an argument and get a pointer to a class as a return value. Then call GetValue() on that class, then increment the value returned."

The expressions listed above are assignment expressions, meaning that the result of these expressions must be an r-value.

The postfix expression form

```
simple-type-name ( expression-list )
```

indicates the invocation of the constructor. If the simple-type-name is a fundamental type, the expression list

must be a single expression, and this expression indicates a cast of the expression's value to the fundamental type. This type of cast expression mimics a constructor. Because this form allows fundamental types and classes to be constructed using the same syntax, this form is especially useful when defining template classes.

The *cast-keyword* is one of `dynamic_cast`, `static_cast` or `reinterpret_cast`. More information may be found in **dynamic\_cast**, **static\_cast** and **reinterpret\_cast**.

The `typeid` operator is considered a postfix expression. See **typeid operator**.

## Formal and actual arguments

Calling programs pass information to called functions in "actual arguments." The called functions access the information using corresponding "formal arguments."

When a function is called, the following tasks are performed:

- All actual arguments (those supplied by the caller) are evaluated. There is no implied order in which these arguments are evaluated, but all arguments are evaluated and all side effects completed prior to entry to the function.
- Each formal argument is initialized with its corresponding actual argument in the expression list. (A formal argument is an argument that is declared in the function header and used in the body of a function.) Conversions are done as if by initialization — both standard and user-defined conversions are performed in converting an actual argument to the correct type. The initialization performed is illustrated conceptually by the following code:

```
void Func( int i ); // Function prototype
...
Func( 7 );          // Execute function call
```

The conceptual initializations prior to the call are:

```
int Temp_i = 7;
Func( Temp_i );
```

Note that the initialization is performed as if using the equal-sign syntax instead of the parentheses syntax. A copy of `i` is made prior to passing the value to the function. (For more information, see [Initializers](#) and [Conversions](#)).

Therefore, if the function prototype (declaration) calls for an argument of type **long**, and if the calling program supplies an actual argument of type `int`, the actual argument is promoted using a standard type conversion to type **long** (see [Standard Conversions](#)).

It is an error to supply an actual argument for which there is no standard or user-defined conversion to the type of the formal argument.

For actual arguments of class type, the formal argument is initialized by calling the class's constructor. (See [Constructors](#) for more about these special class member functions.)

- The function call is executed.

The following program fragment demonstrates a function call:

```

// expre_Formal_and_Actual_Arguments.cpp
void func( long param1, double param2 );

int main()
{
    long i = 1;
    double j = 2;

    // Call func with actual arguments i and j.
    func( i, j );
}

// Define func with formal parameters param1 and param2.
void func( long param1, double param2 )
{
}

```

When `func` is called from `main`, the formal parameter `param1` is initialized with the value of `i` (`i` is converted to type **long** to correspond to the correct type using a standard conversion), and the formal parameter `param2` is initialized with the value of `j` (`j` is converted to type **double** using a standard conversion).

## Treatment of argument types

Formal arguments declared as `const` types cannot be changed within the body of a function. Functions can change any argument that is not of type **const**. However, the change is local to the function and does not affect the actual argument's value unless the actual argument was a reference to an object not of type **const**.

The following functions illustrate some of these concepts:

```

// expre_Treatment_of_Argument_Types.cpp
int func1( const int i, int j, char *c ) {
    i = 7;    // C3892 i is const.
    j = i;    // value of j is lost at return
    *c = 'a' + j;    // changes value of c in calling function
    return i;
}

double& func2( double& d, const char *c ) {
    d = 14.387;    // changes value of d in calling function.
    *c = 'a';    // C3892 c is a pointer to a const object.
    return d;
}

```

## Ellipses and default arguments

Functions can be declared to accept fewer arguments than specified in the function definition, using one of two methods: ellipsis (`...`) or default arguments.

Ellipses denote that arguments may be required but that the number and types are not specified in the declaration. This is normally poor C++ programming practice because it defeats one of the benefits of C++: type safety. Different conversions are applied to functions declared with ellipses than to those functions for which the formal and actual argument types are known:

- If the actual argument is of type **float**, it is promoted to type **double** prior to the function call.
- Any signed or unsigned `char`, `short`, enumerated type, or bit field is converted to either a signed or an unsigned `int` using integral promotion.
- Any argument of class type is passed by value as a data structure; the copy is created by binary copying

instead of by invoking the class's copy constructor (if one exists).

Ellipses, if used, must be declared last in the argument list. For more information about passing a variable number of arguments, see the discussion of [va\\_arg, va\\_start, and va\\_list](#) in the *Run-Time Library Reference*.

For information on default arguments in CLR programming, see [Variable Argument Lists \(...\) \(C++/CLI\)](#).

Default arguments enable you to specify the value an argument should assume if none is supplied in the function call. The following code fragment shows how default arguments work. For more information about restrictions on specifying default arguments, see [Default Arguments](#).

```
// expre_Ellipses_and_Default_Arguments.cpp
// compile with: /EHsc
#include <iostream>

// Declare the function print that prints a string,
// then a terminator.
void print( const char *string,
            const char *terminator = "\n" );

int main()
{
    print( "hello," );
    print( "world!" );

    print( "good morning", ", " );
    print( "sunshine." );
}

using namespace std;
// Define print.
void print( const char *string, const char *terminator )
{
    if( string != NULL )
        cout << string;

    if( terminator != NULL )
        cout << terminator;
}
```

The preceding program declares a function, `print`, that takes two arguments. However, the second argument, `terminator`, has a default value, `"\n"`. In `main`, the first two calls to `print` allow the default second argument to supply a new line to terminate the printed string. The third call specifies an explicit value for the second argument. The output from the program is

```
hello,
world!
good morning, sunshine.
```

## See Also

[Types of Expressions](#)

# Expressions with Unary Operators

12/21/2017 • 1 min to read • [Edit Online](#)

Unary operators act on only one operand in an expression. The unary operators are as follows:

- [Indirection operator \(\\*\)](#)
- [Address-of operator \(&\)](#)
- [Unary plus operator \(+\)](#)
- [Unary negation operator \(-\)](#)
- [Logical negation operator \(!\)](#)
- [One's complement operator \(~\)](#)
- [Prefix increment operator \(++\)](#)
- [Prefix decrement operator \(--\)](#)
- [Cast operator \(\)](#)
- [sizeof operator](#)
- [\\_\\_uuidof operator](#)
- [\\_\\_alignof operator](#)
- [new operator](#)
- [delete operator](#)

These operators have right-to-left associativity. Unary expressions generally involve syntax that precedes a postfix or primary expression.

The following are the possible forms of unary expressions.

- *postfix-expression*
- `[++]` *unary-expression*
- `[--]` *unary-expression*
- *unary-operator cast-expression*
- `[sizeof]` *unary-expression*
- `[sizeof(` *type-name* `)`
- `[decltype(` *expression* `)`
- *allocation-expression*
- *deallocation-expression*

Any *postfix-expression* is considered a *unary-expression*, and because any primary expression is considered a *postfix-expression*, any primary expressions is considered a *unary-expression* also. For more information, see [Postfix Expressions](#) and [Primary Expressions](#).

A *unary-operator* consists of one or more of the following symbols: `* & + - ! ~`

The *cast-expression* is a unary expression with an optional cast to change the type. For more information see [Cast Operator: \(\)](#).

An *expression* can be any expression. For more information, see [Expressions](#).

The *allocation-expression* refers to the `new` operator. The *deallocation-expression* refers to the `delete` operator. For more information, see the links earlier in this topic.

## See Also

[Types of Expressions](#)

# Expressions with Binary Operators

12/21/2017 • 1 min to read • [Edit Online](#)

Binary operators act on two operands in an expression. The binary operators are:

- [Multiplicative operators](#)
  - Multiplication (\*)
  - Division (/)
  - Modulus (%)
- [Additive operators](#)
  - Addition (+)
  - Subtraction (-)
- [Shift operators](#)
  - Right shift (>>)
  - Left shift (<<)
- [Relational and equality operators](#)
  - Less than (<)
  - Greater than (>)
  - Less than or equal to (<=)
  - Greater than or equal to (>=)
  - Equal to (==)
  - Not equal to (!=)
- Bitwise operators
  - [Bitwise AND \(&\)](#)
  - [Bitwise exclusive OR \(^\)](#)
  - [Bitwise inclusive OR \(|\)](#)
- Logical operators
  - [Logical AND \(&&\)](#)
  - [Logical OR \(||\)](#)
- [Assignment operators](#)
  - Assignment (=)
  - Addition assignment (+=)
  - Subtraction assignment (-=)

- Multiplication assignment ( $*=$ )
  - Division assignment ( $/=$ )
  - Modulus assignment ( $\%=$ )
  - Left shift assignment ( $<<=$ )
  - Right shift assignment ( $>>=$ )
  - Bitwise AND assignment ( $\&=$ )
  - Bitwise exclusive OR assignment ( $\^{}=$ )
  - Bitwise inclusive OR assignment ( $\|=$ )
- [Comma Operator \(,\)](#)

## See Also

[Types of Expressions](#)

# C++ Constant Expressions

12/21/2017 • 1 min to read • [Edit Online](#)

A *constant* value is one that doesn't change. C++ provides two keywords to enable you to express the intent that an object is not intended to be modified, and to enforce that intent.

C++ requires constant expressions — expressions that evaluate to a constant — for declarations of:

- Array bounds
- Selectors in case statements
- Bit-field length specification
- Enumeration initializers

The only operands that are legal in constant expressions are:

- Literals
- Enumeration constants
- Values declared as `const` that are initialized with constant expressions
- `sizeof` expressions

Nonintegral constants must be converted (either explicitly or implicitly) to integral types to be legal in a constant expression. Therefore, the following code is legal:

```
const double Size = 11.0;
char chArray[(int)Size];
```

Explicit conversions to integral types are legal in constant expressions; all other types and derived types are illegal except when used as operands to the `sizeof` operator.

The comma operator and assignment operators cannot be used in constant expressions.

## See Also

[Types of Expressions](#)

# Semantics of Expressions

12/21/2017 • 4 min to read • [Edit Online](#)

Expressions are evaluated according to the precedence and grouping of their operators. ([Operator Precedence and Associativity](#) in [Lexical Conventions](#), shows the relationships the C++ operators impose on expressions.)

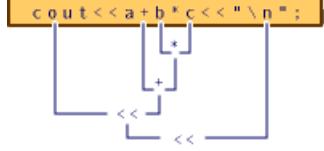
## Order of evaluation

Consider this example:

```
// Order_of_Evaluation.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
int main()
{
    int a = 2, b = 4, c = 9;

    cout << a + b * c << "\n";
    cout << a + (b * c) << "\n";
    cout << (a + b) * c << "\n";
}
```

38  
38  
54

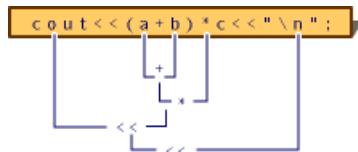


Expression-Evaluation Order

The order in which the expression shown in the above figure is evaluated is determined by the precedence and associativity of the operators:

1. Multiplication (\*) has the highest precedence in this expression; hence the subexpression `b * c` is evaluated first.
2. Addition (+) has the next highest precedence, so `a` is added to the product of `b` and `c`.
3. Left shift (<<) has the lowest precedence in the expression, but there are two occurrences. Because the left-shift operator groups left-to-right, the left subexpression is evaluated first and then the right one.

When parentheses are used to group the subexpressions, they alter the precedence and also the order in which the expression is evaluated, as shown in the following figure.



Expression-Evaluation Order with Parentheses

Expressions such as those in the above figure are evaluated purely for their side effects — in this case, to transfer information to the standard output device.

## Notation in expressions

The C++ language specifies certain compatibilities when specifying operands. The following table shows the types of operands acceptable to operators that require operands of type *type*.

### Operand Types Acceptable to Operators

TYPE EXPECTED	TYPES ALLOWED
<i>type</i>	<code>const type</code> <code>volatile type</code> <code>type&amp;</code> <code>const type&amp;</code> <code>volatile type&amp;</code> <code>volatile const type</code> <code>volatile const type&amp;</code>
<i>type*</i>	<i>type*</i> <code>const type*</code> <code>volatile type*</code> <code>volatile const type*</code>
<code>const type</code>	<i>type</i> <code>const type</code> <code>const type&amp;</code>
<code>volatile type</code>	<i>type</i> <code>volatile type</code> <code>volatile type&amp;</code>

Because the preceding rules can always be used in combination, a `const` pointer to a `volatile` object can be supplied where a pointer is expected.

## Ambiguous expressions

Certain expressions are ambiguous in their meaning. These expressions occur most frequently when an object's value is modified more than once in the same expression. These expressions rely on a particular order of evaluation where the language does not define one. Consider the following example:

```
int i = 7;  
func( i, ++i );
```

The C++ language does not guarantee the order in which arguments to a function call are evaluated. Therefore, in the preceding example, `func` could receive the values 7 and 8, or 8 and 8 for its parameters, depending on whether the parameters are evaluated from left to right or from right to left.

## C++ sequence points (Microsoft Specific)

An expression can modify an object's value only once between consecutive "sequence points."

The C++ language definition does not currently specify sequence points. Microsoft C++ uses the same sequence points as ANSI C for any expression involving C operators and not involving overloaded operators. When

operators are overloaded, the semantics change from operator sequencing to function-call sequencing. Microsoft C++ uses the following sequence points:

- Left operand of the logical AND operator (`&&`). The left operand of the logical AND operator is completely evaluated and all side effects completed before continuing. There is no guarantee that the right operand of the logical AND operator will be evaluated.
- Left operand of the logical OR operator (`||`). The left operand of the logical OR operator is completely evaluated and all side effects completed before continuing. There is no guarantee that the right operand of the logical OR operator will be evaluated.
- Left operand of the comma operator. The left operand of the comma operator is completely evaluated and all side effects completed before continuing. Both operands of the comma operator are always evaluated.
- Function-call operator. The function-call expression and all arguments to a function, including default arguments, are evaluated and all side effects completed prior to entry to the function. There is no specified order of evaluation among the arguments or the function-call expression.
- First operand of the conditional operator. The first operand of the conditional operator is completely evaluated and all side effects completed before continuing.
- The end of a full initialization expression, such as the end of an initialization in a declaration statement.
- The expression in an expression statement. Expression statements consist of an optional expression followed by a semicolon (`;`). The expression is completely evaluated for its side effects.
- The controlling expression in a selection (if or switch) statement. The expression is completely evaluated and all side effects completed before the code dependent on the selection is executed.
- The controlling expression of a while or do statement. The expression is completely evaluated and all side effects completed before any statements in the next iteration of the while or do loop are executed.
- Each of the three expressions of a for statement. Each expression is completely evaluated and all side effects completed before moving to the next expression.
- The expression in a return statement. The expression is completely evaluated and all side effects completed before control returns to the calling function.

## See Also

[Expressions](#)

# Casting

12/21/2017 • 1 min to read • [Edit Online](#)

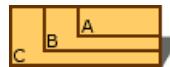
The C++ language provides that if a class is derived from a base class containing virtual functions, a pointer to that base class type can be used to call the implementations of the virtual functions residing in the derived class object. A class containing virtual functions is sometimes called a "polymorphic class."

Since a derived class completely contains the definitions of all the base classes from which it is derived, it is safe to cast a pointer up the class hierarchy to any of these base classes. Given a pointer to a base class, it might be safe to cast the pointer down the hierarchy. It is safe if the object being pointed to is actually of a type derived from the base class. In this case, the actual object is said to be the "complete object." The pointer to the base class is said to point to a "subobject" of the complete object. For example, consider the class hierarchy shown in the following figure.



Class Hierarchy

An object of type `c` could be visualized as shown in the following figure.



Class C with B Subobject and A Subobject

Given an instance of class `c`, there is a `B` subobject and an `A` subobject. The instance of `c`, including the `A` and `B` subobjects, is the "complete object."

Using run-time type information, it is possible to check whether a pointer actually points to a complete object and can be safely cast to point to another object in its hierarchy. The `dynamic_cast` operator can be used to make these types of casts. It also performs the run-time check necessary to make the operation safe.

For conversion of nonpolymorphic types, you can use the `static_cast` operator (this topic explains the difference between static and dynamic casting conversions, and when it is appropriate to use each).

This section covers the following topics:

- [Casting operators](#)
- [Run-time type information](#)

## See Also

[Expressions](#)

# Casting Operators

12/21/2017 • 1 min to read • [Edit Online](#)

There are several casting operators specific to the C++ language. These operators are intended to remove some of the ambiguity and danger inherent in old style C language casts. These operators are:

- [dynamic\\_cast](#) Used for conversion of polymorphic types.
- [static\\_cast](#) Used for conversion of nonpolymorphic types.
- [const\\_cast](#) Used to remove the `const`, `volatile`, and `__unaligned` attributes.
- [reinterpret\\_cast](#) Used for simple reinterpretation of bits.
- [safe\\_cast](#) Used to produce verifiable MSIL.

Use `const_cast` and `reinterpret_cast` as a last resort, since these operators present the same dangers as old style casts. However, they are still necessary in order to completely replace old style casts.

## See Also

[Casting](#)

# dynamic\_cast Operator

12/21/2017 • 7 min to read • [Edit Online](#)

Converts the operand `expression` to an object of type `type-id`.

## Syntax

```
dynamic_cast < type-id > ( expression )
```

## Remarks

The `type-id` must be a pointer or a reference to a previously defined class type or a "pointer to void". The type of `expression` must be a pointer if `type-id` is a pointer, or an l-value if `type-id` is a reference.

See [static\\_cast](#) for an explanation of the difference between static and dynamic casting conversions, and when it is appropriate to use each.

There are two breaking changes in the behavior of `dynamic_cast` in managed code:

- `dynamic_cast` to a pointer to the underlying type of a boxed enum will fail at runtime, returning 0 instead of the converted pointer.
- `dynamic_cast` will no longer throw an exception when `type-id` is an interior pointer to a value type, with the cast failing at runtime. The cast will now return the 0 pointer value instead of throwing.

If `type-id` is a pointer to an unambiguous accessible direct or indirect base class of `expression`, a pointer to the unique subobject of type `type-id` is the result. For example:

```
// dynamic_cast_1.cpp
// compile with: /c
class B { };
class C : public B { };
class D : public C { };

void f(D* pd) {
    C* pc = dynamic_cast<C*>(pd);    // ok: C is a direct base class
                                         // pc points to C subobject of pd
    B* pb = dynamic_cast<B*>(pd);    // ok: B is an indirect base class
                                         // pb points to B subobject of pd
}
```

This type of conversion is called an "upcast" because it moves a pointer up a class hierarchy, from a derived class to a class it is derived from. An upcast is an implicit conversion.

If `type-id` is `void*`, a run-time check is made to determine the actual type of `expression`. The result is a pointer to the complete object pointed to by `expression`. For example:

```

// dynamic_cast_2.cpp
// compile with: /c /GR
class A {virtual void f();};
class B {virtual void f();};

void f() {
    A* pa = new A;
    B* pb = new B;
    void* pv = dynamic_cast<void*>(pa);
    // pv now points to an object of type A

    pv = dynamic_cast<void*>(pb);
    // pv now points to an object of type B
}

```

If `type-id` is not `void*`, a run-time check is made to see if the object pointed to by `expression` can be converted to the type pointed to by `type-id`.

If the type of `expression` is a base class of the type of `type-id`, a run-time check is made to see if `expression` actually points to a complete object of the type of `type-id`. If this is true, the result is a pointer to a complete object of the type of `type-id`. For example:

```

// dynamic_cast_3.cpp
// compile with: /c /GR
class B {virtual void f();};
class D : public B {virtual void f();};

void f() {
    B* pb = new D;    // unclear but ok
    B* pb2 = new B;

    D* pd = dynamic_cast<D*>(pb);    // ok: pb actually points to a D
    D* pd2 = dynamic_cast<D*>(pb2);    // pb2 points to a B not a D
}

```

This type of conversion is called a "downcast" because it moves a pointer down a class hierarchy, from a given class to a class derived from it.

In cases of multiple inheritance, possibilities for ambiguity are introduced. Consider the class hierarchy shown in the following figure.

For CLR types, `dynamic_cast` results in either a no-op if the conversion can be performed implicitly, or an MSIL `isinst` instruction, which performs a dynamic check and returns `nullptr` if the conversion fails.

The following sample uses `dynamic_cast` to determine if a class is an instance of particular type:

```

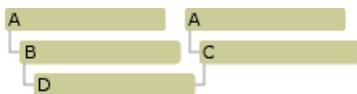
// dynamic_cast_clr.cpp
// compile with: /clr
using namespace System;

void PrintObjectType( Object^o ) {
    if( dynamic_cast<String^>(o) )
        Console::WriteLine("Object is a String");
    else if( dynamic_cast<int>(o) )
        Console::WriteLine("Object is an int");
}

int main() {
    Object^o1 = "hello";
    Object^o2 = 10;

    PrintObjectType(o1);
    PrintObjectType(o2);
}

```



Class Hierarchy Showing Multiple Inheritance

A pointer to an object of type **D** can be safely cast to **B** or **C**. However, if **D** is cast to point to an **A** object, which instance of **A** would result? This would result in an ambiguous casting error. To get around this problem, you can perform two unambiguous casts. For example:

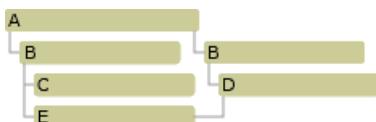
```

// dynamic_cast_4.cpp
// compile with: /c /GR
class A {virtual void f();};
class B {virtual void f();};
class D : public B {virtual void f();};

void f() {
    D* pd = new D;
    B* pb = dynamic_cast<B*>(pd);    // first cast to B
    A* pa2 = dynamic_cast<A*>(pb);    // ok: unambiguous
}

```

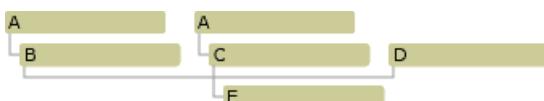
Further ambiguities can be introduced when you use virtual base classes. Consider the class hierarchy shown in the following figure.



Class Hierarchy Showing Virtual Base Classes

In this hierarchy, **A** is a virtual base class. Given an instance of class **E** and a pointer to the **A** subobject, a `dynamic_cast` to a pointer to **B** will fail due to ambiguity. You must first cast back to the complete **E** object, then work your way back up the hierarchy, in an unambiguous manner, to reach the correct **B** object.

Consider the class hierarchy shown in the following figure.



Class Hierarchy Showing Duplicate Base Classes

Given an object of type **E** and a pointer to the **D** subobject, to navigate from the **D** subobject to the left-most

A subobject, three conversions can be made. You can perform a `dynamic_cast` conversion from the D pointer to an E pointer, then a conversion (either `dynamic_cast` or an implicit conversion) from E to B, and finally an implicit conversion from B to A. For example:

```
// dynamic_cast_5.cpp
// compile with: /c /GR
class A {virtual void f();};
class B : public A {virtual void f();};
class C : public A {};
class D {virtual void f();};
class E : public B, public C, public D {virtual void f();};

void f(D* pd) {
    E* pe = dynamic_cast<E*>(pd);
    B* pb = pe;    // upcast, implicit conversion
    A* pa = pb;    // upcast, implicit conversion
}
```

The `dynamic_cast` operator can also be used to perform a "cross cast." Using the same class hierarchy, it is possible to cast a pointer, for example, from the B subobject to the D subobject, as long as the complete object is of type E.

Considering cross casts, it is actually possible to do the conversion from a pointer to D to a pointer to the left-most A subobject in just two steps. You can perform a cross cast from D to B, then an implicit conversion from B to A. For example:

```
// dynamic_cast_6.cpp
// compile with: /c /GR
class A {virtual void f();};
class B : public A {virtual void f();};
class C : public A {};
class D {virtual void f();};
class E : public B, public C, public D {virtual void f();};

void f(D* pd) {
    B* pb = dynamic_cast<B*>(pd);    // cross cast
    A* pa = pb;    // upcast, implicit conversion
}
```

A null pointer value is converted to the null pointer value of the destination type by `dynamic_cast`.

When you use `dynamic_cast < type-id > ( expression )`, if `expression` cannot be safely converted to type `type-id`, the run-time check causes the cast to fail. For example:

```
// dynamic_cast_7.cpp
// compile with: /c /GR
class A {virtual void f();};
class B {virtual void f();};

void f() {
    A* pa = new A;
    B* pb = dynamic_cast<B*>(pa);    // fails at runtime, not safe;
    // B not derived from A
}
```

The value of a failed cast to pointer type is the null pointer. A failed cast to reference type throws a [bad\\_cast Exception](#). If `expression` does not point to or reference a valid object, a `__non_rtti_object` exception is thrown.

See `typeid` for an explanation of the `__non_rtti_object` exception.

## Example

The following sample creates the base class (struct A) pointer, to an object (struct C). This, plus the fact there are virtual functions, enables runtime polymorphism.

The sample also calls a non-virtual function in the hierarchy.

```

// dynamic_cast_8.cpp
// compile with: /GR /EHsc
#include <stdio.h>
#include <iostream>

struct A {
    virtual void test() {
        printf_s("in A\n");
    }
};

struct B : A {
    virtual void test() {
        printf_s("in B\n");
    }

    void test2() {
        printf_s("test2 in B\n");
    }
};

struct C : B {
    virtual void test() {
        printf_s("in C\n");
    }

    void test2() {
        printf_s("test2 in C\n");
    }
};

void Globaltest(A& a) {
    try {
        C &c = dynamic_cast<C&>(a);
        printf_s("in GlobalTest\n");
    }
    catch(std::bad_cast) {
        printf_s("Can't cast to C\n");
    }
}

int main() {
    A *pa = new C;
    A *pa2 = new B;

    pa->test();

    B * pb = dynamic_cast<B *>(pa);
    if (pb)
        pb->test2();

    C * pc = dynamic_cast<C *>(pa2);
    if (pc)
        pc->test2();

    C ConStack;
    Globaltest(ConStack);

    // will fail because B knows nothing about C
    B BonStack;
    Globaltest(BonStack);
}

```

```
in C
test2 in B
in GlobalTest
Can't cast to C
```

## See Also

[Casting Operators](#)

[Keywords](#)

# bad\_cast Exception

12/21/2017 • 1 min to read • [Edit Online](#)

The `bad_cast` exception is thrown by the `dynamic_cast` operator as the result of a failed cast to a reference type.

## Syntax

```
catch (bad_cast)
    statement
```

## Remarks

The interface for `bad_cast` is:

```
class bad_cast : public exception {
public:
    bad_cast(const char * _Message = "bad cast");
    bad_cast(const bad_cast &);
    virtual ~bad_cast();
};
```

The following code contains an example of a failed `dynamic_cast` that throws the `bad_cast` exception.

```
// expre_bad_cast_Exception.cpp
// compile with: /EHsc /GR
#include <typeinfo.h>
#include <iostream>

class Shape {
public:
    virtual void virtualfunc() const {}
};

class Circle: public Shape {
public:
    virtual void virtualfunc() const {}
};

using namespace std;
int main() {
    Shape shape_instance;
    Shape& ref_shape = shape_instance;
    try {
        Circle& ref_circle = dynamic_cast<Circle&>(ref_shape);
    }
    catch (bad_cast b) {
        cout << "Caught: " << b.what();
    }
}
```

The exception is thrown because the object being cast (a `Shape`) is not derived from the specified cast type (`Circle`). To avoid the exception, add these declarations to `main`:

```
Circle circle_instance;
Circle& ref_circle = circle_instance;
```

Then reverse the sense of the cast in the `try` block as follows:

```
Shape& ref_shape = dynamic_cast<Shape&>(ref_circle);
```

## See Also

[dynamic\\_cast Operator](#)

[Keywords](#)

[C++ Exception Handling](#)

# static\_cast Operator

12/21/2017 • 4 min to read • [Edit Online](#)

Converts an *expression* to the type of *type-id*, based only on the types that are present in the expression.

## Syntax

```
static_cast <type-id> ( expression )
```

## Remarks

In standard C++, no run-time type check is made to help ensure the safety of the conversion. In C++/CX, a compile time and runtime check are performed. For more information, see [Casting](#).

The `static_cast` operator can be used for operations such as converting a pointer to a base class to a pointer to a derived class. Such conversions are not always safe.

In general you use `static_cast` when you want to convert numeric data types such as enums to ints or ints to floats, and you are certain of the data types involved in the conversion. `static_cast` conversions are not as safe as `dynamic_cast` conversions, because `static_cast` does no run-time type check, while `dynamic_cast` does. A `dynamic_cast` to an ambiguous pointer will fail, while a `static_cast` returns as if nothing were wrong; this can be dangerous. Although `dynamic_cast` conversions are safer, `dynamic_cast` only works on pointers or references, and the run-time type check is an overhead. For more information, see [dynamic\\_cast Operator](#).

In the example that follows, the line `D* pd2 = static_cast<D*>(pb);` is not safe because `D` can have fields and methods that are not in `B`. However, the line `B* pb2 = static_cast<B*>(pd);` is a safe conversion because `D` always contains all of `B`.

```
// static_cast_Operator.cpp
// compile with: /LD
class B {};

class D : public B {};

void f(B* pb, D* pd) {
    D* pd2 = static_cast<D*>(pb);    // Not safe, D can have fields
                                         // and methods that are not in B.

    B* pb2 = static_cast<B*>(pd);    // Safe conversion, D always
                                         // contains all of B.
}
```

In contrast to `dynamic_cast`, no run-time check is made on the `static_cast` conversion of `pb`. The object pointed to by `pb` may not be an object of type `D`, in which case the use of `*pd2` could be disastrous. For instance, calling a function that is a member of the `D` class, but not the `B` class, could result in an access violation.

The `dynamic_cast` and `static_cast` operators move a pointer throughout a class hierarchy. However, `static_cast` relies exclusively on the information provided in the cast statement and can therefore be unsafe. For example:

```

// static_cast_Operator_2.cpp
// compile with: /LD /GR
class B {
public:
    virtual void Test(){}
};

class D : public B {};

void f(B* pb) {
    D* pd1 = dynamic_cast<D*>(pb);
    D* pd2 = static_cast<D*>(pb);
}

```

If `pb` really points to an object of type `D`, then `pd1` and `pd2` will get the same value. They will also get the same value if `pb == 0`.

If `pb` points to an object of type `B` and not to the complete `D` class, then `dynamic_cast` will know enough to return zero. However, `static_cast` relies on the programmer's assertion that `pb` points to an object of type `D` and simply returns a pointer to that supposed `D` object.

Consequently, `static_cast` can do the inverse of implicit conversions, in which case the results are undefined. It is left to the programmer to verify that the results of a `static_cast` conversion are safe.

This behavior also applies to types other than class types. For instance, `static_cast` can be used to convert from an `int` to a `char`. However, the resulting `char` may not have enough bits to hold the entire `int` value. Again, it is left to the programmer to verify that the results of a `static_cast` conversion are safe.

The `static_cast` operator can also be used to perform any implicit conversion, including standard conversions and user-defined conversions. For example:

```

// static_cast_Operator_3.cpp
// compile with: /LD /GR
typedef unsigned char BYTE;

void f() {
    char ch;
    int i = 65;
    float f = 2.5;
    double dbl;

    ch = static_cast<char>(i);    // int to char
    dbl = static_cast<double>(f);  // float to double
    i = static_cast<BYTE>(ch);
}

```

The `static_cast` operator can explicitly convert an integral value to an enumeration type. If the value of the integral type does not fall within the range of enumeration values, the resulting enumeration value is undefined.

The `static_cast` operator converts a null pointer value to the null pointer value of the destination type.

Any expression can be explicitly converted to type `void` by the `static_cast` operator. The destination `void` type can optionally include the `const`, `volatile`, or `__unaligned` attribute.

The `static_cast` operator cannot cast away the `const`, `volatile`, or `__unaligned` attributes. See [const\\_cast Operator](#) for information on removing these attributes.

Due to the danger of performing unchecked casts on top of a relocating garbage collector, the use of `static_cast` should only be in performance-critical code when you are certain it will work correctly. If you must use `static_cast` in release mode, substitute it with `safe_cast` in your debug builds to ensure success.

## See Also

[Casting Operators](#)

[Keywords](#)

# const\_cast Operator

12/21/2017 • 1 min to read • [Edit Online](#)

Removes the **const**, **volatile**, and **\_unaligned** attribute(s) from a class.

## Syntax

```
const_cast <
type-id
> (
expression
)
```

## Remarks

A pointer to any object type or a pointer to a data member can be explicitly converted to a type that is identical except for the **const**, **volatile**, and **\_unaligned** qualifiers. For pointers and references, the result will refer to the original object. For pointers to data members, the result will refer to the same member as the original (uncast) pointer to data member. Depending on the type of the referenced object, a write operation through the resulting pointer, reference, or pointer to data member might produce undefined behavior.

You cannot use the `const_cast` operator to directly override a constant variable's constant status.

The `const_cast` operator converts a null pointer value to the null pointer value of the destination type.

## Example

```
// expre_const_cast_Operator.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class CCTest {
public:
    void setNumber( int );
    void printNumber() const;
private:
    int number;
};

void CCTest::setNumber( int num ) { number = num; }

void CCTest::printNumber() const {
    cout << "\nBefore: " << number;
    const_cast< CCTest * >( this )->number--;
    cout << "\nAfter: " << number;
}

int main() {
    CCTest X;
    X.setNumber( 8 );
    X.printNumber();
}
```

On the line containing the `const_cast`, the data type of the `this` pointer is `const CCTest *`. The `const_cast` operator changes the data type of the `this` pointer to `CCTest *`, allowing the member `number` to be modified. The cast lasts only for the remainder of the statement in which it appears.

## See Also

[Casting Operators](#)

[Keywords](#)

# reinterpret\_cast Operator

12/21/2017 • 1 min to read • [Edit Online](#)

Allows any pointer to be converted into any other pointer type. Also allows any integral type to be converted into any pointer type and vice versa.

## Syntax

```
reinterpret_cast < type-id > ( expression )
```

## Remarks

Misuse of the `reinterpret_cast` operator can easily be unsafe. Unless the desired conversion is inherently low-level, you should use one of the other cast operators.

The `reinterpret_cast` operator can be used for conversions such as `char*` to `int*`, or `One_class*` to `Unrelated_class*`, which are inherently unsafe.

The result of a `reinterpret_cast` cannot safely be used for anything other than being cast back to its original type. Other uses are, at best, nonportable.

The `reinterpret_cast` operator cannot cast away the `const`, `volatile`, or `_unaligned` attributes. See [const\\_cast Operator](#) for information on removing these attributes.

The `reinterpret_cast` operator converts a null pointer value to the null pointer value of the destination type.

One practical use of `reinterpret_cast` is in a hash function, which maps a value to an index in such a way that two distinct values rarely end up with the same index.

```

#include <iostream>
using namespace std;

// Returns a hash code based on an address
unsigned short Hash( void *p ) {
    unsigned int val = reinterpret_cast<unsigned int>( p );
    return ( unsigned short )( val ^ (val >> 16));
}

using namespace std;
int main() {
    int a[20];
    for ( int i = 0; i < 20; i++ )
        cout << Hash( a + i ) << endl;
}

Output:
64641
64645
64889
64893
64881
64885
64873
64877
64865
64869
64857
64861
64849
64853
64841
64845
64833
64837
64825
64829

```

The `reinterpret_cast` allows the pointer to be treated as an integral type. The result is then bit-shifted and XORed with itself to produce a unique index (unique to a high degree of probability). The index is then truncated by a standard C-style cast to the return type of the function.

## See Also

[Casting Operators](#)

[Keywords](#)

# Run-Time Type Information

12/21/2017 • 1 min to read • [Edit Online](#)

Run-time type information (RTTI) is a mechanism that allows the type of an object to be determined during program execution. RTTI was added to the C++ language because many vendors of class libraries were implementing this functionality themselves. This caused incompatibilities between libraries. Thus, it became obvious that support for run-time type information was needed at the language level.

For the sake of clarity, this discussion of RTTI is almost completely restricted to pointers. However, the concepts discussed also apply to references.

There are three main C++ language elements to run-time type information:

- The `dynamic_cast` operator.

Used for conversion of polymorphic types.

- The `typeid` operator.

Used for identifying the exact type of an object.

- The `type_info` class.

Used to hold the type information returned by the `typeid` operator.

## See Also

[Casting](#)

# bad\_typeid Exception

12/21/2017 • 1 min to read • [Edit Online](#)

The `bad_typeid` exception is thrown by the `typeid` operator when the operand for `typeid` is a NULL pointer.

## Syntax

```
catch (bad_typeid)
statement
```

## Remarks

The interface for `bad_typeid` is:

```
class bad_typeid : public exception
{
public:
    bad_typeid(const char * _Message = "bad typeid");
    bad_typeid(const bad_typeid &);
    virtual ~bad_typeid();
};
```

The following example shows the `typeid` operator throwing a `bad_typeid` exception.

```
// expre_bad_typeid.cpp
// compile with: /EHsc /GR
#include <typeinfo.h>
#include <iostream>

class A{
public:
    // object for class needs vtable
    // for RTTI
    virtual ~A();
};

using namespace std;
int main() {
A* a = NULL;

try {
    cout << typeid(*a).name() << endl; // Error condition
}
catch (bad_typeid){
    cout << "Object is NULL" << endl;
}
}
```

## Output

```
Object is NULL
```

## See Also

[Run-Time Type Information](#)

[Keywords](#)

# type\_info Class

12/21/2017 • 1 min to read • [Edit Online](#)

The **type\_info** class describes type information generated within the program by the compiler. Objects of this class effectively store a pointer to a name for the type. The **type\_info** class also stores an encoded value suitable for comparing two types for equality or collating order. The encoding rules and collating sequence for types are unspecified and may differ between programs.

The `<typeinfo>` header file must be included in order to use the **type\_info** class. The interface for the **type\_info** class is:

```
class type_info {
public:
    virtual ~type_info();
    size_t hash_code() const
    _CRTIMP_PURE bool operator==(const type_info& rhs) const;
    _CRTIMP_PURE bool operator!=(const type_info& rhs) const;
    _CRTIMP_PURE int before(const type_info& rhs) const;
    _CRTIMP_PURE const char* name() const;
    _CRTIMP_PURE const char* raw_name() const;
};
```

You cannot instantiate objects of the **type\_info** class directly, because the class has only a private copy constructor. The only way to construct a (temporary) **type\_info** object is to use the **typeid** operator. Since the assignment operator is also private, you cannot copy or assign objects of class **type\_info**.

**type\_info::hash\_code** defines a hash function suitable for mapping values of type **typeinfo** to a distribution of index values.

The operators `==` and `!=` can be used to compare for equality and inequality with other **type\_info** objects, respectively.

There is no link between the collating order of types and inheritance relationships. Use the **type\_info::before** member function to determine the collating sequence of types. There is no guarantee that **type\_info::before** will yield the same result in different programs or even different runs of the same program. In this manner, **type\_info::before** is similar to the address-of (`&`) operator.

The **type\_info::name** member function returns a **const char\\*** to a null-terminated string representing the human-readable name of the type. The memory pointed to is cached and should never be directly deallocated.

The **type\_info::raw\_name** member function returns a **const char\\*** to a null-terminated string representing the decorated name of the object type. The name is actually stored in its decorated form to save space. Consequently, this function is faster than **type\_info::name** because it doesn't need to undecorate the name. The string returned by the **type\_info::raw\_name** function is useful in comparison operations but is not readable. If you need a human-readable string, use the **type\_info::name** function instead.

Type information is generated for polymorphic classes only if the [/GR \(Enable Run-Time Type Information\)](#) compiler option is specified.

## See Also

[Run-Time Type Information](#)

# Statements (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

C++ statements are the program elements that control how and in what order objects are manipulated. This section includes:

- [Overview](#)
- [Labeled Statements](#)
- Categories of Statements
  - [Expression statements](#). These statements evaluate an expression for its side effects or for its return value.
  - [Null statements](#). These statements can be provided where a statement is required by the C++ syntax but where no action is to be taken.
  - [Compound statements](#). These statements are groups of statements enclosed in curly braces ({}). They can be used wherever a single statement may be used.
  - [Selection statements](#). These statements perform a test; they then execute one section of code if the test evaluates to true (nonzero). They may execute another section of code if the test evaluates to false.
  - [Iteration statements](#). These statements provide for repeated execution of a block of code until a specified termination criterion is met.
  - [Jump statements](#). These statements either transfer control immediately to another location in the function or return control from the function.
  - [Declaration statements](#). Declarations introduce a name into a program. ([Declarations and Definitions](#) provides more detailed information about declarations.)

For information on exception handling statements see [Exception Handling](#).

## See Also

[C++ Language Reference](#)

# Overview of C++ Statements

12/21/2017 • 1 min to read • [Edit Online](#)

C++ statements are executed sequentially, except when an expression statement, a selection statement, an iteration statement, or a jump statement specifically modifies that sequence.

Statements may be of the following types:

```
labeled-statement  
expression-statement  
compound-statement  
selection-statement  
iteration-statement  
jump-statement  
declaration-statement  
try-throw-catch
```

In most cases, the C++ statement syntax is identical to that of ANSI C. The primary difference between the two is that in C, declarations are allowed only at the start of a block; C++ adds the *declaration-statement*, which effectively removes this restriction. This enables you to introduce variables at a point in the program where a precomputed initialization value can be calculated.

Declaring variables inside blocks also allows you to exercise precise control over the scope and lifetime of those variables.

The topics on statements describe the following C++ keywords:

<code>break</code>	<code>else</code>	<code>_if_exists</code>	<code>_try</code>
<code>case</code>	<code>_except</code>	<code>_if_not_exists</code>	<code>try</code>
<code>catch</code>	<code>for</code>	<code>_leave</code>	<code>while</code>
<code>continue</code>	<code>goto</code>	<code>return</code>	
<code>default</code>	<code>_finally</code>	<code>switch</code>	
<code>do</code>	<code>if</code>	<code>throw</code>	

## See Also

[Statements](#)

# Labeled Statements

12/21/2017 • 3 min to read • [Edit Online](#)

Labels are used to transfer program control directly to the specified statement.

```
identifier : statement
case constant-expression : statement
default : statement
```

The scope of a label is the entire function in which it is declared.

## Remarks

There are three types of labeled statements. All use a colon to separate some type of label from the statement. The case and default labels are specific to case statements.

```
#include <iostream>
using namespace std;

void test_label(int x) {

    if (x == 1){
        goto label1;
    }
    goto label2;

label1:
    cout << "in label1" << endl;
    return;

label2:
    cout << "in label2" << endl;
    return;
}

int main() {
    test_label(1); // in label1
    test_label(2); // in label2
}
```

### The `goto` statement

The appearance of an *identifier* label in the source program declares a label. Only a `goto` statement can transfer control to an *identifier* label. The following code fragment illustrates use of the `goto` statement and an *identifier* label:

A label cannot appear by itself but must always be attached to a statement. If a label is needed by itself, place a null statement after the label.

The label has function scope and cannot be redeclared within the function. However, the same name can be used as a label in different functions.

```

// labels_with_goto.cpp
// compile with: /EHsc
#include <iostream>
int main() {
    using namespace std;
    goto Test2;

    cout << "testing" << endl;

Test2:
    cerr << "At Test2 label." << endl;
}

//Output: At Test2 label.

```

## The case statement

Labels that appear after the **case** keyword cannot also appear outside a `switch` statement. (This restriction also applies to the **default** keyword.) The following code fragment shows the correct use of **case** labels:

```

// Sample Microsoft Windows message processing loop.
switch( msg )
{
    case WM_TIMER:      // Process timer event.
        SetClassWord( hWnd, GCW_HICON, ahIcon[nIcon++] );
        ShowWindow( hWnd, SW_SHOWNA );
        nIcon %= 14;
        Yield();
        break;

    case WM_PAINT:
        memset( &ps, 0x00, sizeof(PAINTSTRUCT) );
        hDC = BeginPaint( hWnd, &ps );
        EndPaint( hWnd, &ps );
        break;

    default:
        // This choice is taken for all messages not specifically
        // covered by a case statement.

        return DefWindowProc( hWnd, Message, wParam, lParam );
        break;
}

```

## Labels in the case statement

Labels that appear after the **case** keyword cannot also appear outside a `switch` statement. (This restriction also applies to the **default** keyword.) The following code fragment shows the correct use of **case** labels:

```

// Sample Microsoft Windows message processing loop.
switch( msg )
{
    case WM_TIMER:      // Process timer event.
        SetClassWord( hWnd, GCW_HICON, ahIcon[nIcon++] );
        ShowWindow( hWnd, SW_SHOWNA );
        nIcon %= 14;
        Yield();
        break;

    case WM_PAINT:
        // Obtain a handle to the device context.
        // BeginPaint will send WM_ERASEBKGND if appropriate.

        memset( &ps, 0x00, sizeof(PAINTSTRUCT) );
        hDC = BeginPaint( hWnd, &ps );

        // Inform Windows that painting is complete.

        EndPaint( hWnd, &ps );
        break;

    case WM_CLOSE:
        // Close this window and all child windows.

        KillTimer( hWnd, TIMER1 );
        DestroyWindow( hWnd );
        if ( hWnd == hWndMain )
            PostQuitMessage( 0 ); // Quit the application.
        break;

    default:
        // This choice is taken for all messages not specifically
        // covered by a case statement.

        return DefWindowProc( hWnd, Message, wParam, lParam );
    break;
}

```

## Labels in the goto statement

The appearance of an *identifier* label in the source program declares a label. Only a [goto](#) statement can transfer control to an *identifier* label. The following code fragment illustrates use of the [goto](#) statement and an *identifier* label:

A label cannot appear by itself but must always be attached to a statement. If a label is needed by itself, place a null statement after the label.

The label has function scope and cannot be redeclared within the function. However, the same name can be used as a label in different functions.

```
// labels_with_goto.cpp
// compile with: /EHsc
#include <iostream>
int main() {
    using namespace std;
    goto Test2;

    cout << "testing" << endl;

    Test2:
    cerr << "At Test2 label." << endl;
// At Test2 label.
}
```

## See Also

[Overview of C++ Statements](#)

[switch Statement \(C++\)](#)

# Expression Statement

12/21/2017 • 1 min to read • [Edit Online](#)

Expression statements cause expressions to be evaluated. No transfer of control or iteration takes place as a result of an expression statement.

The syntax for the expression statement is simply

## Syntax

```
[expression ] ;
```

## Remarks

All expressions in an expression statement are evaluated and all side effects are completed before the next statement is executed. The most common expression statements are assignments and function calls. Since the expression is optional, a semicolon alone is considered an empty expression statement, referred to as the [null](#) statement.

## See Also

[Overview of C++ Statements](#)

# Null Statement

12/21/2017 • 1 min to read • [Edit Online](#)

The "null statement" is an expression statement with the *expression* missing. It is useful when the syntax of the language calls for a statement but no expression evaluation. It consists of a semicolon.

Null statements are commonly used as placeholders in iteration statements or as statements on which to place labels at the end of compound statements or functions.

The following code fragment shows how to copy one string to another and incorporates the null statement:

```
// null_statement.cpp
char *myStrCpy( char *Dest, const char *Source )
{
    char *DestStart = Dest;

    // Assign value pointed to by Source to
    // Dest until the end-of-string 0 is
    // encountered.
    while( *Dest++ = *Source++ )
        ;    // Null statement.

    return DestStart;
}

int main()
{
}
```

## See Also

[Expression Statement](#)

# Compound Statements (Blocks)

12/21/2017 • 1 min to read • [Edit Online](#)

A compound statement consists of zero or more statements enclosed in curly braces (`{ }`). A compound statement can be used anywhere a statement is expected. Compound statements are commonly called "blocks."

## Syntax

```
{ [ statement-list ] }
```

## Remarks

The following example uses a compound statement as the *statement* part of the `if` statement (see [The if Statement](#) for details about the syntax):

```
if( Amount > 100 )
{
    cout << "Amount was too large to handle\n";
    Alert();
}
else
{
    Balance -= Amount;
}
```

### NOTE

Because a declaration is a statement, a declaration can be one of the statements in the *statement-list*. As a result, names declared inside a compound statement, but not explicitly declared as static, have local scope and (for objects) lifetime. See [Scope](#) for details about treatment of names with local scope.

## See Also

[Overview of C++ Statements](#)

# Selection Statements (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The C++ selection statements, [if](#) and [switch](#), provide a means to conditionally execute sections of code.

The [\\_\\_if\\_exists](#) and [\\_\\_if\\_not\\_exists](#) statements allow you to conditionally include code depending on the existence of a symbol.

See the individual topics for the syntax for each statement.

## See Also

[Overview of C++ Statements](#)

# if-else Statement (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Controls conditional branching. Statements in the *if-block* are executed only if the *if-expression* evaluates to a non-zero value (or `true`). If the value of *expression* is nonzero, *statement1* and any other statements in the block are executed and the *else-block*, if present, is skipped. If the value of *expression* is zero, then the *if-block* is skipped and the *else-block*, if present, is executed. Expressions that evaluate to non-zero are

- `true`
- a non-null pointer,
- any non-zero arithmetic value, or
- a class type that defines an unambiguous conversion to an arithmetic, boolean or pointer type. (For information about conversions, see [Standard Conversions](#).)

## Syntax

```
if ( expression )
{
    statement1;
    ...
}
else // optional
{
    statement2;
    ...
}

// Visual Studio 2017 version 15.3 and later:
if ( initialization; expression )
{
    statement1;
    ...
}
else // optional
{
    statement2;
    ...
}

// Visual Studio 2017 version 15.3 and later:
if constexpr (expression)
{
    statement1;
    ...
}
else // optional
{
    statement2;
    ...
}
```

## Example

```

// if_else_statement.cpp
#include <iostream>

using namespace std;

class C
{
public:
    void do_somthing(){}
};

void init(C){}
bool is_true() { return true; }
int x = 10;

int main()
{
    if (is_true())
    {
        cout << "b is true!\n"; // executed
    }
    else
    {
        cout << "b is false!\n";
    }

    // no else statement
    if (x == 10)
    {
        x = 0;
    }

    C* c;
    init(c);
    if (c)
    {
        c->do_something();
    }
    else
    {
        cout << "c is null!\n";
    }
}

```

## if statement with an initializer

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): An **if** statement may also contain an expression that declares and initializes a named variable. Use this form of the if-statement when the variable is only needed within the scope of the if-block.

```

## Example
#include <iostream>
#include <mutex>
#include <map>
#include <string>
#include <algorithm>

using namespace std;

map<int, string> m;
mutex mx;
bool shared_flag; // guarded by mx
void unsafe_operation() {}

int main()
{
    if (auto it = m.find(10); it != m.end())
    {
        cout << it->second;
        return 0;
    }

    if (char buf[10]; fgets(buf, 10, stdin))
    {
        m[0] += buf;
    }

    if (lock_guard<mutex> lock(mx); shared_flag)
    {
        unsafe_operation();
        shared_flag = false;
    }

    string s{ "if" };
    if (auto keywords = { "if", "for", "while" }; any_of(keywords.begin(), keywords.end(), [&s](const char* kw) { return s == kw; }))
    {
        cout << "Error! Token must not be a keyword\n";
    }
}

```

In all forms of the **if** statement, *expression*, which can have any value except a structure, is evaluated, including all side effects. Control passes from the **if** statement to the next statement in the program unless one of the *statements* contains a [break](#), [continue](#), or [goto](#).

The **else** clause of an `if...else` statement is associated with the closest previous **if** statement in the same scope that does not have a corresponding **else** statement.

## constexpr if statements

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): In function templates, you can use a **constexpr if** statement to make compile-time branching decisions without having to resort to multiple function overloads. For example, you can write a single function that handles parameter unpacking (no zero-parameter overload is needed):

```
template <class T, class... Rest>
void f(T&& t, Rest&&... r)
{
// handle t
do_something(t);

// handle r conditionally
constexpr if (sizeof...(r))
{
    f(r...);
}
else
{
    g(r...);
}
}
```

## See Also

[Selection Statements](#)

[Keywords](#)

[switch Statement \(C++\)](#)

# `__if_exists` Statement

12/21/2017 • 1 min to read • [Edit Online](#)

The `__if_exists` statement tests whether the specified identifier exists. If the identifier exists, the specified statement block is executed.

## Syntax

```
__if_exists ( identifier ) {
    statements
};
```

### Parameters

PARAMETER	DESCRIPTION
<code>identifier</code>	The identifier whose existence you want to test.
<code>statements</code>	One or more statements to execute if <code>identifier</code> exists.

## Remarks

### Caution

To achieve the most reliable results, use the `__if_exists` statement under the following constraints.

- Apply the `__if_exists` statement to only simple types, not templates.
- Apply the `__if_exists` statement to identifiers both inside or outside a class. Do not apply the `__if_exists` statement to local variables.
- Use the `__if_exists` statement only in the body of a function. Outside of the body of a function, the `__if_exists` statement can test only fully defined types.
- When you test for overloaded functions, you cannot test for a specific form of the overload.

The complement to the `__if_exists` statement is the [`\_\_if\_not\_exists`](#) statement.

## Example

Notice that this example uses templates, which is not advised.

```

// the_if_exists_statement.cpp
// compile with: /EHsc
#include <iostream>

template<typename T>
class X : public T {
public:
    void Dump() {
        std::cout << "In X<T>::Dump()" << std::endl;

        __if_exists(T::Dump) {
            T::Dump();
        }

        __if_not_exists(T::Dump) {
            std::cout << "T::Dump does not exist" << std::endl;
        }
    }
};

class A {
public:
    void Dump() {
        std::cout << "In A::Dump()" << std::endl;
    }
};

class B {};

bool g_bFlag = true;

class C {
public:
    void f(int);
    void f(double);
};

int main() {
    X<A> x1;
    X<B> x2;

    x1.Dump();
    x2.Dump();

    __if_exists(::g_bFlag) {
        std::cout << "g_bFlag = " << g_bFlag << std::endl;
    }

    __if_exists(C::f) {
        std::cout << "C::f exists" << std::endl;
    }

    return 0;
}

```

## Output

```

In X<T>::Dump()
In A::Dump()
In X<T>::Dump()
T::Dump does not exist
g_bFlag = 1
C::f exists

```

## See Also

[Selection Statements](#)

[Keywords](#)

[\\_if\\_not\\_exists Statement](#)

# `__if_not_exists` Statement

12/21/2017 • 1 min to read • [Edit Online](#)

The `__if_not_exists` statement tests whether the specified identifier exists. If the identifier does not exist, the specified statement block is executed.

## Syntax

```
__if_not_exists ( identifier ) {
    statements
};
```

### Parameters

PARAMETER	DESCRIPTION
<code>identifier</code>	The identifier whose existence you want to test.
<code>statements</code>	One or more statements to execute if <code>identifier</code> does not exist.

## Remarks

### Caution

To achieve the most reliable results, use the `__if_not_exists` statement under the following constraints.

- Apply the `__if_not_exists` statement to only simple types, not templates.
- Apply the `__if_not_exists` statement to identifiers both inside or outside a class. Do not apply the `__if_not_exists` statement to local variables.
- Use the `__if_not_exists` statement only in the body of a function. Outside of the body of a function, the `__if_not_exists` statement can test only fully defined types.
- When you test for overloaded functions, you cannot test for a specific form of the overload.

The complement to the `__if_not_exists` statement is the [\\_\\_if\\_exists](#) statement.

## Example

For an example about how to use `__if_not_exists`, see [\\_\\_if\\_exists Statement](#).

## See Also

[Selection Statements](#)

[Keywords](#)

[\\_\\_if\\_exists Statement](#)

# switch Statement (C++)

12/21/2017 • 4 min to read • [Edit Online](#)

Allows selection among multiple sections of code, depending on the value of an integral expression.

## Syntax

```
switch ( init; expression )
case constant-expression : statement
[default : statement]
```

## Remarks

The *expression* must be of an integral type or of a class type for which there is an unambiguous conversion to integral type. Integral promotion is performed as described in [Standard Conversions](#).

The `switch` statement body consists of a series of `case` labels and an optional `default` label. No two constant expressions in `case` statements can evaluate to the same value. The `default` label can appear only once. The labeled statements are not syntactic requirements, but the `switch` statement is meaningless without them. The default statement need not come at the end; it can appear anywhere in the body of the switch statement. A case or default label can only appear inside a switch statement.

The *constant-expression* in each `case` label is converted to the type of *expression* and compared with *expression* for equality. Control passes to the statement whose `case` *constant-expression* matches the value of *expression*. The resulting behavior is shown in the following table.

### Switch Statement Behavior

CONDITION	ACTION
Converted value matches that of the promoted controlling expression.	Control is transferred to the statement following that label.
None of the constants match the constants in the <code>case</code> labels; a <code>default</code> label is present.	Control is transferred to the <code>default</code> label.
None of the constants match the constants in the <code>case</code> labels; <code>default</code> label is not present.	Control is transferred to the statement after the <code>switch</code> statement.

If a matching expression is found, control is not impeded by subsequent `case` or `default` labels. The `break` statement is used to stop execution and transfer control to the statement after the `switch` statement. Without a `break` statement, every statement from the matched `case` label to the end of the `switch`, including the `default`, is executed. For example:

```

// switch_statement1.cpp
#include <stdio.h>

int main() {
    char *buffer = "Any character stream";
    int capa, lettera, nota;
    char c;
    capa = lettera = nota = 0;

    while ( c = *buffer++ ) // Walks buffer until NULL
    {
        switch ( c )
        {
            case 'A':
                capa++;
                break;
            case 'a':
                lettera++;
                break;
            default:
                nota++;
        }
    }
    printf_s( "\nUppercase a: %d\nLowercase a: %d\nTotal: %d\n",
              capa, lettera, (capa + lettera + nota) );
}

```

In the above example, `capa` is incremented if `c` is an uppercase `A`. The `break` statement after `capa++` terminates execution of the `switch` statement body and control passes to the `while` loop. Without the `break` statement, execution would "fall through" to the next labeled statement, so that `lettera` and `nota` would also be incremented. A similar purpose is served by the `break` statement for `case 'a'`. If `c` is a lowercase `a`, `lettera` is incremented and the `break` statement terminates the `switch` statement body. If `c` is not an `a` or `A`, the `default` statement is executed.

**Visual Studio 2017 and later:** (available with [/std:c++17](#)) The `[[fallthrough]]` attribute is specified in the C++17 standard. It can be used in a `switch` statement as a hint to the compiler (or to anyone reading the code) that fall-through behavior is intended. The Visual C++ compiler currently does not warn on fallthrough behavior, so this attribute has no effect compiler behavior. Note that the attribute is applied to an empty statement within the labeled statement; in other words the semicolon is necessary.

```
int main()
{
    int n = 5;
    switch (n)
    {

        case 1:
            a();
            break;
        case 2:
            b();
            d();
            [[fallthrough]]; // I meant to do this!
        case 3:
            c();
            break;
        default:
            d();
            break;
    }

    return 0;
}
```

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): A switch statement may introduce and initialize a variable whose scope is limited to the block of the switch statement:

```
switch (Gadget gadget(args); auto s = gadget.get_status())
{
    case status::good:
        gadget.zip();
        break;
    case status::bad:
        throw BadGadget();
};
```

An inner block of a `switch` statement can contain definitions with initializations as long as they are reachable — that is, not bypassed by all possible execution paths. Names introduced using these declarations have local scope. For example:

```

// switch_statement2.cpp
// C2360 expected
#include <iostream>
using namespace std;
int main(int argc, char *argv[])
{
    switch( tolower( *argv[1] ) )
    {
        // Error. Unreachable declaration.
        char szChEntered[] = "Character entered was: ";

        case 'a' :
        {
            // Declaration of szChEntered OK. Local scope.
            char szChEntered[] = "Character entered was: ";
            cout << szChEntered << "a\n";
        }
        break;

        case 'b' :
        // Value of szChEntered undefined.
        cout << szChEntered << "b\n";
        break;

        default:
        // Value of szChEntered undefined.
        cout << szChEntered << "neither a nor b\n";
        break;
    }
}

```

A `switch` statement can be nested. In such cases, `case` or `default` labels associate with the closest `switch` statement that encloses them.

## Microsoft Specific

Microsoft C does not limit the number of case values in a `switch` statement. The number is limited only by the available memory. ANSI C requires at least 257 case labels be allowed in a `switch` statement.

The default for Microsoft C is that the Microsoft extensions are enabled. Use the `/Za` compiler option to disable these extensions.

### END Microsoft Specific

## See Also

[Selection Statements](#)

[Keywords](#)

# Iteration Statements (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Iteration statements cause statements (or compound statements) to be executed zero or more times, subject to some loop-termination criteria. When these statements are compound statements, they are executed in order, except when either the `break` statement or the `continue` statement is encountered.

C++ provides four iteration statements — `while`, `do`, `for`, and `range-based for`. Each of these iterates until its termination expression evaluates to zero (false), or until loop termination is forced with a `break` statement. The following table summarizes these statements and their actions; each is discussed in detail in the sections that follow.

## Iteration Statements

STATEMENT	EVALUATED AT	INITIALIZATION	INCREMENT
<code>while</code>	Top of loop	No	No
<code>do</code>	Bottom of loop	No	No
<code>for</code>	Top of loop	Yes	Yes
<code>range-based for</code>	Top of loop	Yes	Yes

The statement part of an iteration statement cannot be a declaration. However, it can be a compound statement containing a declaration.

## See Also

[Overview of C++ Statements](#)

# while Statement (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Executes *statement* repeatedly until *expression* evaluates to zero.

## Syntax

```
while ( expression )
    statement
```

## Remarks

The test of *expression* takes place before each execution of the loop; therefore, a `while` loop executes zero or more times. *expression* must be of an integral type, a pointer type, or a class type with an unambiguous conversion to an integral or pointer type.

A `while` loop can also terminate when a `break`, `goto`, or `return` within the statement body is executed. Use `continue` to terminate the current iteration without exiting the `while` loop. `continue` passes control to the next iteration of the `while` loop.

The following code uses a `while` loop to trim trailing underscores from a string:

```
// while_statement.cpp

#include <string.h>
#include <stdio.h>
char *trim( char *szSource )
{
    char *pszEOS = 0;

    // Set pointer to character before terminating NULL
    pszEOS = szSource + strlen( szSource ) - 1;

    // iterate backwards until non '_' is found
    while( (pszEOS >= szSource) && (*pszEOS == '_') )
        *pszEOS-- = '\0';

    return szSource;
}
int main()
{
    char szbuf[] = "12345_____";

    printf_s("\nBefore trim: %s", szbuf);
    printf_s("\nAfter trim: %s\n", trim(szbuf));
}
```

The termination condition is evaluated at the top of the loop. If there are no trailing underscores, the loop never executes.

## See Also

[Iteration Statements](#)

[Keywords](#)

[do-while Statement \(C++\)](#)

[for Statement \(C++\)](#)

[Range-based for Statement \(C++\)](#)

# do-while Statement (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Executes a *statement* repeatedly until the specified termination condition (the *expression*) evaluates to zero.

## Syntax

```
do
    statement
  while ( expression ) ;
```

## Remarks

The test of the termination condition is made after each execution of the loop; therefore, a `do-while` loop executes one or more times, depending on the value of the termination expression. The `do-while` statement can also terminate when a `break`, `goto`, or `return` statement is executed within the statement body.

The *expression* must have arithmetic or pointer type. Execution proceeds as follows:

1. The statement body is executed.
2. Next, *expression* is evaluated. If *expression* is false, the `do-while` statement terminates and control passes to the next statement in the program. If *expression* is true (nonzero), the process is repeated, beginning with step 1.

## Example

The following sample demonstrates the `do-while` statement:

```
// do_while_statement.cpp
#include <stdio.h>
int main()
{
    int i = 0;
    do
    {
        printf_s("\n%d",i++);
    } while (i < 3);
}
```

## See Also

[Iteration Statements](#)

[Keywords](#)

[while Statement \(C++\)](#)

[for Statement \(C++\)](#)

[Range-based for Statement \(C++\)](#)

# for Statement (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Executes a statement repeatedly until the condition becomes false. For information on the range-based for statement, see [Range-based for Statement \(C++\)](#).

## Syntax

```
for ( init-expression ; cond-expression ; loop-expression )
    statement;
```

## Remarks

Use the `for` statement to construct loops that must execute a specified number of times.

The `for` statement consists of three optional parts, as shown in the following table.

### for Loop Elements

SYNTAX NAME	WHEN EXECUTED	DESCRIPTION
<code>init-expression</code>	Before any other element of the <code>for</code> statement, <code>init-expression</code> is executed only once. Control then passes to <code>cond-expression</code> .	Often used to initialize loop indices. It can contain expressions or declarations.
<code>cond-expression</code>	Before execution of each iteration of <code>statement</code> , including the first iteration. <code>statement</code> is executed only if <code>cond-expression</code> evaluates to true (nonzero).	An expression that evaluates to an integral type or a class type that has an unambiguous conversion to an integral type. Normally used to test for loop-termination criteria.
<code>loop-expression</code>	At the end of each iteration of <code>statement</code> . After <code>loop-expression</code> is executed, <code>cond-expression</code> is evaluated.	Normally used to increment loop indices.

The following examples show different ways to use the `for` statement.

```

#include <iostream>
using namespace std;

int main() {
    // The counter variable can be declared in the init-expression.
    for (int i = 0; i < 2; i++ ){
        cout << i;
    }
    // Output: 01
    // The counter variable can be declared outside the for loop.
    int i;
    for (i = 0; i < 2; i++){
        cout << i;
    }
    // Output: 01
    // These for loops are the equivalent of a while loop.
    i = 0;
    while (i < 2){
        cout << i++;
    }
}
// Output: 012

```

`init-expression` and `loop-expression` can contain multiple statements separated by commas. For example:

```

#include <iostream>
using namespace std;

int main(){
    int i, j;
    for ( i = 5, j = 10 ; i + j < 20; i++, j++ ) {
        cout << "i + j = " << (i + j) << '\n';
    }
}
// Output:
i + j = 15
i + j = 17
i + j = 19

```

`loop-expression` can be incremented or decremented, or modified in other ways.

```

#include <iostream>
using namespace std;

int main(){
for (int i = 10; i > 0; i--) {
    cout << i << ' ';
}
// Output: 10 9 8 7 6 5 4 3 2 1
for (int i = 10; i < 20; i = i+2) {
    cout << i << ' ';
}
// Output: 10 12 14 16 18

```

A `for` loop terminates when a `break`, `return`, or `goto` (to a labeled statement outside the `for` loop) within `statement` is executed. A `continue` statement in a `for` loop terminates only the current iteration.

If `cond-expression` is omitted, it is considered true and the `for` loop will not terminate without a `break`, `return`, or `goto` within `statement`.

Although the three fields of the `for` statement are normally used for initialization, testing for termination, and

incrementing, they are not restricted to these uses. For example, the following code prints the numbers 0 through 4. In this case, `statement` is the null statement:

```
#include <iostream>
using namespace std;

int main()
{
    int i;
    for( i = 0; i < 5; cout << i << '\n', i++){
        ;
    }
}
```

## for Loops and the C++ Standard

The C++ standard says that a variable declared in a `for` loop shall go out of scope after the `for` loop ends. For example:

```
for (int i = 0 ; i < 5 ; i++) {
    // do something
}
// i is now out of scope under /Za or /Zc:forScope
```

By default, under `/Ze`, a variable declared in a `for` loop remains in scope until the `for` loop's enclosing scope ends.

`/Zc:forScope` enables standard behavior of variables declared in for loops without needing to specify `/Za`.

It is also possible to use the scoping differences of the `for` loop to redeclare variables under `/Ze` as follows:

```
// for_statement5.cpp
int main(){
    int i = 0;    // hidden by var with same name declared in for loop
    for ( int i = 0 ; i < 3; i++ ) {}

    for ( int i = 0 ; i < 3; i++ ) {}
}
```

This more closely mimics the standard behavior of a variable declared in a `for` loop, which requires variables declared in a `for` loop to go out of scope after the loop is done. When a variable is declared in a `for` loop, the compiler internally promotes it to a local variable in the `for` loop's enclosing scope even if there is already a local variable with the same name.

## See Also

[Iteration Statements](#)

[Keywords](#)

[while Statement \(C++\)](#)

[do-while Statement \(C++\)](#)

[Range-based for Statement \(C++\)](#)

# Range-based for Statement (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Executes `statement` repeatedly and sequentially for each element in `expression`.

## Syntax

```
for ( for-range-declaration : expression )
    statement
```

## Remarks

Use the range-based `for` statement to construct loops that must execute through a "range", which is defined as anything that you can iterate through—for example, `std::vector`, or any other C++ Standard Library sequence whose range is defined by a `begin()` and `end()`. The name that is declared in the `for-range-declaration` portion is local to the `for` statement and cannot be re-declared in `expression` or `statement`. Note that the `auto` keyword is preferred in the `for-range-declaration` portion of the statement.

**New in Visual Studio 2017:** Range-based for loops no longer require that `begin()` and `end()` return objects of the same type. This enables `end()` to return a sentinel object such as used by ranges as defined in the Ranges-V3 proposal. For more information, see [Generalizing the Range-Based For Loop](#) and the [range-v3 library on GitHub](#).

This code shows how to use range-based `for` loops to iterate through an array and a vector:

```

// range-based-for.cpp
// compile by using: cl /EHsc /nologo /W4
#include <iostream>
#include <vector>
using namespace std;

int main()
{
    // Basic 10-element integer array.
    int x[10] = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 };

    // Range-based for loop to iterate through the array.
    for( int y : x ) { // Access by value using a copy declared as a specific type.
        // Not preferred.
        cout << y << " ";
    }
    cout << endl;

    // The auto keyword causes type inference to be used. Preferred.

    for( auto y : x ) { // Copy of 'x', almost always undesirable
        cout << y << " ";
    }
    cout << endl;

    for( auto &y : x ) { // Type inference by reference.
        // Observes and/or modifies in-place. Preferred when modify is needed.
        cout << y << " ";
    }
    cout << endl;

    for( const auto &y : x ) { // Type inference by const reference.
        // Observes in-place. Preferred when no modify is needed.
        cout << y << " ";
    }
    cout << endl;
    cout << "end of integer array test" << endl;
    cout << endl;

    // Create a vector object that contains 10 elements.
    vector<double> v;
    for (int i = 0; i < 10; ++i) {
        v.push_back(i + 0.14159);
    }

    // Range-based for loop to iterate through the vector, observing in-place.
    for( const auto &j : v ) {
        cout << j << " ";
    }
    cout << endl;
    cout << "end of vector test" << endl;
}

```

Here is the output:

1 2 3 4 5 6 7 8 9 10
----------------------

1 2 3 4 5 6 7 8 9 10
----------------------

1 2 3 4 5 6 7 8 9 10
----------------------

1 2 3 4 5 6 7 8 9 10
----------------------

end of integer array test
---------------------------

0.14159 1.14159 2.14159 3.14159 4.14159 5.14159 6.14159 7.14159 8.14159 9.14159
---

end of vector test

A range-based `for` loop terminates when one of these in `statement` is executed: a `break`, `return`, or `goto` to a labeled statement outside the range-based `for` loop. A `continue` statement in a range-based `for` loop terminates only the current iteration.

Keep in mind these facts about range-based `for`:

- Automatically recognizes arrays.
- Recognizes containers that have `.begin()` and `.end()`.
- Uses argument-dependent lookup `begin()` and `end()` for anything else.

## See Also

[auto](#)

[Iteration Statements](#)

[Keywords](#)

[while Statement \(C++\)](#)

[do-while Statement \(C++\)](#)

[for Statement \(C++\)](#)

# Jump Statements (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

A C++ jump statement performs an immediate local transfer of control.

## Syntax

```
break;  
continue;  
return [expression];  
goto identifier;
```

## Remarks

See the following topics for a description of the C++ jump statements.

- [break Statement](#)
- [continue Statement](#)
- [return Statement](#)
- [goto Statement](#)

## See Also

[Overview of C++ Statements](#)

# break Statement (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

The `break` statement ends execution of the nearest enclosing loop or conditional statement in which it appears. Control passes to the statement that follows the end of the statement, if any.

## Syntax

```
break;
```

## Remarks

The `break` statement is used with the conditional `switch` statement and with the `do`, `for`, and `while` loop statements.

In a `switch` statement, the `break` statement causes the program to execute the next statement outside the `switch` statement. Without a `break` statement, every statement from the matched `case` label to the end of the `switch` statement, including the `default` clause, is executed.

In loops, the `break` statement ends execution of the nearest enclosing `do`, `for`, or `while` statement. Control passes to the statement that follows the ended statement, if any.

Within nested statements, the `break` statement ends only the `do`, `for`, `switch`, or `while` statement that immediately encloses it. You can use a `return` or `goto` statement to transfer control from more deeply nested structures.

## Example

The following code shows how to use the `break` statement in a `for` loop.

```
#include <iostream>
using namespace std;

int main()
{
    // An example of a standard for loop
    for (int i = 1; i < 10; i++)
    {
        cout << i << '\n';
        if (i == 4)
            break;
    }

    // An example of a range-based for loop
    int nums []{1, 2, 3, 4, 5, 6, 7, 8, 9, 10};

    for (int i : nums) {
        if (i == 4) {
            break;
        }
        cout << i << '\n';
    }
}
```

In each case:

1  
2  
3

The following code shows how to use `break` in a `while` loop and a `do` loop.

```
#include <iostream>
using namespace std;

int main() {
    int i = 0;

    while (i < 10) {
        if (i == 4) {
            break;
        }
        cout << i << '\n';
        i++;
    }

    i = 0;
    do {
        if (i == 4) {
            break;
        }
        cout << i << '\n';
        i++;
    } while (i < 10);
}
```

In each case:

0123

The following code shows how to use `break` in a switch statement. You must use `break` in every case if you want to handle each case separately; if you do not use `break`, the code execution falls through to the next case.

```
#include <iostream>
using namespace std;

enum Suit{ Diamonds, Hearts, Clubs, Spades };

int main() {

    Suit hand;
    . . .
    // Assume that some enum value is set for hand
    // In this example, each case is handled separately
    switch (hand)
    {
        case Diamonds:
            cout << "got Diamonds \n";
            break;
        case Hearts:
            cout << "got Hearts \n";
            break;
        case Clubs:
            cout << "got Clubs \n";
            break;
        case Spades:
            cout << "got Spades \n";
            break;
        default:
            cout << "didn't get card \n";
    }
    // In this example, Diamonds and Hearts are handled one way, and
    // Clubs, Spades, and the default value are handled another way
    switch (hand)
    {
        case Diamonds:
        case Hearts:
            cout << "got a red card \n";
            break;
        case Clubs:
        case Spades:
        default:
            cout << "didn't get a red card \n";
    }
}
```

## See Also

[Jump Statements](#)

[Keywords](#)

[continue Statement](#)

# continue Statement (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Forces transfer of control to the controlling expression of the smallest enclosing `do`, `for`, or `while` loop.

## Syntax

```
continue;
```

## Remarks

Any remaining statements in the current iteration are not executed. The next iteration of the loop is determined as follows:

- In a `do` or `while` loop, the next iteration starts by reevaluating the controlling expression of the `do` or `while` statement.
- In a `for` loop (using the syntax `for ( init-expr ; cond-expr ; loop-expr )`), the `loop-expr` clause is executed. Then the `cond-expr` clause is reevaluated and, depending on the result, the loop either ends or another iteration occurs.

The following example shows how the `continue` statement can be used to bypass sections of code and begin the next iteration of a loop.

## Example

```
// continue_statement.cpp
#include <stdio.h>
int main()
{
    int i = 0;
    do
    {
        i++;
        printf_s("before the continue\n");
        continue;
        printf("after the continue, should never print\n");
    } while (i < 3);

    printf_s("after the do loop\n");
}
```

```
before the continue
before the continue
before the continue
after the do loop
```

## See Also

[Jump Statements](#)

[Keywords](#)

# return Statement (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Terminates the execution of a function and returns control to the calling function (or to the operating system if you transfer control from the `main` function). Execution resumes in the calling function at the point immediately following the call.

## Syntax

```
return [expression];
```

## Remarks

The `expression` clause, if present, is converted to the type specified in the function declaration, as if an initialization were being performed. Conversion from the type of the expression to the `return` type of the function can create temporary objects. For more information about how and when temporaries are created, see [Temporary Objects](#).

The value of the `expression` clause is returned to the calling function. If the expression is omitted, the return value of the function is undefined. Constructors and destructors, and functions of type `void`, cannot specify an expression in the `return` statement. Functions of all other types must specify an expression in the `return` statement.

When the flow of control exits the block enclosing the function definition, the result is the same as it would be if a `return` statement without an expression had been executed. This is invalid for functions that are declared as returning a value.

A function can have any number of `return` statements.

The following example uses an expression with a `return` statement to obtain the largest of two integers.

## Example

```
// return_statement2.cpp
#include <stdio.h>

int max ( int a, int b )
{
    return ( a > b ? a : b );
}

int main()
{
    int nOne = 5;
    int nTwo = 7;

    printf_s("\n%d is bigger\n", max( nOne, nTwo ) );
}
```

## See Also

[Jump Statements](#)

[Keywords](#)

# goto Statement (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The `goto` statement unconditionally transfers control to the statement labeled by the specified identifier.

## Syntax

```
goto identifier;
```

## Remarks

The labeled statement designated by `identifier` must be in the current function. All `identifier` names are members of an internal namespace and therefore do not interfere with other identifiers.

A statement label is meaningful only to a `goto` statement; otherwise, statement labels are ignored. Labels cannot be redeclared.

It is good programming style to use the `break`, `continue`, and `return` statements instead of the `goto` statement whenever possible. However, because the `break` statement exits from only one level of a loop, you might have to use a `goto` statement to exit a deeply nested loop.

For more information about labels and the `goto` statement, see [Labeled Statements](#) and [Using Labels with the goto Statement](#).

## Example

In this example, a `goto` statement transfers control to the point labeled `stop` when `i` equals 3.

```
// goto_statement.cpp
#include <stdio.h>
int main()
{
    int i, j;

    for ( i = 0; i < 10; i++ )
    {
        printf_s( "Outer loop executing. i = %d\n", i );
        for ( j = 0; j < 2; j++ )
        {
            printf_s( " Inner loop executing. j = %d\n", j );
            if ( i == 3 )
                goto stop;
        }
    }

    // This message does not print:
    printf_s( "Loop exited. i = %d\n", i );

    stop:
    printf_s( "Jumped to stop. i = %d\n", i );
}
```

```
Outer loop executing. i = 0
Inner loop executing. j = 0
Inner loop executing. j = 1
Outer loop executing. i = 1
Inner loop executing. j = 0
Inner loop executing. j = 1
Outer loop executing. i = 2
Inner loop executing. j = 0
Inner loop executing. j = 1
Outer loop executing. i = 3
Inner loop executing. j = 0
Jumped to stop. i = 3
```

## See Also

[Jump Statements](#)

[Keywords](#)

# Transfers of Control

12/21/2017 • 1 min to read • [Edit Online](#)

You can use the `goto` statement or a `case` label in a `switch` statement to specify a program that branches past an initializer. Such code is illegal unless the declaration that contains the initializer is in a block enclosed by the block in which the jump statement occurs.

The following example shows a loop that declares and initializes the objects `total`, `ch`, and `i`. There is also an erroneous `goto` statement that transfers control past an initializer.

```
// transfers_of_control.cpp
// compile with: /W1
// Read input until a nonnumeric character is entered.
int main()
{
    char MyArray[5] = {'2','2','a','c'};
    int i = 0;
    while( 1 )
    {
        int total = 0;

        char ch = MyArray[i++];

        if ( ch >= '0' && ch <= '9' )
        {
            goto Label1;

            int i = ch - '0';
Label1:
            total += i;    // C4700: transfers past initialization of i.
        } // i would be destroyed here if goto error were not present
        else
            // Break statement transfers control out of loop,
            // destroying total and ch.
            break;
    }
}
```

In the preceding example, the `goto` statement tries to transfer control past the initialization of `i`. However, if `i` were declared but not initialized, the transfer would be legal.

The objects `total` and `ch`, declared in the block that serves as the *statement* of the `while` statement, are destroyed when that block is exited using the `break` statement.

# Namespaces (C++)

12/21/2017 • 7 min to read • [Edit Online](#)

A namespace is a declarative region that provides a scope to the identifiers (the names of types, functions, variables, etc) inside it. Namespaces are used to organize code into logical groups and to prevent name collisions that can occur especially when your code base includes multiple libraries. All identifiers at namespace scope are visible to one another without qualification. Identifiers outside the namespace can access the members by using the fully qualified name for each identifier, for example `std::vector<std::string> vec;`, or else by a [using Declaration](#) for a single identifier (`using std::string`), or a [using Directive](#) for all the identifiers in the namespace (`using namespace std;`). Code in header files should always use the fully qualified namespace name.

The following example shows a namespace declaration and three ways that code outside the namespace can access their members.

```
namespace ContosoData
{
    class ObjectManager
    {
        public:
            void DoSomething() {}
    };
    void Func(ObjectManager) {}
}
```

Use the fully qualified name:

```
ContosoData::ObjectManager mgr;
mgr.DoSomething();
ContosoData::Func(mgr);
```

Use a using declaration to bring one identifier into scope:

```
using ContosoData::ObjectManager;
ObjectManager mgr;
mgr.DoSomething();
```

Use a using directive to bring everything in the namespace into scope:

```
using namespace ContosoData;

ObjectManager mgr;
mgr.DoSomething();
Func(mgr);
```

## using directives

The `using` directive allows all the names in a **namespace** to be used without the *namespace-name* as an explicit qualifier. Use a using directive in an implementation file (i.e. \*.cpp) if you are using several different identifiers in a namespace; if you are just using one or two identifiers, then consider a using declaration to only bring those identifiers into scope and not all the identifiers in the namespace. If a local variable has the same name as a namespace variable, the namespace variable is hidden. It is an error to have a namespace variable with the same

name as a global variable.

#### NOTE

A using directive can be placed at the top of a .cpp file (at file scope), or inside a class or function definition.

In general, avoid putting using directives in header files (\*.h) because any file that includes that header will bring everything in the namespace into scope, which can cause name hiding and name collision problems that are very difficult to debug. Always use fully qualified names in a header file. If those names get too long, you can use a namespace alias to shorten them. (See below.)

## Declaring namespaces and namespace members

Typically, you declare a namespace in a header file. If your function implementations are in a separate file, then qualify the function names, as in this example.

```
//contosoData.h
#pragma once
namespace ContosoDataServer
{
    void Foo();
    int Bar();
}
```

Function implementations in contosodata.cpp should use the fully qualified name, even if you place a using directive at the top of the file:

```
#include "contosodata.h"
using namespace ContosoDataServer;

void ContosoDataServer::Foo() // use fully-qualified name here
{
    // no qualification needed for Bar()
    Bar();
}

int ContosoDataServer::Bar(){return 0;}
```

A namespace can be declared in multiple blocks in a single file, and in multiple files. The compiler joins the parts together during preprocessing and the resulting namespace contains all the members declared in all the parts. An example of this is the std namespace which is declared in each of the header files in the standard library.

Members of a named namespace can be defined outside the namespace in which they are declared by explicit qualification of the name being defined. However, the definition must appear after the point of declaration in a namespace that encloses the declaration's namespace. For example:

```
// defining_namespace_members.cpp
// C2039 expected
namespace V {
    void f();
}

void V::f() { } // ok
void V::g() { } // C2039, g() is not yet a member of V

namespace V {
    void g();
}
}
```

This error can occur when namespace members are declared across multiple header files, and you have not included those headers in the correct order.

## The global namespace

If an identifier is not declared in an explicit namespace, it is part of the implicit global namespace. In general, try to avoid making declarations at global scope when possible, except for the entry point [main Function](#), which is required to be in the global namespace. To explicitly qualify a global identifier, use the scope resolution operator with no name, as in `::SomeFunction(x);`. This will differentiate the identifier from anything with the same name in any other namespace, and it will also help to make your code easier for others to understand.

## The std namespace

All C++ standard library types and functions are declared in the `std` namespace or namespaces nested inside `std`.

## Nested namespaces

Namespaces may be nested. An ordinary nested namespace has unqualified access to its parent's members, but the parent members do not have unqualified access to the nested namespace (unless it is declared as inline), as shown in the following example:

```
namespace ContosoDataServer
{
    void Foo();

    namespace Details
    {
        int CountImpl;
        void Bar() { return Foo(); }
    }

    int Baz(int i) { return Details::CountImpl; }
}
```

Ordinary nested namespaces can be used to encapsulate internal implementation details that are not part of the public interface of the parent namespace.

## Inline namespaces (C++ 11)

In contrast to an ordinary nested namespace, members of an inline namespace are treated as members of the

parent namespace. This characteristic enables argument dependent lookup on overloaded functions to work on functions that have overloads in a parent and a nested inline namespace. It also enables you to declare a specialization in a parent namespace for a template that is declared in the inline namespace. The following example shows how external code binds to the inline namespace by default:

```
//Header.h
#include <string>

namespace Test
{
    namespace old_ns
    {
        std::string Func() { return std::string("Hello from old"); }
    }

    inline namespace new_ns
    {
        std::string Func() { return std::string("Hello from new"); }
    }
}

#include "header.h"
#include <string>
#include <iostream>

int main()
{
    using namespace Test;
    using namespace std;

    string s = Func();
    std::cout << s << std::endl; // "Hello from new"
    return 0;
}
```

The following example shows how you can declare a specialization in a parent of a template that is declared in an inline namespace:

```
namespace Parent
{
    inline namespace new_ns
    {
        template <typename T>
        struct C
        {
            T member;
        };
    }
    template<>
    class C<int> {};
}
```

You can use inline namespaces as a versioning mechanism to manage changes to the public interface of a library. For example, you can create a single parent namespace, and encapsulate each version of the interface in its own namespace nested inside the parent. The namespace that holds the most recent or preferred version is qualified as inline, and is therefore exposed as if it were a direct member of the parent namespace. Client code that invokes the Parent::Class will automatically bind to the new code. Clients that prefer to use the older version can still access it by using the fully qualified path to the nested namespace that has that code.

The inline keyword must be applied to the first declaration of the namespace in a compilation unit.

The following example shows two versions of an interface, each in a nested namespace. The v\_20 namespace has

some modification from the `v_10` interface and is marked as inline. Client code that uses the new library and calls `Contoso::Funcs::Add` will invoke the `v_20` version. Code that attempts to call `Contoso::Funcs::Divide` will now get a compile time error. If they really need that function, they can still access the `v_10` version by explicitly calling `Contoso::v_10::Funcs::Divide`.

```
namespace Contoso
{
    namespace v_10
    {
        template <typename T>
        class Funcs
        {
        public:
            Funcs(void);
            T Add(T a, T b);
            T Subtract(T a, T b);
            T Multiply(T a, T b);
            T Divide(T a, T b);
        };
    }

    inline namespace v_20
    {
        template <typename T>
        class Funcs
        {
        public:
            Funcs(void);
            T Add(T a, T b);
            T Subtract(T a, T b);
            T Multiply(T a, T b);
            std::vector<double> Log(double);
            T Accumulate(std::vector<T> nums);
        };
    }
}
```

## Namespace aliases

Namespace names need to be unique, which means that often they should not be too short. If the length of a name makes code difficult to read, or is tedious to type in a header file where using directives can't be used, then you can make a namespace alias which serves as an abbreviation for the actual name. For example:

```
namespace a_very_long_namespace_name { class Foo {}; }
namespace AVLNN = a_very_long_namespace_name;
void Bar(AVLNN::Foo foo){ }
```

## anonymous or unnamed namespaces

You can create an explicit namespace but not give it a name:

```
namespace
{
    int MyFunc(){}
}
```

This is called an unnamed or anonymous namespace and it is useful when you want to make variable declarations invisible to code in other files (i.e. give them internal linkage) without having to create a named namespace. All

code in the same file can see the identifiers in an unnamed namespace but the identifiers, along with the namespace itself, are not visible outside that file—or more precisely outside the translation unit.

## See Also

[Declarations and Definitions](#)

# Enumerations (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

An enumeration is a user-defined type that consists of a set of named integral constants that are known as enumerators.

## NOTE

This article covers the ISO Standard C++ Language `enum` type and the scoped (or strongly-typed) `enum class` type which is introduced in C++11. For information about the `public enum class` or `private enum class` types in C++/CLI and C++/CX, see [enum class](#).

## Syntax

```
// unscoped enum:  
enum [identifier] [: type]  
{enum-list};  
  
// scoped enum:  
enum [class|struct]  
[identifier] [: type]  
{enum-list};
```

```
// Forward declaration of enumerations (C++11):  
enum A : int; // non-scoped enum must have type specified  
enum class B; // scoped enum defaults to int but ...  
enum class C : short; // ... may have any integral underlying type
```

## Parameters

`identifier`

The type name given to the enumeration.

`type`

The underlying type of the enumerators; all enumerators have the same underlying type. May be any integral type.

`enum-list`

Comma-separated list of the enumerators in the enumeration. Every enumerator or variable name in the scope must be unique. However, the values can be duplicated. In a unscoped enum, the scope is the surrounding scope; in a scoped enum, the scope is the `enum-list` itself. In a scoped enum, the list may be empty which in effect defines a new integral type.

`class`

By using this keyword in the declaration, you specify the enum is scoped, and an `identifier` must be provided. You can also use the `struct` keyword in place of `class`, as they are semantically equivalent in this context.

## Enumerator scope

An enumeration provides context to describe a range of values which are represented as named constants and are also called enumerators. In the original C and C++ enum types, the unqualified enumerators are visible

throughout the scope in which the enum is declared. In scoped enums, the enumerator name must be qualified by the enum type name. The following example demonstrates this basic difference between the two kinds of enums:

```
namespace CardGame_Scoped
{
    enum class Suit { Diamonds, Hearts, Clubs, Spades };

    void PlayCard(Suit suit)
    {
        if (suit == Suit::Clubs) // Enumerator must be qualified by enum type
        { /*...*/}
    }
}

namespace CardGame_NonScoped
{
    enum Suit { Diamonds, Hearts, Clubs, Spades };

    void PlayCard(Suit suit)
    {
        if (suit == Clubs) // Enumerator is visible without qualification
        { /*...*/}
    }
}
```

Every name in an enumeration is assigned an integral value that corresponds to its place in the order of the values in the enumeration. By default, the first value is assigned 0, the next one is assigned 1, and so on, but you can explicitly set the value of an enumerator, as shown here:

```
enum Suit { Diamonds = 1, Hearts, Clubs, Spades };
```

The enumerator `Diamonds` is assigned the value `1`. Subsequent enumerators, if they are not given an explicit value, receive the value of the previous enumerator plus one. In the previous example, `Hearts` would have the value `2`, `Clubs` would have `3`, and so on.

Every enumerator is treated as a constant and must have a unique name within the scope where the `enum` is defined (for unscoped enums) or within the enum itself (for scoped enums). The values given to the names do not have to be unique. For example, if the declaration of a unscoped enum `Suit` is this:

```
enum Suit { Diamonds = 5, Hearts, Clubs = 4, Spades };
```

Then the values of `Diamonds`, `Hearts`, `Clubs`, and `Spades` are 5, 6, 4, and 5, respectively. Notice that 5 is used more than once; this is allowed even though it may not be intended. These rules are the same for scoped enums.

## Casting rules

Unscoped enum constants can be implicitly converted to `int`, but an `int` is never implicitly convertible to an enum value. The following example shows what happens if you try to assign `hand` a value that is not a `Suit`:

```
int account_num = 135692;
Suit hand;
hand = account_num; // error C2440: '=' : cannot convert from 'int' to 'Suit'
```

A cast is required to convert an `int` to a scoped or unscoped enumerator. However, you can promote a unscoped enumerator to an integer value without a cast.

```
int account_num = Hearts; //OK if Hearts is in a unscoped enum
```

Using implicit conversions in this way can lead to unintended side-effects. To help eliminate programming errors associated with unscoped enums, scoped enum values are strongly typed. Scoped enumerators must be qualified by the enum type name (identifier) and cannot be implicitly converted, as shown in the following example:

```
namespace ScopedEnumConversions
{
    enum class Suit { Diamonds, Hearts, Clubs, Spades };

    void AttemptConversions()
    {
        Suit hand;
        hand = Clubs; // error C2065: 'Clubs' : undeclared identifier
        hand = Suit::Clubs; //Correct.
        int account_num = 135692;
        hand = account_num; // error C2440: '=' : cannot convert from 'int' to 'Suit'
        hand = static_cast<Suit>(account_num); // OK, but probably a bug!!!

        account_num = Suit::Hearts; // error C2440: '=' : cannot convert from 'Suit' to 'int'
        account_num = static_cast<int>(Suit::Hearts); // OK
    }
}
```

Notice that the line `hand = account_num;` still causes the error that occurs with unscoped enums, as shown earlier. It is allowed with an explicit cast. However, with scoped enums, the attempted conversion in the next statement, `account_num = Suit::Hearts;`, is no longer allowed without an explicit cast.

## Enums with no enumerators

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): By defining an enum (regular or scoped) with an explicit underlying type and no enumerators, you can in effect introduce a new integral type that has no implicit conversion to any other type. By using this type instead of its built-in underlying type, you can eliminate the potential for subtle errors caused by inadvertent implicit conversions.

```
enum class byte : unsigned char { };
```

The new type is an exact copy of the underlying type, and therefore has the same calling convention, which means it can be used across ABIs without any performance penalty. No cast is required when variables of the type are initialized by using direct-list initialization. The following example shows how to initialize enums with no enumerators in various contexts:

```
enum class byte : unsigned char { };

enum class E : int { };
E e1{ 0 };
E e2 = E{ 0 };

struct X
{
    E e{ 0 };
    X() : e{ 0 } { }
};

E* p = new E{ 0 };

void f(E e) {}

int main()
{
    f(E{ 0 });
    byte i{ 42 };
    byte j = byte{ 42 };

    // unsigned char c = j; // C2440: 'initializing': cannot convert from 'byte' to 'unsigned char'
    return 0;
}
```

## See Also

[C Enumeration Declarations](#)

[Keywords](#)

# Unions

12/21/2017 • 9 min to read • [Edit Online](#)

A `union` is a user-defined type in which all members share the same memory location. This means that at any given time a union can contain no more than one object from its list of members. It also means that no matter how many members a union has, it always uses only enough memory to store the largest member.

Unions can be useful for conserving memory when you have lots of objects and/or limited memory. However they require extra care to use correctly because you are responsible for ensuring that you always access the last member that was written to. If any member types have a non-trivial constructor, then you must write additional code to explicitly construct and destroy that member. Before using a union, consider whether the problem you are trying to solve could be better expressed by using a base class and derived classes.

## Syntax

```
union [name] { member-list };
```

### Parameters

`name`

The type name given to the union.

`member-list`

Members that the union can contain. See Remarks.

## Remarks

### Declaring a Union

Begin the declaration of a union with the `union` keyword, and enclose the member list in curly braces:

```
// declaring_a_union.cpp
union RecordType // Declare a simple union type
{
    char    ch;
    int     i;
    long    l;
    float   f;
    double  d;
    int *int_ptr;
};
int main()
{
    RecordType t;
    t.i = 5; // t holds an int
    t.f = 7.25 // t now holds a float
}
```

## Using unions

In the previous example, any code that accesses the union needs to know which member is holding the data. The most common solution to this problem is to enclose the union in a struct along with an additional enum member

that indicates the type of the data currently being stored in the union. This is called a *discriminated union* and the following example shows the basic pattern.

```
#include "stdafx.h"
#include <queue>

using namespace std;

enum class WeatherDataType
{
    Temperature, Wind
};

struct TempData
{
    int StationId;
    time_t time;
    double current;
    double max;
    double min;
};

struct WindData
{
    int StationId;
    time_t time;
    int speed;
    short direction;
};

struct Input
{
    WeatherDataType type;
    union
    {
        TempData temp;
        WindData wind;
    };
};

// Functions that are specific to data types
void Process_Temp(TempData t) {}
void Process_Wind(WindData w) {}

// Container for all the data records
queue<Input> inputs;
void Initialize();

int main(int argc, char* argv[])
{
    Initialize();
    while (!inputs.empty())
    {
        Input i = inputs.front();
        switch (i.type)
        {
            case WeatherDataType::Temperature:
                Process_Temp(i.temp);
                break;
            case WeatherDataType::Wind:
                Process_Wind(i.wind);
                break;
            default:
                break;
        }
        inputs.pop();
    }
}
```

```

        }

    return 0;
}

void Initialize()
{
    Input first, second;
    first.type = WeatherDataType::Temperature;
    first.temp = { 101, 1418855664, 91.8, 108.5, 67.2 };
    inputs.push(first);

    second.type = WeatherDataType::Wind;
    second.wind = { 204, 1418859354, 14, 27 };
    inputs.push(second);
}

```

In the previous example, note that the union in the Input struct has no name. This is an anonymous union and its members can be accessed as if they were direct members of the struct. For more information about anonymous unions, see the section below.

Of course, the previous example shows a problem that could also be solved by using classes that derive from a common base class, and branching your code based on the runtime type of each object in the container. This may result in code that easier to maintain and understand, but it might also be slower than using unions. Also, with a union, you can store completely unrelated types, and dynamically change the type of the value that is stored without changing the type of the union variable itself. Thus you can create a heterogeneous array of MyUnionType whose elements store different values of different types.

Note that the `Input` struct in the preceding example can be easily misused. It is completely up to the user to use the discriminator correctly to access the member that holds the data. You can protect against misuse by making the union private and providing special access functions, as shown in the next example.

## Unrestricted Unions (C++11)

In C++03 and earlier a union can contain non-static data members with class type as long as the type has no user provided constructors, destructors or assignment operators. In C++11, these restrictions are removed. If you include such a member in your union then the compiler will automatically mark any special member functions that are not user provided as deleted. If the union is an anonymous union inside a class or struct, then any special member functions of the class or struct that are not user provided are marked as deleted. The following example shows how to handle the case where one of the members of the union has a member that requires this special treatment:

```

// for MyVariant
#include <crtdbg.h>
#include <new>
#include <utility>

// for sample objects and output
#include <string>
#include <vector>
#include <iostream>

using namespace std;

struct A
{
    A() = default;
    A(int i, const string& str) : num(i), name(str) {}

    int num;
    string name;
    //...
};

```

```
,,  
struct B  
{  
    B() = default;  
    B(int i, const string& str) : num(i), name(str) {}  
  
    int num;  
    string name;  
    vector<int> vec;  
    // ...  
};  
  
enum class Kind { None, A, B, Integer };  
  
#pragma warning(push)  
#pragma warning(disable:4624)  
class MyVariant  
{  
public:  
    MyVariant()  
        : kind_(Kind::None)  
    {}  
  
    MyVariant(Kind kind)  
        : kind_(kind)  
    {}  
    switch (kind_)  
    {  
        case Kind::None:  
            break;  
        case Kind::A:  
            new (&a_) A();  
            break;  
        case Kind::B:  
            new (&b_) B();  
            break;  
        case Kind::Integer:  
            i_ = 0;  
            break;  
        default:  
            _ASSERT(false);  
            break;  
    }  
}  
  
~MyVariant()  
{  
    switch (kind_)  
    {  
        case Kind::None:  
            break;  
        case Kind::A:  
            a_.~A();  
            break;  
        case Kind::B:  
            b_.~B();  
            break;  
        case Kind::Integer:  
            break;  
        default:  
            _ASSERT(false);  
            break;  
    }  
    kind_ = Kind::None;  
}  
  
MyVariant(const MyVariant& other)  
    . kind_(other.kind)
```

```

    . kind_(other.kind_)

{
    switch (kind_)
    {
        case Kind::None:
            break;
        case Kind::A:
            new (&a_) A(other.a_);
            break;
        case Kind::B:
            new (&b_) B(other.b_);
            break;
        case Kind::Integer:
            i_ = other.i_;
            break;
        default:
            _ASSERT(false);
            break;
    }
}

MyVariant(MyVariant&& other)
    : kind_(other.kind_)
{
    switch (kind_)
    {
        case Kind::None:
            break;
        case Kind::A:
            new (&a_) A(move(other.a_));
            break;
        case Kind::B:
            new (&b_) B(move(other.b_));
            break;
        case Kind::Integer:
            i_ = other.i_;
            break;
        default:
            _ASSERT(false);
            break;
    }
    other.kind_ = Kind::None;
}

MyVariant& operator=(const MyVariant& other)
{
    if (&other != this)
    {
        switch (other.kind_)
        {
            case Kind::None:
                this->~MyVariant();
                break;
            case Kind::A:
                *this = other.a_;
                break;
            case Kind::B:
                *this = other.b_;
                break;
            case Kind::Integer:
                *this = other.i_;
                break;
            default:
                _ASSERT(false);
                break;
        }
    }
    return *this;
}

```

```

MyVariant& operator=(MyVariant&& other)
{
    _ASSERT(this != &other);
    switch (other.kind_)
    {
        case Kind::None:
            this->~MyVariant();
            break;
        case Kind::A:
            *this = move(other.a_);
            break;
        case Kind::B:
            *this = move(other.b_);
            break;
        case Kind::Integer:
            *this = other.i_;
            break;
        default:
            _ASSERT(false);
            break;
    }
    other.kind_ = Kind::None;
    return *this;
}

MyVariant(const A& a)
: kind_(Kind::A), a_(a)
{
}

MyVariant(A&& a)
: kind_(Kind::A), a_(move(a))
{
}

MyVariant& operator=(const A& a)
{
    if (kind_ != Kind::A)
    {
        this->~MyVariant();
        new (this) MyVariant(a);
    }
    else
    {
        a_ = a;
    }
    return *this;
}

MyVariant& operator=(A&& a)
{
    if (kind_ != Kind::A)
    {
        this->~MyVariant();
        new (this) MyVariant(move(a));
    }
    else
    {
        a_ = move(a);
    }
    return *this;
}

MyVariant(const B& b)
: kind_(Kind::B), b_(b)
{
}

```

```

MyVariant(B& b)
    : kind_(Kind::B), b_(move(b))
{
}

MyVariant& operator=(const B& b)
{
    if (kind_ != Kind::B)
    {
        this->~MyVariant();
        new (this) MyVariant(b);
    }
    else
    {
        b_ = b;
    }
    return *this;
}

MyVariant& operator=(B&& b)
{
    if (kind_ != Kind::B)
    {
        this->~MyVariant();
        new (this) MyVariant(move(b));
    }
    else
    {
        b_ = move(b);
    }
    return *this;
}

MyVariant(int i)
    : kind_(Kind::Integer), i_(i)
{
}

MyVariant& operator=(int i)
{
    if (kind_ != Kind::Integer)
    {
        this->~MyVariant();
        new (this) MyVariant(i);
    }
    else
    {
        i_ = i;
    }
    return *this;
}

Kind GetKind() const
{
    return kind_;
}

A& GetA()
{
    _ASSERT(kind_ == Kind::A);
    return a_;
}

const A& GetA() const
{
    _ASSERT(kind_ == Kind::A);
    return a_;
}

```

```

B& GetB()
{
    _ASSERT(kind_ == Kind::B);
    return b_;
}

const B& GetB() const
{
    _ASSERT(kind_ == Kind::B);
    return b_;
}

int& GetInteger()
{
    _ASSERT(kind_ == Kind::Integer);
    return i_;
}

const int& GetInteger() const
{
    _ASSERT(kind_ == Kind::Integer);
    return i_;
}

private:
    Kind kind_;
    union
    {
        A a_;
        B b_;
        int i_;
    };
};

#pragma warning (pop)

int main()
{
    A a(1, "Hello from A");
    B b(2, "Hello from B");

    MyVariant mv_1 = a;

    cout << "mv_1 = a: " << mv_1.GetA().name << endl;
    mv_1 = b;
    cout << "mv_1 = b: " << mv_1.GetB().name << endl;
    mv_1 = A(3, "hello again from A");
    cout << R"aaa(mv_1 = A(3, "hello again from A"): )aaa" << mv_1.GetA().name << endl;
    mv_1 = 42;
    cout << "mv_1 = 42: " << mv_1.GetInteger() << endl;

    b.vec = { 10,20,30,40,50 };

    mv_1 = move(b);
    cout << "After move, mv_1 = b: vec.size = " << mv_1.GetB().vec.size() << endl;

    cout << endl << "Press a letter" << endl;
    char c;
    cin >> c;
}

#include <queue>
#include <iostream>
using namespace std;

enum class WeatherDataType
{
    Temperature, Wind
};

struct TempData

```

```

{
    TempData() : StationId(""), time(0), current(0), maxTemp(0), minTemp(0) {}
    TempData(string id, time_t t, double cur, double max, double min)
        : StationId(id), time(t), current(cur), maxTemp(max), minTemp(0) {}
    string StationId;
    time_t time = 0;
    double current;
    double maxTemp;
    double minTemp;
};

struct WindData
{
    int StationId;
    time_t time;
    int speed;
    short direction;
};

struct Input
{
    Input() {}
    Input(const Input&) {}

    ~Input()
    {
        if (type == WeatherDataType::Temperature)
        {
            temp.StationId.~string();
        }
    }

    WeatherDataType type;
    void SetTemp(const TempData& td)
    {
        type = WeatherDataType::Temperature;

        // must use placement new because of string member!
        new(&temp) TempData(td);
    }

    TempData GetTemp()
    {
        if (type == WeatherDataType::Temperature)
            return temp;
        else
            throw logic_error("Can't return TempData when Input holds a WindData");
    }
    void SetWind(WindData wd)
    {
        // Explicitly delete struct member that has a
        // non-trivial constructor
        if (type == WeatherDataType::Temperature)
        {
            temp.StationId.~string();
        }
        wind = wd; //placement new not required.
    }
    WindData GetWind()
    {
        if (type == WeatherDataType::Wind)
        {
            return wind;
        }
        else
            throw logic_error("Can't return WindData when Input holds a TempData");
    }
};

private:

```

```

union
{
    TempData temp;
    WindData wind;
};

};

```

Unions cannot store references. Unions don't support inheritance, therefore a union itself cannot be used as a base class, or inherit from another class, or have virtual functions.

## Initializing unions

You can declare and initialize a union in the same statement by assigning an expression enclosed in braces. The expression is evaluated and assigned to the first field of the union.

```

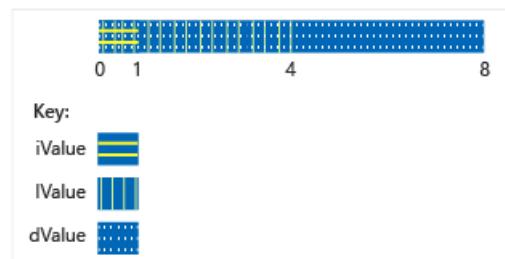
#include <iostream>
using namespace std;

union NumericType
{
    short      iValue;
    long       lValue;
    double     dValue;
};

int main()
{
    union NumericType Values = { 10 }; // iValue = 10
    cout << Values.iValue << endl;
    Values.dValue = 3.1416;
    cout << Values.dValue) << endl;
}
/* Output:
10
3.141600
*/

```

The `NumericType` union is arranged in memory (conceptually) as shown in the following figure.



Storage of Data in NumericType Union

## Anonymous unions

Anonymous unions are unions that are declared without a *class-name* or *declarator-list*.

```
union { member-list }
```

Names declared in an anonymous union are used directly, like nonmember variables. Therefore, the names declared in an anonymous union must be unique in the surrounding scope.

In addition to the restrictions for named unions, anonymous unions are subject to these additional restrictions:

- They must also be declared as **static** if declared in file or namespace scope.
- They can have only public members; private and protected members in anonymous unions generate errors.
- They cannot have member functions.

## See Also

[Classes and Structs](#)

[Keywords](#)

[class](#)

[struct](#)

# Functions (C++)

12/21/2017 • 12 min to read • [Edit Online](#)

A function is a block of code that performs some operation. A function can optionally define input parameters that enable callers to pass arguments into the function. A function can optionally return a value as output. Functions are useful for encapsulating common operations in a single reusable block, ideally with a name that clearly describes what the function does. The following function accepts two integers from a caller and returns their sum; `a` and `b` are *parameters* of type `int`.

```
int sum(int a, int b)
{
    return a + b;
}
```

The function can be *invoked*, or *called*, from any number of places in the program. The values that are passed to the function are the *arguments*, whose types must be compatible with the parameter types in the function definition.

```
int main()
{
    int i = sum(10, 32);
    int j = sum(i, 66);
    cout << "The value of j is" << j << endl; // 108
}
```

There is no practical limit to function length, but good design aims for functions that perform a single well-defined task. Complex algorithms should be broken up into easy-to-understand simpler functions whenever possible.

Functions that are defined at class scope are called member functions. In C++, unlike other languages, a function can also be defined at namespace scope (including the implicit global namespace). Such functions are called *free functions* or *non-member functions*; they are used extensively in the Standard Library.

## Parts of a function declaration

A minimal function *declaration* consists of the return type, function name, and parameter list (which may be empty), along with optional keywords that provide additional instructions to the compiler. The following example is a function declaration:

```
int sum(int a, int b);
```

A function definition consists of a declaration, plus the *body*, which is all the code between the curly braces:

```
int sum(int a, int b)
{
    return a + b;
}
```

A function declaration followed by a semicolon may appear in multiple places in a program. It must appear prior to any calls to that function in each translation unit. The function definition must appear only once in the program, according to the One Definition Rule (ODR).

The required parts of a function declaration are:

1. The return type, which specifies the type of the value that the function returns, or `void` if no value is returned. In C++11, `auto` is a valid return type that instructs the compiler to infer the type from the return statement. In C++14, `decltype(auto)` is also allowed. For more information, see Type Deduction in Return Types below.
2. The function name, which must begin with a letter or underscore and cannot contain spaces. In general, leading underscores in the Standard Library function names indicate private member functions, or non-member functions that are not intended for use by your code.
3. The parameter list, a brace delimited, comma-separated set of zero or more parameters that specify the type and optionally a local name by which the values may be accessed inside the function body.

Optional parts of a function declaration are:

4. `constexpr`, which indicates that the return value of the function is a constant value can be computed at compile time.

```
constexpr float exp(float x, int n)
{
    return n == 0 ? 1 :
        n % 2 == 0 ? exp(x * x, n / 2) :
        exp(x * x, (n - 1) / 2) * x;
};
```

5. Its `linkage` specification, `extern` or `static`.

```
Declare printf with C linkage.
extern "C" int printf( const char *fmt, ... );
```

For more information, see [Program and Linkage](#).

6. `inline`, which instructs the compiler to replace every call to the function with the function code itself. inlining can help performance in scenarios where a function executes quickly and is invoked repeatedly in a performance-critical section of code.

```
inline double Account::GetBalance()
{
    return balance;
}
```

For more information, see [Inline Functions](#).

7. A `noexcept` expression, which specifies whether or not the function can throw an exception. In the following example, the function does not throw an exception if the `is_pod` expression evaluates to `true`.

```
#include <type_traits>

template <typename T>
T copy_object(T& obj) noexcept(std::is_pod<T>) {...}
```

For more information, see [noexcept](#).

8. (Member functions only) The cv-qualifiers, which specify whether the function is `const` or `volatile`.
9. (Member functions only) `virtual`, `override`, or `final`. `virtual` specifies that a function can be

overridden in a derived class. `override` means that a function in a derived class is overriding a virtual function. `final` means a function cannot be overridden in any further derived class. For more information, see [Virtual Functions](#).

10. (member functions only) `static` applied to a member function means that the function is not associated with any object instances of the class.

11. (Non-static member functions only) The ref-qualifier, which specifies to the compiler which overload of a function to choose when the implicit object parameter (`*this`) is an rvalue reference vs. an lvalue reference.

The following figure shows the parts of a function definition. The shaded area is the function body.



Parts of a Function Definition

## Function definitions

Variables declared inside the body are called local variables or locals. They go out of scope when the function exits; therefore, a function should never return a reference to a local!

## Function Templates

A function template is similar to a class template; it generates concrete functions based on the template arguments. In many cases, the template is able to infer the type arguments and therefore it isn't necessary to explicitly specify them.

```
template<typename Lhs, typename Rhs>
auto Add2(const Lhs& lhs, const Rhs& rhs)
{
    return lhs + rhs;
}

auto a = Add2(3.13, 2.895); // a is a double
auto b = Add2(string{ "Hello" }, string{ " World" }); // b is a std::string
```

For more information, see [Function Templates](#)

## Function parameters and arguments

A function has a comma-separated parameter list of zero or more types, each of which has a name by which it can be accessed inside the function body. A function template may specify additional type or value parameters. The caller passes arguments, which are concrete values whose types are compatible with the parameter list.

By default, arguments are passed to the function by value, which means the function receives a copy of the object being passed. For large objects, making a copy can be expensive and is not always necessary. To cause arguments to be passed by reference (specifically lvalue reference), add a reference qualifier to the parameter:

```
void DoSomething(std::string& input){...}
```

When a function modifies an argument that is passed by reference, it modifies the original object, not a local copy. To prevent a function from modifying such an argument, qualify the parameter as `const&`:

```
void DoSomething(const std::string& input){...}
```

**C++ 11:** To explicitly handle arguments that are passed by rvalue-reference OR by lvalue-reference, use a double-ampersand on the parameter to indicate a universal reference:

```
void DoSomething(const std::string&& input){...}
```

A function declared with the single keyword `void` in the parameter declaration list takes no arguments, as long as the keyword `void` is the first and only member of the argument declaration list. Arguments of type `void` elsewhere in the list produce errors. For example:

```
// OK same as GetTickCount()
long GetTickCount( void );
```

Note that, while it is illegal to specify a `void` argument except as outlined here, types derived from type `void` (such as pointers to `void` and arrays of `void`) can appear anywhere the argument declaration list.

## Default Arguments

The last parameter or parameters in a function signature may be assigned a default argument, which means that the caller may leave out the argument when calling the function unless they want to specify some other value.

```
int DoSomething(int num,
    string str,
    Allocator& alloc = defaultAllocator)
{ ... }

// OK both parameters are at end
int DoSomethingElse(int num,
    string str = string{ "Working" },
    Allocator& alloc = defaultAllocator)
{ ... }

// C2548: 'DoMore': missing default parameter for parameter 2
int DoMore(int num = 5, // Not a trailing parameter!
    string str,
    Allocator& = defaultAllocator)
{...}
```

For more information, see [Default Arguments](#).

## Function return types

A function may not return another function, or a built-in array; however it can return pointers to these types, or a *lambda*, which produces a function object. Except for these cases, a function may return a value of any type that is in scope, or it may return no value, in which case the return type is `void`.

### Trailing return types

An "ordinary" return type is located on the left side of the function signature. A *trailing return type* is located on the right most side of the signature and is preceded by the `->` operator. Trailing return types are especially useful in function templates when the type of the return value depends on template parameters.

```
template<typename Lhs, typename Rhs>
auto Add(const Lhs& lhs, const Rhs& rhs) -> decltype(lhs + rhs)
{
    return lhs + rhs;
}
```

When `auto` is used in conjunction with a trailing return type, it just serves as a placeholder for whatever the `decltype` expression produces, and does not itself perform type deduction.

## Function local variables

A variable that is declared inside a function body is called a *local variable* or simply a *local*. Non-static locals are only visible inside the function body and, if they are declared on the stack go out of scope when the function exits. When you construct a local variable and return it by value, the compiler can usually perform the return value optimization to avoid unnecessary copy operations. If you return a local variable by reference, the compiler will issue a warning because any attempt by the caller to use that reference will occur after the local has been destroyed.

Local static objects are destroyed during termination specified by `atexit`. If a static object was not constructed because the program's flow of control bypassed its declaration, no attempt is made to destroy that object.

### Static local variables

In C++ a local variable may be declared as static. The variable is only visible inside the function body, but a single copy of the variable exists for all instances of the function.

### Type deduction in return types (C++14)

In C++14, you can use `auto` to instruct the compiler to infer the return type from the function body without having to provide a trailing return type. Note that `auto` always deduces to a return-by-value. Use `auto&&` to instruct the compiler to deduce a reference.

In this example, `auto` will be deduced as a non-const value copy of the sum of `lhs` and `rhs`.

```
template<typename Lhs, typename Rhs>
auto Add2(const Lhs& lhs, const Rhs& rhs)
{
    return lhs + rhs; //returns a non-const object by value
}
```

Note that `auto` does not preserve the const-ness of the type it deduces. For forwarding functions whose return value needs to preserve the const-ness or ref-ness of its arguments, you can use the `decltype(auto)` keyword, which uses the `decltype` type inference rules and preserves all the type information. `decltype(auto)` may be used as an ordinary return value on the left side, or as a trailing return value.

The following example (based on code from [N3493](#)), shows `decltype(auto)` being used to enable perfect forwarding of function arguments in a return type that isn't known until the template is instantiated.

```

template<typename F, typename Tuple = tuple<T...>, int... I>
decltype(auto) apply_(F&& f, Tuple&& args, index_sequence<I...>)
{
    return std::forward<F>(f)(std::get<I>(std::forward<Tuple>(args))...);
}

template<typename F, typename Tuple = tuple<T...>,
         typename Indices = make_index_sequence<tuple_size<Tuple>::value >>
decltype( auto )
apply(F&& f, Tuple&& args)
{
    return apply_(std::forward<F>(f), std::forward<Tuple>(args), Indices());
}

```

## Returning multiple values from a function

There are various ways to return more than one value from a function:

1. Encapsulate the values in a named class or struct object. Requires the class or struct definition to be visible to the caller:

```

#include <string>
#include <iostream>

using namespace std;

struct S
{
    string name;
    int num;
};

S g()
{
    string t{ "hello" };
    int u{ 42 };
    return { t, u };
}

int main()
{
    S s = g();
    cout << s.name << " " << s.num << endl;
    return 0;
}

```

1. Return a std::tuple or std::pair object:

```

#include <tuple>
#include <string>
#include <iostream>

using namespace std;

tuple<int, string, double> f()
{
    int i{ 108 };
    string s{ "Some text" };
    double d{ .01 };
    return { i,s,d };
}

int main()
{
    auto t = f();
    cout << get<0>(t) << " " << get<1>(t) << " " << get<2>(t) << endl;

    // --or--

    int myval;
    string myname;
    double mydecimal;
    tie(myval, myname, mydecimal) = f();
    cout << myval << " " << myname << " " << mydecimal << endl;

    return 0;
}

```

1. **Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): Use structured bindings. The advantage of structured bindings is that the variables that store the return values are initialized at the same time they are declared, which in some cases can be significantly more efficient. In this statement -- `auto[x, y, z] = f();` -- the brackets introduce and initialize names that are in scope for the entire function block.

```

#include <tuple>
#include <string>
#include <iostream>

using namespace std;

tuple<int, string, double> f()
{
    int i{ 108 };
    string s{ "Some text" };
    double d{ .01 };
    return { i,s,d };
}

struct S
{
    string name;
    int num;
};

S g()
{
    string t{ "hello" };
    int u{ 42 };
    return { t, u };
}

int main()
{
    auto[x, y, z] = f(); // init from tuple
    cout << x << " " << y << " " << z << endl;

    auto[a, b] = g(); // init from POD struct
    cout << a << " " << b << endl;
    return 0;
}

```

1. In addition to using the return value itself, you can "return" values by defining any number of parameters to use pass-by-reference so that the function can modify or initialize the values of objects that the caller provides. For more information, see [Reference-Type Function Arguments](#).

## Function pointers

C++ supports function pointers in the same manner as the C language. However a more type-safe alternative is usually to use a function object.

It is recommended that `typedef` be used to declare an alias for the function pointer type if declaring a function that returns a function pointer type. For example

```

typedef int (*fp)(int);
fp myFunction(char* s); // function returning function pointer

```

If this is not done, the proper syntax for the function declaration may be deduced from the declarator syntax for the function pointer by replacing the identifier (`fp` in the above example) with the functions name and argument list, as follows:

```

int (*myFunction(char* s))(int);

```

The preceding declaration is equivalent to the declaration using `typedef` above.

## See Also

- [Function Overloading](#)
- [Functions with Variable Argument Lists](#)
- [Explicitly Defaulted and Deleted Functions](#)
- [Argument-Dependent Name \(Koenig\) Lookup on Functions](#)
- [Default Arguments](#)
- [Inline Functions](#)

# Functions with Variable Argument Lists (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Function declarations in which the last member of is the ellipsis (...) can take a variable number of arguments. In these cases, C++ provides type checking only for the explicitly declared arguments. You can use variable argument lists when you need to make a function so general that even the number and types of arguments can vary. The family of functions is an example of functions that use variable argument lists. `printf argument-declaration-list`

## Functions with variable arguments

To access arguments after those declared, use the macros contained in the standard include file STDARG.H as described below.

### Microsoft Specific

Microsoft C++ allows the ellipsis to be specified as an argument if the ellipsis is the last argument and the ellipsis is preceded by a comma. Therefore, the declaration `int Func( int i, ... );` is legal, but `int Func( int i ... );` is not.

### END Microsoft Specific

Declaration of a function that takes a variable number of arguments requires at least one placeholder argument, even if it is not used. If this placeholder argument is not supplied, there is no way to access the remaining arguments.

When arguments of type `char` are passed as variable arguments, they are converted to type `int`. Similarly, when arguments of type `float` are passed as variable arguments, they are converted to type `double`. Arguments of other types are subject to the usual integral and floating-point promotions. See [Standard Conversions](#) for more information.

Functions that require variable lists are declared by using the ellipsis (...) in the argument list. Use the types and macros that are described in the STDARG.H include file to access arguments that are passed by a variable list. For more information about these macros, see [va\\_arg](#), [va\\_copy](#), [va\\_end](#), [va\\_start](#) in the documentation for the C Runtime Library.

The following example shows how the macros work together with the type (declared in STDARG.H):

```
// variable_argument_lists.cpp
#include <stdio.h>
#include <stdarg.h>

// Declaration, but not definition, of ShowVar.
void ShowVar( char *szTypes, ... );
int main() {
    ShowVar( "fcsi", 32.4f, 'a', "Test string", 4 );
}

// ShowVar takes a format string of the form
// "ifcs", where each character specifies the
// type of the argument in that position.
//
// i = int
// f = float
// c = char
// s = string (char *)
//
```

```

// Following the format specification is a variable
// list of arguments. Each argument corresponds to
// a format character in the format string to which
// the szTypes parameter points
void ShowVar( char *szTypes, ... ) {
    va_list vl;
    int i;

    // szTypes is the last argument specified; you must access
    // all others using the variable-argument macros.
    va_start( vl, szTypes );

    // Step through the list.
    for( i = 0; szTypes[i] != '\0'; ++i ) {
        union Printable_t {
            int     i;
            float   f;
            char    c;
            char   *s;
        } Printable;
        switch( szTypes[i] ) { // Type to expect.
            case 'i':
                Printable.i = va_arg( vl, int );
                printf_s( "%i\n", Printable.i );
                break;

            case 'f':
                Printable.f = va_arg( vl, double );
                printf_s( "%f\n", Printable.f );
                break;

            case 'c':
                Printable.c = va_arg( vl, char );
                printf_s( "%c\n", Printable.c );
                break;

            case 's':
                Printable.s = va_arg( vl, char * );
                printf_s( "%s\n", Printable.s );
                break;

            default:
                break;
        }
    }
    va_end( vl );
}
//Output:
// 32.400002
// a
// Test string

```

The previous example illustrates these important concepts:

1. You must establish a list marker as a variable of type `va_list` before any variable arguments are accessed. In the previous example, the marker is called `vl`.
2. The individual arguments are accessed by using the `va_arg` macro. You must tell the `va_arg` macro the type of argument to retrieve so that it can transfer the correct number of bytes from the stack. If you specify an incorrect type or a size different from that supplied by the calling program to `va_arg`, the results are unpredictable.
3. You should explicitly cast the result obtained by using the `va_arg` macro to the type that you want.

You must call the macro to terminate variable-argument processing. `va_end`



# Function Overloading

12/21/2017 • 17 min to read • [Edit Online](#)

C++ allows specification of more than one function of the same name in the same scope. These are called overloaded functions and are described in detail in Overloading. Overloaded functions enable programmers to supply different semantics for a function, depending on the types and number of arguments.

For example, a **print** function that takes a string (or **char**) argument performs very different tasks than one that takes an argument of type **\*\*double\***. Overloading permits uniform naming and prevents programmers from having to invent names such as `print_sz` or `print_d`. The following table shows what parts of a function declaration C++ uses to differentiate between groups of functions with the same name in the same scope.

## Overloading Considerations

FUNCTION DECLARATION ELEMENT	USED FOR OVERLOADING?
Function return type	No
Number of arguments	Yes
Type of arguments	Yes
Presence or absence of ellipsis	Yes
Use of <code>typedef</code> names	No
Unspecified array bounds	No
<b>const</b> or <code>volatile</code> (see below)	Yes

Although functions can be distinguished on the basis of return type, they cannot be overloaded on this basis.

`Const` or `volatile` are only used as a basis for overloading if they are used in a class to apply to the **this** pointer for the class, not the function's return type. In other words, overloading applies only if the **const** or `volatile` keyword follows the function's argument list in the declaration.

## Example

The following example illustrates how overloading can be used.

```
// function_overloading.cpp
// compile with: /EHsc
#include <iostream>
#include <math.h>

// Prototype three print functions.
int print( char *s );           // Print a string.
int print( double dvalue );     // Print a double.
int print( double dvalue, int prec ); // Print a double with a
// given precision.
using namespace std;
int main( int argc, char *argv[] )
{
    const double d = 893094.2987;
    cout << d << endl;
    cout << fixed << setprecision( 3 ) << d << endl;
    cout << scientific << d << endl;
}
```

```

if( argc < 2 )
{
// These calls to print invoke print( char *s ).
print( "This program requires one argument." );
print( "The argument specifies the number of" );
print( "digits precision for the second number" );
print( "printed." );
exit(0);
}

// Invoke print( double dvalue ).
print( d );

// Invoke print( double dvalue, int prec );
print( d, atoi( argv[1] ) );
}

// Print a string.
int print( char *s )
{
cout << s << endl;
return cout.good();
}

// Print a double in default precision.
int print( double dvalue )
{
cout << dvalue << endl;
return cout.good();
}

// Print a double in specified precision.
// Positive numbers for precision indicate how many digits
// precision after the decimal point to show. Negative
// numbers for precision indicate where to round the number
// to the left of the decimal point.
int print( double dvalue, int prec )
{
// Use table-lookup for rounding/truncation.
static const double rgPow10[] = {
10E-7, 10E-6, 10E-5, 10E-4, 10E-3, 10E-2, 10E-1, 10E0,
10E1, 10E2, 10E3, 10E4, 10E5, 10E6
};
const int iPowZero = 6;
// If precision out of range, just print the number.
if( prec < -6 || prec > 7 )
return print( dvalue );
// Scale, truncate, then rescale.
dvalue = floor( dvalue / rgPow10[iPowZero - prec] ) *
rgPow10[iPowZero - prec];
cout << dvalue << endl;
return cout.good();
}

```

The preceding code shows overloading of the `print` function in file scope.

The default argument is not considered part of the function type. Therefore, it is not used in selecting overloaded functions. Two functions that differ only in their default arguments are considered multiple definitions rather than overloaded functions.

Default arguments cannot be supplied for overloaded operators.

## Argument Matching

Overloaded functions are selected for the best match of function declarations in the current scope to the arguments supplied in the function call. If a suitable function is found, that function is called. "Suitable" in this

context means one of the following:

- An exact match was found.
- A trivial conversion was performed.
- An integral promotion was performed.
- A standard conversion to the desired argument type exists.
- A user-defined conversion (either conversion operator or constructor) to the desired argument type exists.
- Arguments represented by an ellipsis were found.

The compiler creates a set of candidate functions for each argument. Candidate functions are functions in which the actual argument in that position can be converted to the type of the formal argument.

A set of "best matching functions" is built for each argument, and the selected function is the intersection of all the sets. If the intersection contains more than one function, the overloading is ambiguous and generates an error. The function that is eventually selected is always a better match than every other function in the group for at least one argument. If this is not the case (if there is no clear winner), the function call generates an error.

Consider the following declarations (the functions are marked **Variant 1**, **Variant 2**, and **Variant 3**, for identification in the following discussion):

```
Fraction &Add( Fraction &f, long l );      // Variant 1
Fraction &Add( long l, Fraction &f );      // Variant 2
Fraction &Add( Fraction &f, Fraction &f ); // Variant 3

Fraction F1, F2;
```

Consider the following statement:

```
F1 = Add( F2, 23 );
```

The preceding statement builds two sets:

<b>SET 1: CANDIDATE FUNCTIONS THAT HAVE FIRST ARGUMENT OF TYPE FRACTION</b>	<b>SET 2: CANDIDATE FUNCTIONS WHOSE SECOND ARGUMENT CAN BE CONVERTED TO TYPE INT</b>
Variant 1	Variant 1 ( <code>int</code> can be converted to <code>long</code> using a standard conversion)
Variant 3	

Functions in Set 2 are functions for which there are implicit conversions from actual parameter type to formal parameter type, and among such functions there is a function for which the "cost" of converting the actual parameter type to its formal parameter type is the smallest.

The intersection of these two sets is Variant 1. An example of an ambiguous function call is:

```
F1 = Add( 3, 6 );
```

The preceding function call builds the following sets:

**SET 1: CANDIDATE FUNCTIONS THAT HAVE FIRST ARGUMENT OF TYPE INT**

Variant 2 ( `int` ) can be converted to `long` using a standard conversion)

**SET 2: CANDIDATE FUNCTIONS THAT HAVE SECOND ARGUMENT OF TYPE INT**

Variant 1 ( `int` ) can be converted to `long` using a standard conversion)

Note that the intersection between these two sets is empty. Therefore, the compiler generates an error message.

For argument matching, a function with  $n$  default arguments is treated as  $n+1$  separate functions, each with a different number of arguments.

The ellipsis (...) acts as a wildcard; it matches any actual argument. This can lead to many ambiguous sets, if you do not design your overloaded function sets with extreme care.

**NOTE**

Ambiguity of overloaded functions cannot be determined until a function call is encountered. At that point, the sets are built for each argument in the function call, and you can determine whether an unambiguous overload exists. This means that ambiguities can remain in your code until they are evoked by a particular function call.

## Argument Type Differences

Overloaded functions differentiate between argument types that take different initializers. Therefore, an argument of a given type and a reference to that type are considered the same for the purposes of overloading. They are considered the same because they take the same initializers. For example, `max( double, double )` is considered the same as `max( double &, double & )`. Declaring two such functions causes an error.

For the same reason, function arguments of a type modified by `const` or `volatile` are not treated differently than the base type for the purposes of overloading.

However, the function overloading mechanism can distinguish between references that are qualified by `const` and `volatile` and references to the base type. This makes code such as the following possible:

```
// argument_type_differences.cpp
// compile with: /EHsc /W3
// C4521 expected
#include <iostream>

using namespace std;
class Over {
public:
    Over() { cout << "Over default constructor\n"; }
    Over( Over &o ) { cout << "Over&\n"; }
    Over( const Over &co ) { cout << "const Over&\n"; }
    Over( volatile Over &vo ) { cout << "volatile Over&\n"; }
};

int main() {
    Over o1;           // Calls default constructor.
    Over o2( o1 );    // Calls Over( Over& ). 
    const Over o3;    // Calls default constructor.
    Over o4( o3 );    // Calls Over( const Over& ). 
    volatile Over o5; // Calls default constructor.
    Over o6( o5 );    // Calls Over( volatile Over& ). 
}
```

## Output

```
Over default constructor
Over&
Over default constructor
const Over&
Over default constructor
volatile Over&
```

Pointers to **const** and **volatile** objects are also considered different from pointers to the base type for the purposes of overloading.

## Argument matching and conversions

When the compiler tries to match actual arguments against the arguments in function declarations, it can supply standard or user-defined conversions to obtain the correct type if no exact match can be found. The application of conversions is subject to these rules:

- Sequences of conversions that contain more than one user-defined conversion are not considered.
- Sequences of conversions that can be shortened by removing intermediate conversions are not considered.

The resultant sequence of conversions, if any, is called the best matching sequence. There are several ways to convert an object of type `int` to type `unsigned long` using standard conversions (described in [Standard Conversions](#)):

- Convert from `int` to `long` and then from `long` to `unsigned long`.
- Convert from `int` to `unsigned long`.

The first sequence, although it achieves the desired goal, is not the best matching sequence — a shorter sequence exists.

The following table shows a group of conversions, called trivial conversions, that have a limited effect on determining which sequence is the best matching. The instances in which trivial conversions affect choice of sequence are discussed in the list following the table.

### Trivial Conversions

CONVERT FROM TYPE	CONVERT TO TYPE
<code>type-name</code>	<code>type-name &amp;</code>
<code>type-name &amp;</code>	<code>type-name</code>
<code>type-name []</code>	<code>*type-name**</code>
<code>type-name ( argument-list )</code>	<code>( *type-name ) ( argument-list )</code>
<code>type-name</code>	<code>const type-name</code>
<code>type-name</code>	<code>volatile type-name</code>
<code>*type-name**</code>	<code>const *type-name**</code>
<code>*type-name**</code>	<code>volatile *type-name**</code>

The sequence in which conversions are attempted is as follows:

1. Exact match. An exact match between the types with which the function is called and the types declared in the function prototype is always the best match. Sequences of trivial conversions are classified as exact matches. However, sequences that do not make any of these conversions are considered better than sequences that convert:

- From pointer, to pointer to **const** (`type` `\*` to **const** `type` `\*`).
- From pointer, to pointer to **volatile** (`type` `\*` to **volatile** `type` `\*`).
- From reference, to reference to **const** (`type` `&` to **const** `type` `&`).
- From reference, to reference to **volatile** (`type` `&` to **volatile** `type` `&`).

2. Match using promotions. Any sequence not classified as an exact match that contains only integral promotions, conversions from **float** to **double**, and trivial conversions is classified as a match using promotions. Although not as good a match as any exact match, a match using promotions is better than a match using standard conversions.

3. Match using standard conversions. Any sequence not classified as an exact match or a match using promotions that contains only standard conversions and trivial conversions is classified as a match using standard conversions. Within this category, the following rules are applied:

- Conversion from a pointer to a derived class, to a pointer to a direct or indirect base class is preferable to converting to **void** `\*` or **const void** `\*`.
- Conversion from a pointer to a derived class, to a pointer to a base class produces a better match the closer the base class is to a direct base class. Suppose the class hierarchy is as shown in the following figure.



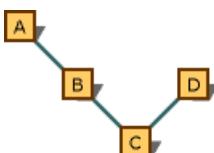
Graph Illustrating Preferred Conversions

Conversion from type `D*` to type `C*` is preferable to conversion from type `D*` to type `B*`. Similarly, conversion from type `D*` to type `B*` is preferable to conversion from type `D*` to type `A*`.

This same rule applies to reference conversions. Conversion from type `D&` to type `C&` is preferable to conversion from type `D&` to type `B&`, and so on.

This same rule applies to pointer-to-member conversions. Conversion from type `T D::*` to type `T C::*` is preferable to conversion from type `T D::*` to type `T B::*`, and so on (where `T` is the type of the member).

The preceding rule applies only along a given path of derivation. Consider the graph shown in the following figure.



Multiple-Inheritance Graph Illustrating Preferred Conversions

Conversion from type `C*` to type `B*` is preferable to conversion from type `C*` to type `A*`. The reason is that they are on the same path, and `B*` is closer. However, conversion from type `C*` to type `D*` is not

preferable to conversion to type `A*`; there is no preference because the conversions follow different paths.

4. Match with user-defined conversions. This sequence cannot be classified as an exact match, a match using promotions, or a match using standard conversions. The sequence must contain only user-defined conversions, standard conversions, or trivial conversions to be classified as a match with user-defined conversions. A match with user-defined conversions is considered a better match than a match with an ellipsis but not as good a match as a match with standard conversions.
5. Match with an ellipsis. Any sequence that matches an ellipsis in the declaration is classified as a match with an ellipsis. This is considered the weakest match.

User-defined conversions are applied if no built-in promotion or conversion exists. These conversions are selected on the basis of the type of the argument being matched. Consider the following code:

```
// argument_matching1.cpp
class UDC
{
public:
    operator int()
    {
        return 0;
    }
    operator long();
};

void Print( int i )
{
};

UDC udc;

int main()
{
    Print( udc );
}
```

The available user-defined conversions for class `UDC` are from type `int` and type `long`. Therefore, the compiler considers conversions for the type of the object being matched: `UDC`. A conversion to `int` exists, and it is selected.

During the process of matching arguments, standard conversions can be applied to both the argument and the result of a user-defined conversion. Therefore, the following code works:

```
void LogToFile( long l );
...
UDC udc;
LogFile( udc );
```

In the preceding example, the user-defined conversion, `operator long`, is invoked to convert `udc` to type `long`. If no user-defined conversion to type `long` had been defined, the conversion would have proceeded as follows: Type `UDC` would have been converted to type `int` using the user-defined conversion. Then the standard conversion from type `int` to type `long` would have been applied to match the argument in the declaration.

If any user-defined conversions are required to match an argument, the standard conversions are not used when evaluating the best match. This is true even if more than one candidate function requires a user-defined conversion; in such a case, the functions are considered equal. For example:

```

// argument_matching2.cpp
// C2668 expected
class UDC1
{
public:
    UDC1( int ); // User-defined conversion from int.
};

class UDC2
{
public:
    UDC2( long ); // User-defined conversion from long.
};

void Func( UDC1 );
void Func( UDC2 );

int main()
{
    Func( 1 );
}

```

Both versions of `Func` require a user-defined conversion to convert type `int` to the class type argument. The possible conversions are:

- Convert from type `int` to type `UDC1` (a user-defined conversion).
- Convert from type `int` to type `long`; then convert to type `UDC2` (a two-step conversion).

Even though the second of these requires a standard conversion, as well as the user-defined conversion, the two conversions are still considered equal.

#### NOTE

User-defined conversions are considered conversion by construction or conversion by initialization (conversion function). Both methods are considered equal when considering the best match.

## Argument matching and the `this` pointer

Class member functions are treated differently, depending on whether they are declared as `static`. Because nonstatic functions have an implicit argument that supplies the `this` pointer, nonstatic functions are considered to have one more argument than static functions; otherwise, they are declared identically.

These nonstatic member functions require that the implied `this` pointer match the object type through which the function is being called, or, for overloaded operators, they require that the first argument match the object on which the operator is being applied. (For more information about overloaded operators, see [Overloaded Operators](#).)

Unlike other arguments in overloaded functions, no temporary objects are introduced and no conversions are attempted when trying to match the `this` pointer argument.

When the `->` member-selection operator is used to access a member function of class `class_name`, the `this` pointer argument has a type of `class_name * const`. If the members are declared as `const` or `volatile`, the types are `const class_name * const` and `volatile class_name * const`, respectively.

The `.` member-selection operator works exactly the same way, except that an implicit `&` (address-of) operator is prefixed to the object name. The following example shows how this works:

```
// Expression encountered in code  
obj.name  
  
// How the compiler treats it  
(&obj)->name
```

The left operand of the `->*` and `.*` (pointer to member) operators are treated the same way as the `.` and `->` (member-selection) operators with respect to argument matching.

## Restrictions

Several restrictions govern an acceptable set of overloaded functions:

- Any two functions in a set of overloaded functions must have different argument lists.
- Overloading functions with argument lists of the same types, based on return type alone, is an error.

### Microsoft Specific

You can overload **operator new** solely on the basis of return type — specifically, on the basis of the memory-model modifier specified.

### END Microsoft Specific

- Member functions cannot be overloaded solely on the basis of one being static and the other nonstatic.
- `typedef` declarations do not define new types; they introduce synonyms for existing types. They do not affect the overloading mechanism. Consider the following code:

```
typedef char * PSTR;  
  
void Print( char *szToPrint );  
void Print( PSTR szToPrint );
```

The preceding two functions have identical argument lists. `PSTR` is a synonym for type `char \*`. In member scope, this code generates an error.

- Enumerated types are distinct types and can be used to distinguish between overloaded functions.
- The types "array of " and "pointer to" are considered identical for the purposes of distinguishing between overloaded functions. This is true only for singly dimensioned arrays. Therefore, the following overloaded functions conflict and generate an error message:

```
void Print( char *szToPrint );  
void Print( char szToPrint[] );
```

For multiply dimensioned arrays, the second and all succeeding dimensions are considered part of the type. Therefore, they are used in distinguishing between overloaded functions:

```
void Print( char szToPrint[] );  
void Print( char szToPrint[][7] );  
void Print( char szToPrint[][9][42] );
```

## Declaration matching

Any two function declarations of the same name in the same scope can refer to the same function, or to two

discrete functions that are overloaded. If the argument lists of the declarations contain arguments of equivalent types (as described in the previous section), the function declarations refer to the same function. Otherwise, they refer to two different functions that are selected using overloading.

Class scope is strictly observed; therefore, a function declared in a base class is not in the same scope as a function declared in a derived class. If a function in a derived class is declared with the same name as a function in the base class, the derived-class function hides the base-class function instead of causing overloading.

Block scope is strictly observed; therefore, a function declared in file scope is not in the same scope as a function declared locally. If a locally declared function has the same name as a function declared in file scope, the locally declared function hides the file-scoped function instead of causing overloading. For example:

```
// declaration_matching1.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
void func( int i )
{
    cout << "Called file-scoped func : " << i << endl;
}

void func( char *sz )
{
    cout << "Called locally declared func : " << sz << endl;
}

int main()
{
    // Declare func local to main.
    extern void func( char *sz );

    func( 3 );    // C2664 Error. func( int ) is hidden.
    func( "s" );
}
```

The preceding code shows two definitions from the function `func`. The definition that takes an argument of type `char *` is local to `main` because of the `extern` statement. Therefore, the definition that takes an argument of type `int` is hidden, and the first call to `func` is in error.

For overloaded member functions, different versions of the function can be given different access privileges. They are still considered to be in the scope of the enclosing class and thus are overloaded functions. Consider the following code, in which the member function `Deposit` is overloaded; one version is public, the other, private.

The intent of this sample is to provide an `Account` class in which a correct password is required to perform deposits. This is accomplished using overloading.

Note that the call to `Deposit` in `Account::Deposit` calls the private member function. This call is correct because `Account::Deposit` is a member function and therefore has access to the private members of the class.

```
// declaration_matching2.cpp
class Account
{
public:
    Account()
    {
    }
    double Deposit( double dAmount, char *szPassword );

private:
    double Deposit( double dAmount )
    {
        return 0.0;
    }
    int Validate( char *szPassword )
    {
        return 0;
    }

};

int main()
{
    // Allocate a new object of type Account.
    Account *pAcct = new Account;

    // Deposit $57.22. Error: calls a private function.
    // pAcct->Deposit( 57.22 );

    // Deposit $57.22 and supply a password. OK: calls a
    // public function.
    pAcct->Deposit( 52.77, "pswd" );
}

double Account::Deposit( double dAmount, char *szPassword )
{
    if ( Validate( szPassword ) )
        return Deposit( dAmount );
    else
        return 0.0;
}
```

## See Also

[Functions \(C++\)](#)

# Explicitly Defaulted and Deleted Functions

12/21/2017 • 6 min to read • [Edit Online](#)

In C++11, defaulted and deleted functions give you explicit control over whether the special member functions are automatically generated. Deleted functions also give you simple language to prevent problematic type promotions from occurring in arguments to functions of all types—special member functions, as well as normal member functions and non-member functions—which would otherwise cause an unwanted function call.

## Benefits of explicitly defaulted and deleted functions

In C++, the compiler automatically generates the default constructor, copy constructor, copy-assignment operator, and destructor for a type if it does not declare its own. These functions are known as the *special member functions*, and they are what make simple user-defined types in C++ behave like structures do in C. That is, you can create, copy, and destroy them without any additional coding effort. C++11 brings move semantics to the language and adds the move constructor and move-assignment operator to the list of special member functions that the compiler can automatically generate.

This is convenient for simple types, but complex types often define one or more of the special member functions themselves, and this can prevent other special member functions from being automatically generated. In practice:

- If any constructor is explicitly declared, then no default constructor is automatically generated.
- If a virtual destructor is explicitly declared, then no default destructor is automatically generated.
- If a move constructor or move-assignment operator is explicitly declared, then:
  - No copy constructor is automatically generated.
  - No copy-assignment operator is automatically generated.
- If a copy constructor, copy-assignment operator, move constructor, move-assignment operator, or destructor is explicitly declared, then:
  - No move constructor is automatically generated.
  - No move-assignment operator is automatically generated.

### NOTE

Additionally, the C++11 standard specifies the following additional rules:

- If a copy constructor or destructor is explicitly declared, then automatic generation of the copy-assignment operator is deprecated.
  - If a copy-assignment operator or destructor is explicitly declared, then automatic generation of the copy constructor is deprecated.

In both cases, Visual Studio continues to automatically generate the necessary functions implicitly, and does not emit a warning.

The consequences of these rules can also leak into object hierarchies. For example, if for any reason a base class fails to have a default constructor that's callable from a deriving class—that is, a `public` or `protected` constructor that takes no parameters—then a class that derives from it cannot automatically generate its own default constructor.

These rules can complicate the implementation of what should be straight-forward, user-defined types and common C++ idioms—for example, making a user-defined type non-copyable by declaring the copy constructor and copy-assignment operator privately and not defining them.

```
struct noncopyable
{
    noncopyable() {};

private:
    noncopyable(const noncopyable&);
    noncopyable& operator=(const noncopyable&);

};
```

Before C++11, this code snippet was the idiomatic form of non-copyable types. However, it has several problems:

- The copy constructor has to be declared privately to hide it, but because it's declared at all, automatic generation of the default constructor is prevented. You have to explicitly define the default constructor if you want one, even if it does nothing.
- Even if the explicitly-defined default constructor does nothing, it's considered non-trivial by the compiler. It's less efficient than an automatically generated default constructor and prevents `noncopyable` from being a true POD type.
- Even though the copy constructor and copy-assignment operator are hidden from outside code, the member functions and friends of `noncopyable` can still see and call them. If they are declared but not defined, calling them causes a linker error.
- Although this is a commonly accepted idiom, the intent is not clear unless you understand all of the rules for automatic generation of the special member functions.

In C++11, the non-copyable idiom can be implemented in a way that is more straightforward.

```
struct noncopyable
{
    noncopyable() =default;
    noncopyable(const noncopyable&) =delete;
    noncopyable& operator=(const noncopyable&) =delete;

};
```

Notice how the problems with the pre-C++11 idiom are resolved:

- Generation of the default constructor is still prevented by declaring the copy constructor, but you can bring it back by explicitly defaulting it.
- Explicitly defaulted special member functions are still considered trivial, so there is no performance penalty, and `noncopyable` is not prevented from being a true POD type.
- The copy constructor and copy-assignment operator are public but deleted. It is a compile-time error to define or call a deleted function.
- The intent is clear to anyone who understands `=default` and `=delete`. You don't have to understand the rules for automatic generation of special member functions.

Similar idioms exist for making user-defined types that are non-movable, that can only be dynamically allocated, or that cannot be dynamically allocated. Each of these idioms have pre-C++11 implementations that suffer similar problems, and that are similarly resolved in C++11 by implementing them in terms of defaulted and deleted special member functions.

## Explicitly defaulted functions

You can default any of the special member functions—to explicitly state that the special member function uses the default implementation, to define the special member function with a non-public access qualifier, or to reinstate a special member function whose automatic generation was prevented by other circumstances.

You default a special member function by declaring it as in this example:

```
struct widget
{
    widget()=default;

    inline widget& operator=(const widget&);

};

inline widget& widget::operator=(const widget&) =default;
```

Notice that you can default a special member function outside the body of a class as long as it's inlinable.

Because of the performance benefits of trivial special member functions, we recommend that you prefer automatically generated special member functions over empty function bodies when you want the default behavior. You can do this either by explicitly defaulting the special member function, or by not declaring it (and also not declaring other special member functions that would prevent it from being automatically generated.)

## Deleted functions

You can delete special member functions as well as normal member functions and non-member functions to prevent them from being defined or called. Deleting of special member functions provides a cleaner way of preventing the compiler from generating special member functions that you don't want. The function must be deleted as it is declared; it cannot be deleted afterwards in the way that a function can be declared and then later defaulted.

```
struct widget
{
    // deleted operator new prevents widget from being dynamically allocated.
    void* operator new(std::size_t) = delete;
};
```

Deleting of normal member function or non-member functions prevents problematic type promotions from causing an unintended function to be called. This works because deleted functions still participate in overload resolution and provide a better match than the function that could be called after the types are promoted. The function call resolves to the more-specific—but deleted—function and causes a compiler error.

```
// deleted overload prevents call through type promotion of float to double from succeeding.
void call_with_true_double_only(float) =delete;
void call_with_true_double_only(double param) { return; }
```

Notice in the preceding sample that calling `call_with_true_double_only` by using a `float` argument would cause a compiler error, but calling `call_with_true_double_only` by using an `int` argument would not; in the `int` case, the argument will be promoted from `int` to `double` and successfully call the `double` version of the function, even though that might not be what's intended. To ensure that any call to this function by using a non-double argument causes a compiler error, you can declare a template version of the function that's deleted.

```
template < typename T >
void call_with_true_double_only(T) =delete; //prevent call through type promotion of any T to double from
succeeding.

void call_with_true_double_only(double param) { return; } // also define for const double, double&, etc. as
needed.
```

# Argument-Dependent Name (Koenig) Lookup on Functions

12/21/2017 • 1 min to read • [Edit Online](#)

The compiler can use argument-dependent name lookup to find the definition of an unqualified function call. Argument-dependent name lookup is also called Koenig lookup. The type of every argument in a function call is defined within a hierarchy of namespaces, classes, structures, unions, or templates. When you specify an unqualified [postfix](#) function call, the compiler searches for the function definition in the hierarchy associated with each argument type.

## Example

In the sample, the compiler notes that function `f()` takes an argument `x`. Argument `x` is of type `A::X`, which is defined in namespace `A`. The compiler searches namespace `A` and finds a definition for function `f()` that takes an argument of type `A::X`.

```
// argument_dependent_name_koenig_lookup_on_functions.cpp
namespace A
{
    struct X
    {
    };
    void f(const X&)
    {
    }
}
int main()
{
    // The compiler finds A::f() in namespace A, which is where
    // the type of argument x is defined. The type of x is A::X.
    A::X x;
    f(x);
}
```

# Default Arguments

12/21/2017 • 2 min to read • [Edit Online](#)

In many cases, functions have arguments that are used so infrequently that a default value would suffice. To address this, the default-argument facility allows for specifying only those arguments to a function that are meaningful in a given call. To illustrate this concept, consider the example presented in [Function Overloading](#).

```
// Prototype three print functions.  
int print( char *s ); // Print a string.  
int print( double dvalue ); // Print a double.  
int print( double dvalue, int prec ); // Print a double with a  
// given precision.
```

In many applications, a reasonable default can be supplied for `prec`, eliminating the need for two functions:

```
// Prototype two print functions.  
int print( char *s ); // Print a string.  
int print( double dvalue, int prec=2 ); // Print a double with a  
// given precision.
```

The implementation of the `print` function is changed slightly to reflect the fact that only one such function exists for type **double**:

```
// default_arguments.cpp  
// compile with: /EHsc /c  
  
// Print a double in specified precision.  
// Positive numbers for precision indicate how many digits  
// precision after the decimal point to show. Negative  
// numbers for precision indicate where to round the number  
// to the left of the decimal point.  
  
#include <iostream>  
#include <math.h>  
using namespace std;  
  
int print( double dvalue, int prec ) {  
    // Use table-lookup for rounding/truncation.  
    static const double rgPow10[] = {  
        10E-7, 10E-6, 10E-5, 10E-4, 10E-3, 10E-2, 10E-1, 10E0,  
        10E1, 10E2, 10E3, 10E4, 10E5, 10E6  
    };  
    const int iPowZero = 6;  
    // If precision out of range, just print the number.  
    if( prec >= -6 && prec <= 7 )  
        // Scale, truncate, then rescale.  
        dvalue = floor( dvalue / rgPow10[iPowZero - prec] ) *  
            rgPow10[iPowZero - prec];  
    cout << dvalue << endl;  
    return cout.good();  
}
```

To invoke the new `print` function, use code such as the following:

```
print( d );      // Precision of 2 supplied by default argument.  
print( d, 0 ); // Override default argument to achieve other  
// results.
```

Note these points when using default arguments:

- Default arguments are used only in function calls where trailing arguments are omitted — they must be the last argument(s). Therefore, the following code is illegal:

```
int print( double dvalue = 0.0, int prec );
```

- A default argument cannot be redefined in later declarations even if the redefinition is identical to the original. Therefore, the following code produces an error:

```
// Prototype for print function.  
int print( double dvalue, int prec = 2 );  
  
...  
  
// Definition for print function.  
int print( double dvalue, int prec = 2 )  
{  
    ...  
}
```

The problem with this code is that the function declaration in the definition redefines the default argument for `prec`.

- Additional default arguments can be added by later declarations.
- Default arguments can be provided for pointers to functions. For example:

```
int (*pShowIntVal)( int i = 0 );
```

## See Also

# Inline Functions (C++)

12/21/2017 • 7 min to read • [Edit Online](#)

A function defined in the body of a class declaration is an inline function.

## Example

In the following class declaration, the `Account` constructor is an inline function. The member functions `GetBalance`, `Deposit`, and `Withdraw` are not specified as **inline** but can be implemented as inline functions.

```
// Inline_Member_Functions.cpp
class Account
{
public:
    Account(double initial_balance) { balance = initial_balance; }
    double GetBalance();
    double Deposit( double Amount );
    double Withdraw( double Amount );
private:
    double balance;
};

inline double Account::GetBalance()
{
    return balance;
}

inline double Account::Deposit( double Amount )
{
    return ( balance += Amount );
}

inline double Account::Withdraw( double Amount )
{
    return ( balance -= Amount );
}
int main()
{}
```

### NOTE

In the class declaration, the functions were declared without the **inline** keyword. The **inline** keyword can be specified in the class declaration; the result is the same.

A given inline member function must be declared the same way in every compilation unit. This constraint causes inline functions to behave as if they were instantiated functions. Additionally, there must be exactly one definition of an inline function.

A class member function defaults to external linkage unless a definition for that function contains the **inline** specifier. The preceding example shows that these functions need not be explicitly declared with the **inline** specifier; using **inline** in the function definition causes it to be an inline function. However, it is illegal to redeclare a function as **inline** after a call to that function.

## Inline, `_inline`, and `_forceinline`

The `inline` and `_inline` specifiers instruct the compiler to insert a copy of the function body into each place the function is called.

The insertion (called inline expansion or inlining) occurs only if the compiler's cost/benefit analysis show it to be profitable. Inline expansion alleviates the function-call overhead at the potential cost of larger code size.

The `_forceinline` keyword overrides the cost/benefit analysis and relies on the judgment of the programmer instead. Exercise caution when using `_forceinline`. Indiscriminate use of `_forceinline` can result in larger code with only marginal performance gains or, in some cases, even performance losses (due to increased paging of a larger executable, for example).

Using inline functions can make your program faster because they eliminate the overhead associated with function calls. Functions expanded inline are subject to code optimizations not available to normal functions.

The compiler treats the inline expansion options and keywords as suggestions. There is no guarantee that functions will be inlined. You cannot force the compiler to inline a particular function, even with the `_forceinline` keyword. When compiling with `/clr`, the compiler will not inline a function if there are security attributes applied to the function.

The **inline** keyword is available only in C++. The `_inline` and `_forceinline` keywords are available in both C and C++. For compatibility with previous versions, **inline** is a synonym for `_inline`.

The **inline** keyword tells the compiler that inline expansion is preferred. However, the compiler can create a separate instance of the function (instantiate) and create standard calling linkages instead of inserting the code inline. Two cases where this can happen are:

- Recursive functions.
- Functions that are referred to through a pointer elsewhere in the translation unit.

These reasons may interfere with inlining, *as may others*, at the discretion of the compiler; you should not depend on the **inline** specifier to cause a function to be inlined.

As with normal functions, there is no defined order of evaluation of the arguments to an inline function. In fact, it could be different from the order in which the arguments are evaluated when passed using normal function call protocol.

The `/Ob` compiler optimization option helps to determine whether inline function expansion actually occurs.

`/LTCG` performs cross-module inlining regardless of whether it was requested in source code.

### Example 1

```
// inline_keyword1.cpp
// compile with: /c
inline int max( int a , int b ) {
    if( a > b )
        return a;
    return b;
}
```

A class's member functions can be declared inline either by using the **inline** keyword or by placing the function definition within the class definition.

### Example 2

```

// inline_keyword2.cpp
// compile with: /EHsc /c
#include <iostream>
using namespace std;

class MyClass {
public:
    void print() { cout << i << ' '; } // Implicitly inline
private:
    int i;
};

```

## Microsoft Specific

The `__inline` keyword is equivalent to **inline**.

Even with `__forceinline`, the compiler cannot inline code in all circumstances. The compiler cannot inline a function if:

- The function or its caller is compiled with `/Ob0` (the default option for debug builds).
- The function and the caller use different types of exception handling (C++ exception handling in one, structured exception handling in the other).
- The function has a variable argument list.
- The function uses inline assembly, unless compiled with `/Og`, `/Ox`, `/O1`, or `/O2`.
- The function is recursive and not accompanied by **#pragma inline\_recursion(on)**. With the pragma, recursive functions are inlined to a default depth of 16 calls. To reduce the inlining depth, use `inline_depth` pragma.
- The function is virtual and is called virtually. Direct calls to virtual functions can be inlined.
- The program takes the address of the function and the call is made via the pointer to the function. Direct calls to functions that have had their address taken can be inlined.
- The function is also marked with the `naked __declspec` modifier.

If the compiler cannot inline a function declared with `__forceinline`, it generates a level 1 warning, except when:

- The function is compiled by using `/Od` or `/Ob0`. No inlining is expected in these cases.
- The function is defined externally, in an included library or another translation unit, or is a virtual call target or indirect call target. The compiler can't identify non-inlined code that it can't find in the current translation unit.

Recursive functions can be substituted inline to a depth specified by the `inline_depth` pragma, up to a maximum of 16 calls. After that depth, recursive function calls are treated as calls to an instance of the function. The depth to which recursive functions are examined by the inline heuristic cannot exceed 16. The `inline_recursion` pragma controls the inline expansion of a function currently under expansion. See the [Inline-Function Expansion \(/Ob\)](#) compiler option for related information.

## END Microsoft Specific

For more information on using the **inline** specifier, see:

- [Inline Class Member Functions](#)
- [Defining Inline C++ Functions with `dllexport` and `dllimport`](#)

## When to use inline functions

Inline functions are best used for small functions such as accessing private data members. The main purpose of these one- or two-line "accessor" functions is to return state information about objects; short functions are sensitive to the overhead of function calls. Longer functions spend proportionately less time in the calling/returning sequence and benefit less from inlining.

A `Point` class can be defined as follows:

```
// when_to_use_inline_functions.cpp
class Point
{
public:
    // Define "accessor" functions as
    // reference types.
    unsigned& x();
    unsigned& y();
private:
    unsigned _x;
    unsigned _y;
};

inline unsigned& Point::x()
{
    return _x;
}
inline unsigned& Point::y()
{
    return _y;
}
int main()
{}
```

Assuming coordinate manipulation is a relatively common operation in a client of such a class, specifying the two accessor functions (`x` and `y` in the preceding example) as **inline** typically saves the overhead on:

- Function calls (including parameter passing and placing the object's address on the stack)
- Preservation of caller's stack frame
- New stack-frame setup
- Return-value communication
- Old stack-frame restore
- Return

## Inline functions vs. macros

Although inline functions are similar to macros (because the function code is expanded at the point of the call at compile time), inline functions are parsed by the compiler, whereas macros are expanded by the preprocessor. As a result, there are several important differences:

- Inline functions follow all the protocols of type safety enforced on normal functions.
- Inline functions are specified using the same syntax as any other function except that they include the **inline** keyword in the function declaration.
- Expressions passed as arguments to inline functions are evaluated once. In some cases, expressions passed as arguments to macros can be evaluated more than once.

The following example shows a macro that converts lowercase letters to uppercase:

```
// inline_functions_macro.c
#include <stdio.h>
#include <conio.h>

#define toupper(a) ((a) >= 'a' && ((a) <= 'z') ? ((a)-('a'-'A')):(a))

int main() {
    char ch;
    printf_s("Enter a character: ");
    ch = toupper( getc(stdin) );
    printf_s( "%c", ch );
}

// Sample Input: xyz
// Sample Output: Z
```

The intent of the expression `toupper(getc(stdin))` is that a character should be read from the console device (`stdin`) and, if necessary, converted to uppercase.

Because of the implementation of the macro, `getc` is executed once to determine whether the character is greater than or equal to "a," and once to determine whether it is less than or equal to "z." If it is in that range, `getc` is executed again to convert the character to uppercase. This means the program waits for two or three characters when, ideally, it should wait for only one.

Inline functions remedy the problem previously described:

```
// inline_functions_inline.cpp
#include <stdio.h>
#include <conio.h>

inline char toupper( char a ) {
    return ((a >= 'a' && a <= 'z') ? a-('a'-'A') : a );
}

int main() {
    printf_s("Enter a character: ");
    char ch = toupper( getc(stdin) );
    printf_s( "%c", ch );
}
```

```
Sample Input: a
Sample Output: A
```

## See Also

[noinline](#)

[auto\\_inline](#)

# Operator Overloading

12/21/2017 • 2 min to read • [Edit Online](#)

The `operator` keyword declares a function specifying what `operator-symbol` means when applied to instances of a class. This gives the operator more than one meaning, or "overloads" it. The compiler distinguishes between the different meanings of an operator by examining the types of its operands.

## Syntax

```
type operator operator-symbol ( parameter-list )
```

## Remarks

You can redefine the function of most built-in operators globally or on a class-by-class basis. Overloaded operators are implemented as functions.

The name of an overloaded operator is `operator x`, where `x` is the operator as it appears in the following table. For example, to overload the addition operator, you define a function called `operator+`. Similarly, to overload the addition/assignment operator, `+=`, define a function called `operator+=`.

### Redefinable Operators

OPERATOR	NAME	TYPE
,	Comma	Binary
!	Logical NOT	Unary
!=	Inequality	Binary
%	Modulus	Binary
%=	Modulus assignment	Binary
&	Bitwise AND	Binary
&	Address-of	Unary
&&	Logical AND	Binary
&=	Bitwise AND assignment	Binary
( )	Function call	—
( )	Cast Operator	Unary
*	Multiplication	Binary

OPERATOR	NAME	TYPE
*	Pointer dereference	Unary
*=	Multiplication assignment	Binary
+	Addition	Binary
+	Unary Plus	Unary
++	Increment <sup>1</sup>	Unary
+=	Addition assignment	Binary
-	Subtraction	Binary
-	Unary negation	Unary
--	Decrement <sup>1</sup>	Unary
-=	Subtraction assignment	Binary
->	Member selection	Binary
->*	Pointer-to-member selection	Binary
/	Division	Binary
/=	Division assignment	Binary
<	Less than	Binary
<<	Left shift	Binary
<<=	Left shift assignment	Binary
<=	Less than or equal to	Binary
=	Assignment	Binary
==	Equality	Binary
>	Greater than	Binary
>=	Greater than or equal to	Binary
>>	Right shift	Binary
>>=	Right shift assignment	Binary

OPERATOR	NAME	TYPE
[ ]	Array subscript	—
^	Exclusive OR	Binary
^=	Exclusive OR assignment	Binary
&#124;	Bitwise inclusive OR	Binary
&#124;=	Bitwise inclusive OR assignment	Binary
&#124;&#124;	Logical OR	Binary
~	One's complement	Unary
delete	Delete	—
new	New	—
conversion operators	conversion operators	Unary

1 Two versions of the unary increment and decrement operators exist: preincrement and postincrement.

See [General Rules for Operator Overloading](#) for more information. The constraints on the various categories of overloaded operators are described in the following topics:

- [Unary Operators](#)
- [Binary Operators](#)
- [Assignment](#)
- [Function Call](#)
- [Subscripting](#)
- [Class-Member Access](#)
- [Increment and Decrement](#).
- [User-Defined Type Conversions](#)

The operators shown in the following table cannot be overloaded. The table includes the preprocessor symbols # and ##.

## Nonredefinable Operators

Operator	Name
.	Member selection
.*	Pointer-to-member selection

<code>::</code>	Scope resolution
<code>? :</code>	Conditional
<code>#</code>	Preprocessor convert to string
<code>##</code>	Preprocessor concatenate

Although overloaded operators are usually called implicitly by the compiler when they are encountered in code, they can be invoked explicitly the same way as any member or nonmember function is called:

```
Point pt;
pt.operator+( 3 ); // Call addition operator to add 3 to pt.
```

## Example

The following example overloads the `+` operator to add two complex numbers and returns the result.

```
// operator_overloading.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

struct Complex {
    Complex( double r, double i ) : re(r), im(i) {}
    Complex operator+( Complex &other );
    void Display( ) { cout << re << ", " << im << endl; }
private:
    double re, im;
};

// Operator overloaded using a member function
Complex Complex::operator+( Complex &other ) {
    return Complex( re + other.re, im + other.im );
}

int main() {
    Complex a = Complex( 1.2, 3.4 );
    Complex b = Complex( 5.6, 7.8 );
    Complex c = Complex( 0.0, 0.0 );

    c = a + b;
    c.Display();
}
```

## Output

```
6.8, 11.2
```

## In this section

1. [General Rules for Operator Overloading](#)
2. [Overloading Unary Operators](#)

[3. Binary Operators](#)

[4. Assignment](#)

[5. Function Call](#)

[6. Subscripting](#)

[7. Member Access](#)

## See Also

[C++ Built-in Operators, Precedence and Associativity](#)

[Keywords](#)

# General Rules for Operator Overloading

12/21/2017 • 2 min to read • [Edit Online](#)

The following rules constrain how overloaded operators are implemented. However, they do not apply to the [new](#) and [delete](#) operators, which are covered separately.

- You cannot define new operators, such as `**`.
- You cannot redefine the meaning of operators when applied to built-in data types.
- Overloaded operators must either be a nonstatic class member function or a global function. A global function that needs access to private or protected class members must be declared as a friend of that class. A global function must take at least one argument that is of class or enumerated type or that is a reference to a class or enumerated type. For example:

```
// rules_for_operator_overloading.cpp
class Point
{
public:
    Point operator<( Point & ); // Declare a member operator
                                // overload.
    // Declare addition operators.
    friend Point operator+( Point&, int );
    friend Point operator+( int, Point& );
};

int main()
{
```

The preceding code sample declares the less-than operator as a member function; however, the addition operators are declared as global functions that have friend access. Note that more than one implementation can be provided for a given operator. In the case of the preceding addition operator, the two implementations are provided to facilitate commutativity. It is just as likely that operators that add a `Point` to a `Point`, `int` to a `Point`, and so on, might be implemented.

- Operators obey the precedence, grouping, and number of operands dictated by their typical use with built-in types. Therefore, there is no way to express the concept "add 2 and 3 to an object of type `Point`," expecting 2 to be added to the `x` coordinate and 3 to be added to the `y` coordinate.
- Unary operators declared as member functions take no arguments; if declared as global functions, they take one argument.
- Binary operators declared as member functions take one argument; if declared as global functions, they take two arguments.
- If an operator can be used as either a unary or a binary operator (`&`, `\`, `**+`, and `**-*`), you can overload each use separately.
- Overloaded operators cannot have default arguments.
- All overloaded operators except assignment (`operator=`) are inherited by derived classes.
- The first argument for member-function overloaded operators is always of the class type of the object for which the operator is invoked (the class in which the operator is declared, or a class derived from that class).

No conversions are supplied for the first argument.

Note that the meaning of any of the operators can be changed completely. That includes the meaning of the address-of (`&`), assignment (`=`), and function-call operators. Also, identities that can be relied upon for built-in types can be changed using operator overloading. For example, the following four statements are usually equivalent when completely evaluated:

```
var = var + 1;  
var += 1;  
var++;  
++var;
```

This identity cannot be relied upon for class types that overload operators. Moreover, some of the requirements implicit in the use of these operators for basic types are relaxed for overloaded operators. For example, the addition/assignment operator, `+=`, requires the left operand to be an l-value when applied to basic types; there is no such requirement when the operator is overloaded.

**NOTE**

For consistency, it is often best to follow the model of the built-in types when defining overloaded operators. If the semantics of an overloaded operator differ significantly from its meaning in other contexts, it can be more confusing than useful.

## See Also

[Operator Overloading](#)

# Overloading Unary Operators

12/21/2017 • 1 min to read • [Edit Online](#)

The unary operators that can be overloaded are the following:

1. `!` (logical NOT)
2. `&` (address-of)
3. `~` (one's complement)
4. `*` (pointer dereference)
5. `+` (unary plus)
6. `-` (unary negation)
7. `++` (increment)
8. `--` (decrement)
9. conversion operators

The postfix increment and decrement operators (`++` and `--`) are treated separately in [Increment and Decrement](#).

Conversion operators are also discussed in a separate topic; see [User-Defined Type Conversions](#).

The following rules are true of all other unary operators. To declare a unary operator function as a nonstatic member, you must declare it in the form:

```
ret-type operator op ()
```

where `ret-type` is the return type and `op` is one of the operators listed in the preceding table.

To declare a unary operator function as a global function, you must declare it in the form:

```
ret-type operator op ( arg )
```

where `ret-type` and `op` are as described for member operator functions and the `arg` is an argument of class type on which to operate.

## NOTE

There is no restriction on the return types of the unary operators. For example, it makes sense for logical NOT (`!`) to return an integral value, but this is not enforced.

## See Also

[Operator Overloading](#)

# Increment and Decrement Operator Overloading (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

The increment and decrement operators fall into a special category because there are two variants of each:

- Preincrement and postincrement
- Prederecrement and postdecrement

When you write overloaded operator functions, it can be useful to implement separate versions for the prefix and postfix versions of these operators. To distinguish between the two, the following rule is observed: The prefix form of the operator is declared exactly the same way as any other unary operator; the postfix form accepts an additional argument of type `int`.

## NOTE

When specifying an overloaded operator for the postfix form of the increment or decrement operator, the additional argument must be of type `int`; specifying any other type generates an error.

The following example shows how to define prefix and postfix increment and decrement operators for the `Point` class:

```

// increment_and_decrement1.cpp
class Point
{
public:
    // Declare prefix and postfix increment operators.
    Point& operator++();           // Prefix increment operator.
    Point operator++(int);         // Postfix increment operator.

    // Declare prefix and postfix decrement operators.
    Point& operator--();           // Prefix decrement operator.
    Point operator--(int);         // Postfix decrement operator.

    // Define default constructor.
    Point() { _x = _y = 0; }

    // Define accessor functions.
    int x() { return _x; }
    int y() { return _y; }
private:
    int _x, _y;
};

// Define prefix increment operator.
Point& Point::operator++()
{
    _x++;
    _y++;
    return *this;
}

// Define postfix increment operator.
Point Point::operator++(int)
{
    Point temp = *this;
    ++*this;
    return temp;
}

// Define prefix decrement operator.
Point& Point::operator--()
{
    _x--;
    _y--;
    return *this;
}

// Define postfix decrement operator.
Point Point::operator--(int)
{
    Point temp = *this;
    --*this;
    return temp;
}
int main()
{
}

```

The same operators can be defined in file scope (globally) using the following function heads:

```

friend Point& operator++( Point& )      // Prefix increment
friend Point& operator++( Point&, int ) // Postfix increment
friend Point& operator--( Point& )      // Prefix decrement
friend Point& operator--( Point&, int ) // Postfix decrement

```

The argument of type `int` that denotes the postfix form of the increment or decrement operator is not commonly used to pass arguments. It usually contains the value 0. However, it can be used as follows:

```
// increment_and_decrement2.cpp
class Int
{
public:
    Int &operator++( int n );
private:
    int _i;
};

Int& Int::operator++( int n )
{
    if( n != 0 )      // Handle case where an argument is passed.
        _i += n;
    else
        _i++;         // Handle case where no argument is passed.
    return *this;
}
int main()
{
    Int i;
    i.operator++( 25 ); // Increment by 25.
}
```

There is no syntax for using the increment or decrement operators to pass these values other than explicit invocation, as shown in the preceding code. A more straightforward way to implement this functionality is to overload the addition/assignment operator (`+=`).

## See Also

[Operator Overloading](#)

# Binary Operators

12/21/2017 • 1 min to read • [Edit Online](#)

The following table shows a list of operators that can be overloaded.

## Redefinable Binary Operators

OPERATOR	NAME
,	Comma
!=	Inequality
%	Modulus
%=	Modulus/assignment
&	Bitwise AND
&&	Logical AND
&=	Bitwise AND/assignment
\*	Multiplication
*=	Multiplication/assignment
+	Addition
+=	Addition/assignment
-	Subtraction
-=	Subtraction/assignment
->	Member selection
->\*	Pointer-to-member selection
/	Division
/=	Division/assignment
<	Less than
<<	Left shift
<<=	Left shift/assignment

OPERATOR	NAME
<code>&lt;=</code>	Less than or equal to
<code>=</code>	Assignment
<code>==</code>	Equality
<code>&gt;</code>	Greater than
<code>&gt;=</code>	Greater than or equal to
<code>&gt;&gt;</code>	Right shift
<code>&gt;&gt;=</code>	Right shift/assignment
<code>^</code>	Exclusive OR
<code>^=</code>	Exclusive OR/assignment
<code> </code>	Bitwise inclusive OR
<code>&amp;#124;=</code>	Bitwise inclusive OR/assignment
<code>&amp;#124;&amp;#124;</code>	Logical OR

To declare a binary operator function as a nonstatic member, you must declare it in the form:

*ret-type* **operator** `op ( arg )`

where *ret-type* is the return type, `op` is one of the operators listed in the preceding table, and `arg` is an argument of any type.

To declare a binary operator function as a global function, you must declare it in the form:

*ret-type* **operator** `op ( arg1 , arg2 )`

where *ret-type* and `op` are as described for member operator functions and `arg1` and `arg2` are arguments. At least one of the arguments must be of class type.

#### NOTE

There is no restriction on the return types of the binary operators; however, most user-defined binary operators return either a class type or a reference to a class type.

## See Also

[Operator Overloading](#)

# Assignment

12/21/2017 • 1 min to read • [Edit Online](#)

The assignment operator (`=`) is, strictly speaking, a binary operator. Its declaration is identical to any other binary operator, with the following exceptions:

- It must be a nonstatic member function. No `operator=` can be declared as a nonmember function.
- It is not inherited by derived classes.
- A default `operator=` function can be generated by the compiler for class types if none exists. (For more information about default `operator=` functions, see [Memberwise Assignment and Initialization](#).)

The following example illustrates how to declare an assignment operator:

```
// assignment.cpp
class Point
{
public:
    Point &operator=( Point & ); // Right side is the argument.
    int _x, _y;
};

// Define assignment operator.
Point &Point::operator=( Point &ptRHS )
{
    _x = ptRHS._x;
    _y = ptRHS._y;

    return *this; // Assignment operator returns left side.
}

int main()
{}
```

Note that the supplied argument is the right side of the expression. The operator returns the object to preserve the behavior of the assignment operator, which returns the value of the left side after the assignment is complete. This allows writing statements such as:

```
pt1 = pt2 = pt3;
```

## See Also

[Operator Overloading](#)

# Function Call (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The function-call operator, invoked using parentheses, is a binary operator.

## Syntax

```
primary-expression ( expression-list )
```

## Remarks

In this context, `primary-expression` is the first operand, and `expression-list`, a possibly empty list of arguments, is the second operand. The function-call operator is used for operations that require a number of parameters. This works because `expression-list` is a list instead of a single operand. The function-call operator must be a nonstatic member function.

The function-call operator, when overloaded, does not modify how functions are called; rather, it modifies how the operator is to be interpreted when applied to objects of a given class type. For example, the following code would usually be meaningless:

```
Point pt;
pt( 3, 2 );
```

Given an appropriate overloaded function-call operator, however, this syntax can be used to offset the `x` coordinate 3 units and the `y` coordinate 2 units. The following code shows such a definition:

```
// function_call.cpp
class Point
{
public:
    Point() { _x = _y = 0; }
    Point &operator()( int dx, int dy )
        { _x += dx; _y += dy; return *this; }
private:
    int _x, _y;
};

int main()
{
    Point pt;
    pt( 3, 2 );
}
```

Note that the function-call operator is applied to the name of an object, not the name of a function.

You can also overload the function call operator using a pointer to a function (rather than the function itself).

```
typedef void(*ptf)();  
void func()  
{  
}  
struct S  
{  
    operator ptf()  
    {  
        return func;  
    }  
};  
  
int main()  
{  
    S s;  
    s(); //operates as s.operator ptf()()  
}
```

## See Also

[Operator Overloading](#)

# Subscripting

12/21/2017 • 1 min to read • [Edit Online](#)

The subscript operator (`[ ]`), like the function-call operator, is considered a binary operator. The subscript operator must be a nonstatic member function that takes a single argument. This argument can be of any type and designates the desired array subscript.

## Example

The following example demonstrates how to create a vector of type `int` that implements bounds checking:

```
// subscripting.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class IntVector {
public:
    IntVector( int cElements );
    ~IntVector() { delete [] _iElements; }
    int& operator[](int nSubscript);
private:
    int *_iElements;
    int _iUpperBound;
};

// Construct an IntVector.
IntVector::IntVector( int cElements ) {
    _iElements = new int[cElements];
    _iUpperBound = cElements;
}

// Subscript operator for IntVector.
int& IntVector::operator[](int nSubscript) {
    static int iErr = -1;

    if( nSubscript >= 0 && nSubscript < _iUpperBound )
        return _iElements[nSubscript];
    else {
        clog << "Array bounds violation." << endl;
        return iErr;
    }
}

// Test the IntVector class.
int main() {
    IntVector v( 10 );
    int i;

    for( i = 0; i <= 10; ++i )
        v[i] = i;

    v[3] = v[9];

    for ( i = 0; i <= 10; ++i )
        cout << "Element: [" << i << "] = " << v[i] << endl;
}
```

```
Array bounds violation.  
Element: [0] = 0  
Element: [1] = 1  
Element: [2] = 2  
Element: [3] = 9  
Element: [4] = 4  
Element: [5] = 5  
Element: [6] = 6  
Element: [7] = 7  
Element: [8] = 8  
Element: [9] = 9  
Array bounds violation.  
Element: [10] = 10
```

## Comments

When `i` reaches 10 in the preceding program, `operator[]` detects that an out-of-bounds subscript is being used and issues an error message.

Note that the function `operator[]` returns a reference type. This causes it to be an l-value, allowing you to use subscripted expressions on either side of assignment operators.

## See Also

[Operator Overloading](#)

# Member Access

12/21/2017 • 1 min to read • [Edit Online](#)

Class member access can be controlled by overloading the member access operator (->). This operator is considered a unary operator in this usage, and the overloaded operator function must be a class member function. Therefore, the declaration for such a function is:

## Syntax

```
class-type *operator->()
```

## Remarks

where *class-type* is the name of the class to which this operator belongs. The member access operator function must be a nonstatic member function.

This operator is used (often in conjunction with the pointer-dereference operator) to implement "smart pointers" that validate pointers prior to dereference or count usage.

The . member access operator cannot be overloaded.

## See Also

[Operator Overloading](#)

# Classes and Structs (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This section introduces C++ classes and structs. The two constructs are identical in C++ except that in structs the default accessibility is public, whereas in classes the default is private.

Classes and structs are the constructs whereby you define your own types. Classes and structs can both contain data members and member functions, which enable you to describe the type's state and behavior.

The following topics are included:

- [class](#)
- [struct](#)
- [Class Member Overview](#)
- [Member Access Control](#)
- [Inheritance](#)
- [Static Members](#)
- [User-Defined Type Conversions](#)
- [Mutable Data Members \(mutable specifier\)](#)
- [Nested Class Declarations](#)
- [Anonymous Class Types](#)
- [Pointers to Members](#)
- [this Pointer](#)
- [C++ Bit Fields](#)

The three class types are structure, class, and union. They are declared using the [struct](#), [class](#), and [union](#) keywords (see [Defining Class Types](#)). The following table shows the differences among the three class types.

For more information on unions, see [Unions](#). For information on managed classes and structs, see [Classes and Structs](#).

## Access Control and Constraints of Structures, Classes and Unions

STRUCTURES	CLASSES	UNIONS
class key is <code>struct</code>	class key is <code>class</code>	class key is <code>union</code>
Default access is public	Default access is private	Default access is public
No usage constraints	No usage constraints	Use only one member at a time

## See Also



# class (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The `class` keyword declares a class type or defines an object of a class type.

## Syntax

```
[template-spec]
    class [ms-decl-spec] [tag [: base-list ]]
{
    member-list
} [declarators];
[ class ] tag declarators;
```

### Parameters

`template-spec`

Optional template specifications. For more information, refer to [Templates](#).

`class`

The `class` keyword.

`ms-decl-spec`

Optional storage-class specification. For more information, refer to the [\\_declspec](#) keyword.

`tag`

The type name given to the class. The tag becomes a reserved word within the scope of the class. The tag is optional. If omitted, an anonymous class is defined. For more information, see [Anonymous Class Types](#).

`base-list`

Optional list of classes or structures this class will derive its members from. See [Base Classes](#) for more information. Each base class or structure name can be preceded by an access specifier ([public](#), [private](#), [protected](#)) and the [virtual](#) keyword. See the member-access table in [Controlling Access to Class Members](#) for more information.

`member-list`

List of class members. Refer to [Class Member Overview](#) for more information.

`declarators`

Declarator list specifying the names of one or more instances of the class type. Declarators may include initializer lists if all data members of the class are `public`. This is more common in structures, whose data members are `public` by default, than in classes. See [Overview of Declarators](#) for more information.

## Remarks

For more information on classes in general, refer to one of the following topics:

- [struct](#)
- [union](#)
- [\\_\\_multiple\\_inheritance](#)

- [\\_\\_single\\_inheritance](#)
- [\\_\\_virtual\\_inheritance](#)

For information on managed classes and structs, see [Classes and Structs](#)

## Example

```
// class.cpp
// compile with: /EHsc
// Example of the class keyword
// Exhibits polymorphism/virtual functions.

#include <iostream>
#include <string>
#define TRUE = 1
using namespace std;

class dog
{
public:
    dog()
    {
        _legs = 4;
        _bark = true;
    }

    void setDogSize(string dogSize)
    {
        _dogSize = dogSize;
    }
    virtual void setEars(string type)      // virtual function
    {
        _earType = type;
    }

private:
    string _dogSize, _earType;
    int _legs;
    bool _bark;

};

class breed : public dog
{
public:
    breed( string color, string size)
    {
        _color = color;
        setDogSize(size);
    }

    string getColor()
    {
        return _color;
    }

    // virtual function redefined
    void setEars(string length, string type)
    {
        _earLength = length;
        _earType = type;
    }

protected:
    string _color, _earLength, _earType;
}:
```

```
int main()
{
    dog mongrel;
    breed labrador("yellow", "large");
    mongrel.setEars("pointy");
    labrador.setEars("long", "floppy");
    cout << "Cody is a " << labrador.getColor() << " labrador" << endl;
}
```

## See Also

[Keywords](#)

[Classes and Structs](#)

# struct (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

The `struct` keyword defines a structure type and/or a variable of a structure type.

## Syntax

```
[template-spec] struct[ms-decl-spec] [tag [: base-list ]]
{
    member-list
} [declarators];
[struct] tag declarators;
```

### Parameters

`template-spec`

Optional template specifications. For more information, refer to [Template Specifications](#).

`struct`

The `struct` keyword.

`ms-decl-spec`

Optional storage-class specification. For more information, refer to the [\\_declspec](#) keyword.

`tag`

The type name given to the structure. The tag becomes a reserved word within the scope of the structure. The tag is optional. If omitted, an anonymous structure is defined. For more information, see [Anonymous Class Types](#).

`base-list`

Optional list of classes or structures this structure will derive its members from. See [Base Classes](#) for more information. Each base class or structure name can be preceded by an access specifier ([public](#), [private](#), [protected](#)) and the [virtual](#) keyword. See the member-access table in [Controlling Access to Class Members](#) for more information.

`member-list`

List of structure members. Refer to [Class Member Overview](#) for more information. The only difference here is that `struct` is used in place of `class`.

`declarators`

Declarator list specifying the names of the class. Declarator lists declare one or more instances of the structure type. Declarators may include initializer lists if all data members of the class are `public`. Initializer lists are common in structures because data members are `public` by default. See [Overview of Declarators](#) for more information.

## Remarks

A structure type is a user-defined composite type. It is composed of fields or members that can have different types.

In C++, a structure is the same as a class except that its members are `public` by default.

For information on managed classes and structs, see [Classes and Structs](#).

# Using a Structure

In C, you must explicitly use the `struct` keyword to declare a structure. In C++, you do not need to use the `struct` keyword after the type has been defined.

You have the option of declaring variables when the structure type is defined by placing one or more comma-separated variable names between the closing brace and the semicolon.

Structure variables can be initialized. The initialization for each variable must be enclosed in braces.

For related information, see [class](#), [union](#), and [enum](#).

## Example

```
#include <iostream>
using namespace std;

struct PERSON { // Declare PERSON struct type
    int age; // Declare member types
    long ss;
    float weight;
    char name[25];
} family_member; // Define object of type PERSON

struct CELL { // Declare CELL bit field
    unsigned short character : 8; // 00000000 ????????
    unsigned short foreground : 3; // 00000??? 00000000
    unsigned short intensity : 1; // 0000?000 00000000
    unsigned short background : 3; // 0??00000 00000000
    unsigned short blink : 1; // ?0000000 00000000
} screen[25][80]; // Array of bit fields

int main() {
    struct PERSON sister; // C style structure declaration
    PERSON brother; // C++ style structure declaration
    sister.age = 13; // assign values to members
    brother.age = 7;
    cout << "sister.age = " << sister.age << '\n';
    cout << "brother.age = " << brother.age << '\n';

    CELL my_cell;
    my_cell.character = 1;
    cout << "my_cell.character = " << my_cell.character;
}
// Output:
// sister.age = 13
// brother.age = 7
// my_cell.character = 1
```

# Class Member Overview

12/21/2017 • 3 min to read • [Edit Online](#)

A class or struct consists of its members. The work that a class does is performed by its member functions. The state that it maintains is stored in its data members. Initialization of members is done by constructors, and cleanup work such as freeing of memory and releasing of resources is done by destructors. In C++11 and later, data members can (and usually should) be initialized at the point of declaration.

## Kinds of class members

The full list of member categories is as follows:

- [Special Member Functions](#).
- [Overview of Member Functions](#).
- [Data members](#) including built-in types and other user defined types.
- Operators
- [Nested Class Declarations](#) and.)
- [Unions](#)
- [Enumerations](#).
- [Bit fields](#).
- [Friends](#).
- [Aliases and typedefs](#).

### NOTE

Friends are included in the preceding list because they are contained in the class declaration. However, they are not true class members, because they are not in the scope of the class.

## Example class declaration

The following example shows a simple class declaration:

```

// TestRun.h

class TestRun
{
    // Start member list.

    //The class interface accessible to all callers.
public:
    // Use compiler-generated default constructor:
    TestRun() = default;
    // Don't generate a copy constructor:
    TestRun(const TestRun&) = delete;
    TestRun(std::string name);
    void DoSomething();
    int Calculate(int a, double d);
    virtual ~TestRun();
    enum class State { Active, Suspended };

    // Accessible to this class and derived classes only.
protected:
    virtual void Initialize();
    virtual void Suspend();
    State GetState();

    // Accessible to this class only.
private:
    // Default brace-initialization of instance members:
    State _state{ State::Suspended };
    std::string _testName{ "" };
    int _index{ 0 };

    // Non-const static member:
    static int _instances;
    // End member list.
};

// Define and initialize static member.
int TestRun::_instances{ 0 };

```

## Member accessibility

The members of a class are declared in the member list. The member list of a class may be divided into any number of **private**, **protected** and **public** sections using keywords known as access specifiers. A colon : must follow the access specifier. These sections need not be contiguous, that is, any of these keywords may appear several times in the member list. The keyword designates the access of all members up until the next access specifier or the closing brace. For more information, see [Member Access Control \(C++\)](#).

## Static members

A data member may be declared as static, which means all objects of the class have access to the same copy of it. A member function may be declared as static, in which case it can only access static data members of the class (and has no *this* pointer). For more information, see [Static Data Members](#).

## Special member functions

Special member functions are functions that are automatically provided by the compiler if you do not specify them in your source code.

1. Default constructor
2. Copy constructor

3. **(C++11)** Move constructor
4. Copy assignment operator
5. **(C++11)** Move assignment operator
6. Destructor

For more information, see [Special Member Functions](#).

## Memberwise initialization

In C++11 and later, non-static member declarators can contain initializers.

```
class CanInit
{
public:
    long num {7};          // OK in C++11
    int k = 9;             // OK in C++11
    static int i = 9; // Error: must be defined and initialized
                      // outside of class declaration.

    // initializes num to 7 and k to 9
    CanInit(){}

    // overwrites original initialized value of num:
    CanInit(int val) : num(val) {}

};

int main()
{}
```

If a member is assigned a value in a constructor, that value overwrites the value with which the member was initialized at the point of declaration.

There is only one shared copy of static data members for all objects of a given class type. Static data members must be defined and can be initialized at file scope. (For more information about static data members, see [Static Data Members](#).) The following example shows how to perform these initializations:

```
// class_members2.cpp
class CanInit2
{
public:
    CanInit2() {} // Initializes num to 7 when new objects of type
                  // CanInit are created.
    long      num {7};
    static int i;
    static int j;
};

// At file scope:

// i is defined at file scope and initialized to 15.
// The initializer is evaluated in the scope of CanInit.
int CanInit2::i = 15;

// The right side of the initializer is in the scope
// of the object being initialized
int CanInit2::j = i;
```

**NOTE**

The class name, `CanInit2`, must precede `i` to specify that the `i` being defined is a member of class `CanInit2`.

## See Also

[Classes and Structs](#)

# Member Access Control (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

Access controls enable you to separate the **public** interface of a class from the **private** implementation details and the **protected** members that are only for use by derived classes. The access specifier applies to all members declared after it until the next access specifier is encountered.

```
class Point
{
public:
    Point( int, int ) // Declare public constructor. ;
    Point(); // Declare public default constructor.
    int &x( int ); // Declare public accessor.
    int &y( int ); // Declare public accessor.

private:           // Declare private state variables.
    int _x;
    int _y;

protected:        // Declare protected function for derived classes only.
    Point ToWindowCoords();
};
```

The default access is **private** in a class, and **public** in a struct or union. Access specifiers in a class can be used any number of times in any order. The allocation of storage for objects of class types is implementation dependent, but members are guaranteed to be assigned successively higher memory addresses between access specifiers.

## Member-Access Control

TYPE OF ACCESS	MEANING
<b>private</b>	Class members declared as <b>private</b> can be used only by member functions and friends (classes or functions) of the class.
<b>protected</b>	Class members declared as <b>protected</b> can be used by member functions and friends (classes or functions) of the class. Additionally, they can be used by classes derived from the class.
<b>public</b>	Class members declared as <b>public</b> can be used by any function.

Access control helps prevent you from using objects in ways they were not intended to be used. This protection is lost when explicit type conversions (casts) are performed.

### NOTE

Access control is equally applicable to all names: member functions, member data, nested classes, and enumerators.

## Access Control in Derived Classes

Two factors control which members of a base class are accessible in a derived class; these same factors control access to the inherited members in the derived class:

- Whether the derived class declares the base class using the **public** access specifier in the *class-head* (*class-head* is described in the Grammar section in [Defining Class Types](#)).
- What the access to the member is in the base class.

The following table shows the interaction between these factors and how to determine base-class member access.

### Member Access in Base Class

PRIVATE	PROTECTED	PUBLIC
Always inaccessible regardless of derivation access	Private in derived class if you use private derivation	Private in derived class if you use private derivation
	Protected in derived class if you use protected derivation	Protected in derived class if you use protected derivation
	Protected in derived class if you use public derivation	Public in derived class if you use public derivation

The following example illustrates this:

```
// access_specifiers_for_base_classes.cpp
class BaseClass
{
public:
    int PublicFunc(); // Declare a public member.
protected:
    int ProtectedFunc(); // Declare a protected member.
private:
    int PrivateFunc(); // Declare a private member.
};

// Declare two classes derived from BaseClass.
class DerivedClass1 : public BaseClass
{
};

class DerivedClass2 : private BaseClass
{
};

int main()
{
}
```

In `DerivedClass1`, the member function `PublicFunc` is a public member and `ProtectedFunc` is a protected member because `BaseClass` is a public base class. `PrivateFunc` is private to `BaseClass`, and it is inaccessible to any derived classes.

In `DerivedClass2`, the functions `PublicFunc` and `ProtectedFunc` are considered private members because `BaseClass` is a private base class. Again, `PrivateFunc` is private to `BaseClass`, and it is inaccessible to any derived classes.

You can declare a derived class without a base-class access specifier. In such a case, the derivation is considered private if the derived class declaration uses the **class** keyword. The derivation is considered public if the derived

class declaration uses the `struct` keyword. For example, the following code:

```
class Derived : Base  
...
```

is equivalent to:

```
class Derived : private Base  
...
```

Similarly, the following code:

```
struct Derived : Base  
...
```

is equivalent to:

```
struct Derived : public Base  
...
```

Note that members declared as having private access are not accessible to functions or derived classes unless those functions or classes are declared using the `friend` declaration in the base class.

A **union** type cannot have a base class.

**NOTE**

When specifying a private base class, it is advisable to explicitly use the `private` keyword so users of the derived class understand the member access.

## Access control and static members

When you specify a base class as `private`, it affects only nonstatic members. Public static members are still accessible in the derived classes. However, accessing members of the base class using pointers, references, or objects can require a conversion, at which time access control is again applied. Consider the following example:

```

// access_control.cpp
class Base
{
public:
    int Print();           // Nonstatic member.
    static int CountOf(); // Static member.
};

// Derived1 declares Base as a private base class.
class Derived1 : private Base
{
};

// Derived2 declares Derived1 as a public base class.
class Derived2 : public Derived1
{
    int ShowCount();      // Nonstatic member.
};

// Define ShowCount function for Derived2.
int Derived2::ShowCount()
{
    // Call static member function CountOf explicitly.
    int cCount = Base::CountOf(); // OK.

    // Call static member function CountOf using pointer.
    cCount = this->CountOf(); // C2247. Conversion of
                             // Derived2 * to Base * not
                             // permitted.

    return cCount;
}

```

In the preceding code, access control prohibits conversion from a pointer to `Derived2` to a pointer to `Base`. The `this` pointer is implicitly of type `Derived2 *`. To select the `Countof` function, `this` must be converted to type `Base *`. Such a conversion is not permitted because `Base` is a private indirect base class to `Derived2`. Conversion to a private base class type is acceptable only for pointers to immediate derived classes. Therefore, pointers of type `Derived1 *` can be converted to type `Base *`.

Note that calling the `Countof` function explicitly, without using a pointer, reference, or object to select it, implies no conversion. Therefore, the call is allowed.

Members and friends of a derived class, `T`, can convert a pointer to `T` to a pointer to a private direct base class of `T`.

## Access to virtual functions

The access control applied to [virtual](#) functions is determined by the type used to make the function call. Overriding declarations of the function do not affect the access control for a given type. For example:

```

// access_to_virtual_functions.cpp
class VFuncBase
{
public:
    virtual int GetState() { return _state; }
protected:
    int _state;
};

class VFuncDerived : public VFuncBase
{
private:
    int GetState() { return _state; }
};

int main()
{
    VFuncDerived vfd;           // Object of derived type.
    VFuncBase *pvfb = &vfd;    // Pointer to base type.
    VFuncDerived *pvfd = &vfd; // Pointer to derived type.

    int State;

    State = pvfb->GetState(); // GetState is public.
    State = pvfd->GetState(); // C2248 error expected; GetState is private;
}

```

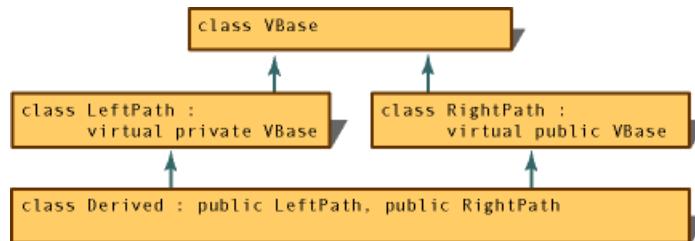
In the preceding example, calling the virtual function `GetState` using a pointer to type `VFuncBase` calls `VFuncDerived::GetState`, and `GetState` is treated as public. However, calling `GetState` using a pointer to type `VFuncDerived` is an access-control violation because `GetState` is declared private in class `VFuncDerived`.

**Caution**

The virtual function `GetState` can be called using a pointer to the base class `VFuncBase`. This does not mean that the function called is the base-class version of that function.

## Access control with multiple inheritance

In multiple-inheritance lattices involving virtual base classes, a given name can be reached through more than one path. Because different access control can be applied along these different paths, the compiler chooses the path that gives the most access. See the following figure.



Access Along Paths of an Inheritance Graph

In the figure, a name declared in class `VBase` is always reached through class `RightPath`. The right path is more accessible because `RightPath` declares `VBase` as a public base class, whereas `LeftPath` declares `VBase` as private.

## See Also

[C++ Language Reference](#)

# friend (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

In some circumstances, it is more convenient to grant member-level access to functions that are not members of a class or to all members in a separate class. Only the class implementer can declare who its friends are. A function or class cannot declare itself as a friend of any class. In a class definition, use the `friend` keyword and the name of a non-member function or other class to grant it access to the private and protected members of your class. In a template definition, a type parameter can be declared as a friend.

## Syntax

```
class friend F  
friend F;
```

## Friend declarations

If you declare a friend function that was not previously declared, that function is exported to the enclosing nonclass scope.

Functions declared in a friend declaration are treated as if they had been declared using the `extern` keyword. (For more information about `extern`, see [Static Storage-Class Specifiers](#).)

Although functions with global scope can be declared as friends prior to their prototypes, member functions cannot be declared as friends before the appearance of their complete class declaration. The following code shows why this fails:

```
class ForwardDeclared; // Class name is known.  
class HasFriends  
{  
    friend int ForwardDeclared::IsAFriend(); // C2039 error expected  
};
```

The preceding example enters the class name `ForwardDeclared` into scope, but the complete declaration — specifically, the portion that declares the function `IsAFriend` — is not known. Therefore, the `friend` declaration in class `HasFriends` generates an error.

Starting in C++11, there are two forms of friend declarations for a class:

```
friend class F;  
friend F;
```

The first form introduces a new class `F` if no existing class by that name was found in the innermost namespace.

**C++11:** The second form does not introduce a new class; it can be used when the class has already been declared, and it must be used when declaring a template type parameter or a `typedef` as a friend.

Use `class friend F` when the referenced type has not yet been declared:

```
namespace NS
{
    class M
    {
        class friend F; // Introduces F but doesn't define it
    };
}
```

```
namespace NS
{
    class M
    {
        friend F; // error C2433: 'NS::F': 'friend' not permitted on data declarations
    };
}
```

In the following example, `friend F` refers to the `F` class that is declared outside the scope of NS.

```
class F {};
namespace NS
{
    class M
    {
        friend F; // OK
    };
}
```

Use `friend F` to declare a template parameter as a friend:

```
template <typename T>
class my_class
{
    friend T;
    //...
};
```

Use `friend F` to declare a typedef as friend:

```
class Foo {};
typedef Foo F;

class G
{
    friend F; // OK
    friend class F // Error C2371 -- redefinition
};
```

To declare two classes that are friends of one another, the entire second class must be specified as a friend of the first class. The reason for this restriction is that the compiler has enough information to declare individual friend functions only at the point where the second class is declared.

#### NOTE

Although the entire second class must be a friend to the first class, you can select which functions in the first class will be friends of the second class.

## friend functions

A `friend` function is a function that is not a member of a class but has access to the class's private and protected members. Friend functions are not considered class members; they are normal external functions that are given special access privileges. Friends are not in the class's scope, and they are not called using the member-selection operators (`.` and `->`) unless they are members of another class. A `friend` function is declared by the class that is granting access. The `friend` declaration can be placed anywhere in the class declaration. It is not affected by the access control keywords.

The following example shows a `Point` class and a friend function, `ChangePrivate`. The `friend` function has access to the private data member of the `Point` object it receives as a parameter.

```
// friend_functions.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class Point
{
    friend void ChangePrivate( Point & );
public:
    Point( void ) : m_i(0) {}
    void PrintPrivate( void ){cout << m_i << endl; }

private:
    int m_i;
};

void ChangePrivate ( Point &i ) { i.m_i++; }

int main()
{
    Point sPoint;
    sPoint.PrintPrivate();
    ChangePrivate(sPoint);
    sPoint.PrintPrivate();
    // Output: 0
    1
}
```

## Class members as friends

Class member functions can be declared as friends in other classes. Consider the following example:

```

// classes_as_friends1.cpp
// compile with: /c
class B;

class A {
public:
    int Func1( B& b );

private:
    int Func2( B& b );
};

class B {
private:
    int _b;

    // A::Func1 is a friend function to class B
    // so A::Func1 has access to all members of B
    friend int A::Func1( B& );
};

int A::Func1( B& b ) { return b._b; }    // OK
int A::Func2( B& b ) { return b._b; }    // C2248

```

In the preceding example, only the function `A::Func1( B& )` is granted friend access to class `B`. Therefore, access to the private member `_b` is correct in `Func1` of class `A` but not in `Func2`.

A `friend` class is a class all of whose member functions are friend functions of a class, that is, whose member functions have access to the other class's private and protected members. Suppose the `friend` declaration in class `B` had been:

```
friend class A;
```

In that case, all member functions in class `A` would have been granted friend access to class `B`. The following code is an example of a friend class:

```

// classes_as_friends2.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class YourClass {
friend class YourOtherClass; // Declare a friend class
public:
    YourClass() : topSecret(0){}
    void printMember() { cout << topSecret << endl; }
private:
    int topSecret;
};

class YourOtherClass {
public:
    void change( YourClass& yc, int x ){yc.topSecret = x;}
};

int main() {
    YourClass yc1;
    YourOtherClass yoc1;
    yc1.printMember();
    yoc1.change( yc1, 5 );
    yc1.printMember();
}

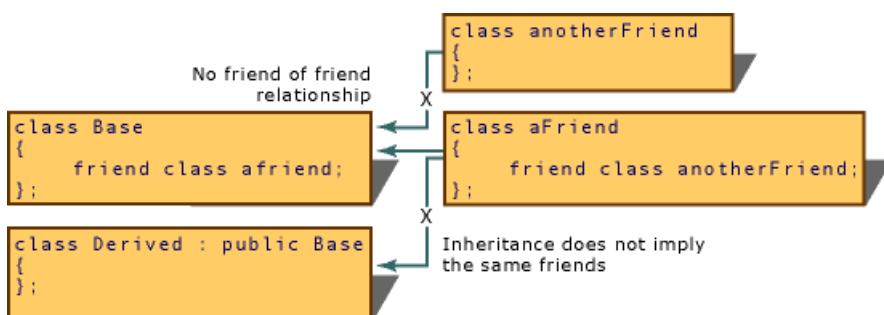
```

Friendship is not mutual unless explicitly specified as such. In the above example, member functions of `YourClass` cannot access the private members of `YourOtherClass`.

A managed type cannot have any friend functions, friend classes, or friend interfaces.

Friendship is not inherited, meaning that classes derived from `YourOtherClass` cannot access `YourClass`'s private members. Friendship is not transitive, so classes that are friends of `YourOtherClass` cannot access `YourClass`'s private members.

The following figure shows four class declarations: `Base`, `Derived`, `aFriend`, and `anotherFriend`. Only class `aFriend` has direct access to the private members of `Base` (and to any members `Base` might have inherited).



Implications of friend Relationship

## Inline friend definitions

Friend functions can be defined inside class declarations. These functions are inline functions, and like member inline functions they behave as though they were defined immediately after all class members have been seen but before the class scope is closed (the end of the class declaration).

Friend functions defined inside class declarations are not considered in the scope of the enclosing class; they are in file scope.

## See Also

## Keywords

# private (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
private:  
    [member-list]  
private base-class
```

## Remarks

When preceding a list of class members, the `private` keyword specifies that those members are accessible only from member functions and friends of the class. This applies to all members declared up to the next access specifier or the end of the class.

When preceding the name of a base class, the `private` keyword specifies that the public and protected members of the base class are private members of the derived class.

Default access of members in a class is private. Default access of members in a structure or union is public.

Default access of a base class is private for classes and public for structures. Unions cannot have base classes.

For related information, see [friend](#), [public](#), [protected](#), and the member-access table in [Controlling Access to Class Members](#).

## /clr Specific

In CLR types, the C++ access specifier keywords (`public`, `private`, and `protected`) can affect the visibility of types and methods with regard to assemblies. For more information, see [Member Access Control](#).

### NOTE

Files compiled with `/LN` are not affected by this behavior. In this case, all managed classes (either public or private) will be visible.

## END /clr Specific

## Example

```
// keyword_private.cpp
class BaseClass {
public:
    // privMem accessible from member function
    int pubFunc() { return privMem; }
private:
    void privMem;
};

class DerivedClass : public BaseClass {
public:
    void usePrivate( int i )
        { privMem = i; } // C2248: privMem not accessible
                          // from derived class
};

class DerivedClass2 : private BaseClass {
public:
    // pubFunc() accessible from derived class
    int usePublic() { return pubFunc(); }
};

int main() {
    BaseClass aBase;
    DerivedClass aDerived;
    DerivedClass2 aDerived2;
    aBase.privMem = 1; // C2248: privMem not accessible
    aDerived.privMem = 1; // C2248: privMem not accessible
                         // in derived class
    aDerived2.pubFunc(); // C2247: pubFunc() is private in
                        // derived class
}
```

## See Also

[Controlling Access to Class Members](#)

[Keywords](#)

# protected (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

## Syntax

```
protected:  
    [member-list]  
protected base-class
```

## Remarks

The `protected` keyword specifies access to class members in the *member-list* up to the next access specifier (`public` or `private`) or the end of the class definition. Class members declared as `protected` can be used only by the following:

- Member functions of the class that originally declared these members.
- Friends of the class that originally declared these members.
- Classes derived with public or protected access from the class that originally declared these members.
- Direct privately derived classes that also have private access to protected members.

When preceding the name of a base class, the `protected` keyword specifies that the public and protected members of the base class are protected members of its derived classes.

Protected members are not as private as `private` members, which are accessible only to members of the class in which they are declared, but they are not as public as `public` members, which are accessible in any function.

Protected members that are also declared as `static` are accessible to any friend or member function of a derived class. Protected members that are not declared as `static` are accessible to friends and member functions in a derived class only through a pointer to, reference to, or object of the derived class.

For related information, see [friend](#), [public](#), [private](#), and the member-access table in [Controlling Access to Class Members](#).

## /clr Specific

In CLR types, the C++ access specifier keywords (`public`, `private`, and `protected`) can affect the visibility of types and methods with regard to assemblies. For more information, see [Member Access Control](#).

### NOTE

Files compiled with `/LN` are not affected by this behavior. In this case, all managed classes (either public or private) will be visible.

## END /clr Specific

## Example

```
// keyword_protected.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class X {
public:
    void setProtMemb( int i ) { m_protMemb = i; }
    void Display() { cout << m_protMemb << endl; }
protected:
    int m_protMemb;
    void Protfunc() { cout << "\nAccess allowed\n"; }
} x;

class Y : public X {
public:
    void useProtfunc() { Protfunc(); }
} y;

int main() {
    // x.m_protMemb;      error, m_protMemb is protected
    x.setProtMemb( 0 ); // OK, uses public access function
    x.Display();
    y.setProtMemb( 5 ); // OK, uses public access function
    y.Display();
    // x.Protfunc();      error, Protfunc() is protected
    y.useProtfunc();    // OK, uses public access function
                        // in derived class
}
```

## See Also

[Controlling Access to Class Members](#)

[Keywords](#)

# public (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Syntax

```
public:  
    [member-list]  
public base-class
```

## Remarks

When preceding a list of class members, the **public** keyword specifies that those members are accessible from any function. This applies to all members declared up to the next access specifier or the end of the class.

When preceding the name of a base class, the **public** keyword specifies that the public and protected members of the base class are public and protected members, respectively, of the derived class.

Default access of members in a class is private. Default access of members in a structure or union is public.

Default access of a base class is private for classes and public for structures. Unions cannot have base classes.

For more information, see [private](#), [protected](#), [friend](#), and the member-access table in [Controlling Access to Class Members](#).

## /clr Specific

In CLR types, the C++ access specifier keywords (**public**, [private](#), and [protected](#)) can affect the visibility of types and methods with regard to assemblies. For more information, see [Member Access Control](#).

### NOTE

Files compiled with [/LN](#) are not affected by this behavior. In this case, all managed classes (either public or private) will be visible.

## END /clr Specific

## Example

```
// keyword_public.cpp
class BaseClass {
public:
    int pubFunc() { return 0; }
};

class DerivedClass : public BaseClass {};

int main() {
    BaseClass aBase;
    DerivedClass aDerived;
    aBase.pubFunc();      // pubFunc() is accessible
                         //   from any function
    aDerived.pubFunc();   // pubFunc() is still public in
                         //   derived class
}
```

## See Also

[Controlling Access to Class Members](#)

[Keywords](#)

# Initializing classes and structs without constructors (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

It is not always necessary to define a constructor for a class, especially ones that are relatively simple. Users can initialize objects of a class or struct by using uniform initialization, as shown in the following example:

```
#include "stdafx.h"
#include <Windows.h>

// No constructor
struct TempData
{
    int StationId;
    time_t time;
    double current;
    double maxTemp;
    double minTemp;
};

// Has a constructor
struct TempData2
{
    TempData2(double minimum, double maximum, double cur, int id, time_t t) :
        minTemp(minimum), maxTemp(maximum), current(cur), stationId(id), time(t) {}
    int stationId;
    time_t time;
    double current;
    double maxTemp;
    double minTemp;
};

int main()
{
    // Member initialization (in order of declaration):
    TempData td{ 45978, GetCurrentTime(), 28.9, 37.0, 16.7 };

    // Default initialization = {0,0,0,0,0}
    TempData td_default{};

    //Error C4700 uninitialized local variable
    TempData td_noInit;

    // Member declaration (in order of ctor parameters)
    TempData2 td2{ 16.7, 37.0, 28.9, 45978, GetCurrentTime() };

    return 0;
}
```

Note that when a class or struct has no constructor, you provide the list elements in the order that the members are declared in the class. If the class has a constructor, provide the elements in the order of the parameters.

## See Also

[Classes and Structs](#)  
[Constructors](#)

# Constructors (C++)

12/21/2017 • 16 min to read • [Edit Online](#)

A constructor is a kind of member function that initializes an instance of its class. A constructor has the same name as the class and no return value. A constructor can have any number of parameters and a class may have any number of overloaded constructors. Constructors may have any accessibility, public, protected or private. If you don't define any constructors, the compiler will generate a default constructor that takes no parameters; you can override this behavior by declaring a default constructor as deleted.

## In this topic

- [Order of Construction](#)
- [Member Lists](#)
- [Explicit Constructors](#)
- [Default Constructors](#)
- [Copy and Move Constructors](#)
- [Explicitly Defaulted and Deleted Constructors](#)
- [Constructors in Derived Classes](#)
- [Virtual Functions in Constructors](#)
- [Constructors and Composite Classes](#)
- [Delegating Constructors](#)
- [Inheriting constructors \(C++11\)](#)
- [Rules for Declaring Constructors](#)
- [Explicitly invoking constructors](#)

## Order of Construction

A constructor performs its work in this order:

1. It calls base class and member constructors in the order of declaration.
2. If the class is derived from virtual base classes, it initializes the object's virtual base pointers.
3. If the class has or inherits virtual functions, it initializes the object's virtual function pointers. Virtual function pointers point to the class's virtual function table to enable correct binding of virtual function calls to code.
4. It executes any code in its function body.

The following example shows the order in which base class and member constructors are called in the constructor for a derived class. First, the base constructor is called, then the base-class members are initialized in the order in which they appear in the class declaration, and then the derived constructor is called.

```

#include <iostream>
using namespace std;

class Contained1 {
public:
    Contained1() {
        cout << "Contained1 constructor." << endl;
    }
};

class Contained2 {
public:
    Contained2() {
        cout << "Contained2 constructor." << endl;
    }
};

class Contained3 {
public:
    Contained3() {
        cout << "Contained3 constructor." << endl;
    }
};

class BaseContainer {
public:
    BaseContainer() {
        cout << "BaseContainer constructor." << endl;
    }
private:
    Contained1 c1;
    Contained2 c2;
};

class DerivedContainer : public BaseContainer {
public:
    DerivedContainer() : BaseContainer() {
        cout << "DerivedContainer constructor." << endl;
    }
private:
    Contained3 c3;
};

int main() {
    DerivedContainer dc;
    int x = 3;
}

```

Here's the output:

```

Contained1 constructor.
Contained2 constructor.
BaseContainer constructor.
Contained3 constructor.
DerivedContainer constructor.

```

If a constructor throws an exception, the order of destruction is the reverse of the order of construction:

1. The code in the body of the constructor function is unwound.
2. Base class and member objects are destroyed, in the reverse order of declaration.
3. If the constructor is non-delegating, all fully-constructed base class objects and members are destroyed. However, because the object itself is not fully constructed, the destructor is not run.

## Member Lists

Initialize class members from constructor arguments by using a member initializer list. This method uses *direct initialization*, which is more efficient than using assignment operators inside the constructor body.

```
class Box {  
public:  
    Box(int width, int length, int height)  
        : m_width(width), m_length(length), m_height(height) // member init list  
    {}  
    int Volume() {return m_width * m_length * m_height; }  
private:  
    int m_width;  
    int m_length;  
    int m_height;  
};
```

Create a Box object:

```
Box b(42, 21, 12);  
cout << "The volume is " << b.Volume();
```

## Explicit Constructors

If a class has a constructor with a single parameter, or if all parameters except one have a default value, the parameter type can be implicitly converted to the class type. For example, if the `Box` class has a constructor like this:

```
Box(int size): m_width(size), m_length(size), m_height(size){}
```

It is possible to initialize a Box like this:

```
Box b = 42;
```

Or pass an int to a function that takes a Box:

```
class ShippingOrder  
{  
public:  
    ShippingOrder(Box b, double postage) : m_box(b), m_postage(postage){}  
  
private:  
    Box m_box;  
    double m_postage;  
}  
//elsewhere...  
ShippingOrder so(42, 10.8);
```

Such conversions can be useful in some cases, but more often they can lead to subtle but serious errors in your code. As a general rule, you should use the `explicit` keyword on a constructor (and user-defined operators) to prevent this kind of implicit type conversion:

```
explicit Box(int size): m_width(size), m_length(size), m_height(size){}
```

When the constructor is explicit, this line causes a compiler error: `ShippingOrder so(42, 10.8);`. For more information, see [User-Defined Type Conversions](#).

## Default Constructors

*Default constructors* have no parameters; they follow slightly different rules:

Default constructors are one of the *special member functions*; if no constructors are declared in a class, the compiler provides a default constructor:

```
class Box {  
    Box(int width, int length, int height)  
        : m_width(width), m_length(length), m_height(height){}  
};  
  
int main(){  
  
    Box box1{}; // call compiler-generated default ctor  
    Box box2;   // call compiler-generated default ctor  
}
```

When you call a default constructor and try to use parentheses, a warning is issued:

```
class myclass{};  
int main(){  
    myclass mc();      // warning C4930: prototyped function not called (was a variable definition intended?)  
}
```

This is an example of the Most Vexing Parse problem. Because the example expression can be interpreted either as the declaration of a function or as the invocation of a default constructor, and because C++ parsers favor declarations over other things, the expression is treated as a function declaration. For more information, see [Most Vexing Parse](#).

If any non-default constructors are declared, the compiler does not provide a default constructor:

```
class Box {  
    Box(int width, int length, int height)  
        : m_width(width), m_length(length), m_height(height){}  
};  
private:  
    int m_width;  
    int m_length;  
    int m_height;  
};  
  
int main(){  
  
    Box box1(1, 2, 3);  
    Box box2{ 2, 3, 4 };  
    Box box4;      // compiler error C2512: no appropriate default constructor available  
}
```

If a class has no default constructor, an array of objects of that class cannot be constructed by using square-bracket syntax alone. For example, given the previous code block, an array of Boxes cannot be declared like this:

```
Box boxes[3]; // compiler error C2512: no appropriate default constructor available
```

However, you can use a set of initializer lists to initialize an array of Boxes:

```
Box boxes[3]{ { 1, 2, 3 }, { 4, 5, 6 }, { 7, 8, 9 } };
```

## Copy and Move Constructors

A *copy constructor* is a special member function that takes as input a reference to an object of the same type, and makes a copy of it. For more information, see [Copy Constructors and Copy Assignment Operators \(C++\)](#). A *move constructor* is also a special member function that moves ownership of an existing object's data to a new variable without copying the original data. For more information, see [Move Constructors and Move Assignment Operators \(C++\)](#).

## Explicitly Defaulted and Deleted Constructors

You can explicitly *default* copy constructors, default constructors, move constructors, copy assignment operators, move assignment operators, and destructors. You can explicitly *delete* all of the special member functions. For more information, see [Explicitly Defaulted and Deleted Functions](#).

## Constructors in Derived Classes

A derived class constructor always calls a base class constructor, so that it can rely on completely constructed base classes before any extra work is done. The base class constructors are called in order of derivation—for example, if ClassA is derived from ClassB, which is derived from ClassC, the ClassC constructor is called first, then the ClassB constructor, then the ClassA constructor.

If a base class does not have a default constructor, you must supply the base class constructor parameters in the derived class constructor:

```

class Box {
public:
    Box(int width, int length, int height){
        m_width = width;
        m_length = length;
        m_height = height;
    }

private:
    int m_width;
    int m_length;
    int m_height;
};

class StorageBox : public Box {
public:
    StorageBox(int width, int length, int height, const string label& ) : Box(width, length, height){
        m_label = label;
    }

private:
    string m_label;
};

int main(){

    const string aLabel = "aLabel";
    StorageBox sb(1, 2, 3, aLabel);
}

```

## Constructors for Classes That Have Multiple Inheritance

If a class is derived from multiple base classes, the base class constructors are invoked in the order in which they are listed in the declaration of the derived class:

```

#include <iostream>
using namespace std;

class BaseClass1 {
public:
    BaseClass1() {
        cout << "BaseClass1 constructor." << endl;
    }
};

class BaseClass2 {
public:
    BaseClass2() {
        cout << "BaseClass2 constructor." << endl;
    }
};

class BaseClass3{
public:
    BaseClass3() {
        cout << "BaseClass3 constructor." << endl;
    }
};

class DerivedClass : public BaseClass1, public BaseClass2, public BaseClass3 {
public:
    DerivedClass() {
        cout << "DerivedClass constructor." << endl;
    }
};

int main() {
    DerivedClass dc;
}

```

You should expect the following output:

```

BaseClass1 constructor.
BaseClass2 constructor.
BaseClass3 constructor.
DerivedClass constructor.

```

## Virtual Functions in Constructors

We recommend that you be careful when you call virtual functions in constructors. Because the base class constructor is always invoked before the derived class constructor, the function that's called in the base constructor is the base class version, not the derived class version. In the following example, constructing a `DerivedClass` causes the `BaseClass` implementation of `print_it()` to execute before the `DerivedClass` constructor causes the `DerivedClass` implementation of `print_it()` to execute:

```

#include <iostream>
using namespace std;

class BaseClass{
public:
    BaseClass(){
        print_it();
    }
    virtual void print_it() {
        cout << "BaseClass print_it" << endl;
    }
};

class DerivedClass : public BaseClass {
public:
    DerivedClass() {
        print_it();
    }
    virtual void print_it(){
        cout << "Derived Class print_it" << endl;
    }
};

int main() {
    DerivedClass dc;
}

```

Here's the output:

```

BaseClass print_it
Derived Class print_it

```

## Constructors and Composite Classes

Classes that contain class-type members are known as *composite classes*. When a class-type member of a composite class is created, the constructor is called before the class's own constructor. When a contained class lacks a default constructor, you must use an initialization list in the constructor of the composite class. In the earlier `StorageBox` example, if you change the type of the `m_label` member variable to a new `Label` class, you must call both the base class constructor and initialize the `m_label` variable in the `StorageBox` constructor:

```

class Label {
public:
    Label(const string& name, const string& address) { m_name = name; m_address = address; }
    string m_name;
    string m_address;
};

class StorageBox : public Box {
public:
    StorageBox(int width, int length, int height, Label label)
        : Box(width, length, height), m_label(label){}
private:
    Label m_label;
};

int main(){
// passing a named Label
Label label1{ "some_name", "some_address" };
StorageBox sb1(1, 2, 3, label1);

// passing a temporary label
StorageBox sb2(3, 4, 5, Label{ "another name", "another address" });

// passing a temporary label as an initializer list
StorageBox sb3(1, 2, 3, {"myname", "myaddress"});
}

```

## Delegating Constructors

A *delegating constructor* calls a different constructor in the same class to do some of the work of initialization. In the following example, the derived class has three constructors—the second constructor delegates to the first, and the third constructor delegates to the second:

```

#include <iostream>
using namespace std;

class ConstructorDestructor {
public:
    ConstructorDestructor() {
        cout << "ConstructorDestructor default constructor." << endl;
    }
    ConstructorDestructor(int int1) {
        cout << "ConstructorDestructor constructor with 1 int." << endl;
    }
    ConstructorDestructor(int int1, int int2) : ConstructorDestructor(int1) {
        cout << "ConstructorDestructor constructor with 2 ints." << endl;

        throw exception();
    }
    ConstructorDestructor(int int1, int int2, int int3) : ConstructorDestructor(int1, int2) {
        cout << "ConstructorDestructor constructor with 3 ints." << endl;
    }
    ~ConstructorDestructor() {
        cout << "ConstructorDestructor destructor." << endl;
    }
};

int main() {
    ConstructorDestructor dc(1, 2, 3);
}

```

Here's the output:

```
ConstructorDestructor constructor with 1 int.  
ConstructorDestructor constructor with 2 ints.  
ConstructorDestructor constructor with 3 ints.
```

The object created by the constructors is fully initialized as soon as any constructor is finished.

`DerivedContainer(int int1)` succeeds, but `DerivedContainer(int int1, int int2)` fails and the destructor is called.

```
class ConstructorDestructor {  
public:  
    ConstructorDestructor() {  
        cout << "ConstructorDestructor default constructor." << endl;  
    }  
    ConstructorDestructor(int int1) {  
        cout << "ConstructorDestructor constructor with 1 int." << endl;  
    }  
    ConstructorDestructor(int int1, int int2) : ConstructorDestructor(int1) {  
        cout << "ConstructorDestructor constructor with 2 ints." << endl;  
        throw exception();  
    }  
    ConstructorDestructor(int int1, int int2, int int3) : ConstructorDestructor(int1, int2) {  
        cout << "ConstructorDestructor constructor with 3 ints." << endl;  
    }  
  
    ~ConstructorDestructor() {  
        cout << "ConstructorDestructor destructor." << endl;  
    }  
};  
  
int main() {  
  
    try {  
        ConstructorDestructor cd{ 1, 2, 3 };  
    }  
    catch (const exception& ex){  
    }  
}
```

Output:

```
ConstructorDestructor constructor with 1 int.  
ConstructorDestructor constructor with 2 ints.  
ConstructorDestructor destructor.
```

For more information, see [Uniform Initialization and Delegating Constructors](#).

## Inheriting constructors (C++11)

A derived class can inherit the constructors from a direct base class by using a using declaration as shown in the following example:

```

#include <iostream>
using namespace std;

class Base
{
public:
    Base() { cout << "Base()" << endl; }
    Base(const Base& other) { cout << "Base(Base&)" << endl; }
    explicit Base(int i) : num(i) { cout << "Base(int)" << endl; }
    explicit Base(char c) : letter(c) { cout << "Base(char)" << endl; }

private:
    int num;
    char letter;
};

class Derived : Base
{
public:
    // Inherit all constructors from Base
    using Base::Base;

private:
    // Can't initialize newMember from Base constructors.
    int newMember{ 0 };
};

int main(int argc, char argv[])
{
    cout << "Derived d1(5) calls: ";
    Derived d1(5);
    cout << "Derived d1('c') calls: ";
    Derived d2('c');
    cout << "Derived d3 = d2 calls: " ;
    Derived d3 = d2;
    cout << "Derived d4 calls: ";
    Derived d4;

    // Keep console open in debug mode:
    cout << endl << "Press Enter to exit.";
    char in[1];
    cin.getline(in, 1);
    return 0;
}

/* Output:
Derived d1(5) calls: Base(int)
Derived d1('c') calls: Base(char)
Derived d3 = d2 calls: Base(Base&)
Derived d4 calls: Base()

Press Enter to exit.

```

The using statement brings into scope all constructors from the base class except those that have an identical signature to constructors in the derived class. In general, it is best to use inheriting constructors when the derived class declares no new data members or constructors.

A class template can inherit all the constructors from a type argument if that type specifies a base class:

```

template< typename T >
class Derived : T {
    using T::T;    // declare the constructors from T
    // ...
};

```

A deriving class cannot inherit from multiple base classes if those base classes have constructors that have an identical signature.

## Rules for Declaring Constructors

A constructor has the same name as its class. Any number of constructors can be declared, subject to the rules of overloaded functions.

The `argument-declaration-list` may be empty.

C++ defines two special kinds of constructors, default and copy constructors, described in the following table.

### Default and Copy Constructors

KIND OF CONSTRUCTION	ARGUMENTS	PURPOSE
Default constructor	Can be called with no arguments	Construct a default object of the class type
Copy constructor	Can accept a single argument of reference to same class type	Copy objects of the class type

Default constructors can be called with no arguments. However, you can declare a default constructor with an argument list, provided all arguments have defaults. Similarly, copy constructors must accept a single argument of reference to the same class type. More arguments can be supplied, provided all subsequent arguments have defaults.

If you do not supply any constructors, the compiler attempts to generate a default constructor. If you do not supply a copy constructor, the compiler attempts to generate one. These compiler-generated constructors are considered public member functions. An error is generated if you specify a copy constructor with a first argument that is an object and not a reference.

A compiler-generated default constructor sets up the object (initializes vftables and vbtables, as described previously), and it calls the default constructors for base classes and members, but it takes no other action. Base class and member constructors are called only if they exist, are accessible, and are unambiguous.

A compiler-generated copy constructor sets up a new object and performs a memberwise copy of the contents of the object to be copied. If base class or member constructors exist, they are called; otherwise, bitwise copying is performed.

If all base and member classes of a class `type` have copy constructors that accept a `const` argument, the compiler-generated copy constructor accepts a single argument of type `const type &`. Otherwise, the compiler-generated copy constructor accepts a single argument of type `type &`.

You can use a constructor to initialize a `const` or `volatile` object, but the constructor itself cannot be declared as `const` or `volatile`. The only legal storage class for a constructor is `inline`; use of any other storage-class modifier, including the `__declspec` keyword, with a constructor causes a compiler error.

The `stdcall` calling convention is used on static member functions and global functions declared with the `__stdcall` keyword, and that do not use a variable argument list. When you use the `__stdcall` keyword on a non-static member function, such as a constructor, the compiler will use the `thiscall` calling convention."

Constructors of base classes are not inherited by derived classes. When an object of derived class type is created, it is constructed starting with the base class components; then it moves to the derived class components. The compiler uses each base class's constructor as that part of the complete object is initialized (except in cases of virtual derivation).

## Explicitly invoking constructors

Constructors can be explicitly called in a program to create objects of a given type. For example, to create two `Point` objects that describe the ends of a line, the following code can be written:

```
DrawLine( Point( 13, 22 ), Point( 87, 91 ) );
```

Two objects of type `Point` are created, passed to the function `DrawLine`, and destroyed at the end of the expression (the function call).

Another context in which a constructor is explicitly called is in an initialization:

```
Point pt = Point( 7, 11 );
```

An object of type `Point` is created and initialized using the constructor that accepts two arguments of type `int`.

Objects that are created by calling constructors explicitly, as in the preceding two examples, are unnamed and have a lifetime of the expression in which they are created. This is discussed in greater detail in [Temporary Objects](#).

It is usually safe to call any member function from within a constructor because the object has been completely set up (virtual tables have been initialized and so on) prior to the execution of the first line of user code. However, it is potentially unsafe for a member function to call a virtual member function for an abstract base class during construction or destruction.

Constructors can call virtual functions. When virtual functions are called, the function invoked is the function defined for the constructor's own class (or inherited from its bases). The following example shows what happens when a virtual function is called from within a constructor:

```

// spec1_calling_virtual_functions.cpp
// compile with: /EHsc
#include <iostream>

using namespace std;
class Base
{
public:
    Base();           // Default constructor.
    virtual void f(); // Virtual member function.
};

Base::Base()
{
    cout << "Constructing Base sub-object\n";
    f();             // Call virtual member function
}                   // from inside constructor.

void Base::f()
{
    cout << "Called Base::f()\n";
}

class Derived : public Base
{
public:
    Derived();       // Default constructor.
    void f();        // Implementation of virtual
}                   // function f for this class.

Derived::Derived()
{
    cout << "Constructing Derived object\n";
}

void Derived::f()
{
    cout << "Called Derived::f()\n";
}

int main()
{
    Derived d;
}

```

When the preceding program is run, the declaration `Derived d` causes the following sequence of events:

1. The constructor for class `Derived` (`Derived::Derived`) is called.
2. Prior to entering the body of the `Derived` class's constructor, the constructor for class `Base` (`Base::Base`) is called.

`Base::Base` calls the function `f`, which is a virtual function. Ordinarily, `Derived::f` would be called because the object `d` is of type `Derived`. Because the `Base::Base` function is a constructor, the object is not yet of the `Derived` type, and `Base::f` is called.

# Copy Constructors and Copy Assignment Operators (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

## NOTE

Starting in C++11, two kinds of assignment are supported in the language: *copy assignment* and *move assignment*. In this article "assignment" means copy assignment unless explicitly stated otherwise. For information about move assignment, see [Move Constructors and Move Assignment Operators \(C++\)](#).

Both the assignment operation and the initialization operation cause objects to be copied.

- **Assignment:** When one object's value is assigned to another object, the first object is copied to the second object. Therefore,

```
Point a, b;  
...  
a = b;
```

causes the value of `b` to be copied to `a`.

- **Initialization:** Initialization occurs when a new object is declared, when arguments are passed to functions by value, or when values are returned from functions by value.

You can define the semantics of "copy" for objects of class type. For example, consider this code:

```
TextFile a, b;  
a.Open( "FILE1.DAT" );  
b.Open( "FILE2.DAT" );  
b = a;
```

The preceding code could mean "copy the contents of FILE1.DAT to FILE2.DAT" or it could mean "ignore FILE2.DAT and make `b` a second handle to FILE1.DAT." You must attach appropriate copying semantics to each class, as follows.

- By using the assignment operator `operator=` together with a reference to the class type as the return type and the parameter that is passed by `const` reference—for example  
`ClassName& operator=(const ClassName& x);`.
- By using the copy constructor.

If you do not declare a copy constructor, the compiler generates a member-wise copy constructor for you. If you do not declare a copy assignment operator, the compiler generates a member-wise copy assignment operator for you. Declaring a copy constructor does not suppress the compiler-generated copy assignment operator, nor vice versa. If you implement either one, we recommend that you also implement the other one so that the meaning of the code is clear.

The copy constructor takes an argument of type `class-name&`, where `class-name` is the name of the class for which the constructor is defined. For example:

```
// spec1_copying_class_objects.cpp
class Window
{
public:
    Window( const Window& ); // Declare copy constructor.
    // ...
};

int main()
{
}
```

#### NOTE

Make the type of the copy constructor's argument `const class-name&` whenever possible. This prevents the copy constructor from accidentally changing the object from which it is copying. It also enables copying from `const` objects.

## Compiler generated copy constructors

Compiler-generated copy constructors, like user-defined copy constructors, have a single argument of type "reference to *class-name*." An exception is when all base classes and member classes have copy constructors declared as taking a single argument of type `const class-name&`. In such a case, the compiler-generated copy constructor's argument is also `const`.

When the argument type to the copy constructor is not `const`, initialization by copying a `const` object generates an error. The reverse is not true: If the argument is `const`, you can initialize by copying an object that is not `const`.

Compiler-generated assignment operators follow the same pattern with regard to `const`. They take a single argument of type `class-name&` unless the assignment operators in all base and member classes take arguments of type `const class-name&`. In this case, the class's generated assignment operator takes a `const` argument.

#### NOTE

When virtual base classes are initialized by copy constructors, compiler-generated or user-defined, they are initialized only once: at the point when they are constructed.

The implications are similar to those of the copy constructor. When the argument type is not `const`, assignment from a `const` object generates an error. The reverse is not true: If a `const` value is assigned to a value that is not `const`, the assignment succeeds.

For more information about overloaded assignment operators, see [Assignment](#).

# Move Constructors and Move Assignment Operators (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

This topic describes how to write a *move constructor* and a move assignment operator for a C++ class. A move constructor enables you to implement move semantics, which can significantly improve the performance of your applications. For more information about move semantics, see [Rvalue Reference Declarator: &&](#).

This topic builds upon the following C++ class, `MemoryBlock`, which manages a memory buffer.

```
// MemoryBlock.h
#pragma once
#include <iostream>
#include <algorithm>

class MemoryBlock
{
public:

    // Simple constructor that initializes the resource.
    explicit MemoryBlock(size_t length)
        : _length(length)
        , _data(new int[length])
    {
        std::cout << "In MemoryBlock(size_t). length = "
            << _length << "." << std::endl;
    }

    // Destructor.
    ~MemoryBlock()
    {
        std::cout << "In ~MemoryBlock(). length = "
            << _length << ".";

        if (_data != nullptr)
        {
            std::cout << " Deleting resource.";
            // Delete the resource.
            delete[] _data;
        }

        std::cout << std::endl;
    }

    // Copy constructor.
    MemoryBlock(const MemoryBlock& other)
        : _length(other._length)
        , _data(new int[other._length])
    {
        std::cout << "In MemoryBlock(const MemoryBlock&). length = "
            << other._length << ". Copying resource." << std::endl;

        std::copy(other._data, other._data + _length, _data);
    }

    // Copy assignment operator.
    MemoryBlock& operator=(const MemoryBlock& other)
    {
        std::cout << "In operator=(const MemoryBlock&). length = "
            << other._length << ". Copying resource." << std::endl;
    }
}
```

```

        if (this != &other)
        {
            // Free the existing resource.
            delete[] _data;

            _length = other._length;
            _data = new int[_length];
            std::copy(other._data, other._data + _length, _data);
        }
        return *this;
    }

    // Retrieves the length of the data resource.
    size_t Length() const
    {
        return _length;
    }

private:
    size_t _length; // The length of the resource.
    int* _data; // The resource.
};


```

The following procedures describe how to write a move constructor and a move assignment operator for the example C++ class.

#### To create a move constructor for a C++ class

1. Define an empty constructor method that takes an rvalue reference to the class type as its parameter, as demonstrated in the following example:

```

MemoryBlock(MemoryBlock&& other)
    : _data(nullptr)
    , _length(0)
{
}

```

2. In the move constructor, assign the class data members from the source object to the object that is being constructed:

```

    _data = other._data;
    _length = other._length;

```

3. Assign the data members of the source object to default values. This prevents the destructor from freeing resources (such as memory) multiple times:

```

    other._data = nullptr;
    other._length = 0;

```

#### To create a move assignment operator for a C++ class

1. Define an empty assignment operator that takes an rvalue reference to the class type as its parameter and returns a reference to the class type, as demonstrated in the following example:

```

MemoryBlock& operator=(MemoryBlock&& other)
{
}

```

2. In the move assignment operator, add a conditional statement that performs no operation if you try to assign the object to itself.

```
if (this != &other)
{
}
```

3. In the conditional statement, free any resources (such as memory) from the object that is being assigned to.

The following example frees the `_data` member from the object that is being assigned to:

```
// Free the existing resource.
delete[] _data;
```

Follow steps 2 and 3 in the first procedure to transfer the data members from the source object to the object that is being constructed:

```
// Copy the data pointer and its length from the
// source object.
_data = other._data;
_length = other._length;

// Release the data pointer from the source object so that
// the destructor does not free the memory multiple times.
other._data = nullptr;
other._length = 0;
```

4. Return a reference to the current object, as shown in the following example:

```
return *this;
```

## Example

The following example shows the complete move constructor and move assignment operator for the `MemoryBlock` class:

```

// Move constructor.
MemoryBlock(MemoryBlock&& other)
    : _data(nullptr)
    , _length(0)
{
    std::cout << "In MemoryBlock(MemoryBlock&&). length = "
        << other._length << ". Moving resource." << std::endl;

    // Copy the data pointer and its length from the
    // source object.
    _data = other._data;
    _length = other._length;

    // Release the data pointer from the source object so that
    // the destructor does not free the memory multiple times.
    other._data = nullptr;
    other._length = 0;
}

// Move assignment operator.
MemoryBlock& operator=(MemoryBlock&& other)
{
    std::cout << "In operator=(MemoryBlock&&). length = "
        << other._length << "." << std::endl;

    if (this != &other)
    {
        // Free the existing resource.
        delete[] _data;

        // Copy the data pointer and its length from the
        // source object.
        _data = other._data;
        _length = other._length;

        // Release the data pointer from the source object so that
        // the destructor does not free the memory multiple times.
        other._data = nullptr;
        other._length = 0;
    }
    return *this;
}

```

## Example

The following example shows how move semantics can improve the performance of your applications. The example adds two elements to a vector object and then inserts a new element between the two existing elements. In Visual C++ 2010, the `vector` class uses move semantics to perform the insertion operation efficiently by moving the elements of the vector instead of copying them.

```

// rvalue-references-move-semantics.cpp
// compile with: /EHsc
#include "MemoryBlock.h"
#include <vector>

using namespace std;

int main()
{
    // Create a vector object and add a few elements to it.
    vector<MemoryBlock> v;
    v.push_back(MemoryBlock(25));
    v.push_back(MemoryBlock(75));

    // Insert a new element into the second position of the vector.
    v.insert(v.begin() + 1, MemoryBlock(50));
}

```

This example produces the following output:

```

In MemoryBlock(size_t). length = 25.
In MemoryBlock(MemoryBlock&&). length = 25. Moving resource.
In ~MemoryBlock(). length = 0.
In MemoryBlock(size_t). length = 75.
In MemoryBlock(MemoryBlock&&). length = 25. Moving resource.
In ~MemoryBlock(). length = 0.
In MemoryBlock(MemoryBlock&&). length = 75. Moving resource.
In ~MemoryBlock(). length = 0.
In MemoryBlock(size_t). length = 50.
In MemoryBlock(MemoryBlock&&). length = 50. Moving resource.
In MemoryBlock(MemoryBlock&&). length = 50. Moving resource.
In operator=(MemoryBlock&&). length = 75.
In operator=(MemoryBlock&&). length = 50.
In ~MemoryBlock(). length = 0.
In ~MemoryBlock(). length = 0.
In ~MemoryBlock(). length = 25. Deleting resource.
In ~MemoryBlock(). length = 50. Deleting resource.
In ~MemoryBlock(). length = 75. Deleting resource.

```

Before Visual C++ 2010, this example produces the following output:

```

In MemoryBlock(size_t). length = 25.
In MemoryBlock(const MemoryBlock&). length = 25. Copying resource.
In ~MemoryBlock(). length = 25. Deleting resource.
In MemoryBlock(size_t). length = 75.
In MemoryBlock(const MemoryBlock&). length = 25. Copying resource.
In ~MemoryBlock(). length = 25. Deleting resource.
In MemoryBlock(const MemoryBlock&). length = 75. Copying resource.
In ~MemoryBlock(). length = 75. Deleting resource.
In MemoryBlock(size_t). length = 50.
In MemoryBlock(const MemoryBlock&). length = 50. Copying resource.
In MemoryBlock(const MemoryBlock&). length = 50. Copying resource.
In operator=(const MemoryBlock&). length = 75. Copying resource.
In operator=(const MemoryBlock&). length = 50. Copying resource.
In ~MemoryBlock(). length = 50. Deleting resource.
In ~MemoryBlock(). length = 50. Deleting resource.
In ~MemoryBlock(). length = 25. Deleting resource.
In ~MemoryBlock(). length = 50. Deleting resource.
In ~MemoryBlock(). length = 75. Deleting resource.

```

The version of this example that uses move semantics is more efficient than the version that does not use move semantics because it performs fewer copy, memory allocation, and memory deallocation operations.

# Robust Programming

To prevent resource leaks, always free resources (such as memory, file handles, and sockets) in the move assignment operator.

To prevent the unrecoverable destruction of resources, properly handle self-assignment in the move assignment operator.

If you provide both a move constructor and a move assignment operator for your class, you can eliminate redundant code by writing the move constructor to call the move assignment operator. The following example shows a revised version of the move constructor that calls the move assignment operator:

```
// Move constructor.  
MemoryBlock(MemoryBlock&& other)  
    : _data(nullptr)  
    , _length(0)  
{  
    *this = std::move(other);  
}
```

The `std::move` function preserves the rvalue property of the `other` parameter.

## See Also

[Rvalue Reference Declarator: &&](#)  
[<utility> move](#)

# Destructors (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

A destructor is a member function that is invoked automatically when the object goes out of scope or is explicitly destroyed by a call to `delete`. A destructor has the same name as the class, preceded by a tilde (`~`). For example, the destructor for class `String` is declared: `~String()`. If you do not define a destructor, the compiler will provide a default one; for many classes this is sufficient. You only need to define a custom destructor when the class stores handles to system resources that need to be released, or pointers that own the memory they point to.

Consider the following declaration of a `String` class:

```
// spec1_destructors.cpp
#include <string.h>

class String {
public:
    String( char *ch ); // Declare constructor
    ~String();          // and destructor.
private:
    char    *_text;
    size_t  sizeOfText;
};

// Define the constructor.
String::String( char *ch ) {
    sizeOfText = strlen( ch ) + 1;

    // Dynamically allocate the correct amount of memory.
    _text = new char[ sizeOfText ];

    // If the allocation succeeds, copy the initialization string.
    if( _text )
        strcpy_s( _text, sizeOfText, ch );
}

// Define the destructor.
String::~String() {
    // Deallocate the memory that was previously reserved
    // for this string.
    if ( _text )
        delete[] _text;
}

int main() {
    String str("The piper in the glen...");
}
```

In the preceding example, the destructor `String::~String` uses the `delete` operator to deallocate the space dynamically allocated for text storage.

## Declaring destructors

Destructors are functions with the same name as the class but preceded by a tilde (`~`)

Several rules govern the declaration of destructors. Destructors:

- Do not accept arguments.

- Do not return a value (or `void`).
- Cannot be declared as **const**, **volatile**, or **static**. However, they can be invoked for the destruction of objects declared as **const**, **volatile**, or **static**.
- Can be declared as **virtual**. Using virtual destructors, you can destroy objects without knowing their type — the correct destructor for the object is invoked using the virtual function mechanism. Note that destructors can also be declared as pure virtual functions for abstract classes.

## Using destructors

Destructors are called when one of the following events occurs:

- A local (automatic) object with block scope goes out of scope.
- An object allocated using the **new** operator is explicitly deallocated using **delete**.
- The lifetime of a temporary object ends.
- A program ends and global or static objects exist.
- The destructor is explicitly called using the destructor function's fully qualified name.

Destructors can freely call class member functions and access class member data.

There are two restrictions on the use of destructors:

- you cannot take its address
- derived classes do not inherit the destructor of their base class.

## Order of destruction

When an object goes out of scope or is deleted, the sequence of events in its complete destruction is as follows:

1. The class's destructor is called, and the body of the destructor function is executed.
2. Destructors for nonstatic member objects are called in the reverse order in which they appear in the class declaration. The optional member initialization list used in construction of these members does not affect the order of construction or destruction.
3. Destructors for nonvirtual base classes are called in the reverse order of declaration.
4. Destructors for virtual base classes are called in the reverse order of declaration.

```

// order_of_destruction.cpp
#include <stdio.h>

struct A1      { virtual ~A1() { printf("A1 dtor\n"); } };
struct A2 : A1 { virtual ~A2() { printf("A2 dtor\n"); } };
struct A3 : A2 { virtual ~A3() { printf("A3 dtor\n"); } };

struct B1      { ~B1() { printf("B1 dtor\n"); } };
struct B2 : B1 { ~B2() { printf("B2 dtor\n"); } };
struct B3 : B2 { ~B3() { printf("B3 dtor\n"); } };

int main() {
    A1 * a = new A3;
    delete a;
    printf("\n");

    B1 * b = new B3;
    delete b;
    printf("\n");

    B3 * b2 = new B3;
    delete b2;
}

Output: A3 dtor
A2 dtor
A1 dtor

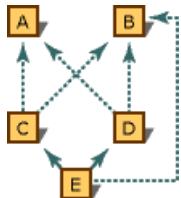
B1 dtor

B3 dtor
B2 dtor
B1 dtor

```

## Virtual base classes

Destructors for virtual base classes are called in the reverse order of their appearance in a directed acyclic graph (depth-first, left-to-right, postorder traversal). the following figure depicts an inheritance graph.



Inheritance Graph Showing Virtual Base Classes

The following lists the class heads for the classes shown in the figure.

```

class A
class B
class C : virtual public A, virtual public B
class D : virtual public A, virtual public B
class E : public C, public D, virtual public B

```

To determine the order of destruction of the virtual base classes of an object of type `E`, the compiler builds a list by applying the following algorithm:

1. Traverse the graph left, starting at the deepest point in the graph (in this case, `E`).
2. Perform leftward traversals until all nodes have been visited. Note the name of the current node.
3. Revisit the previous node (down and to the right) to find out whether the node being remembered is a

virtual base class.

4. If the remembered node is a virtual base class, scan the list to see whether it has already been entered. If it is not a virtual base class, ignore it.
5. If the remembered node is not yet in the list, add it to the bottom of the list.
6. Traverse the graph up and along the next path to the right.
7. Go to step 2.
8. When the last upward path is exhausted, note the name of the current node.
9. Go to step 3.
10. Continue this process until the bottom node is again the current node.

Therefore, for class `E`, the order of destruction is:

11. The nonvirtual base class `E`.
12. The nonvirtual base class `D`.
13. The nonvirtual base class `C`.
14. The virtual base class `B`.
15. The virtual base class `A`.

This process produces an ordered list of unique entries. No class name appears twice. Once the list is constructed, it is walked in reverse order, and the destructor for each of the classes in the list from the last to the first is called.

The order of construction or destruction is primarily important when constructors or destructors in one class rely on the other component being created first or persisting longer — for example, if the destructor for `A` (in the figure shown above) relied on `B` still being present when its code executed, or vice versa.

Such interdependencies between classes in an inheritance graph are inherently dangerous because classes derived later can alter which is the leftmost path, thereby changing the order of construction and destruction.

### Nonvirtual base classes

The destructors for nonvirtual base classes are called in the reverse order in which the base class names are declared. Consider the following class declaration:

```
class MultInherit : public Base1, public Base2  
...
```

In the preceding example, the destructor for `Base2` is called before the destructor for `Base1`.

## Explicit destructor calls

Calling a destructor explicitly is seldom necessary. However, it can be useful to perform cleanup of objects placed at absolute addresses. These objects are commonly allocated using a user-defined `new` operator that takes a placement argument. The `delete` operator cannot deallocate this memory because it is not allocated from the free store (for more information, see [The new and delete Operators](#)). A call to the destructor, however, can perform appropriate cleanup. To explicitly call the destructor for an object, `s`, of class `String`, use one of the following statements:

```
s.String::~String();      // Nonvirtual call  
ps->String::~String();  // Nonvirtual call  
  
s.~String();            // Virtual call  
ps->~String();         // Virtual call
```

The notation for explicit calls to destructors, shown in the preceding, can be used regardless of whether the type defines a destructor. This allows you to make such explicit calls without knowing if a destructor is defined for the type. An explicit call to a destructor where none is defined has no effect.

# Overview of Member Functions

12/21/2017 • 1 min to read • [Edit Online](#)

Member functions are either static or nonstatic. The behavior of static member functions differs from other member functions because static member functions have no implicit **this** argument. Nonstatic member functions have a **this** pointer. Member functions, whether static or nonstatic, can be defined either in or outside the class declaration.

If a member function is defined inside a class declaration, it is treated as an inline function, and there is no need to qualify the function name with its class name. Although functions defined inside class declarations are already treated as inline functions, you can use the **inline** keyword to document code.

An example of declaring a function within a class declaration follows:

```
// overview_of_member_functions1.cpp
class Account
{
public:
    // Declare the member function Deposit within the declaration
    // of class Account.
    double Deposit( double HowMuch )
    {
        balance += HowMuch;
        return balance;
    }
private:
    double balance;
};

int main()
{
}
```

If a member function's definition is outside the class declaration, it is treated as an inline function only if it is explicitly declared as **inline**. In addition, the function name in the definition must be qualified with its class name using the scope-resolution operator (`::`).

The following example is identical to the previous declaration of class `Account`, except that the `Deposit` function is defined outside the class declaration:

```
// overview_of_member_functions2.cpp
class Account
{
public:
    // Declare the member function Deposit but do not define it.
    double Deposit( double HowMuch );
private:
    double balance;
};

inline double Account::Deposit( double HowMuch )
{
    balance += HowMuch;
    return balance;
}

int main()
{
```

#### NOTE

Although member functions can be defined either inside a class declaration or separately, no member functions can be added to a class after the class is defined.

Classes containing member functions can have many declarations, but the member functions themselves must have only one definition in a program. Multiple definitions cause an error message at link time. If a class contains inline function definitions, the function definitions must be identical to observe this "one definition" rule.

# virtual Specifier

12/21/2017 • 1 min to read • [Edit Online](#)

The **virtual** keyword can be applied only to nonstatic class member functions. It signifies that binding of calls to the function is deferred until run time. For more information, see [Virtual Functions](#).

# override Specifier

12/21/2017 • 1 min to read • [Edit Online](#)

You can use the `override` keyword to designate member functions that override a virtual function in a base class.

## Syntax

```
function-declaration override;
```

## Remarks

`override` is context-sensitive and has special meaning only when it's used after a member function declaration; otherwise, it's not a reserved keyword.

## Example

Use `override` to help prevent inadvertent inheritance behavior in your code. The following example shows where, without using `override`, the member function behavior of the derived class may not have been intended. The compiler doesn't emit any errors for this code.

```
class BaseClass
{
    virtual void funcA();
    virtual void funcB() const;
    virtual void funcC(int = 0);
    void funcD();
};

class DerivedClass: public BaseClass
{
    virtual void funcA(); // ok, works as intended

    virtual void funcB(); // DerivedClass::funcB() is non-const, so it does not
                         // override BaseClass::funcB() const and it is a new member function

    virtual void funcC(double = 0.0); // DerivedClass::funcC(double) has a different
                                    // parameter type than BaseClass::funcC(int), so
                                    // DerivedClass::funcC(double) is a new member function
};
```

When you use `override`, the compiler generates errors instead of silently creating new member functions.

```
class BaseClass
{
    virtual void funcA();
    virtual void funcB() const;
    virtual void funcC(int = 0);
    void funcD();
};

class DerivedClass: public BaseClass
{
    virtual void funcA() override; // ok

    virtual void funcB() override; // compiler error: DerivedClass::funcB() does not
                                  // override BaseClass::funcB() const

    virtual void funcC( double = 0.0 ) override; // compiler error:
                                                // DerivedClass::funcC(double) does not
                                                // override BaseClass::funcC(int)

    void funcD() override; // compiler error: DerivedClass::funcD() does not
                          // override the non-virtual BaseClass::funcD()
};
```

To specify that functions cannot be overridden and that classes cannot be inherited, use the [final](#) keyword.

## See Also

[final Specifier](#)

[Keywords](#)

# final Specifier

12/21/2017 • 1 min to read • [Edit Online](#)

You can use the `final` keyword to designate virtual functions that cannot be overridden in a derived class. You can also use it to designate classes that cannot be inherited.

## Syntax

```
function-declaration final;
```

```
class class-name final base-classes
```

## Remarks

`final` is context-sensitive and has special meaning only when it's used after a function declaration or class name; otherwise, it's not a reserved keyword.

When `final` is used in class declarations, `base-classes` is an optional part of the declaration.

## Example

The following example uses the `final` keyword to specify that a virtual function cannot be overridden.

```
class BaseClass
{
    virtual void func() final;
};

class DerivedClass: public BaseClass
{
    virtual void func(); // compiler error: attempting to
                        // override a final function
};
```

For information about how to specify that member functions can be overridden, see [override Specifier](#).

The next example uses the `final` keyword to specify that a class cannot be inherited.

```
class BaseClass final
{
};

class DerivedClass: public BaseClass // compiler error: BaseClass is
                                    // marked as non-inheritable
{
```

## See Also

Keywords

override Specifier

# Inheritance (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This section explains how to use derived classes to produce extensible programs.

## Overview

New classes can be derived from existing classes using a mechanism called "inheritance" (see the information beginning in [Single Inheritance](#)). Classes that are used for derivation are called "base classes" of a particular derived class. A derived class is declared using the following syntax:

```
class Derived : [virtual] [access-specifier] Base
{
    // member list
};

class Derived : [virtual] [access-specifier] Base1,
[virtual] [access-specifier] Base2, . . .
{
    // member list
};
```

After the tag (name) for the class, a colon appears followed by a list of base specifications. The base classes so named must have been declared previously. The base specifications may contain an access specifier, which is one of the keywords **public**, **protected** or **private**. These access specifiers appear before the base class name and apply only to that base class. These specifiers control the derived class's permission to use to members of the base class. See [Member-Access Control](#) for information on access to base class members. If the access specifier is omitted, the access to that base is considered **private**. The base specifications may contain the keyword **virtual** to indicate virtual inheritance. This keyword may appear before or after the access specifier, if any. If virtual inheritance is used, the base class is referred to as a virtual base class.

Multiple base classes can be specified, separated by commas. If a single base class is specified, the inheritance model is [Single inheritance](#). If more than one base class is specified, the inheritance model is called [Multiple inheritance](#),

The following topics are included:

- [Single inheritance](#)
- [Multiple inheritance](#)
- [Multiple base classes](#)
- [Virtual functions](#)
- [Explicit overrides](#)
- [Abstract classes](#)
- [Summary of scope rules](#)

The `_super` and `_interface` keywords are documented in this section.

## See Also

[C++ Language Reference](#)

# Virtual Functions

12/21/2017 • 4 min to read • [Edit Online](#)

A virtual function is a member function that you expect to be redefined in derived classes. When you refer to a derived class object using a pointer or a reference to the base class, you can call a virtual function for that object and execute the derived class's version of the function.

Virtual functions ensure that the correct function is called for an object, regardless of the expression used to make the function call.

Suppose a base class contains a function declared as `virtual` and a derived class defines the same function. The function from the derived class is invoked for objects of the derived class, even if it is called using a pointer or reference to the base class. The following example shows a base class that provides an implementation of the `PrintBalance` function and two derived classes

```
// deriv_VirtualFunctions.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

class Account {
public:
    Account( double d ) { _balance = d; }
    virtual double GetBalance() { return _balance; }
    virtual void PrintBalance() { cerr << "Error. Balance not available for base type." << endl; }
private:
    double _balance;
};

class CheckingAccount : public Account {
public:
    CheckingAccount(double d) : Account(d) {}
    void PrintBalance() { cout << "Checking account balance: " << GetBalance() << endl; }
};

class SavingsAccount : public Account {
public:
    SavingsAccount(double d) : Account(d) {}
    void PrintBalance() { cout << "Savings account balance: " << GetBalance(); }
};

int main() {
    // Create objects of type CheckingAccount and SavingsAccount.
    CheckingAccount *pChecking = new CheckingAccount( 100.00 );
    SavingsAccount *pSavings = new SavingsAccount( 1000.00 );

    // Call PrintBalance using a pointer to Account.
    Account *pAccount = pChecking;
    pAccount->PrintBalance();

    // Call PrintBalance using a pointer to Account.
    pAccount = pSavings;
    pAccount->PrintBalance();
}
```

In the preceding code, the calls to `PrintBalance` are identical, except for the object `pAccount` points to. Because `PrintBalance` is virtual, the version of the function defined for each object is called. The `PrintBalance` function in the derived classes `CheckingAccount` and `SavingsAccount` "override" the function in the base class `Account`.

If a class is declared that does not provide an overriding implementation of the `PrintBalance` function, the default implementation from the base class `Account` is used.

Functions in derived classes override virtual functions in base classes only if their type is the same. A function in a derived class cannot differ from a virtual function in a base class in its return type only; the argument list must differ as well.

When calling a function using pointers or references, the following rules apply:

- A call to a virtual function is resolved according to the underlying type of object for which it is called.
- A call to a nonvirtual function is resolved according to the type of the pointer or reference.

The following example shows how virtual and nonvirtual functions behave when called through pointers:

```
// deriv_VirtualFunctions2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

class Base {
public:
    virtual void NameOf(); // Virtual function.
    void InvokingClass(); // Nonvirtual function.
};

// Implement the two functions.
void Base::NameOf() {
    cout << "Base::NameOf\n";
}

void Base::InvokingClass() {
    cout << "Invoked by Base\n";
}

class Derived : public Base {
public:
    void NameOf(); // Virtual function.
    void InvokingClass(); // Nonvirtual function.
};

// Implement the two functions.
void Derived::NameOf() {
    cout << "Derived::NameOf\n";
}

void Derived::InvokingClass() {
    cout << "Invoked by Derived\n";
}

int main() {
    // Declare an object of type Derived.
    Derived aDerived;

    // Declare two pointers, one of type Derived * and the other
    // of type Base *, and initialize them to point to aDerived.
    Derived *pDerived = &aDerived;
    Base    *pBase   = &aDerived;

    // Call the functions.
    pBase->NameOf();           // Call virtual function.
    pBase->InvokingClass();    // Call nonvirtual function.
    pDerived->NameOf();         // Call virtual function.
    pDerived->InvokingClass(); // Call nonvirtual function.
}
```

## Output

```
Derived::NameOf  
Invoked by Base  
Derived::NameOf  
Invoked by Derived
```

Note that regardless of whether the `NameOf` function is invoked through a pointer to `Base` or a pointer to `Derived`, it calls the function for `Derived`. It calls the function for `Derived` because `NameOf` is a virtual function, and both `pBase` and `pDerived` point to an object of type `Derived`.

Because virtual functions are called only for objects of class types, you cannot declare global or static functions as **virtual**.

The **virtual** keyword can be used when declaring overriding functions in a derived class, but it is unnecessary; overrides of virtual functions are always virtual.

Virtual functions in a base class must be defined unless they are declared using the *pure-specifier*. (For more information about pure virtual functions, see [Abstract Classes](#).)

The virtual function-call mechanism can be suppressed by explicitly qualifying the function name using the scope-resolution operator (`::`). Consider the earlier example involving the `Account` class. To call `PrintBalance` in the base class, use code such as the following:

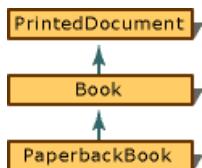
```
CheckingAccount *pChecking = new CheckingAccount( 100.00 );  
  
pChecking->Account::PrintBalance(); // Explicit qualification.  
  
Account *pAccount = pChecking; // Call Account::PrintBalance  
  
pAccount->Account::PrintBalance(); // Explicit qualification.
```

Both calls to `PrintBalance` in the preceding example suppress the virtual function-call mechanism.

# Single Inheritance

12/21/2017 • 4 min to read • [Edit Online](#)

In "single inheritance," a common form of inheritance, classes have only one base class. Consider the relationship illustrated in the following figure.



Simple Single-Inheritance Graph

Note the progression from general to specific in the figure. Another common attribute found in the design of most class hierarchies is that the derived class has a "kind of" relationship with the base class. In the figure, a `Book` is a kind of a `PrintedDocument`, and a `PaperbackBook` is a kind of a `book`.

One other item of note in the figure: `Book` is both a derived class (from `PrintedDocument`) and a base class (`PaperbackBook` is derived from `Book`). A skeletal declaration of such a class hierarchy is shown in the following example:

```
// deriv_SingleInheritance.cpp
// compile with: /LD
class PrintedDocument {};

// Book is derived from PrintedDocument.
class Book : public PrintedDocument {};

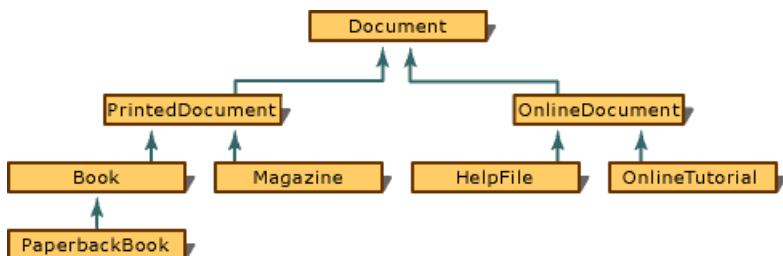
// PaperbackBook is derived from Book.
class PaperbackBook : public Book {};
```

`PrintedDocument` is considered a "direct base" class to `Book`; it is an "indirect base" class to `PaperbackBook`. The difference is that a direct base class appears in the base list of a class declaration and an indirect base does not.

The base class from which each class is derived is declared before the declaration of the derived class. It is not sufficient to provide a forward-referencing declaration for a base class; it must be a complete declaration.

In the preceding example, the access specifier **public** is used. The meaning of public, protected, and private inheritance is described in [Member-Access Control](#).

A class can serve as the base class for many specific classes, as illustrated in the following figure.



Sample of Directed Acyclic Graph

In the diagram shown above, called a "directed acyclic graph" (or "DAG"), some of the classes are base classes for more than one derived class. However, the reverse is not true: there is only one direct base class for any given derived class. The graph in the figure depicts a "single inheritance" structure.

### NOTE

Directed acyclic graphs are not unique to single inheritance. They are also used to depict multiple-inheritance graphs. This topic is covered in [Multiple Inheritance](#).

In inheritance, the derived class contains the members of the base class plus any new members you add. As a result, a derived class can refer to members of the base class (unless those members are redefined in the derived class). The scope-resolution operator ( `::` ) can be used to refer to members of direct or indirect base classes when those members have been redefined in the derived class. Consider this example:

```
// deriv_SingleInheritance2.cpp
// compile with: /EHsc /c
#include <iostream>
using namespace std;
class Document {
public:
    char *Name;    // Document name.
    void PrintNameOf(); // Print name.
};

// Implementation of PrintNameOf function from class Document.
void Document::PrintNameOf() {
    cout << Name << endl;
}

class Book : public Document {
public:
    Book( char *name, long pagecount );
private:
    long PageCount;
};

// Constructor from class Book.
Book::Book( char *name, long pagecount ) {
    Name = new char[ strlen( name ) + 1 ];
    strcpy_s( Name, strlen( Name ), name );
    PageCount = pagecount;
}
```

Note that the constructor for `Book`, ( `Book::Book` ), has access to the data member, `Name`. In a program, an object of type `Book` can be created and used as follows:

```
// Create a new object of type Book. This invokes the
// constructor Book::Book.
Book LibraryBook( "Programming Windows, 2nd Ed", 944 );

...
// Use PrintNameOf function inherited from class Document.
LibraryBook.PrintNameOf();
```

As the preceding example demonstrates, class-member and inherited data and functions are used identically. If the implementation for class `Book` calls for a reimplementaion of the `PrintNameOf` function, the function that belongs to the `Document` class can be called only by using the scope-resolution ( `::` ) operator:

```

// deriv_SingleInheritance3.cpp
// compile with: /EHsc /LD
#include <iostream>
using namespace std;

class Document {
public:
    char *Name;           // Document name.
    void PrintNameOf() {} // Print name.
};

class Book : public Document {
    Book( char *name, long pagecount );
    void PrintNameOf();
    long PageCount;
};

void Book::PrintNameOf() {
    cout << "Name of book: ";
    Document::PrintNameOf();
}

```

Pointers and references to derived classes can be implicitly converted to pointers and references to their base classes if there is an accessible, unambiguous base class. The following code demonstrates this concept using pointers (the same principle applies to references):

```

// deriv_SingleInheritance4.cpp
// compile with: /W3
struct Document {
    char *Name;
    void PrintNameOf() {}
};

class PaperbackBook : public Document {};

int main() {
    Document * DocLib[10]; // Library of ten documents.
    for ( int i = 0 ; i < 10 ; i++)
        DocLib[i] = new Document;
}

```

In the preceding example, different types are created. However, because these types are all derived from the `Document` class, there is an implicit conversion to `Document *`. As a result, `DocLib` is a "heterogeneous list" (a list in which not all objects are of the same type) containing different kinds of objects.

Because the `Document` class has a `PrintNameOf` function, it can print the name of each book in the library, although it may omit some of the information specific to the type of document (page count for `Book`, number of bytes for `HelpFile`, and so on).

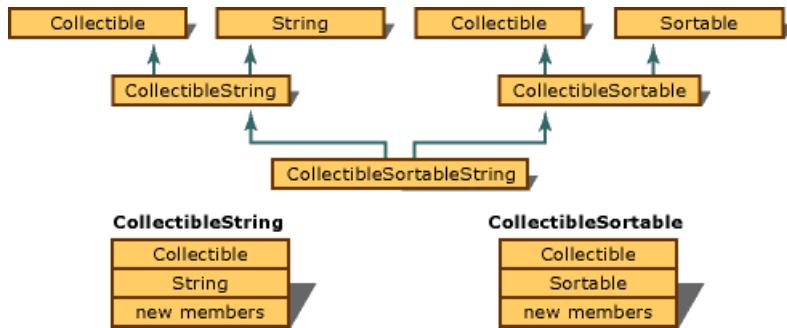
#### **NOTE**

Forcing the base class to implement a function such as `PrintNameOf` is often not the best design. [Virtual Functions](#) offers other design alternatives.

# Base Classes

12/21/2017 • 1 min to read • [Edit Online](#)

The inheritance process creates a new derived class that is made up of the members of the base class(es) plus any new members added by the derived class. In a multiple-inheritance, it is possible to construct an inheritance graph where the same base class is part of more than one of the derived classes. The following figure shows such a graph.



Multiple Instances of a Single Base Class

In the figure, pictorial representations of the components of `CollectibleString` and `CollectibleSortable` are shown. However, the base class, `Collectible`, is in `CollectibleSortableString` through the `CollectibleString` path and the `CollectibleSortable` path. To eliminate this redundancy, such classes can be declared as virtual base classes when they are inherited.

# Multiple Base Classes

12/21/2017 • 7 min to read • [Edit Online](#)

As described in [Multiple Inheritance](#), a class can be derived from more than one base class. In a multiple-inheritance model (where classes are derived from more than one base class), the base classes are specified using the *base-list* grammar element. For example, the class declaration for `CollectionOfBook`, derived from `Collection` and `Book`, can be specified:

```
// deriv_MultipleBaseClasses.cpp
// compile with: /LD
class Collection {
};
class Book {};
class CollectionOfBook : public Book, public Collection {
    // New members
};
```

The order in which base classes are specified is not significant except in certain cases where constructors and destructors are invoked. In these cases, the order in which base classes are specified affects the following:

- The order in which initialization by constructor takes place. If your code relies on the `Book` portion of `CollectionOfBook` to be initialized before the `Collection` part, the order of specification is significant. Initialization takes place in the order the classes are specified in the *base-list*.
- The order in which destructors are invoked to clean up. Again, if a particular "part" of the class must be present when the other part is being destroyed, the order is significant. Destructors are called in the reverse order of the classes specified in the *base-list*.

#### NOTE

The order of specification of base classes can affect the memory layout of the class. Do not make any programming decisions based on the order of base members in memory.

When specifying the *base-list*, you cannot specify the same class name more than once. However, it is possible for a class to be an indirect base to a derived class more than once.

## Virtual base classes

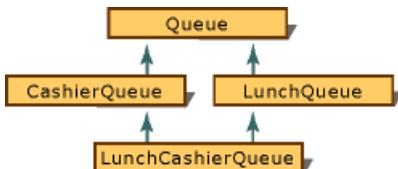
Because a class can be an indirect base class to a derived class more than once, C++ provides a way to optimize the way such base classes work. Virtual base classes offer a way to save space and avoid ambiguities in class hierarchies that use multiple inheritance.

Each nonvirtual object contains a copy of the data members defined in the base class. This duplication wastes space and requires you to specify which copy of the base class members you want whenever you access them.

When a base class is specified as a virtual base, it can act as an indirect base more than once without duplication of its data members. A single copy of its data members is shared by all the base classes that use it as a virtual base.

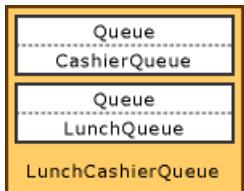
When declaring a virtual base class, the **virtual** keyword appears in the base lists of the derived classes.

Consider the class hierarchy in the following figure, which illustrates a simulated lunch line.



Simulated Lunch-Line Graph

In the figure, `Queue` is the base class for both `CashierQueue` and `LunchQueue`. However, when both classes are combined to form `LunchCashierQueue`, the following problem arises: the new class contains two subobjects of type `Queue`, one from `CashierQueue` and the other from `LunchQueue`. The following figure shows the conceptual memory layout (the actual memory layout might be optimized).



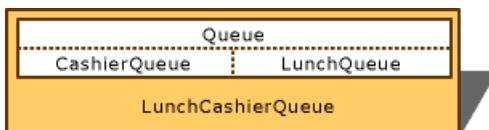
Simulated Lunch-Line Object

Note that there are two `Queue` subobjects in the `LunchCashierQueue` object. The following code declares `Queue` to be a virtual base class:

```

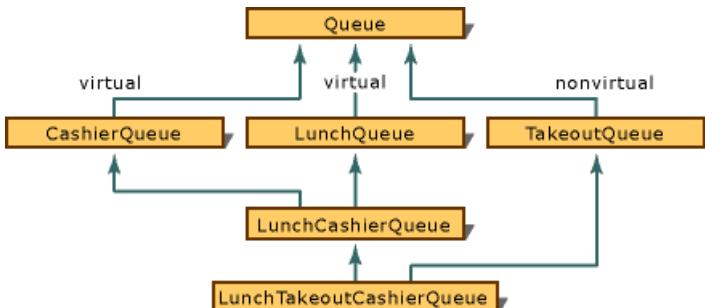
// deriv_VirtualBaseClasses.cpp
// compile with: /LD
class Queue {};
class CashierQueue : virtual public Queue {};
class LunchQueue : virtual public Queue {};
class LunchCashierQueue : public LunchQueue, public CashierQueue {};
  
```

The `virtual` keyword ensures that only one copy of the subobject `Queue` is included (see the following figure).



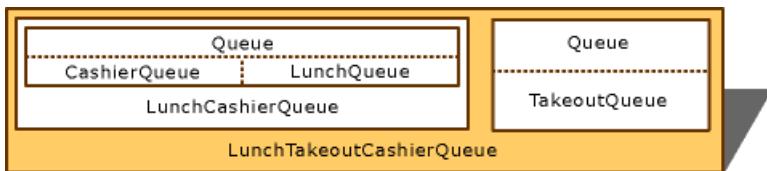
Simulated Lunch-Line Object with Virtual Base Classes

A class can have both a virtual component and a nonvirtual component of a given type. This happens in the conditions illustrated in the following figure.



Virtual and Nonvirtual Components of the Same Class

In the figure, `CashierQueue` and `LunchQueue` use `Queue` as a virtual base class. However, `TakeoutQueue` specifies `Queue` as a base class, not a virtual base class. Therefore, `LunchTakeoutCashierQueue` has two subobjects of type `Queue`: one from the inheritance path that includes `LunchCashierQueue` and one from the path that includes `TakeoutQueue`. This is illustrated in the following figure.



## Object Layout with Virtual and Nonvirtual Inheritance

### NOTE

Virtual inheritance provides significant size benefits when compared with nonvirtual inheritance. However, it can introduce extra processing overhead.

If a derived class overrides a virtual function that it inherits from a virtual base class, and if a constructor or a destructor for the derived base class calls that function using a pointer to the virtual base class, the compiler may introduce additional hidden "vtordisp" fields into the classes with virtual bases. The /vd0 compiler option suppresses the addition of the hidden vtordisp constructor/destructor displacement member. The /vd1 compiler option, the default, enables them where they are necessary. Turn off vtordisps only if you are sure that all class constructors and destructors call virtual functions virtually.

The /vd compiler option affects an entire compilation module. Use the **vtordisp** pragma to suppress and then reenable vtordisp fields on a class-by-class basis:

```
#pragma vtordisp( off )
class GetReal : virtual public { ... };
#pragma vtordisp( on )
```

## Name ambiguities

Multiple inheritance introduces the possibility for names to be inherited along more than one path. The class-member names along these paths are not necessarily unique. These name conflicts are called "ambiguities."

Any expression that refers to a class member must make an unambiguous reference. The following example shows how ambiguities develop:

```
// deriv_NameAmbiguities.cpp
// compile with: /LD
// Declare two base classes, A and B.
class A {
public:
    unsigned a;
    unsigned b();
};

class B {
public:
    unsigned a(); // Note that class A also has a member "a"
    int b();      // and a member "b".
    char c;
};

// Define class C as derived from A and B.
class C : public A, public B {};
```

Given the preceding class declarations, code such as the following is ambiguous because it is unclear whether **b** refers to the **b** in **A** or in **B**:

```
C *pc = new C;  
pc->b();
```

Consider the preceding example. Because the name `a` is a member of both class `A` and class `B`, the compiler cannot discern which `a` designates the function to be called. Access to a member is ambiguous if it can refer to more than one function, object, type, or enumerator.

The compiler detects ambiguities by performing tests in this order:

1. If access to the name is ambiguous (as just described), an error message is generated.
2. If overloaded functions are unambiguous, they are resolved.
3. If access to the name violates member-access permission, an error message is generated. (For more information, see [Member-Access Control](#).)

When an expression produces an ambiguity through inheritance, you can manually resolve it by qualifying the name in question with its class name. To make the preceding example compile properly with no ambiguities, use code such as:

```
C *pc = new C;  
pc->B::a();
```

#### NOTE

When `c` is declared, it has the potential to cause errors when `B` is referenced in the scope of `c`. No error is issued, however, until an unqualified reference to `B` is actually made in `c`'s scope.

## Dominance

It is possible for more than one name (function, object, or enumerator) to be reached through an inheritance graph. Such cases are considered ambiguous with nonvirtual base classes. They are also ambiguous with virtual base classes, unless one of the names "dominates" the others.

A name dominates another name if it is defined in both classes and one class is derived from the other. The dominant name is the name in the derived class; this name is used when an ambiguity would otherwise have arisen, as shown in the following example:

```

// deriv_Dominance.cpp
// compile with: /LD
class A {
public:
    int a;
};

class B : public virtual A {
public:
    int a();
};

class C : public virtual A {};

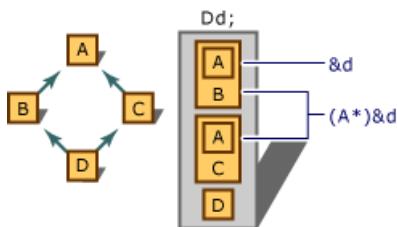
class D : public B, public C {
public:
    D() { a(); } // Not ambiguous. B::a() dominates A::a.
};

```

## Ambiguous conversions

Explicit and implicit conversions from pointers or references to class types can cause ambiguities. The next figure, Ambiguous Conversion of Pointers to Base Classes, shows the following:

- The declaration of an object of type `D`.
- The effect of applying the address-of operator (`&`) to that object. Note that the address-of operator always supplies the base address of the object.
- The effect of explicitly converting the pointer obtained using the address-of operator to the base-class type `A`. Note that coercing the address of the object to type `A*` does not always provide the compiler with enough information as to which subobject of type `A` to select; in this case, two subobjects exist.



Ambiguous Conversion of Pointers to Base Classes

The conversion to type `A*` (pointer to `A`) is ambiguous because there is no way to discern which subobject of type `A` is the correct one. Note that you can avoid the ambiguity by explicitly specifying which subobject you mean to use, as follows:

```

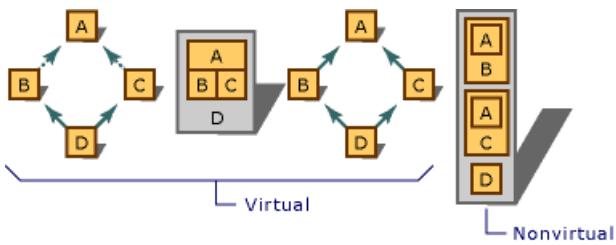
(A *)(B *)&d      // Use B subobject.
(A *)(C *)&d      // Use C subobject.

```

## Ambiguities and virtual base classes

If virtual base classes are used, functions, objects, types, and enumerators can be reached through multiple-inheritance paths. Because there is only one instance of the base class, there is no ambiguity when accessing these names.

The following figure shows how objects are composed using virtual and nonvirtual inheritance.



## Virtual vs. Nonvirtual Derivation

In the figure, accessing any member of class `A` through nonvirtual base classes causes an ambiguity; the compiler has no information that explains whether to use the subobject associated with `B` or the subobject associated with `C`. However, when `A` is specified as a virtual base class, there is no question which subobject is being accessed.

## See Also

[Inheritance](#)

# Explicit Overrides (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

If the same virtual function is declared in two or more [interfaces](#) and if a class is derived from these interfaces, you can explicitly override each virtual function.

For information on explicit overrides in managed code using the new managed syntax, see [Explicit Overrides](#).

## END Microsoft Specific

## Example

The following code example illustrates how to use explicit overrides:

```
// deriv_ExplicitOverrides.cpp
// compile with: /GR
extern "C" int printf_s(const char *, ...);

__interface IMyInt1 {
    void mf1();
    void mf1(int);
    void mf2();
    void mf2(int);
};

__interface IMyInt2 {
    void mf1();
    void mf1(int);
    void mf2();
    void mf2(int);
};

class CMyClass : public IMyInt1, public IMyInt2 {
public:
    void IMyInt1::mf1() {
        printf_s("In CMyClass::IMyInt1::mf1()\n");
    }

    void IMyInt1::mf1(int) {
        printf_s("In CMyClass::IMyInt1::mf1(int)\n");
    }

    void IMyInt1::mf2();
    void IMyInt1::mf2(int);

    void IMyInt2::mf1() {
        printf_s("In CMyClass::IMyInt2::mf1()\n");
    }

    void IMyInt2::mf1(int) {
        printf_s("In CMyClass::IMyInt2::mf1(int)\n");
    }

    void IMyInt2::mf2();
    void IMyInt2::mf2(int);
};

void CMyClass::IMyInt1::mf2() {
    printf_s("In CMyClass::IMyInt1::mf2()\n");
}
```

```
}

void CMyClass::IMyInt1::mf2(int) {
    printf_s("In CMyClass::IMyInt1::mf2(int)\n");
}

void CMyClass::IMyInt2::mf2() {
    printf_s("In CMyClass::IMyInt2::mf2()\n");
}

void CMyClass::IMyInt2::mf2(int) {
    printf_s("In CMyClass::IMyInt2::mf2(int)\n");
}

int main() {
    IMyInt1 *pIMyInt1 = new CMyClass();
    IMyInt2 *pIMyInt2 = dynamic_cast<IMyInt2 *>(pIMyInt1);

    pIMyInt1->mf1();
    pIMyInt1->mf1(1);
    pIMyInt1->mf2();
    pIMyInt1->mf2(2);
    pIMyInt2->mf1();
    pIMyInt2->mf1(3);
    pIMyInt2->mf2();
    pIMyInt2->mf2(4);

    // Cast to a CMyClass pointer so that the destructor gets called
    CMyClass *p = dynamic_cast<CMyClass *>(pIMyInt1);
    delete p;
}
```

```
In CMyClass::IMyInt1::mf1()
In CMyClass::IMyInt1::mf1(int)
In CMyClass::IMyInt1::mf2()
In CMyClass::IMyInt1::mf2(int)
In CMyClass::IMyInt2::mf1()
In CMyClass::IMyInt2::mf1(int)
In CMyClass::IMyInt2::mf2()
In CMyClass::IMyInt2::mf2(int)
```

## See Also

[Inheritance](#)

# Abstract Classes (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Abstract classes act as expressions of general concepts from which more specific classes can be derived. You cannot create an object of an abstract class type; however, you can use pointers and references to abstract class types.

A class that contains at least one pure virtual function is considered an abstract class. Classes derived from the abstract class must implement the pure virtual function or they, too, are abstract classes.

A virtual function is declared as "pure" by using the *pure-specifier* syntax (described in [Class Protocol Implementation](#)). Consider the example presented in [Virtual Functions](#). The intent of class `Account` is to provide general functionality, but objects of type `Account` are too general to be useful. Therefore, `Account` is a good candidate for an abstract class:

```
// deriv_AbstractClasses.cpp
// compile with: /LD
class Account {
public:
    Account( double d );    // Constructor.
    virtual double GetBalance();    // Obtain balance.
    virtual void PrintBalance() = 0;    // Pure virtual function.
private:
    double _balance;
};
```

The only difference between this declaration and the previous one is that `PrintBalance` is declared with the pure specifier (`= 0`).

## Restrictions on abstract classes

Abstract classes cannot be used for:

- Variables or member data
- Argument types
- Function return types
- Types of explicit conversions

Another restriction is that if the constructor for an abstract class calls a pure virtual function, either directly or indirectly, the result is undefined. However, constructors and destructors for abstract classes can call other member functions.

Pure virtual functions can be defined for abstract classes, but they can be called directly only by using the syntax:

`abstract-class-name :: function-name()`

This helps when designing class hierarchies whose base class(es) include pure virtual destructors, because base class destructors are always called in the process of destroying an object. Consider the following example:

```
// Declare an abstract base class with a pure virtual destructor.  
// deriv_RestrictionsonUsingAbstractClasses.cpp  
class base {  
public:  
    base() {}  
    virtual ~base()=0;  
};  
  
// Provide a definition for destructor.  
base::~base() {}  
  
class derived:public base {  
public:  
    derived() {}  
    ~derived(){}  
};  
  
int main() {  
    derived *pDerived = new derived;  
    delete pDerived;  
}
```

When the object pointed to by `pDerived` is deleted, the destructor for class `derived` is called and then the destructor for class `base` is called. The empty implementation for the pure virtual function ensures that at least some implementation exists for the function.

**NOTE**

In the preceding example, the pure virtual function `base::~base` is called implicitly from `derived::~derived`. It is also possible to call pure virtual functions explicitly using a fully qualified member-function name.

## See Also

[Inheritance](#)

# Summary of Scope Rules

12/21/2017 • 3 min to read • [Edit Online](#)

The use of a name must be unambiguous within its scope (up to the point where overloading is determined). If the name denotes a function, the function must be unambiguous with respect to number and type of parameters. If the name remains unambiguous, [member-access](#) rules are applied.

## Constructor initializers

Constructor initializers (described in [Initializing Bases and Members](#)) are evaluated in the scope of the outermost block of the constructor for which they are specified. Therefore, they can use the constructor's parameter names.

## Global names

A name of an object, function, or enumerator is global if it is introduced outside any function or class or prefixed by the global unary scope operator ( `::` ), and if it is not used in conjunction with any of these binary operators:

- Scope-resolution ( `::` )
- Member-selection for objects and references ( `.` )
- Member-selection for pointers ( `->` )

## Qualified names

Names used with the binary scope-resolution operator ( `::` ) are called "qualified names." The name specified after the binary scope-resolution operator must be a member of the class specified on the left of the operator or a member of its base class(es).

Names specified after the member-selection operator ( `.` or `->` ) must be members of the class type of the object specified on the left of the operator or members of its base class(es). Names specified on the right of the member-selection operator ( `->` ) can also be objects of another class type, provided that the left-hand side of `->` is a class object and that the class defines an overloaded member-selection operator ( `->` ) that evaluates to a pointer to some other class type. (This provision is discussed in more detail in [Class Member Access](#).)

The compiler searches for names in the following order, stopping when the name is found:

1. Current block scope if name is used inside a function; otherwise, global scope.
2. Outward through each enclosing block scope, including the outermost function scope (which includes function parameters).
3. If the name is used inside a member function, the class's scope is searched for the name.
4. The class's base classes are searched for the name.
5. The enclosing nested class scope (if any) and its bases are searched. The search continues until the outermost enclosing class scope is searched.
6. Global scope is searched.

However, you can make modifications to this search order as follows:

7. Names preceded by `::` force the search to begin at global scope.

8. Names preceded by the **class**, **struct**, and **union** keywords force the compiler to search only for **class**, **struct**, or **union** names.

9. Names on the left side of the scope-resolution operator (`::`) can be only **class**, **struct**, **namespace**, or **union** names.

If the name refers to a nonstatic member but is used in a static member function, an error message is generated. Similarly, if the name refers to any nonstatic member in an enclosing class, an error message is generated because enclosed classes do not have enclosing-class **this** pointers.

## Function parameter names

Function parameter names in function definitions are considered to be in the scope of the outermost block of the function. Therefore, they are local names and go out of scope when the function is exited.

Function parameter names in function declarations (prototypes) are in local scope of the declaration and go out of scope at the end of the declaration.

Default parameters are in the scope of the parameter for which they are the default, as described in the preceding two paragraphs. However, they cannot access local variables or nonstatic class members. Default parameters are evaluated at the point of the function call, but they are evaluated in the function declaration's original scope. Therefore, the default parameters for member functions are always evaluated in class scope.

## See Also

[Inheritance](#)

# Inheritance Keywords

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

```
class [_single_inheritance] class-name;
class [_multiple_inheritance] class-name;
class [_virtual_inheritance] class-name;
```

where:

*class-name*

The name of the class being declared.

C++ allows you to declare a pointer to a class member prior to the definition of the class. For example:

```
class S;
int S::*p;
```

In the code above, `p` is declared to be a pointer to integer member of class `S`. However, `class S` has not yet been defined in this code; it has only been declared. When the compiler encounters such a pointer, it must make a generalized representation of the pointer. The size of the representation is dependent on the inheritance model specified. There are four ways to specify an inheritance model to the compiler:

- In the IDE under **Pointer-to-member representation**
- At the command line using the `/vmg` switch
- Using the `pointers_to_members` pragma
- Using the inheritance keywords `_single_inheritance`, `_multiple_inheritance`, and `_virtual_inheritance`. This technique controls the inheritance model on a per-class basis.

### NOTE

If you always declare a pointer to a member of a class after defining the class, you don't need to use any of these options.

Declaring a pointer to a member of a class prior to the class definition affects the size and speed of the resulting executable file. The more complex the inheritance used by a class, the greater the number of bytes required to represent a pointer to a member of the class and the larger the code required to interpret the pointer. Single inheritance is least complex, and virtual inheritance is most complex.

If the example above is changed to:

```
class _single_inheritance S;
int S::*p;
```

regardless of command-line options or pragmas, pointers to members of `class S` will use the smallest possible representation.

**NOTE**

The same forward declaration of a class pointer-to-member representation should occur in every translation unit that declares pointers to members of that class, and the declaration should occur before the pointers to members are declared.

**END Microsoft Specific**

## See Also

[Keywords](#)

# virtual (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The `virtual` keyword declares a virtual function or a virtual base class.

## Syntax

```
virtual [type-specifiers] member-function-declarator  
virtual [access-specifier] base-class-name
```

### Parameters

`type-specifiers`

Specifies the return type of the virtual member function.

`member-function-declarator`

Declares a member function.

`access-specifier`

Defines the level of access to the base class, `public`, `protected` or `private`. Can appear before or after the `virtual` keyword.

`base-class-name`

Identifies a previously declared class type.

## Remarks

See [Virtual Functions](#) for more information.

Also see the following keywords: `class`, `private`, `public`, and `protected`.

## See Also

[Keywords](#)

# \_\_super

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Allows you to explicitly state that you are calling a base-class implementation for a function that you are overriding.

## Syntax

```
__super::  
member_function  
();
```

## Remarks

All accessible base-class methods are considered during the overload resolution phase, and the function that provides the best match is the one that is called.

`__super` can only appear within the body of a member function.

`__super` cannot be used with a using declaration. See [using Declaration](#) for more information.

With the introduction of [attributes](#) that inject code, your code might contain one or more base classes whose names you may not know but that contain methods that you wish to call.

## Example

```
// deriv_super.cpp  
// compile with: /c  
struct B1 {  
    void mf(int) {}  
};  
  
struct B2 {  
    void mf(short) {}  
  
    void mf(char) {}  
};  
  
struct D : B1, B2 {  
    void mf(short) {  
        __super::mf(1); // Calls B1::mf(int)  
        __super::mf('s'); // Calls B2::mf(char)  
    }  
};
```

**END Microsoft Specific**

## See Also

[Keywords](#)

# `__interface`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A Visual C++ interface can be defined as follows:

- Can inherit from zero or more base interfaces.
- Cannot inherit from a base class.
- Can only contain public, pure virtual methods.
- Cannot contain constructors, destructors, or operators.
- Cannot contain static methods.
- Cannot contain data members; properties are allowed.

## Syntax

```
modifier  
__interface interface-name {interface-definition};
```

## Remarks

A C++ [class](#) or [struct](#) could be implemented with these rules, but `__interface` enforces them.

For example, the following is a sample interface definition:

```
__interface IMyInterface {  
    HRESULT CommitX();  
    HRESULT get_X(BSTR* pbstrName);  
};
```

For information on managed interfaces, see [interface class](#).

Notice that you do not have to explicitly indicate that the `CommitX` and `get_X` functions are pure virtual. An equivalent declaration for the first function would be:

```
virtual HRESULT CommitX() = 0;
```

`__interface` implies the [novertable](#) `__declspec` modifier.

## Example

The following sample shows how to use properties declared in an interface.

```
// deriv_interface.cpp  
#define _ATL_ATTRIBUTES 1  
#include <atlbase.h>  
#include <atlcom.h>
```

```

#include <string.h>
#include <comdef.h>
#include <stdio.h>

[module(name="test")];

[ object, uuid("00000000-0000-0000-0000-000000000001"), library_block ]
__interface IFace {
    [ id(0) ] int int_data;
    [ id(5) ] BSTR bstr_data;
};

[ coclass, uuid("00000000-0000-0000-0000-000000000002") ]
class MyClass : public IFace {
private:
    int m_i;
    BSTR m_bstr;

public:
    MyClass()
    {
        m_i = 0;
        m_bstr = 0;
    }

    ~MyClass()
    {
        if (m_bstr)
            ::SysFreeString(m_bstr);
    }

    int get_int_data()
    {
        return m_i;
    }

    void put_int_data(int _i)
    {
        m_i = _i;
    }

    BSTR get_bstr_data()
    {
        BSTR bstr = ::SysAllocString(m_bstr);
        return bstr;
    }

    void put_bstr_data(BSTR bstr)
    {
        if (m_bstr)
            ::SysFreeString(m_bstr);
        m_bstr = ::SysAllocString(bstr);
    }
};

int main()
{
    _bstr_t bstr("Testing");
    CoInitialize(NULL);
    CComObject<MyClass>* p;
    CComObject<MyClass>::CreateInstance(&p);
    p->int_data = 100;
    printf_s("p->int_data = %d\n", p->int_data);
    p->bstr_data = bstr;
    printf_s("bstr_data = %S\n", p->bstr_data);
}

```

```
p->int_data = 100  
bstr_data = Testing
```

## END Microsoft Specific

## See Also

[Keywords](#)

[Interface Attributes](#)

# Special member functions

12/21/2017 • 1 min to read • [Edit Online](#)

The *special member functions* are class (or struct) member functions that, in certain cases, the compiler automatically generates for you. These functions are the [default constructor](#), the [destructor](#), the [copy constructor](#) and [copy assignment operator](#), and the [move constructor](#) and [move assignment operator](#). If your class does not define one or more of the special member functions, then the compiler may implicitly declare and define the functions that are used. The compiler-generated implementations are called the *default* special member functions. The compiler does not generate functions if they are not needed.

You can explicitly declare a default special member function by using the `= default` keyword. This causes the compiler to define the function only if needed, in the same way as if the function was not declared at all.

In some cases, the compiler may generate *deleted* special member functions, which are not defined and therefore not callable. This can happen in cases where a call to a particular special member function on a class doesn't make sense, given other properties of the class. To explicitly prevent automatic generation of a special member function, you can declare it as deleted by using the `= delete` keyword.

The compiler generates a *default constructor*, a constructor that takes no arguments, only when you have not declared any other constructor. If you have declared only a constructor that takes parameters, code that attempts to call a default constructor causes the compiler to produce an error message. The compiler-generated default constructor performs simple member-wise [default initialization](#) of the object. Default initialization leaves all member variables in an indeterminate state.

The default destructor performs member-wise destruction of the object. It is virtual only if a base class destructor is virtual.

The default copy and move construction and assignment operations perform member-wise bit-pattern copies or moves of non-static data members. Move operations are only generated when no destructor or move or copy operations are declared. A default copy constructor is only generated when no copy constructor is declared. It is implicitly deleted if a move operation is declared. A default copy assignment operator is generated only when no copy assignment operator is explicitly declared. It is implicitly deleted if a move operation is declared.

## See Also

[C++ Language Reference](#)

# Static Members (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

Classes can contain static member data and member functions. When a data member is declared as **static**, only one copy of the data is maintained for all objects of the class.

Static data members are not part of objects of a given class type. As a result, the declaration of a static data member is not considered a definition. The data member is declared in class scope, but definition is performed at file scope. These static members have external linkage. The following example illustrates this:

```
// static_data_members.cpp
class BufferedOutput
{
public:
    // Return number of bytes written by any object of this class.
    short BytesWritten()
    {
        return bytecount;
    }

    // Reset the counter.
    static void ResetCount()
    {
        bytecount = 0;
    }

    // Static member declaration.
    static long bytecount;
};

// Define bytecount in file scope.
long BufferedOutput::bytecount;

int main()
{
```

In the preceding code, the member `bytecount` is declared in class `BufferedOutput`, but it must be defined outside the class declaration.

Static data members can be referred to without referring to an object of class type. The number of bytes written using `BufferedOutput` objects can be obtained as follows:

```
long nBytes = BufferedOutput::bytecount;
```

For the static member to exist, it is not necessary that any objects of the class type exist. Static members can also be accessed using the member-selection (`.` and `->`) operators. For example:

```
BufferedOutput Console;

long nBytes = Console.bytecount;
```

In the preceding case, the reference to the object (`Console`) is not evaluated; the value returned is that of the static object `bytecount`.

Static data members are subject to class-member access rules, so private access to static data members is allowed only for class-member functions and friends. These rules are described in [Member-Access Control](#). The exception is that static data members must be defined in file scope regardless of their access restrictions. If the data member is to be explicitly initialized, an initializer must be provided with the definition.

The type of a static member is not qualified by its class name. Therefore, the type of `BufferedOutput::bytecount` is `long`.

## See Also

[Classes and Structs](#)

# User-Defined Type Conversions (C++)

12/21/2017 • 10 min to read • [Edit Online](#)

A *conversion* produces a new value of some type from a value of a different type. *Standard conversions* are built into the C++ language and support its built-in types, and you can create *user-defined conversions* to perform conversions to, from, or between user-defined types.

The standard conversions perform conversions between built-in types, between pointers or references to types related by inheritance, to and from void pointers, and to the null pointer. For more information, see [Standard Conversions](#). User-defined conversions perform conversions between user-defined types, or between user-defined types and built-in types. You can implement them as [Conversion constructors](#) or as [Conversion functions](#).

Conversions can either be explicit—when the programmer calls for one type to be converted to another, as in a cast or direct initialization—or implicit—when the language or program calls for a different type than the one given by the programmer.

Implicit conversions are attempted when:

- An argument supplied to a function does not have the same type as the matching parameter.
- The value returned from a function does not have the same type as the function return type.
- An initializer expression does not have the same type as the object it is initializing.
- An expression that controls a conditional statement, looping construct, or switch does not have the result type that's required to control it.
- An operand supplied to an operator does not have the same type as the matching operand-parameter. For built-in operators, both operands must have the same type, and are converted to a common type that can represent both. For more information, see [Standard Conversions](#). For user-defined operators, each operand must have the same type as the matching operand-parameter.

When one standard conversion can't complete an implicit conversion, the compiler can use a user-defined conversion, followed optionally by an additional standard conversion, to complete it.

When two or more user-defined conversions that perform the same conversion are available at a conversion site, the conversion is said to be ambiguous. Such ambiguities are an error because the compiler can't determine which one of the available conversions it should choose. However, it's not an error just to define multiple ways of performing the same conversion because the set of available conversions can be different at different locations in the source code—for example, depending on which header files are included in a source file. As long as only one conversion is available at the conversion site, there is no ambiguity. There are several ways that ambiguous conversions can arise, but the most common ones are:

- Multiple inheritance. The conversion is defined in more than one base class.
- Ambiguous function call. The conversion is defined as a conversion constructor of the target type and as a conversion function of the source type. For more information, see [Conversion functions](#).

You can usually resolve an ambiguity just by qualifying the name of the involved type more fully or by performing an explicit cast to clarify your intent.

Both conversion constructors and conversion functions obey member-access control rules, but the accessibility of the conversions is only considered if and when an unambiguous conversion can be determined. This means that a conversion can be ambiguous even if the access level of a competing

conversion would prevent it from being used. For more information about member accessibility, see [Member Access Control](#).

## The explicit keyword and problems with implicit conversion

By default when you create a user-defined conversion, the compiler can use it to perform implicit conversions. Sometimes this is what you want, but other times the simple rules that guide the compiler in making implicit conversions can lead it to accept code that you don't want it to.

One well-known example of an implicit conversion that can cause problems is the conversion to `bool`. There are many reasons that you might want to create a class type that can be used in a Boolean context—for example, so that it can be used to control an `if` statement or loop—but when the compiler performs a user-defined conversion to a built-in type, the compiler is allowed to apply an additional standard conversion afterwards. The intent of this additional standard conversion is to allow for things like promotion from `short` to `int`, but it also opens the door for less-obvious conversions—for example, from `bool` to `int`, which allows your class type to be used in integer contexts you never intended. This particular problem is known as the *Safe Bool Problem*. This kind of problem is where the `explicit` keyword can help.

The `explicit` keyword tells the compiler that the specified conversion can't be used to perform implicit conversions. If you wanted the syntactic convenience of implicit conversions before the `explicit` keyword was introduced, you had to either accept the unintended consequences that implicit conversion sometimes created or use less-convenient, named conversion functions as a workaround. Now, by using the `explicit` keyword, you can create convenient conversions that can only be used to perform explicit casts or direct initialization, and that won't lead to the kind of problems exemplified by the Safe Bool Problem.

The `explicit` keyword can be applied to conversion constructors since C++98, and to conversion functions since C++11. The following sections contain more information about how to use the `explicit` keyword.

## Conversion constructors

Conversion constructors define conversions from user-defined or built-in types to a user-defined type. The following example demonstrates a conversion constructor that converts from the built-in type `double` to a user-defined type `Money`.

```

#include <iostream>

class Money
{
public:
    Money() : amount{ 0.0 } {};
    Money(double _amount) : amount{ _amount } {};

    double amount;
};

void display_balance(const Money balance)
{
    std::cout << "The balance is: " << balance.amount << std::endl;
}

int main(int argc, char* argv[])
{
    Money payable{ 79.99 };

    display_balance(payable);
    display_balance(49.95);
    display_balance(9.99f);

    return 0;
}

```

Notice that the first call to the function `display_balance`, which takes an argument of type `Money`, doesn't require a conversion because its argument is the correct type. However, on the second call to `display_balance`, a conversion is needed because the type of the argument, a `double` with a value of `49.95`, is not what the function expects. The function can't use this value directly, but because there's a conversion from the type of the argument—`double`—to the type of the matching parameter—`Money`—a temporary value of type `Money` is constructed from the argument and used to complete the function call. In the third call to `display_balance`, notice that the argument is not a `double`, but is instead a `float` with a value of `9.99`—and yet the function call can still be completed because the compiler can perform a standard conversion—in this case, from `float` to `double`—and then perform the user-defined conversion from `double` to `Money` to complete the necessary conversion.

## Declaring conversion constructors

The following rules apply to declaring a conversion constructor:

- The target type of the conversion is the user-defined type that's being constructed.
- Conversion constructors typically take exactly one argument, which is of the source type. However, a conversion constructor can specify additional parameters if each additional parameter has a default value. The source type remains the type of the first parameter.
- Conversion constructors, like all constructors, do not specify a return type. Specifying a return type in the declaration is an error.
- Conversion constructors can be explicit.

## Explicit conversion constructors

By declaring a conversion constructor to be `explicit`, it can only be used to perform direct initialization of an object or to perform an explicit cast. This prevents functions that accept an argument of the class type from also implicitly accepting arguments of the conversion constructor's source type, and prevents the class type from being copy-initialized from a value of the source type. The following example demonstrates how to define an explicit conversion constructor, and the effect it has on what code is well-formed.

```

#include <iostream>

class Money
{
public:
    Money() : amount{ 0.0 } {};
    explicit Money(double _amount) : amount{ _amount } {};

    double amount;
};

void display_balance(const Money balance)
{
    std::cout << "The balance is: " << balance.amount << std::endl;
}

int main(int argc, char* argv[])
{
    Money payable{ 79.99 };           // Legal: direct initialization is explicit.

    display_balance(payable);         // Legal: no conversion required
    display_balance(49.95);          // Error: no suitable conversion exists to convert from double to Money.
    display_balance((Money)9.99f);    // Legal: explicit cast to Money

    return 0;
}

```

In this example, notice that you can still use the explicit conversion constructor to perform direct initialization of `payable`. If instead you were to copy-initialize `Money payable = 79.99;`, it would be an error. The first call to `display_balance` is unaffected because the argument is the correct type. The second call to `display_balance` is an error, because the conversion constructor can't be used to perform implicit conversions. The third call to `display_balance` is legal because of the explicit cast to `Money`, but notice that the compiler still helped complete the cast by inserting an implicit cast from `float` to `double`.

Although the convenience of allowing implicit conversions can be tempting, doing so can introduce hard-to-find bugs. The rule of thumb is to make all conversion constructors explicit except when you're sure that you want a specific conversion to occur implicitly.

## Conversion functions

Conversion functions define conversions from a user-defined type to other types. These functions are sometimes referred to as "cast operators" because they, along with conversion constructors, are called when a value is cast to a different type. The following example demonstrates a conversion function that converts from the user-defined type, `Money`, to a built-in type, `double`:

```

#include <iostream>

class Money
{
public:
    Money() : amount{ 0.0 } {};
    Money(double _amount) : amount{ _amount } {};

    operator double() const { return amount; }
private:
    double amount;
};

void display_balance(const Money balance)
{
    std::cout << "The balance is: " << balance << std::endl;
}

```

Notice that the member variable `amount` is made private and that a public conversion function to type `double` is introduced just to return the value of `amount`. In the function `display_balance`, an implicit conversion occurs when the value of `balance` is streamed to standard output by using the stream insertion operator `<<`. Because no stream-insertion operator is defined for the user-defined type `Money`, but there is one for built-in type `double`, the compiler can use the conversion function from `Money` to `double` to satisfy the stream-insertion operator.

Conversion functions are inherited by derived classes. Conversion functions in a derived class only override an inherited conversion function when they convert to exactly the same type. For example, a user-defined conversion function of the derived class `operator int` does not override—or even influence—a user-defined conversion function of the base class `operator short`, even though the standard conversions define a conversion relationship between `int` and `short`.

## Declaring conversion functions

The following rules apply to declaring a conversion function:

- The target type of the conversion must be declared prior to the declaration of the conversion function. Classes, structures, enumerations, and typedefs cannot be declared within the declaration of the conversion function.

```
operator struct String { char string_storage; }() // illegal
```

- Conversion functions take no arguments. Specifying any parameters in the declaration is an error.
- Conversion functions have a return type that is specified by the name of the conversion function, which is also the name of the conversion's target type. Specifying a return type in the declaration is an error.
- Conversion functions can be virtual.
- Conversion functions can be explicit.

## Explicit conversion functions

When a conversion function is declared to be explicit, it can only be used to perform an explicit cast. This prevents functions that accept an argument of the conversion function's target type from also implicitly accepting arguments of the class type, and prevents instances of the target type from being copy-initialized from a value of the class type. The following example demonstrates how to define an explicit conversion function and the effect it has on what code is well-formed.

```
#include <iostream>

class Money
{
public:
    Money() : amount{ 0.0 } {};
    Money(double _amount) : amount{ _amount } {};

    explicit operator double() const { return amount; }
private:
    double amount;
};

void display_balance(const Money balance)
{
    std::cout << "The balance is: " << (double)balance << std::endl;
}
```

Here the conversion function `operator double` has been made explicit, and an explicit cast to type `double` has been introduced in the function `display_balance` to perform the conversion. If this cast were omitted, the compiler would be unable to locate a suitable stream-insertion operator `<<` for type `Money` and an error would occur.

# Mutable Data Members (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This keyword can only be applied to non-static and non-const data members of a class. If a data member is declared `mutable`, then it is legal to assign a value to this data member from a **const** member function.

## Syntax

```
mutable member-variable-declaration;
```

## Remarks

For example, the following code will compile without error because `m_accessCount` has been declared to be `mutable`, and therefore can be modified by `GetFlag` even though `GetFlag` is a const member function.

```
// mutable.cpp
class X
{
public:
    bool GetFlag() const
    {
        m_accessCount++;
        return m_flag;
    }
private:
    bool m_flag;
    mutable int m_accessCount;
};

int main()
{}
```

## See Also

[Keywords](#)

# Nested Class Declarations

12/21/2017 • 4 min to read • [Edit Online](#)

A class can be declared within the scope of another class. Such a class is called a "nested class." Nested classes are considered to be within the scope of the enclosing class and are available for use within that scope. To refer to a nested class from a scope other than its immediate enclosing scope, you must use a fully qualified name.

The following example shows how to declare nested classes:

```
// nested_class_declarations.cpp
class BufferedIO
{
public:
    enum IOError { None, Access, General };

    // Declare nested class BufferedInput.
    class BufferedInput
    {
public:
    int read();
    int good()
    {
        return _inputerror == None;
    }
private:
    IOError _inputerror;
};

    // Declare nested class BufferedOutput.
    class BufferedOutput
    {
        // Member list
    };
};

int main()
{}
```

`BufferedIO::BufferedInput` and `BufferedIO::BufferedOutput` are declared within `BufferedIO`. These class names are not visible outside the scope of class `BufferedIO`. However, an object of type `BufferedIO` does not contain any objects of types `BufferedInput` or `BufferedOutput`.

Nested classes can directly use names, type names, names of static members, and enumerators only from the enclosing class. To use names of other class members, you must use pointers, references, or object names.

In the preceding `BufferedIO` example, the enumeration `IOError` can be accessed directly by member functions in the nested classes, `BufferedIO::BufferedInput` or `BufferedIO::BufferedOutput`, as shown in function `good`.

## NOTE

Nested classes declare only types within class scope. They do not cause contained objects of the nested class to be created. The preceding example declares two nested classes but does not declare any objects of these class types.

An exception to the scope visibility of a nested class declaration is when a type name is declared together with a forward declaration. In this case, the class name declared by the forward declaration is visible outside the enclosing

class, with its scope defined to be the smallest enclosing non-class scope. For example:

```
// nested_class_declarations_2.cpp
class C
{
public:
    typedef class U u_t; // class U visible outside class C scope
    typedef class V {} v_t; // class V not visible outside class C
};

int main()
{
    // okay, forward declaration used above so file scope is used
    U* pu;

    // error, type name only exists in class C scope
    u_t* pu2; // C2065

    // error, class defined above so class C scope
    V* pv; // C2065

    // okay, fully qualified name
    C::V* pv2;
}
```

## Access privilege in nested classes

Nesting a class within another class does not give special access privileges to member functions of the nested class. Similarly, member functions of the enclosing class have no special access to members of the nested class.

## Member functions in nested classes

Member functions declared in nested classes can be defined in file scope. The preceding example could have been written:

```

// member_functions_in_nested_classes.cpp
class BufferedIO
{
public:
    enum IOError { None, Access, General };
    class BufferedInput
    {
public:
    int read(); // Declare but do not define member
    int good(); // functions read and good.
private:
    IOError _inputerror;
};

class BufferedOutput
{
    // Member list.
};

};

// Define member functions read and good in
// file scope.
int BufferedIO::BufferedInput::read()
{
    return(1);
}

int BufferedIO::BufferedInput::good()
{
    return _inputerror == None;
}

int main()
{
}

```

In the preceding example, the *qualified-type-name* syntax is used to declare the function name. The declaration:

```
BufferedIO::BufferedInput::read()
```

means "the `read` function that is a member of the `BufferedInput` class that is in the scope of the `BufferedIO` class." Because this declaration uses the *qualified-type-name* syntax, constructs of the following form are possible:

```

typedef BufferedIO::BufferedInput BIO_INPUT;
int BIO_INPUT::read()
```

The preceding declaration is equivalent to the previous one, but it uses a `typedef` name in place of the class names.

## Friend functions in nested classes

Friend functions declared in a nested class are considered to be in the scope of the nested class, not the enclosing class. Therefore, the friend functions gain no special access privileges to members or member functions of the enclosing class. If you want to use a name that is declared in a nested class in a friend function and the friend function is defined in file scope, use qualified type names as follows:

```

// friend_functions_and_nested_classes.cpp

#include <string.h>

enum
{
    sizeOfMessage = 255
};

char *rgszMessage[sizeOfMessage];

class BufferedIO
{
public:
    class BufferedInput
    {
public:
        friend int GetExtendedErrorStatus();
        static char *message;
        static int messageSize;
        int iMsgNo;
    };
};

char *BufferedIO::BufferedInput::message;
int BufferedIO::BufferedInput::messageSize;

int GetExtendedErrorStatus()
{
    int iMsgNo = 1; // assign arbitrary value as message number

    strcpy_s( BufferedIO::BufferedInput::message,
              BufferedIO::BufferedInput::messageSize,
              rgszMessage[iMsgNo] );

    return iMsgNo;
}

int main()
{
}

```

The following code shows the function `GetExtendedErrorStatus` declared as a friend function. In the function, which is defined in file scope, a message is copied from a static array into a class member. Note that a better implementation of `GetExtendedErrorStatus` is to declare it as:

```
int GetExtendedErrorStatus( char *message )
```

With the preceding interface, several classes can use the services of this function by passing a memory location where they want the error message copied.

## See Also

[Classes and Structs](#)

# Anonymous Class Types

12/21/2017 • 1 min to read • [Edit Online](#)

Classes can be anonymous — that is, they can be declared without an *identifier*. This is useful when you replace a class name with a `typedef` name, as in the following:

```
typedef struct
{
    unsigned x;
    unsigned y;
} POINT;
```

## NOTE

The use of anonymous classes shown in the previous example is useful for preserving compatibility with existing C code. In some C code, the use of `typedef` in conjunction with anonymous structures is prevalent.

Anonymous classes are also useful when you want a reference to a class member to appear as though it were not contained in a separate class, as in the following:

```
struct PTValue
{
    POINT ptLoc;
    union
    {
        int iValue;
        long lValue;
    };
};

PTValue ptv;
```

In the preceding code, `iValue` can be accessed using the object member-selection operator (`.`) as follows:

```
int i = ptv.iValue;
```

Anonymous classes are subject to certain restrictions. (For more information about anonymous unions, see [Unions](#).) Anonymous classes:

- Cannot have a constructor or destructor.
- Cannot be passed as arguments to functions (unless type checking is defeated using ellipses).
- Cannot be returned as return values from functions.

## Anonymous structs

### Microsoft Specific

A Microsoft C extension allows you to declare a structure variable within another structure without giving it a name. These nested structures are called anonymous structures. C++ does not allow anonymous structures.

You can access the members of an anonymous structure as if they were members in the containing structure.

```
// anonymous_structures.c
#include <stdio.h>

struct phone
{
    int areacode;
    long number;
};

struct person
{
    char name[30];
    char gender;
    int age;
    int weight;
    struct phone; // Anonymous structure; no name needed
} Jim;

int main()
{
    Jim.number = 1234567;
    printf_s("%d\n", Jim.number);
}
//Output: 1234567
```

**END Microsoft Specific**

# Pointers to Members

12/21/2017 • 3 min to read • [Edit Online](#)

Declarations of pointers to members are special cases of pointer declarations. They are declared using the following sequence:

```
[storage-class-specifiers] [cv-qualifiers] type-specifiers [ms-modifier]qualified-name ::* [cv-qualifiers]  
identifier  
[= & qualified-name :: member-name];
```

## 1. The declaration specifier:

- An optional storage class specifier.
- Optional **const** and/or **volatile** specifiers.
- The type specifier: the name of a type. This is the type of the member to be pointed to, not the class.

## 2. The declarator:

- An optional Microsoft specific modifier. For more information, see [Microsoft-Specific Modifiers](#).
- The qualified name of the class containing the members to be pointed to.
- The **::** operator.
- The **\*** operator.
- Optional **const** and/or **volatile** specifiers.
- The identifier naming the pointer to member.
- An optional initializer:

The **=** operator.

The **&** operator.

The qualified name of the class.

The **::** operator.

The name of a nonstatic member of the class of the appropriate type.

As always, multiple declarators (and any associated initializers) are allowed in a single declaration.

A pointer to a member of a class differs from a normal pointer because it has type information for the type of the member and for the class to which the member belongs. A normal pointer identifies (has the address of) only a single object in memory. A pointer to a member of a class identifies that member in any instance of the class. The following example declares a class, **Window**, and some pointers to member data.

```

// pointers_to_members1.cpp
class Window
{
public:
    Window();                                // Default constructor.
    Window( int x1, int y1,                  // Constructor specifying
            int x2, int y2 );                // window size.
    bool SetCaption( const char *szTitle ); // Set window caption.
    const char *GetCaption();                // Get window caption.
    char *szWinCaption;                     // Window caption.
};

// Declare a pointer to the data member szWinCaption.
char * Window::* pwCaption = &Window::szWinCaption;
int main()
{
}

```

In the preceding example, `pwCaption` is a pointer to any member of class `Window` that has type `char`. The type of `pwCaption` is `char * Window::*`. The next code fragment declares pointers to the `SetCaption` and `GetCaption` member functions.

```

const char * (Window::*pfnwGC)() = &Window::GetCaption;
bool (Window::*pfnwSC)( const char * ) = &Window::SetCaption;

```

The pointers `pfnwGC` and `pfnwSC` point to `GetCaption` and `SetCaption` of the `Window` class, respectively. The code copies information to the window caption directly using the pointer to member `pwCaption`:

```

Window wMainWindow;
Window *pwChildWindow = new Window;
char *szUntitled = "Untitled - ";
int cUntitledLen = strlen( szUntitled );

strcpy_s( wMainWindow.*pwCaption, cUntitledLen, szUntitled );
(wMainWindow.*pwCaption)[cUntitledLen - 1] = '1';      //same as
//wMainWindow.SzWinCaption [cUntitledLen - 1] = '1';
strcpy_s( pwChildWindow->*pwCaption, cUntitledLen, szUntitled );
(pwChildWindow->*pwCaption)[cUntitledLen - 1] = '2'; //same as //pwChildWindow->szWinCaption[cUntitledLen - 1]
= '2';

```

The difference between the `.*` and `->*` operators (the pointer-to-member operators) is that the `.*` operator selects members given an object or object reference, while the `->*` operator selects members through a pointer. (For more about these operators, see [Expressions with Pointer-to-Member Operators](#).)

The result of the pointer-to-member operators is the type of the member — in this case, `char *`.

The following code fragment invokes the member functions `GetCaption` and `SetCaption` using pointers to members:

```
// Allocate a buffer.
enum {
    sizeOfBuffer = 100
};
char szCaptionBase[sizeOfBuffer];

// Copy the main window caption into the buffer
// and append " [View 1]".
strcpy_s( szCaptionBase, sizeOfBuffer, (wMainWindow.*pfnwGC)() );
strcat_s( szCaptionBase, sizeOfBuffer, " [View 1]" );
// Set the child window's caption.
(pwChildWindow->*pfnwSC)( szCaptionBase );
```

## Restrictions on Pointers to Members

The address of a static member is not a pointer to a member. It is a regular pointer to the one instance of the static member. Because only one instance of a static member exists for all objects of a given class, the ordinary address-of (**&**) and dereference (\*) operators can be used.

## Pointers to Members and Virtual Functions

Invoking a virtual function through a pointer-to-member function works as if the function had been called directly; the correct function is looked up in the v-table and invoked.

The key to virtual functions working, as always, is invoking them through a pointer to a base class. (For more information about virtual functions, see [Virtual Functions](#).)

The following code shows how to invoke a virtual function through a pointer-to-member function:

```
// virtual_functions.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

class Base
{
public:
virtual void Print();
};

void (Base ::* bfnPrint)() = &Base :: Print;
void Base :: Print()
{
cout << "Print function for class Base\n";
}

class Derived : public Base
{
public:
void Print(); // Print is still a virtual function.
};

void Derived :: Print()
{
cout << "Print function for class Derived\n";
}

int main()
{
    Base    *bPtr;
    Base    bObject;
Derived dObject;
bPtr = &bObject; // Set pointer to address of bObject.
(bPtr->*bfnPrint)();
bPtr = &dObject; // Set pointer to address of dObject.
(bPtr->*bfnPrint)();
}

//Output: Print function for class Base
Print function for class Derived
```

## See Also

# this Pointer

12/21/2017 • 3 min to read • [Edit Online](#)

The **this** pointer is a pointer accessible only within the nonstatic member functions of a **class**, **struct**, or **union** type. It points to the object for which the member function is called. Static member functions do not have a **this** pointer.

## Syntax

```
this  
this->member-identifier
```

## Remarks

An object's **this** pointer is not part of the object itself; it is not reflected in the result of a **sizeof** statement on the object. Instead, when a nonstatic member function is called for an object, the address of the object is passed by the compiler as a hidden argument to the function. For example, the following function call:

```
myDate.setMonth( 3 );
```

can be interpreted this way:

```
setMonth( &myDate, 3 );
```

The object's address is available from within the member function as the **this** pointer. Most uses of **this** are implicit. It is legal, though unnecessary, to explicitly use **this** when referring to members of the class. For example:

```
void Date::setMonth( int mn )  
{  
    month = mn;           // These three statements  
    this->month = mn;     // are equivalent  
    (*this).month = mn;  
}
```

The expression **\*this** is commonly used to return the current object from a member function:

```
return *this;
```

The **this** pointer is also used to guard against self-reference:

```
if (&Object != this) {  
// do not execute in cases of self-reference
```

#### **NOTE**

Because the **this** pointer is nonmodifiable, assignments to **this** are not allowed. Earlier implementations of C++ allowed assignments to **this**.

Occasionally, the **this** pointer is used directly — for example, to manipulate self-referential data structures, where the address of the current object is required.

## Example

```

// this_pointer.cpp
// compile with: /EHsc

#include <iostream>
#include <string.h>

using namespace std;

class Buf
{
public:
    Buf( char* szBuffer, size_t sizeOfBuffer );
    Buf& operator=( const Buf & );
    void Display() { cout << buffer << endl; }

private:
    char*   buffer;
    size_t  sizeOfBuffer;
};

Buf::Buf( char* szBuffer, size_t sizeOfBuffer )
{
    sizeOfBuffer++; // account for a NULL terminator

    buffer = new char[ sizeOfBuffer ];
    if (buffer)
    {
        strcpy_s( buffer, sizeOfBuffer, szBuffer );
        sizeOfBuffer = sizeOfBuffer;
    }
}

Buf& Buf::operator=( const Buf &otherbuf )
{
    if( &otherbuf != this )
    {
        if (buffer)
            delete [] buffer;

        sizeOfBuffer = strlen( otherbuf.buffer ) + 1;
        buffer = new char[sizeOfBuffer];
        strcpy_s( buffer, sizeOfBuffer, otherbuf.buffer );
    }
    return *this;
}

int main()
{
    Buf myBuf( "my buffer", 10 );
    Buf yourBuf( "your buffer", 12 );

    // Display 'my buffer'
    myBuf.Display();

    // assignment operator
    myBuf = yourBuf;

    // Display 'your buffer'
    myBuf.Display();
}

```

```

my buffer
your buffer

```

# Type of the `this` pointer

The `this` pointer's type can be modified in the function declaration by the `const` and `volatile` keywords. To declare a function as having the attributes of one or more of these keywords, add the keyword(s) after the function argument list.

Consider this example:

```
// type_of_this_pointer1.cpp
class Point
{
    unsigned X() const;
};

int main()
{
}
```

The preceding code declares a member function, `X`, in which the `this` pointer is treated as a `const` pointer to a `const` object. Combinations of *cv-mod-list* options can be used, but they always modify the object pointed to by `this`, not the `this` pointer itself. Therefore, the following declaration declares function `X`; the `this` pointer is a `const` pointer to a `const` object:

```
// type_of_this_pointer2.cpp
class Point
{
    unsigned X() const;
};

int main()
{
}
```

The type of `this` in a member function is described by the following syntax, where *cv-qualifier-list* is determined from the member functions declarator and can be `const` or `volatile` (or both), and *class-type* is the name of the class:

*[cv-qualifier-list] class-type \* const this*

In other words, `this` is always a `const` pointer; it cannot be reassigned. The `const` or `volatile` qualifiers used in the member function declaration apply to the class instance pointed to by `this` in the scope of that function.

The following table explains more about how these modifiers work.

## Semantics of this Modifiers

MODIFIER	MEANING
<code>const</code>	Cannot change member data; cannot invoke member functions that are not <code>const</code> .
<code>volatile</code>	Member data is loaded from memory each time it is accessed; disables certain optimizations.

It is an error to pass a `const` object to a member function that is not `const`. Similarly, it is an error to pass a `volatile` object to a member function that is not `volatile`.

Member functions declared as `const` cannot change member data — in such functions, the `this` pointer is a pointer to a `const` object.

**NOTE**

Constructors and destructors cannot be declared as **const** or **volatile**. They can, however, be invoked on **const** or **volatile** objects.

## See Also

[Keywords](#)

# C++ Bit Fields

12/21/2017 • 2 min to read • [Edit Online](#)

Classes and structures can contain members that occupy less storage than an integral type. These members are specified as bit fields. The syntax for bit-field *member-declarator* specification follows:

## Syntax

```
declarator : constant-expression
```

## Remarks

The (optional) `declarator` is the name by which the member is accessed in the program. It must be an integral type (including enumerated types). The *constant-expression* specifies the number of bits the member occupies in the structure. Anonymous bit fields — that is, bit-field members with no identifier — can be used for padding.

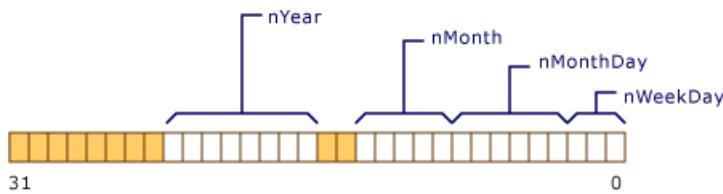
### NOTE

An unnamed bit field of width 0 forces alignment of the next bit field to the next `type` boundary, where `type` is the type of the member.

The following example declares a structure that contains bit fields:

```
// bit_fields1.cpp
// compile with: /LD
struct Date {
    unsigned short nWeekDay : 3;      // 0..7  (3 bits)
    unsigned short nMonthDay : 6;     // 0..31 (6 bits)
    unsigned short nMonth : 5;       // 0..12 (5 bits)
    unsigned short nYear : 8;        // 0..100 (8 bits)
};
```

The conceptual memory layout of an object of type `Date` is shown in the following figure.



Memory Layout of Date Object

Note that `nYear` is 8 bits long and would overflow the word boundary of the declared type, `unsigned short`. Therefore, it is begun at the beginning of a new `unsigned short`. It is not necessary that all bit fields fit in one object of the underlying type; new units of storage are allocated, according to the number of bits requested in the declaration.

## Microsoft Specific

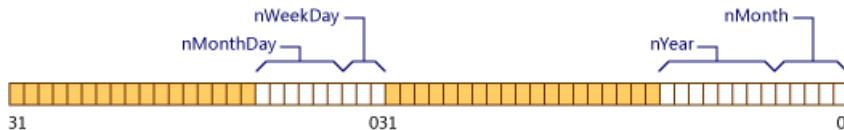
The ordering of data declared as bit fields is from low to high bit, as shown in the figure above.

## END Microsoft Specific

If the declaration of a structure includes an unnamed field of length 0, as shown in the following example,

```
// bit_fields2.cpp
// compile with: /LD
struct Date {
    unsigned nWeekDay : 3;      // 0..7  (3 bits)
    unsigned nMonthDay : 6;     // 0..31 (6 bits)
    unsigned         : 0;       // Force alignment to next boundary.
    unsigned nMonth : 5;       // 0..12 (5 bits)
    unsigned nYear  : 8;       // 0..100 (8 bits)
};
```

the memory layout is as shown in the following figure.



Layout of Date Object with Zero-Length Bit Field

The underlying type of a bit field must be an integral type, as described in [Fundamental Types](#).

## Restrictions on bit fields

The following list details erroneous operations on bit fields:

1. Taking the address of a bit field.
2. Initializing a reference with a bit field.

## See Also

[Classes and Structs](#)

# Lambda Expressions in C++

12/21/2017 • 12 min to read • [Edit Online](#)

In C++11 and later, a lambda expression—often called a *lambda*—is a convenient way of defining an anonymous function object (a *closure*) right at the location where it is invoked or passed as an argument to a function. Typically lambdas are used to encapsulate a few lines of code that are passed to algorithms or asynchronous methods. This article defines what lambdas are, compares them to other programming techniques, describes their advantages, and provides a basic example.

## Related Topics

- [Lambda expressions vs. function objects](#)
- [Working with lambda expressions](#)
- [constexpr lambda expressions](#)

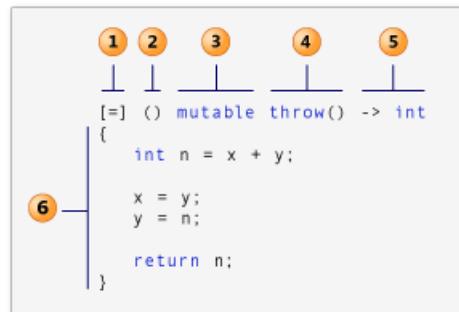
## Parts of a Lambda Expression

The ISO C++ Standard shows a simple lambda that is passed as the third argument to the `std::sort()` function:

```
#include <algorithm>
#include <cmath>

void abssort(float* x, unsigned n) {
    std::sort(x, x + n,
              // Lambda expression begins
              [] (float a, float b) {
                  return (std::abs(a) < std::abs(b));
              } // end of lambda expression
    );
}
```

This illustration shows the parts of a lambda:



1. *capture clause* (Also known as the *lambda-introducer* in the C++ specification.)
2. *parameter list* Optional. (Also known as the *lambda declarator*)
3. *mutable specification* Optional.
4. *exception-specification* Optional.
5. *trailing-return-type* Optional.
6. *lambda body*)

## Capture Clause

A lambda can introduce new variables in its body (in **C++14**), and it can also access, or *capture*, variables from the surrounding scope. A lambda begins with the capture clause (*lambda-introducer* in the Standard syntax), which specifies which variables are captured, and whether the capture is by value or by reference. Variables that have the ampersand (&) prefix are accessed by reference and variables that do not have it are accessed by value.

An empty capture clause, [ ], indicates that the body of the lambda expression accesses no variables in the enclosing scope.

You can use the default capture mode (*capture-default* in the Standard syntax) to indicate how to capture any outside variables that are referenced in the lambda: [&] means all variables that you refer to are captured by reference, and [=] means they are captured by value. You can use a default capture mode, and then specify the opposite mode explicitly for specific variables. For example, if a lambda body accesses the external variable `total` by reference and the external variable `factor` by value, then the following capture clauses are equivalent:

```
[&total, factor]
[factor, &total]
[&, factor]
[factor, &]
[=, &total]
[&total, =]
```

Only variables that are mentioned in the lambda are captured when a capture-default is used.

If a capture clause includes a capture-default &, then no `identifier` in a `capture` of that capture clause can have the form `& identifier`. Likewise, if the capture clause includes a capture-default =, then no `capture` of that capture clause can have the form `= identifier`. An identifier or `this` cannot appear more than once in a capture clause. The following code snippet illustrates some examples.

```
struct S { void f(int i); };

void S::f(int i) {
    [&, i]{};           // OK
    [&, &i]{};         // ERROR: i preceded by & when & is the default
    [=, this]{};        // ERROR: this when = is the default
    [=, *this]{ };     // OK: captures this by value. See below.
    [i, i]{};          // ERROR: i repeated
}
```

A capture followed by an ellipsis is a pack expansion, as shown in this [variadic template](#) example:

```
template<class... Args>
void f(Args... args) {
    auto x = [args...] { return g(args...); };
    x();
}
```

To use lambda expressions in the body of a class method, pass the `this` pointer to the capture clause to provide access to the methods and data members of the enclosing class.

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): The `this` pointer may be captured by value by specifying `*this` in the capture clause. Capture by value means that the entire *closure*, which is the anonymous function object that encapsulates the lambda expression, is copied to every call site where the lambda is invoked. Capture by value is useful when the lambda will execute in parallel or asynchronous operations, especially on certain hardware architectures such as NUMA.

For an example that shows how to use lambda expressions with class methods, see "Example: Using a Lambda Expression in a Method" in [Examples of Lambda Expressions](#).

When you use the capture clause, we recommend that you keep these points in mind, particularly when you use lambdas with multithreading:

- Reference captures can be used to modify variables outside, but value captures cannot. (`mutable` allows copies to be modified, but not originals.)
- Reference captures reflect updates to variables outside, but value captures do not.
- Reference captures introduce a lifetime dependency, but value captures have no lifetime dependencies. This is especially important when the lambda runs asynchronously. If you capture a local by reference in an `async` lambda, that local will very possibly be gone by the time the lambda runs, resulting in an access violation at run time.

### Generalized capture (C++ 14)

In C++ 14, you can introduce and initialize new variables in the capture clause, without the need to have those variables exist in the lambda function's enclosing scope. The initialization can be expressed as any arbitrary expression; the type of the new variable is deduced from the type produced by the expression. One benefit of this feature is that in C++ 14 you can capture move-only variables (such as `std::unique_ptr`) from the surrounding scope and use them in a lambda.

```
pNums = make_unique<vector<int>>(nums);
//...
auto a = [ptr = move(pNums)]()
{
    // use ptr
};
```

### Parameter List

In addition to capturing variables, a lambda can accept input parameters. A parameter list (*lambda declarator* in the Standard syntax) is optional and in most aspects resembles the parameter list for a function.

```
auto y = [] (int first, int second)
{
    return first + second;
};
```

In C++ 14, if the parameter type is generic, you can use the `auto` keyword as the type specifier. This tells the compiler to create the function call operator as a template. Each instance of `auto` in a parameter list is equivalent to a distinct type parameter.

```
auto y = [] (auto first, auto second)
{
    return first + second;
};
```

A lambda expression can take another lambda expression as its argument. For more information, see "Higher-Order Lambda Expressions" in the topic [Examples of Lambda Expressions](#).

Because a parameter list is optional, you can omit the empty parentheses if you do not pass arguments to the lambda expression and its lambda-declarator does not contain *exception-specification*, *trailing-return-type*, or `mutable`.

### Mutable Specification

Typically, a lambda's function call operator is const-by-value, but use of the `mutable` keyword cancels this out. It does not produce mutable data members. The mutable specification enables the body of a lambda expression to modify variables that are captured by value. Some of the examples later in this article show how to use `mutable`.

## Exception Specification

You can use the `noexcept` exception specification to indicate that the lambda expression does not throw any exceptions. As with ordinary functions, the Visual C++ compiler generates warning [C4297](#) if a lambda expression declares the `noexcept` exception specification and the lambda body throws an exception, as shown here:

```
// throw_lambda_expression.cpp
// compile with: /W4 /EHsc
int main() // C4297 expected
{
    []() noexcept { throw 5; }();
}
```

For more information, see [Exception Specifications \(throw\)](#).

## Return Type

The return type of a lambda expression is automatically deduced. You don't have to use the `auto` keyword unless you specify a *trailing-return-type*. The *trailing-return-type* resembles the return-type part of an ordinary method or function. However, the return type must follow the parameter list, and you must include the trailing-return-type keyword `->` before the return type.

You can omit the return-type part of a lambda expression if the lambda body contains just one return statement or the expression does not return a value. If the lambda body contains one return statement, the compiler deduces the return type from the type of the return expression. Otherwise, the compiler deduces the return type to be `void`. Consider the following example code snippets that illustrate this principle.

```
auto x1 = [](int i){ return i; }; // OK: return type is int
auto x2 = []{ return{ 1, 2 }; }; // ERROR: return type is void, deducing
                                // return type from braced-init-list is not valid
```

A lambda expression can produce another lambda expression as its return value. For more information, see "Higher-Order Lambda Expressions" in [Examples of Lambda Expressions](#).

## Lambda Body

The lambda body (*compound-statement* in the Standard syntax) of a lambda expression can contain anything that the body of an ordinary method or function can contain. The body of both an ordinary function and a lambda expression can access these kinds of variables:

- Captured variables from the enclosing scope, as described previously.
- Parameters
- Locally-declared variables
- Class data members, when declared inside a class and `this` is captured
- Any variable that has static storage duration—for example, global variables

The following example contains a lambda expression that explicitly captures the variable `n` by value and implicitly captures the variable `m` by reference:

```
// captures_lambda_expression.cpp
// compile with: /W4 /EHsc
#include <iostream>
using namespace std;

int main()
{
    int m = 0;
    int n = 0;
    [&, n] (int a) mutable { m = ++n + a; }(4);
    cout << m << endl << n << endl;
}
```

```
5
0
```

Because the variable `n` is captured by value, its value remains `0` after the call to the lambda expression. The `mutable` specification allows `n` to be modified within the lambda.

Although a lambda expression can only capture variables that have automatic storage duration, you can use variables that have static storage duration in the body of a lambda expression. The following example uses the `generate` function and a lambda expression to assign a value to each element in a `vector` object. The lambda expression modifies the static variable to generate the value of the next element.

```
void fillVector(vector<int>& v)
{
    // A local static variable.
    static int nextValue = 1;

    // The lambda expression that appears in the following call to
    // the generate function modifies and uses the local static
    // variable nextValue.
    generate(v.begin(), v.end(), [] { return nextValue++; });
    //WARNING: this is not thread-safe and is shown for illustration only
}
```

For more information, see [generate](#).

The following code example uses the function from the previous example, and adds an example of a lambda expression that uses the C++ Standard Library algorithm `generate_n`. This lambda expression assigns an element of a `vector` object to the sum of the previous two elements. The `mutable` keyword is used so that the body of the lambda expression can modify its copies of the external variables `x` and `y`, which the lambda expression captures by value. Because the lambda expression captures the original variables `x` and `y` by value, their values remain `1` after the lambda executes.

```
// compile with: /W4 /EHsc
#include <algorithm>
#include <iostream>
#include <vector>
#include <string>

using namespace std;

template <typename C> void print(const string& s, const C& c) {
    cout << s;

    for (const auto& e : c) {
        cout << e << " ";
    }
}
```

```

    cout << endl;
}

void fillVector(vector<int>& v)
{
    // A local static variable.
    static int nextValue = 1;

    // The lambda expression that appears in the following call to
    // the generate function modifies and uses the local static
    // variable nextValue.
    generate(v.begin(), v.end(), [] { return nextValue++; });
    //WARNING: this is not thread-safe and is shown for illustration only
}

int main()
{
    // The number of elements in the vector.
    const int elementCount = 9;

    // Create a vector object with each element set to 1.
    vector<int> v(elementCount, 1);

    // These variables hold the previous two elements of the vector.
    int x = 1;
    int y = 1;

    // Sets each element in the vector to the sum of the
    // previous two elements.
    generate_n(v.begin() + 2,
               elementCount - 2,
               [=]() mutable throw() -> int { // lambda is the 3rd parameter
                   // Generate current value.
                   int n = x + y;
                   // Update previous two values.
                   x = y;
                   y = n;
                   return n;
               });
    print("vector v after call to generate_n() with lambda: ", v);

    // Print the local variables x and y.
    // The values of x and y hold their initial values because
    // they are captured by value.
    cout << "x: " << x << " y: " << y << endl;

    // Fill the vector with a sequence of numbers
    fillVector(v);
    print("vector v after 1st call to fillVector(): ", v);
    // Fill the vector with the next sequence of numbers
    fillVector(v);
    print("vector v after 2nd call to fillVector(): ", v);
}

```

```

vector v after call to generate_n() with lambda: 1 1 2 3 5 8 13 21 34
x: 1 y: 1
vector v after 1st call to fillVector(): 1 2 3 4 5 6 7 8 9
vector v after 2nd call to fillVector(): 10 11 12 13 14 15 16 17 18

```

For more information, see [generate\\_n](#).

## constexpr lambda expressions

**Visual Studio 2017 version 15.3 and later** (available with `/std:c++17`): A lambda expression may be declared as

`constexpr` or used in a constant expression when the initialization of each data member that it captures or introduces is allowed within a constant expression.

```
int y = 32;
auto answer = [y]() constexpr
{
    int x = 10;
    return y + x;
};

constexpr int Increment(int n)
{
    return [n] { return n + 1; }();
}
```

A lambda is implicitly `constexpr` if its result satisfies the requirements of a `constexpr` function:

```
auto answer = [](int n)
{
    return 32 + n;
};

constexpr int response = answer(10);
```

If a lambda is implicitly or explicitly `constexpr`, conversion to a function pointer produces a `constexpr` function:

```
auto Increment = [](int n)
{
    return n + 1;
};

constexpr int(*inc)(int) = Increment;
```

## Microsoft-Specific

Lambdas are not supported in the following common language runtime (CLR) managed entities: `ref class`, `ref struct`, `value class`, or `value struct`.

If you are using a Microsoft-specific modifier such as `_declspec`, you can insert it into a lambda expression immediately after the `parameter-declaration-clause`—for example:

```
auto Sqr = [](int t) __declspec(code_seg("PagedMem")) -> int { return t*t; };
```

To determine whether a modifier is supported by lambdas, see the article about it in the [Microsoft-Specific Modifiers](#) section of the documentation.

In addition to C++11 Standard lambda functionality, Visual Studio supports stateless lambdas, which are omni-convertible to function pointers that use arbitrary calling conventions.

## See Also

[C++ Language Reference](#)

[Function Objects in the C++ Standard Library](#)

[Function Call](#)

[for\\_each](#)

# Lambda Expression Syntax

12/21/2017 • 4 min to read • [Edit Online](#)

This article demonstrates the syntax and structural elements of lambda expressions. For a description of lambda expressions, see [Lambda Expressions](#).

## Function Objects vs. Lambdas

When you write code, you probably use function pointers and function objects to solve problems and perform calculations, especially when you use [C++ Standard Library algorithms](#). Function pointers and function objects each have advantages and disadvantages—for example, function pointers have minimal syntactic overhead but do not retain state within a scope, and function objects can maintain state but require the syntactic overhead of a class definition.

A lambda combines the benefits of function pointers and function objects and avoids their disadvantages. Like a function objects, a lambda is flexible and can maintain state, but unlike a function object, its compact syntax doesn't require an explicit class definition. By using lambdas, you can write code that's less cumbersome and less prone to errors than the code for an equivalent function object.

The following examples compare the use of a lambda to the use of a function object. The first example uses a lambda to print to the console whether each element in a `vector` object is even or odd. The second example uses a function object to accomplish the same task.

## Example 1: Using a Lambda

This example passes a lambda to the `for_each` function. The lambda prints a result that states whether each element in a `vector` object is even or odd.

### Code

```

// even_lambda.cpp
// compile with: cl /EHsc /nologo /W4 /MTd
#include <algorithm>
#include <iostream>
#include <vector>
using namespace std;

int main()
{
    // Create a vector object that contains 10 elements.
    vector<int> v;
    for (int i = 1; i < 10; ++i) {
        v.push_back(i);
    }

    // Count the number of even numbers in the vector by
    // using the for_each function and a lambda.
    int evenCount = 0;
    for_each(v.begin(), v.end(), [&evenCount] (int n) {
        cout << n;
        if (n % 2 == 0) {
            cout << " is even " << endl;
            ++evenCount;
        } else {
            cout << " is odd " << endl;
        }
    });

    // Print the count of even numbers to the console.
    cout << "There are " << evenCount
        << " even numbers in the vector." << endl;
}

```

## Output

```

1 is odd
2 is even
3 is odd
4 is even
5 is odd
6 is even
7 is odd
8 is even
9 is odd
There are 4 even numbers in the vector.

```

## Comments

In the example, the third argument to the `for_each` function is a lambda. The `[&evenCount]` part specifies the capture clause of the expression, `(int n)` specifies the parameter list, and remaining part specifies the body of the expression.

## Example 2: Using a Function Object

Sometimes a lambda would be too unwieldy to extend much further than the previous example. The next example uses a function object instead of a lambda, together with the `for_each` function, to produce the same results as Example 1. Both examples store the count of even numbers in a `vector` object. To maintain the state of the operation, the `FunctorClass` class stores the `m_evenCount` variable by reference as a member variable. To perform the operation, `FunctorClass` implements the function-call operator, `operator()`. The Visual C++ compiler generates code that is comparable in size and performance to the lambda code in Example 1. For a basic problem like the one in this article, the simpler lambda design is probably better than the function-object design. However,

if you think that the functionality might require significant expansion in the future, then use a function object design so that code maintenance will be easier.

For more information about the `operator()`, see [Function Call](#). For more information about the `for_each` function, see [for\\_each](#).

## Code

```
// even_functor.cpp
// compile with: /EHsc
#include <algorithm>
#include <iostream>
#include <vector>
using namespace std;

class FunctorClass
{
public:
    // The required constructor for this example.
    explicit FunctorClass(int& evenCount)
        : m_evenCount(evenCount) { }

    // The function-call operator prints whether the number is
    // even or odd. If the number is even, this method updates
    // the counter.
    void operator()(int n) const {
        cout << n;

        if (n % 2 == 0) {
            cout << " is even " << endl;
            ++m_evenCount;
        } else {
            cout << " is odd " << endl;
        }
    }

private:
    // Default assignment operator to silence warning C4512.
    FunctorClass& operator=(const FunctorClass&);

    int& m_evenCount; // the number of even variables in the vector.
};

int main()
{
    // Create a vector object that contains 10 elements.
    vector<int> v;
    for (int i = 1; i < 10; ++i) {
        v.push_back(i);
    }

    // Count the number of even numbers in the vector by
    // using the for_each function and a function object.
    int evenCount = 0;
    for_each(v.begin(), v.end(), FunctorClass(evenCount));

    // Print the count of even numbers to the console.
    cout << "There are " << evenCount
        << " even numbers in the vector." << endl;
}
```

## Output

```
1 is odd
2 is even
3 is odd
4 is even
5 is odd
6 is even
7 is odd
8 is even
9 is odd
There are 4 even numbers in the vector.
```

## See Also

[Lambda Expressions](#)

[Examples of Lambda Expressions](#)

[generate](#)

[generate\\_n](#)

[for\\_each](#)

[Exception Specifications \(throw\)](#)

[Compiler Warning \(level 1\) C4297](#)

[Microsoft-Specific Modifiers](#)

# Examples of Lambda Expressions

12/21/2017 • 9 min to read • [Edit Online](#)

This article shows how to use lambda expressions in your programs. For an overview of lambda expressions, see [Lambda Expressions](#). For more information about the structure of a lambda expression, see [Lambda Expression Syntax](#).

## Declaring Lambda Expressions

### Example 1

Because a lambda expression is typed, you can assign it to an `auto` variable or to a `function` object, as shown here:

### Code

```
// declaring_lambda_expressions1.cpp
// compile with: /EHsc /W4
#include <functional>
#include <iostream>

int main()
{
    using namespace std;

    // Assign the lambda expression that adds two numbers to an auto variable.
    auto f1 = [](int x, int y) { return x + y; };

    cout << f1(2, 3) << endl;

    // Assign the same lambda expression to a function object.
    function<int(int, int)> f2 = [](int x, int y) { return x + y; };

    cout << f2(3, 4) << endl;
}
```

### Output

```
5
7
```

### Remarks

For more information, see `auto`, `function Class`, and `Function Call`.

Although lambda expressions are most often declared in the body of a function, you can declare them anywhere that you can initialize a variable.

### Example 2

The Visual C++ compiler binds a lambda expression to its captured variables when the expression is declared instead of when the expression is called. The following example shows a lambda expression that captures the local variable `i` by value and the local variable `j` by reference. Because the lambda expression captures `i` by value, the reassignment of `i` later in the program does not affect the result of the expression. However, because the lambda expression captures `j` by reference, the reassignment of `j` does affect the result of the expression.

## Code

```
// declaring_lambda_expressions2.cpp
// compile with: /EHsc /W4
#include <functional>
#include <iostream>

int main()
{
    using namespace std;

    int i = 3;
    int j = 5;

    // The following lambda expression captures i by value and
    // j by reference.
    function<int (void)> f = [i, &j] { return i + j; };

    // Change the values of i and j.
    i = 22;
    j = 44;

    // Call f and print its result.
    cout << f() << endl;
}
```

## Output

```
47
```

[\[In This Article\]](#)

## Calling Lambda Expressions

You can call a lambda expression immediately, as shown in the next code snippet. The second snippet shows how to pass a lambda as an argument to C++ Standard Library algorithms such as `find_if`.

### Example 1

This example declares a lambda expression that returns the sum of two integers and calls the expression immediately with the arguments `5` and `4`:

## Code

```
// calling_lambda_expressions1.cpp
// compile with: /EHsc
#include <iostream>

int main()
{
    using namespace std;
    int n = [] (int x, int y) { return x + y; }(5, 4);
    cout << n << endl;
}
```

## Output

```
9
```

### Example 2

This example passes a lambda expression as an argument to the `find_if` function. The lambda expression returns `true` if its parameter is an even number.

## Code

```
// calling_lambda_expressions2.cpp
// compile with: /EHsc /W4
#include <list>
#include <algorithm>
#include <iostream>

int main()
{
    using namespace std;

    // Create a list of integers with a few initial elements.
    list<int> numbers;
    numbers.push_back(13);
    numbers.push_back(17);
    numbers.push_back(42);
    numbers.push_back(46);
    numbers.push_back(99);

    // Use the find_if function and a lambda expression to find the
    // first even number in the list.
    const list<int>::const_iterator result =
        find_if(numbers.begin(), numbers.end(),[](int n) { return (n % 2) == 0; });

    // Print the result.
    if (result != numbers.end()) {
        cout << "The first even number in the list is " << *result << endl;
    } else {
        cout << "The list contains no even numbers." << endl;
    }
}
```

## Output

```
The first even number in the list is 42.
```

## Remarks

For more information about the `find_if` function, see [find\\_if](#). For more information about the C++ Standard Library functions that perform common algorithms, see [<algorithm>](#).

[\[In This Article\]](#)

## Nesting Lambda Expressions

### Example

You can nest a lambda expression inside another one, as shown in this example. The inner lambda expression multiplies its argument by 2 and returns the result. The outer lambda expression calls the inner lambda expression with its argument and adds 3 to the result.

## Code

```
// nesting_lambda_expressions.cpp
// compile with: /EHsc /W4
#include <iostream>

int main()
{
    using namespace std;

    // The following lambda expression contains a nested lambda
    // expression.
    int timestwoplusthree = [](int x) { return [](int y) { return y * 2; }(x) + 3; }(5);

    // Print the result.
    cout << timestwoplusthree << endl;
}
```

## Output

```
13
```

## Remarks

In this example, `[](int y) { return y * 2; }` is the nested lambda expression.

[[In This Article](#)]

# Higher-Order Lambda Functions

## Example

Many programming languages support the concept of a *higher-order function*. A higher-order function is a lambda expression that takes another lambda expression as its argument or returns a lambda expression. You can use the [function](#) class to enable a C++ lambda expression to behave like a higher-order function. The following example shows a lambda expression that returns a `function` object and a lambda expression that takes a `function` object as its argument.

## Code

```

// higher_order_lambda_expression.cpp
// compile with: /EHsc /W4
#include <iostream>
#include <functional>

int main()
{
    using namespace std;

    // The following code declares a lambda expression that returns
    // another lambda expression that adds two numbers.
    // The returned lambda expression captures parameter x by value.
    auto addtwointegers = [](int x) -> function<int(int)> {
        return [=](int y) { return x + y; };
    };

    // The following code declares a lambda expression that takes another
    // lambda expression as its argument.
    // The lambda expression applies the argument z to the function f
    // and multiplies by 2.
    auto higherorder = [](const function<int(int)>& f, int z) {
        return f(z) * 2;
    };

    // Call the lambda expression that is bound to higherorder.
    auto answer = higherorder(addtwointegers(7), 8);

    // Print the result, which is (7+8)*2.
    cout << answer << endl;
}

```

## Output

30

[\[In This Article\]](#)

## Using a Lambda Expression in a Function

### Example

You can use lambda expressions in the body of a function. The lambda expression can access any function or data member that the enclosing function can access. You can explicitly or implicitly capture the `this` pointer to provide access to functions and data members of the enclosing class.

**Visual Studio 2017 version 15.3 and later** (available with [/std:c++17](#)): Capture `this` by value (`[*this]`) when the lambda will be used in asynchronous or parallel operations where the code might execute after the original object goes out of scope.

You can use the `this` pointer explicitly in a function, as shown here:

```

// capture "this" by reference
void ApplyScale(const vector<int>& v) const
{
    for_each(v.begin(), v.end(),
        [this](int n) { cout << n * _scale << endl; });
}

// capture "this" by value (Visual Studio 2017 version 15.3 and later)
void ApplyScale2(const vector<int>& v) const
{
    for_each(v.begin(), v.end(),
        [*this](int n) { cout << n * _scale << endl; });
}

```

You can also capture the `this` pointer implicitly:

```

void ApplyScale(const vector<int>& v) const
{
    for_each(v.begin(), v.end(),
        [=](int n) { cout << n * _scale << endl; });
}

```

The following example shows the `scale` class, which encapsulates a scale value.

```

// function_lambda_expression.cpp
// compile with: /EHsc /W4
#include <algorithm>
#include <iostream>
#include <vector>

using namespace std;

class Scale
{
public:
    // The constructor.
    explicit Scale(int scale) : _scale(scale) {}

    // Prints the product of each element in a vector object
    // and the scale value to the console.
    void ApplyScale(const vector<int>& v) const
    {
        for_each(v.begin(), v.end(), [=](int n) { cout << n * _scale << endl; });
    }

private:
    int _scale;
};

int main()
{
    vector<int> values;
    values.push_back(1);
    values.push_back(2);
    values.push_back(3);
    values.push_back(4);

    // Create a Scale object that scales elements by 3 and apply
    // it to the vector object. Does not modify the vector.
    Scale s(3);
    s.ApplyScale(values);
}

```

## Output

```
3  
6  
9  
12
```

## Remarks

The `ApplyScale` function uses a lambda expression to print the product of the scale value and each element in a `vector` object. The lambda expression implicitly captures `this` so that it can access the `_scale` member.

[\[In This Article\]](#)

# Using Lambda Expressions with Templates

## Example

Because lambda expressions are typed, you can use them with C++ templates. The following example shows the `negate_all` and `print_all` functions. The `negate_all` function applies the unary `operator-` to each element in the `vector` object. The `print_all` function prints each element in the `vector` object to the console.

## Code

```
// template_lambda_expression.cpp  
// compile with: /EHsc  
#include <vector>  
#include <algorithm>  
#include <iostream>  
  
using namespace std;  
  
// Negates each element in the vector object. Assumes signed data type.  
template <typename T>  
void negate_all(vector<T>& v)  
{  
    for_each(v.begin(), v.end(), [](T& n) { n = -n; });  
}  
  
// Prints to the console each element in the vector object.  
template <typename T>  
void print_all(const vector<T>& v)  
{  
    for_each(v.begin(), v.end(), [](const T& n) { cout << n << endl; });  
}  
  
int main()  
{  
    // Create a vector of signed integers with a few elements.  
    vector<int> v;  
    v.push_back(34);  
    v.push_back(-43);  
    v.push_back(56);  
  
    print_all(v);  
    negate_all(v);  
    cout << "After negate_all(): " << endl;  
    print_all(v);  
}
```

## Output

```
34
-43
56
After negate_all():
-34
43
-56
```

## Remarks

For more information about C++ templates, see [Templates](#).

[\[In This Article\]](#)

# Handling Exceptions

## Example

The body of a lambda expression follows the rules for both structured exception handling (SEH) and C++ exception handling. You can handle a raised exception in the body of a lambda expression or defer exception handling to the enclosing scope. The following example uses the `for_each` function and a lambda expression to fill a `vector` object with the values of another one. It uses a `try / catch` block to handle invalid access to the first vector.

## Code

```
// eh_lambda_expression.cpp
// compile with: /EHsc /W4
#include <vector>
#include <algorithm>
#include <iostream>
using namespace std;

int main()
{
    // Create a vector that contains 3 elements.
    vector<int> elements(3);

    // Create another vector that contains index values.
    vector<int> indices(3);
    indices[0] = 0;
    indices[1] = -1; // This is not a valid subscript. It will trigger an exception.
    indices[2] = 2;

    // Use the values from the vector of index values to
    // fill the elements vector. This example uses a
    // try/catch block to handle invalid access to the
    // elements vector.
    try
    {
        for_each(indices.begin(), indices.end(), [&](int index) {
            elements.at(index) = index;
        });
    }
    catch (const out_of_range& e)
    {
        cerr << "Caught '" << e.what() << "'." << endl;
    };
}
```

## Output

```
Caught 'invalid vector<T> subscript'.
```

## Remarks

For more information about exception handling, see [Exception Handling](#).

[[In This Article](#)]

# Using Lambda Expressions with Managed Types (C++/CLI)

## Example

The capture clause of a lambda expression cannot contain a variable that has a managed type. However, you can pass an argument that has a managed type to the parameter list of a lambda expression. The following example contains a lambda expression that captures the local unmanaged variable `ch` by value and takes a `System.String` object as its parameter.

## Code

```
// managed_lambda_expression.cpp
// compile with: /clr
using namespace System;

int main()
{
    char ch = '!'; // a local unmanaged variable

    // The following lambda expression captures local variables
    // by value and takes a managed String object as its parameter.
    [=](String ^s) {
        Console::WriteLine(s + Convert::ToChar(ch));
    }("Hello");
}
```

## Output

```
Hello!
```

## Remarks

You can also use lambda expressions with the STL/CLR library. For more information, see [STL/CLR Library Reference](#).

### IMPORTANT

Lambdas are not supported in these common language runtime (CLR) managed entities: `ref class`, `ref struct`, `value class`, and `value struct`.

[[In This Article](#)]

## See Also

[Lambda Expressions](#)

[Lambda Expression Syntax](#)

`auto`

`function Class`

`find_if`

[`<algorithm>`](#)  
[`Function Call`](#)  
[`Templates`](#)  
[`Exception Handling`](#)  
[`STL/CLR Library Reference`](#)

# constexpr Lambda Expressions in C++

12/21/2017 • 1 min to read • [Edit Online](#)

**Visual Studio 2017 version 15.3 and later** (available with `/std:c++17`): A lambda expression may be declared as

`constexpr` or used in a constant expression when the initialization of each data member that it captures or introduces is allowed within a constant expression.

```
int y = 32;
auto answer = [y]() constexpr
{
    int x = 10;
    return y + x;
};

constexpr int Increment(int n)
{
    return [n] { return n + 1; }();
}
```

A lambda is implicitly `constexpr` if its result satisfies the requirements of a `constexpr` function:

```
auto answer = [](int n)
{
    return 32 + n;
};

constexpr int response = answer(10);
```

If a lambda is implicitly or explicitly `constexpr`, and you convert it to a function pointer, the resulting function is also `constexpr`:

```
auto Increment = [](int n)
{
    return n + 1;
};

constexpr int(*inc)(int) = Increment;
```

## See Also

[C++ Language Reference](#)

[Function Objects in the C++ Standard Library](#)

[Function Call](#)

[for\\_each](#)

# Arrays (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

An array is a collection of like objects. The simplest case of an array is a vector, which may be declared by the following sequence:

```
decl-specifier identifier [ constant-expression ]
decl-specifier identifier []
decl-specifier identifier [] [ constant-expression ] . .
decl-specifier identifier [ constant-expression ]
[ constant-expression ] . . .
```

1. The declaration specifier:

- An optional storage class specifier.
- Optional **const** and/or **volatile** specifiers.
- The type name of the elements of the array.

1. The declarator:

- The identifier.
- A constant expression of integral type enclosed in brackets, **[]**. If multiple dimensions are declared using additional brackets, the constant expression may be omitted on the first set of brackets.
- Optional additional brackets enclosing constant expressions.

1. An optional initializer. See [Initializers](#).

The number of elements in the array is given by the constant expression. The first element in the array is the 0th element, and the last element is the  $(n-1)$  element, where  $n$  is the number of elements the array can contain. The *constant-expression* must be of an integral type and must be greater than 0. A zero-sized array is legal only when the array is the last field in a **struct** or **union** and when the Microsoft extensions (/Ze) are enabled.

The following example shows how to define an array at run time:

```

// arrays.cpp
// compile with: /EHsc
#include <iostream>

int main() {
    using namespace std;
    int size = 3, i = 0;

    int* myarr = new int[size];

    for (i = 0 ; i < size ; i++)
        myarr[i] = 10;

    for (i = 0 ; i < size ; i++)
        printf_s("myarr[%d] = %d\n", i, myarr[i]);

    delete [] myarr;
}

```

Arrays are derived types and can therefore be constructed from any other derived or fundamental type except functions, references, and `void`.

Arrays constructed from other arrays are multidimensional arrays. These multidimensional arrays are specified by placing multiple bracketed constant expressions in sequence. For example, consider this declaration:

```
int i2[5][7];
```

It specifies an array of type `int`, conceptually arranged in a two-dimensional matrix of five rows and seven columns, as shown in the following figure:

0, 0	0, 1	0, 2	0, 3	0, 4	0, 5	0, 6
1, 0	1, 1	1, 2	1, 3	1, 4	1, 5	1, 6
2, 0	2, 1	2, 2	2, 3	2, 4	2, 5	2, 6
3, 0	3, 1	3, 2	3, 3	3, 4	3, 5	3, 6
4, 0	4, 1	4, 2	4, 3	4, 4	4, 5	4, 6

Conceptual Layout of Multidimensional Array

In declarations of multidimensioned arrays that have an initializer list (as described in [Initializers](#)), the constant expression that specifies the bounds for the first dimension can be omitted. For example:

```

// arrays2.cpp
// compile with: /c
const int cMarkets = 4;
// Declare a float that represents the transportation costs.
double TransportCosts[][][cMarkets] = {
    { 32.19, 47.29, 31.99, 19.11 },
    { 11.29, 22.49, 33.47, 17.29 },
    { 41.97, 22.09, 9.76, 22.55 }
};

```

The preceding declaration defines an array that is three rows by four columns. The rows represent factories and the columns represent markets to which the factories ship. The values are the transportation costs from the factories to the markets. The first dimension of the array is left out, but the compiler fills it in by examining the initializer.

Topics in this section:

- [Using Arrays](#)
- [Arrays in Expressions](#)

- Interpretation of Subscript Operator
- Indirection on Array Types
- Ordering of C++ Arrays

## Example

The technique of omitting the bounds specification for the first dimension of a multidimensional array can also be used in function declarations as follows:

```
// multidimensional_arrays.cpp
// compile with: /EHsc
// arguments: 3
#include <limits> // Includes DBL_MAX
#include <iostream>

const int cMkts = 4, cFacts = 2;

// Declare a float that represents the transportation costs
double TransportCosts[][cMkts] = {
    { 32.19, 47.29, 31.99, 19.11 },
    { 11.29, 22.49, 33.47, 17.29 },
    { 41.97, 22.09, 9.76, 22.55 }
};

// Calculate size of unspecified dimension
const int cFactories = sizeof(TransportCosts) /
    sizeof( double[cMkts] );

double FindMinToMkt( int Mkt, double myTransportCosts[][cMkts], int mycFacts);

using namespace std;

int main( int argc, char *argv[] ) {
    double MinCost;

    if (argc[1] == 0) {
        cout << "You must specify the number of markets." << endl;
        exit(0);
    }
    MinCost = FindMinToMkt( *argv[1] - '0', TransportCosts, cFacts);
    cout << "The minimum cost to Market " << argv[1] << " is: "
        << MinCost << "\n";
}

double FindMinToMkt(int Mkt, double myTransportCosts[][cMkts], int mycFacts) {
    double MinCost = DBL_MAX;

    for( int i = 0; i < cFacts; ++i )
        MinCost = (MinCost < TransportCosts[i][Mkt]) ?
            MinCost : TransportCosts[i][Mkt];

    return MinCost;
}
```

The minimum cost to Market 3 is: 17.29

## Comments

The function `FindMinToMkt` is written such that adding new factories does not require any code changes, just a recompilation.

## See Also

# Using Arrays (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

You can access individual elements of an array by using the array subscript operator (`[ ]`). If a one-dimensional array is used in an expression that has no subscript, the array name evaluates to a pointer to the first element in the array.

```
// using_arrays.cpp
int main() {
    char chArray[10];
    char *pch = chArray;    // Evaluates to a pointer to the first element.
    char   ch = chArray[0]; // Evaluates to the value of the first element.
    ch = chArray[3];      // Evaluates to the value of the fourth element.
}
```

When you use multidimensional arrays, you can use various combinations in expressions.

```
// using_arrays_2.cpp
// compile with: /EHsc /W1
#include <iostream>
using namespace std;
int main() {
    double multi[4][4][3];    // Declare the array.
    double (*p2multi)[3];
    double (*p1multi);

    cout << multi[3][2][2] << "\n";    // C4700 Use three subscripts.
    p2multi = multi[3];                // Make p2multi point to
                                       // fourth "plane" of multi.
    p1multi = multi[3][2];            // Make p1multi point to
                                       // fourth plane, third row
                                       // of multi.
}
```

In the preceding code, `multi` is a three-dimensional array of type `double`. The `p2multi` pointer points to an array of type `double` of size three. In this example, the array is used with one, two, and three subscripts. Although it is more common to specify all subscripts, as in the `cout` statement, it is sometimes useful to select a specific subset of array elements, as shown in the statements that follow `cout`.

## See Also

[Arrays](#)

# Initializing Arrays

12/21/2017 • 1 min to read • [Edit Online](#)

If a class has a constructor, arrays of that class are initialized by a constructor. If there are fewer items in the initializer list than elements in the array, the default constructor is used for the remaining elements. If no default constructor is defined for the class, the initializer list must be complete — that is, there must be one initializer for each element in the array.

Consider the `Point` class that defines two constructors:

```
// initializing_arrays1.cpp
class Point
{
public:
    Point() // Default constructor.
    {
    }
    Point( int, int ) // Construct from two ints
    {
    }
};

// An array of Point objects can be declared as follows:
Point aPoint[3] = {
    Point( 3, 3 ) // Use int, int constructor.
};

int main()
{}
```

The first element of `aPoint` is constructed using the constructor `Point( int, int )`; the remaining two elements are constructed using the default constructor.

Static member arrays (whether **const** or not) can be initialized in their definitions (outside the class declaration). For example:

```
// initializing_arrays2.cpp
class WindowColors
{
public:
    static const char *rgszWindowPartList[7];
};

const char *WindowColors::rgszWindowPartList[7] = {
    "Active Title Bar", "Inactive Title Bar", "Title Bar Text",
    "Menu Bar", "Menu Bar Text", "Window Background", "Frame"  };
int main()
{}
```

# Arrays in Expressions

12/21/2017 • 1 min to read • [Edit Online](#)

When an identifier of an array type appears in an expression other than `sizeof`, address-of (`&`), or initialization of a reference, it is converted to a pointer to the first array element. For example:

```
char szError1[] = "Error: Disk drive not ready.";
char *psz = szError1;
```

The pointer `psz` points to the first element of the array `szError1`. Note that arrays, unlike pointers, are not modifiable l-values. Therefore, the following assignment is illegal:

```
szError1 = psz;
```

## See Also

[Arrays](#)

# Interpretation of Subscript Operator

12/21/2017 • 1 min to read • [Edit Online](#)

Like other operators, the subscript operator (`[ ]`) can be redefined by the user. The default behavior of the subscript operator, if not overloaded, is to combine the array name and the subscript using the following method:

`*((array-name) + (subscript))`

As in all addition that involves pointer types, scaling is performed automatically to adjust for the size of the type. Therefore, the resultant value is not *subscript* bytes from the origin of *array-name*; rather, it is the *subscript<sup>n</sup>*th element of the array. (For more information about this conversion, see [Additive Operators](#).)

Similarly, for multidimensional arrays, the address is derived using the following method:

```
((  
array-name) + (  
subscript1 max2 * max3...maxn) + subscript2 * max3...maxn) . . . + subscriptn))
```

## See Also

[Arrays](#)

# Indirection on Array Types

12/21/2017 • 1 min to read • [Edit Online](#)

Use of the indirection operator (`\`) on an  $n$ -dimensional array type yields an  $n-1$  dimensional array. If  $n$  is 1, a scalar (or array element) is yielded.

## See Also

[Arrays](#)

# Ordering of C++ Arrays

12/21/2017 • 1 min to read • [Edit Online](#)

C++ arrays are stored in row-major order. Row-major order means the last subscript varies the fastest.

## See Also

[Arrays](#)

# References (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

A reference, like a pointer, stores the address of an object that is located elsewhere in memory. Unlike a pointer, a reference after it is initialized cannot be made to refer to a different object or set to null. There are two kinds of references: lvalue references which refer to a named variable and rvalue references which refer to a [temporary object](#). The & operator signifies an lvalue reference and the && operator signifies either an rvalue reference, or a universal reference (either rvalue or lvalue) depending on the context.

References may be declared using the following syntax:

```
[storage-class-specifiers] [cv-qualifiers] type-specifiers  
[ms-modifier] declarator [= expression];
```

Any valid declarator specifying a reference may be used. Unless the reference is a reference to function or array type, the following simplified syntax applies:

```
[storage-class-specifiers] [cv-qualifiers] type-specifiers [& or &&]  
[cv-qualifiers] identifier [= expression];
```

References are declared using the following sequence:

1. The declaration specifiers:

- An optional storage class specifier.
- Optional **const** and/or **volatile** qualifiers.
- The type specifier: the name of a type.
- 1. The declarator:
- An optional Microsoft specific modifier. For more information, see [Microsoft-Specific Modifiers](#).
- The & operator or && operator.
- Optional **const** and/or **volatile** qualifiers.
- The identifier.
  1. An optional initializer.

The more complex declarator forms for pointers to arrays and functions also apply to references to arrays and functions, see [pointers](#) and [declarators](#).

Multiple declarators and initializers may appear in a comma-separated list following a single declaration specifier. For example:

```
int &i;  
int &i, &j;
```

References, pointers and objects may be declared together:

```
int &ref, *ptr, k;
```

A reference holds the address of an object, but behaves syntactically like an object.

In the following program, notice that the name of the object, `Today`, and the reference to the object, `TodayRef`, can be used identically in programs:

## Example

```
// references.cpp
#include <stdio.h>
struct S {
    short i;
};

int main() {
    S s;      // Declare the object.
    S& SRef = s;  // Declare the reference.
    s.i = 3;

    printf_s("%d\n", s.i);
    printf_s("%d\n", SRef.i);

    SRef.i = 4;
    printf_s("%d\n", s.i);
    printf_s("%d\n", SRef.i);
}
```

```
3
3
4
4
```

## Comment

Topics in this section:

- [Reference-Type Function Arguments](#)
- [Reference-Type Function Returns](#)
- [References to Pointers](#)

# Lvalue Reference Declarator: &

12/21/2017 • 1 min to read • [Edit Online](#)

Holds the address of an object but behaves syntactically like an object.

## Syntax

```
type-id & cast-expression
```

## Remarks

You can think of an lvalue reference as another name for an object. An lvalue reference declaration consists of an optional list of specifiers followed by a reference declarator. A reference must be initialized and cannot be changed.

Any object whose address can be converted to a given pointer type can also be converted to the similar reference type. For example, any object whose address can be converted to type `char *` can also be converted to type `char &`.

Do not confuse reference declarations with use of the [address-of operator](#). When the `& identifier` is preceded by a type, such as `int` or `char`, `identifier` is declared as a reference to the type. When `& identifier` is not preceded by a type, the usage is that of the address-of operator.

## Example

The following example demonstrates the reference declarator by declaring a `Person` object and a reference to that object. Because `rFriend` is a reference to `myFriend`, updating either variable changes the same object.

```
// reference_declarator.cpp
// compile with: /EHsc
// Demonstrates the reference declarator.
#include <iostream>
using namespace std;

struct Person
{
    char* Name;
    short Age;
};

int main()
{
    // Declare a Person object.
    Person myFriend;

    // Declare a reference to the Person object.
    Person& rFriend = myFriend;

    // Set the fields of the Person object.
    // Updating either variable changes the same object.
    myFriend.Name = "Bill";
    rFriend.Age = 40;

    // Print the fields of the Person object to the console.
    cout << rFriend.Name << " is " << myFriend.Age << endl;
}
```

```
Bill is 40
```

## See Also

[References](#)

[Reference-Type Function Arguments](#)

[Reference-Type Function Returns](#)

[References to Pointers](#)

# Rvalue Reference Declarator: &&

1/3/2018 • 12 min to read • [Edit Online](#)

Holds a reference to an rvalue expression.

## Syntax

```
type-id && cast-expression
```

## Remarks

Rvalue references enable you to distinguish an lvalue from an rvalue. Lvalue references and rvalue references are syntactically and semantically similar, but they follow somewhat different rules. For more information about lvalues and rvalues, see [Lvalues and Rvalues](#). For more information about lvalue references, see [Lvalue Reference Declarator: &](#).

The following sections describe how rvalue references support the implementation of *move semantics* and *perfect forwarding*.

## Move Semantics

Rvalue references support the implementation of *move semantics*, which can significantly increase the performance of your applications. Move semantics enables you to write code that transfers resources (such as dynamically allocated memory) from one object to another. Move semantics works because it enables resources to be transferred from temporary objects that cannot be referenced elsewhere in the program.

To implement move semantics, you typically provide a *move constructor*, and optionally a move assignment operator (`operator=`), to your class. Copy and assignment operations whose sources are rvalues then automatically take advantage of move semantics. Unlike the default copy constructor, the compiler does not provide a default move constructor. For more information about how to write a move constructor and how to use it in your application, see [Move Constructors and Move Assignment Operators \(C++\)](#).

You can also overload ordinary functions and operators to take advantage of move semantics. Visual C++ 2010 introduces move semantics into the C++ Standard Library. For example, the `string` class implements operations that perform move semantics. Consider the following example that concatenates several strings and prints the result:

```
// string_concatenation.cpp
// compile with: /EHsc
#include <iostream>
#include <string>
using namespace std;

int main()
{
    string s = string("h") + "e" + "ll" + "o";
    cout << s << endl;
}
```

Before Visual C++ 2010, each call to `operator+` allocates and returns a new temporary `string` object (an rvalue).

`operator+` cannot append one string to the other because it does not know whether the source strings are lvalues or rvalues. If the source strings are both lvalues, they might be referenced elsewhere in the program and therefore must not be modified. By using rvalue references, `operator+` can be modified to take rvalues, which cannot be referenced elsewhere in the program. Therefore, `operator+` can now append one string to another. This can significantly reduce the number of dynamic memory allocations that the `string` class must perform. For more information about the `string` class, see [basic\\_string Class](#).

Move semantics also helps when the compiler cannot use Return Value Optimization (RVO) or Named Return Value Optimization (NRVO). In these cases, the compiler calls the move constructor if the type defines it. For more information about Named Return Value Optimization, see [Named Return Value Optimization in Visual C++ 2005](#).

To better understand move semantics, consider the example of inserting an element into a `vector` object. If the capacity of the `vector` object is exceeded, the `vector` object must reallocate memory for its elements and then copy each element to another memory location to make room for the inserted element. When an insertion operation copies an element, it creates a new element, calls the copy constructor to copy the data from the previous element to the new element, and then destroys the previous element. Move semantics enables you to move objects directly without having to perform expensive memory allocation and copy operations.

To take advantage of move semantics in the `vector` example, you can write a move constructor to move data from one object to another.

For more information about the introduction of move semantics into the C++ Standard Library in Visual C++ 2010, see [C++ Standard Library](#).

## Perfect Forwarding

Perfect forwarding reduces the need for overloaded functions and helps avoid the forwarding problem. The *forwarding problem* can occur when you write a generic function that takes references as its parameters and it passes (or *forwards*) these parameters to another function. For example, if the generic function takes a parameter of type `const T&`, then the called function cannot modify the value of that parameter. If the generic function takes a parameter of type `T&`, then the function cannot be called by using an rvalue (such as a temporary object or integer literal).

Ordinarily, to solve this problem, you must provide overloaded versions of the generic function that take both `T&` and `const T&` for each of its parameters. As a result, the number of overloaded functions increases exponentially with the number of parameters. Rvalue references enable you to write one version of a function that accepts arbitrary arguments and forwards them to another function as if the other function had been called directly.

Consider the following example that declares four types, `W`, `X`, `Y`, and `Z`. The constructor for each type takes a different combination of `const` and non-`const` lvalue references as its parameters.

```

struct W
{
    W(int&, int&) {}
};

struct X
{
    X(const int&, int&) {}
};

struct Y
{
    Y(int&, const int&) {}
};

struct Z
{
    Z(const int&, const int&) {}
};

```

Suppose you want to write a generic function that generates objects. The following example shows one way to write this function:

```

template <typename T, typename A1, typename A2>
T* factory(A1& a1, A2& a2)
{
    return new T(a1, a2);
}

```

The following example shows a valid call to the `factory` function:

```

int a = 4, b = 5;
W* pw = factory<W>(a, b);

```

However, the following example does not contain a valid call to the `factory` function because `factory` takes lvalue references that are modifiable as its parameters, but it is called by using rvalues:

```

Z* pz = factory<Z>(2, 2);

```

Ordinarily, to solve this problem, you must create an overloaded version of the `factory` function for every combination of `A&` and `const A&` parameters. Rvalue references enable you to write one version of the `factory` function, as shown in the following example:

```

template <typename T, typename A1, typename A2>
T* factory(A1&& a1, A2&& a2)
{
    return new T(std::forward<A1>(a1), std::forward<A2>(a2));
}

```

This example uses rvalue references as the parameters to the `factory` function. The purpose of the `std::forward` function is to forward the parameters of the factory function to the constructor of the template class.

The following example shows the `main` function that uses the revised `factory` function to create instances of the `W`, `X`, `Y`, and `Z` classes. The revised `factory` function forwards its parameters (either lvalues or rvalues) to the appropriate class constructor.

```

int main()
{
    int a = 4, b = 5;
    W* pw = factory<W>(a, b);
    X* px = factory<X>(2, b);
    Y* py = factory<Y>(a, 2);
    Z* pz = factory<Z>(2, 2);

    delete pw;
    delete px;
    delete py;
    delete pz;
}

```

## Additional Properties of Rvalue References

**You can overload a function to take an lvalue reference and an rvalue reference.**

By overloading a function to take a `const` lvalue reference or an rvalue reference, you can write code that distinguishes between non-modifiable objects (lvalues) and modifiable temporary values (rvalues). You can pass an object to a function that takes an rvalue reference unless the object is marked as `const`. The following example shows the function `f`, which is overloaded to take an lvalue reference and an rvalue reference. The `main` function calls `f` with both lvalues and an rvalue.

```

// reference-overload.cpp
// Compile with: /EHsc
#include <iostream>
using namespace std;

// A class that contains a memory resource.
class MemoryBlock
{
    // TODO: Add resources for the class here.
};

void f(const MemoryBlock&)
{
    cout << "In f(const MemoryBlock&). This version cannot modify the parameter." << endl;
}

void f(MemoryBlock&&)
{
    cout << "In f(MemoryBlock&&). This version can modify the parameter." << endl;
}

int main()
{
    MemoryBlock block;
    f(block);
    f(MemoryBlock());
}

```

This example produces the following output:

```

In f(const MemoryBlock&). This version cannot modify the parameter.
In f(MemoryBlock&&). This version can modify the parameter.

```

In this example, the first call to `f` passes a local variable (an lvalue) as its argument. The second call to `f` passes a temporary object as its argument. Because the temporary object cannot be referenced elsewhere in the

program, the call binds to the overloaded version of `f` that takes an rvalue reference, which is free to modify the object.

### The compiler treats a named rvalue reference as an lvalue and an unnamed rvalue reference as an rvalue.

When you write a function that takes an rvalue reference as its parameter, that parameter is treated as an lvalue in the body of the function. The compiler treats a named rvalue reference as an lvalue because a named object can be referenced by several parts of a program; it would be dangerous to allow multiple parts of a program to modify or remove resources from that object. For example, if multiple parts of a program try to transfer resources from the same object, only the first part will successfully transfer the resource.

The following example shows the function `g`, which is overloaded to take an lvalue reference and an rvalue reference. The function `f` takes an rvalue reference as its parameter (a named rvalue reference) and returns an rvalue reference (an unnamed rvalue reference). In the call to `g` from `f`, overload resolution selects the version of `g` that takes an lvalue reference because the body of `f` treats its parameter as an lvalue. In the call to `g` from `main`, overload resolution selects the version of `g` that takes an rvalue reference because `f` returns an rvalue reference.

```
// named-reference.cpp
// Compile with: /EHsc
#include <iostream>
using namespace std;

// A class that contains a memory resource.
class MemoryBlock
{
    // TODO: Add resources for the class here.
};

void g(const MemoryBlock&)
{
    cout << "In g(const MemoryBlock&)." << endl;
}

void g(MemoryBlock&&)
{
    cout << "In g(MemoryBlock&&)." << endl;
}

MemoryBlock&& f(MemoryBlock&& block)
{
    g(block);
    return block;
}

int main()
{
    g(f(MemoryBlock()));
}
```

This example produces the following output:

```
In g(const MemoryBlock&).
In g(MemoryBlock&&).
```

In this example, the `main` function passes an rvalue to `f`. The body of `f` treats its named parameter as an lvalue. The call from `f` to `g` binds the parameter to an lvalue reference (the first overloaded version of `g`).

- You can cast an lvalue to an rvalue reference.

The C++ Standard Library `std::move` function enables you to convert an object to an rvalue reference to that object. Alternatively, you can use the `static_cast` keyword to cast an lvalue to an rvalue reference, as shown in the following example:

```
// cast-reference.cpp
// Compile with: /EHsc
#include <iostream>
using namespace std;

// A class that contains a memory resource.
class MemoryBlock
{
    // TODO: Add resources for the class here.
};

void g(const MemoryBlock&)
{
    cout << "In g(const MemoryBlock&)." << endl;
}

void g(MemoryBlock&&)
{
    cout << "In g(MemoryBlock&&)." << endl;
}

int main()
{
    MemoryBlock block;
    g(block);
    g(static_cast<MemoryBlock&&>(block));
}
```

This example produces the following output:

```
In g(const MemoryBlock&).
In g(MemoryBlock&&).
```

### Function templates deduce their template argument types and then use reference collapsing rules.

It is common to write a function template that passes (or *forwards*) its parameters to another function. It is important to understand how template type deduction works for function templates that take rvalue references.

If the function argument is an rvalue, the compiler deduces the argument to be an rvalue reference. For example, if you pass an rvalue reference to an object of type `x` to a template function that takes type `T&&` as its parameter, template argument deduction deduces `T` to be `x`. Therefore, the parameter has type `x&&`. If the function argument is an lvalue or `const` lvalue, the compiler deduces its type to be an lvalue reference or `const` lvalue reference of that type.

The following example declares one structure template and then specializes it for various reference types. The `print_type_and_value` function takes an rvalue reference as its parameter and forwards it to the appropriate specialized version of the `S::print` method. The `main` function demonstrates the various ways to call the `S::print` method.

```
// template-type-deduction.cpp
// Compile with: /EHsc
#include <iostream>
#include <string>
using namespace std;

template<typename T> struct S;
```

```

// The following structures specialize S by
// lvalue reference (T&), const lvalue reference (const T&),
// rvalue reference (T&&), and const rvalue reference (const T&&).
// Each structure provides a print method that prints the type of
// the structure and its parameter.

template<typename T> struct S<T&> {
    static void print(T& t)
    {
        cout << "print<T&>: " << t << endl;
    }
};

template<typename T> struct S<const T&> {
    static void print(const T& t)
    {
        cout << "print<const T&>: " << t << endl;
    }
};

template<typename T> struct S<T&&> {
    static void print(T&& t)
    {
        cout << "print<T&&>: " << t << endl;
    }
};

template<typename T> struct S<const T&&> {
    static void print(const T&& t)
    {
        cout << "print<const T&&>: " << t << endl;
    }
};

// This function forwards its parameter to a specialized
// version of the S type.
template <typename T> void print_type_and_value(T&& t)
{
    S<T&&>::print(std::forward<T>(t));
}

// This function returns the constant string "fourth".
const string fourth() { return string("fourth"); }

int main()
{
    // The following call resolves to:
    // print_type_and_value<string&>(string& && t)
    // Which collapses to:
    // print_type_and_value<string&>(string& t)
    string s1("first");
    print_type_and_value(s1);

    // The following call resolves to:
    // print_type_and_value<const string&>(const string& && t)
    // Which collapses to:
    // print_type_and_value<const string&>(const string& t)
    const string s2("second");
    print_type_and_value(s2);

    // The following call resolves to:
    // print_type_and_value<string&&>(string&& t)
    print_type_and_value(string("third"));

    // The following call resolves to:
    // print_type_and_value<const string&&>(const string&& t)
    print_type_and_value(fourth());
}

```

This example produces the following output:

```
print<T&>; first
print<const T&>; second
print<T&&>; third
print<const T&&>; fourth
```

To resolve each call to the `print_type_and_value` function, the compiler first performs template argument deduction. The compiler then applies reference collapsing rules when it substitutes the deduced template arguments for the parameter types. For example, passing the local variable `s1` to the `print_type_and_value` function causes the compiler to produce the following function signature:

```
print_type_and_value<string&>(string& && t)
```

The compiler uses reference collapsing rules to reduce the signature to the following:

```
print_type_and_value<string&>(string& t)
```

This version of the `print_type_and_value` function then forwards its parameter to the correct specialized version of the `S::print` method.

The following table summarizes the reference collapsing rules for template argument type deduction:

Expanded type	Collapsed type
<code>T&amp; &amp;</code>	<code>T&amp;</code>
<code>T&amp; &amp;&amp;</code>	<code>T&amp;</code>
<code>T&amp;&amp; &amp;</code>	<code>T&amp;</code>
<code>T&amp;&amp; &amp;&amp;</code>	<code>T&amp;&amp;</code>

Template argument deduction is an important element of implementing perfect forwarding. The section Perfect Forwarding, which is presented earlier in this topic, describes perfect forwarding in more detail.

## Summary

Rvalue references distinguish lvalues from rvalues. They can help you improve the performance of your applications by eliminating the need for unnecessary memory allocations and copy operations. They also enable you to write one version of a function that accepts arbitrary arguments and forwards them to another function as if the other function had been called directly.

## See Also

[Expressions with Unary Operators](#)

[Lvalue Reference Declarator: &](#)

[Lvalues and Rvalues](#)

[Move Constructors and Move Assignment Operators \(C++\)](#)

[C++ Standard Library](#)



# Reference-Type Function Arguments

12/21/2017 • 1 min to read • [Edit Online](#)

It is often more efficient to pass references, rather than large objects, to functions. This allows the compiler to pass the address of the object while maintaining the syntax that would have been used to access the object. Consider the following example that uses the `Date` structure:

```
// reference_type_function_arguments.cpp
struct Date
{
short DayOfWeek;
short Month;
short Day;
short Year;
};

// Create a Julian date of the form DDDYYYY
// from a Gregorian date.
long JulianFromGregorian( Date& GDate )
{
static int cDaysInMonth[] = {
31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31
};
long JDate = 0;
// Add in days for months already elapsed.
for ( int i = 0; i < GDate.Month - 1; ++i )
JDate += cDaysInMonth[i];
// Add in days for this month.
JDate += GDate.Day;

// Check for leap year.
if ( GDate.Year % 100 != 0 && GDate.Year % 4 == 0 )
JDate++;
// Add in year.
JDate *= 10000;
JDate += GDate.Year;

return JDate;
}

int main()
{
```

The preceding code shows that members of a structure passed by reference are accessed using the member-selection operator `(.)` instead of the pointer member-selection operator `(->)`.

Although arguments passed as reference types observe the syntax of non-pointer types, they retain one important characteristic of pointer types: they are modifiable unless declared as **const**. Because the intent of the preceding code is not to modify the object `GDate`, a more appropriate function prototype is:

```
long JulianFromGregorian( const Date& GDate );
```

This prototype guarantees that the function `JulianFromGregorian` will not change its argument.

Any function prototyped as taking a reference type can accept an object of the same type in its place because there is a standard conversion from *typename* to *typename&*.

## See Also

[References](#)

# Reference-Type Function Returns

12/21/2017 • 2 min to read • [Edit Online](#)

Functions can be declared to return a reference type. There are two reasons to make such a declaration:

- The information being returned is a large enough object that returning a reference is more efficient than returning a copy.
- The type of the function must be an l-value.
- The referred-to object will not go out of scope when the function returns.

Just as it can be more efficient to pass large objects *to* functions by reference, it also can be more efficient to return large objects *from* functions by reference. Reference-return protocol eliminates the necessity of copying the object to a temporary location prior to returning.

Reference-return types can also be useful when the function must evaluate to an l-value. Most overloaded operators fall into this category, particularly the assignment operator. Overloaded operators are covered in [Overloaded Operators](#).

## Example

Consider the `Point` example:

```

// refType_function_returns.cpp
// compile with: /EHsc

#include <iostream>
using namespace std;

class Point
{
public:
// Define "accessor" functions as
// reference types.
unsigned& x();
unsigned& y();

private:
// Note that these are declared at class scope:
unsigned obj_x;
unsigned obj_y;
};

unsigned& Point :: x()
{
return obj_x;
}

unsigned& Point :: y()
{
return obj_y;
}

int main()
{
Point ThePoint;
// Use x() and y() as l-values.
ThePoint.x() = 7;
ThePoint.y() = 9;

// Use x() and y() as r-values.
cout << "x = " << ThePoint.x() << "\n"
<< "y = " << ThePoint.y() << "\n";
}

```

## Output

```

x = 7
y = 9

```

Notice that the functions `x` and `y` are declared as returning reference types. These functions can be used on either side of an assignment statement.

Note also that in `main`, `ThePoint` object remains in scope, and therefore its reference members are still alive and can be safely accessed.

Declarations of reference types must contain initializers except in the following cases:

- Explicit `extern` declaration
- Declaration of a class member
- Declaration within a class
- Declaration of an argument to a function or the return type for a function

## Caution returning address of local

If you declare an object at local scope, that object will be destroyed when the function returns. If the function returns a reference to that object, that reference will probably cause an access violation at runtime if the caller attempts to use the null reference.

```
// C4172 means Don't do this!!!
Foo& GetFoo()
{
    Foo f;
    ...
    return f;
} // f is destroyed here
```

The compiler issues a warning in this case: `warning C4172: returning address of local variable or temporary`. In simple programs it is possible that occasionally no access violation will occur if the reference is accessed by the caller before the memory location is overwritten. This is due to sheer luck. Heed the warning.

## See Also

[References](#)

# References to Pointers

12/21/2017 • 2 min to read • [Edit Online](#)

References to pointers can be declared in much the same way as references to objects. Declaring a reference to a pointer yields a modifiable value that is used like a normal pointer.

## Example

The following code samples illustrate the difference between using a pointer to a pointer and a reference to a pointer.

Functions `Add1` and `Add2` are functionally equivalent (although they are not called the same way). The difference is that `Add1` uses double indirection whereas `Add2` uses the convenience of a reference to a pointer.

```
// references_to_pointers.cpp
// compile with: /EHsc

#include <iostream>
#include <string>

// C++ Standard Library namespace
using namespace std;

enum {
    sizeOfBuffer = 132
};

// Define a binary tree structure.
struct BTree {
    char *szText;
    BTree *Left;
    BTree *Right;
};

// Define a pointer to the root of the tree.
BTree *btRoot = 0;

int Add1( BTree **Root, char *szToAdd );
int Add2( BTree*& Root, char *szToAdd );
void PrintTree( BTree* btRoot );

int main( int argc, char *argv[] ) {
    // Usage message
    if( argc < 2 ) {
        cerr << "Usage: Refptr [1 | 2]" << "\n";
        cerr << "\nwhere:\n";
        cerr << "1 uses double indirection\n";
        cerr << "2 uses a reference to a pointer.\n";
        cerr << "\nInput is from stdin.\n";
        return 1;
    }

    char *szBuf = new char[sizeOfBuffer];
    if (szBuf == NULL) {
        cerr << "Out of memory!\n";
        return -1;
    }

    // Read a text file from the standard input device and
    // build a binary tree.
```

```

//while( !cin.eof() )
{
    cin.get( szBuf, sizeOfBuffer, '\n' );
    cin.get();

    if ( strlen( szBuf ) ) {
        switch ( *argv[1] ) {
            // Method 1: Use double indirection.
            case '1':
                Add1( &btRoot, szBuf );
                break;
            // Method 2: Use reference to a pointer.
            case '2':
                Add2( btRoot, szBuf );
                break;
            default:
                cerr << "Illegal value "
                    << *argv[1]
                    << "' supplied for add method.\n"
                    << "Choose 1 or 2.\n";
                return -1;
        }
    }
}

// Display the sorted list.
PrintTree( btRoot );
}

// PrintTree: Display the binary tree in order.
void PrintTree( BTREE* MybtRoot ) {
    // Traverse the left branch of the tree recursively.
    if ( btRoot->Left )
        PrintTree( btRoot->Left );

    // Print the current node.
    cout << btRoot->szText << "\n";

    // Traverse the right branch of the tree recursively.
    if ( btRoot->Right )
        PrintTree( btRoot->Right );
}

// Add1: Add a node to the binary tree.
//      Uses double indirection.
int Add1( BTREE **Root, char *szToAdd ) {
    if ( (*Root) == 0 ) {
        (*Root) = new BTREE;
        (*Root)->Left = 0;
        (*Root)->Right = 0;
        (*Root)->szText = new char[strlen( szToAdd ) + 1];
        strcpy_s((*Root)->szText, (strlen( szToAdd ) + 1), szToAdd );
        return 1;
    }
    else {
        if ( strcmp( (*Root)->szText, szToAdd ) > 0 )
            return Add1( &((*Root)->Left), szToAdd );
        else
            return Add1( &((*Root)->Right), szToAdd );
    }
}

// Add2: Add a node to the binary tree.
//      Uses reference to pointer
int Add2( BTREE*& Root, char *szToAdd ) {
    if ( Root == 0 ) {
        Root = new BTREE;
        Root->Left = 0;
        Root->Right = 0;
        Root->szText = new char[strlen( szToAdd ) + 1];
    }
}

```

```
    strcpy_s( Root->szText, (strlen( szToAdd ) + 1), szToAdd );
    return 1;
}
else {
    if ( strcmp( Root->szText, szToAdd ) > 0 )
        return Add2( Root->Left, szToAdd );
    else
        return Add2( Root->Right, szToAdd );
}
}
```

Usage: Refptr [1 | 2]

where:

1 uses double indirection

2 uses a reference to a pointer.

Input is from stdin.

## See Also

[References](#)

# Pointers (C++)

12/21/2017 • 4 min to read • [Edit Online](#)

Pointers are declared using the following sequence.

```
[storage-class-specifiers] [cv-qualifiers] type-specifiers  
[ms-modifier] declarator ;
```

where any valid pointer declarator may be used for `declarator`. The syntax for a simple pointer declarator is as follows:

```
* [cv-qualifiers] identifier [= expression]
```

1. The declaration specifiers:

- An optional storage class specifier. For more information, see [Specifiers](#).
- An optional `const` or `volatile` keyword applying to the type of the object to be pointed to.
- The type specifier: the name of a type representing the type of the object to be pointed to.
  1. The declarator:
    - An optional Microsoft specific modifier. For more information, see [Microsoft-Specific Modifiers](#).
    - The `*` operator.
    - An optional `const` or `volatile` keyword applying to the pointer itself.
    - The identifier.
    - An optional initializer.

The declarator for a pointer to function looks like this:

```
(* [cv-qualifiers] identifier )( argument-list ) [cv-qualifiers]  
[exception specification] [= expression];
```

- For an array of pointers, the syntax looks like this:

```
* identifier [ [ constant-expression ] ]
```

- Multiple declarators and their initializers may appear together in a single declaration in a comma separated list following the declaration specifier.

A simple example of a pointer declaration is:

```
char *pch;
```

The preceding declaration specifies that `pch` points to an object of type `char`.

A more complex example is

```
static unsigned int * const ptr;
```

The preceding declaration specifies that `ptr` is a constant pointer to an object of type `unsigned int` with static storage duration.

The next example shows how multiple pointers are declared and initialized:

```
static int *p = &i, *q = &j;
```

In the preceding example, pointers `p` and `q` both point to objects of type `int` and are initialized to the addresses of `i` and `j` respectively. The storage class specifier `static` applies to both pointers.

## Example

```
// pointer.cpp
// compile with: /EHsc
#include <iostream>
int main() {
    int i = 1, j = 2; // local variables on the stack
    int *p;

    // a pointer may be assigned to "point to" the value of
    // another variable using the & (address of) operator
    p = & j;

    // since j was on the stack, this address will be somewhere
    // on the stack. Pointers are printed in hex format using
    // %p and conventionally marked with 0x.
    printf_s("0x%p\n", p);

    // The * (indirection operator) can be read as "the value
    // pointed to by".
    // Since p is pointing to j, this should print "2"
    printf_s("0x%p %d\n", p, *p);

    // changing j will change the result of the indirection
    // operator on p.
    j = 7;
    printf_s("0x%p %d\n", p, *p);

    // The value of j can also be changed through the pointer
    // by making an assignment to the dereferenced pointer
    *p = 10;
    printf_s("j is %d\n", j); // j is now 10

    // allocate memory on the heap for an integer,
    // initialize to 5
    p = new int(5);

    // print the pointer and the object pointed to
    // the address will be somewhere on the heap
    printf_s("0x%p %d\n", p, *p);

    // free the memory pointed to by p
    delete p;

    // At this point, dereferencing p with *p would trigger
    // a runtime access violation.

    // Pointer arithmetic may be done with an array declared
    // on the stack or allocated on the heap with new.
    // The increment operator takes into account the size
    // of the elements.
```

```

// or the objects pointed to.
p = new int[5];
for (i = 0; i < 5; i++, p++) {
    *p = i * 10;
    printf_s("0x%p %d\n", p, *p);
}

// A common expression seen is dereferencing in combination
// with increment or decrement operators, as shown here.
// The indirection operator * takes precedence over the
// increment operator ++.
// These are particularly useful in manipulating char arrays.
char s1[4] = "cat";
char s2[4] = "dog";
char* p1 = s1;
char* p2 = s2;

// the following is a string copy operation
while (*p1++ = *p2++);

// s2 was copied into s1, so now they are both equal to "dog"
printf_s("%s %s", s1, s2);
}

```

```

0x0012FEC8
0x0012FEC8 2
0x0012FEC8 7
j is 10
0x00320850 5
0x00320850 0
0x00320854 10
0x00320858 20
0x0032085C 30
0x00320860 40
dog dog

```

## Example

Another example illustrates the use of pointers in data structures; in this case, a linked list.

```

// pointer_linkedlist.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

struct NewNode {
    NewNode() : node(0){}
    int i;
    NewNode * node;
};

void WalkList(NewNode * ptr) {
    if (ptr != 0) {
        int i = 1;
        while (ptr->node != 0 ) {
            cout << "node " << i++ << " = " << ptr->i << endl;
            ptr = ptr->node;
        }
        cout << "node " << i++ << " = " << ptr->i << endl;
    }
}

void AddNode(NewNode ** ptr) {
    NewNode * walker = 0;
    NewNode * MyNewNode = new NewNode;
    cout << "enter a number: " << endl;
    cin >> MyNewNode->i;

    if (*ptr == 0)
        *ptr = MyNewNode;
    else {
        walker = *ptr;
        while (walker->node != 0)
            walker = walker->node;

        walker->node = MyNewNode;
    }
}

int main() {
    char ans = ' ';
    NewNode * ptr = 0;
    do {
        cout << "a (add node)  d (display list)  q (quit)" << endl;
        cin >> ans;
        switch (ans) {
        case 'a':
            AddNode(&ptr);
            break;
        case 'd':
            WalkList(ptr);
            break;
        }
    } while (ans != 'q');
}

```

```
a  
45  
d  
a  
789  
d  
qa (add node) d (display list) q (quit)  
enter a number:  
a (add node) d (display list) q (quit)  
node 1 = 45  
a (add node) d (display list) q (quit)  
enter a number:  
a (add node) d (display list) q (quit)  
node 1 = 45  
node 2 = 789  
a (add node) d (display list) q (quit)
```

## See Also

[Indirection Operator: \\*](#)

[Address-of Operator: &](#)

# const and volatile Pointers

12/21/2017 • 3 min to read • [Edit Online](#)

The `const` and `volatile` keywords change how pointers are treated. The `const` keyword specifies that the pointer cannot be modified after initialization; the pointer is protected from modification thereafter.

The `volatile` keyword specifies that the value associated with the name that follows can be modified by actions other than those in the user application. Therefore, the `volatile` keyword is useful for declaring objects in shared memory that can be accessed by multiple processes or global data areas used for communication with interrupt service routines.

When a name is declared as `volatile`, the compiler reloads the value from memory each time it is accessed by the program. This dramatically reduces the possible optimizations. However, when the state of an object can change unexpectedly, it is the only way to ensure predictable program performance.

To declare the object pointed to by the pointer as `const` or `volatile`, use a declaration of the form:

```
const char *cpch;
volatile char *vpch;
```

To declare the value of the pointer — that is, the actual address stored in the pointer — as `const` or `volatile`, use a declaration of the form:

```
char * const pchc;
char * volatile pchv;
```

The C++ language prevents assignments that would allow modification of an object or pointer declared as `const`. Such assignments would remove the information that the object or pointer was declared with, thereby violating the intent of the original declaration. Consider the following declarations:

```
const char cch = 'A';
char ch = 'B';
```

Given the preceding declarations of two objects (`cch`, of type `const char`, and `ch`, of type `char`), the following declaration/initializations are valid:

```
const char *pch1 = &cch;
const char *const pch4 = &cch;
const char *pch5 = &ch;
char *pch6 = &ch;
char *const pch7 = &ch;
const char *const pch8 = &ch;
```

The following declaration/initializations are erroneous.

```
char *pch2 = &cch; // Error
char *const pch3 = &cch; // Error
```

The declaration of `pch2` declares a pointer through which a constant object might be modified and is therefore

disallowed. The declaration of `pch3` specifies that the `pointer` is constant, not the object; the declaration is disallowed for the same reason the `pch2` declaration is disallowed.

The following eight assignments show assigning through pointer and changing of pointer value for the preceding declarations; for now, assume that the initialization was correct for `pch1` through `pch8`.

```
*pch1 = 'A'; // Error: object declared const
pch1 = &ch; // OK: pointer not declared const
*pch2 = 'A'; // OK: normal pointer
pch2 = &ch; // OK: normal pointer
*pch3 = 'A'; // OK: object not declared const
pch3 = &ch; // Error: pointer declared const
*pch4 = 'A'; // Error: object declared const
pch4 = &ch; // Error: pointer declared const
```

Pointers declared as `volatile`, or as a mixture of `const` and `volatile`, obey the same rules.

Pointers to `const` objects are often used in function declarations as follows:

```
errno_t strcpy_s( char *strDestination, size_t numberOfElements, const char *strSource );
```

The preceding statement declares a function, `strcpy_s`, where two of the three arguments are of type pointer to `char`. Because the arguments are passed by reference and not by value, the function would be free to modify both `strDestination` and `strSource` if `strSource` were not declared as `const`. The declaration of `strSource` as `const` assures the caller that `strSource` cannot be changed by the called function.

#### NOTE

Because there is a standard conversion from `typename \*` to `const typename \*`, it is legal to pass an argument of type `char \*` to `strcpy_s`. However, the reverse is not true; no implicit conversion exists to remove the `const` attribute from an object or pointer.

A `const` pointer of a given type can be assigned to a pointer of the same type. However, a pointer that is not `const` cannot be assigned to a `const` pointer. The following code shows correct and incorrect assignments:

```
// const_pointer.cpp
int *const cpObject = 0;
int *pObject;

int main() {
    pObject = cpObject;
    cpObject = pObject; // C3892
}
```

The following sample shows how to declare an object as `const` if you have a pointer to a pointer to an object.

```
// const_pointer2.cpp
struct X {
    X(int i) : m_i(i) { }
    int m_i;
};

int main() {
    // correct
    const X cx(10);
    const X * pcx = &cx;
    const X ** ppcx = &pcx;

    // also correct
    X const cx2(20);
    X const * pcx2 = &cx2;
    X const ** ppcx2 = &pcx2;
}
```

## See Also

[Pointers](#)

# new and delete Operators

12/21/2017 • 6 min to read • [Edit Online](#)

C++ supports dynamic allocation and deallocation of objects using the `new` and `delete` operators. These operators allocate memory for objects from a pool called the free store. The `new` operator calls the special function `operator new`, and the `delete` operator calls the special function `operator delete`.

The `new` function in the C++ Standard Library supports the behavior specified in the C++ standard, which is to throw a `std::bad_alloc` exception if the memory allocation fails. If you still want the non-throwing version of `new`, link your program with `nothrownew.obj`. However, when you link with `nothrownew.obj`, the default `operator new` in the C++ Standard Library no longer functions.

For a list of the library files that comprise the C Runtime Library and the C++ Standard Library, see [CRT Library Features](#).

## The new operator

When a statement such as the following is encountered in a program, it translates into a call to the function

`operator new :`

```
char *pch = new char[BUFFER_SIZE];
```

If the request is for zero bytes of storage, `operator new` returns a pointer to a distinct object (that is, repeated calls to `operator new` return different pointers). If there is insufficient memory for the allocation request, `operator new` throws a `std::bad_alloc` exception, or returns `nullptr` if you have linked in non-throwing `operator new` support.

You can write a routine that attempts to free memory and retry the allocation; see [\\_set\\_new\\_handler](#) for more information. For more details on the recovery scheme, see the Handling insufficient memory section of this topic.

The two scopes for `operator new` functions are described in the following table.

### Scope for operator new Functions

OPERATOR	SCOPE
<code>::operator new</code>	Global
<code>class-name ::operator new</code>	Class

The first argument to `operator new` must be of type `size_t` (a type defined in `STDDEF.H`), and the return type is always `void *`.

The global `operator new` function is called when the `new` operator is used to allocate objects of built-in types, objects of class type that do not contain user-defined `operator new` functions, and arrays of any type. When the `new` operator is used to allocate objects of a class type where an `operator new` is defined, that class's `operator new` is called.

An `operator new` function defined for a class is a static member function (which cannot, therefore, be virtual) that hides the global `operator new` function for objects of that class type. Consider the case where `new` is used to allocate and set memory to a given value:

```

// spec1_the_operator_new_function1.cpp
#include <malloc.h>
#include <memory.h>

class Blanks
{
public:
    Blanks(){}
    void *operator new( size_t stAllocateBlock, char chInit );
};

void *Blanks::operator new( size_t stAllocateBlock, char chInit )
{
    void *pvTemp = malloc( stAllocateBlock );
    if( pvTemp != 0 )
        memset( pvTemp, chInit, stAllocateBlock );
    return pvTemp;
}

// For discrete objects of type Blanks, the global operator new function
// is hidden. Therefore, the following code allocates an object of type
// Blanks and initializes it to 0xa5
int main()
{
    Blanks *a5 = new(0xa5) Blanks;
    return a5 != 0;
}

```

The argument supplied in parentheses to **new** is passed to `Blanks::operator new` as the `chInit` argument. However, the global **operator new** function is hidden, causing code such as the following to generate an error:

```
Blanks *SomeBlanks = new Blanks;
```

In Visual C++ 5.0 and earlier, nonclass types and all arrays (regardless of whether they were of **class** type) allocated using the **new** operator always used the global **operator new** function.

Beginning with Visual C++ 5.0, the compiler supports member array **new** and **delete** operators in a class declaration. For example:

```

// spec1_the_operator_new_function2.cpp
class MyClass
{
public:
    void * operator new[] (size_t)
    {
        return 0;
    }
    void operator delete[] (void*)
    {
    }
};

int main()
{
    MyClass *pMyClass = new MyClass[5];
    delete [] pMyClass;
}

```

## Handling insufficient memory

Testing for failed memory allocation can be done with code such as the following:

```
// insufficient_memory_conditions.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
#define BIG_NUMBER 100000000
int main() {
    int *pI = new int[BIG_NUMBER];
    if( pI == 0x0 ) {
        cout << "Insufficient memory" << endl;
        return -1;
    }
}
```

There is another ways to handle failed memory allocation requests: write a custom recovery routine to handle such a failure, then register your function by calling the [\\_set\\_new\\_handler](#) run-time function.

## The delete operator

Memory that is dynamically allocated using the **new** operator can be freed using the **delete** operator. The **delete** operator calls the **operator delete** function, which frees memory back to the available pool. Using the **delete** operator also causes the class destructor (if there is one) to be called.

There are global and class-scoped **operator delete** functions. Only one **operator delete** function can be defined for a given class; if defined, it hides the global **operator delete** function. The global **operator delete** function is always called for arrays of any type.

The global **operator delete** function. Two forms exist for the global **operator delete** and class-member **operator delete** functions:

```
void operator delete( void * );
void operator delete( void *, size_t );
```

Only one of the preceding two forms can be present for a given class. The first form takes a single argument of type **void \**, which contains a pointer to the object to deallocate. The second form—sized deallocation—takes two arguments, the first of which is a pointer to the memory block to deallocate and the second of which is the number of bytes to deallocate. The return type of both forms is **void** (**operator delete**\* cannot return a value).

The intent of the second form is to speed up searching for the correct size category of the object to be deleted, which is often not stored near the allocation itself and likely uncached; the second form is particularly useful when an **operator delete** function from a base class is used to delete an object of a derived class.

The **operator delete** function is static; therefore, it cannot be virtual. The **operator delete** function obeys access control, as described in [Member-Access Control](#).

The following example shows user-defined **operator new** and **operator delete** functions designed to log allocations and deallocations of memory:

```

// spec1_the_operator_delete_function1.cpp
// compile with: /EHsc
// arguments: 3
#include <iostream>
using namespace std;

int fLogMemory = 0;      // Perform logging (0=no; nonzero=yes)?
int cBlocksAllocated = 0; // Count of blocks allocated.

// User-defined operator new.
void *operator new( size_t stAllocateBlock ) {
    static int fInOpNew = 0; // Guard flag.

    if ( fLogMemory && !fInOpNew ) {
        fInOpNew = 1;
        clog << "Memory block " << ++cBlocksAllocated
            << " allocated for " << stAllocateBlock
            << " bytes\n";
        fInOpNew = 0;
    }
    return malloc( stAllocateBlock );
}

// User-defined operator delete.
void operator delete( void *pvMem ) {
    static int fInOpDelete = 0; // Guard flag.
    if ( fLogMemory && !fInOpDelete ) {
        fInOpDelete = 1;
        clog << "Memory block " << cBlocksAllocated--
            << " deallocated\n";
        fInOpDelete = 0;
    }

    free( pvMem );
}

int main( int argc, char *argv[] ) {
    fLogMemory = 1; // Turn logging on
    if( argc > 1 )
        for( int i = 0; i < atoi( argv[1] ); ++i ) {
            char *pMem = new char[10];
            delete[] pMem;
        }
    fLogMemory = 0; // Turn logging off.
    return cBlocksAllocated;
}

```

The preceding code can be used to detect "memory leakage" — that is, memory that is allocated on the free store but never freed. To perform this detection, the global **new** and **delete** operators are redefined to count allocation and deallocation of memory.

Beginning with Visual C++ 5.0, the compiler supports member array **new** and **delete** operators in a class declaration. For example:

```
// spec1_the_operator_delete_function2.cpp
// compile with: /c
class X {
public:
    void * operator new[] (size_t) {
        return 0;
    }
    void operator delete[] (void*) {}
};

void f() {
    X *pX = new X[5];
    delete [] pX;
}
```

# Exception Handling in Visual C++

12/21/2017 • 1 min to read • [Edit Online](#)

An exception is an error condition, possibly outside the program's control, that prevents the program from continuing along its regular execution path. Certain operations, including object creation, file input/output, and function calls made from other modules, are all potential sources of exceptions even when your program is running correctly. Robust code anticipates and handles exceptions.

To detect logic errors within a single program or module, use assertions rather than exceptions (see [Using Assertions](#)).

Visual C++ supports three kinds of exception handling:

- [C++ exception handling](#)

For most C++ programs, you should use C++ exception handling, which is type-safe and ensures that object destructors are invoked during stack unwinding.

- [Structured exception handling](#)

Windows supplies its own exception mechanism, called SEH. It is not recommended for C++ or MFC programming. Use SEH only in non-MFC C programs.

- [MFC exceptions](#)

Since version 3.0, MFC has used C++ exceptions but still supports its older exception handling macros, which are similar to C++ exceptions in form. Although these macros are not recommended for new programming, they are still supported for backward compatibility. In programs that already use the macros, you can freely use C++ exceptions as well. During preprocessing, the macros evaluate to the exception handling keywords defined in the Visual C++ implementation of the C++ language as of Visual C++ version 2.0. You can leave existing exception macros in place while you begin to use C++ exceptions.

Use the [/EH](#) compiler option to specify the type of exception handling to use in a project; C++ exception handling is the default. Do not mix the error handling mechanisms; for example, do not use C++ exceptions with structured exception handling. Using C++ exception handling makes your code more portable, and it allows you to handle exceptions of any type. For more information about the drawbacks of structured exception handling, see [Structured Exception Handling](#). For advice about mixing MFC macros and C++ exceptions, see [Exceptions: Using MFC Macros and C++ Exceptions](#).

For information on handling exceptions in CLR applications, see [Exception Handling](#).

For information about exception handling on x64 processors, see [Exception Handling \(x64\)](#).

## See Also

[C++ Language Reference](#)

# C++ Exception Handling

12/21/2017 • 1 min to read • [Edit Online](#)

The C++ language provides built-in support for throwing and catching exceptions. When programming in C++, you should almost always use the built-in C++ exception support as described in this section.

To enable C++ exception handling in your code, use [/EHsc](#).

## In This Section

This discussion on C++ exception handling includes:

- [The try, catch, and throw Statements](#)
- [How Catch Blocks are Evaluated](#)
- [Exceptions and Stack Unwinding](#)
- [Exception Specifications](#)
- [noexcept](#)
- [Unhandled C++ Exceptions](#)
- [Mixing C \(Structured\) and C++ Exceptions](#)

## Support for Earlier MFC Exceptions

As of version 4.0, MFC began using the C++ exception handling mechanism. Although you are encouraged to use C++ exception handling in new code, MFC version 4.0 and later retains the macros from previous versions of MFC so that old code will not be broken. The macros and the new mechanism can be combined as well. For information on mixing macros and C++ exception handling and on converting old code to use the new mechanism, see the articles [Exceptions: Using MFC Macros and C++ Exceptions](#) and [Exceptions: Converting from MFC Exception Macros](#). The older MFC exception macros, if you still use them, evaluate to C++ exception keywords. See [Exceptions: Changes to Exception Macros in Version 3.0](#).

## See Also

[Exception Handling](#)

# try, throw, and catch Statements (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

To implement exception handling in C++, you use `try`, `throw`, and `catch` expressions.

First, use a `try` block to enclose one or more statements that might throw an exception.

A `throw` expression signals that an exceptional condition—often, an error—has occurred in a `try` block. You can use an object of any type as the operand of a `throw` expression. Typically, this object is used to communicate information about the error. In most cases, we recommend that you use the `std::exception` class or one of the derived classes that are defined in the standard library. If one of those is not appropriate, we recommend that you derive your own exception class from `std::exception`.

To handle exceptions that may be thrown, implement one or more `catch` blocks immediately following a `try` block. Each `catch` block specifies the type of exception it can handle.

This example shows a `try` block and its handlers. Assume that `GetNetworkResource()` acquires data over a network connection and that the two exception types are user-defined classes that derive from `std::exception`. Notice that the exceptions are caught by `const` reference in the `catch` statement. We recommend that you throw exceptions by value and catch them by `const` reference.

## Example

```
MyData md;
try {
    // Code that could throw an exception
    md = GetNetworkResource();
}
catch (const networkIOException& e) {
    // Code that executes when an exception of type
    // networkIOException is thrown in the try block
    // ...
    // Log error message in the exception object
    cerr << e.what();
}
catch (const myDataFormatException& e) {
    // Code that handles another exception type
    // ...
    cerr << e.what();
}

// The following syntax shows a throw expression
MyData GetNetworkResource()
{
    // ...
    if (IOSuccess == false)
        throw networkIOException("Unable to connect");
    // ...
    if (readError)
        throw myDataFormatException("Format error");
    // ...
}
```

## Remarks

The code after the `try` clause is the guarded section of code. The `throw` expression *throws*—that is, raises—an exception. The code block after the `catch` clause is the exception handler. This is the handler that *catches* the exception that's thrown if the types in the `throw` and `catch` expressions are compatible. For a list of rules that govern type-matching in `catch` blocks, see [How Catch Blocks are Evaluated](#). If the `catch` statement specifies an ellipsis (...) instead of a type, the `catch` block handles every type of exception. When you compile with the /EHc option, these can include C structured exceptions and system-generated or application-generated asynchronous exceptions such as memory protection, divide-by-zero, and floating-point violations. Because `catch` blocks are processed in program order to find a matching type, an ellipsis handler must be the last handler for the associated `try` block. Use `catch(...)` with caution; do not allow a program to continue unless the catch block knows how to handle the specific exception that is caught. Typically, a `catch(...)` block is used to log errors and perform special cleanup before program execution is stopped.

A `throw` expression that has no operand re-throws the exception currently being handled. We recommend this form when re-throwing the exception, because this preserves the original exception's polymorphic type information. Such an expression should only be used in a `catch` handler or in a function that's called from a `catch` handler. The re-thrown exception object is the original exception object, not a copy.

```
try {
    throw CSomeOtherException();
}
catch(...) {
    // Catch all exceptions - dangerous!!!
    // Respond (perhaps only partially) to the exception, then
    // re-throw to pass the exception to some other handler
    // ...
    throw;
}
```

## See Also

[C++ Exception Handling](#)

[Keywords](#)

[Unhandled C++ Exceptions](#)

[\\_\\_uncaught\\_exception](#)

# How Catch Blocks are Evaluated (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

C++ enables you to throw exceptions of any type, although in general it is recommended to throw types that are derived from std::exception. A C++ exception can be caught by a **catch** handler that specifies the same type as the thrown exception, or by a handler that can catch any type of exception.

If the type of thrown exception is a class, which also has a base class (or classes), it can be caught by handlers that accept base classes of the exception's type, as well as references to bases of the exception's type. Note that when an exception is caught by a reference, it is bound to the actual thrown exception object; otherwise, it is a copy (much the same as an argument to a function).

When an exception is thrown, it may be caught by the following types of **catch** handlers:

- A handler that can accept any type (using the ellipsis syntax).
- A handler that accepts the same type as the exception object; because it is a copy, **const** and **volatile** modifiers are ignored.
- A handler that accepts a reference to the same type as the exception object.
- A handler that accepts a reference to a **const** or **volatile** form of the same type as the exception object.
- A handler that accepts a base class of the same type as the exception object; since it is a copy, **const** and **volatile** modifiers are ignored. The **catch** handler for a base class must not precede the **catch** handler for the derived class.
- A handler that accepts a reference to a base class of the same type as the exception object.
- A handler that accepts a reference to a **const** or **volatile** form of a base class of the same type as the exception object.
- A handler that accepts a pointer to which a thrown pointer object can be converted via standard pointer conversion rules.

The order in which **catch** handlers appear is significant, because handlers for a given **try** block are examined in order of their appearance. For example, it is an error to place the handler for a base class before the handler for a derived class. After a matching **catch** handler is found, subsequent handlers are not examined. As a result, an ellipsis **catch** handler must be the last handler for its **try** block. For example:

```
// ...
try
{
    // ...
}
catch( ... )
{
    // Handle exception here.
}
// Error: the next two handlers are never examined.
catch( const char * str )
{
    cout << "Caught exception: " << str << endl;
}
catch( CExcptClass E )
{
    // Handle CExcptClass exception here.
}
```

In this example, the ellipsis **catch** handler is the only handler that is examined.

## See Also

[C++ Exception Handling](#)

# Exceptions and Stack Unwinding in C++

12/21/2017 • 4 min to read • [Edit Online](#)

In the C++ exception mechanism, control moves from the throw statement to the first catch statement that can handle the thrown type. When the catch statement is reached, all of the automatic variables that are in scope between the throw and catch statements are destroyed in a process that is known as *stack unwinding*. In stack unwinding, execution proceeds as follows:

1. Control reaches the `try` statement by normal sequential execution. The guarded section in the `try` block is executed.
2. If no exception is thrown during execution of the guarded section, the `catch` clauses that follow the `try` block are not executed. Execution continues at the statement after the last `catch` clause that follows the associated `try` block.
3. If an exception is thrown during execution of the guarded section or in any routine that the guarded section calls either directly or indirectly, an exception object is created from the object that is created by the `throw` operand. (This implies that a copy constructor may be involved.) At this point, the compiler looks for a `catch` clause in a higher execution context that can handle an exception of the type that is thrown, or for a `catch` handler that can handle any type of exception. The `catch` handlers are examined in order of their appearance after the `try` block. If no appropriate handler is found, the next dynamically enclosing `try` block is examined. This process continues until the outermost enclosing `try` block is examined.
4. If a matching handler is still not found, or if an exception occurs during the unwinding process but before the handler gets control, the predefined run-time function `terminate` is called. If an exception occurs after the exception is thrown but before the unwind begins, `terminate` is called.
5. If a matching `catch` handler is found, and it catches by value, its formal parameter is initialized by copying the exception object. If it catches by reference, the parameter is initialized to refer to the exception object. After the formal parameter is initialized, the process of unwinding the stack begins. This involves the destruction of all automatic objects that were fully constructed—but not yet destructed—between the beginning of the `try` block that is associated with the `catch` handler and the throw site of the exception. Destruction occurs in reverse order of construction. The `catch` handler is executed and the program resumes execution after the last handler—that is, at the first statement or construct that is not a `catch` handler. Control can only enter a `catch` handler through a thrown exception, never through a `goto` statement or a `case` label in a `switch` statement.

## Stack Unwinding Example

The following example demonstrates how the stack is unwound when an exception is thrown. Execution on the thread jumps from the throw statement in `c` to the catch statement in `main`, and unwinds each function along the way. Notice the order in which the `Dummy` objects are created and then destroyed as they go out of scope. Also notice that no function completes except `main`, which contains the catch statement. Function `A` never returns from its call to `B()`, and `B` never returns from its call to `C()`. If you uncomment the definition of the `Dummy` pointer and the corresponding delete statement, and then run the program, notice that the pointer is never deleted. This shows what can happen when functions do not provide an exception guarantee. For more information, see How to: Design for Exceptions. If you comment out the catch statement, you can observe what happens when a program terminates because of an unhandled exception.

```
#include <string>
```

```

#include <iostream>
using namespace std;

class MyException{};
class Dummy
{
public:
    Dummy(string s) : MyName(s) { PrintMsg("Created Dummy:"); }
    Dummy(const Dummy& other) : MyName(other.MyName){ PrintMsg("Copy created Dummy:"); }
    ~Dummy(){ PrintMsg("Destroyed Dummy:"); }
    void PrintMsg(string s) { cout << s << endl; }
    string MyName;
    int level;
};

void C(Dummy d, int i)
{
    cout << "Entering FunctionC" << endl;
    d.MyName = " C";
    throw MyException();

    cout << "Exiting FunctionC" << endl;
}

void B(Dummy d, int i)
{
    cout << "Entering FunctionB" << endl;
    d.MyName = "B";
    C(d, i + 1);
    cout << "Exiting FunctionB" << endl;
}

void A(Dummy d, int i)
{
    cout << "Entering FunctionA" << endl;
    d.MyName = " A" ;
    // Dummy* pd = new Dummy("new Dummy"); //Not exception safe!!!
    B(d, i + 1);
    // delete pd;
    cout << "Exiting FunctionA" << endl;
}

int main()
{
    cout << "Entering main" << endl;
    try
    {
        Dummy d(" M");
        A(d,1);
    }
    catch (MyException& e)
    {
        cout << "Caught an exception of type: " << typeid(e).name() << endl;
    }

    cout << "Exiting main." << endl;
    char c;
    cin >> c;
}

/* Output:
   Entering main
   Created Dummy: M
   Copy created Dummy: M
   Entering FunctionA
   Copy created Dummy: A
   Entering FunctionB
   Copy created Dummy: B
   Entering FunctionC

```

```
Destroyed Dummy: C
Destroyed Dummy: B
Destroyed Dummy: A
Destroyed Dummy: M
Caught an exception of type: class MyException
Exiting main.
```

```
*/
```

# Exception Specifications (throw, noexcept) (C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Exception specifications are a C++ language feature that indicate the programmer's intent about the exception types that can be propagated by a function. You can specify that a function may or may not exit by an exception by using an *exception specification*. The compiler can use this information to optimize calls to the function, and to terminate the program if an unexpected exception escapes the function. There are two kinds of exception specification. The *noexcept specification* is new in C++11. It specifies whether the set of potential exceptions that can escape the function is empty. The *dynamic exception specification*, or `throw(optional_type_list)` specification, is deprecated in C++11 and is only partially supported by Visual Studio. This exception specification was designed to provide summary information about what exceptions can be thrown out of a function, but in practice it was found to be problematic. The one dynamic exception specification that did prove to be somewhat useful was the unconditional `throw()` specification. For example, the function declaration:

```
void MyFunction(int i) throw();
```

tells the compiler that the function does not throw any exceptions. It is the equivalent to using `__declspec(nothrow)`. Its use is considered optional.

In the ISO C++11 Standard, the `noexcept` operator was introduced as a replacement. It is supported in Visual Studio 2015 and later. Whenever possible, use a `noexcept` expression to specify whether a function might throw exceptions. For example, use this function declaration instead of the one above:

```
void MyFunction(int i) noexcept;
```

While Visual C++ fully supports the `noexcept` expression, it departs from the ISO C++ Standard in its implementation of dynamic exception specifications. The following table summarizes the Visual C++ implementation of exception specifications:

EXCEPTION SPECIFICATION	MEANING
<code>noexcept</code> <code>noexcept(true)</code> <code>throw()</code>	The function does not throw an exception. However, if an exception is thrown out of a function marked <code>throw()</code> , the Visual C++ compiler calls <code>std::terminate</code> , not <code>std::unexpected</code> . See <a href="#">std::unexpected</a> for more information. If a function is marked <code>noexcept</code> , <code>noexcept(true)</code> , or <code>throw()</code> , the Visual C++ compiler assumes that the function does not throw C++ exceptions and generates code accordingly. Because code optimizations might be performed by the C++ compiler based on the assumption that the function does not throw any C++ exceptions, if a function does throw an exception, the program may not execute correctly.
<code>noexcept(false)</code> <code>throw(...)</code> No specification	The function can throw an exception of any type.

EXCEPTION SPECIFICATION	MEANING
<code>throw(type)</code>	The function can throw an exception of type <code>type</code> . In Visual C++, this syntax is accepted, but it is interpreted as <code>noexcept(false)</code> .

If exception handling is used in an application, there must be a function in the call stack that handles thrown exceptions before they exit the outer scope of a function marked `noexcept`, `noexcept(true)`, or `throw()`. If any functions called between the one that throws an exception and the one that handles the exception are specified as `noexcept`, `noexcept(true)`, or `throw()`, the program is terminated when the noexcept function propagates the exception.

The exception behavior of a function depends on the following factors:

- Whether you are compiling the function under C or C++.
- Which `/EH` compiler option you use.
- Whether you explicitly specify the exception specification.

Explicit exception specifications are not allowed on C functions. A C function is assumed not to throw exceptions under `/EHsc`, and may throw structured exceptions under `/EHs`, `/EHa`, or `/EHac`.

The following table summarizes whether a C++ function may potentially throw under various compiler exception handling options:

FUNCTION	/EHSC	/EHS	/EHA	/EHAC
C++ function with no exception specification	Yes	Yes	Yes	Yes
C++ function with <code>noexcept</code> , <code>noexcept(true)</code> , or <code>throw()</code> exception specification	No	No	Yes	Yes
C++ function with <code>noexcept(false)</code> , <code>throw(...)</code> , or <code>throw(type)</code> exception specification	Yes	Yes	Yes	Yes

## Example

```

// exception_specification.cpp
// compile with: /EHs
#include <stdio.h>

void handler() {
    printf_s("in handler\n");
}

void f1(void) throw(int) {
    printf_s("About to throw 1\n");
    if (1)
        throw 1;
}

void f5(void) throw() {
    try {
        f1();
    }
    catch(...) {
        handler();
    }
}

// invalid, doesn't handle the int exception thrown from f1()
// void f3(void) throw() {
//     f1();
// }

void __declspec(nothrow) f2(void) {
    try {
        f1();
    }
    catch(int) {
        handler();
    }
}

// only valid if compiled without /EHc
// /EHc means assume extern "C" functions don't throw exceptions
extern "C" void f4(void);
void f4(void) {
    f1();
}

int main() {
    f2();

    try {
        f4();
    }
    catch(...) {
        printf_s("Caught exception from f4\n");
    }
    f5();
}

```

```

About to throw 1
in handler
About to throw 1
Caught exception from f4
About to throw 1
in handler

```

## See Also

[try, throw, and catch Statements \(C++\)](#)

[C++ Exception Handling](#)

# noexcept (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

**C++11:** Specifies whether a function might throw exceptions.

## Syntax

*noexcept-expression:*

**noexcept**  
**noexcept(** *constant-expression* **)**

## Parameters

*constant-expression*

A constant expression of type `bool` that represents whether the set of potential exception types is empty. The unconditional version is equivalent to `noexcept(true)`.

## Remarks

A *noexcept expression* is a kind of *exception specification*, a suffix to a function declaration that represents a set of types that might be matched by an exception handler for any exception that exits a function. Unary conditional operator `noexcept(` *constant\_expression* `)` where *constant\_expression* yields `true`, and its unconditional synonym `noexcept`, specify that the set of potential exception types that can exit a function is empty. That is, the function never throws an exception and never allows an exception to be propagated outside its scope. The operator `noexcept(` *constant\_expression* `)` where *constant\_expression* yields `false`, or the absence of an exception specification (other than for a destructor or deallocation function), indicates that the set of potential exceptions that can exit the function is the set of all types.

Mark a function as `noexcept` only if all the functions that it calls, either directly or indirectly, are also `noexcept` or `const`. The compiler does not necessarily check every code path for exceptions that might bubble up to a `noexcept` function. If an exception does exit the outer scope of a function marked `noexcept`, `std::terminate` is invoked immediately, and there is no guarantee that destructors of any in-scope objects will be invoked. Use `noexcept` instead of the dynamic exception specifier `throw`, which is deprecated in C++11 and later and not fully implemented in Visual Studio. We recommended you apply `noexcept` to any function that never allows an exception to propagate up the call stack. When a function is declared `noexcept`, it enables the compiler to generate more efficient code in several different contexts.

## Example

A template function that copies its argument might be declared `noexcept` on the condition that the object being copied is a plain old data type (POD). Such a function could be declared like this:

```
#include <type_traits>

template <typename T>
T copy_object(const T& obj) noexcept(std::is_pod<T>)
{
    // ...
}
```

## See Also

[C++ Exception Handling Exception Specifications \(throw, noexcept\)](#)

# Unhandled C++ Exceptions

12/21/2017 • 1 min to read • [Edit Online](#)

If a matching handler (or ellipsis **catch** handler) cannot be found for the current exception, the predefined `terminate` run-time function is called. (You can also explicitly call `terminate` in any of your handlers.) The default action of `terminate` is to call `abort`. If you want `terminate` to call some other function in your program before exiting the application, call the `set_terminate` function with the name of the function to be called as its single argument. You can call `set_terminate` at any point in your program. The `terminate` routine always calls the last function given as an argument to `set_terminate`.

## Example

The following example throws a `char *` exception, but does not contain a handler designated to catch exceptions of type `char *`. The call to `set_terminate` instructs `terminate` to call `term_func`.

```
// exceptions_Unhandled_Exceptions.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;
void term_func() {
    cout << "term_func was called by terminate." << endl;
    exit( -1 );
}
int main() {
    try
    {
        set_terminate( term_func );
        throw "Out of memory!" // No catch handler for this exception
    }
    catch( int )
    {
        cout << "Integer exception raised." << endl;
    }
    return 0;
}
```

## Output

```
term_func was called by terminate.
```

The `term_func` function should terminate the program or current thread, ideally by calling `exit`. If it doesn't, and instead returns to its caller, `abort` is called.

## See Also

[C++ Exception Handling](#)

# Mixing C (Structured) and C++ Exceptions

12/21/2017 • 1 min to read • [Edit Online](#)

If you want to write more portable code, using structured exception handling in a C++ program is not recommended. However, you may sometimes want to compile with **/EHs** and mix structured exceptions and C++ source code, and need some facility for handling both kinds of exceptions. Because a structured exception handler has no concept of objects or typed exceptions, it cannot handle exceptions thrown by C++ code; however, C++ **catch** handlers can handle structured exceptions. As such, C++ exception handling syntax (**try**, **throw**, **catch**) is not accepted by the C compiler, but structured exception handling syntax (**\_try**, **\_except**, **\_finally**) is supported by the C++ compiler.

See [\\_set\\_se\\_translator](#) for information on handling structured exceptions as C++ exceptions.

If you mix structured and C++ exceptions, note the following:

1. C++ exceptions and structured exceptions cannot be mixed within the same function.
2. Termination handlers (**\_finally** blocks) are always executed, even during unwinding after an exception is thrown.
3. C++ exception handling can catch and preserve unwind semantics in all modules compiled with the **/EH** compiler option (this option enables unwind semantics).
4. There may be some situations in which destructor functions are not called for all objects. For example, if a structured exception occurs while attempting to make a function call through an uninitialized function pointer, and that function takes as parameters objects that were constructed before the call, those objects will not have their destructors called during stack unwind.

## What do you want to know more about?

- [Using setjmp or longjmp in C++ programs](#)
- [Differences between SEH and C++ EH](#)

## See Also

[C++ Exception Handling](#)

# Using setjmp/longjmp

12/21/2017 • 1 min to read • [Edit Online](#)

When `setjmp` and `longjmp` are used together, they provide a way to execute a non-local `goto`. They are typically used to pass execution control to error-handling or recovery code in a previously called routine without using the standard calling or return conventions.

**Caution**

However, because `setjmp` and `longjmp` do not support C++ object semantics, and because they might degrade performance by preventing optimization on local variables, we recommend that you do not use them in C++ programs. We recommend that you use `try / catch` constructs instead.

If you decide to use `setjmp / longjmp` in a C++ program, also include `SETJMP.H` or `SETJMPEX.H` to assure correct interaction between the functions and C++ exception handling. If you use `/EH` to compile, destructors for local objects are called during the stack unwind. If you use `/EHs` to compile, and one of your functions calls a function that uses `nothrow` and the function that uses `nothrow` calls `longjmp`, then the destructor unwind might not occur, depending on the optimizer.

In portable code, when a non-local `goto` that calls `longjmp` is executed, correct destruction of frame-based objects might be unreliable.

## See Also

[Mixing C \(Structured\) and C++ Exceptions](#)

# Exception Handling Differences

12/21/2017 • 4 min to read • [Edit Online](#)

The major difference between structured exception handling and C++ exception handling is that the C++ exception handling model deals in types, while the C structured exception handling model deals with exceptions of one type — specifically, `unsigned int`. That is, C exceptions are identified by an unsigned integer value, whereas C++ exceptions are identified by data type. When an exception is raised in C, each possible handler executes a filter that examines the C exception context and determines whether to accept the exception, pass it to some other handler, or ignore it. When an exception is thrown in C++, it may be of any type.

A second difference is that the C structured exception handling model is referred to as "asynchronous" in that exceptions occur secondary to the normal flow of control. The C++ exception handling mechanism is fully "synchronous," which means that exceptions occur only when they are thrown.

If a C exception is raised in a C++ program, it can be handled by a structured exception handler with its associated filter or by a C++ **catch** handler, whichever is dynamically nearer to the exception context. For example, the following C++ program raises a C exception inside a C++ **try** context:

## Example

```
// exceptions_Exception_Handling_Differences.cpp
// compile with: /EHs
#include <iostream>

using namespace std;
void SEHFunc( void );

int main() {
    try {
        SEHFunc();
    }
    catch( ... ) {
        cout << "Caught a C exception." << endl;
    }
}

void SEHFunc() {
    __try {
        int x, y = 0;
        x = 5 / y;
    }
    __finally {
        cout << "In finally." << endl;
    }
}
```

```
In finally.
Caught a C exception.
```

## C Exception Wrapper Class

In a simple example like the above, the C exception can be caught only by an ellipsis (...) **catch** handler. No information about the type or nature of the exception is communicated to the handler. While this method works, in

some cases you may need to define a transformation between the two exception handling models so that each C exception is associated with a specific class. To do this, you can define a C exception "wrapper" class, which can be used or derived from in order to attribute a specific class type to a C exception. By doing so, each C exception can be handled by a C++ **catch** handler more separately than in the preceding example.

Your wrapper class might have an interface consisting of some member functions that determine the value of the exception, and that access the extended exception context information provided by the C exception model. You might also want to define a default constructor and a constructor that accepts an `unsigned int` argument (to provide for the underlying C exception representation), and a bitwise copy constructor. The following is a possible implementation of a C exception wrapper class:

```
// exceptions_Exception_Handling_Differences2.cpp
// compile with: /c
class SE_Exception {
private:
    SE_Exception() {}
    SE_Exception( SE_Exception& ) {}
    unsigned int nSE;
public:
    SE_Exception( unsigned int n ) : nSE( n ) {}
    ~SE_Exception() {}
    unsigned int getSeNumber() {
        return nSE;
    }
};
```

To use this class, you install a custom C exception translation function that is called by the internal exception handling mechanism each time a C exception is thrown. Within your translation function, you can throw any typed exception (perhaps an `SE_Exception` type, or a class type derived from `SE_Exception`) that can be caught by an appropriate matching C++ **catch** handler. The translation function can simply return, which indicates that it did not handle the exception. If the translation function itself raises a C exception, `terminate` is called.

To specify a custom translation function, call the `_set_se_translator` function with the name of your translation function as its single argument. The translation function that you write is called once for each function invocation on the stack that has **try** blocks. There is no default translation function; if you do not specify one by calling `_set_se_translator`, the C exception can only be caught by an ellipsis **catch** handler.

## Example

For example, the following code installs a custom translation function, and then raises a C exception that is wrapped by the `SE_Exception` class:

```

// exceptions_Exception_Handling_Differences3.cpp
// compile with: /EHs
#include <stdio.h>
#include <eh.h>
#include <windows.h>

class SE_Exception {
private:
    SE_Exception() {}
    unsigned int nSE;

public:
    SE_Exception( SE_Exception& e ) : nSE(e.nSE) {}
    SE_Exception(unsigned int n) : nSE(n) {}
    ~SE_Exception() {}
    unsigned int getSeNumber() { return nSE; }
};

void SEFunc() {
    __try {
        int x, y = 0;
        x = 5 / y;
    }
    __finally {
        printf_s( "In finally\n" );
    }
}

void trans_func( unsigned int u, _EXCEPTION_POINTERS* pExp ) {
    printf_s( "In trans_func.\n" );
    throw SE_Exception( u );
}

int main() {
    _set_se_translator( trans_func );
    try {
        SEFunc();
    }
    catch( SE_Exception e ) {
        printf_s( "Caught a __try exception with SE_Exception.\n" );
        printf_s( "nSE = 0x%08x\n", e.getSeNumber() );
    }
}

```

```

In trans_func.
In finally
Caught a __try exception with SE_Exception.
nSE = 0xc0000094

```

## See Also

[Mixing C \(Structured\) and C++ Exceptions](#)

# Structured Exception Handling (C/C++)

12/21/2017 • 3 min to read • [Edit Online](#)

Although Windows and Visual C++ support structured exception handling (SEH), we recommend that you use ISO-standard C++ exception handling because it makes code more portable and flexible. Nevertheless, in existing code or for particular kinds of programs, you still might have to use SEH.

## Microsoft Specific

### Grammar

*try-except-statement* :

`__try compound-statement`

`__except ( expression ) compound-statement`

## Remarks

With SEH, you can ensure that resources such as memory blocks and files are correctly released if execution unexpectedly terminates. You can also handle specific problems—for example, insufficient memory—by using concise structured code that does not rely on `goto` statements or elaborate testing of return codes.

The try-except and try-finally statements referred to in this article are Microsoft extensions to the C language. They support SEH by enabling applications to gain control of a program after events that would otherwise terminate execution. Although SEH works with C++ source files, it's not specifically designed for C++. If you use SEH in a C++ program that you compile by using the `/EH` option—together with certain modifiers—destructors for local objects are called but other execution behavior might not be what you expect. (For an illustration, see the example later in this article.) In most cases, instead of SEH we recommend that you use ISO-standard [C++ exception handling](#), which Visual C++ also supports. By using C++ exception handling, you can ensure that your code is more portable, and you can handle exceptions of any type.

If you have C modules that use SEH, you can mix them with C++ modules that use C++ exception handling. For information, see [Exception Handling Differences](#).

There are two SEH mechanisms:

- [Exception handlers](#), which can respond to or dismiss the exception.
- [Termination handlers](#), which are called when an exception causes termination in a block of code.

These two kinds of handlers are distinct, but are closely related through a process known as "unwinding the stack." When an exception occurs, Windows looks for the most recently installed exception handler that is currently active. The handler can do one of three things:

- Fail to recognize the exception and pass control to other handlers.
- Recognize the exception but dismiss it.
- Recognize the exception and handle it.

The exception handler that recognizes the exception may not be in the function that was running when the exception occurred. In some cases, it may be in a function much higher on the stack. The currently running function and all other functions on the stack frame are terminated. During this process, the stack

is "unwound;" that is, local variables of terminated functions—unless they are `static`—are cleared from the stack.

As it unwinds the stack, the operating system calls any termination handlers that you've written for each function. By using a termination handler, you can clean up resources that otherwise would remain open because of an abnormal termination. If you've entered a critical section, you can exit in the termination handler. If the program is going to shut down, you can perform other housekeeping tasks such as closing and removing temporary files.

For more information, see:

- [Writing an Exception Handler](#)
- [Writing a Termination Handler](#)
- [Using Structured Exception Handling with C++](#)

## Example

As stated earlier, destructors for local objects are called if you use SEH in a C++ program and compile it by using the **/EH** option with certain modifiers—for example, **/EHsc** and **/EHfa**. However, the behavior during execution may not be what you expect if you are also using C++ exceptions. The following example demonstrates these behavioral differences.

```

#include <stdio.h>
#include <Windows.h>
#include <exception>

class TestClass
{
public:
    ~TestClass()
    {
        printf("Destroying TestClass!\r\n");
    }
};

__declspec(noinline) void TestCPPEX()
{
#ifdef CPPEX
    printf("Throwing C++ exception\r\n");
    throw std::exception("");
#else
    printf("Triggering SEH exception\r\n");
    volatile int *pInt = 0x00000000;
    *pInt = 20;
#endif
}

__declspec(noinline) void TestExceptions()
{
    TestClass d;
    TestCPPEX();
}

int main()
{
    __try
    {
        TestExceptions();
    }
    __except(EXCEPTION_EXECUTE_HANDLER)
    {
        printf("Executing SEH __except block\r\n");
    }

    return 0;
}

```

If you use **/EHsc** to compile this code but the local test control **CPPEX** is undefined, there is no execution of the **TestClass** destructor and the output looks like this:

```

Triggering SEH exception
Executing SEH __except block

```

If you use **/EHsc** to compile the code and **CPPEX** is defined by using **/DCPPEX** (so that a C++ exception is thrown), the **TestClass** destructor executes and the output looks like this:

```

Throwing C++ exception
Destroying TestClass!
Executing SEH __except block

```

If you use **/EHs** to compile the code, the **TestClass** destructor executes regardless of whether the exception was thrown by using **std::throw** or by using SEH to trigger the exception (**CPPEX** defined or not). The output looks like this:

```
Throwing C++ exception
Destroying TestClass!
Executing SEH __except block
```

For more information, see [/EH \(Exception Handling Model\)](#).

**END Microsoft Specific**

## See Also

[Exception Handling](#)

[Keywords](#)

[<exception>](#)

[Errors and Exception Handling](#)

[Structured Exception Handling \(Windows\)](#)

# Writing an Exception Handler

12/21/2017 • 1 min to read • [Edit Online](#)

Exception handlers are typically used to respond to specific errors. You can use the exception-handling syntax to filter out all exceptions other than those you know how to handle. Other exceptions should be passed to other handlers (possibly in the run-time library or the operating system) written to look for those specific exceptions.

Exception handlers use the try-except statement.

## What do you want to know more about?

- [The try-except statement](#)
- [Writing an exception filter](#)
- [Raising software exceptions](#)
- [Hardware exceptions](#)
- [Restrictions on exception handlers](#)

## See Also

[Structured Exception Handling \(C/C++\)](#)

# try-except Statement

12/21/2017 • 4 min to read • [Edit Online](#)

## Microsoft Specific

The **try-except** statement is a Microsoft extension to the C and C++ languages that supports structured exception handling.

## Syntax

```
_try
{
// guarded code
}
_except ( expression )
{
// exception handler code
}
```

## Remarks

The **try-except** statement is a Microsoft extension to the C and C++ languages that enables target applications to gain control when events that normally terminate program execution occur. Such events are called *exceptions*, and the mechanism that deals with exceptions is called *structured exception handling* (SEH).

For related information, see the [try-finally statement](#).

Exceptions can be either hardware-based or software-based. Even when applications cannot completely recover from hardware or software exceptions, structured exception handling makes it possible to display error information and trap the internal state of the application to help diagnose the problem. This is especially useful for intermittent problems that cannot be reproduced easily.

### NOTE

Structured exception handling works with Win32 for both C and C++ source files. However, it is not specifically designed for C++. You can ensure that your code is more portable by using C++ exception handling. Also, C++ exception handling is more flexible, in that it can handle exceptions of any type. For C++ programs, it is recommended that you use the C++ exception-handling mechanism ([try, catch, and throw statements](#)).

The compound statement after the `_try` clause is the body or guarded section. The compound statement after the `_except` clause is the exception handler. The handler specifies a set of actions to be taken if an exception is raised during execution of the body of the guarded section. Execution proceeds as follows:

1. The guarded section is executed.
2. If no exception occurs during execution of the guarded section, execution continues at the statement after the `_except` clause.
3. If an exception occurs during execution of the guarded section or in any routine the guarded section calls, the `_except expression` (called the *filter expression*) is evaluated and the value determines how the exception is handled. There are three values:

**EXCEPTION\_CONTINUE\_EXECUTION (-1)** Exception is dismissed. Continue execution at the point where the exception occurred.

**EXCEPTION\_CONTINUE\_SEARCH (0)** Exception is not recognized. Continue to search up the stack for a handler, first for containing **try-except** statements, then for handlers with the next highest precedence.

**EXCEPTION\_EXECUTE\_HANDLER (1)** Exception is recognized. Transfer control to the exception handler by executing the `__except` compound statement, then continue execution after the `__except` block.

Because the `__except` expression is evaluated as a C expression, it is limited to a single value, the conditional-expression operator, or the comma operator. If more extensive processing is required, the expression can call a routine that returns one of the three values listed above.

Each application can have its own exception handler.

It is not valid to jump into a `__try` statement, but valid to jump out of one. The exception handler is not called if a process is terminated in the middle of executing a **try-except** statement.

For more information, see Knowledge Base article Q315937 : HOW TO: Trap Stack Overflow in a Visual C++ Application.

## The `_leave` Keyword

The `_leave` keyword is valid only within the guarded section of a **try-except** statement, and its effect is to jump to the end of the guarded section. Execution continues at the first statement after the exception handler.

A `goto` statement can also jump out of the guarded section, and it does not degrade performance as it does in a **try-finally** statement because stack unwinding does not occur. However, we recommend that you use the `_leave` keyword rather than a `goto` statement because you are less likely to make a programming mistake if the guarded section is large or complex.

### Structured Exception Handling Intrinsic Functions

Structured exception handling provides two intrinsic functions that are available to use with the **try-except** statement: `GetExceptionCode` and `GetExceptionInformation`.

`GetExceptionCode` returns the code (a 32-bit integer) of the exception.

The intrinsic function `GetExceptionInformation` returns a pointer to a structure containing additional information about the exception. Through this pointer, you can access the machine state that existed at the time of a hardware exception. The structure is as follows:

```
typedef struct _EXCEPTION_POINTERS {
    PEXCEPTION_RECORD ExceptionRecord;
    PCONTEXT ContextRecord;
} EXCEPTION_POINTERS, *PEXCEPTION_POINTERS;
```

The pointer types `PEXCEPTION_RECORD` and `PCONTEXT` are defined in the include file WINNT.H, and `_EXCEPTION_RECORD` and `_CONTEXT` are defined in the include file EXCPT.H

You can use `GetExceptionCode` within the exception handler. However, you can use `GetExceptionInformation` only within the exception filter expression. The information it points to is generally on the stack and is no longer available when control is transferred to the exception handler.

The intrinsic function `AbnormalTermination` is available within a termination handler. It returns 0 if the body of the **try-finally** statement terminates sequentially. In all other cases, it returns 1.

EXCPT.H defines some alternate names for these intrinsics:

`GetExceptionCode` is equivalent to `_exception_code`

`GetExceptionInformation` is equivalent to `_exception_info`

`AbnormalTermination` is equivalent to `_abnormal_termination`

## Example

```
// exceptions_try_except_Statement.cpp
// Example of try-except and try-finally statements
#include <stdio.h>
#include <windows.h> // for EXCEPTION_ACCESS_VIOLATION
#include <excpt.h>

int filter(unsigned int code, struct _EXCEPTION_POINTERS *ep)
{
    puts("in filter.");
    if (code == EXCEPTION_ACCESS_VIOLATION)
    {
        puts("caught AV as expected.");
        return EXCEPTION_EXECUTE_HANDLER;
    }
    else
    {
        puts("didn't catch AV, unexpected.");
        return EXCEPTION_CONTINUE_SEARCH;
    };
}

int main()
{
    int* p = 0x00000000;    // pointer to NULL
    puts("hello");
    __try
    {
        puts("in try");
        __try
        {
            puts("in try");
            *p = 13;      // causes an access violation exception,
        }
        __finally
        {
            puts("in finally. termination: ");
            puts(AbnormalTermination() ? "\tabnormal" : "\tnormal");
        }
    }
    __except(filter(GetExceptionCode(), GetExceptionInformation()))
    {
        puts("in except");
    }
    puts("world");
}
```

## Output

```
hello
in try
in try
in filter.
caught AV as expected.
in finally. termination:
    abnormal
in except
world
```

## END Microsoft Specific

## See Also

- [Writing an Exception Handler](#)
- [Structured Exception Handling \(C/C++\)](#)
- [Keywords](#)

# Writing an Exception Filter

12/21/2017 • 2 min to read • [Edit Online](#)

You can handle an exception either by jumping to the level of the exception handler or by continuing execution. Instead of using the exception handler code to handle the exception and falling through, you can use *filter* to clean up the problem and then, by returning -1, resume normal flow without clearing the stack.

## NOTE

Some exceptions cannot be continued. If *filter* evaluates to -1 for such an exception, the system raises a new exception. When you call [RaiseException](#), you determine whether the exception will continue.

For example, the following code uses a function call in the *filter* expression: this function handles the problem and then returns -1 to resume normal flow of control:

```
// exceptions_Writing_an_Exception_Filter.cpp
#include <windows.h>
int main() {
    int Eval_Exception( int );

    __try {}
    __except ( Eval_Exception( GetExceptionCode( )) ) {
        ;
    }

    void ResetVars( int ) {}
    int Eval_Exception ( int n_except ) {
        if ( n_except != STATUS_INTEGER_OVERFLOW &&
            n_except != STATUS_FLOAT_OVERFLOW ) // Pass on most exceptions
        return EXCEPTION_CONTINUE_SEARCH;

        // Execute some code to clean up problem
        ResetVars( 0 ); // initializes data to 0
        return EXCEPTION_CONTINUE_EXECUTION;
    }
}
```

It is a good idea to use a function call in the *filter* expression whenever *filter* needs to do anything complex. Evaluating the expression causes execution of the function, in this case, `Eval_Exception`.

Note the use of `GetExceptionCode` to determine the exception. You must call this function inside the filter itself. `Eval_Exception` cannot call `GetExceptionCode`, but it must have the exception code passed to it.

This handler passes control to another handler unless the exception is an integer or floating-point overflow. If it is, the handler calls a function (`ResetVars` is only an example, not an API function) to reset some global variables. *Statement-block-2*, which in this example is empty, can never be executed because `Eval_Exception` never returns `EXCEPTION_EXECUTE_HANDLER` (1).

Using a function call is a good general-purpose technique for dealing with complex filter expressions. Two other C language features that are useful are:

- The conditional operator
- The comma operator

The conditional operator is frequently useful, because it can be used to check for a specific return code and then return one of two different values. For example, the filter in the following code recognizes the exception only if the exception is `STATUS_INTEGER_OVERFLOW` :

```
_except( GetExceptionCode() == STATUS_INTEGER_OVERFLOW ? 1 : 0 ) {
```

The purpose of the conditional operator in this case is mainly to provide clarity, because the following code produces the same results:

```
_except( GetExceptionCode() == STATUS_INTEGER_OVERFLOW ) {
```

The conditional operator is more useful in situations where you might want the filter to evaluate to -1, `EXCEPTION_CONTINUE_EXECUTION`.

The comma operator enables you to perform multiple, independent operations inside a single expression. The effect is roughly that of executing multiple statements and then returning the value of the last expression. For example, the following code stores the exception code in a variable and then tests it:

```
_except( nCode = GetExceptionCode(), nCode == STATUS_INTEGER_OVERFLOW )
```

## See Also

[Writing an Exception Handler](#)

[Structured Exception Handling \(C/C++\)](#)

# Raising Software Exceptions

12/21/2017 • 2 min to read • [Edit Online](#)

Some of the most common sources of program errors are not flagged as exceptions by the system. For example, if you attempt to allocate a memory block but there is insufficient memory, the run-time or API function does not raise an exception but returns an error code.

However, you can treat any condition as an exception by detecting that condition in your code and then reporting it by calling the [RaiseException](#) function. By flagging errors this way, you can bring the advantages of structured exception handling to any kind of run-time error.

To use structured exception handling with errors:

- Define your own exception code for the event.
- Call **RaiseException** when you detect a problem.
- Use exception-handling filters to test for the exception code you defined.

The WINERROR.H file shows the format for exception codes. To make sure that you do not define a code that conflicts with an existing exception code, set the third most significant bit to 1. The four most-significant bits should be set as shown in the following table.

BITS	RECOMMENDED BINARY SETTING	DESCRIPTION
31-30	11	These two bits describe the basic status of the code: 11 = error, 00 = success, 01 = informational, 10 = warning.
29	1	Client bit. Set to 1 for user-defined codes.
28	0	Reserved bit. (Leave set to 0.)

You can set the first two bits to a setting other than 11 binary if you want, although the "error" setting is appropriate for most exceptions. The important thing to remember is to set bits 29 and 28 as shown in the previous table.

The resulting error code should therefore have the highest four bits set to hexadecimal E. For example, the following definitions define exception codes that do not conflict with any Windows exception codes. (You may, however, need to check which codes are used by third-party DLLs.)

```
#define STATUS_INSUFFICIENT_MEM      0xE0000001
#define STATUS_FILE_BAD_FORMAT      0xE0000002
```

After you have defined an exception code, you can use it to raise an exception. For example, the following code raises the STATUS\_INSUFFICIENT\_MEM exception in response to a memory allocation problem:

```
lpstr = _malloc( nBufferSize );
if (lpstr == NULL)
    RaiseException( STATUS_INSUFFICIENT_MEM, 0, 0, 0);
```

If you want to simply raise an exception, you can set the last three parameters to 0. The three last parameters are useful for passing additional information and setting a flag that prevents handlers from continuing execution. See the [RaiseException](#) function in the Windows SDK for more information.

In your exception-handling filters, you can then test for the codes you've defined. For example:

```
_try {
    ...
}
_except (GetExceptionCode() == STATUS_INSUFFICIENT_MEM ||
          GetExceptionCode() == STATUS_FILE_BAD_FORMAT )
```

## See Also

[Writing an Exception Handler](#)

[Structured Exception Handling \(C/C++\)](#)

# Hardware Exceptions

12/21/2017 • 1 min to read • [Edit Online](#)

Most of the standard exceptions recognized by the operating system are hardware-defined exceptions. Windows recognizes a few low-level software exceptions, but these are usually best handled by the operating system.

Windows maps the hardware errors of different processors to the exception codes in this section. In some cases, a processor may generate only a subset of these exceptions. Windows preprocesses information about the exception and issues the appropriate exception code.

The hardware exceptions recognized by Windows are summarized in the following table:

EXCEPTION CODE	CAUSE OF EXCEPTION
<b>STATUS_ACCESS_VIOLATION</b>	Reading or writing to an inaccessible memory location.
<b>STATUS_BREAKPOINT</b>	Encountering a hardware-defined breakpoint; used only by debuggers.
<b>STATUS_DATATYPE_MISALIGNMENT</b>	Reading or writing to data at an address that is not properly aligned; for example, 16-bit entities must be aligned on 2-byte boundaries. (Not applicable to Intel 80x86 processors.)
<b>STATUS_FLOAT_DIVIDE_BY_ZERO</b>	Dividing floating-point type by 0.0.
<b>STATUS_FLOAT_OVERFLOW</b>	Exceeding maximum positive exponent of floating-point type.
<b>STATUS_FLOAT_UNDERFLOW</b>	Exceeding magnitude of lowest negative exponent of floating-point type.
<b>STATUS_FLOATING_RESEVERED_OPERAND</b>	Using a reserved floating-point format (invalid use of format).
<b>STATUS_ILLEGAL_INSTRUCTION</b>	Attempting to execute an instruction code not defined by the processor.
<b>STATUS_PRIVILEGED_INSTRUCTION</b>	Executing an instruction not allowed in current machine mode.
<b>STATUS_INTEGER_DIVIDE_BY_ZERO</b>	Dividing an integer type by 0.
<b>STATUS_INTEGER_OVERFLOW</b>	Attempting an operation that exceeds the range of the integer.
<b>STATUS_SINGLE_STEP</b>	Executing one instruction in single-step mode; used only by debuggers.

Many of the exceptions listed in the previous table are intended to be handled by debuggers, the operating system, or other low-level code. With the exception of integer and floating-point errors, your code should not handle these errors. Thus, you should usually use the exception-handling filter to ignore exceptions (evaluate to 0). Otherwise, you may prevent lower-level mechanisms from responding appropriately. You can, however, take appropriate precautions against the potential effect of these low-level errors by [writing termination handlers](#).

## See Also

[Writing an Exception Handler](#)

[Structured Exception Handling \(C/C++\)](#)

# Restrictions on Exception Handlers

12/21/2017 • 1 min to read • [Edit Online](#)

The principal limitation to using exception handlers in code is that you cannot use a `goto` statement to jump into a `__try` statement block. Instead, you must enter the statement block through normal flow of control. You can jump out of a `__try` statement block and nest exception handlers as you choose.

## See Also

[Writing an Exception Handler](#)

[Structured Exception Handling \(C/C++\)](#)

# Writing a Termination Handler

12/21/2017 • 1 min to read • [Edit Online](#)

Unlike an exception handler, a termination handler is always executed, regardless of whether the protected block of code terminated normally. The sole purpose of the termination handler should be to ensure that resources, such as memory, handles, and files, are properly closed regardless of how a section of code finishes executing.

Termination handlers use the try-finally statement.

## What do you want to know more about?

- [The try-finally statement](#)
- [Cleaning up resources](#)
- [Timing of actions in exception handling](#)
- [Restrictions on termination handlers](#)

## See Also

[Structured Exception Handling \(C/C++\)](#)

# try-finally Statement

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

The following syntax describes the `try-finally` statement:

```
_try {  
    // guarded code  
}  
_finally {  
    // termination code  
}
```

## Grammar

*try-finally-statement:*

`_try compound-statement`

`_finally compound-statement`

The `try-finally` statement is a Microsoft extension to the C and C++ languages that enables target applications to guarantee execution of cleanup code when execution of a block of code is interrupted. Cleanup consists of such tasks as deallocating memory, closing files, and releasing file handles. The `try-finally` statement is especially useful for routines that have several places where a check is made for an error that could cause premature return from the routine.

For related information and a code sample, see [try-except Statement](#). For more information on structured exception handling in general, see [Structured Exception Handling](#). For more information on handling exceptions in managed applications, see [Exception Handling under /clr](#).

### NOTE

Structured exception handling works with Win32 for both C and C++ source files. However, it is not specifically designed for C++. You can ensure that your code is more portable by using C++ exception handling. Also, C++ exception handling is more flexible, in that it can handle exceptions of any type. For C++ programs, it is recommended that you use the C++ exception-handling mechanism ([try, catch, and throw statements](#)).

The compound statement after the `_try` clause is the guarded section. The compound statement after the `_finally` clause is the termination handler. The handler specifies a set of actions that execute when the guarded section is exited, regardless of whether the guarded section is exited by an exception (abnormal termination), or by standard fall through (normal termination).

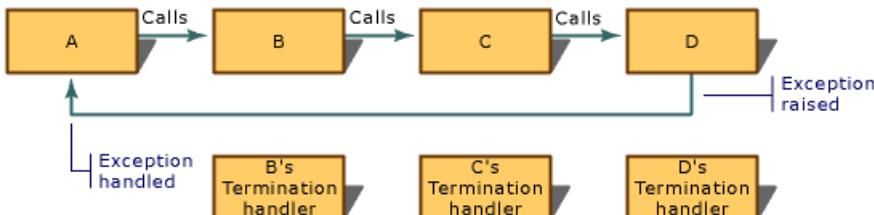
Control reaches a `_try` statement by simple sequential execution (fall through). When control enters the `_try`, its associated handler becomes active. If the flow of control reaches the end of the try block, execution proceeds as follows:

1. The termination handler is invoked.
2. When the termination handler completes, execution continues after the `_finally` statement. Regardless of how the guarded section ends (for example, via a `goto` out of the guarded body or a `return` statement), the termination handler is executed `before` the flow of control moves out of the guarded section.

A `__finally` statement does not block searching for an appropriate exception handler.

If an exception occurs in the `__try` block, the operating system must find a handler for the exception or the program will fail. If a handler is found, any and all `__finally` blocks are executed and execution resumes in the handler.

For example, suppose a series of function calls links function A to function D, as shown in the following figure. Each function has one termination handler. If an exception is raised in function D and handled in A, the termination handlers are called in this order as the system unwinds the stack: D, C, B.



Order of Termination-Handler Execution

#### NOTE

The behavior of try-finally is different from some other languages that support the use of **finally**, such as C#. A single `__try` may have either, but not both, of `__finally` and `__except`. If both are to be used together, an outer try-except statement must enclose the inner try-finally statement. The rules specifying when each block executes are also different.

## The `_leave` Keyword

The `_leave` keyword is valid only within the guarded section of a `try-finally` statement, and its effect is to jump to the end of the guarded section. Execution continues at the first statement in the termination handler.

A `goto` statement can also jump out of the guarded section, but it degrades performance because it invokes stack unwinding. The `_leave` statement is more efficient because it does not cause stack unwinding.

## Abnormal Termination

Exiting a `try-finally` statement using the `longjmp` run-time function is considered abnormal termination. It is illegal to jump into a `__try` statement, but legal to jump out of one. All `__finally` statements that are active between the point of departure (normal termination of the `__try` block) and the destination (the `__except` block that handles the exception) must be run. This is called a local unwind.

If a `try` block is prematurely terminated for any reason, including a jump out of the block, the system executes the associated **finally** block as a part of the process of unwinding the stack. In such cases, the `AbnormalTermination` function returns TRUE if called from within the **finally** block; otherwise, it returns FALSE.

The termination handler is not called if a process is killed in the middle of executing a `try-finally` statement.

## END Microsoft Specific

## See Also

- [Writing a Termination Handler](#)
- [Structured Exception Handling \(C/C++\)](#)
- [Keywords](#)
- [Termination-Handler Syntax](#)

# Cleaning up Resources

12/21/2017 • 1 min to read • [Edit Online](#)

During termination-handler execution, you may not know which resources are actually allocated before the termination handler was called. It is possible that the `__try` statement block was interrupted before all resources were allocated, so that not all resources were opened.

Therefore, to be safe, you should check to see which resources are actually open before proceeding with termination-handling cleanup. A recommended procedure is to:

1. Initialize handles to NULL.
2. In the `__try` statement block, allocate resources. Handles are set to positive values as the resource is allocated.
3. In the `__finally` statement block, release each resource whose corresponding handle or flag variable is nonzero or not NULL.

## Example

For example, the following code uses a termination handler to close three files and a memory block that were allocated in the `__try` statement block. Before cleaning up a resource, the code first checks to see if the resource was allocated.

```
// exceptions_Cleaning_up_Resources.cpp
#include <stdlib.h>
#include <malloc.h>
#include <stdio.h>
#include <windows.h>

void fileOps() {
    FILE *fp1 = NULL,
        *fp2 = NULL,
        *fp3 = NULL;
    LPVOID lpvoid = NULL;
    errno_t err;

    __try {
        lpvoid = malloc( BUFSIZ );

        err = fopen_s(&fp1, "ADDRESS.DAT", "w+" );
        err = fopen_s(&fp2, "NAMES.DAT", "w+" );
        err = fopen_s(&fp3, "CARS.DAT", "w+" );
    }
    __finally {
        if ( fp1 )
            fclose( fp1 );
        if ( fp2 )
            fclose( fp2 );
        if ( fp3 )
            fclose( fp3 );
        if ( lpvoid )
            free( lpvoid );
    }
}

int main() {
    fileOps();
}
```

## See Also

- [Writing a Termination Handler](#)
- [Structured Exception Handling \(C/C++\)](#)

# Timing of Exception Handling: A Summary

12/21/2017 • 1 min to read • [Edit Online](#)

A termination handler is executed no matter how the `__try` statement block is terminated. Causes include jumping out of the `__try` block, a `longjmp` statement that transfers control out of the block, and unwinding the stack due to exception handling.

## NOTE

Visual C++ supports two forms of the `setjmp` and `longjmp` statements. The fast version bypasses termination handling but is more efficient. To use this version, include the file `SETJMP.H`. The other version supports termination handling as described in the previous paragraph. To use this version, include the file `SETJMPEX.H`. The increase in performance of the fast version depends on hardware configuration.

The operating system executes all termination handlers in the proper order before any other code can be executed, including the body of an exception handler.

When the cause for interruption is an exception, the system must first execute the filter portion of one or more exception handlers before deciding what to terminate. The order of events is:

1. An exception is raised.
2. The system looks at the hierarchy of active exception handlers and executes the filter of the handler with highest precedence; this is the exception handler most recently installed and most deeply nested, in terms of blocks and function calls.
3. If this filter passes control (returns 0), the process continues until a filter is found that does not pass control.
4. If this filter returns -1, execution continues where the exception was raised, and no termination takes place.
5. If the filter returns 1, the following events occur:
  - The system unwinds the stack, clearing all stack frames between the currently executing code (where the exception was raised) and the stack frame that contains the exception handler gaining control.
  - As the stack is unwound, each termination handler on the stack is executed.
  - The exception handler itself is executed.
  - Control passes to the line of code after the end of this exception handler.

## See Also

[Writing a Termination Handler](#)  
[Structured Exception Handling \(C/C++\)](#)

# Restrictions on Termination Handlers

12/21/2017 • 1 min to read • [Edit Online](#)

You cannot use a `goto` statement to jump into a `__try` statement block or a `__finally` statement block. Instead, you must enter the statement block through normal flow of control. (You can, however, jump out of a `__try` statement block.) Also, you cannot nest an exception handler or termination handler inside a `__finally` block.

In addition, some kinds of code permitted in a termination handler produce questionable results, so you should use them with caution, if at all. One is a `goto` statement that jumps out of a `__finally` statement block. If the block is executing as part of normal termination, nothing unusual happens. But if the system is unwinding the stack, that unwinding stops, and the current function gains control as if there were no abnormal termination.

A `return` statement inside a `__finally` statement block presents roughly the same situation. Control returns to the immediate caller of the function containing the termination handler. If the system was unwinding the stack, this process is halted, and the program proceeds as if there had been no exception raised.

## See Also

[Writing a Termination Handler](#)  
[Structured Exception Handling \(C/C++\)](#)

# Using Structured Exception Handling with C++

12/21/2017 • 1 min to read • [Edit Online](#)

Structured exception handling described in these articles works with both C and C++ source files. However, it is not specifically designed for C++ and is not recommended. You can ensure that your code is more portable by using C++ exception handling. Also, the C++ exception handling mechanism is more flexible, in that it can handle exceptions of any type.

Microsoft C++ now supports the C++ exception handling model, based on the ANSI C++ Standard. This mechanism automatically handles destruction of local objects during stack unwind. If you are writing fault-tolerant C++ code, and you want to implement exception handling, it is strongly recommended that you use C++ exception handling, rather than structured exception handling. (Note that while the C++ compiler supports structured exception handling constructs as described in these articles, the standard C compiler does not support the C++ exception handling syntax.) For detailed information about C++ exception handling, see [C++ Exception Handling](#) and the *Annotated C++ Reference Manual* by Margaret Ellis and Bjarne Stroustrup.

## See Also

[Structured Exception Handling \(C/C++\)](#)

# Transporting Exceptions Between Threads

12/21/2017 • 10 min to read • [Edit Online](#)

Visual C++ supports *transporting an exception* from one thread to another. Transporting exceptions enables you to catch an exception in one thread and then make the exception appear to be thrown in a different thread. For example, you can use this feature to write a multithreaded application where the primary thread handles all the exceptions thrown by its secondary threads. Transporting exceptions is useful mostly to developers who create parallel programming libraries or systems. To implement transporting exceptions, Visual C++ provides the `exception_ptr` type and the `current_exception`, `rethrow_exception`, and `make_exception_ptr` functions.

## Syntax

```
namespace std
{
    typedef unspecified exception_ptr;
    exception_ptr current_exception();
    void rethrow_exception(exception_ptr p);
    template<class E>
        exception_ptr make_exception_ptr(E e) noexcept;
}
```

### Parameters

PARAMETER	DESCRIPTION
<code>unspecified</code>	An unspecified internal class that is used to implement the <code>exception_ptr</code> type.
<code>p</code>	An <code>exception_ptr</code> object that references an exception.
<code>E</code>	A class that represents an exception.
<code>e</code>	An instance of the parameter <code>E</code> class.

## Return Value

The `current_exception` function returns an `exception_ptr` object that references the exception that is currently in progress. If no exception is in progress, the function returns an `exception_ptr` object that is not associated with any exception.

The `make_exception_ptr` function returns an `exception_ptr` object that references the exception specified by the `e` parameter.

## Remarks

### Scenario

Imagine that you want to create an application that can scale to handle a variable amount of work. To achieve this objective, you design a multithreaded application where an initial, primary thread creates as many secondary threads as it needs in order to do the job. The secondary threads help the primary thread to manage resources, to

balance loads, and to improve throughput. By distributing the work, the multithreaded application performs better than a single-threaded application.

However, if a secondary thread throws an exception, you want the primary thread to handle it. This is because you want your application to handle exceptions in a consistent, unified manner regardless of the number of secondary threads.

## Solution

To handle the previous scenario, the C++ Standard supports transporting an exception between threads. If a secondary thread throws an exception, that exception becomes the *current exception*. By analogy to the real world, the current exception is said to be *in flight*. The current exception is in flight from the time it is thrown until the exception handler that catches it returns.

The secondary thread can catch the current exception in a `catch` block, and then call the `current_exception` function to store the exception in an `exception_ptr` object. The `exception_ptr` object must be available to the secondary thread and to the primary thread. For example, the `exception_ptr` object can be a global variable whose access is controlled by a mutex. The term *transport an exception* means an exception in one thread can be converted to a form that can be accessed by another thread.

Next, the primary thread calls the `rethrow_exception` function, which extracts and then throws the exception from the `exception_ptr` object. When the exception is thrown, it becomes the current exception in the primary thread. That is, the exception appears to originate in the primary thread.

Finally, the primary thread can catch the current exception in a `catch` block and then process it or throw it to a higher level exception handler. Or, the primary thread can ignore the exception and allow the process to end.

Most applications do not have to transport exceptions between threads. However, this feature is useful in a parallel computing system because the system can divide work among secondary threads, processors, or cores. In a parallel computing environment, a single, dedicated thread can handle all the exceptions from the secondary threads and can present a consistent exception-handling model to any application.

For more information about the C++ Standards committee proposal, search the Internet for document number N2179, titled "Language Support for Transporting Exceptions between Threads".

## Exception-Handling Models and Compiler Options

Your application's exception-handling model determines whether it can catch and transport an exception. Visual C++ supports three models that can handle C++ exceptions, structured exception handling (SEH) exceptions, and common language runtime (CLR) exceptions. Use the `/EH` and `/clr` compiler options to specify your application's exception-handling model.

Only the following combination of compiler options and programming statements can transport an exception. Other combinations either cannot catch exceptions, or can catch but cannot transport exceptions.

- The `/EHs` compiler option and the `catch` statement can transport SEH and C++ exceptions.
- The `/EHs`, `/EHsc`, and `/EHsc` compiler options and the `catch` statement can transport C++ exceptions.
- The `/CLR` or `/CLR:pure` compiler option and the `catch` statement can transport C++ exceptions. The `/CLR` compiler options imply specification of the `/EHs` option. Note that the compiler does not support transporting managed exceptions. This is because managed exceptions, which are derived from the `System.Exception class`, are already objects that you can move between threads by using the facilities of the common language runtime.

## IMPORTANT

We recommend that you specify the `/EHsc` compiler option and catch only C++ exceptions. You expose yourself to a security threat if you use the `/EHa` or `/CLR` compiler option and a `catch` statement with an ellipsis *exception-declaration* (`catch(...)`). You probably intend to use the `catch` statement to capture a few specific exceptions. However, the `catch(...)` statement captures all C++ and SEH exceptions, including unexpected ones that should be fatal. If you ignore or mishandle an unexpected exception, malicious code can use that opportunity to undermine the security of your program.

# Usage

The following sections describe how to transport exceptions by using the `exception_ptr` type, and the `current_exception`, `rethrow_exception`, and `make_exception_ptr` functions.

## `exception_ptr` Type

Use an `exception_ptr` object to reference the current exception or an instance of a user-specified exception. In the Microsoft implementation, an exception is represented by an `EXCEPTION_RECORD` structure. Each `exception_ptr` object includes an exception reference field that points to a copy of the `EXCEPTION_RECORD` structure that represents the exception.

When you declare an `exception_ptr` variable, the variable is not associated with any exception. That is, its exception reference field is NULL. Such an `exception_ptr` object is called a *null exception\_ptr*.

Use the `current_exception` or `make_exception_ptr` function to assign an exception to an `exception_ptr` object. When you assign an exception to an `exception_ptr` variable, the variable's exception reference field points to a copy of the exception. If there is insufficient memory to copy the exception, the exception reference field points to a copy of a `std::bad_alloc` exception. If the `current_exception` or `make_exception_ptr` function cannot copy the exception for any other reason, the function calls the `terminate` function to exit the current process.

Despite its name, an `exception_ptr` object is not itself a pointer. It does not obey pointer semantics and cannot be used with the pointer member access (`->`) or indirection (`*`) operators. The `exception_ptr` object has no public data members or member functions.

## Comparisons:

You can use the equal (`==`) and not-equal (`!=`) operators to compare two `exception_ptr` objects. The operators do not compare the binary value (bit pattern) of the `EXCEPTION_RECORD` structures that represent the exceptions. Instead, the operators compare the addresses in the exception reference field of the `exception_ptr` objects. Consequently, a null `exception_ptr` and the NULL value compare as equal.

## `current_exception` Function

Call the `current_exception` function in a `catch` block. If an exception is in flight and the `catch` block can catch the exception, the `current_exception` function returns an `exception_ptr` object that references the exception. Otherwise, the function returns a null `exception_ptr` object.

## Details:

The `current_exception` function captures the exception that is in flight regardless of whether the `catch` statement specifies an *exception-declaration* statement.

The destructor for the current exception is called at the end of the `catch` block if you do not rethrow the exception. However, even if you call the `current_exception` function in the destructor, the function returns an `exception_ptr` object that references the current exception.

Successive calls to the `current_exception` function return `exception_ptr` objects that refer to different copies of the

current exception. Consequently, the objects compare as unequal because they refer to different copies, even though the copies have the same binary value.

## SEH Exceptions:

If you use the `/EHa` compiler option, you can catch an SEH exception in a C++ `catch` block. The `current_exception` function returns an `exception_ptr` object that references the SEH exception. And the `rethrow_exception` function throws the SEH exception if you call it with the transported `exception_ptr` object as its argument.

The `current_exception` function returns a null `exception_ptr` if you call it in an SEH `__finally` termination handler, an `__except` exception handler, or the `__except` filter expression.

A transported exception does not support nested exceptions. A nested exception occurs if another exception is thrown while an exception is being handled. If you catch a nested exception, the `EXCEPTION_RECORD.ExceptionRecord` data member points to a chain of `EXCEPTION_RECORD` structures that describe the associated exceptions. The `current_exception` function does not support nested exceptions because it returns an `exception_ptr` object whose `ExceptionRecord` data member is zeroed out.

If you catch an SEH exception, you must manage the memory referenced by any pointer in the `EXCEPTION_RECORD.ExceptionInformation` data member array. You must guarantee that the memory is valid during the lifetime of the corresponding `exception_ptr` object, and that the memory is freed when the `exception_ptr` object is deleted.

You can use structured exception (SE) translator functions together with the transport exceptions feature. If an SEH exception is translated to a C++ exception, the `current_exception` function returns an `exception_ptr` that references the translated exception instead of the original SEH exception. The `rethrow_exception` function subsequently throws the translated exception, not the original exception. For more information about SE translator functions, see [\\_set\\_se\\_translator](#).

## rethrow\_exception Function

After you store a caught exception in an `exception_ptr` object, the primary thread can process the object. In your primary thread, call the `rethrow_exception` function together with the `exception_ptr` object as its argument. The `rethrow_exception` function extracts the exception from the `exception_ptr` object and then throws the exception in the context of the primary thread. If the `p` parameter of the `rethrow_exception` function is a null `exception_ptr`, the function throws `std::bad_exception`.

The extracted exception is now the current exception in the primary thread, and you can handle it as you would any other exception. If you catch the exception, you can handle it immediately or use a `throw` statement to send it to a higher level exception handler. Otherwise, do nothing and let the default system exception handler terminate your process.

## make\_exception\_ptr Function

The `make_exception_ptr` function takes an instance of a class as its argument and then returns an `exception_ptr` that references the instance. Usually, you specify an `exception class` object as the argument to the `make_exception_ptr` function, although any class object can be the argument.

Calling the `make_exception_ptr` function is equivalent to throwing a C++ exception, catching it in a `catch` block, and then calling the `current_exception` function to return an `exception_ptr` object that references the exception. The Microsoft implementation of the `make_exception_ptr` function is more efficient than throwing and then catching an exception.

An application typically does not require the `make_exception_ptr` function, and we discourage its use.

## Example

The following example transports a standard C++ exception and a custom C++ exception from one thread to

another.

```
// transport_exception.cpp
// compile with: /EHsc /MD
#include <windows.h>
#include <stdio.h>
#include <exception>
#include <stdexcept>

using namespace std;

// Define thread-specific information.
#define THREADCOUNT 2
exception_ptr aException[THREADCOUNT];
int           aArg[THREADCOUNT];

DWORD WINAPI ThrowExceptions( LPVOID );

// Specify a user-defined, custom exception.
// As a best practice, derive your exception
// directly or indirectly from std::exception.
class myException : public std::exception {
};

int main()
{
    HANDLE aThread[THREADCOUNT];
    DWORD ThreadID;

    // Create secondary threads.
    for( int i=0; i < THREADCOUNT; i++ )
    {
        aArg[i] = i;
        aThread[i] = CreateThread(
            NULL,          // Default security attributes.
            0,             // Default stack size.
            (LPTHREAD_START_ROUTINE) ThrowExceptions,
            (LPVOID) &aArg[i], // Thread function argument.
            0,             // Default creation flags.
            &ThreadID); // Receives thread identifier.

        if( aThread[i] == NULL )
        {
            printf("CreateThread error: %d\n", GetLastError());
            return -1;
        }
    }

    // Wait for all threads to terminate.
    WaitForMultipleObjects(THREADCOUNT, aThread, TRUE, INFINITE);
    // Close thread handles.
    for( int i=0; i < THREADCOUNT; i++ ) {
        CloseHandle(aThread[i]);
    }

    // Rethrow and catch the transported exceptions.
    for ( int i = 0; i < THREADCOUNT; i++ ) {
        try {
            if (aException[i] == NULL) {
                printf("exception_ptr %d: No exception was transported.\n", i);
            }
            else {
                rethrow_exception( aException[i] );
            }
        }
        catch( const invalid_argument & ) {
            printf("exception_ptr %d: Caught an invalid_argument exception.\n", i);
        }
        catch( const myException & ) {
            printf("exception_ptr %d: Caught a myException exception.\n", i);
        }
    }
}
```

```
        }
    }
}

// Each thread throws an exception depending on its thread
// function argument, and then ends.
DWORD WINAPI ThrowExceptions( LPVOID lpParam )
{
    int x = *((int*)lpParam);
    if (x == 0) {
        try {
            // Standard C++ exception.
            // This example explicitly throws invalid_argument exception.
            // In practice, your application performs an operation that
            // implicitly throws an exception.
            throw invalid_argument("A C++ exception.");
        }
        catch ( const invalid_argument & ) {
            aException[x] = current_exception();
        }
    }
    else {
        // User-defined exception.
        aException[x] = make_exception_ptr( myException() );
    }
    return TRUE;
}
```

```
exception_ptr 0: Caught an invalid_argument exception.
exception_ptr 1: Caught a myException exception.
```

## Requirements

**Header:** <exception>

## See Also

[Exception Handling](#)

[/EH \(Exception Handling Model\)](#)

[/clr \(Common Language Runtime Compilation\)](#)

# Assertion and User-Supplied Messages (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

The C++ language supports three error handling mechanisms that help you debug your application: the [#error directive](#), the [static\\_assert](#) keyword, and the [assert Macro, \\_assert, \\_wassert](#) macro. All three mechanisms issue error messages, and two also test software assertions. A software assertion specifies a condition that you expect to be true at a particular point in your program. If a compile time assertion fails, the compiler issues a diagnostic message and a compilation error. If a run-time assertion fails, the operating system issues a diagnostic message and closes your application.

## Remarks

The lifetime of your application consists of a preprocessing, compile, and run time phase. Each error handling mechanism accesses debug information that is available during one of these phases. To debug effectively, select the mechanism that provides appropriate information about that phase:

- The [#error directive](#) is in effect at preprocessing time. It unconditionally emits a user-specified message and causes the compilation to fail with an error. The message can contain text that is manipulated by preprocessor directives but any resulting expression is not evaluated.
- The [static\\_assert](#) declaration is in effect at compile time. It tests a software assertion that is represented by a user-specified integral expression that can be converted to a Boolean. If the expression evaluates to zero (false), the compiler issues the user-specified message and the compilation fails with an error.

The `static_assert` declaration is especially useful for debugging templates because template arguments can be included in the user-specified expression.

- The [assert Macro, \\_assert, \\_wassert](#) macro is in effect at run time. It evaluates a user-specified expression, and if the result is zero, the system issues a diagnostic message and closes your application. Many other macros, such as `_ASSERT` and `_ASSERTE`, resemble this macro but issue different system-defined or user-defined diagnostic messages.

## See Also

[#error Directive \(C/C++\)](#)  
[assert Macro, \\_assert, \\_wassert](#)  
[\\_ASSERT, \\_ASERTE, \\_ASSERT\\_EXPR Macros](#)  
[static\\_assert](#)  
[\\_STATIC\\_ASSERT Macro](#)  
[Templates](#)

# static\_assert

12/21/2017 • 2 min to read • [Edit Online](#)

Tests a software assertion at compile time. If the specified constant expression is `false`, the compiler displays the specified message, if one is provided, and the compilation fails with error C2338; otherwise, the declaration has no effect.

## Syntax

```
static_assert( constant-expression, string-literal );  
  
**Visual Studio 2017 and later:**  
static_assert( constant-expression );
```

### Parameters

PARAMETER	DESCRIPTION
<code>constant-expression</code>	An integral constant expression that can be converted to a Boolean.  If the evaluated expression is zero (false), the <code>string-literal</code> parameter is displayed and the compilation fails with an error. If the expression is nonzero (true), the <code>static_assert</code> declaration has no effect.
<code>string-literal</code>	An message that is displayed if the <code>constant-expression</code> parameter is zero. The message is a string of characters in the <a href="#">base character set</a> of the compiler; that is, not <a href="#">multibyte or wide characters</a> .

## Remarks

The `constant-expression` parameter of a `static_assert` declaration represents a *software assertion*. A software assertion specifies a condition that you expect to be true at a particular point in your program. If the condition is true, the `static_assert` declaration has no effect. If the condition is false, the assertion fails, the compiler displays the message in `string-literal` parameter, and the compilation fails with an error. In Visual Studio 2017 and later, the `string-literal` parameter is optional.

The `static_assert` declaration tests a software assertion at compile time. In contrast, the [assert Macro, \\_assert, \\_wassert](#) macro tests a software assertion at run time and incurs a run time cost in space or time. The `static_assert` declaration is especially useful for debugging templates because template arguments can be included in the `constant-expression` parameter.

The compiler examines the `static_assert` declaration for syntax errors when the declaration is encountered. The compiler evaluates the `constant-expression` parameter immediately if it does not depend on a template parameter. Otherwise, the compiler evaluates the `constant-expression` parameter when the template is instantiated. Consequently, the compiler might issue a diagnostic message once when the declaration is encountered, and again when the template is instantiated.

You can use the `static_assert` keyword at namespace, class, or block scope. (The `static_assert` keyword is

technically a declaration, even though it does not introduce new name into your program, because it can be used at namespace scope.)

## Description

In the following example, the `static_assert` declaration has namespace scope. Because the compiler knows the size of type `void *`, the expression is evaluated immediately.

## Example

```
static_assert(sizeof(void *) == 4, "64-bit code generation is not supported.");
```

## Description

In the following example, the `static_assert` declaration has class scope. The `static_assert` verifies that a template parameter is a *plain old data* (POD) type. The compiler examines the `static_assert` declaration when it is declared, but does not evaluate the `constant-expression` parameter until the `basic_string` class template is instantiated in `main()`.

## Example

```
#include <type_traits>
#include <iostream>
namespace std {
template <class CharT, class Traits = std::char_traits<CharT> >
class basic_string {
    static_assert(std::is_pod<CharT>::value,
                 "Template argument CharT must be a POD type in class template basic_string");
    // ...
};

struct NonPOD {
    NonPOD(const NonPOD &)
    virtual ~NonPOD() {}
};

int main()
{
    std::basic_string<char> bs;
```

## Description

In the following example, the `static_assert` declaration has block scope. The `static_assert` verifies that the size of the `VMPage` structure is equal to the virtual memory pagesize of the system.

## Example

```
#include <sys/param.h> // defines PAGESIZE
class VMMClient {
public:
    struct VMPage { // ...
    };
    int check_pagesize() {
        static_assert(sizeof(VMPage) == PAGESIZE,
            "Struct VMPage must be the same size as a system virtual memory page.");
        // ...
    }
    // ...
};
```

## See Also

[Assertion and User-Supplied Messages \(C++\)](#)

[#error Directive \(C/C++\)](#)

[assert Macro, \\_assert, \\_wassert](#)

[Templates](#)

[ASCII Character Set](#)

[Declarations and Definitions](#)

# Templates (C++)

12/21/2017 • 6 min to read • [Edit Online](#)

Templates are the basis for generic programming in C++. As a strongly-typed language, C++ requires all variables to have a specific type, either explicitly declared by the programmer or deduced by the compiler. However, many data structures and algorithms look the same no matter what type they are operating on. Templates enable you to define the operations of a class or function, and let the user specify what concrete types those operations should work on.

## Defining and using templates

A template is a construct that generates an ordinary type or function at compile time based on arguments the user supplies for the template parameters. For example, you can define a function template like this:

```
template <typename T>
T minimum(const T& lhs, const T& rhs)
{
    return lhs < rhs ? lhs : rhs;
}
```

The above code describes a template for a generic function with a single type parameter `T`, whose return value and call parameters (`lhs` and `rhs`) are all of this type. You can name a type parameter anything you like, but by convention single upper case letters are most commonly used. `T` is a template parameter; the `typename` keyword says that this parameter is a placeholder for a type. When the function is called, the compiler will replace every instance of `T` with the concrete type argument that is either specified by the user or deduced by the compiler. The process in which the compiler generates a class or function from a template is referred to as *template instantiation*; `minimum<int>` is an instantiation of the template `minimum<T>`.

Elsewhere, a user can declare an instance of the template that is specialized for `int`. Assume that `get_a()` and `get_b()` are functions that return an `int`:

```
int a = get_a();
int b = get_b();
int i = minimum<int>(a, b);
```

However, because this is a function template and the compiler can deduce the type of `T` from the arguments `a` and `b`, you can call it just like an ordinary function:

```
int i = minimum(a, b);
```

When the compiler encounters that last statement, it generates a new function in which every occurrence of `T` in the template is replaced with `int`:

```
int minimum(const int& lhs, const int& rhs)
{
    return lhs < rhs ? lhs : rhs;
}
```

The rules for how the compiler performs type deduction in function templates are based on the rules for

ordinary functions. For more information, see [Overload Resolution of Function Template Calls](#).

## Type parameters

In the `minimum` template above, note that the type parameter `T` is not qualified in any way until it is used in the function call parameters, where the `const` and `reference` qualifiers are added.

There is no practical limit to the number of type parameters. Separate multiple parameters by commas:

```
template <typename T, typename U, typename V> class Foo{};
```

The keyword `class` is equivalent to `typename` in this context. You can express the previous example as:

```
template <class T, class U, class V> class Foo{};
```

You can use the ellipses operator (...) to define a template that takes an arbitrary number of zero or more type parameters:

```
template<typename... Arguments> class vtclass;

vtclass< > vtinstance1;
vtclass<int> vtinstance2;
vtclass<float, bool> vtinstance3;
```

Any built-in or user-defined type can be used as a type argument. For example, you can use `std::vector` in the Standard Library to store ints, doubles, strings, `MyClass`, `const MyClass*`, `MyClass&`. The primary restriction when using templates is that a type argument must support any operations that are applied to the type parameters. For example, if we call `minimum` using `MyClass` as in this example:

```
class MyClass
{
public:
    int num;
    std::wstring description;
};

int main()
{
    MyClass mc1 {1, L"hello"};
    MyClass mc2 {2, L"goodbye"};
    auto result = minimum(mc1, mc2); // Error! C2678
}
```

A compiler error will be generated because `MyClass` does not provide an overload for the `<` operator.

There is no inherent requirement that the type arguments for any particular template all belong to the same object hierarchy, although you can define a template that enforces such a restriction. You can combine object-oriented techniques with templates; for example, you can store a `Derived*` in a `vector<Base*>`. Note that the arguments must be pointers

```

vector<MyClass*> vec;
MyDerived d(3, L"back again", time(0));
vec.push_back(&d);

// or more realistically:
vector<shared_ptr<MyClass>> vec2;
vec2.push_back(make_shared<MyDerived>());

```

The basic requirements that vector and other standard library containers impose on elements of `T` is that `T` be copy-assignable and copy-constructible.

## Non-type parameters

Unlike generic types in other languages such as C# and Java, C++ templates support non-type parameters, also called value parameters. For example, you can provide a constant integral value to specify the length of an array, as with this example that is similar to the `std::array` class in the Standard Library:

```

template<typename T, size_t L>
class MyArray
{
    T arr[L];
public:
    MyArray() { ... }
};

```

Note the syntax in the template declaration. The `size_t` value is passed in as a template argument at compile time and must be constant or a `constexpr` expression. You use it like this:

```
MyArray<MyClass*, 10> arr;
```

Other kinds of values including pointers and references can be passed in as non-type parameters. For example, you can pass in a pointer to a function or function object to customize some operation inside the template code.

## Templates as template parameters

A template can be a template parameter. In this example, `MyClass2` has two template parameters: a typename parameter `T` and a template parameter `Arr`:

```

template<typename T, template<typename U, int I> class Arr>
class MyClass2
{
    T t; //OK
    Arr<T, 10> a;
    U u; //Error. U not in scope
};

```

Because the `Arr` parameter itself has no body, its parameter names are not needed. In fact, it is an error to refer to `Arr`'s typename or class parameter names from within the body of `MyClass2`. For this reason, `Arr`'s type parameter names can be omitted, as shown in this example:

```
template<typename T, template<typename, int> class Arr>
class MyClass2
{
    T t; //OK
    Arr<T, 10> a;
};
```

## Default template arguments

Class and function templates can have default arguments. When a template has a default argument you can leave it unspecified when you use it. For example, the std::vector template has a default argument for the allocator:

```
template <class T, class Allocator = allocator<T>> class vector;
```

In most cases the default std::allocator class is acceptable, so you use a vector like this:

```
vector<int> myInts;
```

But if necessary you can specify a custom allocator like this:

```
vector<int, MyAllocator> ints;
```

For multiple template arguments, all arguments after the first default argument must have default arguments.

When using a template whose parameters are all defaulted, use empty angle brackets:

```
template<typename A = int, typename B = double>
class Bar
{
    //...
};

int main()
{
    Bar<> bar; // use all default type arguments
}
```

## Template specialization

In some cases, it isn't possible or desirable for a template to define exactly the same code for any type. For example, you might wish to define a code path to be executed only if the type argument is a pointer, or a std::wstring, or a type derived from a particular base class. In such cases you can define a *specialization* of the template for that particular type. When a user instantiates the template with that type, the compiler uses the specialization to generate the class, and for all other types, the compiler chooses the more general template. Specializations in which all parameters are specialized are *complete specializations*. If only some of the parameters are specialized, it is called a *partial specialization*.

```
template <typename K, typename V>
class MyMap{/*...*/};

// partial specialization for string keys
template<typename V>
class MyMap<string, V> {/*...*/};
...
MyMap<int, MyClass> classes; // uses original template
MyMap<string, MyClass> classes2; // uses the partial specialization
```

A template can have any number of specializations as long as each specialized type parameter is unique. Only class templates may be partially specialized. All complete and partial specializations of a template must be declared in the same namespace as the original template.

For more information, see [Template Specialization](#).

# typename

12/21/2017 • 1 min to read • [Edit Online](#)

In template definitions, provides a hint to the compiler that an unknown identifier is a type. In template parameter lists, is used to specify a type parameter.

## Syntax

```
typename identifier;
```

## Remarks

This keyword must be used if a name in a template definition is a qualified name that is dependent on a template argument; it is optional if the qualified name is not dependent. For more information, see [Templates and Name Resolution](#).

**typename** can be used by any type anywhere in a template declaration or definition. It is not allowed in the base class list, unless as a template argument to a template base class.

```
template <class T>
class C1 : typename T::InnerType // Error - typename not allowed.
{};
template <class T>
class C2 : A<typename T::InnerType> // typename OK.
{}
```

The **typename** keyword can also be used in place of **class** in template parameter lists. For example, the following statements are semantically equivalent:

```
template<class T1, class T2>...
template<typename T1, typename T2>...
```

## Example

```
// typename.cpp
template<class T> class X
{
    typename T::Y m_y; // treat Y as a type
};

int main()
{
}
```

## See Also

[Templates](#)  
[Keywords](#)

# Class Templates

12/21/2017 • 6 min to read • [Edit Online](#)

This topic describes rules that are specific to C++ class templates.

## Member functions of class templates

Member functions can be defined inside or outside of a class template. They are defined like function templates if defined outside the class template.

```
// member_function_templates1.cpp
template<class T, int i> class MyStack
{
    T* pStack;
    T StackBuffer[i];
    static const int cItems = i * sizeof(T);
public:
    MyStack( void );
    void push( const T item );
    T& pop( void );
};

template< class T, int i > MyStack< T, i >::MyStack( void )
{
};

template< class T, int i > void MyStack< T, i >::push( const T item )
{
};

template< class T, int i > T& MyStack< T, i >::pop( void )
{
};

int main()
{
}
```

Note that just as with any template class member function, the definition of the class's constructor member function includes the template argument list twice.

Member functions can themselves be function templates, specifying additional parameters, as in the following example.

```

// member_templates.cpp
template<typename T>
class X
{
public:
    template<typename U>
    void mf(const U &u);
};

template<typename T> template <typename U>
void X<T>::mf(const U &u)
{
}

int main()
{
}

```

## Nested class templates

Templates can be defined within classes or class templates, in which case they are referred to as member templates. Member templates that are classes are referred to as nested class templates. Member templates that are functions are discussed in [Member Function Templates](#).

Nested class templates are declared as class templates inside the scope of the outer class. They can be defined inside or outside of the enclosing class.

The following code demonstrates a nested class template inside an ordinary class.

```

// nested_class_template1.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

class X
{
    template <class T>
    struct Y
    {
        T m_t;
        Y(T t): m_t(t) { }
    };

    Y<int> yInt;
    Y<char> yChar;

public:
    X(int i, char c) : yInt(i), yChar(c) { }
    void print()
    {
        cout << yInt.m_t << " " << yChar.m_t << endl;
    }
};

int main()
{
    X x(1, 'a');
    x.print();
}

```

```
// nested_class_template2.cpp
```

```

// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T>
class X
{
    template <class U> class Y
    {
        U* u;
    public:
        Y();
        U& Value();
        void print();
        ~Y();
    };
    Y<int> y;
public:
    X(T t) { y.Value() = t; }
    void print() { y.print(); }
};

template <class T>
template <class U>
X<T>::Y<U>::Y()
{
    cout << "X<T>::Y<U>::Y()" << endl;
    u = new U();
}

template <class T>
template <class U>
U& X<T>::Y<U>::Value()
{
    return *u;
}

template <class T>
template <class U>
void X<T>::Y<U>::print()
{
    cout << this->Value() << endl;
}

template <class T>
template <class U>
X<T>::Y<U>::~Y()
{
    cout << "X<T>::Y<U>::~Y()" << endl;
    delete u;
}

int main()
{
    X<int>* xi = new X<int>(10);
    X<char>* xc = new X<char>('c');
    xi->print();
    xc->print();
    delete xi;
    delete xc;
}

//Output:
X<T>::Y<U>::Y()
X<T>::Y<U>::Y()
10
99
X<T>::Y<U>::~Y()

```

```
X<T>::Y<U>::~Y()
```

Local classes are not allowed to have member templates.

## Template friends

Class templates can have [friends](#). A class or class template, function, or function template can be a friend to a template class. Friends can also be specializations of a class template or function template, but not partial specializations.

In the following example, a friend function is defined as a function template within the class template. This code produces a version of the friend function for every instantiation of the template. This construct is useful if your friend function depends on the same template parameters as the class does.

```
// template_friend1.cpp
// compile with: /EHsc

#include <iostream>
using namespace std;

template <class T> class Array {
    T* array;
    int size;

public:
    Array(int sz): size(sz) {
        array = new T[size];
        memset(array, 0, size * sizeof(T));
    }

    Array(const Array& a) {
        size = a.size;
        array = new T[size];
        memcpy_s(array, a.array, sizeof(T));
    }

    T& operator[](int i) {
        return *(array + i);
    }

    int Length() { return size; }

    void print() {
        for (int i = 0; i < size; i++)
            cout << *(array + i) << " ";
        cout << endl;
    }
};

template<class T>
friend Array<T>* combine(Array<T>& a1, Array<T>& a2);
};

template<class T>
Array<T>* combine(Array<T>& a1, Array<T>& a2) {
    Array<T>* a = new Array<T>(a1.size + a2.size);
    for (int i = 0; i < a1.size; i++)
        (*a)[i] = *(a1.array + i);

    for (int i = 0; i < a2.size; i++)
        (*a)[i + a1.size] = *(a2.array + i);

    return a;
}
```

```

int main() {
    Array<char> alpha1(26);
    for (int i = 0 ; i < alpha1.Length() ; i++)
        alpha1[i] = 'A' + i;

    alpha1.print();

    Array<char> alpha2(26);
    for (int i = 0 ; i < alpha2.Length() ; i++)
        alpha2[i] = 'a' + i;

    alpha2.print();
    Array<char>*alpha3 = combine(alpha1, alpha2);
    alpha3->print();
    delete alpha3;
}

//Output:
A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z
A B C D E F G H I J K L M N O P Q R S T U V W X Y Z a b c d e f g h i j k l m n o p q r s t u v w x y z

```

The next example involves a friend that has a template specialization. A function template specialization is automatically a friend if the original function template is a friend.

It is also possible to declare only the specialized version of the template as the friend, as the comment before the friend declaration in the following code indicates. If you do this, you must put the definition of the friend template specialization outside of the template class.

```

// template_friend2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T>
class Array;

template <class T>
void f(Array<T>& a);

template <class T> class Array
{
    T* array;
    int size;

public:
    Array(int sz): size(sz)
    {
        array = new T[size];
        memset(array, 0, size * sizeof(T));
    }
    Array(const Array& a)
    {
        size = a.size;
        array = new T[size];
        memcpy_s(array, a.array, sizeof(T));
    }
    T& operator[](int i)
    {
        return *(array + i);
    }
    int Length()
    {
        return size;
    }
    void print()
    {

```

```

        for (int i = 0; i < size; i++)
    {
        cout << *(array + i) << " ";
    }
    cout << endl;
}

// If you replace the friend declaration with the int-specific
// version, only the int specialization will be a friend.
// The code in the generic f will fail
// with C2248: 'Array<T>::size' :
// cannot access private member declared in class 'Array<T>'.
//friend void f<int>(Array<int>& a);

friend void f<>(Array<T>& a);
};

// f function template, friend of Array<T>
template <class T>
void f(Array<T>& a)
{
    cout << a.size << " generic" << endl;
}

// Specialization of f for int arrays
// will be a friend because the template f is a friend.
template<> void f(Array<int>& a)
{
    cout << a.size << " int" << endl;
}

int main()
{
    Array<char> ac(10);
    f(ac);

    Array<int> a(10);
    f(a);
}
//Output:
10 generic
10 int

```

The next example shows a friend class template declared within a class template. The class template is then used as the template argument for the friend class. Friend class templates must be defined outside of the class template in which they are declared. Any specializations or partial specializations of the friend template are also friends of the original class template.

```

// template_friend3.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T>
class X
{
private:
    T* data;
    void InitData(int seed) { data = new T(seed); }
public:
    void print() { cout << *data << endl; }
    template <class U> friend class Factory;
};

template <class U>
class Factory
{
public:
    U* GetNewObject(int seed)
    {
        U* pu = new U;
        pu->InitData(seed);
        return pu;
    }
};

int main()
{
    Factory< X<int> > XintFactory;
    X<int>* x1 = XintFactory.GetNewObject(65);
    X<int>* x2 = XintFactory.GetNewObject(97);

    Factory< X<char> > XcharFactory;
    X<char>* x3 = XcharFactory.GetNewObject('A');
    X<char>* x4 = XcharFactory.GetNewObject('a');

    x1->print();
    x2->print();
    x3->print();
    x4->print();
}
//Output:
65
97
A
a

```

## Reuse of Template Parameters

Template parameters can be reused in the template parameter list. For example, the following code is allowed:

```
// template_specifications2.cpp

class Y
{
};

template<class T, T* pT> class X1
{
};

template<class T1, class T2 = T1> class X2
{
};

Y aY;

X1<Y, &aY> x1;
X2<int> x2;

int main()
{
}
```

## See Also

[Templates](#)

# Template Friends

12/21/2017 • 4 min to read • [Edit Online](#)

Class templates can have [friends](#). A class or class template, function, or function template can be a friend to a template class. Friends can also be specializations of a class template or function template, but not partial specializations.

**C++ 11:** A type parameter can be declared as a friend by using the form `friend T;`.

```
template <typename T>
class my_class
{
    friend T;
    //...
};
```

## Example

In the following example, a friend function is defined as a function template within the class template. This code produces a version of the friend function for every instantiation of the template. This construct is useful if your friend function depends on the same template parameters as the class does.

```
// template_friend1.cpp
// compile with: /EHsc

#include <iostream>
using namespace std;

template <class T> class Array {
    T* array;
    int size;

public:
    Array(int sz): size(sz) {
        array = new T[size];
        memset(array, 0, size * sizeof(T));
    }

    Array(const Array& a) {
        size = a.size;
        array = new T[size];
        memcpy_s(array, a.array, sizeof(T));
    }

    T& operator[](int i) {
        return *(array + i);
    }

    int Length() { return size; }

    void print() {
        for (int i = 0; i < size; i++)
            cout << *(array + i) << " ";

        cout << endl;
    }
};

template<class T>
```

```

friend Array<T>* combine(Array<T>& a1, Array<T>& a2);
};

template<class T>
Array<T>* combine(Array<T>& a1, Array<T>& a2) {
    Array<T>* a = new Array<T>(a1.size + a2.size);
    for (int i = 0; i < a1.size; i++)
        (*a)[i] = *(a1.array + i);

    for (int i = 0; i < a2.size; i++)
        (*a)[i + a1.size] = *(a2.array + i);

    return a;
}

int main() {
    Array<char> alpha1(26);
    for (int i = 0 ; i < alpha1.Length() ; i++)
        alpha1[i] = 'A' + i;

    alpha1.print();

    Array<char> alpha2(26);
    for (int i = 0 ; i < alpha2.Length() ; i++)
        alpha2[i] = 'a' + i;

    alpha2.print();
    Array<char>*alpha3 = combine(alpha1, alpha2);
    alpha3->print();
    delete alpha3;
}

```

```

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z
A B C D E F G H I J K L M N O P Q R S T U V W X Y Z a b c d e f g h i j k l m n o p q r s t u v w x y z

```

## Example

The next example involves a friend that has a template specialization. A function template specialization is automatically a friend if the original function template is a friend.

It is also possible to declare only the specialized version of the template as the friend, as the comment before the friend declaration in the following code indicates. If you do this, you must put the definition of the friend template specialization outside of the template class.

```

// template_friend2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T>
class Array;

template <class T>
void f(Array<T>& a);

template <class T> class Array
{
    T* array;
    int size;

public:
    Array(int sz): size(sz)

```

```

{
    array = new T[size];
    memset(array, 0, size * sizeof(T));
}
Array(const Array& a)
{
    size = a.size;
    array = new T[size];
    memcpy_s(array, a.array, sizeof(T));
}
T& operator[](int i)
{
    return *(array + i);
}
int Length()
{
    return size;
}
void print()
{
    for (int i = 0; i < size; i++)
    {
        cout << *(array + i) << " ";
    }
    cout << endl;
}
// If you replace the friend declaration with the int-specific
// version, only the int specialization will be a friend.
// The code in the generic f will fail
// with C2248: 'Array<T>::size' :
// cannot access private member declared in class 'Array<T>'.
//friend void f<int>(Array<int>& a);

friend void f<>(Array<T>& a);
};

// f function template, friend of Array<T>
template <class T>
void f(Array<T>& a)
{
    cout << a.size << " generic" << endl;
}

// Specialization of f for int arrays
// will be a friend because the template f is a friend.
template<> void f(Array<int>& a)
{
    cout << a.size << " int" << endl;
}

int main()
{
    Array<char> ac(10);
    f(ac);

    Array<int> a(10);
    f(a);
}

```

```

10 generic
10 int

```

## Example

The next example shows a friend class template declared within a class template. The class template is then used as

the template argument for the friend class. Friend class templates must be defined outside of the class template in which they are declared. Any specializations or partial specializations of the friend template are also friends of the original class template.

```
// template_friend3.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T>
class X
{
private:
    T* data;
    void InitData(int seed) { data = new T(seed); }
public:
    void print() { cout << *data << endl; }
    template <class U> friend class Factory;
};

template <class U>
class Factory
{
public:
    U* GetNewObject(int seed)
    {
        U* pu = new U;
        pu->InitData(seed);
        return pu;
    }
};

int main()
{
    Factory< X<int> > XintFactory;
    X<int>* x1 = XintFactory.GetNewObject(65);
    X<int>* x2 = XintFactory.GetNewObject(97);

    Factory< X<char> > XcharFactory;
    X<char>* x3 = XcharFactory.GetNewObject(65);
    X<char>* x4 = XcharFactory.GetNewObject(97);
    x1->print();
    x2->print();
    x3->print();
    x4->print();
}
```

```
65
97
A
a
```

## See Also

[Default Arguments](#)

# Function Templates

12/21/2017 • 1 min to read • [Edit Online](#)

Class templates define a family of related classes that are based on the type arguments passed to the class upon instantiation. Function templates are similar to class templates but define a family of functions. With function templates, you can specify a set of functions that are based on the same code but act on different types or classes. The following function template swaps two items:

```
// function_templates1.cpp
template< class T > void MySwap( T& a, T& b ) {
    T c(a);
    a = b;
    b = c;
}
int main() {
```

This code defines a family of functions that swap the values of the arguments. From this template, you can generate functions that will swap **int** and **long** types and also user-defined types. **MySwap** will even swap classes if the class's copy constructor and assignment operator are properly defined.

In addition, the function template will prevent you from swapping objects of different types, because the compiler knows the types of the **a** and **b** parameters at compile time.

Although this function could be performed by a nontemplated function, using void pointers, the template version is typesafe. Consider the following calls:

```
int j = 10;
int k = 18;
CString Hello = "Hello, Windows!";
MySwap( j, k );           //OK
MySwap( j, Hello );      //error
```

The second **MySwap** call triggers a compile-time error, because the compiler cannot generate a **MySwap** function with parameters of different types. If void pointers were used, both function calls would compile correctly, but the function would not work properly at run time.

Explicit specification of the template arguments for a function template is allowed. For example:

```
// function_templates2.cpp
template<class T> void f(T) {}
int main(int j) {
    f<char>(j);    // Generate the specialization f(char).
    // If not explicitly specified, f(int) would be deduced.
}
```

When the template argument is explicitly specified, normal implicit conversions are done to convert the function argument to the type of the corresponding function template parameters. In the above example, the compiler will convert **char j** to type **int**.

## See Also

[Templates](#)

[Function Template Instantiation](#)

[Explicit Instantiation](#)

[Explicit Specialization of Function Templates](#)

# Function Template Instantiation

12/21/2017 • 1 min to read • [Edit Online](#)

When a function template is first called for each type, the compiler creates an instantiation. Each instantiation is a version of the templated function specialized for the type. This instantiation will be called every time the function is used for the type. If you have several identical instantiations, even in different modules, only one copy of the instantiation will end up in the executable file.

Conversion of function arguments is allowed in function templates for any argument and parameter pair where the parameter does not depend on a template argument.

Function templates can be explicitly instantiated by declaring the template with a particular type as an argument. For example, the following code is allowed:

```
// function_template_instantiation.cpp
template<class T> void f(T) { }

// Instantiate f with the explicitly specified template.
// argument 'int'
//
template void f<int> (int);

// Instantiate f with the deduced template argument 'char'.
template void f(char);
int main()
{}
```

## See Also

[Function Templates](#)

# Explicit Instantiation

12/21/2017 • 1 min to read • [Edit Online](#)

You can use explicit instantiation to create an instantiation of a templated class or function without actually using it in your code. Because this is useful when you are creating library (.lib) files that use templates for distribution, uninstantiated template definitions are not put into object (.obj) files.

This code explicitly instantiates `MyStack` for `int` variables and six items:

```
template class MyStack<int, 6>;
```

This statement creates an instantiation of `MyStack` without reserving any storage for an object. Code is generated for all members.

The next line explicitly instantiates only the constructor member function:

```
template MyStack<int, 6>::MyStack( void );
```

You can explicitly instantiate function templates by using a specific type argument to re-declare them, as shown in the example in [Function Template Instantiation](#).

You can use the `extern` keyword to prevent the automatic instantiation of members. For example:

```
extern template class MyStack<int, 6>;
```

Similarly, you can mark specific members as being external and not instantiated:

```
extern template MyStack<int, 6>::MyStack( void );
```

You can use the `extern` keyword to keep the compiler from generating the same instantiation code in more than one object module. You must instantiate the template function by using the specified explicit template parameters in at least one linked module if the function is called, or you will get a linker error when the program is built.

## NOTE

The `extern` keyword in the specialization only applies to member functions defined outside of the body of the class. Functions defined inside the class declaration are considered inline functions and are always instantiated.

## See Also

[Function Templates](#)

# Explicit Specialization of Function Templates

12/21/2017 • 1 min to read • [Edit Online](#)

With a function template, you can define special behavior for a specific type by providing an explicit specialization (override) of the function template for that type. For example:

```
template<> void MySwap(double a, double b);
```

This declaration enables you to define a different function for **double** variables. Like non-template functions, standard type conversions (such as promoting a variable of type **float** to **double**) are applied.

## Example

```
// explicit_specialization.cpp
template<class T> void f(T t)
{
};

// Explicit specialization of f with 'char' with the
// template argument explicitly specified:
//
template<> void f<char>(char c)
{
}

// Explicit specialization of f with 'double' with the
// template argument deduced:
//
template<> void f(double d)
{
}
int main()
{
}
```

## See Also

[Function Templates](#)

# Partial Ordering of Function Templates (C++)

12/21/2017 • 2 min to read • [Edit Online](#)

Multiple function templates that match the argument list of a function call can be available. C++ defines a partial ordering of function templates to specify which function should be called. The ordering is partial because there can be some templates that are considered equally specialized.

The compiler chooses the most specialized template function available from the possible matches. For example, if a function template takes a type **T**, and another function template taking **T\*** is available, the **T\*** version is said to be more specialized and is preferred over the generic **T** version whenever the argument is a pointer type, even though both would be allowable matches.

Use the following process to determine if one function template candidate is more specialized:

1. Consider two function templates, T1 and T2.
2. Replace the parameters in T1 with a hypothetical unique type X.
3. With the parameter list in T1, see if T2 is a valid template for that parameter list. Ignore any implicit conversions.
4. Repeat the same process with T1 and T2 reversed.
5. If one template is a valid template argument list for the other template, but the converse is not true, then that template is considered to be less specialized than the other template. If both templates using the previous step form valid arguments for each other, then they are considered to be equally specialized, and an ambiguous call results when you attempt to use them.
6. Using these rules:
  - a. A template specialization for a specific type is more specialized than one taking a generic type argument.
  - b. A template taking only **T\*** is more specialized than one taking only **T**, because a hypothetical type **X\*** is a valid argument for a **T** template argument, but **X** is not a valid argument for a **T\*** template argument.
  - c. **const T** is more specialized than **T**, because **const X** is a valid argument for a **T** template argument, but **X** is not a valid argument for a **const T** template argument.
  - d. **const T\*** is more specialized than **T\***, because **const X\*** is a valid argument for a **T\*** template argument, but **X\*** is not a valid argument for a **const T\*** template argument.

## Example

The following sample works as specified in the standard:

```
// partial_ordering_of_function_templates.cpp
// compile with: /EHsc
#include <iostream>

extern "C" int printf(const char*,...);

template <class T> void f(T) {
    printf_s("Less specialized function called\n");
}

template <class T> void f(T*) {
    printf_s("More specialized function called\n");
}

template <class T> void f(const T*) {
    printf_s("Even more specialized function for const T*\n");
}

int main() {
    int i =0;
    const int j = 0;
    int *pi = &i;
    const int *cpi = &j;

    f(i);    // Calls less specialized function.
    f(pi);   // Calls more specialized function.
    f(cpi); // Calls even more specialized function.
    // Without partial ordering, these calls would be ambiguous.
}
```

## Output

```
Less specialized function called
More specialized function called
Even more specialized function for const T*
```

## See Also

[Function Templates](#)

# Member Function Templates

12/21/2017 • 1 min to read • [Edit Online](#)

The term member template refers to both member function templates and nested class templates. Member function templates are template functions that are members of a class or class template.

Member functions can be function templates in several contexts. All functions of class templates are generic but are not referred to as member templates or member function templates. If these member functions take their own template arguments, they are considered to be member function templates.

## Example

Member function templates of nontemplate or template classes are declared as function templates with their template parameters.

```
// member_function_templates.cpp
struct X
{
    template <class T> void mf(T* t) {}
};

int main()
{
    int i;
    X* x = new X();
    x->mf(&i);
}
```

## Example

The following example shows a member function template of a template class.

```
// member_function_templates2.cpp
template<typename T>
class X
{
public:
    template<typename U>
    void mf(const U &u)
    {
    }
};

int main()
{}
```

## Example

```
// defining_member_templates_outside_class.cpp
template<typename T>
class X
{
public:
    template<typename U>
    void mf(const U &u);
};

template<typename T> template <typename U>
void X<T>::mf(const U &u)
{
}

int main()
{
}
```

## Example

Local classes are not allowed to have member templates.

Member template functions cannot be virtual functions and cannot override virtual functions from a base class when they are declared with the same name as a base class virtual function.

The following example shows a templated user-defined conversion:

```
// templated_user_defined_conversions.cpp
template <class T>
struct S
{
    template <class U> operator S<U>()
    {
        return S<U>();
    }
};

int main()
{
    S<int> s1;
    S<long> s2 = s1; // Convert s1 using UDC and copy constructs S<long>.
}
```

## See Also

[Function Templates](#)

# Template Specialization (C++)

12/21/2017 • 5 min to read • [Edit Online](#)

Class templates can be partially specialized, and the resulting class is still a template. Partial specialization allows template code to be partially customized for specific types in situations, such as:

- A template has multiple types and only some of them need to be specialized. The result is a template parameterized on the remaining types.
- A template has only one type, but a specialization is needed for pointer, reference, pointer to member, or function pointer types. The specialization itself is still a template on the type pointed to or referenced.

## Example

```
// partial_specialization_of_class_templates.cpp
template <class T> struct PTS {
    enum {
        IsPointer = 0,
        IsPointerToDataMember = 0
    };
};

template <class T> struct PTS<T*> {
    enum {
        IsPointer = 1,
        IsPointerToDataMember = 0
    };
};

template <class T, class U> struct PTS<T U::*> {
    enum {
        IsPointer = 0,
        IsPointerToDataMember = 1
    };
};

struct S{};

extern "C" int printf_s(const char*,...);

int main() {
    S s, *pS;
    int S::*ptm;
    printf_s("PTS<S>::IsPointer == %d PTS<S>::IsPointerToDataMember == %d\n",
            PTS<S>::IsPointer, PTS<S>:: IsPointerToDataMember);
    printf_s("PTS<S*>::IsPointer == %d PTS<S*>::IsPointerToDataMember ==%d\n"
            , PTS<S*>::IsPointer, PTS<S*>:: IsPointerToDataMember);
    printf_s("PTS<int S::*>::IsPointer == %d PTS"
            "<int S::*>::IsPointerToDataMember == %d\n",
            PTS<int S::*>::IsPointer, PTS<int S::*>::
            IsPointerToDataMember);
}
```

```
PTS<S>::IsPointer == 0 PTS<S>::IsPointerToDataMember == 0
PTS<S*>::IsPointer == 1 PTS<S*>::IsPointerToDataMember ==0
PTS<int S::*>::IsPointer == 0 PTS<int S::*>::IsPointerToDataMember == 1
```

## Example

If you have a template collection class that takes any type **T**, you can create a partial specialization that takes any pointer type **T\***. The following code demonstrates a collection class template **Bag** and a partial specialization for pointer types in which the collection dereferences the pointer types before copying them to the array. The collection then stores the values that are pointed to. With the original template, only the pointers themselves would have been stored in the collection, leaving the data vulnerable to deletion or modification. In this special pointer version of the collection, code to check for a null pointer in the **add** method is added.

```
// partial_specialization_of_class_templates2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

// Original template collection class.
template <class T> class Bag {
    T* elem;
    int size;
    int max_size;

public:
    Bag() : elem(0), size(0), max_size(1) {}
    void add(T t) {
        T* tmp;
        if (size + 1 >= max_size) {
            max_size *= 2;
            tmp = new T [max_size];
            for (int i = 0; i < size; i++)
                tmp[i] = elem[i];
            tmp[size++] = t;
            delete[] elem;
            elem = tmp;
        }
        else
            elem[size++] = t;
    }

    void print() {
        for (int i = 0; i < size; i++)
            cout << elem[i] << " ";
        cout << endl;
    }
};

// Template partial specialization for pointer types.
// The collection has been modified to check for NULL
// and store types pointed to.
template <class T> class Bag<T*> {
    T* elem;
    int size;
    int max_size;

public:
    Bag() : elem(0), size(0), max_size(1) {}
    void add(T* t) {
        T* tmp;
        if (t == NULL) { // Check for NULL
            cout << "Null pointer!" << endl;
            return;
        }

        if (size + 1 >= max_size) {
            max_size *= 2;
            tmp = new T [max_size];
            for (int i = 0; i < size; i++)
                tmp[i] = elem[i];
            elem = tmp;
        }
    }
}
```

```

        tmp[i] = *t;
        tmp[size++] = *t; // Dereference
        delete[] elem;
        elem = tmp;
    }
    else
        elem[size++] = *t; // Dereference
}

void print() {
    for (int i = 0; i < size; i++)
        cout << elem[i] << " ";
    cout << endl;
}
};

int main() {
    Bag<int> xi;
    Bag<char> xc;
    Bag<int*> xp; // Uses partial specialization for pointer types.

    xi.add(10);
    xi.add(9);
    xi.add(8);
    xi.print();

    xc.add('a');
    xc.add('b');
    xc.add('c');
    xc.print();

    int i = 3, j = 87, *p = new int[2];
    *p = 8;
    *(p + 1) = 100;
    xp.add(&i);
    xp.add(&j);
    xp.add(p);
    xp.add(p + 1);
    p = NULL;
    xp.add(p);
    xp.print();
}
}

```

```

10 9 8
a b c
Null pointer!
3 87 8 100

```

## Example

The following example defines a template class that takes pairs of any two types and then defines a partial specialization of that template class specialized so that one of the types is `int`. The specialization defines an additional sort method that implements a simple bubble sort based on the integer.

```

// partial_specialization_of_class_templates3.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class Key, class Value> class Dictionary {
    Key* keys;
    Value* values;
    int size;
    int max_size;

```

```

public:
    Dictionary(int initial_size) : size(0) {
        max_size = 1;
        while (initial_size >= max_size)
            max_size *= 2;
        keys = new Key[max_size];
        values = new Value[max_size];
    }
    void add(Key key, Value value) {
        Key* tmpKey;
        Value* tmpVal;
        if (size + 1 >= max_size) {
            max_size *= 2;
            tmpKey = new Key [max_size];
            tmpVal = new Value [max_size];
            for (int i = 0; i < size; i++) {
                tmpKey[i] = keys[i];
                tmpVal[i] = values[i];
            }
            tmpKey[size] = key;
            tmpVal[size] = value;
            delete[] keys;
            delete[] values;
            keys = tmpKey;
            values = tmpVal;
        }
        else {
            keys[size] = key;
            values[size] = value;
        }
        size++;
    }

    void print() {
        for (int i = 0; i < size; i++)
            cout << "{" << keys[i] << ", " << values[i] << "}" << endl;
    }
};

// Template partial specialization: Key is specified to be int.
template <class Value> class Dictionary<int, Value> {
    int* keys;
    Value* values;
    int size;
    int max_size;
public:
    Dictionary(int initial_size) : size(0) {
        max_size = 1;
        while (initial_size >= max_size)
            max_size *= 2;
        keys = new int[max_size];
        values = new Value[max_size];
    }
    void add(int key, Value value) {
        int* tmpKey;
        Value* tmpVal;
        if (size + 1 >= max_size) {
            max_size *= 2;
            tmpKey = new int [max_size];
            tmpVal = new Value [max_size];
            for (int i = 0; i < size; i++) {
                tmpKey[i] = keys[i];
                tmpVal[i] = values[i];
            }
            tmpKey[size] = key;
            tmpVal[size] = value;
            delete[] keys;
            delete[] values;
            keys = tmpKey;
        }
    }
}

```

```

        values = tmpVal;
    }
    else {
        keys[size] = key;
        values[size] = value;
    }
    size++;
}

void sort() {
    // Sort method is defined.
    int smallest = 0;
    for (int i = 0; i < size - 1; i++) {
        for (int j = i; j < size; j++) {
            if (keys[j] < keys[smallest])
                smallest = j;
        }
        swap(keys[i], keys[smallest]);
        swap(values[i], values[smallest]);
    }
}

void print() {
    for (int i = 0; i < size; i++)
        cout << "{" << keys[i] << ", " << values[i] << "}" << endl;
}
};

int main() {
    Dictionary<char*, char*>* dict = new Dictionary<char*, char*>(10);
    dict->print();
    dict->add("apple", "fruit");
    dict->add("banana", "fruit");
    dict->add("dog", "animal");
    dict->print();

    Dictionary<int, char*>* dict_specialized = new Dictionary<int, char*>(10);
    dict_specialized->print();
    dict_specialized->add(100, "apple");
    dict_specialized->add(101, "banana");
    dict_specialized->add(103, "dog");
    dict_specialized->add(89, "cat");
    dict_specialized->print();
    dict_specialized->sort();
    cout << endl << "Sorted list:" << endl;
    dict_specialized->print();
}
}

```

```

{apple, fruit}
{banana, fruit}
{dog, animal}
{100, apple}
{101, banana}
{103, dog}
{89, cat}

Sorted list:
{89, cat}
{100, apple}
{101, banana}
{103, dog}

```

# Templates and Name Resolution

12/21/2017 • 1 min to read • [Edit Online](#)

In template definitions, there are three types of names.

- Locally declared names, including the name of the template itself and any names declared inside the template definition.
- Names from the enclosing scope outside the template definition.
- Names that depend in some way on the template arguments, referred to as dependent names.

While the first two names also pertain to class and function scopes, special rules for name resolution are required in template definitions to deal with the added complexity of dependent names. This is because the compiler knows little about these names until the template is instantiated, because they could be totally different types depending on which template arguments are used. Nondependent names are looked up according to the usual rules and at the point of definition of the template. These names, being independent of the template arguments, are looked up once for all template specializations. Dependent names are not looked up until the template is instantiated and are looked up separately for each specialization.

A type is dependent if it depends on the template arguments. Specifically, a type is dependent if it is:

- The template argument itself:

```
T
```

- A qualified name with a qualification including a dependent type:

```
T::myType
```

- A qualified name if the unqualified part identifies a dependent type:

```
N::T
```

- A const or volatile type for which the base type is a dependent type:

```
const T
```

- A pointer, reference, array, or function pointer type based on a dependent type:

```
T *, T &, T [10], T (*)()
```

- An array whose size is based on a template parameter:

```
template <int arg> class X {  
    int x[arg] ; // dependent type  
}
```

- a template type constructed from a template parameter:

```
T<int>, MyTemplate<T>
```

## Type Dependence and Value Dependence

Names and expressions dependent on a template parameter are categorized as type dependent or value dependent, depending on whether the template parameter is a type parameter or a value parameter. Also, any identifiers declared in a template with a type dependent on the template argument are considered value dependent, as is a integral or enumeration type initialized with a value-dependent expression.

Type-dependent and value-dependent expressions are expressions that involve variables that are type dependent or value dependent. These expressions can have semantics that differ, depending on the parameters used for the template.

## See Also

[Templates](#)

# Name Resolution for Dependent Types

12/21/2017 • 2 min to read • [Edit Online](#)

Use **typename** for qualified names in template definitions to tell the compiler that the given qualified name identifies a type. For more information, see [typename](#).

```
// template_name_resolution1.cpp
#include <stdio.h>
template <class T> class X
{
public:
    void f(typename T::myType* mt) {}
};

class Yarg
{
public:
    struct myType { };
};

int main()
{
    X<Yarg> x;
    x.f(new Yarg::myType());
    printf("Name resolved by using typename keyword.");
}
```

Name resolved by using typename keyword.

Name lookup for dependent names examines names from both the context of the template definition—in the following example, this context would find `myFunction(char)`—and the context of the template instantiation. In the following example, the template is instantiated in `main`; therefore, the `MyNamespace::myFunction` is visible from the point of instantiation and is picked as the better match. If `MyNamespace::myFunction` were renamed, `myFunction(char)` would be called instead.

All names are resolved as if they were dependent names. Nevertheless, we recommend that you use fully qualified names if there is any possible conflict.

```

// template_name_resolution2.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

void myFunction(char)
{
    cout << "Char myFunction" << endl;
}

template <class T> class Class1
{
public:
    Class1(T i)
    {
        // If replaced with myFunction(1), myFunction(char)
        // will be called
        myFunction(i);
    }
};

namespace MyNamespace
{
    void myFunction(int)
    {
        cout << "Int MyNamespace::myFunction" << endl;
    }
};

using namespace MyNamespace;

int main()
{
    Class1<int>* c1 = new Class1<int>(100);
}

```

## Output

```
Int MyNamespace::myFunction
```

## Template Disambiguation

Visual C++ in Visual Studio 2012 enforces the C++98/03/11 standard rules for disambiguation with the "template" keyword. In the following example, Visual C++ 2010 would accept both the nonconforming lines and the conforming lines. Visual C++ in Visual Studio 2012 accepts only the conforming lines.

```

#include <iostream>
#include <ostream>
#include <typeinfo>
using namespace std;

template <typename T> struct Allocator {
    template <typename U> struct Rebind {
        typedef Allocator<U> Other;
    };
};

template <typename X, typename AY> struct Container {
    #if defined(NONCONFORMANT)
        typedef typename AY::Rebind<X>::Other AX; // nonconformant
    #elif defined(CONFORMANT)
        typedef typename AY::template Rebind<X>::Other AX; // conformant
    #else
        #error Define NONCONFORMANT or CONFORMANT.
    #endif
};

int main() {
    cout << typeid(Container<int, Allocator<float>>::AX).name() << endl;
}

```

Conformance with the disambiguation rules is required because, by default, C++ assumes that `AY::Rebind` isn't a template, and so the compiler interprets the following "`<`" as a less-than. It has to know that `Rebind` is a template so that it can correctly parse "`<`" as an angle bracket.

## See Also

[Name Resolution](#)

# Name Resolution for Locally Declared Names

12/21/2017 • 3 min to read • [Edit Online](#)

The template's name itself can be referred to with or without the template arguments. In the scope of a class template, the name itself refers to the template. In the scope of a template specialization or partial specialization, the name alone refers to the specialization or partial specialization. Other specializations or partial specializations of the template can also be referenced, with the appropriate template arguments.

## Example

The following code shows that the class template's name A is interpreted differently in the scope of a specialization or partial specialization.

```
// template_name_resolution3.cpp
// compile with: /c
template <class T> class A {
    A* a1;    // A refers to A<T>
    A<int>* a2; // A<int> refers to a specialization of A.
    A<T*>* a3; // A<T*> refers to the partial specialization A<T*>.
};

template <class T> class A<T*> {
    A* a4; // A refers to A<T*>.
};

template<> class A<int> {
    A* a5; // A refers to A<int>.
};
```

## Example

In the case of a name conflict between a template parameter and another object, the template parameter can or cannot be hidden. The following rules will help determine precedence.

The template parameter is in scope from the point where it first appears until the end of the class or function template. If the name appears again in the template argument list or in the list of base classes, it refers to the same type. In standard C++, no other name that is identical to the template parameter can be declared in the same scope. A Microsoft extension allows the template parameter to be redefined in the scope of the template. The following example shows using the template parameter in the base specification of a class template.

```
// template_name_resolution4.cpp
// compile with: /EHsc
template <class T>
class Base1 {};

template <class T>
class Derived1 : Base1<T> {};

int main() {
    // Derived1<int> d;
}
```

## Example

When defining a template's member functions outside the class template, a different template parameter name can be used. If the template member function definition uses a different name for the template parameter than the declaration does, and the name used in the definition conflicts with another member of the declaration, the member in the template declaration takes precedence.

```
// template_name_resolution5.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

template <class T> class C {
public:
    struct Z {
        Z() { cout << "Z::Z()" << endl; }
    };
    void f();
};

template <class Z>
void C<Z>::f() {
    // Z refers to the struct Z, not to the template arg;
    // Therefore, the constructor for struct Z will be called.
    Z z;
}

int main() {
    C<int> c;
    c.f();
}
```

```
Z::Z()
```

## Example

When defining a template function or member function outside the namespace in which the template was declared, the template argument takes precedence over the names of other members of the namespace.

```
// template_name_resolution6.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

namespace NS {
    void g() { cout << "NS::g" << endl; }

    template <class T> struct C {
        void f();
        void g() { cout << "C<T>::g" << endl; }
    };
}

template <class T>
void NS::C<T>::f() {
    g(); // C<T>::g, not NS::g
}

int main() {
    NS::C<int> c;
    c.f();
}
```

## Example

In definitions that are outside of the template class declaration, if a template class has a base class that does not depend on a template argument and if the base class or one of its members has the same name as a template argument, then the base class or member name hides the template argument.

```
// template_name_resolution7.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

struct B {
    int i;
    void print() { cout << "Base" << endl; }
};

template <class T, int i> struct C : public B {
    void f();
};

template <class B, int i>
void C<B, i>::f() {
    B b; // Base class b, not template argument.
    b.print();
    i = 1; // Set base class's i to 1.
}

int main() {
    C<int, 1> c;
    c.f();
    cout << c.i << endl;
}
```

```
Base
1
```

## See Also

[Name Resolution](#)

# Overload Resolution of Function Template Calls

12/21/2017 • 1 min to read • [Edit Online](#)

A function template can overload nontemplate functions of the same name. In this scenario, function calls are resolved by first using template argument deduction to instantiate the function template with a unique specialization. If template argument deduction fails, the other function overloads are considered to resolve the call. These other overloads, also known as the candidate set, include nontemplate functions and other instantiated function templates. If template argument deduction succeeds, then the generated function is compared with the other functions to determine the best match, following the rules for overload resolution. For more information, see [Function Overloading](#).

## Example

If a nontemplate function is an equally good match to a template function, the nontemplate function is chosen (unless the template arguments were explicitly specified), as in the call `f(1, 1)` in the following example.

```
// template_name_resolution9.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

void f(int, int) { cout << "f(int, int)" << endl; }
void f(char, char) { cout << "f(char, char)" << endl; }

template <class T1, class T2>
void f(T1, T2)
{
    cout << "void f(T1, T2)" << endl;
};

int main()
{
    f(1, 1);    // Equally good match; choose the nontemplate function.
    f('a', 1); // Chooses the template function.
    f<int, int>(2, 2); // Template arguments explicitly specified.
}
```

```
f(int, int)
void f(T1, T2)
void f(T1, T2)
```

## Example

The next example illustrates that the exactly matching template function is preferred if the nontemplate function requires a conversion.

```
// template_name_resolution10.cpp
// compile with: /EHsc
#include <iostream>
using namespace std;

void f(int, int) { cout << "f(int, int)" << endl; }

template <class T1, class T2>
void f(T1, T2)
{
    cout << "void f(T1, T2)" << endl;
};

int main()
{
    long l = 0;
    int i = 0;
    // Call the template function f(long, int) because f(int, int)
    // would require a conversion from long to int.
    f(l, i);
}
```

```
void f(T1, T2)
```

## See Also

[Name Resolution](#)

[typename](#)

# Source code organization (C++ Templates)

9/22/2017 • 3 min to read • [Edit Online](#)

When defining a class template, you must organize the source code in such a way that the member definitions are visible to the compiler when it needs them. You have the choice of using the *inclusion model* or the *explicit instantiation* model. In the inclusion model, you include the member definitions in every file that uses a template. This approach is simplest and provides maximum flexibility in terms of what concrete types can be used with your template. Its disadvantage is that it can increase compilation times. The impact can be significant if a project and/or the included files themselves are large. With the explicit instantiation approach, the template itself instantiates concrete classes or class members for specific types. This approach can speed up compilation times, but it limits usage to only those classes that the template implementer has enabled ahead of time. In general, we recommend that you use the inclusion model unless the compilation times become a problem.

## Background

Templates are not like ordinary classes in the sense that the compiler does not generate object code for a template or any of its members. There is nothing to generate until the template is instantiated with concrete types. When the compiler encounters a template instantiation such as `MyClass<int> mc;` and no class with that signature exists yet, it generates a new class. It also attempts to generate code for any member functions that are used. If those definitions are in a file that is not #included, directly or indirectly, in the .cpp file that is being compiled, the compiler can't see them. From the compiler's point of view, this isn't necessarily an error because the functions may be defined in another translation unit, in which case the linker will find them. If the linker does not find that code, it raises an **unresolved external** error.

## The inclusion model

The simplest and most common way to make template definitions visible throughout a translation unit, is to put the definitions in the header file itself. Any .cpp file that uses the template simply has to #include the header. This is the approach used in the Standard Library.

```

#ifndef MYARRAY
#define MYARRAY
#include <iostream>

template<typename T, size_t N>
class MyArray
{
    T arr[N];
public:
    // Full definitions:
    MyArray(){}
    void Print()
    {
        for (const auto v : arr)
        {
            std::cout << v << " , ";
        }
    }

    T& operator[](int i)
    {
        return arr[i];
    }
};

#endif

```

With this approach, the compiler has access to the complete template definition and can instantiate templates on-demand for any type. It is simple and relatively easy to maintain. However, the inclusion model does have a cost in terms of compilation times. This cost can be significant in large programs, especially if the template header itself #includes other headers. Every .cpp file that #includes the header will get its own copy of the function templates and all the definitions. The linker will generally be able to sort things out so that you do not end up with multiple definitions for a function, but it takes time to do this work. In smaller programs that extra compilation time is probably not significant.

## The explicit instantiation model

If the inclusion model is not viable for your project, and you know definitively the set of types that will be used to instantiate a template, then you can separate out the template code into an .h and .cpp file, and in the .cpp file explicitly instantiate the templates. This will cause object code to be generated that the compiler will see when it encounters user instantiations.

You create an explicit instantiation by using the keyword `template` followed by the signature of the entity you want to instantiate. This can be a type or a member. If you explicitly instantiate a type, all members are instantiated.

```
template MyArray<double, 5>;
```

```

//MyArray.h
#ifndef MYARRAY
#define MYARRAY

template<typename T, size_t N>
class MyArray
{
    T arr[N];
public:
    MyArray();
    void Print();
    T& operator[](int i);
};

#endif

//MyArray.cpp
#include "stdafx.h"
#include <iostream>
#include "MyArray.h"

using namespace std;

template<typename T, size_t N>
MyArray<T,N>::MyArray(){}

template<typename T, size_t N>
void MyArray<T,N>::Print()
{
    for(const auto v : arr)
    {
        cout << v << " ";
    }
    cout << endl;
}

template MyArray<double, 5>;template MyArray<string, 5>;

```

In the previous example, the explicit instantiations are at the bottom of the .cpp file. A `MyArray` may be used only for **double** or **string** types.

#### NOTE

In C++11 the **export** keyword was deprecated in the context of template definitions. In practical terms this has little impact because most compilers never supported it.

# Event Handling

12/21/2017 • 1 min to read • [Edit Online](#)

Event handling is primarily supported for COM classes (C++ classes that implement COM objects, typically using ATL classes or the `coclass` attribute). For more information, see [Event Handling in COM](#).

Event handling is also supported for native C++ classes (C++ classes that do not implement COM objects), however, that support is deprecated and will be removed in a future release. For more information, see [Event Handling in Native C++](#).

Event handling supports single- and multithreaded usage and protects data from simultaneous multithread access. It also allows you to derive subclasses from event source or receiver classes and support extended event sourcing/receiving in the derived class.

Visual C++ includes attributes and keywords for declaring events and event handlers. The event attributes and keywords can be used in CLR programs and in native C++ programs.

TOPIC	DESCRIPTION
<a href="#">event_source</a>	Creates an event source.
<a href="#">event_receiver</a>	Creates an event receiver (sink).
<a href="#">_event</a>	Declares an event.
<a href="#">_raise</a>	Emphasizes the call site of an event.
<a href="#">_hook</a>	Associates a handler method with an event.
<a href="#">_unhook</a>	Dissociates a handler method from an event.

## See Also

[C++ Language Reference](#)

[Keywords](#)

[Event Handling Samples](#)

# `__event`

12/21/2017 • 3 min to read • [Edit Online](#)

Declares an event.

## Syntax

```
__event
method-declarator
;
__event __interface interface-specifier;
__event member-declarator;
```

## Remarks

The keyword `__event` can be applied to a method declaration, an interface declaration, or a data member declaration. However, you cannot use the `__event` keyword to qualify a member of a nested class.

Depending on whether your event source and receiver are native C++, COM, or managed (.NET Framework), you can use the following constructs as events:

NATIVE C++	COM	MANAGED (.NET FRAMEWORK)
Method	—	method
—	interface	—
—	—	data member

Use `__hook` in an event receiver to associate a handler method with an event method. Note that after you create an event with the `__event` keyword, all event handlers subsequently hooked to that event will be called when the event is called.

An `__event` method declaration cannot have a definition; a definition is implicitly generated, so the event method can be called as if it were any ordinary method.

### NOTE

A templated class or struct cannot contain events.

## Native Events

Native events are methods. The return type is typically `HRESULT` or `void`, but can be any integral type, including an `enum`. When an event uses an integral return type, an error condition is defined when an event handler returns a nonzero value, in which case the event being raised will call the other delegates.

```
// Examples of native C++ events:  
__event void OnDoubleClick();  
__event HRESULT OnClick(int* b, char* s);
```

See [Event Handling in Native C++](#) for sample code.

## COM Events

COM events are interfaces. The parameters of a method in an event source interface should be **in** parameters (but this is not rigorously enforced), because an **out** parameter is not useful when multicasting. A level 1 warning will be issued if you use an **out** parameter.

The return type is typically `HRESULT` or `void`, but can be any integral type, including `enum`. When an event uses an integral return type and an event handler returns a nonzero value, it is an error condition, in which case the event being raised aborts calls to the other delegates. Note that the compiler will automatically mark an event source interface as a [source](#) in the generated IDL.

The `__interface` keyword is always required after `__event` for a COM event source.

```
// Example of a COM event:  
__event __interface IEvent1;
```

See [Event Handling in COM](#) for sample code.

## Managed Events

For information on coding events in the new syntax, see [event](#).

Managed events are data members or methods. When used with an event, the return type of a delegate must be compliant with the [Common Language Specification](#). The return type of the event handler must match the return type of the delegate. For more information on delegates, see [Delegates and Events](#). If a managed event is a data member, its type must be a pointer to a delegate.

In the .NET Framework, you can treat a data member as if it were a method itself (that is, the `Invoke` method of its corresponding delegate). You must predefine the delegate type for declaring a managed event data member. In contrast, a managed event method implicitly defines the corresponding managed delegate if it is not already defined. For example, you can declare an event value such as `OnClick` as an event as follows:

```
// Examples of managed events:  
__event ClickEventHandler* OnClick; // data member as event  
__event void OnClick(String* s); // method as event
```

When implicitly declaring a managed event, you can specify add and remove accessors that will be called when event handlers are added or removed. You can also define the method that calls (raises) the event from outside the class.

## Example: Native Events

```
// EventHandling_Native_Event.cpp
// compile with: /c
[event_source(native)]
class CSource {
public:
    __event void MyEvent(int nValue);
};
```

## Example: COM Events

```
// EventHandling_COM_Event.cpp
// compile with: /c
#define _ATL_ATTRIBUTES 1
#include <atlbase.h>
#include <atlcom.h>

[ module(dll, name="EventSource", uuid="6E46B59E-89C3-4c15-A6D8-B8A1CEC98830") ];

[ dual, uuid("00000000-0000-0000-0000-000000000002") ]
__interface IEventSource {
    [id(1)] HRESULT MyEvent();
};

[ coclass, uuid("00000000-0000-0000-0000-000000000003"), event_source(com) ]
class CSource : public IEventSource {
public:
    __event __interface IEventSource;
    HRESULT FireEvent() {
        __raise MyEvent();
        return S_OK;
    }
};
```

## See Also

[Keywords](#)

[Event Handling](#)

[event\\_source](#)

[event\\_receiver](#)

[\\_hook](#)

[\\_unhook](#)

[\\_raise](#)

# `_hook`

12/21/2017 • 2 min to read • [Edit Online](#)

Associates a handler method with an event.

## Syntax

```
long __hook(
    &SourceClass::EventMethod,
    source,
    &ReceiverClass::HandlerMethod
    [, receiver = this]
);
long __hook(
    interface,
    source
);
```

### Parameters

`& SourceClass :: EventMethod`

A pointer to the event method to which you hook the event handler method:

- Native C++ events: `SourceClass` is the event source class and `EventMethod` is the event.
- COM events: `SourceClass` is the event source interface and `EventMethod` is one of its methods.
- Managed events: `SourceClass` is the event source class and `EventMethod` is the event.

`interface`

The interface name being hooked to `receiver`, only for COM event receivers in which the `layout_dependent` parameter of the `event_receiver` attribute is `true`.

`source`

A pointer to an instance of the event source. Depending on the code `type` specified in `event_receiver`, `source` can be one of the following:

- A native event source object pointer.
- An **IUnknown**-based pointer (COM source).
- A managed object pointer (for managed events).

`& ReceiverClass :: HandlerMethod`

A pointer to the event handler method to be hooked to an event. The handler is specified as a method of a class or a reference to the same; if you do not specify the class name, `_hook` assumes the class to be that in which it is called.

- Native C++ events: `ReceiverClass` is the event receiver class and `HandlerMethod` is the handler.
- COM events: `ReceiverClass` is the event receiver interface and `HandlerMethod` is one of its handlers.
- Managed events: `ReceiverClass` is the event receiver class and `HandlerMethod` is the handler.

`receiver` (optional)

A pointer to an instance of the event receiver class. If you do not specify a receiver, the default is the

receiver class or structure in which `__hook` is called.

## Usage

Can be used in any function scope, including main, outside the event receiver class.

## Remarks

Use the intrinsic function `__hook` in an event receiver to associate or hook a handler method with an event method. The specified handler is then called when the source raises the specified event. You can hook several handlers to a single event or hook several events to a single handler.

There are two forms of `__hook`. You can use the first (four-argument) form in most cases, specifically, for COM event receivers in which the *layout\_dependent* parameter of the `event_receiver` attribute is **false**.

In these cases you do not need to hook all methods in an interface before firing an event on one of the methods; only the method handling the event needs to be hooked. You can use the second (two-argument) form of `__hook` only for a COM event receiver in which *layout\_dependent=true*.

`__hook` returns a long value. A nonzero return value indicates that an error has occurred (managed events throw an exception).

The compiler checks for the existence of an event and that the event signature agrees with the delegate signature.

With the exception of COM events, `__hook` and `__unhook` can be called outside the event receiver.

An alternative to using `__hook` is to use the `+=` operator.

For information on coding managed events in the new syntax, see [event](#).

### NOTE

A templated class or struct cannot contain events.

## Example

See [Event Handling in Native C++](#) and [Event Handling in COM](#) for samples.

## See Also

[Keywords](#)

[Event Handling](#)

[event\\_source](#)

[event\\_receiver](#)

[\\_\\_unhook](#)

[\\_\\_raise](#)

# `_raise`

12/21/2017 • 1 min to read • [Edit Online](#)

Emphasizes the call site of an event.

## Syntax

```
_raise  
method-declarator  
;
```

## Remarks

From managed code, an event can only be raised from within the class where it is defined. See [event](#) for more information.

The keyword `_raise` causes an error to be emitted if you call a non-event.

### NOTE

A templated class or struct cannot contain events.

## Example

```
// EventHandlingRef_raise.cpp  
struct E {  
    __event void func1();  
    void func1(int) {}  
  
    void func2() {}  
  
    void b() {  
        __raise func1();  
        __raise func1(1); // C3745: 'int Event::bar(int)':  
                           // only an event can be 'raised'  
        __raise func2(); // C3745  
    }  
};  
  
int main() {  
    E e;  
    __raise e.func1();  
    __raise e.func1(1); // C3745  
    __raise e.func2(); // C3745  
}
```

## See Also

[Keywords](#)

[Event Handling](#)

[Component Extensions for Runtime Platforms](#)



# `_unhook`

12/21/2017 • 2 min to read • [Edit Online](#)

Dissociates a handler method from an event.

## Syntax

```
long __unhook(
    &SourceClass::EventMethod,
    source,
    &ReceiverClass::HandlerMethod
    [, receiver = this]
);
long __unhook(
    interface,
    source
);
long __unhook(
    source
);
```

### Parameters

**& SourceClass** :: *EventMethod*

A pointer to the event method from which you unhook the event handler method:

- Native C++ events: *SourceClass* is the event source class and *EventMethod* is the event.
- COM events: *SourceClass* is the event source interface and *EventMethod* is one of its methods.
- Managed events: *SourceClass* is the event source class and *EventMethod* is the event.

**interface**

The interface name being unhooked from **receiver**, only for COM event receivers in which the *layout\_dependent* parameter of the [event\\_receiver](#) attribute is **true**.

**source**

A pointer to an instance of the event source. Depending on the code **type** specified in [event\\_receiver](#), *source* can be one of the following:

- A native event source object pointer.
- An **IUnknown**-based pointer (COM source).
- A managed object pointer (for managed events).

**& ReceiverClass** :: **HandlerMethod**

A pointer to the event handler method to be unhooked from an event. The handler is specified as a method of a class or a reference to the same; if you do not specify the class name, `_unhook` assumes the class to be that in which it is called.

- Native C++ events: *ReceiverClass* is the event receiver class and **HandlerMethod** is the handler.
- COM events: *ReceiverClass* is the event receiver interface and **HandlerMethod** is one of its handlers.
- Managed events: *ReceiverClass* is the event receiver class and **HandlerMethod** is the handler.

`receiver` (optional)

A pointer to an instance of the event receiver class. If you do not specify a receiver, the default is the receiver class or structure in which `__unhook` is called.

## Usage

Can be used in any function scope, including main, outside the event receiver class.

## Remarks

Use the intrinsic function `__unhook` in an event receiver to dissociate or "unhook" a handler method from an event method.

There are three forms of `__unhook`. You can use the first (four-argument) form in most cases. You can use the second (two-argument) form of `__unhook` only for a COM event receiver; this unhooks the entire event interface. You can use the third (one-argument) form to unhook all delegates from the specified source.

A nonzero return value indicates that an error has occurred (managed events will throw an exception).

If you call `__unhook` on an event and event handler that are not already hooked, it will have no effect.

At compile time, the compiler verifies that the event exists and does parameter type checking with the specified handler.

With the exception of COM events, `__hook` and `__unhook` can be called outside the event receiver.

An alternative to using `__unhook` is to use the `-=` operator.

For information on coding managed events in the new syntax, see [event](#).

### NOTE

A templated class or struct cannot contain events.

## Example

See [Event Handling in Native C++](#) and [Event Handling in COM](#) for samples.

## See Also

[Keywords](#)

[event\\_source](#)

[event\\_receiver](#)

[\\_\\_event](#)

[\\_\\_hook](#)

[\\_\\_raise](#)

# Event Handling in Native C++

12/21/2017 • 1 min to read • [Edit Online](#)

In native C++ event handling, you set up an event source and event receiver using the `event_source` and `event_receiver` attributes, respectively, specifying `type = native`. These attributes allow the classes to which they are applied to fire events and handle events in a native, non-COM context.

## Declaring Events

In an event source class, use the `_event` keyword on a method declaration to declare the method as an event. Make sure to declare the method, but do not define it; to do so will generate a compiler error, because the compiler defines the method implicitly when it is made into an event. Native events can be methods with zero or more parameters. The return type can be void or any integral type.

## Defining Event Handlers

In an event receiver class, you define event handlers, which are methods with signatures (return types, calling conventions, and arguments) that match the event that they will handle.

## Hooking Event Handlers to Events

Also in an event receiver class, you use the intrinsic function `_hook` to associate events with event handlers and `_unhook` to dissociate events from event handlers. You can hook several events to an event handler, or several event handlers to an event.

## Firing Events

To fire an event, simply call the method declared as an event in the event source class. If handlers have been hooked to the event, the handlers will be called.

### Native C++ Event Code

The following example shows how to fire an event in native C++. To compile and run the example, refer to the comments in the code.

## Example

### Code

```

// evh_native.cpp
#include <stdio.h>

[event_source(native)]
class CSource {
public:
    __event void MyEvent(int nValue);
};

[event_receiver(native)]
class CReceiver {
public:
    void MyHandler1(int nValue) {
        printf_s("MyHandler1 was called with value %d.\n", nValue);
    }

    void MyHandler2(int nValue) {
        printf_s("MyHandler2 was called with value %d.\n", nValue);
    }

    void hookEvent(CSource* pSource) {
        __hook(&CSource::MyEvent, pSource, &CReceiver::MyHandler1);
        __hook(&CSource::MyEvent, pSource, &CReceiver::MyHandler2);
    }

    void unhookEvent(CSource* pSource) {
        __unhook(&CSource::MyEvent, pSource, &CReceiver::MyHandler1);
        __unhook(&CSource::MyEvent, pSource, &CReceiver::MyHandler2);
    }
};

int main() {
    CSource source;
    CReceiver receiver;

    receiver.hookEvent(&source);
    __raise source.MyEvent(123);
    receiver.unhookEvent(&source);
}

```

## Output

```

MyHandler2 was called with value 123.
MyHandler1 was called with value 123.

```

## See Also

[Event Handling](#)

# Event Handling in COM

12/21/2017 • 4 min to read • [Edit Online](#)

In COM event handling, you set up an event source and event receiver using the `event_source` and `event_receiver` attributes, respectively, specifying `type =com`. These attributes inject the appropriate code for custom, dispatch, and dual interfaces to allow the classes to which they are applied to fire events and handle events through COM connection points.

## Declaring Events

In an event source class, use the `_event` keyword on an interface declaration to declare that interface's methods as events. The events of that interface are fired when you call them as interface methods. Methods on event interfaces can have zero or more parameters (which should all be `in` parameters). The return type can be void or any integral type.

## Defining Event Handlers

In an event receiver class, you define event handlers, which are methods with signatures (return types, calling conventions, and arguments) that match the event that they will handle. For COM events, calling conventions do not have to match; see [Layout Dependent COM Events](#) below for details.

## Hooking Event Handlers to Events

Also in an event receiver class, you use the intrinsic function `_hook` to associate events with event handlers and `_unhook` to dissociate events from event handlers. You can hook several events to an event handler, or several event handlers to an event.

### NOTE

Typically, there are two techniques to allow a COM event receiver to access event source interface definitions. The first, as shown below, is to share a common header file. The second is to use `#import` with the `embedded_idl` import qualifier, so that the event source type library is written to the .tlh file with the attribute-generated code preserved.

## Firing Events

To fire an event, simply call a method in the interface declared with the `_event` keyword in the event source class. If handlers have been hooked to the event, the handlers will be called.

### COM Event Code

The following example shows how to fire an event in a COM class. To compile and run the example, refer to the comments in the code.

```

// evh_server.h
#pragma once

[ dual, uuid("00000000-0000-0000-0000-000000000001") ]
__interface IEvents {
    [id(1)] HRESULT MyEvent([in] int value);
};

[ dual, uuid("00000000-0000-0000-0000-000000000002") ]
__interface IEventSource {
    [id(1)] HRESULT FireEvent();
};

class DECLSPEC_UUID("530DF3AD-6936-3214-A83B-27B63C7997C4") CSource;

```

And then the server:

```

// evh_server.cpp
// compile with: /LD
// post-build command: Regsvr32.exe /s evh_server.dll
#define _ATL_ATTRIBUTES 1
#include <atlbase.h>
#include <atlcom.h>
#include "evh_server.h"

[ module(dll, name="EventSource", uuid="6E46B59E-89C3-4c15-A6D8-B8A1CEC98830") ];

[coclass, event_source(com), uuid("530DF3AD-6936-3214-A83B-27B63C7997C4")]
class CSource : public IEventSource {
public:
    __event __interface IEvents;

    HRESULT FireEvent() {
        __raise MyEvent(123);
        return S_OK;
    }
};

```

And then the client:

```

// evh_client.cpp
// compile with: /link /OPT:NOREF
#define _ATL_ATTRIBUTES 1
#include <atibase.h>
#include <atlcom.h>
#include <stdio.h>
#include "evh_server.h"

[ module(name="EventReceiver") ];

[ event_receiver(com) ]
class CReceiver {
public:
    HRESULT MyHandler1(int nValue) {
        printf_s("MyHandler1 was called with value %d.\n", nValue);
        return S_OK;
    }

    HRESULT MyHandler2(int nValue) {
        printf_s("MyHandler2 was called with value %d.\n", nValue);
        return S_OK;
    }

    void HookEvent(IEventSource* pSource) {
        __hook(&IEvents::MyEvent, pSource, &CReceiver::MyHandler1);
        __hook(&IEvents::MyEvent, pSource, &CReceiver::MyHandler2);
    }

    void UnhookEvent(IEventSource* pSource) {
        __unhook(&IEvents::MyEvent, pSource, &CReceiver::MyHandler1);
        __unhook(&IEvents::MyEvent, pSource, &CReceiver::MyHandler2);
    }
};

int main() {
    // Create COM object
    CoInitialize(NULL);
    {
        IEventSource* pSource = 0;
        HRESULT hr = CoCreateInstance(__uuidof(CSource), NULL, CLSCTX_ALL, __uuidof(IEventSource), (void**) &pSource);
        if (FAILED(hr)) {
            return -1;
        }

        // Create receiver and fire event
        CReceiver receiver;
        receiver.HookEvent(pSource);
        pSource->FireEvent();
        receiver.UnhookEvent(pSource);
    }
    CoUninitialize();
    return 0;
}

```

## Output

```

MyHandler1 was called with value 123.
MyHandler2 was called with value 123.

```

## Layout Dependent COM Events

Layout dependency is only an issue for COM programming. In native and managed event handling, the signatures

(return type, calling convention, and arguments) of the handlers must match their events, but the handler names do not have to match their events.

However, in COM event handling, when you set the *layout\_dependent* parameter of **event\_receiver** to **true**, the name and signature matching is enforced. This means that the names and signatures of the handlers in the event receiver must exactly match the names and signatures of the events to which they are hooked.

When *layout\_dependent* is set to **false**, the calling convention and storage class (virtual, static, and so on) can be mixed and matched between the firing event method and the hooking methods (its delegates). It is slightly more efficient to have *layout\_dependent=true*.

For example, suppose `IEventSource` is defined to have the following methods:

```
[id(1)] HRESULT MyEvent1([in] int value);
[id(2)] HRESULT MyEvent2([in] int value);
```

Assume the event source has the following form:

```
[coclass, event_source(com)]
class CSource : public IEventSource {
public:
    __event __interface IEvents;

    HRESULT FireEvent() {
        MyEvent1(123);
        MyEvent2(123);
        return S_OK;
    }
};
```

Then, in the event receiver, any handler hooked to a method in `IEventSource` must match its name and signature, as follows:

```
[coclass, event_receiver(com, true)]
class CReceiver {
public:
    HRESULT MyEvent1(int nValue) { // name and signature matches MyEvent1
        ...
    }
    HRESULT MyEvent2(E c, char* pc) { // signature doesn't match MyEvent2
        ...
    }
    HRESULT MyHandler1(int nValue) { // name doesn't match MyEvent1 (or 2)
        ...
    }
    void HookEvent(IEventSource* pSource) {
        __hook(IFace, pSource); // Hooks up all name-matched events
                               // under layout_dependent = true
        __hook(&IFace::MyEvent1, pSource, &CReceive::MyEvent1); // valid
        __hook(&IFace::MyEvent2, pSource, &CSink::MyEvent2); // not valid
        __hook(&IFace::MyEvent1, pSource, &CSink:: MyHandler1); // not valid
    }
};
```

## See Also

[Event Handling](#)

# Microsoft-Specific Modifiers

12/21/2017 • 1 min to read • [Edit Online](#)

This section describes Microsoft-specific extensions to C++ in the following areas:

- [Based addressing](#), the practice of using a pointer as a base from which other pointers can be offset
- [Function calling conventions](#)
- Extended storage-class attributes declared with the `_declspec` keyword
- The `_w64` keyword

Many of the Microsoft-specific keywords can be used to modify declarators to form derived types. For more information about declarators, see [Declarators](#).

## Microsoft-Specific Keywords

KEYWORD	MEANING	USED TO FORM DERIVED TYPES?
<code>_based</code>	The name that follows declares a 32-bit offset to the 32-bit base contained in the declaration.	Yes
<code>_cdecl</code>	The name that follows uses the C naming and calling conventions.	Yes
<code>_declspec</code>	The name that follows specifies a Microsoft-specific storage-class attribute.	No
<code>_fastcall</code>	The name that follows declares a function that uses registers, when available, instead of the stack for argument passing.	Yes
<code>_restrict</code>	Similar to <code>_declspec(restrict)</code> , but for use on variables.	No
<code>_stdcall</code>	The name that follows specifies a function that observes the standard calling convention.	Yes
<code>_w64</code>	Marks a data type as being larger on a 64-bit compiler.	No
<code>_unaligned</code>	Specifies that a pointer to a type or other data is not aligned..	No
<code>_vectorcall</code>	The name that follows declares a function that uses registers, including SSE registers, when available, instead of the stack for argument passing.	Yes

## See Also

[C++ Language Reference](#)

# Based Addressing

12/21/2017 • 1 min to read • [Edit Online](#)

This section includes the following topics:

- [\\_\\_based Grammar](#)
- [Based Pointers](#)

## See Also

[Microsoft-Specific Modifiers](#)

# \_based Grammar

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Based addressing is useful when you need precise control over the segment in which objects are allocated (static and dynamic based data).

The only form of based addressing acceptable in 32-bit and 64-bit compilations is "based on a pointer" that defines a type that contains a 32-bit or 64-bit displacement to a 32-bit or 64-bit base or based on `void`.

## Grammar

*based-range-modifier:*

**\_based(** *base-expression* **)**

*base-expression:*

*based-variable**based-abstract-declarator**segment-name**segment-cast*

*based-variable:*

*identifier*

*based-abstract-declarator:*

*abstract-declarator*

*base-type:*

*type-name*

**END Microsoft Specific**

## See Also

[Based Pointers](#)

# Based Pointers (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `__based` keyword allows you to declare pointers based on pointers (pointers that are offsets from existing pointers).

## Syntax

```
type __based( base ) declarator
```

## Remarks

Pointers based on pointer addresses are the only form of the `__based` keyword valid in 32-bit or 64-bit compilations. For the Microsoft 32-bit C/C++ compiler, a based pointer is a 32-bit offset from a 32-bit pointer base. A similar restriction holds for 64-bit environments, where a based pointer is a 64-bit offset from the 64-bit base.

One use for pointers based on pointers is for persistent identifiers that contain pointers. A linked list that consists of pointers based on a pointer can be saved to disk, then reloaded to another place in memory, with the pointers remaining valid. For example:

```
// based_pointers1.cpp
// compile with: /c
void *vpBuffer;
struct llist_t {
    void __based( vpBuffer ) *vpData;
    struct llist_t __based( vpBuffer ) *llNext;
};
```

The pointer `vpBuffer` is assigned the address of memory allocated at some later point in the program. The linked list is relocated relative to the value of `vpBuffer`.

### NOTE

Persisting identifiers containing pointers can also be accomplished by using [memory-mapped files](#).

When dereferencing a based pointer, the base must be either explicitly specified or implicitly known through the declaration.

For compatibility with previous versions, `_based` is a synonym for `__based`.

## Example

The following code demonstrates changing a based pointer by changing its base.

```
// based_pointers2.cpp
// compile with: /EHsc
#include <iostream>

int a1[] = { 1,2,3 };
int a2[] = { 10,11,12 };
int *pBased;

typedef int __based(pBased) * pBasedPtr;

using namespace std;
int main() {
    pBased = &a1[0];
    pBasedPtr pb = 0;

    cout << *pb << endl;
    cout << *(pb+1) << endl;

    pBased = &a2[0];

    cout << *pb << endl;
    cout << *(pb+1) << endl;
}
```

```
1
2
10
11
```

## See Also

[Keywords](#)  
[alloc\\_text](#)

# Calling Conventions

12/21/2017 • 1 min to read • [Edit Online](#)

The Visual C/C++ compiler provides several different conventions for calling internal and external functions. Understanding these different approaches can help you debug your program and link your code with assembly-language routines.

The topics on this subject explain the differences between the calling conventions, how arguments are passed, and how values are returned by functions. They also discuss naked function calls, an advanced feature that enables you to write your own prolog and epilog code.

For information on calling conventions for x64 processors, see [Calling Convention](#).

## Topics in this section

- [Argument Passing and Naming Conventions](#) (`__cdecl`, `__stdcall`, `__fastcall`, and others)
- [Calling Example: Function Prototype and Call](#)
- [Using naked function calls to write custom prolog/epilog code](#)
- [Floating Point Coprocessor and Calling Conventions](#)
- [Obsolete calling conventions](#)

## See Also

[Microsoft-Specific Modifiers](#)

# Argument Passing and Naming Conventions

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The Visual C++ compilers allow you to specify conventions for passing arguments and return values between functions and callers. Not all conventions are available on all supported platforms, and some conventions use platform-specific implementations. In most cases, keywords or compiler switches that specify an unsupported convention on a particular platform are ignored, and the platform default convention is used.

On x86 platforms, all arguments are widened to 32 bits when they are passed. Return values are also widened to 32 bits and returned in the EAX register, except for 8-byte structures, which are returned in the EDX:EAX register pair. Larger structures are returned in the EAX register as pointers to hidden return structures. Parameters are pushed onto the stack from right to left. Structures that are not PODs will not be returned in registers.

The compiler generates prolog and epilog code to save and restore the ESI, EDI, EBX, and EBP registers, if they are used in the function.

### NOTE

When a struct, union, or class is returned from a function by value, all definitions of the type need to be the same, else the program may fail at runtime.

For information on how to define your own function prolog and epilog code, see [Naked Function Calls](#).

For information about the default calling conventions in code that targets x64 platforms, see [Overview of x64 Calling Conventions](#). For information about calling convention issues in code that targets ARM platforms, see [Common Visual C++ ARM Migration Issues](#).

The following calling conventions are supported by the Visual C/C++ compiler.

KEYWORD	STACK CLEANUP	PARAMETER PASSING
<code>_cdecl</code>	Caller	Pushes parameters on the stack, in reverse order (right to left)
<code>_clrcall</code>	n/a	Load parameters onto CLR expression stack in order (left to right).
<code>_stdcall</code>	Callee	Pushes parameters on the stack, in reverse order (right to left)
<code>_fastcall</code>	Callee	Stored in registers, then pushed on stack
<code>_thiscall</code>	Callee	Pushed on stack; <code>this</code> pointer stored in ECX
<code>_vectorcall</code>	Callee	Stored in registers, then pushed on stack in reverse order (right to left)

For related information, see [Obsolete Calling Conventions](#).

**END Microsoft Specific**

## See Also

[Calling Conventions](#)

# `_cdecl`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

`_cdecl` is the default calling convention for C and C++ programs. Because the stack is cleaned up by the caller, it can do **vararg** functions. The `_cdecl` calling convention creates larger executables than `_stdcall`, because it requires each function call to include stack cleanup code. The following list shows the implementation of this calling convention.

ELEMENT	IMPLEMENTATION
Argument-passing order	Right to left.
Stack-maintenance responsibility	Calling function pops the arguments from the stack.
Name-decoration convention	Underscore character (_) is prefixed to names, except when <code>_cdecl</code> functions that use C linkage are exported.
Case-translation convention	No case translation performed.

### NOTE

For related information, see [Decorated Names](#).

Place the `_cdecl` modifier before a variable or a function name. Because the C naming and calling conventions are the default, the only time you must use `_cdecl` in x86 code is when you have specified the **/Gv** (vectorcall), **/Gz** (stdcall), or **/Gr** (fastcall) compiler option. The **/Gd** compiler option forces the `_cdecl` calling convention.

On ARM and x64 processors, `_cdecl` is accepted but typically ignored by the compiler. By convention on ARM and x64, arguments are passed in registers when possible, and subsequent arguments are passed on the stack. In x64 code, use `_cdecl` to override the **/Gv** compiler option and use the default x64 calling convention.

For non-static class functions, if the function is defined out-of-line, the calling convention modifier does not have to be specified on the out-of-line definition. That is, for class non-static member methods, the calling convention specified during declaration is assumed at the point of definition. Given this class definition:

```
struct CMyClass {  
    void __cdecl mymethod();  
};
```

this:

```
void CMyClass::mymethod() { return; }
```

is equivalent to this:

```
void __cdecl CMyClass::mymethod() { return; }
```

## Example

In the following example, the compiler is instructed to use C naming and calling conventions for the `system` function.

```
// Example of the __cdecl keyword on function
int __cdecl system(const char *);

// Example of the __cdecl keyword on function pointer
typedef BOOL (__cdecl *funcname_ptr)(void * arg1, const char * arg2, DWORD flags, ...);
```

## See Also

[Argument Passing and Naming Conventions](#)

[Keywords](#)

# `_clrcall`

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

Specifies that a function can only be called from managed code. Use `_clrcall` for all virtual functions that will only be called from managed code. However this calling convention cannot be used for functions that will be called from native code.

Use `_clrcall` to improve performance when calling from a managed function to a virtual managed function or from managed function to managed function through pointer.

Entry points are separate, compiler-generated functions. If a function has both native and managed entry points, one of them will be the actual function with the function implementation. The other function will be a separate function (a thunk) that calls into the actual function and lets the common language runtime perform PInvoke. When marking a function as `_clrcall`, you indicate the function implementation must be MSIL and that the native entry point function will not be generated.

When taking the address of a native function if `_clrcall` is not specified, the compiler uses the native entry point. `_clrcall` indicates that the function is managed and there is no need to go through the transition from managed to native. In that case the compiler uses the managed entry point.

When **/clr** (not **/clr:pure** or **/clr:safe**) is used and `_clrcall` is not used, taking the address of a function always returns the address of the native entry point function. When `_clrcall` is used, the native entry point function is not created, so you get the address of the managed function, not an entry point thunk function. For more information, see [Double Thunking](#). The **/clr:pure** and **/clr:safe** compiler options are deprecated in Visual Studio 2015.

[/clr \(Common Language Runtime Compilation\)](#) implies that all functions and function pointers are `_clrcall` and the compiler will not permit a function inside the compiland to be marked anything other than `_clrcall`. When **/clr:pure** is used, `_clrcall` can only be specified on function pointers and external declarations.

You can directly call `_clrcall` functions from existing C++ code that was compiled by using **/clr** as long as that function has an MSIL implementation. `_clrcall` functions cannot be called directly from functions that have inline asm and call CPU-specific intrinsics, for example, even if those functions are compiled with **/clr**.

`_clrcall` function pointers are only meant to be used in the application domain in which they were created. Instead of passing `_clrcall` function pointers across application domains, use [CrossAppDomainDelegate](#). For more information, see [Application Domains and Visual C++](#).

## Example

Note that when a function is declared with `_clrcall`, code will be generated when needed; for example, when function is called.

```

// clrcall2.cpp
// compile with: /clr
using namespace System;
int __clrcall Func1() {
    Console::WriteLine("in Func1");
    return 0;
}

// Func1 hasn't been used at this point (code has not been generated),
// so runtime returns the address of a stub to the function
int (__clrcall *pf)() = &Func1;

// code calls the function, code generated at difference address
int i = pf(); // comment this line and comparison will pass

int main() {
    if (&Func1 == pf)
        Console::WriteLine("&Func1 == pf, comparison succeeds");
    else
        Console::WriteLine("&Func1 != pf, comparison fails");

    // even though comparison fails, stub and function call are correct
    pf();
    Func1();
}

```

```

in Func1
&Func1 != pf, comparison fails
in Func1
in Func1

```

## Example

The following sample shows that you can define a function pointer, such that you declare that the function pointer will only be invoked from managed code. This allows the compiler to directly call the managed function and avoid the native entry point (double thunk issue).

```

// clrcall3.cpp
// compile with: /clr
void Test() {
    System::Console::WriteLine("in Test");
}

int main() {
    void (*pTest)() = &Test;
    (*pTest)();

    void (__clrcall *pTest2)() = &Test;
    (*pTest2)();
}

```

## See Also

[Argument Passing and Naming Conventions](#)

[Keywords](#)

# `_stdcall`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_stdcall` calling convention is used to call Win32 API functions. The callee cleans the stack, so the compiler makes **vararg** functions `_cdecl`. Functions that use this calling convention require a function prototype.

## Syntax

```
return-type __stdcall function-name[(argument-list)]
```

## Remarks

The following list shows the implementation of this calling convention.

ELEMENT	IMPLEMENTATION
Argument-passing order	Right to left.
Argument-passing convention	By value, unless a pointer or reference type is passed.
Stack-maintenance responsibility	Called function pops its own arguments from the stack.
Name-decoration convention	An underscore (_) is prefixed to the name. The name is followed by the at sign (@) followed by the number of bytes (in decimal) in the argument list. Therefore, the function declared as <code>int func( int a, double b )</code> is decorated as follows: <code>_func@12</code>
Case-translation convention	None

The `/Gz` compiler option specifies `_stdcall` for all functions not explicitly declared with a different calling convention.

Functions declared using the `_stdcall` modifier return values the same way as functions declared using `_cdecl`.

On ARM and x64 processors, `_stdcall` is accepted and ignored by the compiler; on ARM and x64 architectures, by convention, arguments are passed in registers when possible, and subsequent arguments are passed on the stack.

For non-static class functions, if the function is defined out-of-line, the calling convention modifier does not have to be specified on the out-of-line definition. That is, for class non-static member methods, the calling convention specified during declaration is assumed at the point of definition. Given this class definition,

```
struct CMyClass {  
    void __stdcall mymethod();  
};
```

this

```
void CMyClass::mymethod() { return; }
```

is equivalent to this

```
void __stdcall CMyClass::mymethod() { return; }
```

## Example

In the following example, use of **\_\_stdcall** results in all **WINAPI** function types being handled as a standard call:

```
// Example of the __stdcall keyword
#define WINAPI __stdcall
// Example of the __stdcall keyword on function pointer
typedef BOOL (__stdcall *funcname_ptr)(void * arg1, const char * arg2, DWORD flags, ...);
```

## See Also

[Argument Passing and Naming Conventions](#)

[Keywords](#)

# `_fastcall`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_fastcall` calling convention specifies that arguments to functions are to be passed in registers, when possible. This calling convention only applies to the x86 architecture. The following list shows the implementation of this calling convention.

ELEMENT	IMPLEMENTATION
Argument-passing order	The first two DWORD or smaller arguments that are found in the argument list from left to right are passed in ECX and EDX registers; all other arguments are passed on the stack from right to left.
Stack-maintenance responsibility	Called function pops the arguments from the stack.
Name-decoration convention	At sign (@) is prefixed to names; an at sign followed by the number of bytes (in decimal) in the parameter list is suffixed to names.
Case-translation convention	No case translation performed.

### NOTE

Future compiler versions may use different registers to store parameters.

Using the `/Gr` compiler option causes each function in the module to compile as `_fastcall` unless the function is declared by using a conflicting attribute, or the name of the function is `main`.

The `_fastcall` keyword is accepted and ignored by the compilers that target ARM and x64 architectures; on an x64 chip, by convention, the first four arguments are passed in registers when possible, and additional arguments are passed on the stack. For more information, see [Overview of x64 Calling Conventions](#). On an ARM chip, up to four integer arguments and eight floating-point arguments may be passed in registers, and additional arguments are passed on the stack.

For non-static class functions, if the function is defined out-of-line, the calling convention modifier does not have to be specified on the out-of-line definition. That is, for class non-static member methods, the calling convention specified during declaration is assumed at the point of definition. Given this class definition:

```
struct CMyClass {  
    void __fastcall mymethod();  
};
```

this:

```
void CMyClass::mymethod() { return; }
```

is equivalent to this:

```
void __fastcall CMyClass::mymethod() { return; }
```

## Example

In the following example, the function `DeleteAggrWrapper` is passed arguments in registers:

```
// Example of the __fastcall keyword
#define FASTCALL __fastcall

void FASTCALL DeleteAggrWrapper(void* pWrapper);
// Example of the __fastcall keyword on function pointer
typedef BOOL (__fastcall *funcname_ptr)(void * arg1, const char * arg2, DWORD flags, ...);
```

## END Microsoft Specific

## See Also

[Argument Passing and Naming Conventions](#)

[Keywords](#)

# `_thiscall`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `_thiscall` calling convention is used on member functions and is the default calling convention used by C++ member functions that do not use variable arguments. Under `_thiscall`, the callee cleans the stack, which is impossible for `vararg` functions. Arguments are pushed on the stack from right to left, with the `this` pointer being passed via register ECX, and not on the stack, on the x86 architecture.

One reason to use `_thiscall` is in classes whose member functions use `_clrcall` by default. In that case, you can use `_thiscall` to make individual member functions callable from native code.

When compiling with `/clr:pure`, all functions and function pointers are `_clrcall` unless specified otherwise. The `/clr:pure` and `/clr:safe` compiler options are deprecated in Visual Studio 2015.

In releases before Visual C++ 2005, the `thiscall` calling convention could not be explicitly specified in a program, because `thiscall` was not a keyword.

`vararg` member functions use the `_cdecl` calling convention. All function arguments are pushed on the stack, with the `this` pointer placed on the stack last.

Because this calling convention applies only to C++, there is no C name decoration scheme.

On ARM and x64 machines, `_thiscall` is accepted and ignored by the compiler.

For non-static class functions, if the function is defined out-of-line, the calling convention modifier does not have to be specified on the out-of-line definition. That is, for class non-static member methods, the calling convention specified during declaration is assumed at the point of definition.

## END Microsoft Specific

## See Also

[Argument Passing and Naming Conventions](#)

# `_vectorcall`

12/21/2017 • 12 min to read • [Edit Online](#)

## Microsoft Specific

The `_vectorcall` calling convention specifies that arguments to functions are to be passed in registers, when possible. `_vectorcall` uses more registers for arguments than `_fastcall` or the default [x64 calling convention](#) use.

The `_vectorcall` calling convention is only supported in native code on x86 and x64 processors that include Streaming SIMD Extensions 2 (SSE2) and above. Use `_vectorcall` to speed functions that pass several floating-point or SIMD vector arguments and perform operations that take advantage of the arguments loaded in registers. The following list shows the features that are common to the x86 and x64 implementations of `_vectorcall`. The differences are explained later in this article.

ELEMENT	IMPLEMENTATION
C name-decoration convention	Function names are suffixed with two "at" signs (@@) followed by the number of bytes (in decimal) in the parameter list.
Case-translation convention	No case translation is performed.

Using the `/Gv` compiler option causes each function in the module to compile as `_vectorcall` unless the function is a member function, is declared with a conflicting calling convention attribute, uses a `vararg` variable argument list, or has the name `main`.

You can pass three kinds of arguments by register in `_vectorcall` functions: *integer type* values, *vector type* values, and *homogeneous vector aggregate* (HVA) values.

An integer type satisfies two requirements: it fits in the native register size of the processor—for example, 4 bytes on an x86 machine or 8 bytes on an x64 machine—and it's convertible to an integer of register length and back again without changing its bit representation. For example, any type that can be promoted to `int` on x86 (`long long` on x64)—for example, a `char` or `short`—or that can be cast to `int` (`long long` on x64) and back to its original type without change is an integer type. Integer types include pointer, reference, and `struct` or `union` types of 4 bytes (8 bytes on x64) or less. On x64 platforms, larger `struct` and `union` types are passed by reference to memory allocated by the caller; on x86 platforms, they are passed by value on the stack.

A vector type is either a floating-point type—for example, a `float` or `double`—or an SIMD vector type—for example, `_m128` or `_m256`.

An HVA type is a composite type of up to four data members that have identical vector types. An HVA type has the same alignment requirement as the vector type of its members. This is an example of an HVA `struct` definition that contains three identical vector types and has 32-byte alignment:

```
typedef struct {
    _m256 x;
    _m256 y;
    _m256 z;
} hva3;      // 3 element HVA type on _m256
```

Declare your functions explicitly with the `_vectorcall` keyword in header files to allow separately compiled code to link without errors. Functions must be prototyped to use `_vectorcall`, and can't use a `vararg` variable length argument list.

A member function may be declared by using the `__vectorcall` specifier. The hidden `this` pointer is passed by register as the first integer type argument.

On ARM machines, `__vectorcall` is accepted and ignored by the compiler.

For non-static class member functions, if the function is defined out-of-line, the calling convention modifier does not have to be specified on the out-of-line definition. That is, for class non-static members, the calling convention specified during declaration is assumed at the point of definition. Given this class definition:

```
struct MyClass {  
    void __vectorcall mymethod();  
};
```

this:

```
void MyClass::mymethod() { return; }
```

is equivalent to this:

```
void __vectorcall MyClass::mymethod() { return; }
```

The `__vectorcall` calling convention modifier must be specified when a pointer to a `__vectorcall` function is created. The next example creates a `typedef` for a pointer to a `__vectorcall` function that takes four `double` arguments and returns an `__m256` value:

```
typedef __m256 (__vectorcall * vcfnptr)(double, double, double, double);
```

## `__vectorcall` convention on x64

The `__vectorcall` calling convention on x64 extends the standard x64 calling convention to take advantage of additional registers. Both integer type arguments and vector type arguments are mapped to registers based on position in the argument list. HVA arguments are allocated to unused vector registers.

When any of the first four arguments in order from left to right are integer type arguments, they are passed in the register that corresponds to that position—RCX, RDX, R8, or R9. A hidden `this` pointer is treated as the first integer type argument. When an HVA argument in one of the first four arguments can't be passed in the available registers, a reference to caller-allocated memory is passed in the corresponding integer type register instead. Integer type arguments after the fourth parameter position are passed on the stack.

When any of the first six arguments in order from left to right are vector type arguments, they are passed by value in SSE vector registers 0 to 5 according to argument position. Floating-point and `__m128` types are passed in XMM registers, and `__m256` types are passed in YMM registers. This differs from the standard x64 calling convention, because the vector types are passed by value instead of by reference, and additional registers are used. The shadow stack space allocated for vector type arguments is fixed at 8 bytes, and the `/homeparams` option does not apply. Vector type arguments in the seventh and later parameter positions are passed on the stack by reference to memory allocated by the caller.

After registers are allocated for vector arguments, the data members of HVA arguments are allocated, in ascending order, to unused vector registers XMM0 to XMM5 (or YMM0 to YMM5, for `__m256` types), as long as there are enough registers available for the entire HVA. If not enough registers are available, the HVA argument is passed by reference to memory allocated by the caller. The stack shadow space for an HVA argument is fixed at 8 bytes with undefined content. HVA arguments are assigned to registers in order from left to right in the parameter list, and

may be in any position. HVA arguments in one of the first four argument positions that are not assigned to vector registers are passed by reference in the integer register that corresponds to that position. HVA arguments passed by reference after the fourth parameter position are pushed on the stack.

Results of `__vectorcall` functions are returned by value in registers when possible. Results of integer type, including structs or unions of 8 bytes or less, are returned by value in RAX. Vector type results are returned by value in XMM0 or YMM0, depending on size. HVA results have each data element returned by value in registers XMM0:XMM3 or YMM0:YMM3, depending on element size. Result types that don't fit in the corresponding registers are returned by reference to memory allocated by the caller.

The stack is maintained by the caller in the x64 implementation of `__vectorcall`. The caller prolog and epilog code allocates and clears the stack for the called function. Arguments are pushed on the stack from right to left, and shadow stack space is allocated for arguments passed in registers.

Examples:

```
// crt_vc64.c
// Build for amd64 with: cl /arch:AVX /W3 /FAs crt_vc64.c
// This example creates an annotated assembly listing in
// crt_vc64.asm.

#include <intrin.h>
#include <xmmmintrin.h>

typedef struct {
    __m128 array[2];
} hva2;      // 2 element HVA type on __m128

typedef struct {
    __m256 array[4];
} hva4;      // 4 element HVA type on __m256

// Example 1: All vectors
// Passes a in XMM0, b in XMM1, c in YMM2, d in XMM3, e in YMM4.
// Return value in XMM0.
__m128 __vectorcall
example1(__m128 a, __m128 b, __m256 c, __m128 d, __m256 e) {
    return d;
}

// Example 2: Mixed int, float and vector parameters
// Passes a in RCX, b in XMM1, c in R8, d in XMM3, e in YMM4,
// f in XMM5, g pushed on stack.
// Return value in YMM0.
__m256 __vectorcall
example2(int a, __m128 b, int c, __m128 d, __m256 e, float f, int g) {
    return e;
}

// Example 3: Mixed int and HVA parameters
// Passes a in RCX, c in R8, d in R9, and e pushed on stack.
// Passes b by element in [XMM0:XMM1];
// b's stack shadow area is 8-bytes of undefined value.
// Return value in XMM0.
__m128 __vectorcall example3(int a, hva2 b, int c, int d, int e) {
    return b.array[0];
}

// Example 4: Discontiguous HVA
// Passes a in RCX, b in XMM1, d in XMM3, and e is pushed on stack.
// Passes c by element in [YMM0,YMM2,YMM4,YMM5], discontiguous because
// vector arguments b and d were allocated first.
// Shadow area for c is an 8-byte undefined value.
// Return value in XMM0.
float __vectorcall example4(int a, float b, hva4 c, __m128 d, int e) {
```

```

    return b;
}

// Example 5: Multiple HVA arguments
// Passes a in RCX, c in R8, e pushed on stack.
// Passes b in [XMM0:XMM1], d in [YMM2:YMM5], each with
// stack shadow areas of an 8-byte undefined value.
// Return value in RAX.
int __vectorcall example5(int a, hva2 b, int c, hva4 d, int e) {
    return c + e;
}

// Example 6: HVA argument passed by reference, returned by register
// Passes a in [XMM0:XMM1], b passed by reference in RDX, c in YMM2,
// d in [XMM3:XMM4].
// Register space was insufficient for b, but not for d.
// Return value in [YMM0:YMM3].
hva4 __vectorcall example6(hva2 a, hva4 b, __m256 c, hva2 d) {
    return b;
}

int __cdecl main( void )
{
    hva4 h4;
    hva2 h2;
    int i;
    float f;
    __m128 a, b, d;
    __m256 c, e;

    a = b = d = _mm_set1_ps(3.0f);
    c = e = _mm256_set1_ps(5.0f);
    h2.array[0] = _mm_set1_ps(6.0f);
    h4.array[0] = _mm256_set1_ps(7.0f);

    b = example1(a, b, c, d, e);
    e = example2(1, b, 3, d, e, 6.0f, 7);
    d = example3(1, h2, 3, 4, 5);
    f = example4(1, 2.0f, h4, d, 5);
    i = example5(1, h2, 3, h4, 5);
    h4 = example6(h2, h4, c, h2);
}

```

## \_\_vectorcall convention on x86

The `__vectorcall` calling convention follows the `__fastcall` convention for 32-bit integer type arguments, and takes advantage of the SSE vector registers for vector type and HVA arguments.

The first two integer type arguments found in the parameter list from left to right are placed in ECX and EDX, respectively. A hidden `this` pointer is treated as the first integer type argument, and is passed in ECX. The first six vector type arguments are passed by value through SSE vector registers 0 to 5, in the XMM or YMM registers, depending on argument size.

The first six vector type arguments in order from left to right are passed by value in SSE vector registers 0 to 5. Floating-point and `__m128` types are passed in XMM registers, and `__m256` types are passed in YMM registers. No shadow stack space is allocated for vector type arguments passed by register. The seventh and subsequent vector type arguments are passed on the stack by reference to memory allocated by the caller. The limitation of compiler error [C2719](#) does not apply to these arguments.

After registers are allocated for vector arguments, the data members of HVA arguments are allocated in ascending order to unused vector registers XMM0 to XMM5 (or YMM0 to YMM5, for `__m256` types), as long as there are enough registers available for the entire HVA. If not enough registers are available, the HVA argument is passed on the stack by reference to memory allocated by the caller. No stack shadow space for an HVA argument is allocated.

HVA arguments are assigned to registers in order from left to right in the parameter list, and may be in any position.

Results of `__vectorcall` functions are returned by value in registers when possible. Results of integer type, including structs or unions of 4 bytes or less, are returned by value in EAX. Integer type structs or unions of 8 bytes or less are returned by value in EDX:EAX. Vector type results are returned by value in XMM0 or YMM0, depending on size. HVA results have each data element returned by value in registers XMM0:XMM3 or YMM0:YMM3, depending on element size. Other result types are returned by reference to memory allocated by the caller.

The x86 implementation of `__vectorcall` follows the convention of arguments pushed on the stack from right to left by the caller, and the called function clears the stack just before it returns. Only arguments that are not placed in registers are pushed on the stack.

Examples:

```
// crt_vc86.c
// Build for x86 with: cl /arch:AVX /W3 /FAs crt_vc86.c
// This example creates an annotated assembly listing in
// crt_vc86.asm.

#include <intrin.h>
#include <xmmmintrin.h>

typedef struct {
    __m128 array[2];
} hva2;      // 2 element HVA type on __m128

typedef struct {
    __m256 array[4];
} hva4;      // 4 element HVA type on __m256

// Example 1: All vectors
// Passes a in XMM0, b in XMM1, c in YMM2, d in XMM3, e in YMM4.
// Return value in XMM0.
__m128 __vectorcall
example1(__m128 a, __m128 b, __m256 c, __m128 d, __m256 e) {
    return d;
}

// Example 2: Mixed int, float and vector parameters
// Passes a in ECX, b in XMM0, c in EDX, d in XMM1, e in YMM2,
// f in XMM3, g pushed on stack.
// Return value in YMM0.
__m256 __vectorcall
example2(int a, __m128 b, int c, __m128 d, __m256 e, float f, int g) {
    return e;
}

// Example 3: Mixed int and HVA parameters
// Passes a in ECX, c in EDX, d and e pushed on stack.
// Passes b by element in [XMM0:XMM1].
// Return value in XMM0.
__m128 __vectorcall example3(int a, hva2 b, int c, int d, int e) {
    return b.array[0];
}

// Example 4: HVA assigned after vector types
// Passes a in ECX, b in XMM0, d in XMM1, and e in EDX.
// Passes c by element in [YMM2:YMM5].
// Return value in XMM0.
float __vectorcall example4(int a, float b, hva4 c, __m128 d, int e) {
    return b;
}

// Example 5: Multiple HVA arguments
```

```

// Passes a in ECX, c in EDX, e pushed on stack.
// Passes b in [XMM0:XMM1], d in [YMM2:YMM5].
// Return value in EAX.
int __vectorcall example5(int a, hva2 b, int c, hva4 d, int e) {
    return c + e;
}

// Example 6: HVA argument passed by reference, returned by register
// Passes a in [XMM1:XMM2], b passed by reference in ECX, c in YMM0,
// d in [XMM3:XMM4].
// Register space was insufficient for b, but not for d.
// Return value in [YMM0:YMM3].
hva4 __vectorcall example6(hva2 a, hva4 b, __m256 c, hva2 d) {
    return b;
}

int __cdecl main( void )
{
    hva4 h4;
    hva2 h2;
    int i;
    float f;
    __m128 a, b, d;
    __m256 c, e;

    a = b = d = _mm_set1_ps(3.0f);
    c = e = _mm256_set1_ps(5.0f);
    h2.array[0] = _mm_set1_ps(6.0f);
    h4.array[0] = _mm256_set1_ps(7.0f);

    b = example1(a, b, c, d, e);
    e = example2(1, b, 3, d, e, 6.0f, 7);
    d = example3(1, h2, 3, 4, 5);
    f = example4(1, 2.0f, h4, d, 5);
    i = example5(1, h2, 3, h4, 5);
    h4 = example6(h2, h4, c, h2);
}

```

## End Microsoft Specific

## See Also

[Argument Passing and Naming Conventions](#)  
[Keywords](#)

# Calling Example: Function Prototype and Call

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The following example shows the results of making a function call using various calling conventions.

This example is based on the following function skeleton. Replace `calltype` with the appropriate calling convention.

```
void    calltype MyFunc( char c, short s, int i, double f );
.
.
.
void    MyFunc( char c, short s, int i, double f )
{
.
.
.
}
.
.
.
MyFunc ('x', 12, 8192, 2.7183);
```

For more information, see [Results of Calling Example](#).

**END Microsoft Specific**

## See Also

[Calling Conventions](#)

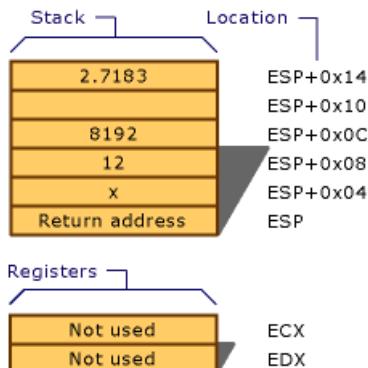
# Results of Calling Example

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

### cdecl

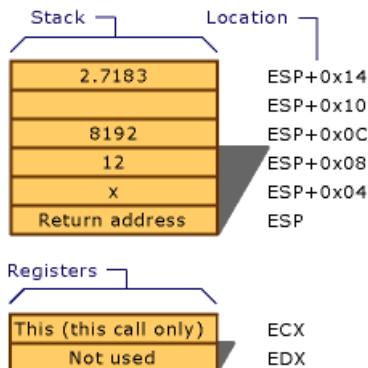
The C decorated function name is "\_MyFunc."



The \_cdecl calling convention

### stdcall and thiscall

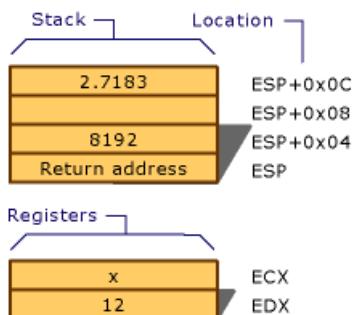
The C decorated name (`__stdcall`) is "\_MyFunc@20." The C++ decorated name is proprietary.



The \_stdcall and \_thiscall calling conventions

### fastcall

The C decorated name (`__fastcall`) is "@MyFunc@20." The C++ decorated name is proprietary.



The \_fastcall calling convention

**END Microsoft Specific**

## See Also

[Calling Example: Function Prototype and Call](#)

# Naked Function Calls

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Functions declared with the `naked` attribute are emitted without prolog or epilog code, enabling you to write your own custom prolog/epilog sequences using the [inline assembler](#). Naked functions are provided as an advanced feature. They enable you to declare a function that is being called from a context other than C/C++, and thus make different assumptions about where parameters are, or which registers are preserved. Examples include routines such as interrupt handlers. This feature is particularly useful for writers of virtual device drivers (VxDs).

## What do you want to know more about?

- [naked](#)
- [Rules and Limitations for Naked Functions](#)
- [Considerations for Writing Prolog/Epilog Code](#)

**END Microsoft Specific**

## See Also

[Calling Conventions](#)

# Rules and Limitations for Naked Functions

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The following rules and limitations apply to naked functions:

- The `return` statement is not permitted.
- Structured Exception Handling and C++ Exception Handling constructs are not permitted because they must unwind across the stack frame.
- For the same reason, any form of `setjmp` is prohibited.
- Use of the `_alloca` function is prohibited.
- To ensure that no initialization code for local variables appears before the prolog sequence, initialized local variables are not permitted at function scope. In particular, the declaration of C++ objects is not permitted at function scope. There may, however, be initialized data in a nested scope.
- Frame pointer optimization (the /Oy compiler option) is not recommended, but it is automatically suppressed for a naked function.
- You cannot declare C++ class objects at the function lexical scope. You can, however, declare objects in a nested block.
- The `naked` keyword is ignored when compiling with `/clr`.
- For `__fastcall` naked functions, whenever there is a reference in C/C++ code to one of the register arguments, the prolog code should store the values of that register into the stack location for that variable.  
For example:

```
// nkdfastcl.cpp
// compile with: /c
// processor: x86
__declspec(naked) int __fastcall power(int i, int j) {
    // calculates i^j, assumes that j >= 0

    // prolog
    __asm {
        push ebp
        mov ebp, esp
        sub esp, __LOCAL_SIZE
        // store ECX and EDX into stack locations allocated for i and j
        mov i, ecx
        mov j, edx
    }

    {
        int k = 1;    // return value
        while (j-- > 0)
            k *= i;
        __asm {
            mov eax, k
        };
    }

    // epilog
    __asm {
        mov esp, ebp
        pop ebp
        ret
    }
}
```

## END Microsoft Specific

## See Also

[Naked Function Calls](#)

# Considerations for Writing Prolog/Epilog Code

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Before writing your own prolog and epilog code sequences, it is important to understand how the stack frame is laid out. It is also useful to know how to use the **\_LOCAL\_SIZE** symbol.

### Stack Frame Layout

This example shows the standard prolog code that might appear in a 32-bit function:

```
push    ebp      ; Save ebp
mov     ebp, esp ; Set stack frame pointer
sub     esp, localbytes ; Allocate space for locals
push    <registers> ; Save registers
```

The `localbytes` variable represents the number of bytes needed on the stack for local variables, and the `<registers>` variable is a placeholder that represents the list of registers to be saved on the stack. After pushing the registers, you can place any other appropriate data on the stack. The following is the corresponding epilog code:

```
pop    <registers> ; Restore registers
mov     esp, ebp   ; Restore stack pointer
pop    ebp        ; Restore ebp
ret               ; Return from function
```

The stack always grows down (from high to low memory addresses). The base pointer (`ebp`) points to the pushed value of `ebp`. The locals area begins at `ebp-4`. To access local variables, calculate an offset from `ebp` by subtracting the appropriate value from `ebp`.

### **\_LOCAL\_SIZE**

The compiler provides a symbol, **\_LOCAL\_SIZE**, for use in the inline assembler block of function prolog code. This symbol is used to allocate space for local variables on the stack frame in custom prolog code.

The compiler determines the value of **\_LOCAL\_SIZE**. Its value is the total number of bytes of all user-defined local variables and compiler-generated temporary variables. **\_LOCAL\_SIZE** can be used only as an immediate operand; it cannot be used in an expression. You must not change or redefine the value of this symbol. For example:

```
mov    eax, _LOCAL_SIZE      ; Immediate operand--Okay
mov    eax, [ebp - _LOCAL_SIZE] ; Error
```

The following example of a naked function containing custom prolog and epilog sequences uses the **\_LOCAL\_SIZE** symbol in the prolog sequence:

```
// the__local_size_symbol.cpp
// processor: x86
__declspec ( naked ) int main() {
    int i;
    int j;

    __asm {      /* prolog */
        push    ebp
        mov     ebp, esp
        sub     esp, __LOCAL_SIZE
    }

    /* Function body */
    __asm {      /* epilog */
        mov     esp, ebp
        pop    ebp
        ret
    }
}
```

## END Microsoft Specific

## See Also

[Naked Function Calls](#)

# Floating Point Coprocessor and Calling Conventions

12/21/2017 • 1 min to read • [Edit Online](#)

If you are writing assembly routines for the floating point coprocessor, you must preserve the floating point control word and clean the coprocessor stack unless you are returning a **float** or **double** value (which your function should return in ST(0)).

## See Also

[Calling Conventions](#)

# Obsolete Calling Conventions

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `__pascal`, `__fortran`, and `__syscall` calling conventions are no longer supported. You can emulate their functionality by using one of the supported calling conventions and appropriate linker options.

WINDOWS.H now supports the **WINAPI** macro, which translates to the appropriate calling convention for the target. Use **WINAPI** where you previously used **PASCAL** or `__far __pascal`.

**END Microsoft Specific**

## See Also

[Argument Passing and Naming Conventions](#)

# restrict (C++ AMP)

12/21/2017 • 2 min to read • [Edit Online](#)

The restriction specifier can be applied to function and lambda declarations. It enforces restrictions on the code in the function and on the behavior of the function in applications that use the C++ Accelerated Massive Parallelism (C++ AMP) runtime.

## NOTE

For information about the `restrict` keyword that is part of the `__declspec` storage-class attributes, see [restrict](#).

The `restrict` clause takes the following forms:

CLAUSE	DESCRIPTION
<code>restrict(cpu)</code>	The function can use the full C++ language. Only other functions that are declared by using <code>restrict(cpu)</code> functions can call the function.
<code>restrict(amp)</code>	The function can only use the subset of the C++ language that C++ AMP can accelerate.
A sequence of <code>restrict(cpu)</code> and <code>restrict(amp)</code> .	<p>The function must adhere to the limitations of both <code>restrict(cpu)</code> and <code>restrict(amp)</code>. The function can be called by functions that are declared by using <code>restrict(cpu)</code>, <code>restrict(amp)</code>, <code>restrict(cpu, amp)</code>, or <code>restrict(amp, cpu)</code>.</p> <p>The form <code>restrict(A) restrict(B)</code> can be written as <code>restrict(A,B)</code>.</p>

## Remarks

The `restrict` keyword is a contextual keyword. The restriction specifiers, `cpu` and `amp` are not reserved words. The list of specifiers is not extensible. A function that does not have a `restrict` clause is the same as a function that has the `restrict(cpu)` clause.

A function that has the `restrict(amp)` clause has the following limitations:

- The function can call only functions that have the `restrict(amp)` clause.
- The function must be inlinable.
- The function can declare only `int`, `unsigned int`, `float`, and `double` variables, and classes and structures that contain only these types. `bool` is also allowed, but it must be 4-byte-aligned if you use it in a compound type.
- Lambda functions cannot capture by reference and cannot capture pointers.
- References and single-indirection pointers are supported only as local variables, function arguments, and return types.

- The following are not allowed:
  - Recursion.
  - Variables declared with the [volatile](#) keyword.
  - Virtual functions.
  - Pointers to functions.
  - Pointers to member functions.
  - Pointers in structures.
  - Pointers to pointers.
  - `goto` statements.
  - Labeled statements.
  - `try`, `catch`, or `throw` statements.
  - Global variables.
  - Static variables. Use [tile\\_static Keyword](#) instead.
  - `dynamic_cast` casts.
  - The `typeid` operator.
  - `asm` declarations.
  - Varargs.

For a discussion of function limitations, see [restrict\(amp\) Restrictions](#).

## Example

The following example shows how to use the `restrict(amp)` clause.

```
void functionAmp() restrict(amp) {}
void functionNonAmp() {}

void callFunctions() restrict(amp)
{
    // int is allowed.
    int x;
    // long long int is not allowed in an amp-restricted function. This generates a compiler error.
    // long long int y;

    // Calling an amp-restricted function is allowed.
    functionAmp();

    // Calling a non-amp-restricted function is not allowed.
    // functionNonAmp();
}
```

## See Also

[C++ AMP \(C++ Accelerated Massive Parallelism\)](#)

# tile\_static Keyword

12/21/2017 • 3 min to read • [Edit Online](#)

The `tile_static` keyword is used to declare a variable that can be accessed by all threads in a tile of threads. The lifetime of the variable starts when execution reaches the point of declaration and ends when the kernel function returns. For more information on using tiles, see [Using Tiles](#).

The `tile_static` keyword has the following limitations:

- It can be used only on variables that are in a function that has the `restrict(amp)` modifier.
- It cannot be used on variables that are pointer or reference types.
- A `tile_static` variable cannot have an initializer. Default constructors and destructors are not invoked automatically.
- The value of an uninitialized `tile_static` variable is undefined.
- If a `tile_static` variable is declared in a call graph that is rooted by a non-tiled call to `parallel_for_each`, a warning is generated and the behavior of the variable is undefined.

## Example

The following example shows how a `tile_static` variable can be used to accumulate data across several threads in a tile.

```
// Sample data:  
int sampledata[] = {  
    2, 2, 9, 7, 1, 4,  
    4, 4, 8, 8, 3, 4,  
    1, 5, 1, 2, 5, 2,  
    6, 8, 3, 2, 7, 2};  
  
// The tiles:  
// 2 2      9 7      1 4  
// 4 4      8 8      3 4  
//  
// 1 5      1 2      5 2  
// 6 8      3 2      7 2  
  
// Averages:  
int averagedata[] = {  
    0, 0, 0, 0, 0, 0,  
    0, 0, 0, 0, 0, 0,  
    0, 0, 0, 0, 0, 0,  
    0, 0, 0, 0, 0, 0,  
};  
  
array_view<int, 2> sample(4, 6, sampledata);  
array_view<int, 2> average(4, 6, averagedata);  
  
parallel_for_each(  
    // Create threads for sample.extent and divide the extent into 2 x 2 tiles.  
    sample.extent.tile<2,2>(),  
    [=](tiled_index<2,2> idx) restrict(amp)  
    {  
        // Create a 2 x 2 array to hold the values in this tile.  
        tile_static int nums[2][2];  
        // Copy the values for the tile into the 2 x 2 array.  
    }
```



```
// Output.  
// 3 3 8 8 3 3  
// 3 3 8 8 3 3  
// 5 5 2 2 4 4  
// 5 5 2 2 4 4
```

## See Also

[Microsoft-Specific Modifiers](#)

[C++ AMP Overview](#)

[parallel\\_for\\_each Function \(C++ AMP\)](#)

[Walkthrough: Matrix Multiplication](#)

# `__declspec`

12/21/2017 • 2 min to read • [Edit Online](#)

## Microsoft Specific

The extended attribute syntax for specifying storage-class information uses the `__declspec` keyword, which specifies that an instance of a given type is to be stored with a Microsoft-specific storage-class attribute listed below. Examples of other storage-class modifiers include the `static` and `extern` keywords. However, these keywords are part of the ANSI specification of the C and C++ languages, and as such are not covered by extended attribute syntax. The extended attribute syntax simplifies and standardizes Microsoft-specific extensions to the C and C++ languages.

## Grammar

*decl-specifier:*

`__declspec ( [extended-decl-modifier-seq] )`

*extended-decl-modifier-seq:*

`extended-decl-modifieropt`

*extended-decl-modifier extended-decl-modifier-seq*

*extended-decl-modifier:*

`align( # )`

`allocate(" segname ")`

`appdomain`

`code_seg(" segname ")`

`deprecated`

`dllimport`

`dllexport`

`jitintrinsic`

`naked`

`noalias`

`noinline`

`noreturn`

`nothrow`

`novtable`

`process`

`property( { get= get_func_name| ,put= put_func_name} )`

`restrict`

```
safebuffers  
selectany  
thread  
uuid(" ComObjectGUID ")
```

White space separates the declaration modifier sequence. Examples appear in later sections.

Extended attribute grammar supports these Microsoft-specific storage-class attributes: [align](#), [allocate](#), [appdomain](#), [code\\_seg](#), [deprecated](#), [dllexport](#), [dllimport](#), [jitintrinsic](#), [naked](#), [noalias](#), [noinline](#), [noreturn](#), [nothrow](#), [novtable](#), [process](#), [restrict](#), [safebuffers](#), [selectany](#), and [thread](#). It also supports these COM-object attributes: [property](#) and [uuid](#).

The `code_seg`, `dllexport`, `dllimport`, `naked`, `noalias`, `nothrow`, `property`, `restrict`, `selectany`, `thread`, and `uuid` storage-class attributes are properties only of the declaration of the object or function to which they are applied. The `thread` attribute affects data and objects only. The `naked` attribute affects functions only. The `dllimport` and `dllexport` attributes affect functions, data, and objects. The `property`, `selectany`, and `uuid` attributes affect COM objects.

The `__declspec` keywords should be placed at the beginning of a simple declaration. The compiler ignores, without warning, any `__declspec` keywords placed after \* or & and in front of the variable identifier in a declaration.

A `__declspec` attribute specified in the beginning of a user-defined type declaration applies to the variable of that type. For example:

```
__declspec(dllexport) class X {} varX;
```

In this case, the attribute applies to `varX`. A `__declspec` attribute placed after the `class` or `struct` keyword applies to the user-defined type. For example:

```
class __declspec(dllexport) X {};
```

In this case, the attribute applies to `X`.

The general guideline for using the `__declspec` attribute for simple declarations is as follows:

```
decl-specifier-seq  
declarator-list;
```

The *decl-specifier-seq* should contain, among other things, a base type (e.g. `int`, `float`, a `typedef`, or a class name), a storage class (e.g. `static`, `extern`), or the `__declspec` extension. The *init-declarator-list* should contain, among other things, the pointer part of declarations. For example:

```
__declspec(selectany) int * pi1 = 0; //OK, selectany & int both part of decl-specifier  
int __declspec(selectany) * pi2 = 0; //OK, selectany & int both part of decl-specifier  
int * __declspec(selectany) pi3 = 0; //ERROR, selectany is not part of a declarator
```

The following code declares an integer thread local variable and initializes it with a value:

```
// Example of the __declspec keyword  
__declspec( thread ) int tls_i = 1;
```

## END Microsoft Specific

## See Also

[Keywords](#)

[C Extended Storage-Class Attributes](#)

# align (C++)

12/21/2017 • 8 min to read • [Edit Online](#)

In Visual Studio 2015 and later, use the C++11 standard  `specifier to control alignment. For more information, see Alignment.`

## Microsoft Specific

Use `__declspec(align(#))` to precisely control the alignment of user-defined data (for example, static allocations or automatic data in a function).

## Syntax

```
__declspec( align( # ) ) declarator
```

## Remarks

Writing applications that use the latest processor instructions introduces some new constraints and issues. In particular, many new instructions require that data must be aligned to 16-byte boundaries. Additionally, by aligning frequently used data to the cache line size of a specific processor, you improve cache performance. For example, if you define a structure whose size is less than 32 bytes, you may want to align it to 32 bytes to make sure that objects of that structure type are efficiently cached.

# is the alignment value. Valid entries are integer powers of two from 1 to 8192 (bytes), such as 2, 4, 8, 16, 32, or 64. `declarator` is the data that you are declaring as aligned.

For information about how to return a value of type `size_t` that is the alignment requirement of the type, see [\\_alignof](#). For information about how to declare unaligned pointers when targeting 64-bit processors, see [\\_unaligned](#).

You can use `__declspec(align(#))` when you define a `struct`, `union`, or `class`, or when you declare a variable.

The compiler does not guarantee or attempt to preserve the alignment attribute of data during a copy or data transform operation. For example, `memcpy` can copy a struct declared with `__declspec(align(#))` to any location. Note that ordinary allocators—for example, `malloc`, C++ `operator new`, and the Win32 allocators—return memory that is usually not sufficiently aligned for `__declspec(align(#))` structures or arrays of structures. To guarantee that the destination of a copy or data transformation operation is correctly aligned, use `_aligned_malloc`, or write your own allocator.

You cannot specify alignment for function parameters. When data that has an alignment attribute is passed by value on the stack, its alignment is controlled by the calling convention. If data alignment is important in the called function, copy the parameter into correctly aligned memory before use.

Without `__declspec(align(#))`, Visual C++ generally aligns data on natural boundaries based on the target processor and the size of the data, up to 4-byte boundaries on 32-bit processors, and 8-byte boundaries on 64-bit processors. Data in classes or structures is aligned in the class or structure at the minimum of its natural alignment and the current packing setting (from `#pragma pack` or the `/Zp` compiler option).

This example demonstrates the use of `__declspec(align(#))`:

```
__declspec(align(32)) struct Str1{
    int a, b, c, d, e;
};
```

This type now has a 32-byte alignment attribute. This means that all static and automatic instances start on a 32-byte boundary. Additional structure types declared with this type as a member preserve this type's alignment attribute, that is, any structure with `Str1` as an element will have an alignment attribute of at least 32.

Note that `sizeof(struct Str1)` is equal to 32. This implies that if an array of `Str1` objects is created, and the base of the array is 32-byte aligned, each member of the array is also 32-byte aligned. To create an array whose base is correctly aligned in dynamic memory, use `_aligned_malloc`, or write your own allocator.

The `sizeof` value for any structure is the offset of the final member, plus that member's size, rounded up to the nearest multiple of the largest member alignment value or the whole structure alignment value, whichever is larger.

The compiler uses these rules for structure alignment:

- Unless overridden with `__declspec(align(#))`, the alignment of a scalar structure member is the minimum of its size and the current packing.
- Unless overridden with `__declspec(align(#))`, the alignment of a structure is the maximum of the individual alignments of its member(s).
- A structure member is placed at an offset from the start of its parent structure which is the smallest multiple of its alignment greater than or equal to the offset of the end of the previous member.
- The size of a structure is the smallest multiple of its alignment greater than or equal to the offset of the end of its last member.

`__declspec(align(#))` can only increase alignment restrictions.

For more information, see:

- [align Examples](#)
- [Defining New Types with \\_\\_declspec\(align\(#\)\)](#)
- [Aligning Data in Thread Local Storage](#)
- [How align Works with Data Packing](#)
- [Examples of Structure Alignment \(x64 specific\)](#)

## align Examples

The following examples show how `__declspec(align(#))` affects the size and alignment of data structures. The examples assume the following definitions:

```
#define CACHE_LINE 32
#define CACHE_ALIGN __declspec(align(CACHE_LINE))
```

In this example, the `s1` structure is defined by using `__declspec(align(32))`. All uses of `s1` for a variable definition or in other type declarations are 32-byte aligned. `sizeof(struct s1)` returns 32, and `s1` has 16 padding bytes following the 16 bytes required to hold the four integers. Each `int` member requires 4-byte alignment, but the alignment of the structure itself is declared to be 32. Therefore, the overall alignment is 32.

```

struct CACHE_ALIGN S1 { // cache align all instances of S1
    int a, b, c, d;
};

struct S1 s1; // s1 is 32-byte cache aligned

```

In this example, `sizeof(struct S2)` returns 16, which is exactly the sum of the member sizes, because that is a multiple of the largest alignment requirement (a multiple of 8).

```

__declspec(align(8)) struct S2 {
    int a, b, c, d;
};

```

In the following example, `sizeof(struct S3)` returns 64.

```

struct S3 {
    struct S1 s1; // S3 inherits cache alignment requirement
                   // from S1 declaration
    int a;        // a is now cache aligned because of s1
                   // 28 bytes of trailing padding
};

```

In this example, notice that `a` has the alignment of its natural type, in this case, 4 bytes. However, `s1` must be 32-byte aligned. Twenty-eight bytes of padding follow `a`, so that `s1` starts at offset 32. `S4` then inherits the alignment requirement of `s1`, because it is the largest alignment requirement in the structure. `sizeof(struct S4)` returns 64.

```

struct S4 {
    int a;
    // 28 bytes padding
    struct S1 s1; // S4 inherits cache alignment requirement of S1
};

```

The following three variable declarations also use `__declspec(align(#))`. In each case, the variable must be 32-byte aligned. In the case of the array, the base address of the array, not each array member, is 32-byte aligned. The `sizeof` value for each array member is not affected when you use `__declspec(align(#))`.

```

CACHE_ALIGN int i;
CACHE_ALIGN int array[128];
CACHE_ALIGN struct s2 s;

```

To align each member of an array, code such as this should be used:

```

typedef CACHE_ALIGN struct { int a; } S5;
S5 array[10];

```

In this example, notice that aligning the structure itself and aligning the first element have the same effect:

```
CACHE_ALIGN struct S6 {
    int a;
    int b;
};

struct S7 {
    CACHE_ALIGN int a;
    int b;
};
```

`S6` and `S7` have identical alignment, allocation, and size characteristics.

In this example, the alignment of the starting addresses of `a`, `b`, `c`, and `d` are 4, 1, 4, and 1, respectively.

```
void fn() {
    int a;
    char b;
    long c;
    char d[10]
}
```

The alignment when memory is allocated on the heap depends on which allocation function is called. For example, if you use `malloc`, the result depends on the operand size. If  $arg \geq 8$ , the memory returned is 8 byte aligned. If  $arg < 8$ , the alignment of the memory returned is the first power of 2 less than  $arg$ . For example, if you use `malloc(7)`, the alignment is 4 bytes.

## Defining New Types with `_declspec(align(#))`

You can define a type with an alignment characteristic.

For example, you can define a `struct` with an alignment value this way:

```
struct aType {int a; int b;};
typedef __declspec(align(32)) struct aType bType;
```

Now, `aType` and `bType` are the same size (8 bytes) but variables of type `bType` will be 32-byte aligned.

## Aligning Data in Thread Local Storage

Static thread-local storage (TLS) created with the `__declspec(thread)` attribute and put in the TLS section in the image works for alignment exactly like normal static data. To create TLS data, the operating system allocates memory the size of the TLS section and respects the TLS section alignment attribute.

This example shows various ways to place aligned data into thread local storage.

```

// put an aligned integer in TLS
__declspec(thread) __declspec(align(32)) int a;

// define an aligned structure and put a variable of the struct type
// into TLS
__declspec(thread) __declspec(align(32)) struct F1 { int a; int b; } a;

// create an aligned structure
struct CACHE_ALIGN S9 {
    int a;
    int b;
};

// put a variable of the structure type into TLS
__declspec(thread) struct S9 a;

```

## How align Works with Data Packing

The **/Zp** compiler option and the `pack` pragma have the effect of packing data for structure and union members. This example shows how **/Zp** and `__declspec(align(#))` work together:

```

struct S {
    char a;
    short b;
    double c;
    CACHE_ALIGN double d;
    char e;
    double f;
};

```

The following table lists the offset of each member under a variety of **/Zp** (or #pragma `pack`) values, showing how the two interact.

VARIABLE	/ZP1	/ZP2	/ZP4	/ZP8
a	0	0	0	0
b	1	2	2	2
c	3	4	4	8
d	32	32	32	32
e	40	40	40	40
f	41	42	44	48
sizeof(S)	64	64	64	64

For more information, see [/Zp \(Struct Member Alignment\)](#).

The offset of an object is based on the offset of the previous object and the current packing setting, unless the object has a `__declspec(align(#))` attribute, in which case the alignment is based on the offset of the previous object and the `__declspec(align(#))` value for the object.

**END Microsoft Specific**

## See Also

[\\_declspec](#)

[Overview of ARM ABI Conventions](#)

[Overview of x64 Calling Conventions](#)

# allocate

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The **allocate** declaration specifier names a data segment in which the data item will be allocated.

## Syntax

```
__declspec(allocate(  
segname  
")) declarator
```

## Remarks

The name *segname* must be declared using one of the following pragmas:

- [code\\_seg](#)
- [const\\_seg](#)
- [data\\_seg](#)
- [init\\_seg](#)
- [section](#)

## Example

```
// allocate.cpp  
#pragma section("mycode", read)  
__declspec(allocate("mycode")) int i = 0;  
  
int main() {  
}
```

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# appdomain

12/21/2017 • 2 min to read • [Edit Online](#)

Specifies that each application domain of your managed application should have its own copy of a particular global variable or static member variable. See [Application Domains and Visual C++](#) for more information.

Every application domain has its own copy of a per-appdomain variable. A constructor of an appdomain variable is executed when an assembly is loaded into an application domain, and the destructor is executed when the application domain is unloaded.

If you want all application domains within a process in the common language runtime to share a global variable, use the `__declspec(process)` modifier. `__declspec(process)` is in effect by default under `/clr` and `__declspec(appdomain)` is in effect by default under `/clr:pure`. `__declspec(appdomain)` is enforced under `/clr:safe`. The `/clr:pure` and `/clr:safe` compiler options are deprecated in Visual Studio 2015.

`__declspec(appdomain)` is only valid when one of the `/clr` compiler options is used. Only a global variable, static member variable, or a static local variable can be marked with `__declspec(appdomain)`. It is an error to apply `__declspec(appdomain)` to static members of managed types because they always have this behavior.

Using `__declspec(appdomain)` is similar to using [Thread Local Storage \(TLS\)](#). Threads have their own storage, as do application domains. Using `__declspec(appdomain)` ensures the global variable has its own storage in each application domain created for this application.

There are limitations to mixing the use of per process and per appdomain variables; see [process](#) for more information.

For example, at program start up, all per-process variables are initialized, then all per-appdomain variables are initialized. Therefore when a per-process variable is being initialized, it cannot depend on the value of any per-application domain variable. It is bad practice to mix the use (assignment) of per appdomain and per process variables.

For information on how to call a function in a specific application domain, see [call\\_in\\_appdomain Function](#).

## Example

```
// declspec_appdomain.cpp
// compile with: /clr
#include <stdio.h>
using namespace System;

class CGlobal {
public:
    CGlobal(bool bProcess) {
        Counter = 10;
        m_bProcess = bProcess;
        Console::WriteLine("__declspec({0}) CGlobal::CGlobal constructor", m_bProcess ? (String^)"process" :
(String^)"appdomain");
    }

    ~CGlobal() {
        Console::WriteLine("__declspec({0}) CGlobal::~CGlobal destructor", m_bProcess ? (String^)"process" :
(String^)"appdomain");
    }

    int Counter;
}
```

```

private:
    bool m_bProcess;
};

__declspec(process) CGlobal process_global = CGlobal(true);
__declspec(appdomain) CGlobal appdomain_global = CGlobal(false);

value class Functions {
public:
    static void change() {
        ++appdomain_global.Counter;
    }

    static void display() {
        Console::WriteLine("process_global value in appdomain '{0}': {1}",
                           AppDomain::CurrentDomain->FriendlyName,
                           process_global.Counter);

        Console::WriteLine("appdomain_global value in appdomain '{0}': {1}",
                           AppDomain::CurrentDomain->FriendlyName,
                           appdomain_global.Counter);
    }
};

int main() {
    AppDomain^ defaultDomain = AppDomain::GetCurrentDomain();
    AppDomain^ domain = AppDomain::CreateDomain("Domain 1");
    AppDomain^ domain2 = AppDomain::CreateDomain("Domain 2");
    CrossAppDomainDelegate^ changeDelegate = gcnew CrossAppDomainDelegate(&Functions::change);
    CrossAppDomainDelegate^ displayDelegate = gcnew CrossAppDomainDelegate(&Functions::display);

    // Print the initial values of appdomain_global in all appdomains.
    Console::WriteLine("Initial value");
    defaultDomain->DoCallBack(displayDelegate);
    domain->DoCallBack(displayDelegate);
    domain2->DoCallBack(displayDelegate);

    // Changing the value of appdomain_global in the domain and domain2
    // appdomain_global value in "default" appdomain remain unchanged
    process_global.Counter = 20;
    domain->DoCallBack(changeDelegate);
    domain2->DoCallBack(changeDelegate);
    domain2->DoCallBack(changeDelegate);

    // Print values again
    Console::WriteLine("Changed value");
    defaultDomain->DoCallBack(displayDelegate);
    domain->DoCallBack(displayDelegate);
    domain2->DoCallBack(displayDelegate);

    AppDomain::Unload(domain);
    AppDomain::Unload(domain2);
}

```

```
_declspec(process) CGlobal::CGlobal constructor
_declspec(appdomain) CGlobal::CGlobal constructor
Initial value
process_global value in appdomain 'declspec_appdomain.exe': 10
appdomain_global value in appdomain 'declspec_appdomain.exe': 10
_declspec(appdomain) CGlobal::CGlobal constructor
process_global value in appdomain 'Domain 1': 10
appdomain_global value in appdomain 'Domain 1': 10
_declspec(appdomain) CGlobal::CGlobal constructor
process_global value in appdomain 'Domain 2': 10
appdomain_global value in appdomain 'Domain 2': 10
Changed value
process_global value in appdomain 'declspec_appdomain.exe': 20
appdomain_global value in appdomain 'declspec_appdomain.exe': 10
process_global value in appdomain 'Domain 1': 20
appdomain_global value in appdomain 'Domain 1': 11
process_global value in appdomain 'Domain 2': 20
appdomain_global value in appdomain 'Domain 2': 12
_declspec(appdomain) CGlobal::~CGlobal destructor
_declspec(appdomain) CGlobal::~CGlobal destructor
_declspec(appdomain) CGlobal::~CGlobal destructor
_declspec(process) CGlobal::~CGlobal destructor
```

## See Also

[\\_declspec](#)

[Keywords](#)

# code\_seg (\_declspec)

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

The `code_seg` declaration attribute names an executable text segment in the .obj file in which the object code for the function or class member functions will be stored.

## Syntax

```
_declspec(code_seg("segname")) declarator
```

## Remarks

The `_declspec(code_seg(...))` attribute enables the placement of code into separate named segments that can be paged or locked in memory individually. You can use this attribute to control the placement of instantiated templates and compiler-generated code.

A *segment* is a named block of data in an .obj file that is loaded into memory as a unit. A *text segment* is a segment that contains executable code. The term *section* is often used interchangeably with segment.

Object code that's generated when `declarator` is defined is put in the text segment specified by `segname`, which is a narrow-string literal. The name `segname` does not have to be specified in a `section` pragma before it can be used in a declaration. By default, when no `code_seg` is specified, object code is put in a segment named .text. A `code_seg` attribute overrides any existing `#pragma code_seg` directive. A `code_seg` attribute applied to a member function overrides any `code_seg` attribute applied to the enclosing class.

If an entity has a `code_seg` attribute, all declarations and definitions of the same entity must have identical `code_seg` attributes. If a base-class has a `code_seg` attribute, derived classes must have the same attribute.

When a `code_seg` attribute is applied to a namespace-scope function or a member function, the object code for that function is put in the specified text segment. When this attribute is applied to a class, all member functions of the class and nested classes—this includes compiler-generated special member functions—are put in the specified segment. Locally defined classes—for example, classes defined in a member function body—do not inherit the `code_seg` attribute of the enclosing scope.

When a `code_seg` attribute is applied to a template class or template function, all implicit specializations of the template are put in the specified segment. Explicit or partial specializations do not inherit the `code_seg` attribute from the primary template. You may specify the same or a different `code_seg` attribute on the specialization. A `code_seg` attribute can't be applied to an explicit template instantiation.

By default, compiler-generated code such as a special member function is put in the .text segment. The `#pragma code_seg` directive does not override this default. Use the `code_seg` attribute on the class, class template, or function template to control where compiler-generated code is put.

Lambdas inherit `code_seg` attributes from their enclosing scope. To specify a segment for a lambda, apply a `code_seg` attribute after the parameter-declaration clause and before any mutable or exception specification, any trailing return-type specification, and the lambda body. For more information, see [Lambda Expression Syntax](#). This example defines a lambda in a segment named PagedMem:

```
auto Sqr = [](int t) __declspec(code_seg("PagedMem")) -> int { return t*t; };
```

Be careful when you put specific member functions—especially virtual member functions—in different segments. If you define a virtual function in a derived class that resides in a paged segment when the base class method resides in a non-paged segment, other base class methods or user code may assume that invoking the virtual method will not trigger a page fault.

## Example

This example shows how a `code_seg` attribute controls segment placement when implicit and explicit template specialization is used:

```
// code_seg.cpp
// Compile: cl /EHsc /W4 code_seg.cpp

// Base template places object code in Segment_1 segment
template<class T>
class __declspec(code_seg("Segment_1")) Example
{
public:
    virtual void VirtualMemberFunction(T /*arg*/) {}

// bool specialization places code in default .text segment
template<>
class Example<bool>
{
public:
    virtual void VirtualMemberFunction(bool /*arg*/) {}

// int specialization places code in Segment_2 segment
template<>
class __declspec(code_seg("Segment_2")) Example<int>
{
public:
    virtual void VirtualMemberFunction(int /*arg*/) {}

// Compiler warns and ignores __declspec(code_seg("Segment_3"))
// in this explicit specialization
__declspec(code_seg("Segment_3")) Example<short>; // C4071

int main()
{
    // implicit double specialization uses base template's
    // __declspec(code_seg("Segment_1")) to place object code
    Example<double> doubleExample{};
    doubleExample.VirtualMemberFunction(3.14L);

    // bool specialization places object code in default .text segment
    Example<bool> boolExample{};
    boolExample.VirtualMemberFunction(true);

    // int specialization uses __declspec(code_seg("Segment_2"))
    // to place object code
    Example<int> intExample{};
    intExample.VirtualMemberFunction(42);
}
```

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# deprecated (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

This topic is about the Microsoft-specific deprecated declspec declaration. For information about the C++14 `[[deprecated]]` attribute, and guidance on when to use that attribute vs. the Microsoft-specific declspec or pragma, see [C++ Standard Attributes](#).

With the exceptions noted below, the **deprecated** declaration offers the same functionality as the [deprecated](#) pragma:

- The **deprecated** declaration lets you specify particular forms of function overloads as deprecated, whereas the pragma form applies to all overloaded forms of a function name.
- The **deprecated** declaration lets you specify a message that will display at compile time. The text of the message can be from a macro.
- Macros can only be marked as deprecated with the **deprecated** pragma.

If the compiler encounters the use of a deprecated identifier or the standard `[[deprecated]]` attribute, a [C4996](#) warning is thrown.

## Example

The following sample shows how to mark functions as deprecated, and how to specify a message that will be displayed at compile time, when the deprecated function is used.

```
// deprecated.cpp
// compile with: /W3
#define MY_TEXT "function is deprecated"
void func1(void) {}
__declspec(deprecated) void func1(int) {}
__declspec(deprecated("** this is a deprecated function **")) void func2(int) {}
__declspec(deprecated(MY_TEXT)) void func3(int) {}

int main() {
    func1();
    func1(1);    // C4996
    func2(1);    // C4996
    func3(1);    // C4996
}
```

## Example

The following sample shows how to mark classes as deprecated, and how to specify a message that will be displayed at compile time, when the deprecated class is used.

```
// deprecate_class.cpp
// compile with: /W3
struct __declspec(deprecated) X {
    void f(){}
};

struct __declspec(deprecated("** X2 is deprecated **")) X2 {
    void f(){}
};

int main() {
    X x;      // C4996
    X2 x2;    // C4996
}
```

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# \_\_declspec( dllexport ) and \_\_declspec( dllimport )

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The `__declspec( dllexport )` and `__declspec( dllimport )` storage-class attributes are Microsoft-specific extensions to the C and C++ languages. You can use them to export and import functions, data, and objects to or from a DLL.

## Syntax

```
__declspec( dllimport ) declarator  
__declspec( dllexport ) declarator
```

## Remarks

These attributes explicitly define the DLL's interface to its client, which can be the executable file or another DLL. Declaring functions as `__declspec( dllexport )` eliminates the need for a module-definition (.def) file, at least with respect to the specification of exported functions. The `__declspec( dllexport )` attribute replaces the `__declspec( export )` keyword.

If a class is marked `__declspec( dllexport )`, any specializations of class templates in the class hierarchy are implicitly marked as `__declspec( dllexport )`. This means that class templates are explicitly instantiated and the class's members must be defined.

`__declspec( dllexport )` of a function exposes the function with its decorated name. For C++ functions, this includes name mangling. For C functions or functions that are declared as `extern "C"`, this includes platform-specific decoration that's based on the calling convention. For information on name decoration in C/C++ code, see [Decorated Names](#). No name decoration is applied to exported C functions or C++ `extern "C"` functions using the `__cdecl` calling convention.

To export an undecorated name, you can link by using a Module Definition (.def) file that defines the undecorated name in an EXPORTS section. For more information, see [EXPORTS](#). Another way to export an undecorated name is to use a `#pragma comment(linker, "/export:alias=decorated_name")` directive in the source code.

When you declare `__declspec( dllexport )` or `__declspec( dllimport )`, you must use [extended attribute syntax](#) and the `__declspec` keyword.

## Example

```
// Example of the __declspec( dllimport ) and __declspec( dllexport ) class attributes  
__declspec( dllimport ) int i;  
__declspec( dllexport ) void func();
```

Alternatively, to make your code more readable, you can use macro definitions:

```
#define DllImport  __declspec( dllimport )
#define DllExport   __declspec( dllexport )

DllExport void func();
DllExport int i = 10;
DllImport int j;
DllExport int n;
```

For more information, see:

- [Definitions and Declarations](#)
- [Defining Inline C++ Functions with dllexport and dllimport](#)
- [General Rules and Limitations](#)
- [Using dllimport and dllexport in C++ Classes](#)

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

## Definitions and Declarations (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

# Microsoft Specific

The DLL interface refers to all items (functions and data) that are known to be exported by some program in the system; that is, all items that are declared as `dllimport` or `dllexport`. All declarations included in the DLL interface must specify either the `dllimport` or `dllexport` attribute. However, the definition must specify only the `dllexport` attribute. For example, the following function definition generates a compiler error:

```
__declspec( dllexport ) int func() {    // Error; dllexport
                                         // prohibited on definition.
    return 1;
}
```

This code also generates an error:

```
_declspec( dllimport ) int i = 10; // Error; this is a definition.
```

However, this is correct syntax:

```
_declspec( dllexport ) int i = 10; // Okay--export definition
```

The use of `dllexport` implies a definition, while `dllimport` implies a declaration. You must use the `extern` keyword with `dllexport` to force a declaration; otherwise, a definition is implied. Thus, the following examples are correct:

```
#define DllImport  __declspec( dllexport )
#define DllExport   __declspec( dllimport )

extern DllExport int k; // These are both correct and imply a
DllImport int j;        // declaration.
```

The following examples clarify the preceding:

**END Microsoft Specific**

## See Also

[dllexport, dllimport](#)

# Defining Inline C++ Functions with `dllexport` and `dllimport`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

You can define as inline a function with the `dllexport` attribute. In this case, the function is always instantiated and exported, whether or not any module in the program references the function. The function is presumed to be imported by another program.

You can also define as inline a function declared with the `dllimport` attribute. In this case, the function can be expanded (subject to /Ob specifications), but never instantiated. In particular, if the address of an inline imported function is taken, the address of the function residing in the DLL is returned. This behavior is the same as taking the address of a non-inline imported function.

These rules apply to inline functions whose definitions appear within a class definition. In addition, static local data and strings in inline functions maintain the same identities between the DLL and client as they would in a single program (that is, an executable file without a DLL interface).

Exercise care when providing imported inline functions. For example, if you update the DLL, don't assume that the client will use the changed version of the DLL. To ensure that you are loading the proper version of the DLL, rebuild the DLL's client as well.

### END Microsoft Specific

## See Also

[dllexport, dllimport](#)

# General Rules and Limitations

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

- If you declare a function or object without the **dllimport** or **dllexport** attribute, the function or object is not considered part of the DLL interface. Therefore, the definition of the function or object must be present in that module or in another module of the same program. To make the function or object part of the DLL interface, you must declare the definition of the function or object in the other module as **dllexport**. Otherwise, a linker error is generated.

If you declare a function or object with the **dllexport** attribute, its definition must appear in some module of the same program. Otherwise, a linker error is generated.

- If a single module in your program contains both **dllimport** and **dllexport** declarations for the same function or object, the **dllexport** attribute takes precedence over the **dllimport** attribute. However, a compiler warning is generated. For example:

```
__declspec( dllimport ) int i;
__declspec( dllexport ) int i; // Warning; inconsistent;
                           // dllexport takes precedence.
```

- In C++, you can initialize a globally declared or static local data pointer or with the address of a data object declared with the **dllimport** attribute, which generates an error in C. In addition, you can initialize a static local function pointer with the address of a function declared with the **dllimport** attribute. In C, such an assignment sets the pointer to the address of the DLL import thunk (a code stub that transfers control to the function) rather than the address of the function. In C++, it sets the pointer to the address of the function. For example:

```
__declspec( dllimport ) void func1( void );
__declspec( dllimport ) int i;

int *pi = &i;                                // Error in C
static void ( *pf )( void ) = &func1;          // Address of thunk in C,
                                             // function in C++

void func2()
{
    static int *pi = &i;                      // Error in C
    static void ( *pf )( void ) = &func1; // Address of thunk in C,
                                         // function in C++
}
```

However, because a program that includes the **dllexport** attribute in the declaration of an object must provide the definition for that object somewhere in the program, you can initialize a global or local static function pointer with the address of a **dllexport** function. Similarly, you can initialize a global or local static data pointer with the address of a **dllexport** data object. For example, the following code does not generate errors in C or C++:

```

__declspec( dllexport ) void func1( void );
__declspec( dllexport ) int i;

int *pi = &i;                                // Okay
static void ( *pf )( void ) = &func1;          // Okay

void func2()
{
    static int *pi = &i;                      // Okay
    static void ( *pf )( void ) = &func1; // Okay
}

```

- If you apply `dllexport` to a regular class that has a base class that is not marked as `dllexport`, the compiler will generate C4275.

The compiler generates the same warning if the base class is a specialization of a class template. To work around this, mark the base-class with `dllexport`. The problem with a specialization of a class template is where to place the `__declspec(dllexport)`; you are not allowed to mark the class template. Instead, explicitly instantiate the class template and mark this explicit instantiation with `dllexport`. For example:

```

template class __declspec(dllexport) B<int>;
class __declspec(dllexport) D : public B<int> {
// ...

```

This workaround fails if the template argument is the deriving class. For example:

```

class __declspec(dllexport) D : public B<D> {
// ...

```

Because this is common pattern with templates, the compiler changed the semantics of `dllexport` when it is applied to a class that has one or more base-classes and when one or more of the base classes is a specialization of a class template. In this case, the compiler implicitly applies `dllexport` to the specializations of class templates. You can do the following and not get a warning:

```

class __declspec(dllexport) D : public B<D> {
// ...

```

## END Microsoft Specific

## See Also

[dllexport](#), [dllimport](#)

# Using `dllimport` and `dllexport` in C++ Classes

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

You can declare C++ classes with the `dllimport` or `dllexport` attribute. These forms imply that the entire class is imported or exported. Classes exported this way are called exportable classes.

The following example defines an exportable class. All its member functions and static data are exported:

```
#define DllExport __declspec( dllexport )

class DllExport C {
    int i;
    virtual int func( void ) { return 1; }
};
```

Note that explicit use of the `dllimport` and `dllexport` attributes on members of an exportable class is prohibited.

## `dllexport` Classes

When you declare a class `dllexport`, all its member functions and static data members are exported. You must provide the definitions of all such members in the same program. Otherwise, a linker error is generated. The one exception to this rule applies to pure virtual functions, for which you need not provide explicit definitions. However, because a destructor for an abstract class is always called by the destructor for the base class, pure virtual destructors must always provide a definition. Note that these rules are the same for nonexportable classes.

If you export data of class type or functions that return classes, be sure to export the class.

## `dllimport` Classes

When you declare a class `dllimport`, all its member functions and static data members are imported. Unlike the behavior of `dllimport` and `dllexport` on nonclass types, static data members cannot specify a definition in the same program in which a `dllimport` class is defined.

## Inheritance and Exportable Classes

All base classes of an exportable class must be exportable. If not, a compiler warning is generated. Moreover, all accessible members that are also classes must be exportable. This rule permits a `dllexport` class to inherit from a `dllimport` class, and a `dllimport` class to inherit from a `dllexport` class (though the latter is not recommended). As a rule, everything that is accessible to the DLL's client (according to C++ access rules) should be part of the exportable interface. This includes private data members referenced in inline functions.

## Selective Member Import/Export

Because member functions and static data within a class implicitly have external linkage, you can declare them with the `dllimport` or `dllexport` attribute, unless the entire class is exported. If the entire class is imported or exported, the explicit declaration of member functions and data as `dllimport` or `dllexport` is prohibited. If you declare a static data member within a class definition as `dllexport`, a definition must occur somewhere within the same program (as with nonclass external linkage).

Similarly, you can declare member functions with the **dllimport** or **dllexport** attributes. In this case, you must provide a **dllexport** definition somewhere within the same program.

It is worthwhile to note several important points regarding selective member import and export:

- Selective member import/export is best used for providing a version of the exported class interface that is more restrictive; that is, one for which you can design a DLL that exposes fewer public and private features than the language would otherwise allow. It is also useful for fine-tuning the exportable interface: when you know that the client, by definition, is unable to access some private data, you need not export the entire class.
- If you export one virtual function in a class, you must export all of them, or at least provide versions that the client can use directly.
- If you have a class in which you are using selective member import/export with virtual functions, the functions must be in the exportable interface or defined inline (visible to the client).
- If you define a member as **dllexport** but do not include it in the class definition, a compiler error is generated. You must define the member in the class header.
- Although the definition of class members as **dllimport** or **dllexport** is permitted, you cannot override the interface specified in the class definition.
- If you define a member function in a place other than the body of the class definition in which you declared it, a warning is generated if the function is defined as **dllexport** or **dllimport** (if this definition differs from that specified in the class declaration).

#### END Microsoft Specific

## See Also

[dllexport, dllimport](#)

# jitintrinsic

12/21/2017 • 1 min to read • [Edit Online](#)

Marks the function as significant to the 64-bit common language runtime. This is used on certain functions in Microsoft-provided libraries.

## Syntax

```
__declspec(jitintrinsic)
```

## Remarks

`jitintrinsic` adds a MODOPT ([IsJitIntrinsic](#)) to a function signature.

Users are discouraged from using this `__declspec` modifier, as unexpected results can occur.

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# naked (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

For functions declared with the `naked` attribute, the compiler generates code without prolog and epilog code. You can use this feature to write your own prolog/epilog code sequences using inline assembler code. Naked functions are particularly useful in writing virtual device drivers. Note that the `naked` attribute is only valid on x86 and ARM, and is not available on x64.

## Syntax

```
__declspec(naked) declarator
```

## Remarks

Because the `naked` attribute is only relevant to the definition of a function and is not a type modifier, naked functions must use extended attribute syntax and the `__declspec` keyword.

The compiler cannot generate an inline function for a function marked with the `naked` attribute, even if the function is also marked with the `_forceinline` keyword.

The compiler issues an error if the `naked` attribute is applied to anything other than the definition of a non-member method.

## Examples

This code defines a function with the `naked` attribute:

```
__declspec( naked ) int func( formal_parameters ) {}
```

Or, alternately:

```
#define Naked __declspec( naked )
Naked int func( formal_parameters ) {}
```

The `naked` attribute affects only the nature of the compiler's code generation for the function's prolog and epilog sequences. It does not affect the code that is generated for calling such functions. Thus, the `naked` attribute is not considered part of the function's type, and function pointers cannot have the `naked` attribute. Furthermore, the `naked` attribute cannot be applied to a data definition. For example, this code sample generates an error:

```
__declspec( naked ) int i;      // Error--naked attribute not
                                // permitted on data declarations.
```

The `naked` attribute is relevant only to the definition of the function and cannot be specified in the function's prototype. For example, this declaration generates a compiler error:

```
__declspec( naked ) int func(); // Error--naked attribute not
                                // permitted on function declarations
```

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

[Naked Function Calls](#)

# noalias

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

`noalias` means that a function call does not modify or reference visible global state and only modifies the memory pointed to *directly* by pointer parameters (first-level indirections).

If a function is annotated as `noalias`, the optimizer can assume that, in addition to the parameters themselves, only first-level indirections of pointer parameters are referenced or modified inside the function. The visible global state is the set of all data that is not defined or referenced outside of the compilation scope, and their address is not taken. The compilation scope is all source files ([/LTCG \(Link-time Code Generation\)](#) builds) or a single source file (non-[/LTCG](#) build).

## Example

The following sample demonstrates using `__declspec(restrict)` and `__declspec(noalias)`. Normally, memory returned from `malloc` is `restrict` because the CRT headers are decorated appropriately.

However, in this example, the pointers `mempool` and `memptr` are global so the compiler has no assurance that the memory is not subject to aliasing. Decorating the functions that return pointers with `__declspec(restrict)` tells the compiler that the memory pointed to by the return value is not aliased.

Decorating the function in the example that accesses memory with `__declspec(noalias)` tells the compiler that this function does not interfere with the global state except through the pointers in its parameter list.

```

// declspec_noalias.c
#include <stdio.h>
#include <stdlib.h>

#define M 800
#define N 600
#define P 700

float * mempool, * memptr;

__declspec(restrict) float * ma(int size)
{
    float * retval;
    retval = memptr;
    memptr += size;
    return retval;
}

__declspec(restrict) float * init(int m, int n)
{
    float * a;
    int i, j;
    int k=1;

    a = ma(m * n);
    if (!a) exit(1);
    for (i=0; i<m; i++)
        for (j=0; j<n; j++)
            a[i*n+j] = 0.1/k++;
    return a;
}

__declspec(noalias) void multiply(float * a, float * b, float * c)
{
    int i, j, k;

    for (j=0; j<P; j++)
        for (i=0; i<M; i++)
            for (k=0; k<N; k++)
                c[i * P + j] =
                    a[i * N + k] *
                    b[k * P + j];
}

int main()
{
    float * a, * b, * c;

    mempool = (float *) malloc(sizeof(float) * (M*N + N*P + M*P));

    if (!mempool)
    {
        puts("ERROR: Malloc returned null");
        exit(1);
    }

    memptr = mempool;
    a = init(M, N);
    b = init(N, P);
    c = init(M, P);

    multiply(a, b, c);
}

```

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# noinline

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

**`_declspec(noinline)`** tells the compiler to never inline a particular member function (function in a class).

It may be worthwhile to not inline a function if it is small and not critical to the performance of your code. That is, if the function is small and not likely to be called often, such as a function that handles an error condition.

Keep in mind that if a function is marked `noinline`, the calling function will be smaller and thus, itself a candidate for compiler inlining.

```
class X {  
    __declspec(noinline) int mbrfunc() {  
        return 0;  
    } // will not inline  
};
```

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

[inline, \\_\\_inline, \\_\\_forceinline](#)

# noreturn

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

This `__declspec` attribute tells the compiler that a function does not return. As a consequence, the compiler knows that the code following a call to a `__declspec(noreturn)` function is unreachable.

If the compiler finds a function with a control path that does not return a value, it generates a warning (C4715) or error message (C2202). If the control path cannot be reached due to a function that never returns, you can use `__declspec(noreturn)` to prevent this warning or error.

### NOTE

Adding `__declspec(noreturn)` to a function that is expected to return can result in undefined behavior.

## Example

In the following sample, the `else` clause does not contain a return statement. Declaring `fatal` as `__declspec(noreturn)` avoids an error or warning message.

```
// noreturn2.cpp
__declspec(noreturn) extern void fatal () {}

int main() {
    if(1)
        return 1;
    else if(0)
        return 0;
    else
        fatal();
}
```

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# nothrow (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A `__declspec` extended attribute which can be used in the declaration of functions.

## Syntax

```
return-type __declspec(nothrow) [call-convention] function-name ([argument-list])
```

## Remarks

This attribute tells the compiler that the declared function and the functions it calls never throw an exception. With the synchronous exception handling model, now the default, the compiler can eliminate the mechanics of tracking the lifetime of certain unwindable objects in such a function, and significantly reduce the code size. Given the following preprocessor directive, the three function declarations below are equivalent:

```
#define WINAPI __declspec(nothrow) __stdcall

void WINAPI f1();
void __declspec(nothrow) __stdcall f2();
void __stdcall f3() throw();
```

Using `void __declspec(nothrow) __stdcall f2();` has the advantage that you can use an API definition, such as that illustrated by the `#define` statement, to easily specify `nothrow` on a set of functions. The third declaration `, void __stdcall f3() throw();` is the syntax defined by the C++ standard.

## END Microsoft Specific

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# novtable

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

This is a `__declspec` extended attribute.

This form of `__declspec` can be applied to any class declaration, but should only be applied to pure interface classes, that is, classes that will never be instantiated on their own. The `__declspec` stops the compiler from generating code to initialize the vptr in the constructor(s) and destructor of the class. In many cases, this removes the only references to the vtable that are associated with the class and, thus, the linker will remove it. Using this form of `__declspec` can result in a significant reduction in code size.

If you attempt to instantiate a class marked with `novtable` and then access a class member, you will receive an access violation (AV).

## Example

```
// novtable.cpp
#include <stdio.h>

struct __declspec(novtable) X {
    virtual void mf();
};

struct Y : public X {
    void mf() {
        printf_s("In Y\n");
    }
};

int main() {
    // X *pX = new X();
    // pX->mf();    // Causes a runtime access violation.

    Y *pY = new Y();
    pY->mf();
}
```

In Y

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# process

12/21/2017 • 1 min to read • [Edit Online](#)

Specifies that your managed application process should have a single copy of a particular global variable, static member variable, or static local variable shared across all application domains in the process. This is primarily intended to be used when compiling with **/clr:pure**, because under **/clr:pure** global and static variables are per application domain, by default. The **/clr:pure** and **/clr:safe** compiler options are deprecated in Visual Studio 2015. When compiling with **/clr**, global and static variables are per process by default (do not need to use

```
__declspec(process)
```

Only a global variable, a static member variable, or a static local variable of native type can be marked with

```
__declspec(process)
```

When compiling with **/clr:pure**, variables marked as per process must also be declared as `const`. This ensures that per process variables are not changed in one application domain, and giving unexpected results in another application domain. The primary intended use of `__declspec(process)` is to enable compile time initialization of a global variable, static member variable, or static local variable under **/clr:pure**.

`process` is only valid when compiling with **/clr** or **/clr:pure** and is not valid when compiling with **/clr:safe**.

If you want each application domain to have its own copy of a global variable, use [appdomain](#).

See [Application Domains and Visual C++](#) for more information.

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# property (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

This attribute can be applied to non-static "virtual data members" in a class or structure definition. The compiler treats these "virtual data members" as data members by changing their references into function calls.

## Syntax

```
__declspec( property( get=get_func_name ) ) declarator  
__declspec( property( put=put_func_name ) ) declarator  
__declspec( property( get=get_func_name, put=put_func_name ) ) declarator
```

## Remarks

When the compiler sees a data member declared with this attribute on the right of a member-selection operator ("." or "->"), it converts the operation to a **get** or **put** function, depending on whether such an expression is an l-value or an r-value. In more complicated contexts, such as "`+ =`", a rewrite is performed by doing both **get** and **put**.

This attribute can also be used in the declaration of an empty array in a class or structure definition. For example:

```
__declspec(property(get=GetX, put=PutX)) int x[];
```

The above statement indicates that `x[]` can be used with one or more array indices. In this case, `i=p->x[a][b]` will be turned into `i=p->GetX(a, b)`, and `p->x[a][b] = i` will be turned into `p->PutX(a, b, i);`

## END Microsoft Specific

## Example

```
// declspec_property.cpp  
struct S {  
    int i;  
    void putprop(int j) {  
        i = j;  
    }  
  
    int getprop() {  
        return i;  
    }  
  
    __declspec(property(get = getprop, put = putprop)) int the_prop;  
};  
  
int main() {  
    S s;  
    s.the_prop = 5;  
    return s.the_prop;  
}
```

## See Also

[\\_declspec](#)

[Keywords](#)

# restrict

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Applied to a function declaration or definition that returns a pointer type and tells the compiler that the function returns an object that will not be aliased with any other pointers.

## Syntax

```
__declspec(restrict) return_type f();
```

## Remarks

The compiler will propagate `__declspec(restrict)`. For example, the CRT `malloc` function is decorated with `__declspec(restrict)` and therefore, pointers initialized to memory locations with `malloc` are also implied to not be aliased.

The compiler does not check that the pointer is actually not aliased. It is the developer's responsibility to ensure the program does not alias a pointer marked with the `restrict __declspec` modifier.

For similar semantics on variables, see [\\_\\_restrict](#).

## Example

See [noalias](#) for an example using `restrict`.

For information about the restrict keyword that is part of C++ AMP, see [restrict \(C++ AMP\)](#).

## END Microsoft Specific

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# safebuffers

1/3/2018 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Tells the compiler not to insert buffer overrun security checks for a function.

## Syntax

```
_declspec( safebuffers )
```

## Remarks

The **/GS** compiler option causes the compiler to test for buffer overruns by inserting security checks on the stack. The types of data structures that are eligible for security checks are described in [/GS \(Buffer Security Check\)](#). For more information about buffer overrun detection, see [Compiler Security Checks In Depth](#) on the MSDN Web site.

An expert manual code review or external analysis might determine that a function is safe from a buffer overrun. In that case, you can suppress security checks for a function by applying the `_declspec(safebuffers)` keyword to the function declaration.

### Caution

Buffer security checks provide important security protection and have a negligible affect on performance. Therefore, we recommend that you do not suppress them, except in the rare case where the performance of a function is a critical concern and the function is known to be safe.

## Inline Functions

A *primary function* can use an [inlining](#) keyword to insert a copy of a *secondary function*. If the `_declspec(safebuffers)` keyword is applied to a function, buffer overrun detection is suppressed for that function. However, inlining affects the `_declspec(safebuffers)` keyword in the following ways.

Suppose the **/GS** compiler option is specified for both functions, but the primary function specifies the `_declspec(safebuffers)` keyword. The data structures in the secondary function make it eligible for security checks, and the function does not suppress those checks. In this case:

- Specify the [\\_forceinline](#) keyword on the secondary function to force the compiler to inline that function regardless of compiler optimizations.
- Because the secondary function is eligible for security checks, security checks are also applied to the primary function even though it specifies the `_declspec(safebuffers)` keyword.

## Example

The following code shows how to use the `_declspec(safebuffers)` keyword.

```
// compile with: /c /GS
typedef struct {
    int x[20];
} BUFFER;
static int checkBuffers() {
    BUFFER cb;
    // Use the buffer...
    return 0;
};
static __declspec(safebuffers)
int noCheckBuffers() {
    BUFFER ncb;
    // Use the buffer...
    return 0;
}
int wmain() {
    checkBuffers();
    noCheckBuffers();
    return 0;
}
```

## END Microsoft Specific

## See Also

[\\_\\_declspec](#)

[Keywords](#)

[inline, \\_\\_inline, \\_\\_forceinline](#)

[strict\\_gs\\_check](#)

# selectany

12/21/2017 • 2 min to read • [Edit Online](#)

## Microsoft Specific

Tells the compiler that the declared global data item (variable or object) is a pick-any COMDAT (a packaged function).

## Syntax

```
__declspec( selectany ) declarator
```

## Remarks

At link time, if multiple definitions of a COMDAT are seen, the linker picks one and discards the rest. If the linker option [/OPT:REF](#) (Optimizations) is selected, then COMDAT elimination will occur to remove all the unreferenced data items in the linker output.

Constructors and assignment by global function or static methods in the declaration do not create a reference and will not prevent /OPT:REF elimination. Side effects from such code should not be depended on when no other references to the data exist.

For dynamically initialized, global objects, `selectany` will discard an unreferenced object's initialization code, as well.

A global data item can normally be initialized only once in an EXE or DLL project. `selectany` can be used in initializing global data defined by headers, when the same header appears in more than one source file. `selectany` is available in both the C and C++ compilers.

### NOTE

`selectany` can only be applied to the actual initialization of global data items that are externally visible.

## Example

This code shows how to use the `selectany` attribute:

```

//Correct - x1 is initialized and externally visible
__declspec(selectany) int x1=1;

//Incorrect - const is by default static in C++, so
//x2 is not visible externally (This is OK in C, since
//const is not by default static in C)
const __declspec(selectany) int x2 =2;

//Correct - x3 is extern const, so externally visible
extern const __declspec(selectany) int x3=3;

//Correct - x4 is extern const, so it is externally visible
extern const int x4;
const __declspec(selectany) int x4=4;

//Incorrect - __declspec(selectany) is applied to the uninitialized
//declaration of x5
extern __declspec(selectany) int x5;

// OK: dynamic initialization of global object
class X {
public:
X(int i){i++;};
int i;
};

__declspec(selectany) X x(1);

```

## Example

This code shows how to use the `selectany` attribute to ensure data COMDAT folding when you also use the `/OPT:ICF` linker option. Note that data must be marked with `selectany` and placed in a **const** (readonly) section. You must explicitly specify the read-only section.

```

// selectany2.cpp
// in the following lines, const marks the variables as read only
__declspec(selectany) extern const int ix = 5;
__declspec(selectany) extern const int jx = 5;
int main() {
    int ij;
    ij = ix + jx;
}

```

## END Microsoft Specific

## See Also

[\\_\\_declspec](#)  
[Keywords](#)

# thread

12/21/2017 • 3 min to read • [Edit Online](#)

## Microsoft Specific

The **thread** extended storage-class modifier is used to declare a thread local variable. For the portable equivalent in C++11 and later, use the [thread\\_local](#) storage class specifier for portable code. On Windows **thread\_local** is implemented with `__declspec(thread)`.

## Syntax

```
__declspec( thread ) declarator
```

## Remarks

Thread Local Storage (TLS) is the mechanism by which each thread in a multithreaded process allocates storage for thread-specific data. In standard multithreaded programs, data is shared among all threads of a given process, whereas thread local storage is the mechanism for allocating per-thread data. For a complete discussion of threads, see [Multithreading](#).

Declarations of thread local variables must use [extended attribute syntax](#) and the `__declspec` keyword with the **thread** keyword. For example, the following code declares an integer thread local variable and initializes it with a value:

```
__declspec( thread ) int tls_i = 1;
```

When using thread-local variables in dynamically-loaded libraries, you need to be aware of factors that can cause a thread-local variable to not be initialized correctly:

1) If the variable is initialized with a function call (including constructors), this function will only be called for the thread that caused the binary/DLL to load into the process, and for those threads that started after the binary/DLL was loaded. The initialization functions are not called for any other thread that was already running when the DLL was loaded. Dynamic initialization occurs on the DllMain call for DLL\_THREAD\_ATTACH, but the DLL never gets that message if the DLL isn't in the process when the thread starts.

2) Thread-local variables that are initialized statically with constant values are generally initialized properly on all threads. However, as of December 2017 there is a known conformance issue in the Microsoft C++ compiler whereby `constexpr` variables receive dynamic rather than static initialization.

Note: Both of these issues are expected to be fixed in future updates of the compiler.

Additionally, you must observe these guidelines when declaring thread local objects and variables:

- You can apply the **thread** attribute only to class and data declarations and definitions; **thread** cannot be used on function declarations or definitions.
- You can specify the **thread** attribute only on data items with static storage duration. This includes global data objects (both **static** and **extern**), local static objects, and static data members of classes. You cannot declare automatic data objects with the **thread** attribute.
- You must use the **thread** attribute for the declaration and the definition of a thread local object, whether the declaration and definition occur in the same file or separate files.

- You cannot use the **thread** attribute as a type modifier.
- Because the declaration of objects that use the **thread** attribute is permitted, these two examples are semantically equivalent:

```
// declspec_thread_2.cpp
// compile with: /LD
__declspec( thread ) class B {
public:
    int data;
} BObject; // BObject declared thread local.

class B2 {
public:
    int data;
};
__declspec( thread ) B2 BObject2; // BObject2 declared thread local.
```

- Standard C permits initialization of an object or variable with an expression involving a reference to itself, but only for objects of nonstatic extent. Although C++ normally permits such dynamic initialization of an object with an expression involving a reference to itself, this type of initialization is not permitted with thread local objects. For example:

```
// declspec_thread_3.cpp
// compile with: /LD
#define Thread __declspec( thread )
int j = j; // Okay in C++; C error
Thread int tls_i = sizeof( tls_i ); // Okay in C and C++
```

Note that a **sizeof** expression that includes the object being initialized does not constitute a reference to itself and is allowed in C and C++.

#### **END Microsoft Specific**

## See Also

[\\_\\_declspec](#)  
[Keywords](#)  
[Thread Local Storage \(TLS\)](#)

# uuid (C++)

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The compiler attaches a GUID to a class or structure declared or defined (full COM object definitions only) with the `uuid` attribute.

## Syntax

```
__declspec( uuid("ComObjectGUID") ) declarator
```

## Remarks

The `uuid` attribute takes a string as its argument. This string names a GUID in normal registry format with or without the {} delimiters. For example:

```
struct __declspec(uuid("00000000-0000-0000-c000-000000000046")) IUnknown;
struct __declspec(uuid("{00020400-0000-0000-c000-000000000046}") IDispatch;
```

This attribute can be applied in a redeclaration. This allows the system headers to supply the definitions of interfaces such as **IUnknown**, and the redeclaration in some other header (such as COMDEF.H) to supply the GUID.

The keyword `_uuidof` can be applied to retrieve the constant GUID attached to a user-defined type.

**END Microsoft Specific**

## See Also

[\\_\\_declspec](#)

[Keywords](#)

# `_restrict`

12/21/2017 • 1 min to read • [Edit Online](#)

Like the `_declspec ( restrict )` modifier, the `_restrict` keyword indicates that a symbol is not aliased in the current scope. The `_restrict` keyword differs from the `_declspec ( restrict )` modifier in the following ways:

- The `_restrict` keyword is valid only on variables, and `_declspec ( restrict )` is only valid on function declarations and definitions.
- `_restrict` is similar to `restrict` from the C99 spec, but `_restrict` can be used in C++ or C programs.
- When `_restrict` is used, the compiler will not propagate the no-alias property of a variable. That is, if you assign a `_restrict` variable to a non-`_restrict` variable, the compiler will still allow the non-`_restrict` variable to be aliased. This is different from the behavior of the `restrict` keyword from the C99 specification.

Generally, if you affect the behavior of an entire function, it is better to use `_declspec ( restrict )` than the keyword.

In Visual Studio 2015 and later, `_restrict` can be used on C++ references.

## NOTE

When used on a variable that also has the `volatile` keyword, `volatile` will take precedence.

## Example

```
// _restrict_keyword.c
// compile with: /LD
// In the following function, declare a and b as disjoint arrays
// but do not have same assurance for c and d.
void sum2(int n, int * __restrict a, int * __restrict b,
          int * c, int * d) {
    int i;
    for (i = 0; i < n; i++) {
        a[i] = b[i] + c[i];
        c[i] = b[i] + d[i];
    }
}

// By marking union members as __restrict, tell compiler that
// only z.x or z.y will be accessed in any given scope.
union z {
    int * __restrict x;
    double * __restrict y;
};
```

## See Also

[Keywords](#)

# `_sptr`, `_uptr`

12/21/2017 • 2 min to read • [Edit Online](#)

## Microsoft Specific

Use the `_sptr` or `_uptr` modifier on a 32-bit pointer declaration to specify how the compiler converts a 32-bit pointer to a 64-bit pointer. A 32-bit pointer is converted, for example, when it is assigned to a 64-bit pointer variable or is dereferenced on a 64-bit platform.

Microsoft documentation for support of 64-bit platforms sometimes refers to the most significant bit of a 32-bit pointer as the sign bit. By default, the compiler uses sign extension to convert a 32-bit pointer to a 64-bit pointer. That is, the least significant 32 bits of the 64-bit pointer are set to the value of the 32-bit pointer and the most significant 32 bits are set to the value of the sign bit of the 32-bit pointer. This conversion yields correct results if the sign bit is 0, but not if the sign bit is 1. For example, the 32-bit address 0xFFFFFFFF yields the equivalent 64-bit address 0x000000007FFFFFFF, but the 32-bit address 0x80000000 is incorrectly changed to 0xFFFFFFFF80000000.

The `_sptr`, or signed pointer, modifier specifies that a pointer conversion set the most significant bits of a 64-bit pointer to the sign bit of the 32-bit pointer. The `_uptr`, or unsigned pointer, modifier specifies that a conversion set the most significant bits to zero. The following declarations show the `_sptr` and `_uptr` modifiers used with two unqualified pointers, two pointers qualified with the `_ptr32` type, and a function parameter.

```
int * __sptr psp;
int * __uptr pup;
int * __ptr32 __sptr psp32;
int * __ptr32 __uptr pup32;
void MyFunction(char * __uptr __ptr32 myValue);
```

Use the `_sptr` and `_uptr` modifiers with pointer declarations. Use the modifiers in the position of a [pointer type qualifier](#), which means the modifier must follow the asterisk. You cannot use the modifiers with [pointers to members](#). The modifiers do not affect non-pointer declarations.

## Example

The following example declares 32-bit pointers that use the `_sptr` and `_uptr` modifiers, assigns each 32-bit pointer to a 64-bit pointer variable, and then displays the hexadecimal value of each 64-bit pointer. The example is compiled with the native 64-bit compiler and is executed on a 64-bit platform.

```

// sptr_uptr.cpp
// processor: x64
#include "stdio.h"

int main()
{
    void *      __ptr64 p64;
    void *      __ptr32 p32d; //default signed pointer
    void * __sptr __ptr32 p32s; //explicit signed pointer
    void * __uptr __ptr32 p32u; //explicit unsigned pointer

    // Set the 32-bit pointers to a value whose sign bit is 1.
    p32d = reinterpret_cast<void *>(0x87654321);
    p32s = p32d;
    p32u = p32d;

    // The printf() function automatically displays leading zeroes with each 32-bit pointer. These are unrelated
    // to the __sptr and __uptr modifiers.
    printf("Display each 32-bit pointer (as an unsigned 64-bit pointer):\n");
    printf("p32d:      %p\n", p32d);
    printf("p32s:      %p\n", p32s);
    printf("p32u:      %p\n", p32u);

    printf("\nDisplay the 64-bit pointer created from each 32-bit pointer:\n");
    p64 = p32d;
    printf("p32d: p64 = %p\n", p64);
    p64 = p32s;
    printf("p32s: p64 = %p\n", p64);
    p64 = p32u;
    printf("p32u: p64 = %p\n", p64);
    return 0;
}

```

```

Display each 32-bit pointer (as an unsigned 64-bit pointer):
p32d:      000000087654321
p32s:      000000087654321
p32u:      000000087654321

```

```

Display the 64-bit pointer created from each 32-bit pointer:
p32d: p64 = FFFFFFFF87654321
p32s: p64 = FFFFFFFF87654321
p32u: p64 = 000000087654321

```

## END Microsoft Specific

## See Also

[Microsoft-Specific Modifiers](#)

# `_unaligned`

12/21/2017 • 1 min to read • [Edit Online](#)

When you declare a pointer with the `_unaligned` modifier, the compiler assumes that the pointer addresses data that is not aligned. Consequently, for an application that targets an Itanium Processor Family (IPF) computer, the compiler generates code that reads the unaligned data one byte at a time.

## Remarks

The `_unaligned` modifier is valid for the x64 and Itanium compilers but affects only applications that target an IPF computer. This modifier describes the alignment of the addressed data only; the pointer itself is assumed to be aligned.

The Itanium processor generates an alignment fault when it accesses misaligned data, and the time to process the fault weakens performance. Use the `_unaligned` modifier to cause the processor to read data one byte at a time and avoid the fault. This modifier is not required for x64 applications because the x64 processor handles misaligned data without faulting.

For more information about alignment, see:

- [align](#)
- [\\_alignof Operator](#)
- [pack](#)
- [/Zp \(Struct Member Alignment\)](#)
- [Examples of Structure Alignment](#)

## Example

```
// unaligned_keyword.cpp
// compile with: /c
// processor: x64 IPF
#include <stdio.h>
int main() {
    char buf[100];

    int __unaligned *p1 = (int*)(&buf[37]);
    int *p2 = (int *)p1;

    *p1 = 0;    // ok

    __try {
        *p2 = 0; // throws an exception
    }
    __except(1) {
        puts("exception");
    }
}
```

## See Also

[Keywords](#)



# `_w64`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

This keyword is obsolete. In versions of Visual Studio earlier than Visual Studio 2013, this lets you mark variables, so that when you compile with `/Wp64` the compiler will report any warnings that would be reported if you were compiling with a 64-bit compiler.

## Syntax

`type __w64 identifier`

### Parameters

`type`

One of the three types that could cause problems in code being ported from a 32-bit to a 64-bit compiler: `int`, `long`, or a pointer.

`identifier`

The identifier for the variable you are creating.

## Remarks

### IMPORTANT

The `/Wp64` compiler option and `__w64` keyword are deprecated in Visual Studio 2010 and Visual Studio 2013 and removed starting in Visual Studio 2013. If you use the `/Wp64` compiler option on the command line, the compiler issues [Command-Line Warning D9002](#). The `__w64` keyword is silently ignored. Instead of using this option and keyword to detect 64-bit portability issues, use a Visual C++ compiler that targets a 64-bit platform. For more information, see [Configure Visual C++ for 64-bit, x64 targets](#).

Any typedef that has `__w64` on it must be 32 bits on x86 and 64 bits on x64.

To detect portability issues by using versions of the Visual C++ compiler earlier than Visual Studio 2010, the `__w64` keyword should be specified on any typedefs that change size between 32 bit and 64 bit platforms. For any such type, `__w64` must appear only on the 32-bit definition of the typedef.

The `__w64` keyword is ignored if the compilation does not use `/Wp64`.

For more information about porting to 64-bit, see the following topics:

- [Compiler Options](#)
- [Porting 32-Bit Code to 64-Bit Code](#)
- [Configure Visual C++ for 64-bit, x64 targets](#)

## Example

```
// __w64.cpp
// compile with: /W3 /WP64
typedef int Int_32;
#ifndef _WIN64
typedef __int64 Int_Native;
#else
typedef int __w64 Int_Native;
#endif

int main() {
    Int_32 i0 = 5;
    Int_Native i1 = 10;
    i0 = i1;    // C4244 64-bit int assigned to 32-bit int

    // char __w64 c;  error, cannot use __w64 on char
}
```

## See Also

[Keywords](#)

# func

12/21/2017 • 1 min to read • [Edit Online](#)

**(C++11)** The predefined identifier `__func__` is implicitly defined as a string that contains the unqualified and unadorned name of the enclosing function. `__func__` is mandated by the C++ standard and is not a Microsoft extension.

## Syntax

```
__func__
```

## Return Value

Returns a null-terminated const char array of characters that contains the function name.

## Example

```
#include <string>
#include <iostream>

namespace Test
{
    struct Foo
    {
        static void DoSomething(int i, std::string s)
        {
            std::cout << __func__ << std::endl; // Output: DoSomething
        }
    };
}

int main()
{
    Test::Foo::DoSomething(42, "Hello");

    return 0;
}
```

## Requirements

C++11

# Compiler COM Support

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The Visual C++ compiler can directly read component object model (COM) type libraries and translate the contents into C++ source code that can be included in the compilation. Language extensions are available to facilitate COM programming on the client side.

By using the [#import preprocessor directive](#), the compiler can read a type library and convert it into a C++ header file that describes the COM interfaces as classes. A set of `#import` attributes is available for user control of the content for the resulting type library header files.

You can use the `_declspec` extended attribute `uuid` to assign a globally unique identifier (GUID) to a COM object. The keyword `_uuidof` can be used to extract the GUID associated with a COM object. Another `_declspec` attribute, `property`, can be used to specify the `get` and `set` methods for a data member of a COM object.

A set of COM support global functions and classes is provided to support the **VARIANT** and `BSTR` types, implement smart pointers, and encapsulate the error object thrown by `_com_raise_error`:

- [Compiler COM Global Functions](#)
- [\\_bstr\\_t](#)
- [\\_com\\_error](#)
- [\\_com\\_ptr\\_t](#)
- [\\_variant\\_t](#)

## END Microsoft Specific

## See Also

[Compiler COM Support Classes](#)  
[Compiler COM Global Functions](#)

# Compiler COM Global Functions

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

The following routines are available:

FUNCTION	DESCRIPTION
<a href="#">_com_raise_error</a>	Throws a <a href="#">_com_error</a> in response to a failure.
<a href="#">_set_com_error_handler</a>	Replaces the default function that is used for COM error-handling.
<a href="#">ConvertBSTRToString</a>	Converts a <code>BSTR</code> value to a <code>char \*</code> .
<a href="#">ConvertStringToBSTR</a>	Converts a <code>char \*</code> value to a <code>BSTR</code> .

## END Microsoft Specific

## See Also

[Compiler COM Support Classes](#)

[Compiler COM Support](#)

# \_com\_raise\_error

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Throws a [\\_com\\_error](#) in response to a failure.

## Syntax

```
void __stdcall _com_raise_error(
    HRESULT hr,
    IErrorInfo* perrinfo = 0
);
```

### Parameters

`hr`

`HRESULT` information.

`perrinfo`

`IErrorInfo` object.

## Remarks

`_com_raise_error`, which is defined in comdef.h, can be replaced by a user-written version of the same name and prototype. This could be done if you want to use `#import` but do not want to use C++ exception handling. In that case, a user version of `_com_raise_error` might decide to do a `longjmp` or display a message box and halt. The user version should not return, though, because the compiler COM support code does not expect it to return.

You can also use [\\_set\\_com\\_error\\_handler](#) to replace the default error-handling function.

By default, `_com_raise_error` is defined as follows:

```
void __stdcall _com_raise_error(HRESULT hr, IErrorInfo* perrinfo) {
    throw _com_error(hr, perrinfo);
}
```

## END Microsoft Specific

## Requirements

**Header:** comdef.h

**Lib:** If the **wchar\_t is Native Type** compiler option is on, use comsuppw.lib or comsuppwd.lib. If **wchar\_t is Native Type** is off, use comsupp.lib. For more information, see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#).

## See Also

[Compiler COM Global Functions](#)  
[\\_set\\_com\\_error\\_handler](#)

# ConvertStringToBSTR

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Converts a `char \*` value to a `BSTR`.

## Syntax

```
BSTR __stdcall ConvertStringToBSTR(
    const char* pSrc
)
```

### Parameters

`pSrc`

A `char \*` variable.

## Example

```
// ConvertStringToBSTR.cpp
#include <comutil.h>
#include <stdio.h>

#pragma comment(lib, "comsuppw.lib")
#pragma comment(lib, "kernel32.lib")

int main() {
    char* lpszText = "Test";
    printf_s("char * text: %s\n", lpszText);

    BSTR bstrText = _com_util::ConvertStringToBSTR(lpszText);
    wprintf_s(L"BSTR text: %s\n", bstrText);

    SysFreeString(bstrText);
}
```

```
char * text: Test
BSTR text: Test
```

## END Microsoft Specific

## Requirements

**Header:** comutil.h

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Global Functions](#)

# ConvertBSTRToString

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Converts a `BSTR` value to a `char \*`.

## Syntax

```
char* __stdcall ConvertBSTRToString(  
    BSTR pSrc  
) ;
```

### Parameters

`pSrc`

A BSTR variable.

## Remarks

`ConvertBSTRToString` allocates a string you must delete.

## Example

```
// ConvertBSTRToString.cpp  
#include <comutil.h>  
#include <stdio.h>  
  
#pragma comment(lib, "comsuppw.lib")  
  
int main() {  
    BSTR bstrText = ::SysAllocString(L"Test");  
    wprintf_s(L"BSTR text: %s\n", bstrText);  
  
    char* lpszText2 = _com_util::ConvertBSTRToString(bstrText);  
    printf_s("char * text: %s\n", lpszText2);  
  
    SysFreeString(bstrText);  
    delete[] lpszText2;  
}
```

```
BSTR text: Test  
char * text: Test
```

## END Microsoft Specific

## Requirements

**Header:** comutil.h.

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Global Functions](#)

# \_set\_com\_error\_handler

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Replaces the default function that is used for COM error-handling.

## Syntax

```
void __stdcall _set_com_error_handler(
    void (__stdcall *pHandler)(
        HRESULT hr,
        IErrorInfo* perrinfo
    )
);
```

### Parameters

`pHandler`

Pointer to the replacement function.

`hr`

`HRESULT` information.

`perrinfo`

`IErrorInfo` object.

## Remarks

By default, `_com_raise_error` handles all COM errors. You can change this behavior by using `_set_com_error_handler` to call your own error-handling function.

The replacement function must have a signature that is equivalent to that of `_com_raise_error`.

## Example

```

// _set_com_error_handler.cpp
// compile with /EHsc
#include <stdio.h>
#include <comdef.h>
#include <comutil.h>

// Importing ado.dll to attempt to establish an ado connection.
// Not related to _set_com_error_handler
#import "C:\Program Files\Common Files\System\ado\msado15.dll" no_namespace rename("EOF", "adoEOF")

void __stdcall _My_com_raise_error(HRESULT hr, IErrorInfo* perrinfo)
{
    throw "Unable to establish the connection!";
}

int main()
{
    _set_com_error_handler(_My_com_raise_error);
    _bstr_t bstrEmpty(L"");
    _ConnectionPtr Connection = NULL;
    try
    {
        Connection.CreateInstance(__uuidof(Connection));
        Connection->Open(bstrEmpty, bstrEmpty, bstrEmpty, 0);
    }
    catch(char* errorMessage)
    {
        printf("Exception raised: %s\n", errorMessage);
    }

    return 0;
}

```

Exception raised: Unable to establish the connection!

## Requirements

**Header:** comdef.h

**Lib:** If the **wchar\_t is Native Type** compiler option is on, use comsuppw.lib or comsuppwd.lib. If **wchar\_t is Native Type** is off, use comsupp.lib. For more information, see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#).

## See Also

[Compiler COM Global Functions](#)

# Compiler COM Support Classes

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Standard classes are used to support some of the COM types. The classes are defined in comdef.h and the header files generated from the type library.

CLASS	PURPOSE
<code>_bstr_t</code>	Wraps the <code>BSTR</code> type to provide useful operators and methods.
<code>_com_error</code>	Defines the error object thrown by <code>_com_raise_error</code> in most failures.
<code>_com_ptr_t</code>	Encapsulates COM interface pointers, and automates the required calls to <code>AddRef</code> , <code>Release</code> , and <code>QueryInterface</code> .
<code>_variant_t</code>	Wraps the <b>VARIANT</b> type to provide useful operators and methods.

**END Microsoft Specific**

## See Also

- [Compiler COM Support](#)
- [Compiler COM Global Functions](#)
- [C++ Language Reference](#)

# \_bstr\_t Class

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A `_bstr_t` object encapsulates the [BSTR data type](#). The class manages resource allocation and deallocation through function calls to **SysAllocString** and **SysFreeString** and other `BSTR` APIs when appropriate. The `_bstr_t` class uses reference counting to avoid excessive overhead.

## Construction

<code>_bstr_t</code>	Constructs a <code>_bstr_t</code> object.
----------------------	---

## Operations

<code>Assign</code>	Copies a <code>BSTR</code> into the <code>BSTR</code> wrapped by a <code>_bstr_t</code> .
<code>Attach</code>	Links a <code>_bstr_t</code> wrapper to a <code>BSTR</code> .
<code>copy</code>	Constructs a copy of the encapsulated <code>BSTR</code> .
<code>Detach</code>	Returns the <code>BSTR</code> wrapped by a <code>_bstr_t</code> and detaches the <code>BSTR</code> from the <code>_bstr_t</code> .
<code>GetAddress</code>	Points to the <code>BSTR</code> wrapped by a <code>_bstr_t</code> .
<code>GetBSTR</code>	Points to the beginning of the <code>BSTR</code> wrapped by the <code>_bstr_t</code> .
<code>length</code>	Returns the number of characters in the <code>_bstr_t</code> .

## Operators

<code>operator =</code>	Assigns a new value to an existing <code>_bstr_t</code> object.
<code>operator +=</code>	Appends characters to the end of the <code>_bstr_t</code> object.
<code>operator +</code>	Concatenates two strings.
<code>operator !</code>	Checks if the encapsulated <code>BSTR</code> is a <b>NULL</b> string.
<code>operator ==, !=, &lt;, &gt;, &lt;=, &gt;=</code>	Compares two <code>_bstr_t</code> objects.
<code>operator wchar_t*   char*</code>	Extract the pointers to the encapsulated Unicode or multibyte <code>BSTR</code> object.

**END Microsoft Specific**

## Requirements

**Header:** comutil.h

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Support Classes](#)

# `_bstr_t` Member Functions

12/21/2017 • 1 min to read • [Edit Online](#)

For information about `_bstr_t` member functions, see [\\_bstr\\_t Class](#).

## See Also

[\\_bstr\\_t Class](#)

# \_bstr\_t::Assign

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Copies a `BSTR` into the `BSTR` wrapped by a `_bstr_t`.

## Syntax

```
void Assign(  
    BSTR s  
>);
```

### Parameters

`s`

A `BSTR` to be copied into the `BSTR` wrapped by a `_bstr_t`.

## Remarks

`Assign` does a binary copy, which means the entire length of the `BSTR` is copied, regardless of content.

## Example

```

// _bstr_t_Assign.cpp

#include <comdef.h>
#include <stdio.h>

int main()
{
    // creates a _bstr_t wrapper
    _bstr_t bstrWrapper;

    // creates BSTR and attaches to it
    bstrWrapper = "some text";
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));

    // bstrWrapper releases its BSTR
    BSTR bstr = bstrWrapper.Detach();
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));
    // "some text"
    wprintf_s(L"bstr = %s\n", bstr);

    bstrWrapper.Attach(SysAllocString(OLESTR("SysAllocatedString")));
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));

    // assign a BSTR to our _bstr_t
    bstrWrapper.Assign(bstr);
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));

    // done with BSTR, do manual cleanup
    SysFreeString(bstr);

    // resuse bstr
    bstr= SysAllocString(OLESTR("Yet another string"));
    // two wrappers, one BSTR
    _bstr_t bstrWrapper2 = bstrWrapper;

    *bstrWrapper.GetAddress() = bstr;

    // bstrWrapper and bstrWrapper2 do still point to BSTR
    bstr = 0;
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));
    wprintf_s(L"bstrWrapper2 = %s\n",
             static_cast<wchar_t*>(bstrWrapper2));

    // new value into BSTR
    _snwprintf_s(bstrWrapper.GetBSTR(), 100, bstrWrapper.length(),
                 L"changing BSTR");
    wprintf_s(L"bstrWrapper = %s\n",
             static_cast<wchar_t*>(bstrWrapper));
    wprintf_s(L"bstrWrapper2 = %s\n",
             static_cast<wchar_t*>(bstrWrapper2));
}

```

```
bstrWrapper = some text
bstrWrapper = (null)
bstr = some text
bstrWrapper = SysAllocedString
bstrWrapper = some text
bstrWrapper = Yet another string
bstrWrapper2 = some text
bstrWrapper = changing BSTR
bstrWrapper2 = some text
```

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# \_bstr\_t::Attach

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Links a `_bstr_t` wrapper to a `BSTR`.

## Syntax

```
void Attach(  
    BSTR s  
) ;
```

### Parameters

`s`

A `BSTR` to be associated with, or assigned to, the `_bstr_t` variable.

## Remarks

If the `_bstr_t` was previously attached to another `BSTR`, the `_bstr_t` will clean up the `BSTR` resource, if no other `_bstr_t` variables are using the `BSTR`.

## Example

See [\\_bstr\\_t::Assign](#) for an example using **Attach**.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::_bstr_t`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Constructs a `_bstr_t` object.

## Syntax

```
_bstr_t( ) throw( );
_bstr_t(
    const _bstr_t& s1
) throw( );
_bstr_t(
    const char* s2
);
_bstr_t(
    const wchar_t* s3
);
_bstr_t(
    const _variant_t& var
);
_bstr_t(
    BSTR bstr,
    bool fCopy
);
```

### Parameters

`s1`

A `_bstr_t` object to be copied.

`s2`

A multibyte string.

`s3`

A Unicode string

`var`

A `_variant_t` object.

`bstr`

An existing `BSTR` object.

`fCopy`

If `false`, the `bstr` argument is attached to the new object without making a copy by calling `SysAllocString`.

## Remarks

The following table describes the `_bstr_t` constructors.

CONSTRUCTOR	DESCRIPTION
<code>_bstr_t( )</code>	Constructs a default <code>_bstr_t</code> object that encapsulates a null <code>BSTR</code> object.

CONSTRUCTOR	DESCRIPTION
<code data-bbox="165 175 450 204"><code>_bstr_t( _bstr_t&amp; s1 )</code></code>	<p>Constructs a <code data-bbox="822 175 1044 204"><code>_bstr_t</code></code> object as a copy of another.</p> <p>This is a <i>shallow copy</i>, which increments the reference count of the encapsulated <code data-bbox="1028 271 1091 300"><code>BSTR</code></code> object instead of creating a new one.</p>
<code data-bbox="165 399 418 428"><code>_bstr_t( char* s2 )</code></code>	<p>Constructs a <code data-bbox="822 399 1044 428"><code>_bstr_t</code></code> object by calling <code data-bbox="1218 399 1377 428"><code>SysAllocString</code></code> to create a new <code data-bbox="1028 435 1091 464"><code>BSTR</code></code> object and then encapsulates it.</p> <p>This constructor first performs a multibyte to Unicode conversion.</p>
<code data-bbox="165 624 450 653"><code>_bstr_t( wchar_t* s3 )</code></code>	<p>Constructs a <code data-bbox="822 624 1044 653"><code>_bstr_t</code></code> object by calling <code data-bbox="1218 624 1377 653"><code>SysAllocString</code></code> to create a new <code data-bbox="1028 660 1091 689"><code>BSTR</code></code> object and then encapsulates it.</p>
<code data-bbox="165 720 498 750"><code>_bstr_t( _variant_t&amp; var )</code></code>	<p>Constructs a <code data-bbox="822 720 1044 750"><code>_bstr_t</code></code> object from a <code data-bbox="1218 720 1377 750"><code>_variant_t</code></code> object by first retrieving a <code data-bbox="1028 756 1091 786"><code>BSTR</code></code> object from the encapsulated VARIANT object.</p>
<code data-bbox="165 871 609 900"><code>_bstr_t( BSTR bstr, bool fCopy )</code></code>	<p>Constructs a <code data-bbox="822 871 1044 900"><code>_bstr_t</code></code> object from an existing <code data-bbox="1298 871 1361 900"><code>BSTR</code></code> (as opposed to a <code data-bbox="1028 907 1091 936"><code>wchar_t*</code></code> string). If <code data-bbox="1218 907 1282 936"><code>fCopy</code></code> is false, the supplied <code data-bbox="1028 943 1091 972"><code>BSTR</code></code> is attached to the new object without making a new copy with <code data-bbox="1028 979 1187 1008"><code>SysAllocString</code></code>.</p> <p>This constructor is used by wrapper functions in the type library headers to encapsulate and take ownership of a <code data-bbox="1377 1064 1440 1093"><code>BSTR</code></code> that is returned by an interface method.</p>

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)  
[\\_variant\\_t Class](#)

# \_bstr\_t::copy

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Constructs a copy of the encapsulated `BSTR`.

## Syntax

```
BSTR copy(  
    bool fCopy = true  
) const;
```

### Parameters

`fCopy`

If `true`, `copy` returns a copy of the contained `BSTR`, otherwise `copy` returns the actual `BSTR`.

## Remarks

Returns a newly allocated copy of the encapsulated `BSTR` object.

## Example

```
STDMETHODIMP CAAlertMsg::get_ConnectionStr(BSTR *pVal){ // m_bsConStr is _bstr_t  
    *pVal = m_bsConStr.copy();  
}
```

**END Microsoft Specific**

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::Detach`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Returns the `BSTR` wrapped by a `_bstr_t` and detaches the `BSTR` from the `_bstr_t`.

## Syntax

```
BSTR Detach( ) throw;
```

## Return Value

The `BSTR` wrapped by the `_bstr_t`.

## Example

See [\\_bstr\\_t::Assign](#) for a example using **Detach**.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::GetAddress`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Frees any existing string and returns the address of a newly allocated string.

## Syntax

```
BSTR* GetAddress( );
```

## Return Value

A pointer to the `BSTR` wrapped by the `_bstr_t`.

## Remarks

`GetAddress` affects all `_bstr_t` objects that share a `BSTR`. More than one `_bstr_t` can share a `BSTR` through the use of the copy constructor and and `operator=`.

## Example

See `_bstr_t::Assign` for a example using `GetAddress`.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::GetBSTR`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Points to the beginning of the `BSTR` wrapped by the `_bstr_t`.

## Syntax

```
BSTR& GetBSTR( );
```

## Return Value

The beginning of the `BSTR` wrapped by the `_bstr_t`.

## Remarks

`GetBSTR` affects all `_bstr_t` objects that share a `BSTR`. More than one `_bstr_t` can share a `BSTR` through the use of the copy constructor and and `operator=`.

## Example

See [\\_bstr\\_t::Assign](#) for an example using `GetBSTR`.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# \_bstr\_t::length

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Returns the number of characters in the `_bstr_t`, not including the terminating null, of the encapsulated `BSTR`.

## Syntax

```
unsigned int length( ) const throw( );
```

## Remarks

**END Microsoft Specific**

## See Also

[\\_bstr\\_t Class](#)

# \_bstr\_t Operators

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_bstr_t` operators, see [\\_bstr\\_t Class](#).

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::operator =`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Assigns a new value to an existing `_bstr_t` object.

## Syntax

```
_bstr_t& operator=(  
    const _bstr_t& s1  
) throw ( );  
_bstr_t& operator=(  
    const char* s2  
)  
_bstr_t& operator=(  
    const wchar_t* s3  
)  
_bstr_t& operator=(  
    const _variant_t& var  
)
```

### Parameters

`s1`

A `_bstr_t` object to be assigned to an existing `_bstr_t` object.

`s2`

A multibyte string to be assigned to an existing `_bstr_t` object.

`s3`

A Unicode string to be assigned to an existing `_bstr_t` object.

`var`

A `_variant_t` object to be assigned to an existing `_bstr_t` object.

## END Microsoft Specific

## Example

See [\\_bstr\\_t::Assign](#) for an example of using `operator=`.

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::operator +=, +`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Appends characters to the end of the `_bstr_t` object or concatenates two strings.

## Syntax

```
_bstr_t& operator+=(  
    const _bstr_t& s1  
)  
_bstr_t operator+(  
    const _bstr_t& s1  
)  
friend _bstr_t operator+(  
    const char* s2,  
    const _bstr_t& s1  
)  
friend _bstr_t operator+(  
    const wchar_t* s3,  
    const _bstr_t& s1  
)
```

### Parameters

`s1`

A `_bstr_t` object.

`s2`

A multibyte string.

`s3`

A Unicode string.

## Remarks

These operators perform string concatenation:

- **`operator+=( s1 )`** Appends the characters in the encapsulated `BSTR` of `s1` to the end of this object's encapsulated `BSTR`.
- **`operator+( s1 )`** Returns the new `_bstr_t` that is formed by concatenating this object's `BSTR` with that of `s1`.
- **`operator+( s2 | s1 )`** Returns a new `_bstr_t` that is formed by concatenating a multibyte string `s2`, converted to Unicode, with the `BSTR` encapsulated in `s1`.
- **`operator+( s3 , s1 )`** Returns a new `_bstr_t` that is formed by concatenating a Unicode string `s3` with the `BSTR` encapsulated in `s1`.

**END Microsoft Specific**

## See Also

`_bstr_t` Class

# \_bstr\_t::operator !

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Checks if the encapsulated `BSTR` is a **NULL** string.

## Syntax

```
bool operator!( ) const throw( );
```

## Return Value

It returns **true** if yes, **false** if not.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# \_bstr\_t Relational Operators

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Compares two `_bstr_t` objects.

## Syntax

```
    bool operator!( ) const throw( );
bool operator==( 
    const _bstr_t& str
) const throw( );
bool operator!=(
    const _bstr_t& str
) const throw( );
bool operator<(
    const _bstr_t& str
) const throw( );
bool operator>(
    const _bstr_t& str
) const throw( );
bool operator<=( 
    const _bstr_t& str
) const throw( );
bool operator>=( 
    const _bstr_t& str
) const throw( );
```

## Remarks

These operators compare two `_bstr_t` objects lexicographically. The operators return **true** if the comparisons hold, otherwise return **false**.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# `_bstr_t::wchar_t, _bstr_t::char`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Returns the BSTR characters as a narrow or wide character array.

## Syntax

```
operator const wchar_t*( ) const throw( );
operator wchar_t*( ) const throw( );
operator const char*( ) const;
operator char*( ) const;
```

## Remarks

These operators can be used to extract the character data that is encapsulated by the `BSTR` object. Assigning a new value to the returned pointer does not modify the original BSTR data.

## END Microsoft Specific

## See Also

[\\_bstr\\_t Class](#)

# \_com\_error Class

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A `_com_error` object represents an exception condition detected by the error-handling wrapper functions in the header files generated from the type library or by one of the COM support classes. The `_com_error` class encapsulates the `HRESULT` error code and any associated `IErrorInfo` interface object.

## Construction

<code>_com_error</code>	Constructs a <code>_com_error</code> object.
-------------------------	--

## Operators

<code>operator =</code>	Assigns an existing <code>_com_error</code> object to another.
-------------------------	--

## Extractor Functions

<code>Error</code>	Retrieves the <code>HRESULT</code> passed to the constructor.
<code>ErrorInfo</code>	Retrieves the <code>IErrorInfo</code> object passed to the constructor.
<code>WCode</code>	Retrieves the 16-bit error code mapped into the encapsulated <code>HRESULT</code> .

## IErrorInfo Functions

<code>Description</code>	Calls <code>IErrorInfo::GetDescription</code> function.
<code>HelpContext</code>	Calls <code>IErrorInfo::GetHelpContext</code> function.
<code>HelpFile</code>	Calls <code>IErrorInfo::GetHelpFile</code> function
<code>Source</code>	Calls <code>IErrorInfo::GetSource</code> function.
<code>GUID</code>	Calls <code>IErrorInfo::GetGUID</code> function.

## Format Message Extractor

<code>ErrorMessage</code>	Retrieves the string message for <code>HRESULT</code> stored in the <code>_com_error</code> object.
---------------------------	---

## ExepInfo.wCode to HRESULT Mappers

<a href="#">HRESULTToWCode</a>	Maps 32-bit <code>HRESULT</code> to 16-bit <code>wCode</code> .
<a href="#">WCodeToHRESULT</a>	Maps 16-bit <code>wCode</code> to 32-bit <code>HRESULT</code> .

## END Microsoft Specific

## Requirements

**Header:** comdef.h

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Support Classes](#)

[IErrorInfo Interface](#)

# `_com_error` Member Functions

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_com_error` member functions, see [\\_com\\_error Class](#).

## See Also

[\\_com\\_error Class](#)

# \_com\_error::\_com\_error

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Constructs a `_com_error` object.

## Syntax

```
_com_error(
    HRESULT hr,
    IErrorInfo* perrinfo = NULL,
    bool fAddRef=false
) throw( );
_com_error(
    const _com_error& that
) throw( );
```

### Parameters

`hr`

`HRESULT` information.

`perrinfo`

`IErrorInfo` object.

### `bool fAddRef=false`

Causes the constructor to call AddRef on a non-null `IErrorInfo` interface. This provides for correct reference counting in the common case where ownership of the interface is passed into the `_com_error` object, such as:

```
throw _com_error(hr, perrinfo);
```

If you do not want your code to transfer ownership to the `_com_error` object, and the `AddRef` is required to offset the `Release` in the `_com_error` destructor, construct the object as follows:

```
_com_error err(hr, perrinfo, true);
```

`that`

An existing `_com_error` object.

## Remarks

The first constructor creates a new object given an `HRESULT` and optional `IErrorInfo` object. The second creates a copy of an existing `_com_error` object.

## END Microsoft Specific

## See Also

[\\_com\\_error Class](#)

# \_com\_error::Description

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls **IErrorInfo::GetDescription** function.

## Syntax

```
_bstr_t Description( ) const;
```

## Return Value

Returns the result of **IErrorInfo::GetDescription** for the **IErrorInfo** object recorded within the **\_com\_error** object. The resulting **BSTR** is encapsulated in a **\_bstr\_t** object. If no **IErrorInfo** is recorded, it returns an empty **\_bstr\_t**.

## Remarks

Calls the **IErrorInfo::GetDescription** function and retrieves **IErrorInfo** recorded within the **\_com\_error** object. Any failure while calling the **IErrorInfo::GetDescription** method is ignored.

## END Microsoft Specific

## See Also

[\\_com\\_error Class](#)

# \_com\_error::Error

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Retrieves the `HRESULT` passed to the constructor.

## Syntax

```
HRESULT Error( ) const throw( );
```

## Return Value

Raw `HRESULT` item passed into the constructor.

## Remarks

Retrieves the encapsulated `HRESULT` item in a `_com_error` object.

## END Microsoft Specific

## See Also

[\\_com\\_error Class](#)

# `_com_error::ErrorInfo`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Retrieves the **IErrorInfo** object passed to the constructor.

## Syntax

```
IErrorInfo * ErrorInfo( ) const throw( );
```

## Return Value

Raw **IErrorInfo** item passed into the constructor.

## Remarks

Retrieves the encapsulated **IErrorInfo** item in a `_com_error` object, or **NULL** if no **IErrorInfo** item is recorded. The caller must call **Release** on the returned object when finished using it.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_error::ErrorMessage

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Retrieves the string message for `HRESULT` stored in the `_com_error` object.

## Syntax

```
const TCHAR * ErrorMessage( ) const throw( );
```

## Return Value

Returns the string message for the `HRESULT` recorded within the `_com_error` object. If the `HRESULT` is a mapped 16-bit `wCode`, then a generic message "`IDispatch error #<wCode>`" is returned. If no message is found, then a generic message "`Unknown error #<HRESULT>`" is returned. The returned string is either a Unicode or multibyte string, depending on the state of the `_UNICODE` macro.

## Remarks

Retrieves the appropriate system message text for `HRESULT` recorded within the `_com_error` object. The system message text is obtained by calling the Win32 [FormatMessage](#) function. The string returned is allocated by the [FormatMessage](#) API, and it is released when the `_com_error` object is destroyed.

## END Microsoft Specific

## See Also

[\\_com\\_error Class](#)

# \_com\_error::GUID

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls **IErrorInfo::GetGUID** function.

## Syntax

```
GUID GUID( ) const throw( );
```

## Return Value

Returns the result of **IErrorInfo::GetGUID** for the **IErrorInfo** object recorded within the **\_com\_error** object. If no **IErrorInfo** object is recorded, it returns **GUID\_NULL**.

## Remarks

Any failure while calling the **IErrorInfo::GetGUID** method is ignored.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_error::HelpContext

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls **IErrorInfo::GetHelpContext** function.

## Syntax

```
DWORD HelpContext( ) const throw( );
```

## Return Value

Returns the result of **IErrorInfo::GetHelpContext** for the **IErrorInfo** object recorded within the `_com_error` object. If no **IErrorInfo** object is recorded, it returns a zero.

## Remarks

Any failure while calling the **IErrorInfo::GetHelpContext** method is ignored.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_error::HelpFile

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls **IErrorInfo::GetHelpFile** function.

## Syntax

```
_bstr_t HelpFile() const;
```

## Return Value

Returns the result of **IErrorInfo::GetHelpFile** for the **IErrorInfo** object recorded within the `_com_error` object. The resulting BSTR is encapsulated in a `_bstr_t` object. If no **IErrorInfo** is recorded, it returns an empty `_bstr_t`.

## Remarks

Any failure while calling the **IErrorInfo::GetHelpFile** method is ignored.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_error::HRESULTToWCode

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Maps 32-bit `HRESULT` to 16-bit `wCode`.

## Syntax

```
static WORD HRESULTToWCode(  
    HRESULT hr  
) throw( );
```

### Parameters

`hr`

The 32-bit `HRESULT` to be mapped to 16-bit `wCode`.

## Return Value

16-bit `wCode` mapped from the 32-bit `HRESULT`.

## Remarks

See [\\_com\\_error::WCode](#) for more information.

## END Microsoft Specific

## See Also

[\\_com\\_error::WCode](#)

[\\_com\\_error::WCodeToHRESULT](#)

[\\_com\\_error Class](#)

# \_com\_error::Source

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls **IErrorInfo::GetSource** function.

## Syntax

```
_bstr_t Source() const;
```

## Return Value

Returns the result of **IErrorInfo::GetSource** for the **IErrorInfo** object recorded within the **\_com\_error** object. The resulting BSTR is encapsulated in a **\_bstr\_t** object. If no **IErrorInfo** is recorded, it returns an empty **\_bstr\_t**.

## Remarks

Any failure while calling the **IErrorInfo::GetSource** method is ignored.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_error::WCode

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Retrieves the 16-bit error code mapped into the encapsulated `HRESULT`.

## Syntax

```
WORD WCode( ) const throw();
```

## Return Value

If the `HRESULT` is within the range 0x80040200 to 0x8004FFFF, the **WCode** method returns the `HRESULT` minus 0x80040200; otherwise, it returns zero.

## Remarks

The **WCode** method is used to undo a mapping that happens in the COM support code. The wrapper for a **dispinterface** property or method calls a support routine that packages the arguments and calls **IDispatch::Invoke**. Upon return, if a failure `HRESULT` of `DISP_E_EXCEPTION` is returned, the error information is retrieved from the **EXCEPINFO** structure passed to **IDispatch::Invoke**. The error code can either be a 16-bit value stored in the `wCode` member of the **EXCEPINFO** structure or a full 32-bit value in the `scode` member of the **EXCEPINFO** structure. If a 16-bit `wCode` is returned, it must first be mapped to a 32-bit failure `HRESULT`.

## END Microsoft Specific

## See Also

[\\_com\\_error::HRESULTToWCode](#)  
[\\_com\\_error::WCodeToHRESULT](#)  
[\\_com\\_error Class](#)

# \_com\_error::WCodeToHRESULT

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Maps 16-bit `wCode` to 32-bit `HRESULT`.

## Syntax

```
static HRESULT WCodeToHRESULT(
WORD wCode
) throw( );
```

### Parameters

`wCode`

The 16-bit `wCode` to be mapped to 32-bit `HRESULT`.

## Return Value

32-bit `HRESULT` mapped from the 16-bit `wCode`.

## Remarks

See the [WCode](#) member function.

## END Microsoft Specific

## See Also

[\\_com\\_error::WCode](#)

[\\_com\\_error::HRESULTToWCode](#)

[\\_com\\_error Class](#)

# \_com\_error Operators

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_com_error` operators, see [\\_com\\_error Class](#).

## See Also

[\\_com\\_error Class](#)

# `_com_error::operator =`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Assigns an existing `_com_error` object to another.

## Syntax

```
_com_error& operator = (
    const _com_error& that
) throw ( );
```

### Parameters

`that`

A `_com_error` object.

**END Microsoft Specific**

## See Also

[\\_com\\_error Class](#)

# \_com\_ptr\_t Class

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A `_com_ptr_t` object encapsulates a COM interface pointer and is called a "smart" pointer. This template class manages resource allocation and deallocation through function calls to the **IUnknown** member functions: `QueryInterface`, `AddRef`, and `Release`.

A smart pointer is usually referenced by the `typedef` definition provided by the `_COM_SMARTPTR_TYPEDEF` macro. This macro takes an interface name and the IID and declares a specialization of `_com_ptr_t` with the name of the interface plus a suffix of `Ptr`. For example:

```
_COM_SMARTPTR_TYPEDEF(IMyInterface, __uuidof(IMyInterface));
```

declares the `_com_ptr_t` specialization **IMyInterfacePtr**.

A set of [function templates](#), not members of this template class, support comparisons with a smart pointer on the right side of the comparison operator.

## Construction

<code>_com_ptr_t</code>	Constructs a <code>_com_ptr_t</code> object.
-------------------------	--

## Low-Level Operations

<a href="#">AddRef</a>	Calls the <code>AddRef</code> member function of <b>IUnknown</b> on the encapsulated interface pointer.
<a href="#">Attach</a>	Encapsulates a raw interface pointer of this smart pointer's type.
<a href="#">CreateInstance</a>	Creates a new instance of an object given a <b>CLSID</b> or <b>ProgID</b> .
<a href="#">Detach</a>	Extracts and returns the encapsulated interface pointer.
<a href="#">GetActiveObject</a>	Attaches to an existing instance of an object given a <b>CLSID</b> or <b>ProgID</b> .
<a href="#">GetInterfacePtr</a>	Returns the encapsulated interface pointer.
<a href="#">QueryInterface</a>	Calls the <code>QueryInterface</code> member function of <b>IUnknown</b> on the encapsulated interface pointer.
<a href="#">Release</a>	Calls the <code>Release</code> member function of <b>IUnknown</b> on the encapsulated interface pointer.

## Operators

<code>operator =</code>	Assigns a new value to an existing <code>_com_ptr_t</code> object.
<code>operators ==, !=, &lt;, &gt;, &lt;=, &gt;=</code>	Compare the smart pointer object to another smart pointer, raw interface pointer, or <b>NULL</b> .
<code>Extractors</code>	Extract the encapsulated COM interface pointer.

## END Microsoft Specific

## Requirements

**Header:** comip.h

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zc:wchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Support Classes](#)

# `_com_ptr_t` Member Functions

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_com_ptr_t` member functions, see [\\_com\\_ptr\\_t Class](#).

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::\_com\_ptr\_t

12/21/2017 • 2 min to read • [Edit Online](#)

## Microsoft Specific

Constructs a `_com_ptr_t` object.

## Syntax

```
// Default constructor.  
// Constructs a NULL smart pointer.  
_com_ptr_t() throw();  
  
// Constructs a NULL smart pointer. The NULL argument must be zero.  
_com_ptr_t(  
    int null  
);  
  
// Constructs a smart pointer as a copy of another instance of the  
// same smart pointer. AddRef is called to increment the reference  
// count for the encapsulated interface pointer.  
_com_ptr_t(  
    const _com_ptr_t& cp  
) throw();  
  
// Move constructor (Visual Studio 2015 Update 3 and later)  
_com_ptr_t(_com_ptr_t&& cp) throw();  
  
// Constructs a smart pointer from a raw interface pointer of this  
// smart pointer's type. If fAddRef is true, AddRef is called  
// to increment the reference count for the encapsulated  
// interface pointer. If fAddRef is false, this constructor  
// takes ownership of the raw interface pointer without calling AddRef.  
_com_ptr_t(  
    Interface* pInterface,  
    bool fAddRef  
) throw();  
  
// Construct pointer for a _variant_t object.  
// Constructs a smart pointer from a _variant_t object. The  
// encapsulated VARIANT must be of type VT_DISPATCH or VT_UNKNOWN, or  
// it can be converted into one of these two types. If QueryInterface  
// fails with an E_NOINTERFACE error, a NULL smart pointer is  
// constructed.  
_com_ptr_t(  
    const _variant_t& varSrc  
);  
  
// Constructs a smart pointer given the CLSID of a coclass. This  
// function calls CoCreateInstance, by the member function  
// CreateInstance, to create a new COM object and then queries for  
// this smart pointer's interface type. If QueryInterface fails with  
// an E_NOINTERFACE error, a NULL smart pointer is constructed.  
explicit _com_ptr_t(  
    const CLSID& clsid,  
    IUnknown* pOuter = NULL,  
    DWORD dwClсхContext = CLSCTX_ALL  
);  
  
// Calls CoCreateClass with provided CLSID retrieved from string.  
explicit _com_ptr_t(
```

```

LPCWSTR str,
IUnknown* pOuter = NULL,
DWORD dwClsContext = CLSCTX_ALL
);

// Constructs a smart pointer given a multibyte character string that
// holds either a CLSID (starting with "{") or a ProgID. This function
// calls CoCreateInstance, by the member function CreateInstance, to
// create a new COM object and then queries for this smart pointer's
// interface type. If QueryInterface fails with an E_NOINTERFACE error,
// a NULL smart pointer is constructed.
explicit _com_ptr_t(
    LPCSTR str,
    IUnknown* pOuter = NULL,
    DWORD dwClsContext = CLSCTX_ALL
);

// Saves the interface.
template<>
_com_ptr_t(
    Interface* pInterface
) throw();

// Make sure correct ctor is called
template<>
_com_ptr_t(
    LPSTR str
);

// Make sure correct ctor is called
template<>
_com_ptr_t(
    LPWSTR str
);

// Constructs a smart pointer from a different smart pointer type or
// from a different raw interface pointer. QueryInterface is called to
// find an interface pointer of this smart pointer's type. If
// QueryInterface fails with an E_NOINTERFACE error, a NULL smart
// pointer is constructed.
template<typename _OtherIID>
_com_ptr_t(
    const _com_ptr_t<_OtherIID>& p
);

// Constructs a smart-pointer from any IUnknown-based interface pointer.
template<typename _InterfaceType>
_com_ptr_t(
    _InterfaceType* p
);

// Disable conversion using _com_ptr_t* specialization of
// template<typename _InterfaceType> _com_ptr_t(_InterfaceType* p)
template<>
explicit _com_ptr_t(
    _com_ptr_t* p
);

```

## Parameters

`pInterface`

A raw interface pointer.

`fAddRef`

If **true**, `AddRef` is called to increment the reference count of the encapsulated interface pointer.

A `_com_ptr_t` object.

`p`

A raw interface pointer, its type being different from the smart pointer type of this `_com_ptr_t` object.

`varSrc`

A `_variant_t` object.

`clsid`

The **CLSID** of a coclass.

`dwClsContext`

Context for running executable code.

`lpcStr`

A multibyte string that holds either a **CLSID** (starting with "{") or a **ProgID**.

`pOuter`

The outer unknown for [aggregation](#).

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::AddRef

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls the `AddRef` member function of **IUnknown** on the encapsulated interface pointer.

## Syntax

```
void AddRef( );
```

## Remarks

Calls `IUnknown::AddRef` on the encapsulated interface pointer, raising an `E_POINTER` error if the pointer is **NULL**.

## END Microsoft Specific

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::Attach

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Encapsulates a raw interface pointer of this smart pointer's type.

## Syntax

```
void Attach(  
    Interface* pInterface  
) throw( );  
  
void Attach(  
    Interface* pInterface,  
    bool fAddRef  
) throw( );
```

### Parameters

`pInterface`

A raw interface pointer.

`fAddRef`

If it is **true**, then `AddRef` is called. If it is **false**, the `_com_ptr_t` object takes ownership of the raw interface pointer without calling `AddRef`.

## Remarks

- **Attach( `pInterface` )** `AddRef` is not called. The ownership of the interface is passed to this `_com_ptr_t` object. **Release** is called to decrement the reference count for the previously encapsulated pointer.
- **Attach( `pInterface` , `fAddRef` )** If `fAddRef` is **true**, `AddRef` is called to increment the reference count for the encapsulated interface pointer. If `fAddRef` is **false**, this `_com_ptr_t` object takes ownership of the raw interface pointer without calling `AddRef`. **Release** is called to decrement the reference count for the previously encapsulated pointer.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::CreateInstance

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Creates a new instance of an object given a **CLSID** or **ProgID**.

## Syntax

```
HRESULT CreateInstance(
    const CLSID& rclsid,
    IUnknown* pOuter=NULL,
    DWORD dwClsContext = CLSCTX_ALL
) throw( );
HRESULT CreateInstance(
    LPCWSTR clsidString,
    IUnknown* pOuter=NULL,
    DWORD dwClsContext = CLSCTX_ALL
) throw( );
HRESULT CreateInstance(
    LPCSTR clsidStringA,
    IUnknown* pOuter=NULL,
    DWORD dwClsContext = CLSCTX_ALL
) throw( );
```

### Parameters

`rclsid`

The **CLSID** of an object.

`clsidString`

A Unicode string that holds either a **CLSID** (starting with "{") or a **ProgID**.

`clsidStringA`

A multibyte string, using the ANSI code page, that holds either a **CLSID** (starting with "{") or a **ProgID**.

`dwClsContext`

Context for running executable code.

`pOuter`

The outer unknown for [aggregation](#).

## Remarks

These member functions call `CoCreateInstance` to create a new COM object and then queries for this smart pointer's interface type. The resulting pointer is then encapsulated within this `_com_ptr_t` object. **Release** is called to decrement the reference count for the previously encapsulated pointer. This routine returns the `HRESULT` to indicate success or failure.

- **CreateInstance(** `rclsid` **,** `dwClsContext` **)** Creates a new running instance of an object given a **CLSID**.
- **CreateInstance(** `clsidString` **,** `dwClsContext` **)** Creates a new running instance of an object given a Unicode string that holds either a **CLSID** (starting with "{") or a **ProgID**.
- **CreateInstance(** `clsidStringA` **,** `dwClsContext` **)** Creates a new running instance of an object given a

multibyte character string that holds either a **CLSID** (starting with "{") or a **ProgID**. Calls [MultiByteToWideChar](#), which assumes that the string is in the ANSI code page rather than an OEM code page.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::Detach

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Extracts and returns the encapsulated interface pointer.

## Syntax

```
Interface* Detach( ) throw( );
```

## Remarks

Extracts and returns the encapsulated interface pointer, and then clears the encapsulated pointer storage to **NULL**. This removes the interface pointer from encapsulation. It is up to you to call **Release** on the returned interface pointer.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::GetActiveObject

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Attaches to an existing instance of an object given a **CLSID** or **ProgID**.

## Syntax

```
HRESULT GetActiveObject(
    const CLSID& rclsid
) throw( );
HRESULT GetActiveObject(
    LPCWSTR clsidString
) throw( );
HRESULT GetActiveObject(
    LPCSTR clsidStringA
) throw( );
```

### Parameters

`rclsid`

The **CLSID** of an object.

`clsidString`

A Unicode string that holds either a **CLSID** (starting with "{") or a **ProgID**.

`clsidStringA`

A multibyte string, using the ANSI code page, that holds either a **CLSID** (starting with "{") or a **ProgID**.

## Remarks

These member functions call `GetActiveObject` to retrieve a pointer to a running object that has been registered with OLE and then queries for this smart pointer's interface type. The resulting pointer is then encapsulated within this `_com_ptr_t` object. **Release** is called to decrement the reference count for the previously encapsulated pointer. This routine returns the `HRESULT` to indicate success or failure.

- **GetActiveObject( `rclsid` )** Attaches to an existing instance of an object given a **CLSID**.
- **GetActiveObject( `clsidString` )** Attaches to an existing instance of an object given a Unicode string that holds either a **CLSID** (starting with "{") or a **ProgID**.
- **GetActiveObject( `clsidStringA` )** Attaches to an existing instance of an object given a multibyte character string that holds either a **CLSID** (starting with "{") or a **ProgID**. Calls `MultiByteToWideChar`, which assumes that the string is in the ANSI code page rather than an OEM code page.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::GetInterfacePtr

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Returns the encapsulated interface pointer.

## Syntax

```
Interface* GetInterfacePtr( ) const throw( );
Interface*& GetInterfacePtr() throw();
```

## Remarks

Returns the encapsulated interface pointer, which may be **NULL**.

## END Microsoft Specific

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::QueryInterface

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls the `QueryInterface` member function of **IUnknown** on the encapsulated interface pointer.

## Syntax

```
template<typename _InterfaceType> HRESULT QueryInterface (
    const IID& iid,
    _InterfaceType*& p
) throw( );
template<typename _InterfaceType> HRESULT QueryInterface (
    const IID& iid,
    _InterfaceType** p
) throw( );
```

### Parameters

`iid`

**IID** of an interface pointer.

`p`

Raw interface pointer.

## Remarks

Calls **IUnknown::QueryInterface** on the encapsulated interface pointer with the specified **IID** and returns the resulting raw interface pointer in `p`. This routine returns the `HRESULT` to indicate success or failure.

## END Microsoft Specific

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t::Release

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Calls the **Release** member function of **IUnknown** on the encapsulated interface pointer.

## Syntax

```
void Release( );
```

## Remarks

Calls `IUnknown::Release` on the encapsulated interface pointer, raising an `E_POINTER` error if this interface pointer is **NULL**.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# `_com_ptr_t` Operators

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_com_ptr_t` operators, see [\\_com\\_ptr\\_t Class](#).

## See Also

[\\_com\\_ptr\\_t Class](#)

# `_com_ptr_t::operator =`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Assigns a new value to an existing `_com_ptr_t` object.

## Syntax

```
template<typename _OtherIID>
_com_ptr_t& operator=(
    const _com_ptr_t<_OtherIID>& p
);

// Sets a smart pointer to be a different smart pointer of a different
// type or a different raw interface pointer. QueryInterface is called
// to find an interface pointer of this smart pointer's type, and
// Release is called to decrement the reference count for the previously
// encapsulated pointer. If QueryInterface fails with an E_NOINTERFACE,
// a NULL smart pointer results.
template<typename _InterfaceType>
_com_ptr_t& operator=(
    _InterfaceType* p
);

// Encapsulates a raw interface pointer of this smart pointer's type.
// AddRef is called to increment the reference count for the encapsulated
// interface pointer, and Release is called to decrement the reference
// count for the previously encapsulated pointer.
template<> _com_ptr_t&
operator=(
    Interface* pInterface
) throw();

// Sets a smart pointer to be a copy of another instance of the same
// smart pointer of the same type. AddRef is called to increment the
// reference count for the encapsulated interface pointer, and Release
// is called to decrement the reference count for the previously
// encapsulated pointer.
_com_ptr_t& operator=(
    const _com_ptr_t& cp
) throw();

// Sets a smart pointer to NULL. The NULL argument must be a zero.
_com_ptr_t& operator=(
    int null
);
// Sets a smart pointer to be a _variant_t object. The encapsulated
// VARIANT must be of type VT_DISPATCH or VT_UNKNOWN, or it can be
// converted to one of these two types. If QueryInterface fails with an
// E_NOINTERFACE error, a NULL smart pointer results.
_com_ptr_t& operator=(
    const _variant_t& varSrc
);
```

## Remarks

Assigns an interface pointer to this `_com_ptr_t` object.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t Relational Operators

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Compare the smart pointer object to another smart pointer, raw interface pointer, or **NULL**.

## Syntax

```
template<typename _OtherIID>
bool operator==(const _com_ptr_t<_OtherIID>& p);

template<typename _OtherIID>
bool operator==(const _com_ptr_t<_OtherIID>& p);

template<typename _InterfaceType>
bool operator==(const _InterfaceType* p);

template<>
bool operator==(const Interface* p);

template<>
bool operator==(const _com_ptr_t& p) throw();

template<>
bool operator==(const _com_ptr_t& p) throw();

bool operator==(int null);
```

```
template<typename _OtherIID>
bool operator!=(
    const _com_ptr_t<_OtherIID>& p
);

template<typename _OtherIID>
bool operator!=(
    _com_ptr_t<_OtherIID>& p
);

template<typename _InterfaceType>
bool operator!=(
    _InterfaceType* p
);

bool operator!=(  
    int null  
);
```

```
template<typename _OtherIID>
bool operator<(
    const _com_ptr_t<_OtherIID>& p
);

template<typename _OtherIID>
bool operator<(
    _com_ptr_t<_OtherIID>& p
);

template<typename _InterfaceType>
bool operator<(
    _InterfaceType* p
);
```

```
template<typename _OtherIID>
bool operator>(
    const _com_ptr_t<_OtherIID>& p
);

template<typename _OtherIID>
bool operator>(_com_ptr_t<  
    _OtherIID>& p
);

template<typename _InterfaceType>
bool operator>(  
    _InterfaceType* p
);
```

```
template<typename _OtherIID>
bool operator<=(
    const _com_ptr_t<_OtherIID>& p
);

template<typename _OtherIID>
bool operator<=(
    _com_ptr_t<_OtherIID>& p
);

template<typename _InterfaceType>
bool operator<=(
    _InterfaceType* p
);
```

```
template<typename _OtherIID>
bool operator>=(
    const _com_ptr_t<_OtherIID>& p
);

template<typename _OtherIID>
bool operator>=(
    _com_ptr_t<_OtherIID>& p
);

template<typename _InterfaceType>
bool operator>=(
    _InterfaceType* p
);
```

## Remarks

Compares a smart pointer object to another smart pointer, raw interface pointer, or **NULL**. Except for the **NULL** pointer tests, these operators first query both pointers for **IUnknown**, and compare the results.

### END Microsoft Specific

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_com\_ptr\_t Extractors

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Extract the encapsulated COM interface pointer.

## Syntax

```
operator Interface*( ) const throw( );
operator Interface&( ) const;
Interface& operator*( ) const;
Interface* operator->( ) const;
Interface** operator&( ) throw( );
operator bool( ) const throw( );
```

## Remarks

- **operator Interface\\*** Returns the encapsulated interface pointer, which may be **NULL**.
- **operator Interface&** Returns a reference to the encapsulated interface pointer, and issues an error if the pointer is **NULL**.
- **operator\\*** Allows a smart pointer object to act as though it were the actual encapsulated interface when dereferenced.
- **operator->** Allows a smart pointer object to act as though it were the actual encapsulated interface when dereferenced.
- **operator&** Releases any encapsulated interface pointer, replacing it with **NULL**, and returns the address of the encapsulated pointer. This allows the smart pointer to be passed by address to a function that has an **out** parameter through which it returns an interface pointer.
- **operator bool** Allows a smart pointer object to be used in a conditional expression. This operator returns **true** if the pointer is not **NULL**.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# Relational Function Templates

12/21/2017 • 1 min to read • [Edit Online](#)

**Microsoft Specific**

Syntax

```

        template<typename _InterfaceType> bool operator==(

int NULL,
_com_ptr_t<_InterfaceType>& p
);

template<typename _Interface,
typename _InterfacePtr> bool operator==(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

template<typename _Interface> bool operator!=(

int NULL,
_com_ptr_t<_Interface>& p

);

template<typename _Interface,
typename _InterfacePtr> bool operator!=(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

template<typename _Interface> bool operator<(

int NULL,
_com_ptr_t<_Interface>& p

);

template<typename _Interface,
typename _InterfacePtr> bool operator<(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

template<typename _Interface> bool operator>(

int NULL,
_com_ptr_t<_Interface>& p

);

template<typename _Interface,
typename _InterfacePtr> bool operator>(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

template<typename _Interface> bool operator<=(

int NULL,
_com_ptr_t<_Interface>& p

);

template<typename _Interface,
typename _InterfacePtr> bool operator<=(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

template<typename _Interface> bool operator>=(

int NULL,
_com_ptr_t<_Interface>& p

);

template<typename _Interface,
typename _InterfacePtr> bool operator>=(

_Interface* i,
_com_ptr_t<_InterfacePtr>& p

);

```

#### Parameters

*i*

A raw interface pointer.

p

A smart pointer.

#### Remarks

These function templates allow comparison with a smart pointer on the right side of the comparison operator.

These are not member functions of `_com_ptr_t`.

**END Microsoft Specific**

## See Also

[\\_com\\_ptr\\_t Class](#)

# \_variant\_t Class

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

A `_variant_t` object encapsulates the `VARIANT` data type. The class manages resource allocation and deallocation and makes function calls to **VariantInit** and **VariantClear** as appropriate.

## Construction

<code>_variant_t</code>	Constructs a <code>_variant_t</code> object.
-------------------------	--

## Operations

<code>Attach</code>	Attaches a <b>VARIANT</b> object into the <code>_variant_t</code> object.
<code>Clear</code>	Clears the encapsulated <b>VARIANT</b> object.
<code>ChangeType</code>	Changes the type of the <code>_variant_t</code> object to the indicated <b>VARTYPE</b> .
<code>Detach</code>	Detaches the encapsulated <b>VARIANT</b> object from this <code>_variant_t</code> object.
<code>SetString</code>	Assigns a string to this <code>_variant_t</code> object.

## Operators

<code>Operator =</code>	Assigns a new value to an existing <code>_variant_t</code> object.
<code>operator ==, !=</code>	Compare two <code>_variant_t</code> objects for equality or inequality.
<code>Extractors</code>	Extract data from the encapsulated <b>VARIANT</b> object.

## END Microsoft Specific

## Requirements

**Header:** comutil.h

**Lib:** comsuppw.lib or comsuppwd.lib (see [/Zcwchar\\_t \(wchar\\_t Is Native Type\)](#) for more information)

## See Also

[Compiler COM Support Classes](#)

# \_variant\_t Member Functions

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_variant_t` member functions, see [\\_variant\\_t Class](#).

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::variant\_t

12/21/2017 • 4 min to read • [Edit Online](#)

## Microsoft Specific

Constructs a `_variant_t` object.

## Syntax

```
_variant_t( ) throw( );

_variant_t(
    const VARIANT& varSrc
);

_variant_t(
    const VARIANT* pVarSrc
);

_variant_t(
    const _variant_t& var_t_src
);

_variant_t(
    VARIANT& varSrc,
    bool fCopy
);

_variant_t(
    short sSrc,
    VARTYPE vtSrc = VT_I2
);

_variant_t(
    long lSrc,
    VARTYPE vtSrc = VT_I4
);

_variant_t(
    float fltSrc
) throw( );

_variant_t(
    double dblSrc,
    VARTYPE vtSrc = VT_R8
);

_variant_t(
    const CY& cySrc
) throw( );

_variant_t(
    const _bstr_t& bstrSrc
);

_variant_t(
    const wchar_t *wstrSrc
);

_variant_t(
```

```

    const char* strSrc
);

_variant_t(
    IDispatch* pDispSrc,
    bool fAddRef = true
) throw( );

_variant_t(
    bool bSrc
) throw( );

_variant_t(
    IUnknown* pIUnknownSrc,
    bool fAddRef = true
) throw( );

_variant_t(
    const DECIMAL& decSrc
) throw( );

_variant_t(
    BYTE bSrc
) throw( );

variant_t(
    char cSrc
) throw();

_variant_t(
    unsigned short usSrc
) throw();

_variant_t(
    unsigned long ulSrc
) throw();

_variant_t(
    int iSrc
) throw();

_variant_t(
    unsigned int uiSrc
) throw();

_variant_t(
    __int64 i8Src
) throw();

_variant_t(
    unsigned __int64 ui8Src
) throw();

```

#### Parameters

*varSrc*

A **VARIANT** object to be copied into the new `_variant_t` object.

*pVarSrc*

Pointer to a **VARIANT** object to be copied into the new `_variant_t` object.

*var\_t\_Src*

A `_variant_t` object to be copied into the new `_variant_t` object.

*fCopy*

If false, the supplied **VARIANT** object is attached to the new `_variant_t` object without making a new copy by

## **VariantCopy**

*lSrc, sSrc*

An integer value to be copied into the new `_variant_t` object.

`vtSrc`

The **VARTYPE** for the new `_variant_t` object.

*fltSrc, dblSrc*

A numerical value to be copied into the new `_variant_t` object.

`cySrc`

A **CY** object to be copied into the new `_variant_t` object.

`bstrSrc`

A `_bstr_t` object to be copied into the new `_variant_t` object.

*strSrc, wstrSrc*

A string to be copied into the new `_variant_t` object.

`bSrc`

A `bool` value to be copied into the new `_variant_t` object.

`pIUnknownSrc`

COM interface pointer to a **VT\_UNKNOWN** object to be encapsulated into the new `_variant_t` object.

`pDispSrc`

COM interface pointer to a **VT\_DISPATCH** object to be encapsulated into the new `_variant_t` object.

`decSrc`

A **DECIMAL** value to be copied into the new `_variant_t` object.

`bSrc`

A **BYTE** value to be copied into the new `_variant_t` object.

`cSrc`

A `char` value to be copied into the new `_variant_t` object.

*usSrc*

A **unsigned short** value to be copied into the new `_variant_t` object.

*ulSrc*

A `unsigned long` value to be copied into the new `_variant_t` object.

`iSrc`

An `int` value to be copied into the new `_variant_t` object.

*uiSrc*

An `unsigned int` value to be copied into the new `_variant_t` object.

*i8Src*

An `_int64` value to be copied into the new `_variant_t` object.

*ui8Src*

An `unsigned _int64` value to be copied into the new `_variant_t` object.

## Remarks

- `_variant_t()` Constructs an empty `_variant_t` object, `VT_EMPTY`.

- `_variant_t(VARIANT& varSrc)` Constructs a `_variant_t` object from a copy of the **VARIANT** object. The variant type is retained.
- `_variant_t(VARIANT\* pVarSrc)` Constructs a `_variant_t` object from a copy of the **VARIANT** object. The variant type is retained.
- `_variant_t(_variant_t& var_t_Src)` Constructs a `_variant_t` object from another `_variant_t` object. The variant type is retained.
- `_variant_t(VARIANT& varSrc, bool fCopy)` Constructs a `_variant_t` object from an existing **VARIANT** object. If `fCopy` is **false**, the **VARIANT** object is attached to the new object without making a copy.
- `_variant_t(short sSrc, VARTYPE vtSrc = VT_I2)` Constructs a `_variant_t` object of type `VT_I2` or `VT_BOOL` from a **short** integer value. Any other **VARTYPE** results in an `E_INVALIDARG` error.
- `_variant_t(long lSrc, VARTYPE vtSrc = VT_I4)` Constructs a `_variant_t` object of type `VT_I4`, `VT_BOOL`, or `VT_ERROR` from a **long** integer value. Any other **VARTYPE** results in an `E_INVALIDARG` error.
- `_variant_t(float fltSrc)` Constructs a `_variant_t` object of type `VT_R4` from a **float** numerical value.
- `_variant_t(double dblSrc, VARTYPE vtSrc = VT_R8)` Constructs a `_variant_t` object of type `VT_R8` or `VT_DATE` from a **double** numerical value. Any other **VARTYPE** results in an `E_INVALIDARG` error.
- `_variant_t(CY& cySrc)` Constructs a `_variant_t` object of type `VT_CY` from a **CY** object.
- `_variant_t(_bstr_t& bstrSrc)` Constructs a `_variant_t` object of type `VT_BSTR` from a `_bstr_t` object. A new `BSTR` is allocated.
- `_variant_t(wchar_t\* wstrSrc)` Constructs a `_variant_t` object of type `VT_BSTR` from a Unicode string. A new `BSTR` is allocated.
- `_variant_t(char\* strSrc)` Constructs a `_variant_t` object of type `VT_BSTR` from a string. A new `BSTR` is allocated.
- `_variant_t(bool bSrc)` Constructs a `_variant_t` object of type `VT_BOOL` from a `bool` value.
- `_variant_t(IUnknown\* pIUnknownSrc, bool fAddRef = true)` Constructs a `_variant_t` object of type `VT_UNKNOWN` from a COM interface pointer. If `fAddRef` is **true**, then `AddRef` is called on the supplied interface pointer to match the call to **Release** that will occur when the `_variant_t` object is destroyed. It is up to you to call **Release** on the supplied interface pointer. If `fAddRef` is **false**, this constructor takes ownership of the supplied interface pointer; do not call **Release** on the supplied interface pointer.
- `_variant_t(IDispatch\* pDispSrc, bool fAddRef = true)` Constructs a `_variant_t` object of type `VT_DISPATCH` from a COM interface pointer. If `fAddRef` is **true**, then `AddRef` is called on the supplied interface pointer to match the call to **Release** that will occur when the `_variant_t` object is destroyed. It is up to you to call **Release** on the supplied interface pointer. If `fAddRef` is **false**, this constructor takes ownership of the supplied interface pointer; do not call **Release** on the supplied interface pointer.
- `_variant_t(DECIMAL& decSrc)` Constructs a `_variant_t` object of type `VT_DECIMAL` from a **DECIMAL** value.
- `_variant_t(BYTE bSrc)` Constructs a `_variant_t` object of type `VT_UI1` from a **BYTE** value.

#### END Microsoft Specific

## See Also

[\\_variant\\_t Class](#)

# `_variant_t::Attach`

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Attaches a **VARIANT** object into the `_variant_t` object.

## Syntax

```
void Attach(  
    VARIANT& varSrc  
>;
```

### Parameters

*varSrc*

A **VARIANT** object to be attached to this `_variant_t` object.

## Remarks

Takes ownership of the **VARIANT** by encapsulating it. This member function releases any existing encapsulated **VARIANT**, then copies the supplied **VARIANT**, and sets its **VARTYPE** to `VT_EMPTY` to make sure its resources can only be released by the `_variant_t` destructor.

## END Microsoft Specific

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::Clear

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Clears the encapsulated **VARIANT** object.

## Syntax

```
void Clear( );
```

## Remarks

Calls **VariantClear** on the encapsulated **VARIANT** object.

**END Microsoft Specific**

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::ChangeType

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Changes the type of the `_variant_t` object to the indicated **VARTYPE**.

## Syntax

```
void ChangeType(  
    VARTYPE vartype,  
    const _variant_t* pSrc = NULL  
) ;
```

### Parameters

`vartype`

The **VARTYPE** for this `_variant_t` object.

`pSrc`

A pointer to the `_variant_t` object to be converted. If this value is **NULL**, conversion is done in place.

## Remarks

This member function converts a `_variant_t` object into the indicated **VARTYPE**. If `pSrc` is **NULL**, the conversion is done in place, otherwise this `_variant_t` object is copied from `pSrc` and then converted.

## END Microsoft Specific

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::Detach

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Detaches the encapsulated **VARIANT** object from this `_variant_t` object.

## Syntax

```
VARIANT Detach( );
```

## Return Value

The encapsulated **VARIANT**.

## Remarks

Extracts and returns the encapsulated **VARIANT**, then clears this `_variant_t` object without destroying it. This member function removes the **VARIANT** from encapsulation and sets the **VARTYPE** of this `_variant_t` object to `VT_EMPTY`. It is up to you to release the returned **VARIANT** by calling the [VariantClear](#) function.

## END Microsoft Specific

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::SetString

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Assigns a string to this `_variant_t` object.

## Syntax

```
void SetString(  
    const char* pSrc  
) ;
```

### Parameters

`pSrc`

Pointer to the character string.

## Remarks

Converts an ANSI character string to a Unicode `BSTR` string and assigns it to this `_variant_t` object.

## END Microsoft Specific

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t Operators

12/21/2017 • 1 min to read • [Edit Online](#)

For information about the `_variant_t` operators, see [\\_variant\\_t Class](#).

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t::operator =

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

### Syntax

```
_variant_t& operator=(  
    const VARIANT& varSrc  
)  
  
_variant_t& operator=(  
    const VARIANT* pVarSrc  
)  
  
_variant_t& operator=(  
    const _variant_t& var_t_src  
)  
  
_variant_t& operator=(  
    short sSrc  
)  
  
_variant_t& operator=(  
    long lSrc  
)  
  
_variant_t& operator=(  
    float fltSrc  
)  
  
_variant_t& operator=(  
    double dblSrc  
)  
  
_variant_t& operator=(  
    const CY& cySrc  
)  
  
_variant_t& operator=(  
    const _bstr_t& bstrSrc  
)  
  
_variant_t& operator=(  
    const wchar_t* wstrSrc  
)  
  
_variant_t& operator=(  
    const char* strSrc  
)  
  
_variant_t& operator=(  
    IDispatch* pDispSrc  
)  
  
_variant_t& operator=(  
    bool bSrc  
)  
  
_variant_t& operator=(
```

```

    IUnknown* pSrc
);

_variant_t& operator=(
    const DECIMAL& decSrc
);

_variant_t& operator=(
    BYTE bSrc
);

_variant_t& operator=(
    char cSrc
);

_variant_t& operator=(
    unsigned short usSrc
);

_variant_t& operator=(
    unsigned long ulSrc
);

_variant_t& operator=(
    int iSrc
);

_variant_t& operator=(
    unsigned int uiSrc
);

_variant_t& operator=(
    __int64 i8Src
);

_variant_t& operator=(
    unsigned __int64 ui8Src
);

```

## Remarks

The operator assigns a new value to the `_variant_t` object:

- **`operator=( varSrc )`** Assigns an existing **VARIANT** to a `_variant_t` object.
- **`operator=( pVarSrc )`** Assigns an existing **VARIANT** to a `_variant_t` object.
- **`operator=( var_t_Src )`** Assigns an existing `_variant_t` object to a `_variant_t` object.
- **`operator=( sSrc )`** Assigns a **short** integer value to a `_variant_t` object.
- **`operator=( lSrc )`** Assigns a **long** integer value to a `_variant_t` object.
- **`operator=( fltSrc )`** Assigns a **float** numerical value to a `_variant_t` object.
- **`operator=( dblSrc )`** Assigns a **double** numerical value to a `_variant_t` object.
- **`operator=( cySrc )`** Assigns a **CY** object to a `_variant_t` object.
- **`operator=( bstrSrc )`** Assigns a **BSTR** object to a `_variant_t` object.
- **`operator=( wstrSrc )`** Assigns a Unicode string to a `_variant_t` object.
- **`operator=( strSrc )`** Assigns a multibyte string to a `_variant_t` object.

- **operator=(** `bSrc` **)** Assigns a `bool` value to a `_variant_t` object.
- **operator=(** `pDispSrc` **)** Assigns a `VT_DISPATCH` object to a `_variant_t` object.
- **operator=(** `pIUnknownSrc` **)** Assigns a `VT_UNKNOWN` object to a `_variant_t` object.
- **operator=(** `decSrc` **)** Assigns a `DECIMAL` value to a `_variant_t` object.
- **operator=(** `bSrc` **)** Assigns a `BYTE` value to a `_variant_t` object.

**END Microsoft Specific**

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t Relational Operators

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Compare two `_variant_t` objects for equality or inequality.

## Syntax

```
bool operator==(  
    const VARIANT& varSrc  
) const;  
bool operator==(  
    const VARIANT* pSrc  
) const;  
bool operator!=(  
    const VARIANT& varSrc  
) const;  
bool operator!=(  
    const VARIANT* pSrc  
) const;
```

### Parameters

`varSrc`

A **VARIANT** to be compared with the `_variant_t` object.

`pSrc`

Pointer to the **VARIANT** to be compared with the `_variant_t` object.

## Return Value

Returns **true** if comparison holds, **false** if not.

## Remarks

Compares a `_variant_t` object with a **VARIANT**, testing for equality or inequality.

**END Microsoft Specific**

## See Also

[\\_variant\\_t Class](#)

# \_variant\_t Extractors

12/21/2017 • 1 min to read • [Edit Online](#)

## Microsoft Specific

Extract data from the encapsulated **VARIANT** object.

## Syntax

```
operator short( ) const;
operator long( ) const;
operator float( ) const;
operator double( ) const;
operator CY( ) const;
operator _bstr_t( ) const;
operator IDispatch*( ) const;
operator bool( ) const;
operator IUnknown*( ) const;
operator DECIMAL( ) const;
operator BYTE( ) const;
operator VARIANT() const throw();
operator char() const;
operator unsigned short() const;
operator unsigned long() const;
operator int() const;
operator unsigned int() const;
operator __int64() const;
operator unsigned __int64() const;
```

## Remarks

Extracts raw data from an encapsulated **VARIANT**. If the **VARIANT** is not already the proper type, **VariantChangeType** is used to attempt a conversion, and an error is generated upon failure:

- **operator short()** Extracts a **short** integer value.
- **operator long()** Extracts a **long** integer value.
- **operator float()** Extracts a **float** numerical value.
- **operator double()** Extracts a **double** integer value.
- **operator CY()** Extracts a **CY** object.
- **operator bool()** Extracts a **bool** value.
- **operator DECIMAL()** Extracts a **DECIMAL** value.
- **operator BYTE()** Extracts a **BYTE** value.
- **operator \_bstr\_t()** Extracts a string, which is encapsulated in a **\_bstr\_t** object.
- **operator IDispatch\*()** Extracts a dispinterface pointer from an encapsulated **VARIANT**. **AddRef** is called on the resulting pointer, so it is up to you to call **Release** to free it.
- **operator IUnknown\*()** Extracts a COM interface pointer from an encapsulated **VARIANT**. **AddRef** is called on the resulting pointer, so it is up to you to call **Release** to free it.

**END Microsoft Specific**

## See Also

[\\_variant\\_t Class](#)

# Microsoft Extensions

12/21/2017 • 1 min to read • [Edit Online](#)

*asm-statement:*

`_asm assembly-instruction ;opt`

`_asm { assembly-instruction-list } ;opt`

*assembly-instruction-list:*

`assembly-instruction ;opt`

`assembly-instruction ; assembly-instruction-list ;opt`

*ms-modifier-list:*

`ms-modifier ms-modifier-list opt`

*ms-modifier:*

`_cdecl`

`_fastcall`

`_stdcall`

`_syscall` (reserved for future implementations)

`_oldcall` (reserved for future implementations)

`_unaligned` (reserved for future implementations)

*based-modifier*

*based-modifier:*

`_based ( based-type )`

*based-type:*

*name*

# Nonstandard Behavior

12/21/2017 • 2 min to read • [Edit Online](#)

The following sections list some of the places where the Visual C++ implementation of C++ does not comply with the C++ standard. The section numbers given below refer to the section numbers in the C++ 11 standard (ISO/IEC 14882:2011(E)).

The list of compiler limits that differ from those defined in the C++ standard is given in [Compiler Limits](#).

## Covariant Return Types

Virtual base classes are not supported as covariant return types when the virtual function has a variable number of arguments. This does not comply with section 10.3, paragraph 7 of the C++ ISO specification. The following sample does not compile, giving compiler error [C2688](#)

```
// CovariantReturn.cpp
class A
{
    virtual A* f(int c, ...); // remove ...
};

class B : virtual A
{
    B* f(int c, ...); // C2688 remove ...
};
```

## Binding Nondependent Names in Templates

The Visual C++ compiler does not currently support binding nondependent names when initially parsing a template. This does not comply with section 14.6.3 of the C++ ISO specification. This can cause overloads declared after the template (but before the template is instantiated) to be seen.

```
#include <iostream>
using namespace std;

namespace N {
    void f(int) { cout << "f(int)" << endl;}
}

template <class T> void g(T) {
    N::f('a'); // calls f(char), should call f(int)
}

namespace N {
    void f(char) { cout << "f(char)" << endl;}
}

int main() {
    g('c');
}
// Output: f(char)
```

## Function Exception Specifiers

Function exception specifiers other than `throw()` are parsed but not used. This does not comply with section 15.4 of the ISO C++ specification. For example:

```
void f() throw(int); // parsed but not used
void g() throw();    // parsed and used
```

For more information on exception specifications, see [Exception Specifications](#).

## char\_traits::eof()

The C++ standard states that `char_traits::eof` must not correspond to a valid `char_type` value. The Visual C++ compiler enforces this constraint for type `char`, but not for type `wchar_t`. This does not comply with the requirement in Table 62 in section 12.1.1 of the C++ ISO specification. The example below demonstrates this.

```
#include <iostream>

int main()
{
    using namespace std;

    char_traits<char>::int_type int2 = char_traits<char>::eof();
    cout << "The eof marker for char_traits<char> is: " << int2 << endl;

    char_traits<wchar_t>::int_type int3 = char_traits<wchar_t>::eof();
    cout << "The eof marker for char_traits<wchar_t> is: " << int3 << endl;
}
```

## Storage Location of Objects

The C++ standard (section 1.8 paragraph 6) requires complete C++ objects to have unique storage locations. However with Visual C++, there are cases where types without data members will share a storage location with other types for the lifetime of the object.

# Compiler Limits

12/21/2017 • 1 min to read • [Edit Online](#)

The C++ standard recommends limits for various language constructs. The following is a list of cases where the Visual C++ compiler does not implement the recommended limits. The first number is the limit that is established in the ISO C++ 11 standard (INCITS/ISO/IEC 14882-2011[2012], Annex B) and the second number is the limit implemented by Visual C++:

- Nesting levels of compound statements, iteration control structures, and selection control structures - C++ standard: 256, Visual C++ compiler: depends on the combination of statements that are nested, but generally between 100 and 110.
- Parameters in one macro definition - C++ standard: 256, Visual C++ compiler: 127.
- Arguments in one macro invocation - C++ standard: 256, Visual C++ compiler 127.
- Characters in a character string literal or wide string literal (after concatenation) - C++ standard: 65536, Visual C++ compiler: 65535 single-byte characters, including the `null` terminator, and 32767 double-byte characters, including the `null` terminator.
- Levels of nested class, structure, or union definitions in a single `struct-declaration-list` - C++ standard: 256, Visual C++ compiler: 16.
- Member initializers in a constructor definition - C++ standard: 6144, Visual C++ compiler: at least 6144.
- Scope qualifications of one identifier - C++ standard: 256, Visual C++ compiler: 127.
- Nested `extern` specifications - C++ standard: 1024, Visual C++ compiler: 9 (not counting the implicit `extern` specification in global scope, or 10, if you count the implicit `extern` specification in global scope..
- Template arguments in a template declaration - C++ standard: 1024, Visual C++ compiler: 2046.

## See Also

[Nonstandard Behavior](#)